

From Tools to Teammates: A Comprehensive Survey of Human-AI Collaboration in Scientific Writing and Research Workflows

1 Conceptual Foundations and the Evolution of Collaborative Paradigms

1.1 From Passive Tools to Active Teammates: A Historical Trajectory

1.1 The Evolution from Tool to Teammate: A Historical Trajectory

The historical trajectory of human-AI interaction reveals a profound paradigmatic shift, moving away from the conventional “human factors” perspective that views artificial intelligence merely as highly capable tools toward an emerging conception of AI as collaborative teammates and joint cognitive partners. For decades, the dominant framework in engineering psychology and human-computer interaction was rooted in the “tool” metaphor, where AI systems were designed to augment human intelligence by executing specific, well-defined tasks under strict human supervision. In this traditional view, the human remains the sole active agent, while the machine serves as a passive instrument to accomplish increasingly sophisticated objectives [1]. While foundational, this perspective increasingly appears limiting as AI capabilities evolve from simple automation to complex, adaptive behaviors that mimic human cognitive processes, setting the stage for more integrated socio-technical models.

As the sophistication of intelligent entities has surged, particularly with the advent of large pre-trained models and generative agents, a second “camp” of thought has gained significant traction. This perspective argues that it is no longer appropriate to classify these advanced systems solely as tools; rather, they should be conceptualized as “teammates” capable of engaging in sustained, dynamic collaboration. This shift necessitates a fundamental re-examination of research paradigms, drawing heavily from literature on human-human teaming to inform the design of human-AI systems [1]. The transition to viewing AI as a teammate implies a move from a unidirectional command-and-control structure to a bidirectional partnership where the AI possesses a degree of agency, allowing it to initiate actions, offer unsolicited advice, and adapt to the human partner’s changing needs. This evolution is critical for achieving mission effectiveness in complex domains where the sheer volume of data and the speed of decision-making exceed human cognitive limits alone, bridging the gap toward the symbiotic fusion described in later paradigms.

However, the discourse has not settled simply on replacing the “tool” metaphor with the “teammate” metaphor. A third, more nuanced perspective is emerging, which posits that both previous metaphors may be flawed or insufficient for capturing the full complexity of modern human-AI interactions. This emerging view emphasizes concepts such as “joint activity,” “joint cognitive activity,” or “human-AI joint cognitive systems” (HAIJCS). In this framework, the focus shifts from the individual attributes of the human or the machine to the synergistic properties of the coupled system itself [1]. The human-AI system is characterized as a joint cognitive ecosystem where humans and intelligent agents work as teammates, leveraging their unique strengths to overcome the limitations inherent in either entity acting alone [2]. This approach aligns with the principles of Human-Centered AI (HCAI), ensuring that while AI functions as a teammate, the ultimate control and responsibility remain with the human, thereby preventing the erosion of human agency [3]. This conceptualization of a unified “ecosystem” provides the necessary theoretical groundwork for defining specific architectures like Hybrid Intelligence and the Centaur model.

The historical progression also reflects a change in the nature of the tasks assigned to AI. Initially, AI was deployed for classification tasks and low-level service jobs where consciousness or deep intuition

was not required. However, as the field advances toward tasks requiring creativity, empathy, and metathinking, the requirement for AI to act as a conscious or quasi-conscious partner becomes more pressing [4]. The vision of AI collaborators, once a staple of science fiction, is now a tangible research goal where artificial agents assist human partners with special talents, understanding the nuances of collaboration and communication that define skilled human interaction [5]. This aspiration moves beyond the intermediate human collaboration tools of the 1960s, which augmented intelligence but lacked true collaborative intelligence, toward robust AIs that can learn, share knowledge, and collaborate resiliently [5].

Furthermore, the evolution from tools to teammates is driven by the necessity to manage uncertainty and novelty. In environments characterized by rapid change, such as future command and control operations or clinical decision-making, static tools are insufficient. Instead, scalable interactive machine learning and resilient human-AI teams are required to foster adaptability and efficiency [6]. The recognition that AI can function as a coequal partner in clinical settings, for instance, highlights the potential to improve treatment quality by assisting overworked professionals, provided that the tension between technical adoption and human-centered collaboration is effectively managed [7].

Despite these advancements, challenges remain in fully realizing the teammate paradigm. There is an ongoing debate regarding whether AI can truly function as a teammate without contradicting the human-centered approach, with concerns that the “teaming” paradigm might lead to humans losing control over AI systems [3]. Additionally, the perception of team cognition often differs between human-only teams and mixed human-AI teams, with mixed teams sometimes reporting lower perceived cognition despite achieving higher performance metrics [8]. This dichotomy underscores the complexity of the transition; while performance may improve with AI teammates, the psychological and social dynamics of the team undergo significant transformation, highlighting the need for rigorous definitions of interaction modes.

Ultimately, the historical trajectory signifies a move toward “human-AI joint cognitive ecosystems,” where the overall performance depends on the optimal collaboration and design across multiple human-AI systems within an intelligent sociotechnical environment [2]. This evolution demands new interaction patterns that support clear communication, trustworthiness, and genuine collaboration, moving beyond the simplistic, one-off interactions that have dominated the era of AI as a mere tool [9]. As we stand at this juncture, the definition of the human-AI relationship continues to expand, promising a future where humans and machines engage in deep, critical, and symbiotic partnerships. It is within this evolving landscape that precise conceptualizations, such as Hybrid Intelligence and the Centaur paradigm, become essential to structure these collaborations effectively.

1.2 Defining Hybrid Intelligence and the Centaur Paradigm

1.2 Defining Hybrid Intelligence and the Centaur Paradigm

Building upon the historical trajectory from passive tools to active teammates and joint cognitive ecosystems outlined in the previous section, it becomes imperative to establish precise conceptual definitions for the resulting socio-technical entities. The terms “Hybrid Intelligence” and the “Centaur” model have emerged as the dominant paradigms for describing systems where human and machine capabilities are inextricably linked. However, these terms are often used interchangeably in popular discourse, obscuring critical distinctions in their operational mechanisms and theoretical underpinnings. A rigorous definition is essential to guide the design of future scientific workflows, moving beyond the vague notion of “human-in-the-loop” to a structured understanding of how

bounded rationality and algorithmic precision can be synthesized to achieve superior outcomes. These specific conceptualizations provide the necessary vocabulary to operationalize the “joint cognitive ecosystem” described earlier, serving as the bridge between abstract theory and the analytical framework introduced in the subsequent section.

Hybrid Intelligence Systems (HIS) represent a broad conceptual framework wherein humans and AI systems are integrated to create robust, reliable, and ethically grounded decision-making entities. Unlike traditional automation, which seeks to replace human intervention, Hybrid Intelligence is predicated on the complementary strengths of both agents. The core thesis of this paradigm is that the combination of human intuition, contextual awareness, and ethical reasoning with the computational speed, scalability, and pattern recognition of AI yields a performance level that neither could attain in isolation [10]. This concept challenges the historical dichotomy between statistical methods and human cognition, arguing instead for a balanced coexistence where the limitations of one are offset by the capabilities of the other. For instance, while machines excel in processing vast datasets and identifying quantitative correlations, they frequently struggle with “soft information,” extreme uncertainty, and the interpretation of nuanced social cues—domains where human intelligence remains the gold standard [11]. Consequently, a Hybrid Intelligence System is not merely a tool used by a human, but a unified entity where the interaction dynamics are designed to leverage information asymmetry and capability asymmetry to generate Complementary Team Performance (CTP) [12].

Within this broader landscape of Hybrid Intelligence, the “Centaur” model offers a more specific and potent instantiation of human-AI collaboration. Originating from the realm of competitive chess, where teams comprising a human player and an AI engine consistently outperformed both solo human grandmasters and standalone supercomputers, the Centaur paradigm posits that the optimal decision-maker is a symbiotic fusion of half-human and half-AI agency [13]. In a Centaur system, the AI is not a passive calculator awaiting commands but an active partner capable of recognizing the goals, constraints, and cognitive bounds of its human counterpart. The defining characteristic of the Centaur is its ability to navigate the trade-offs inherent in partially observable tasks, where the AI must align its beliefs with the human’s mental model to ensure its suggestions are acceptable and actionable [14]. This alignment is crucial; without it, the friction between human intuition and algorithmic logic can degrade performance rather than enhance it, undermining the “teammate” status established in the evolving historical context.

The theoretical foundation of the Centaur model rests on the concept of bounded rationality. Humans are inherently limited by cognitive resources, biases, and the inability to process infinite information streams, leading to decisions that are “good enough” rather than globally optimal. Conversely, AI systems, while computationally powerful, often lack the contextual grounding and flexible adaptability required for complex, open-ended problems. The Centaur paradigm seeks to extend the bounds of human rationality by augmenting it with machine intelligence, effectively creating a decision-maker that operates with expanded cognitive horizons [15]. This extension is not automatic; it requires a deliberate architectural design where the AI models the human’s behavior—not as an idealized rational agent, but as a noisily rational or risk-aware entity subject to specific biases and heuristics [16]. By accurately modeling these granular human behaviors, the AI can tailor its assistance to mitigate specific human blind spots, such as over-confidence or risk aversion, thereby nudging the joint system toward better decisions [14].

A key distinction between the general concept of Hybrid Intelligence and the specific Centaur model lies in the degree of coupling and the distribution of authority—dimensions that will be rigorously defined in the following subsection. While Hybrid Intelligence encompasses a wide spectrum of

interactions, the Centaur model specifically implies a tight coupling where the human and AI function as a single cognitive unit with shared objectives. In this configuration, the AI acts as a “co-pilot” that dynamically delegates tasks based on real-time assessments of who—human or machine—is better suited to handle a specific sub-task at any given moment [17]. This dynamic delegation is central to the Centaur’s success, as it allows the system to optimize for team utility rather than individual agent accuracy. Research indicates that the most accurate AI is not necessarily the best teammate; instead, an AI optimized for predictability and complementarity with human limitations often yields higher team performance [18].

The implications of these definitions extend far beyond gaming or theoretical exercises. In scientific research, adopting a Centaur mindset implies shifting from viewing AI as a mere instrument for data analysis to treating it as a collaborative partner in hypothesis generation and experimental design. This requires acknowledging that the “intelligence” of the system emerges from the interaction itself, not just from the sum of its parts. As noted in recent literature, the failure to properly conceptualize these systems can lead to fragmented implementations where human factors are overlooked, undermining the potential for true synergy [19]. Therefore, defining Hybrid Intelligence and the Centaur paradigm with precision is a prerequisite for developing the next generation of scientific tools that can truly amplify human creativity and rigor. Having established these foundational models, the discussion now turns to a tri-dimensional framework designed to characterize and measure the specific topology of agency, coupling, and authority within these collaborative structures.

1.3 Dimensions of Collaboration: Agency, Coupling, and Authority

1.3 A Tri-Dimensional Framework for Characterizing Collaboration

Having established the conceptual distinctions between Hybrid Intelligence and the Centaur paradigm, it is necessary to operationalize these abstract models into a rigorous analytical framework. Relying on vague analogies or isolated case studies is insufficient for dissecting the complex interplay between human operators and artificial agents in scientific workflows. Instead, a unified conceptual structure is required to characterize the specific topology of any given human-AI system. This subsection introduces such a framework, structured around three critical, interdependent dimensions: **agency**, **coupling**, and **directive authority**. These dimensions provide a granular vocabulary for characterizing the degree of integration and control within joint cognitive systems, moving beyond the binary distinction of “human-in-the-loop” versus “fully autonomous” toward a nuanced understanding of the collaborative paradigms introduced in the previous section.

The first dimension, **agency**, addresses the fundamental question of who initiates action within the system, directly reflecting the shift from static automation to the dynamic partnerships described in the Centaur model. In traditional automation, agency is static; the machine executes predefined routines only when triggered by a human. However, in modern collaborative systems, agency becomes dynamic and distributed. A robust analysis must distinguish between the initiation of tasks, the generation of sub-goals, and the execution of specific actions. Recent literature suggests that effective collaboration often requires a shift from rigid role assignment to fluid agency, where both human and AI can propose actions based on their respective strengths. For instance, frameworks analyzing human-agent teaming emphasize the need for agents to possess social capabilities that allow them to negotiate roles and initiate interactions as equal partners rather than passive tools [20]. Furthermore, the concept of “mixed-initiative” interaction highlights that optimal perfor-

mance often arises when control can be transferred seamlessly between human and machine based on situational context and capability assessments [21]. Therefore, agency is not merely a property of the agent but a relational state defined by the capacity of both parties to influence the trajectory of the workflow, a prerequisite for the “active partner” status central to the Centaur paradigm.

The second dimension, **coupling**, describes the tightness of the integration between the human and the AI, quantifying the “inextricable link” posited in Hybrid Intelligence definitions. This spectrum ranges from loose coupling, where the AI operates as a distinct backend service consulted occasionally, to tight coupling, where the human and AI function as a symbiotic unit with continuous, real-time feedback loops. Tight coupling often necessitates a shared mental model, where the human maintains an accurate understanding of the AI’s state, and the AI possesses a “theory of mind” regarding the human’s intent and cognitive load. Research into joint cognitive systems posits that high-performance teams require a “symbiotic understanding,” where each partner holds an internal representation of the other’s intentions, facilitating highly coordinated action [22]. In loosely coupled systems, interaction may be transactional and discrete, whereas tightly coupled systems demand continuous alignment, often achieved through mechanisms like predictive coding and active inference to regulate the complexity of interaction [23]. The degree of coupling directly impacts the bandwidth of communication required and the potential for friction; too tight a coupling without adequate transparency can lead to cognitive overload, while too loose a coupling may result in missed opportunities for synergistic problem-solving, thereby failing to realize the Complementary Team Performance discussed in the following section.

The third dimension, **directive authority**, determines who holds the final control and bears the ultimate responsibility for outcomes, addressing the ethical and operational boundaries of the Centaur model. This is distinct from agency; an AI may have the agency to propose a hypothesis or execute a maneuver, but the directive authority resides with the entity empowered to veto or finalize that action. The preservation of meaningful human control is a central ethical and operational challenge, particularly in high-stakes domains. It has been argued that for human control to be “meaningful,” there must be explicit links between the actions of AI agents and the humans who are aware of their moral responsibility, ensuring that responsibility attribution remains commensurate with the human’s ability to intervene [24]. However, the quantification of responsibility reveals that simply keeping a human “in the loop” does not guarantee effective oversight; if the system design obscures the causal chain of events, human comparative responsibility can diminish even when formal authority is retained [25]. Consequently, directive authority must be supported by designs that ensure humans have sufficient causal power and epistemic access to the system’s operations to exercise their authority effectively [26].

These three dimensions—agency, coupling, and authority—are not independent variables but interact dynamically to define the collaborative topology of a scientific workflow. A system might feature high AI agency (initiating complex data analyses) but loose coupling (presenting results only upon request) and retained human authority (final approval of conclusions). Conversely, a self-driving laboratory might exhibit tight coupling and high AI agency in experimental execution, yet remain under strict human directive authority regarding the scope of inquiry. To navigate this design space, researchers have proposed unified sets of dimensions to describe these elaborate processes, arguing that existing models often fail to capture the full scope of collaboration paradigms [27]. By mapping specific applications onto this three-dimensional framework, we can better identify mismatches—for example, where high coupling exists without corresponding mutual understanding, leading to automation bias—or opportunities for optimizing the division of labor. Ultimately, this framework serves as a foundational tool for designing hybrid intelligence systems

that are not only efficient but also aligned with human values and capable of sustaining meaningful human oversight [19]. Through the careful calibration of agency, coupling, and authority, we can transition from viewing AI as mere instrumentation to recognizing it as a genuine teammate, setting the stage for achieving the Complementary Team Performance explored in the subsequent theoretical analysis.

1.4 Theoretical Foundations of Complementarity and Joint Performance

1.4 Theoretical Foundations of Complementarity and Joint Performance

Building upon the tri-dimensional framework of agency, coupling, and directive authority established in the previous section, the ultimate aspiration of integrating artificial intelligence into scientific workflows is to achieve a synergistic outcome where the combined performance of the human-AI dyad exceeds the capabilities of either entity operating in isolation. This phenomenon is formally defined as Complementary Team Performance (CTP) [12]. Achieving CTP represents the “holy grail” of hybrid intelligence, yet empirical evidence suggests it remains elusive in many practical applications [12; 28]. Theoretical frameworks posit that CTP is not an automatic consequence of deploying high-accuracy models or merely configuring the appropriate levels of agency and coupling; rather, it is contingent upon the existence and effective exploitation of specific sources of complementarity potential [12; 29]. Central to this theoretical foundation is the distinction between two primary drivers of synergy: information asymmetry and capability asymmetry [12]. Understanding the mechanics of these asymmetries is critical for designing collaborative systems that transcend the limitations of standalone human or machine cognition, moving from structural definitions of collaboration to the functional realization of joint success.

Information asymmetry serves as a fundamental source of complementarity potential, arising from the divergent access to contextual data between human and artificial agents [29]. In real-world decision-making environments, humans often possess unique, situated, or tacit knowledge that is unavailable to the AI system, while the AI may have access to vast historical datasets or pattern recognition capabilities beyond human reach. Theoretical models suggest that CTP is realized when humans can leverage their unique contextual information to adjust, override, or refine AI recommendations [29]. Research indicates that when teams are designed to explicitly account for this asymmetry, allowing humans to inject local context into the decision loop, the resulting joint performance can surpass the ceiling of individual performance [30]. However, realizing this potential requires more than just the presence of asymmetric information; it demands mechanisms that enable humans to recognize when their private information contradicts the AI’s statistical inference and the agency to act upon it. Without such mechanisms, the team fails to capitalize on the informational diversity that defines their complementary nature, rendering the structural coupling described in Section 1.3 ineffective.

Parallel to information asymmetry is the concept of capability asymmetry, which stems from the inherent heterogeneity in the cognitive and computational strengths of humans and machines [12]. While AI systems excel at processing high-dimensional data, identifying statistical regularities, and maintaining consistent performance across repetitive tasks, humans demonstrate superior adaptability, causal reasoning, and the ability to navigate scenarios characterized by extreme uncertainty or “soft” information that resists quantification [10; 31]. Theoretical work emphasizes that CTP is most likely to emerge in domains where neither the human nor the algorithm holds a dominant

performance advantage across all instances [32]. Instead, the optimal strategy involves a dynamic allocation of tasks based on the specific difficulty of the instance relative to the distinct capability profiles of the team members. For instance, in complex diagnostic tasks, AI might handle cases with clear statistical signatures, while deferring ambiguous or novel cases to human experts who can apply intuitive leaps and cross-domain knowledge [33]. This routing approach, often formalized through learning complementary policies, aims to maximize the expected utility of the team by strategically deferring decisions to the agent best suited for the specific context, thereby mitigating the weaknesses of each partner [32; 34].

Despite the clear theoretical promise, achieving CTP in practice proves challenging. A significant body of literature highlights that simply pairing a human with a highly accurate AI does not guarantee superior team outcomes [18]. In many experimental settings, the presence of AI explanations or recommendations has been shown to increase human reliance on the system without necessarily improving the accuracy of the final decision, particularly when the AI is incorrect [28; 35]. This phenomenon underscores a critical gap in current implementations: the failure to optimize AI systems specifically for teamwork rather than for standalone accuracy. Traditional machine learning objectives focus on minimizing global loss functions, implicitly assuming autonomous operation. However, theoretical arguments propose that AI systems intended for collaboration should be trained to optimize team utility, potentially sacrificing marginal gains in standalone accuracy to become more predictable, interpretable, or aligned with human cognitive styles [18]. By modeling the human partner’s error patterns and cognitive biases during the training phase, AI agents can learn to complement human reasoning more effectively, focusing their efforts on problem instances that are difficult for humans while seeking human input for cases where the algorithm is uncertain [33; 36].

Furthermore, the realization of CTP is deeply intertwined with the concept of appropriate reliance, which acts as the mediating variable between complementarity potential and actual joint performance [37]. Appropriate reliance requires humans to possess an accurate mental model of the AI’s capabilities and limitations, enabling them to discern when to trust the algorithm and when to rely on their own judgment. Theoretical frameworks suggest that human learning plays a pivotal role in calibrating this reliance over time [37]. As humans interact with AI partners, they engage in a co-learning process where they refine their understanding of the system’s behavior, ideally leading to a state where the team consistently outperforms individual baselines [38]. However, this process is fragile; factors such as automation bias, where humans over-defer to the machine, or algorithmic aversion, where humans unjustifiably reject correct AI advice, can erode the benefits of complementarity. Consequently, the theoretical foundation of CTP extends beyond the static properties of the agents to include the dynamic interaction protocols, feedback loops, and interface designs that facilitate the effective exchange of information and the calibration of trust [39; 40].

In summary, the theoretical landscape of human-AI collaboration posits that Complementary Team Performance is an achievable but non-trivial outcome dependent on the strategic alignment of information and capability asymmetries. It requires a shift from viewing AI as a standalone oracle to designing it as a complementary partner whose strengths are explicitly tuned to offset human weaknesses, and vice versa. Yet, the successful exploitation of these asymmetries hinges on a deeper cognitive prerequisite: the ability of both partners to accurately model each other’s intentions and limitations. Without this shared understanding, the mechanisms for leveraging asymmetry described above cannot function optimally. This necessity sets the stage for the subsequent discussion on Cognitive Alignment and Mutual Theory of Mind, which explores how bidirectional mental modeling serves as the essential bridge between static complementarity theories and the emergence

of truly symbiotic autonomous systems. Future advancements in this domain rely on developing rigorous metrics to quantify complementarity potential and creating adaptive systems that can dynamically navigate the trade-offs between human intuition and machine precision to unlock the full synergistic potential of hybrid intelligence [31; 41].

1.5 Cognitive Alignment and Mutual Theory of Mind in Joint Systems

1.5 Cognitive Alignment and Mutual Theory of Mind in Joint Systems

Building upon the theoretical foundations of complementarity and the strategic exploitation of information and capability asymmetries discussed in the previous section, the transition from viewing artificial intelligence as a passive tool to recognizing it as an active teammate necessitates a fundamental re-evaluation of the cognitive architectures underpinning human-AI interaction. While Section 1.4 established that Complementary Team Performance (CTP) relies on the effective exchange of asymmetric strengths, Section 1.5 argues that realizing this potential requires **cognitive alignment**: a state where human and machine partners share a coherent understanding of goals, constraints, and situational contexts. Central to achieving this alignment is the concept of **Mutual Theory of Mind (MToM)**, a bidirectional cognitive capability wherein both the human and the AI maintain dynamic, evolving mental models of each other’s beliefs, intentions, and capabilities. Without MToM, the mechanisms for leveraging asymmetry described previously fail; communication breakdowns become inevitable as the AI misinterprets user needs due to a flawed understanding of human intent, while the human develops incorrect assumptions about the AI’s limitations or operational logic [42]. This section explores the theoretical imperatives of MToM, the mechanisms for achieving cognitive alignment, and the critical role of trust calibration in sustaining these joint systems, serving as a necessary bridge between static complementarity theories and the emerging paradigms of symbiotic autonomy.

The foundation of MToM lies in the ability to attribute mental states to others, a cognitive faculty humans employ naturally but which must be explicitly engineered into artificial agents to facilitate the “appropriate reliance” highlighted in prior theoretical frameworks. In traditional human-human teams, Theory of Mind (ToM) allows individuals to infer desires and beliefs to anticipate actions and resolve conflicts. When applied to human-AI teams, this concept expands into a dual requirement: the AI must possess a computational ToM to model the human, and the human must develop a “Theory of AI’s Mind” (ToAIM) to accurately predict the system’s behavior [43]. Research indicates that humans often struggle to integrate information from AI teammates effectively, particularly when communication loads are high or when cognitive biases lead to the underweighting of ambiguous but useful information [44]. Therefore, fostering a robust MToM is not merely about improving the AI’s social intelligence but also about scaffolding the human’s ability to understand the machine’s decision-making processes, thereby closing the gap between potential and actual complementarity.

Achieving cognitive alignment requires more than static explanations; it demands dynamic, context-aware modeling of nested beliefs that can adapt to the fluid nature of scientific workflows. Recent advancements propose frameworks where agents utilize epistemic logic to represent discrepancies in beliefs about feasible plans, allowing them to adapt or communicate to resolve these gaps in real-time [45]. For instance, in scenarios where shared mental models cannot be guaranteed, such as ad-hoc teaming, agents must infer the human partner’s willingness to comply and adjust their

intervention strategies accordingly. Studies have shown that mixed-initiative collaborations, where the robot infers human suboptimality and adapts its behavior via online Bayesian methods, significantly enhance both objective team performance and subjective measures of trust [46]. This adaptability is crucial because rigid adherence to optimal planning by the AI can alienate human partners who operate under different cognitive constraints or possess unique contextual knowledge, effectively nullifying the benefits of information asymmetry.

A critical component of MToM is the management of trust, which serves as the lubricant for collaborative friction and directly influences the calibration of reliance discussed in Section 1.4. Trust in human-AI teams is not a static trait but a dynamic variable that evolves through interaction phases and is heavily influenced by the transparency of the AI’s mental modeling. The Collaborative Human-AI Trust (CHAI-T) framework posits that trust management must account for task specificity and the temporal evolution of team dynamics, moving beyond simple antecedents like system reliability [39]. However, the pursuit of trust through transparency carries ethical risks. If an AI models the human’s mental state too accurately, it may exploit this knowledge to manipulate human behavior toward a “greater good,” raising profound questions about autonomy and acceptable symbiosis [47]. Consequently, the design of MToM systems must balance the utility of predictive modeling with the preservation of human agency, ensuring that implicit guidance does not devolve into covert control [48].

Furthermore, the efficacy of MToM is contingent upon the quality of the mental models being exchanged, a challenge that current technologies are only beginning to address. While large language models (LLMs) have demonstrated impressive capabilities in various domains, empirical evaluations suggest they still lag behind humans in robust social intelligence and genuine Theory of Mind reasoning [49]. Benchmarks such as ToMBench reveal that even advanced models struggle with tasks requiring the inference of false beliefs or complex social dynamics, indicating that current AI lacks the deep, intuitive understanding necessary for seamless cognitive alignment [50]. To bridge this gap, researchers are exploring hybrid approaches, such as combining Bayesian inverse planning with language models, to create systems that can better extract unified representations of human goals from multimodal data [51]. Additionally, the concept of “belief resonance,” where inferred mental states of one agent directly influence the predictive beliefs of another, offers a promising pathway for efficient, automatic coordination without the need for explicit, computationally expensive reasoning at every step [52].

Ultimately, the realization of effective human-AI joint systems depends on the co-evolution of human and machine mental models, a process termed “co-learning” that emphasizes mutual understanding, benefits, and growth, ensuring that both entities adapt to each other over time [38]. As AI systems become more autonomous, the necessity for humans to cultivate a precise ToAIM becomes paramount; training humans to recognize AI quirks and failure modes can be more effective than providing internal model states that may increase cognitive overhead without improving prediction accuracy [43]. By prioritizing the development of bidirectional mental modeling and addressing the asymmetries in cognitive capabilities, the scientific community can move toward joint systems where cognitive alignment is not an aspiration but a foundational architectural principle. This evolution sets the stage for the next frontier in collaboration: the emergence of Symbiotic Autonomous Systems and Embodied Agents, where the boundaries between distinct mental models begin to dissolve into integrated, bio-inspired cognitive structures.

1.6 Emerging Paradigms: Symbiotic Autonomous Systems and Embodied Agents

1.6 Emerging Paradigms: Symbiotic Autonomous Systems and Embodied Agents

Building directly upon the foundational principle of cognitive alignment and the bidirectional mental modeling discussed in Section 1.5, the evolution of human-AI collaboration is now transcending the traditional model of distinct entities interacting through discrete interfaces. Where Mutual Theory of Mind (MToM) sought to bridge the gap between separate cognitive architectures, emerging frontier concepts propose a future where the boundaries between human and machine cognition become increasingly porous. This shift is epitomized by the emergence of Symbiotic Autonomous Systems (SAS) and Embodied AI, paradigms that propose bio-brain-social-inspired structures and grounded physical-virtual amalgamations to enable seamless, coherent human-machine cognition. Unlike earlier iterations of automation which viewed AI as a tool for specific tasks, or even the active teammate model reliant on explicit communication, these emerging frameworks envision a state where intelligent systems function as coherent cognitive partners within hybrid societies, characterized by autonomous collective intelligence enabled by the symbiosis of human and machine interactions [53]. The theoretical underpinning of SAS suggests a transition from simple cooperation to a state of mutual dependence and enhancement, where the system adopts heterogeneously synergized structures to augment human capability with autonomous machine intelligence, aiming for the next generation of trustworthy, mission-critical intelligent systems [53].

Central to this paradigm shift is the concept of embodiment, which posits that intelligence cannot be fully realized without a physical or virtual body that interacts dynamically with an environment, thereby grounding the abstract “mental models” of Section 1.5 in tangible reality. The call for Embodied AI represents a fundamental step beyond the static learning capabilities of Large Language Models, emphasizing a theoretical framework based on cognitive architectures that integrate perception, action, memory, and learning as essential components [54]. This perspective aligns with the principle of active inference, suggesting that true artificial general intelligence (AGI) must be capable of seamless communication, collaboration, and coexistence with humans within real-world environments [54]. The integration of foundation models into agent actions gives rise to “Agent AI,” a holistic intelligence that spans robotics, gaming, and healthcare, challenging reductionist approaches by functioning as cohesive wholes rather than isolated algorithms [55]. These agents are designed to perceive visual stimuli and language inputs to produce meaningful embodied actions, thereby mitigating the hallucinations common in large foundation models by grounding outputs in environmentally correct contexts [56].

The realization of these symbiotic and embodied systems often relies on novel architectural innovations that bridge the gap between biological inspiration and computational execution, effectively operationalizing the “co-learning” dynamics previously theorized. For instance, the concept of using Large Language Models as a robotic brain, known as LLM-Brain, unifies egocentric memory and control through natural language dialogues, enabling robots to perform complex tasks like active exploration and embodied question answering in closed-loop multi-round interactions [57]. Furthermore, the development of “Robot Instinct” introduces a layered architecture that harmonizes the high-level intelligence of LLMs with innate, survival-essential responses such as safety assurance and obstacle avoidance, ensuring that autonomous systems can operate safely across diverse environments [58]. This bio-inspired approach extends to the hardware level, where neuromorphic computing and spiking neural networks mimic the energy efficiency and adaptability of biological brains, offering a pathway to embodied neuromorphic AI that can learn and adapt in real-time [59].

Beyond physical robotics, the scope of embodiment extends into virtual and mixed realities, creating “symmetrical reality” frameworks that accommodate dual perceptual centers for both humans and AI agents, thus expanding the domain of MToM into shared digital spaces. This unified representation allows for the consolidation of physical-virtual amalgamations, facilitating coexistence and active assistance services where AI agents collaborate with humans in both physical and virtual worlds [60]. The role of the metaverse is also pivotal in this context, serving as a calibrated symbolic interface where feedback loops between AGI and human users can facilitate the emergence of stable, embodied general intelligence [61]. In such environments, the distinction between the agent and the environment blurs, as seen in concepts like “embodied swarm robots,” where collectives of robots dynamically alter their shape and density to serve as functional body parts, leveraging human brain plasticity to achieve intuitive interaction [62].

The cognitive implications of these paradigms are profound, suggesting a move toward “hybrid active inference” where machine learning parts integrate into human cognition by explaining away multi-modal sensory measurements alongside brain signals, effectively dissolving the rigid separation of mental models described in prior sections [63]. This integration fosters a joint embodiment where human action and perception are treated as the machine’s own, potentially accelerating emergent properties within artificial intelligence [63]. Moreover, the development of introspective agents that consider their own abilities in the context of their environment mirrors biological adaptation, reframing strategy as one tool among many for succeeding in complex environments [64]. As these systems evolve, they promise to transform scientific workflows not merely by automating tasks, but by creating symbiotic partnerships where the collective intelligence of the human-machine dyad exceeds the sum of its parts, fundamentally reshaping the landscape of research and discovery [65]. Ultimately, the convergence of SAS and Embodied AI signals a future where scientific inquiry is conducted by hybrid entities capable of autonomous development, lifelong learning, and deep contextual understanding, anchored in both physical reality and virtual simulation [66].

2 AI-Augmented Ideation, Literature Synthesis, and Knowledge Discovery

2.1 Zero-Shot and Fine-Tuned Hypothesis Generation

2.1 Zero-Shot and Fine-Tuned Hypothesis Generation

The capacity of Large Language Models (LLMs) to function not merely as repositories of existing knowledge but as active generators of novel scientific hypotheses represents a foundational shift in the research workflow. Historically, the “Eureka moment” of hypothesis generation has been the exclusive domain of human researchers, driven by painstaking data analysis and intuitive leaps across disciplinary boundaries. However, the exponential growth of scientific literature has created information barriers that slow the pace of discovery, necessitating tools capable of synthesizing interdisciplinary insights at scale. Recent investigations have formally explored whether LLMs can propose valid, testable scientific hypotheses, moving beyond simple text completion to genuine ideation. In zero-shot settings, where models are tasked with generating hypotheses without task-specific training examples, LLMs have demonstrated a surprising ability to propose untrained yet validated hypotheses derived from testing literature that was not present in their training data [67]. This capability suggests that the vast, global knowledge embedded within these models can be leveraged to break down silos between disciplines, fostering a new wave of scientific discovery that transcends the limitations of human information processing capacities.

The mechanism behind successful zero-shot hypothesis generation often involves a counter-intuitive

relationship with model uncertainty. While traditional AI applications strive to minimize uncertainty to ensure factual accuracy, the generative phase of scientific discovery may benefit from increased uncertainty. Research indicates that increasing model uncertainty can facilitate the generation of a broader range of candidate hypotheses, potentially enhancing the model’s ability to propose novel ideas in zero-shot scenarios [67]. By allowing the model to explore low-probability token sequences, researchers can catalyze the formation of connections between disparate fields that a strictly deterministic model might suppress. This approach aligns with findings that LLMs can surpass human experts in predicting experimental outcomes in fields like neuroscience, particularly when the models are tuned on domain-specific literature, suggesting that the integration of noisy, interrelated findings allows AI to forecast novel results more effectively than humans alone [68].

Building upon these zero-shot foundations, fine-tuned approaches offer a pathway to refine hypothesis generation for specific domains or data contexts, introducing a layer of optimization that precedes the complex social refinement seen in later stages. When provided with labeled examples or specific datasets, LLMs can be guided to generate hypotheses that are not only novel but also optimized for predictive performance. One notable methodology involves an iterative process where initial hypotheses are generated from a small number of examples and subsequently updated to improve quality, utilizing reward functions inspired by multi-armed bandits to manage the trade-off between exploration and exploitation [69]. This iterative refinement has been shown to yield hypotheses that enable significantly better predictive performance in classification tasks compared to standard few-shot prompting, outperforming even supervised learning models on challenging real-world datasets. Furthermore, these generated hypotheses have been found to not only corroborate human-verified theories but also uncover entirely new insights, demonstrating the potential for AI to act as a co-discoverer rather than just an assistant.

However, the deployment of LLMs for hypothesis generation—whether zero-shot or fine-tuned—is not without significant challenges, primarily concerning the reliability of the generated content and the model’s awareness of its own knowledge limitations. A critical prerequisite for trustworthy hypothesis generation is the model’s ability to distinguish between what it knows and what it does not know. Studies evaluating LLM self-knowledge reveal that while models possess an intrinsic capacity to identify unanswerable questions, there remains a considerable gap between their proficiency and human capabilities in recognizing the limits of their parametric knowledge [70]. In scenarios where models lack relevant knowledge, they often fail to express uncertainty consistently, instead generating plausible but factually incorrect “hallucinations.” This tendency is exacerbated in fine-tuned models if the training process forces the model to complete sentences regardless of knowledge gaps, leading to the fabrication of non-existent facts [71].

To mitigate these risks and establish a reliable baseline for subsequent collaborative workflows, recent advancements have focused on uncertainty-aware frameworks and refusal-aware instruction tuning. Techniques such as Refusal-Aware Instruction Tuning (R-Tuning) train models to refrain from responding to questions that fall outside their parametric knowledge, effectively teaching the model a “meta-skill” of uncertainty estimation that generalizes across tasks [71]. Additionally, methods that leverage uncertainty quantification, such as analyzing logit output values or employing sampling-based approaches, allow models to autonomously recognize when they are unsure and filter out high-uncertainty responses [72; 73]. These mechanisms are crucial for ensuring that the hypotheses proposed by AI are grounded in reality and that the system can signal when a proposed idea is speculative due to a lack of evidentiary support rather than genuine novelty.

Furthermore, the integration of external knowledge sources and structured reasoning frameworks enhances the robustness of hypothesis generation, setting the stage for more sophisticated multi-

agent interactions. Approaches that combine LLMs with Knowledge Graphs (KGs) allow for the retrieval of task-relevant subgraphs to support reasoning, reducing hallucinations and improving the traceability of the generated hypotheses [74]. Similarly, causal reasoning capabilities in LLMs enable them to identify background causal contexts and generate causal graphs, offering a complementary approach to traditional statistical methods [75]. By synthesizing these diverse capabilities—zero-shot creativity, fine-tuned precision, uncertainty awareness, and structured reasoning—the scientific community is establishing the essential single-agent competencies required for the next evolution in research: the transition from solitary hypothesis generation to the dynamic, iterative refinement enabled by multi-agent collaborative frameworks.

2.2 Multi-Agent Frameworks for Iterative Idea Refinement

2.2 Multi-Agent Frameworks for Iterative Idea Refinement

Building upon the foundational single-agent competencies established in the previous section—ranging from zero-shot creativity and fine-tuned optimization to uncertainty-aware refusal—the scientific workflow now faces the challenge of transforming isolated hypotheses into robust, rigorously vetted research programs. While Section 2.1 detailed how individual models can generate novel ideas and recognize their own knowledge limitations, these solitary agents often struggle with the nuanced, multi-step reasoning required to validate complex scientific concepts without external guidance. To bridge this gap, the field is transitioning toward multi-agent frameworks designed to emulate the complex, iterative social processes inherent in human scientific discovery. These architectures move beyond simple prompt-response dynamics, establishing dynamic loops where multiple AI agents interact to simulate peer review, brainstorming sessions, and critical debate, thereby refining research problems, methods, and experimental designs through structured collaboration.

A foundational strategy in this domain involves emulating the academic peer review process to enhance the reasoning capabilities of Large Language Models (LLMs). By orchestrating interactions where multiple models correct and critique one another, significant improvements in accuracy and logical coherence can be achieved, effectively mitigating the hallucination risks identified in single-agent systems. In such frameworks, each agent independently constructs a solution or research idea, provides reviews on the solutions of others, and assigns confidence levels to their assessments. Upon receiving these peer reviews, agents revise their initial proposals, creating a cycle of iterative refinement. Empirical studies have demonstrated that this collaboration approach delivers superior accuracy across various reasoning tasks compared to existing single-agent methods, highlighting that the exchange of feedback is more effective than mere solution sharing [76]. This mechanism mirrors the human scientific endeavor, where the friction between differing perspectives often leads to more robust conclusions.

Expanding on the concept of peer review, advanced frameworks like ResearchAgent have been developed to automate the generation of research ideas by integrating literature connectivity with iterative feedback loops. Starting with a core paper, these agents augment their knowledge by connecting information over an academic graph and retrieving entities from concept-based stores. Crucially, these systems leverage multiple “ReviewingAgents” that provide iterative feedback on generated problems, methods, and experiment designs. These reviewing agents are instantiated with models aligned to human preferences, ensuring that the evaluation criteria derive from actual human judgments rather than abstract algorithmic metrics. Experimental validation across multiple disciplines has showcased the effectiveness of such agents in generating novel, clear, and valid

research ideas, effectively bridging the gap between vast literature corpora and actionable scientific hypotheses [77].

The creative potential of these multi-agent societies is further amplified by diversifying agent personas and traits. Research indicates that LLM agents can manifest human-like social behaviors, such as conformity and consensus reaching, when organized into specific “societies” with distinct traits and thinking patterns. By fabricating agents with varying characteristics—such as being “easy-going” or “overconfident”—and employing different collaboration strategies like debate or reflection, these systems can optimize efficiency and outperform top-tier single-agent approaches. This social psychology-informed design suggests that cognitive synergy, similar to that found in human teams, can be unleashed within LLMs by dynamically simulating different personas to tackle complex tasks [78] [79]. Furthermore, the introduction of competitive dynamics, as seen in Collaborative Competitive Agents (CCA), allows generator agents to learn from successful executions of others through a discriminator’s feedback, enhancing the robustness and quality of generated outputs in intricate tasks [80].

In the context of specific scientific domains, such as biomedical discovery, the vision of “AI scientists” involves collaborative agents that integrate machine learning tools with experimental platforms. These agents are proficient in self-assessment and planning, using structured memory for continual learning to navigate hypothesis spaces. Rather than removing humans from the loop, these multi-agent systems combine human creativity with the AI’s ability to execute repetitive tasks and analyze large datasets, empowering researchers to focus on high-level strategic decisions [81]. Similarly, in design and creative fields, frameworks like DesignGPT simulate the roles of different positions within a design company, allowing human designers to collaborate with specialized agents in natural language. This approach has been shown to improve designer performance compared to using separate AI tools, underscoring the value of integrating domain knowledge into multi-agent workflows [82].

However, the deployment of these frameworks is not without challenges, particularly regarding interaction design and the risk of homogenization. Studies on research question co-creation have found that while depth-first generation may seem more creative during the task, breadth-first approaches are often perceived as more creative and trustworthy upon completion. Additionally, introducing processing delays can allow users to reflect on multiple ideas simultaneously, enhancing their sense of control and the quantity of generated ideas [83]. Moreover, the risk of “epistemic monoculture” looms if these agents are trained on homogeneous benchmarks, potentially disincentivizing exploration beyond established paradigms [84]. Therefore, effective multi-agent frameworks must incorporate mechanisms for diversity, such as leveraging behavioral diversity techniques to generate a wider range of ideas, ensuring that the collective intelligence of the agent team surpasses that of a homogeneous group [85].

Ultimately, multi-agent frameworks for iterative idea refinement represent a paradigm shift towards symbiotic human-AI research ecosystems, serving as the critical intermediate step between raw hypothesis generation and evidence synthesis. By structuring interactions that mimic the critical, social, and iterative nature of human science, these systems offer a pathway to accelerate discovery while maintaining rigorous standards of validity and novelty. Yet, the efficacy of these collaborative agents is fundamentally dependent on the quality and comprehensiveness of the evidence they access. As these agents refine ideas, they require a robust, up-to-date foundation of existing literature to validate their claims—a requirement that leads directly to the necessity of automating Systematic Literature Reviews, the focus of the following section. Future advancements will likely focus on refining the social dynamics between agents, improving the alignment of agent goals with human

values, and developing robust evaluation metrics that capture the synergistic value added by these collaborative digital teams.

2.3 Automating Systematic Literature Reviews with Generative Agents

2.3 Automating Systematic Literature Reviews with Generative Agents

While multi-agent frameworks described in the preceding section excel at iterative idea refinement and hypothesis generation, their efficacy is fundamentally dependent on the quality and comprehensiveness of the underlying evidence base. Systematic Literature Reviews (SLRs) serve as the cornerstone of this evidence, providing rigorous syntheses of existing studies to answer specific research questions. However, the traditional execution of an SLR is a notoriously labor-intensive and time-consuming endeavor, often requiring months of manual effort to search, screen, extract, and synthesize data from thousands of documents. The exponential growth of scientific production has exacerbated this bottleneck, prompting a paradigmatic shift toward automation. Moving beyond the static machine learning classifiers of early attempts, the advent of Large Language Models (LLMs) and generative agents has catalyzed the development of end-to-end automated systems capable of managing the entire SLR lifecycle. These generative agents function not merely as passive tools but as active collaborators that can generate search strategies, perform inclusion/exclusion filtering, summarize abstracts, and extract data with a degree of autonomy previously unattainable, effectively bridging the gap between raw literature and the structured inputs required for the collaborative agents discussed previously.

The initial phase of any SLR involves constructing complex Boolean search queries to retrieve relevant literature from diverse databases. Historically, this step required significant domain expertise to balance precision and recall. Recent advancements demonstrate that generative AI can effectively alleviate this burden. For instance, research indicates that models like ChatGPT are capable of generating effective Boolean queries that achieve high search precision, although they may sometimes trade off recall, making them particularly valuable for rapid reviews where time is a constraint [86]. Furthermore, novel frameworks have been introduced to automate the generation of search strings based on user input topics, streamlining the retrieval process and reducing the reliance on manual query formulation [87]. These systems often integrate semantic search capabilities and natural language processing to download and organize PDF files, creating a seamless pipeline from query inception to document acquisition [88].

Once the candidate literature is retrieved, the screening phase—determining which studies meet inclusion and exclusion criteria—represents the most resource-intensive component of an SLR. Generative agents have shown remarkable proficiency in automating title and abstract screening. Studies evaluating GPT-4 in a “human-out-of-the-loop” approach reveal that, when equipped with highly reliable prompts, the model can achieve “almost perfect” performance in screening full-text literature, rivaling human accuracy in specific contexts [89]. Similarly, instruction-tuned models such as Bio-SIEVE have demonstrated superior generalization across medical domains compared to both traditional classifiers and general-purpose chatbots, effectively performing abstract screening based on detailed selection criteria [90]. To further enhance efficiency, hybrid learning approaches and active learning pipelines, such as ASReview, combine human feedback with machine learning to prioritize relevant documents, significantly reducing the number of records requiring manual review while maintaining high sensitivity [91; 92].

Beyond screening, generative agents are increasingly tasked with data extraction and synthesis. Multi-agent architectures have been designed to autonomously summarize abstracts and conduct

thorough analyses of selected papers against predefined research questions, effectively mimicking the cognitive workflow of human reviewers [87]. In the biomedical domain, frameworks leveraging distant supervision and natural language processing have been developed to automatically identify and extract specific data elements from primary publications, achieving high recall rates and drastically reducing the manual annotation burden [93]. Moreover, the integration of knowledge graphs with LLMs offers a promising avenue for classifying and visualizing review results, enabling continuous updates and deeper reasoning over the synthesized evidence [94]. This specific integration foreshadows the more profound Graph-LLM synergies discussed in the following section, where structured ontologies are used not just for organization, but for causal reasoning and hallucination mitigation.

Despite these advancements, the deployment of fully automated SLR systems faces significant fidelity challenges, primarily concerning hallucinations and the potential erosion of methodological rigor. Human evaluations of machine-generated literature reviews reveal that while some outputs are comparable to human-written reviews, they often suffer from hallucinations and a lack of detailed information [95]. Experts in medical systematic reviews have expressed valid concerns regarding confidently composed but inaccurate LLM outputs, which could lead to dangerous misinformation if left unchecked [96]. Consequently, the consensus in the field suggests that while LLMs can rival human performance in specific tasks under reliable prompting, substantial caution is required. The current trajectory favors a human-in-the-loop paradigm where generative agents handle the heavy lifting of screening and preliminary synthesis, while human researchers retain oversight for critical validation, bias mitigation, and final interpretation. This symbiotic approach ensures that the efficiency gains of automation do not come at the cost of scientific integrity, paving the way for more dynamic and up-to-date evidence syntheses that can fuel the next generation of AI-assisted discovery [97; 98].

2.4 Synthesizing Knowledge through Graph-LLM Integration

2.4 Synthesizing Knowledge through Graph-LLM Integration

Building upon the automated evidence synthesis described in the preceding section, the integration of Large Language Models (LLMs) with Knowledge Graphs (KGs) represents a critical evolution in scientific workflows: moving from the aggregation of literature to the construction of structured, verifiable, and interpretable reasoning frameworks. While the generative agents discussed previously excel at screening and summarizing vast corpora, they remain susceptible to hallucinations and often lack the explicit logical scaffolding required for rigorous causal inference. Purely parametric LLMs possess remarkable fluency but struggle with factual grounding and traceability, whereas Knowledge Graphs offer repositories of explicit facts and ontological relationships but lack the flexibility to generate natural language insights or reason over unseen concepts. The synergy between these two modalities creates a “Large Knowledge Model” architecture, where the generative capabilities of LLMs are anchored by the factual precision of KGs, directly addressing the fidelity challenges identified in automated systematic reviews and enabling robust scientific discovery [99].

A primary advantage of this Graph-LLM integration is the substantial mitigation of hallucinations through strict factual grounding, a necessary refinement following the concerns raised regarding autonomous SLR systems. Traditional retrieval-augmented generation (RAG) systems often retrieve unstructured text chunks that may contain conflicting or noisy information, perpetuating the risks of inaccurate synthesis. In contrast, KG-enhanced frameworks utilize structured triples to verify facts prior to generation. For instance, the “MindMap” framework leverages knowledge graphs

to elicit the reasoning pathways of LLMs, forcing the model to align its internal logic with an external ontology, which significantly improves performance in medical question-answering tasks and reduces the generation of unsupported claims [100]. Similarly, the “Knowledge Graph-based Retrofitting” (KGR) framework introduces an autonomous verification loop where the LLM’s initial draft responses are retrofitted against factual statements extracted from a KG, effectively mitigating hallucinations during complex reasoning processes without requiring manual intervention [101]. This mechanism is paramount in scientific writing, ensuring that the integrity of cited facts is maintained even as automation scales.

Beyond verification, structured ontological graphs significantly enhance the capacity for causal reasoning and hypothesis generation, areas where standard LLMs often falter due to their reliance on statistical correlation rather than causal mechanisms. By combining LLMs with causal knowledge graphs, researchers can automate the extraction of causal relationships from the vast literature corpora processed in the previous stage. A notable approach involves using LLMs to extract causal relation pairs from tens of thousands of psychology articles to construct a specialized causal graph, which is then used to generate novel hypotheses that rival those conceived by human experts [102]. This hybrid approach allows for the discovery of interdisciplinary insights that might remain hidden in siloed textual data. The “From Query Tools to Causal Architects” framework further demonstrates this potential by integrating LLM expertise on existing causal mechanisms with data-driven statistical analysis, creating a robust baseline for causal structure learning that surpasses traditional methods [103].

The integration also facilitates complex, multi-step reasoning over interconnected scientific concepts, overcoming the limitations of linear text processing. Standard Chain-of-Thought (CoT) prompting often fails when the reasoning path requires navigating a dense web of entity relationships found in bibliographic or domain-specific graphs. To address this, frameworks like “Graph Chain-of-Thought” (Graph-CoT) encourage LLMs to reason iteratively on the graph structure itself, engaging in a cycle of reasoning, interaction, and graph execution to answer complex queries [104]. Similarly, the “Reasoning on Graphs” (RoG) method generates faithful plans grounded in KG relation paths, which are then used to retrieve valid reasoning trajectories, ensuring that the LLM’s output is both interpretable and factually consistent [105]. These methods transform the LLM from a passive text generator into an active agent capable of navigating the topological structure of scientific knowledge, providing a structured foundation that contrasts with the unstructured qualitative scaling discussed in the following section.

Moreover, this synergy enables the effective handling of incomplete knowledge scenarios, a common challenge in emerging research fields where data is sparse and formal ontologies are not yet fully developed. The “Generate-on-Graph” (GoG) approach treats the LLM not only as an agent to explore existing graphs but also as a knowledge source to generate new factual triples when the provided KG is incomplete, effectively simulating the human ability to infer missing links based on general knowledge [106]. Additionally, frameworks like “HyKGE” utilize hypothesis outputs from LLMs to extend exploration directions within a KG, optimizing the retrieval of diverse and relevant knowledge for accurate medical responses [107]. This dynamic interaction ensures that the system remains adaptable even when the underlying evidence base is evolving.

Finally, the interpretability of scientific synthesis is greatly enhanced by this integration, offering a solution to the “black box” nature of pure LLM reasoning. Graph-LLM systems can provide explicit evidence paths, allowing researchers to trace every synthesized claim back to its ontological roots. The “Leveraging Structured Information for Explainable Multi-hop Question Answering and Reasoning” study highlights that extracting semantic structures (graphs) provides grounded

explanations that are preferred by humans over generated reasoning chains, thereby increasing trust in AI-assisted research [108]. By anchoring generative outputs in the rigid structure of knowledge graphs, the scientific community can achieve a new standard of transparency. This structured foundation for factual verification and causal reasoning serves as a crucial bridge, contrasting with the more fluid, interpretive challenges of scaling qualitative inquiry in the humanities and social sciences explored in the subsequent section.

2.5 Scaling Qualitative Analysis and Interdisciplinary Synthesis

2.5 Scaling Qualitative Inquiry in Humanities and Social Sciences

While the integration of Large Language Models (LLMs) with Knowledge Graphs discussed in the previous section provides a structured foundation for factual verification and causal reasoning in the hard sciences, the application of LLMs in the humanities and social sciences represents a distinct paradigmatic shift: moving from automation tools to powerful instruments for scaling qualitative analysis. Traditionally, qualitative inquiry—encompassing tasks such as thematic analysis, content coding, and discourse interpretation—has been constrained by the labor-intensive nature of manual data processing. The sheer volume of unstructured text available in digital archives, interview transcripts, and open-ended survey responses often exceeds the capacity of individual researchers or small teams to analyze with depth and rigor. However, the emergence of LLMs offers an unprecedented opportunity to augment human expertise with machine scalability, enabling scholars to tackle datasets of magnitudes previously unmanageable while maintaining, and in some cases enhancing, methodological transparency [109].

A primary application of this scaling capability is found in the domain of qualitative coding, where researchers assign labels to textual data to identify patterns and themes. Recent empirical studies have demonstrated that state-of-the-art LLMs can achieve intercoder reliability comparable to human experts when provided with well-structured codebooks and appropriate prompting strategies. For instance, research utilizing GPT-4 has shown excellent agreement with human-derived gold standards on complex socio-historical coding tasks, particularly when the model is prompted to employ chain-of-thought reasoning to justify its categorization decisions [110]. This suggests that the cognitive load of interpreting dense, paragraph-long passages can be effectively distributed between human scholars, who define the theoretical frameworks, and AI agents, who execute the initial labeling at scale. Furthermore, frameworks that combine expert-drafted codebooks with LLM inference have achieved fair to substantial agreement in deductive coding scenarios, validating the utility of these models as “research assistants” rather than replacements for human judgment [111].

Beyond simple categorization, LLMs are proving instrumental in detecting semantic shifts and analyzing ideological scales within vast textual corpora. The ability of these models to navigate the “conceptual clutter” inherent in social scientific measurement allows for the flexible extraction of diffuse concepts such as political ideology or power dynamics. By instructing LLMs to furnish numeric scores directly based on contextual nuances, researchers can elicit ideological scales that align closely with established methods and human judgment, thereby facilitating large-scale longitudinal studies of discourse evolution [112]. This capacity extends to the analysis of historical texts and smaller languages, where LLMs have been successfully deployed to detect lexical semantic changes and infer social networks, tasks that previously required months of team effort and complex computational pipelines [109].

The synthesis of interdisciplinary concepts also benefits significantly from LLM-assisted workflows. In fields ranging from history to public policy, researchers are leveraging mixed-methods approaches

where unsupervised topic modeling informs the creation of prompts for LLMs, which then perform nuanced thematic and sentiment analysis on stakeholder interviews. Such human-computer interactive approaches have demonstrated congruence with expert coding exceeding 96% for broader themes, significantly outperforming traditional Natural Language Processing methods and offering new perspectives on policy gaps [113]. Moreover, the application of LLMs in emulating inductive thematic analysis suggests that these models can partially reproduce the interpretive steps proposed by human analysts, identifying latent meanings in semi-structured interviews that might otherwise remain obscured in large datasets [114].

However, the scaling of qualitative analysis via LLMs introduces significant challenges and risks that parallel the concerns regarding hallucination and bias discussed in the context of autonomous discovery. Critical investigations have revealed that LLM annotations can introduce systematic biases that are not random with respect to subject characteristics, potentially leading to misleading inferences if not carefully audited [115]. For example, models may struggle to accurately represent the perspectives of specific demographic groups, such as underestimating climate change concerns among Black Americans, highlighting the necessity for meticulous conditioning and bias assessment when employing LLMs for survey simulation or population modeling [116]. Furthermore, the concept of “algorithmic fidelity”—the degree to which LLM-generated outputs mirror real human populations—remains a contentious issue. Studies indicate that while key themes generated by “silicon participants” may resemble those of humans, the structure and tone of their responses often differ markedly, raising questions about the generalizability of findings derived solely from AI-generated data [117]. Additionally, the inherent limitations of LLMs in portraying diverse identity groups suggest that they cannot fully replace human participants in research where positionality and social identity are central to the inquiry [118].

To mitigate these risks and ensure methodological rigor, the field is moving toward “prompt science,” a systematic approach to prompt construction that emphasizes replicability, objectivity, and transparency akin to traditional codebook development. This involves multi-phase verification processes where researchers deliberate and document their prompting strategies to remove subjectivity [119]. Visual analytics tools are also being developed to assist non-expert users in refining prompts through perturbation and testing, thereby reducing the cognitive overhead associated with optimizing LLM interactions for specific qualitative tasks [120]. Ultimately, the most effective workflows appear to be those that adopt a “human-in-the-loop” or “LLM-in-the-loop” philosophy, where the AI handles the heavy lifting of data processing and initial coding, allowing human researchers to focus on higher-order interpretive tasks, theory building, and the critical validation of machine-generated insights [121]. This symbiotic relationship ensures that the scalability of AI augments rather than erodes the depth and nuance that define rigorous qualitative research, setting the stage for the broader discussion on benchmarking, validation, and the limits of autonomous discovery that follows.

2.6 Benchmarking, Validation, and the Limits of Autonomous Discovery

2.6 Benchmarking, Validation, and the Limits of Autonomous Discovery

Building upon the “human-in-the-loop” and “prompt science” methodologies established for scaling qualitative inquiry in the humanities, the transition of artificial intelligence from a supportive assistant to an active collaborator in hypothesis generation and literature synthesis demands an even more rigorous framework for evaluation. As AI systems increasingly participate in the scientific ideation process, the imperative to rigorously assess their outputs becomes paramount. The

integration of Large Language Models (LLMs) and generative agents introduces a new layer of complexity to scientific validation, as current evaluation frameworks often rely on traditional task-oriented metrics that are insufficient for assessing the nuanced quality of AI-generated research ideas. These ideas require high levels of novelty, feasibility, and logical coherence—dimensions that simple accuracy metrics fail to capture. Consequently, the field is grappling with a significant “AI Knowledge Gap,” where the proliferation of unique AI systems outpaces the number of studies characterizing their actual behavior and reliability [122]. This disparity necessitates the development of robust benchmarking standards that move beyond superficial correctness to capture the multifaceted nature of scientific discovery.

A primary challenge in benchmarking autonomous discovery lies in the performance gaps observed in complex reasoning tasks. While AI models have demonstrated superhuman capabilities in specific, bounded domains, their ability to maintain logical consistency across multi-step scientific arguments remains fragile. When models operate at a level where human experts can no longer serve as reliable proxies for ground truth, traditional evaluation methods fail. To address this, researchers have proposed frameworks based on consistency checks, which aim to surface logical inconsistencies in decision-making even when the absolute correctness of a hypothesis is difficult to verify [123]. For instance, an AI might generate a scientifically plausible hypothesis that contradicts its own previous assertions or violates fundamental physical laws when tested against varying parameters. Furthermore, the evaluation of AI capabilities must evolve from purely task-based assessments to ability-oriented evaluations that consider cognitive dimensions similar to those used in human and animal psychology [124]. Without such rigorous, ability-focused benchmarking, there is a significant risk of overestimating the readiness of AI systems for autonomous scientific roles.

The risks of algorithmic bias and hallucination pose perhaps the most critical threats to the integrity of AI-augmented research, echoing but amplifying the concerns raised regarding qualitative coding in social sciences. Generative AI tools, trained on vast and often uncensored corpora, are susceptible to producing “context-induced algorithmic bias,” where subtle variations in prompting and context lead to unintended and potentially harmful outputs [125]. These models can inadvertently reinforce existing societal stereotypes or scientific misconceptions present in their training data, leading to a homogenization of scientific thought rather than the diversification of ideas. Moreover, the phenomenon of hallucination—where models generate confident but factually incorrect statements—remains a pervasive issue. In high-stakes scientific contexts, hallucinations can manifest as fabricated citations, non-existent experimental results, or logically flawed causal mechanisms. Recent efforts have introduced automated scalable frameworks, such as the HypoTermQA dataset, to benchmark the hallucination tendencies of LLMs by generating tasks related to hypothetical phenomena, revealing that even state-of-the-art models exhibit significant error rates in these specialized contexts [126].

Compounding these technical limitations is the problem of “groundless truth” in human-AI decision-making. Predictive models often optimize for proxy labels that are readily available in datasets rather than the latent constructs of genuine scientific interest, such as disease status or true causal relationships [127]. This reliance on simplistic proxies can lead to valid-looking but scientifically vacuous hypotheses that fail to address the core phenomena under investigation. The stochastic nature of AI models further complicates verification, making it difficult to reproduce results or certify the safety and reliability of autonomous discovery systems [128]. Consequently, the notion of fully autonomous discovery remains elusive; the current technological landscape suggests that AI functions best not as a replacement for human intuition, but as a component within a broader, human-in-the-loop ecosystem.

The critical role of human validation cannot be overstated in ensuring scientific integrity, reinforcing the symbiotic relationship highlighted in previous sections. Empirical studies indicate that while AI can generate code or draft hypotheses, substantial human intervention is required to validate accuracy and ensure logical soundness [129]. The “human-in-the-loop” paradigm serves as a necessary checkpoint to mitigate automation bias, where users might otherwise over-rely on algorithmic suggestions without sufficient critical scrutiny [130]. Effective validation strategies involve iterative feedback loops where human experts refine AI outputs, correct hallucinations, and inject domain-specific context that models lack. Approaches such as “Deliberative AI” encourage active discussion between humans and machines, allowing for the elicitation of conflicting opinions and the refinement of decisions through dimension-level analysis, thereby improving appropriate reliance and task performance [131].

Furthermore, the development of hallucination-minimized solutions requires a structured journey involving prototyping, scaling, and continuous evolution via human feedback, particularly in domains requiring high precision like finance and healthcare [132]. Auditing large language models using human-in-the-loop methodologies has also proven effective in exposing inconsistencies and biases that automated tests might miss [133]. Ultimately, while AI offers transformative potential for accelerating ideation and synthesizing knowledge, the limits of autonomous discovery are currently defined by the need for rigorous, human-centered validation frameworks. The future of scientific collaboration lies not in the pursuit of full autonomy, but in the design of symbiotic systems where human expertise guides AI creativity, ensuring that the resulting scientific outputs are novel, reliable, and ethically sound.

3 Transforming Scientific Writing, Authorship, and Peer Review

3.1 The Shift from Material Production to Critical Integration in Drafting

The integration of generative artificial intelligence (GenAI) into academic writing signifies a profound paradigmatic shift, moving the scholarly process away from traditional linear material production toward a dynamic, iterative cycle of critical integration. Historically, academic writing was conceptualized as a sequential endeavor involving distinct phases of planning, translation, and revision, where the primary labor resided in the manual generation of text. However, the advent of large language models (LLMs) has disrupted this trajectory, fundamentally altering the nature of authorship. The researcher’s role is increasingly transitioning from that of a primary producer of raw textual content to a curator, verifier, and synthesizer of AI-generated outputs. This transformation redefines the writing workflow not as an act of creation from scratch, but as a complex process of sensemaking and knowledge transformation where human intellect is applied to the refinement and validation of machine-proposed drafts.

This shift is best understood through the lens of “critical integration,” where the value added by the human author lies less in the initial formulation of sentences and more in the rigorous evaluation and contextual grounding of AI contributions. Recent scholarship suggests that generative AI capabilities are reshaping knowledge work by shifting the focus from material production to critical integration, emphasizing the importance of the creative and curatorial voice of the user over the mere generation of text [134]. In this new model, the drafting process becomes a collaborative dialogue rather than a solitary monologue. The AI acts as a rapid prototyper of ideas, generating outlines, expanding arguments, or synthesizing literature, while the human researcher engages in a higher-order cognitive task of assessing factual accuracy, logical coherence, and alignment with the specific nuances of the research domain.

The necessity for this critical stance is underscored by empirical evidence regarding the limitations of current models in producing deep, original scientific content. Studies comparing AI-generated texts with human-written manuscripts reveal that while LLMs excel at merging existing knowledge and producing fluent prose, they often struggle with original scientific contributions and may exhibit a “shallow synthesis” of related works without significant human intervention [135]. Furthermore, evaluations of various chatbots indicate that while they can achieve high quantitative accuracy in factual retrieval, they frequently fail to replicate the intricate originality required for high-impact humanities and scientific research [136]. Consequently, the efficacy of this collaborative model hinges on the researcher’s ability to engage in “prompt literacy” and iterative refinement, possessing the skills to generate precise inputs, interpret ambiguous outputs, and recursively guide the AI toward desired scholarly standards [137].

This dynamic cycle of critical integration also introduces new challenges regarding the cognitive processes involved in writing. Traditional models of writing distinguish between “exploration” (planning and idea generation) and “exploitation” (translating ideas into text). AI tools have the potential to accelerate the exploitation phase significantly, yet there is a risk that over-reliance on these tools may atrophy the exploratory phases essential for deep learning and creativity. Experimental evidence suggests that complete reliance on AI for writing tasks can lead to significant reductions in learning outcomes and comprehension accuracy, highlighting the danger of passive consumption of AI-generated text [138]. Therefore, effective human-AI collaboration requires a balanced approach where the researcher actively engages in the “knowledge-telling” and “knowledge transformation” processes, using AI as a scaffold rather than a substitute for cognitive effort [139].

Moreover, the shift toward critical integration demands a re-evaluation of how writing quality and authenticity are assessed. The emergence of “hybrid” texts, co-authored by humans and machines, complicates the detection of authorship and necessitates new frameworks for evaluating scholarly integrity. Researchers must now navigate a landscape where the boundary between human and machine contribution is increasingly blurred, requiring transparent reporting mechanisms such as “PaperCard” or “Cardwriter” to disclose the extent of machine assistance [140; 141]. The integrity of the scientific record depends on the human author’s ability to maintain accountability, ensuring that the final manuscript reflects not just the statistical probability of the next token, but a verified, reasoned, and ethically sound argument.

Ultimately, the transformation of academic writing from material production to critical integration represents both an opportunity and a responsibility. It offers the potential to alleviate the communication burden, accelerate the dissemination of discoveries, and empower authors to focus on high-level conceptual synthesis [142]. However, realizing these benefits requires a fundamental change in scholarly habits. Researchers must cultivate a stance of skeptical engagement, treating AI outputs as provisional drafts that require rigorous verification against primary sources and logical scrutiny. By adopting this critical integration model, the scientific community can harness the efficiency of generative AI while preserving the depth, originality, and accountability that are the hallmarks of robust scientific inquiry. The future of scientific writing will likely be defined not by who writes the most, but by who can most effectively curate, verify, and synthesize the vast generative capacities of artificial intelligence into coherent, trustworthy knowledge, thereby laying the groundwork for the advanced prompt literacies and interactive co-authorship models discussed in the subsequent section.

3.2 Prompt Literacy and Interactive Co-Authorship Models

3.2 Prompt Literacy and Interactive Co-Authorship Models

Building upon the paradigmatic shift from material production to critical integration outlined previously, the operational realization of human-AI collaboration hinges on the development of “prompt literacy” and the adoption of interactive co-authorship models. As the researcher’s role evolves into that of a curator and verifier, the ability to effectively direct AI systems becomes a foundational competency rather than a mere technical skill. Prompt literacy in scientific research is defined as the sophisticated capacity to generate precise inputs, critically interpret stochastic outputs, and iteratively refine dialogue strategies to achieve rigorous scholarly outcomes. This literacy encompasses a cyclical process of drafting, refining, and reflecting, enabling users to manifest imagined outcomes while recognizing the inherently collaborative nature of the interaction [137]. Without this proficiency, researchers risk engaging in superficial interactions that yield shallow syntheses or, more critically, inadvertently propagate factual hallucinations, thereby undermining the very integrity they seek to preserve.

The acquisition of prompt literacy is neither intuitive nor automatic; it represents a novel creative skill that demands deliberate practice and structural understanding. Research indicates that possessing domain knowledge does not automatically translate to the ability to craft effective prompts, as researchers often lack the specific vocabulary required to constrain generative models appropriately [143]. In scientific writing, this gap can result in outputs that are grammatically fluent yet scientifically vacuous or factually incorrect. Consequently, effective co-authorship requires a transition from passive acceptance to “critical interaction,” where the researcher actively verifies suggestions and engages deeply with the tool’s limitations [144]. This critical stance is indispensable, given that generative models remain prone to hallucinations and context-induced biases that can compromise the validity of scientific arguments if left unchecked [125].

To address these challenges, the scientific writing workflow is evolving toward interactive co-authorship models characterized by mixed-initiative dialogues. Unlike traditional static tools, modern generative systems facilitate dynamic exchanges where both human and AI can initiate content changes, fostering a fluid and responsive writing environment [145]. In this conversational loop, the researcher poses hypotheses or outlines, the AI generates draft content or counter-arguments, and the researcher refines the trajectory based on critical evaluation. This iterative prompting strategy is vital for enhancing the truthfulness and accuracy of the final output; studies demonstrate that while naive prompting methods can undermine reliability, structured approaches involving step-by-step refinement markedly improve the calibration and factual grounding of generated text [146].

The efficacy of these co-authorship models is further mediated by the design of the interaction interfaces. Advanced frameworks, such as conversational Explainable AI (XAI), empower researchers to query the rationale behind AI-generated suggestions, thereby fostering deeper understanding and trust. Systems equipped with multifaceted explanations and context-aware drill-down capabilities allow users to scrutinize the AI’s reasoning, ensuring alignment with scientific standards [147]. However, the level of scaffolding provided must be carefully balanced; while high levels of support can boost productivity, particularly for less experienced writers, they may simultaneously erode the user’s sense of text ownership and satisfaction [148]. Thus, adaptive systems that tailor support to the researcher’s expertise and the specific phase of the writing process are essential for optimizing the human-AI partnership.

Furthermore, sophisticated prompt engineering serves as a crucial mechanism for managing the inherent trade-off between novelty and usefulness in creative scientific inquiry. Over-emphasizing novelty can precipitate hallucinations, whereas an exclusive focus on usefulness may yield memorized, unoriginal content. By employing advanced strategies such as chain-of-thought reasoning or persona-based simulation, researchers can guide AI models to generate content that is both innovative and grounded in established evidence [149]. Specifically, techniques like structured chain-of-thought prompting have been shown to increase faithfulness to grounding documents, thereby reducing the incidence of fabrication in content generation [150].

Ultimately, the future of scientific writing rests on a symbiotic relationship between human critical thinking and AI generative power, contingent upon the researcher’s proficiency in prompt literacy and the availability of tools that support transparent, iterative workflows. This relationship necessitates a shift from viewing AI as a mere text generator to treating it as a collaborative partner requiring careful guidance, continuous monitoring, and rigorous validation. By cultivating these literacies and adopting robust co-authorship models, the scientific community can harness the efficiency of generative AI while safeguarding the epistemic foundations of research [142]. Nevertheless, this enhanced agency introduces complex questions regarding the psychological sense of ownership and the attribution of credit. Phenomena such as the “AI Ghostwriter Effect,” where users feel disconnected from the authorship of AI-assisted text despite self-declaring as authors, highlight the tension between collaborative efficiency and personal accountability [151]. These emerging challenges regarding agency and credit assignment naturally lead to the critical examination of authorship definitions and transparent disclosure frameworks discussed in the following section.

3.3 The Crisis of Authorship Attribution and Credit Assignment

3.3 Authorship, Accountability, and Transparent Disclosure Frameworks

While the preceding section established prompt literacy and interactive co-authorship as the operational mechanics of human-AI collaboration, these practices inevitably precipitate a profound crisis regarding the fundamental concepts of authorship and credit assignment. Traditionally, authorship has served as a proxy for intellectual responsibility, accountability, and the tangible contribution of human agency to the creation of knowledge. However, the advent of Large Language Models (LLMs) capable of drafting coherent arguments, synthesizing literature, and even proposing novel hypotheses challenges the ontological boundaries of what it means to be an “author.” This transition from the *how* of collaboration to the *who* of attribution necessitates a critical evaluation of the philosophical and legal complexities arising from this shift, examining whether AI systems can ever merit co-authorship status and analyzing emerging frameworks designed to restore transparency through standardized disclosure mechanisms.

At the heart of this authorship crisis lies a philosophical re-evaluation of creativity and originality. Critics often characterize generative AI outputs as mere “stochastic parrots” or randomized plagiarism, suggesting that these systems lack the intentionality required for true authorship. Conversely, alternative perspectives argue that creativity is not an intrinsic property of an object but rather a property of a process, an author, or a viewer, heavily dependent on community definitions [134]. As AI tools transition from passive editors to active collaborators—facilitated by the mixed-initiative dialogues discussed previously—the nature of knowledge work shifts from material production to critical integration. In this new paradigm, the human researcher’s primary role evolves into that of a curator and verifier, raising difficult questions about how credit should be allocated when the “writing” is performed by an algorithm [134]. Furthermore, empirical studies indicate that while

AI assistance boosts productivity and confidence, it may simultaneously erode the user’s sense of ownership and accountability, creating a psychological disconnect between the researcher and the final text [152].

Legally, the landscape surrounding these hybrid workflows is equally fragmented and uncertain. Current intellectual property frameworks in many jurisdictions, including the United States, generally deny copyright protection to works generated solely by non-human entities, asserting that authorship is inherently a human endeavor [153]. This legal stance creates a precarious situation for researchers who rely heavily on AI; if significant portions of a manuscript are AI-generated, the work might fall into the public domain or face challenges regarding ownership and liability. The debate extends to the training data itself, with ongoing disputes over whether the use of copyrighted works to train generative models constitutes “fair use” or misappropriation, a conflict that directly impacts the legitimacy of the resulting outputs [154]. Moreover, the complexity of proving causation in cases of defamation or product liability involving AI-generated content further complicates the assignment of legal responsibility, suggesting that existing legal systems are ill-equipped to handle the nuances of algorithmic agency [155].

In response to these ambiguities, the scientific community has largely moved to prohibit listing AI systems as co-authors, recognizing that machines cannot assume ethical responsibility or accountability for the research findings. A bibliometric analysis of top academic publishers and journals reveals that the vast majority explicitly forbid AI authorship, yet there remains substantial heterogeneity in how the use of such tools must be disclosed [156]. To address this lack of standardization and ensure research integrity, several frameworks have been proposed to mandate transparent reporting of machine assistance, serving as a crucial bridge between the collaborative workflows of Section 3.2 and the forensic detection challenges outlined in the subsequent section.

One prominent proposal is the “PaperCard” framework, which serves as a structured documentation method for human authors to declare the specific roles AI played in their writing process [140]. This approach moves beyond binary declarations of use, encouraging researchers to detail the extent of AI involvement in tasks ranging from idea generation to text polishing. Similarly, the concept of “AI Usage Cards” has been introduced as a standardized reporting mechanism grounded in a three-dimensional model of transparency, integrity, and accountability [157]. These cards aim to provide a common language for the research community to trace, compare, and question various forms of AI usage, thereby supporting the development of accepted community norms. The utility of such cards lies in their ability to make the “black box” of AI interaction visible, allowing reviewers and readers to assess the provenance of the ideas and text presented.

Complementing these card-based systems, other scholars have advocated for “Three Disclaimers for Safe Disclosure,” offering an intuitive interface for authors to generate short reports on their generative AI usage [141]. These initiatives collectively argue that while AI cannot be an author, its contribution must be meticulously documented to prevent the erosion of trust in scientific literature. The European AI Act also underscores the necessity of such transparency, with Article 52 mandating disclosure obligations for AI-generated content, although the specific interpretation of these obligations in academic contexts remains a subject of ongoing debate [158].

Ultimately, resolving the crisis of authorship requires a shift from viewing authorship as a static label to seeing it as a dynamic spectrum of contribution and responsibility. As generative AI capabilities continue to expand, the academic community must adopt robust governance frameworks that balance the efficiency gains of AI assistance with the imperative of maintaining human accountability. By implementing standardized disclosure tools like PaperCards and AI Usage Cards,

the scientific enterprise can navigate the transition toward hybrid intelligence while preserving the core values of honesty, attribution, and integrity that underpin scholarly communication. However, as the following section will demonstrate, even with robust disclosure policies, the technical reality of distinguishing between human and machine contributions in these hybrid texts presents a formidable forensic challenge, rendering simple declarations insufficient without verifiable detection mechanisms.

3.4 Forensic Challenges in Detecting Hybrid and AI-Generated Texts

3.4 Forensic Challenges in Distinguishing Human, Machine, and Hybrid Authorship

While the preceding section established transparent disclosure frameworks like PaperCards and AI Usage Cards as essential governance mechanisms, these policy-based solutions rely heavily on the honesty and accuracy of human reporting. However, as the following section will detail regarding algorithmic bias, the integrity of scientific discourse depends not only on declared intent but also on the verifiable provenance of content. This creates a critical gap: without robust technical means to validate disclosures, the scientific community faces a complex forensic challenge in distinguishing between pure human authorship, fully machine-generated text, and the increasingly prevalent hybrid texts resulting from iterative human-AI collaboration. As researchers integrate generative AI into their workflows, the traditional binary classification of text as either “human” or “AI” becomes insufficient. The reality of modern scientific production often involves a mosaic of authorship where humans select, edit, and refine AI-generated segments, creating a continuum of contribution that confounds traditional detection mechanisms. This subsection surveys the fragility of current detection tools, the specific difficulties posed by hybrid texts, and the limitations of emerging solutions like watermarking and self-detection, illustrating why simple declarations are insufficient without verifiable detection mechanisms.

The primary hurdle in this forensic landscape is the diminishing distinguishability between human and machine output. Early detectors relied on statistical anomalies or perplexity scores, but these methods degrade rapidly as models improve. Recent evaluations indicate that state-of-the-art detection tools often fail to reliably differentiate between human-written scientific content and AI-generated text, particularly when the AI is instructed to mimic human stylistic nuances [159]. In academic settings, where precision and specific domain knowledge are paramount, detectors have shown a tendency to produce high false positive rates, potentially accusing honest researchers of misconduct, or conversely, failing to identify sophisticated AI usage [160]. The problem is exacerbated by the fact that human annotators themselves struggle to identify AI-generated content, with studies showing that humans are often only slightly better than random guessing when faced with high-quality deepfake texts [161].

A critical and under-explored dimension of this crisis is the detection of hybrid texts, which directly challenges the assumptions of the disclosure frameworks discussed previously. Unlike purely synthetic content, hybrid texts involve seamless transitions between human and machine authorship within a single document. Research utilizing realistic datasets of human-AI collaborative writing reveals that detecting AI-generated sentences within these hybrid contexts is exceptionally challenging [162]. The difficulty arises from several factors: human editors often modify AI output to align with personal preferences, obscuring the original stylistic signatures; frequent shifts in authorship between neighboring sentences create noise for segmentation algorithms; and the resulting segments are often too short to provide sufficient stylistic cues for reliable classification [162]. Consequently, standard detectors that assume uniform authorship across a document fail to identify the specific

boundaries where AI contribution begins and ends, necessitating new approaches focused on boundary detection rather than whole-document classification [163]. This inability to pinpoint specific AI contributions undermines the granularity required by frameworks like the “Three Disclaimers,” which ask authors to specify the extent of AI involvement in various tasks.

Furthermore, the robustness of existing detection frameworks is severely compromised by adversarial attacks, rendering many automated verification attempts futile. Simple techniques such as paraphrasing, translation, or minor perturbations can effectively evade most current detectors. Studies have demonstrated that recursive paraphrasing attacks can break a wide range of detection systems, including those employing watermarking schemes and neural network-based classifiers, while maintaining the semantic integrity and quality of the text [164]. Advanced paraphrasing models, capable of reordering content and controlling lexical diversity, have been shown to drop detection accuracy to near-zero levels for several prominent detectors [165]. Even more concerning is the vulnerability of detectors to “human-written perturbations,” where minor, naturalistic changes to spelling or syntax—common in real-world editing—can fool systems trained on clean data [166]. The ease with which these attacks can be executed, sometimes compromising models in seconds, highlights the insufficiency of current defensive measures [167].

In response to these vulnerabilities, the research community has proposed watermarking and self-detection as potential solutions, though both face significant limitations that prevent them from serving as standalone guarantees of integrity. Watermarking involves embedding invisible patterns into generated text to facilitate identification. However, theoretical and empirical analyses suggest that watermarks are fragile; they can be removed or forged through post-processing attacks that add imperceptible perturbations, effectively erasing the watermark while preserving content quality [168]. Moreover, watermarking schemes can be susceptible to spoofing attacks, where adversaries manipulate inputs to cause detectors to misclassify human text as AI-generated, posing reputational risks to authors [164]. The effectiveness of watermarks also varies significantly across different models and tasks, with some conditional generation tasks suffering from performance degradation when watermarks are applied [169].

Self-detection, the capability of an LLM to identify its own output, offers another avenue but yields inconsistent results. Empirical studies reveal vast disparities in self-detection capabilities among major models; for instance, while some models demonstrate high accuracy in recognizing their own generations, others appear entirely incapable of doing so [170]. This inconsistency renders self-detection unreliable as a standalone forensic tool. Additionally, the “Counter Turing Test” framework demonstrates that as models scale in size, their outputs become increasingly indistinguishable from human text, leading to a higher “AI Detectability Index” that correlates with reduced detectability [171].

Ultimately, the current forensic landscape is characterized by an arms race where detection methods lag behind generation and evasion capabilities. While some specialized approaches, such as leveraging intrinsic dimensionality of text embeddings or retrieving semantically similar generations from a database, show promise in specific scenarios [172; 165], no universal solution exists. The complexity of hybrid authorship, combined with the susceptibility of detectors to adversarial manipulation and the limitations of watermarking, suggests that ensuring research integrity will require a shift from purely automated detection toward transparent reporting frameworks and human-in-the-loop verification processes. This forensic uncertainty sets the stage for the next layer of concern: even if authorship could be perfectly distinguished, the widespread reliance on these models introduces deeper risks of algorithmic bias, homogenization, and the erosion of originality in scientific thought, as explored in the following section.

3.5 Algorithmic Bias, Homogenization, and the Erosion of Originality

3.5 Algorithmic Bias, Homogenization, and the Erosion of Originality

While the preceding section highlighted the forensic difficulties in distinguishing human from machine authorship, the integration of Large Language Models (LLMs) into scientific writing introduces a deeper, more insidious challenge: the potential for algorithmic bias to distort knowledge production, the homogenization of scientific discourse, and the erosion of originality. As researchers increasingly delegate drafting and editing tasks to generative AI, the distinct voices that characterize diverse scientific communities risk being smoothed into a standardized, algorithmic average. This phenomenon represents not merely a stylistic concern but a fundamental shift in the epistemic landscape of science, where the very nature of knowledge production may be altered by the inherent biases and convergence tendencies of the models themselves.

A primary concern is the emergence of context-induced algorithmic bias, a subtle yet pervasive distortion that arises when generative AI tools are used to compose research manuscripts. Unlike earlier AI models trained on finite, specific datasets where biases were often human-recognizable, modern LLMs are trained on vast, imperceptibly large corpora that embed complex societal and cognitive biases. When researchers utilize these tools, they inadvertently introduce a new type of bias that stems from the interaction between the user’s prompts and the model’s probabilistic generation processes [125]. This context-induced bias is particularly dangerous because it operates under the guise of objective assistance, potentially reinforcing systemic inequalities or skewing scientific interpretations without the author’s awareness. Furthermore, LLMs have been shown to harbor deep-rooted biases against various socio-demographic groups, which can surface unexpectedly when the model is assigned specific personas or tasked with reasoning, leading to erroneous presumptions and performance drops that compromise the integrity of the generated content [173].

Beyond individual instances of bias, the widespread adoption of a limited set of foundation models threatens to create an “algorithmic monoculture” within the scientific community. When multiple researchers and institutions rely on the same underlying models—such as those powering ChatGPT or similar systems—the diversity of thought and expression in scientific literature faces the risk of outcome homogenization. Research indicates that when decision-makers share components like training data or specific models, the resulting outputs tend to become increasingly homogeneous, potentially institutionalizing systemic exclusion and reinscribing social hierarchies within the literature [174]. In the context of writing, this manifests as a reduction in content diversity; studies have found that co-writing with feedback-tuned LLMs results in statistically significant reductions in lexical and content diversity, making the writings of different authors more similar to one another [175]. This convergence towards a central distribution of language and ideas could stifle the unconventional phrasing and novel theoretical frameworks that often drive scientific breakthroughs.

The erosion of originality is further exacerbated by the tendency of LLMs to generate outputs that mimic well-considered opinions while lacking genuine human reasoning, a phenomenon that challenges the essence of authorship and democratic discourse [176]. As models are optimized for fluency and alignment with human feedback, they may prioritize safe, consensus-driven narratives over risky, innovative hypotheses. This dynamic creates a feedback loop where the “center” of the knowledge distribution is repeatedly reinforced, leading to a process known as “knowledge collapse,” where public beliefs and scientific outputs drift further from the truth due to the recursive reliance on AI-generated content [177]. The risk is that science becomes a reflection of the model’s training data rather than a dynamic exploration of the unknown, with the “writing style” gap between AI and humans revealing that AI-generated content often lacks the depth and argumentative logistics

of high-quality human writing [178].

Moreover, the suppression of diverse writing styles entails a loss of cognitive diversity, not just aesthetic variety. Different cultural and disciplinary backgrounds often express scientific concepts through unique linguistic structures and rhetorical strategies. However, LLMs have been observed to portray socially subordinate groups as more homogeneous, narrowing the range of human experience described in text and potentially perpetuating stereotypes [179]. If scientific writing becomes dominated by these homogenized outputs, the field risks losing the rich tapestry of perspectives necessary for robust peer review and critical analysis. The danger is compounded by the fact that users often rely heavily on LM generations regardless of their certainty, and the models themselves are reluctant to express uncertainty, creating an illusion of definitive knowledge that discourages critical scrutiny [180].

Ultimately, while generative AI offers tools to enhance efficiency, uncritical reliance poses a threat to the vibrancy of scientific inquiry. The potential for “laissez-faire” harms, where open-ended prompting leads to the subordination and stereotyping of minoritized identities in generated text, underscores the need for vigilance [181]. To preserve the integrity and diversity of scientific discourse, the community must recognize that the convenience of AI assistance comes with the cost of potential homogenization. Strategies must be developed to ensure that AI serves as a catalyst for diverse human creativity rather than a filter that standardizes scientific thought into a singular, algorithmic voice. These concerns regarding the quality and diversity of AI-assisted output directly inform the evolving challenges in peer review, where reviewers must now evaluate not only the scientific merit of a manuscript but also navigate the complexities of algorithmic bias and the authenticity of the presented arguments.

3.6 Evolving Peer Review Processes and AI-Assisted Evaluation

3.6 Evolving Peer Review Processes and AI-Assisted Evaluation

The concerns regarding algorithmic bias, homogenization, and the erosion of originality outlined in the preceding section directly inform the evolving challenges in peer review, where the integrity of scientific discourse is now contingent on the ability to detect and evaluate AI-mediated contributions. The integration of artificial intelligence into the peer review ecosystem represents a paradigm shift, positioning AI not merely as a passive tool but as an active participant with a dual role: it is both a subject of intense scrutiny regarding the integrity of manuscripts—particularly those suffering from the “algorithmic monoculture” and lack of uncertainty expression previously discussed—and a potent instrument for assisting human reviewers in an era of information overload. As the volume of scientific publications grows exponentially, the traditional peer review system faces a crisis of scalability, characterized by reviewer fatigue, inconsistent quality, and significant delays. In response, the research community is increasingly exploring the potential of Large Language Models (LLMs) to augment, and in some cases, partially automate the evaluation process. However, this transition is fraught with complex ethical, epistemic, and practical challenges that necessitate a careful re-evaluation of the norms governing scholarly communication.

Recent pilot studies have begun to empirically assess the capability of generative AI to produce high-quality peer reviews, offering a potential counterbalance to the efficiency losses caused by the need for deeper forensic analysis of manuscripts. For instance, investigations into the use of GPT-4 have provided initial evidence that AI-generated reviews can achieve a level of helpfulness comparable to human reviewers in specific contexts, particularly in identifying structural issues and summarizing core ideas [182]. These studies suggest that AI can effectively address resource

constraints by generating first-pass reviews that cover a broader range of paper aspects than typical human reviews, although they may currently lack the constructiveness and nuanced critical insight required for final decision-making [183]. Furthermore, experimental evaluations in specialized domains such as cybersecurity have demonstrated that machine learning models, particularly those utilizing Doc2Vec approaches, can predict review outcomes with high accuracy, sometimes outperforming general-purpose LLMs like ChatGPT [184]. These findings indicate a promising trajectory where AI serves as a preliminary filter or a supportive co-reviewer, potentially reducing the cognitive load on human experts who must now also guard against the homogenized arguments and hidden biases described earlier.

Beyond generating full reviews, AI is being deployed to scaffold the reviewing process, particularly for novice researchers who may lack familiarity with community standards or the specific rhetorical strategies needed to identify subtle algorithmic distortions. Systems like ReviewFlow have been developed to provide contextual reflections, in-situ knowledge support, and synthesis tools that guide inexperienced reviewers in annotating submissions and structuring their critiques [185]. Such intelligent scaffolding helps align novice outputs with community expectations, thereby improving the comprehensiveness and quality of reviews while mitigating the variability inherent in human evaluation. Additionally, preference learning frameworks are being utilized to assist area chairs and program committees in making final acceptance decisions. By formalizing the decision-making process as a ranking problem that incorporates both review scores and textual content, these systems aim to alleviate bias and noise, offering a more robust mechanism for evaluating the multi-faceted quality of scientific work [186].

However, the automation of evaluation raises profound ethical concerns regarding the legitimacy and transparency of the peer review process, echoing the fears of “knowledge collapse” and the loss of cognitive diversity. The legitimacy of AI-driven review hinges on its alignment with the scientific ethos, encompassing moral and epistemic norms that define appropriate conduct within the scholarly community [187]. Critics argue that while AI can enhance efficiency, it risks compromising the social and cultural functions of peer review, which serve not only to validate knowledge but also to foster community dialogue and mentorship. There is a notable “norm-counternorm continuum” wherein the acceptability of AI intervention is shaped by institutional logics and existing regulatory mechanisms. Moreover, the reliance on AI introduces risks of algorithmic bias and the potential homogenization of scientific discourse, where diverse perspectives might be suppressed in favor of statistically probable but less innovative critiques, effectively creating a feedback loop that reinforces the very biases identified in the writing phase.

The impact of AI-written manuscripts further complicates the reviewer’s task, creating a dynamic arms race between generation and detection. As authors increasingly utilize AI to draft manuscripts, reviewers must remain vigilant against hallucinated references, superficial arguments, and fabricated data, issues that are prevalent in current generative models [188]. This shift places a new burden on reviewers to act as forensic auditors of AI-generated content, potentially increasing rather than decreasing their workload if adequate detection tools are not available. The fragility of current detection mechanisms and the complexity of identifying authorship in hybrid human-AI texts exacerbate these challenges, threatening the credibility of the publication record and underscoring the urgent need for the standardized reporting frameworks discussed in the following section.

To navigate these complexities, the community is exploring incentive structures and governance models to ensure responsible AI adoption, laying the groundwork for the broader regulatory ecosystems to come. Proposals include the implementation of rating systems and repeated endogenous matching rules to incentivize high-effort reviewing and mitigate moral hazard in peer review systems

[189]. Additionally, there is a growing call for financial compensation and professional recognition to treat reviewing as a formal professional service, thereby avoiding a “tragedy of the commons” where the quality of review collapses under the weight of increased submission volumes [190]. Ultimately, the future of peer review likely lies in a hybrid model where AI handles routine screening, consistency checks, and preliminary summarization, while human experts retain authority over novelty assessment, ethical judgment, and the final validation of scientific contributions. This symbiotic relationship requires rigorous benchmarking, transparent disclosure of AI usage, and a continuous dialogue on the ethical boundaries of algorithmic assistance in the stewardship of scientific knowledge, directly motivating the move toward standardized governance and responsible reporting frameworks.

3.7 Towards Standardized Governance and Responsible Reporting Frameworks

3.7 Towards Standardized Governance and Responsible Reporting Frameworks

Building upon the hybrid peer review models and ethical challenges outlined in the previous section, the rapid integration of generative artificial intelligence into scientific workflows has precipitated a critical juncture in the broader governance of research integrity. As Large Language Models (LLMs) and other generative tools transition from experimental aids to standard components of the research lifecycle, the scientific community currently faces a fragmented landscape of policies regarding their use, disclosure, and authorship. While individual publishers and institutions have begun to draft ad hoc guidelines, the absence of a unified global framework risks creating a “Babel Tower Effect,” where conflicting standards lead to confusion rather than the desired standardization of ethical practices [191]. This subsection synthesizes current global efforts to establish consensus-driven governance, contrasting them with the disparate regulations that currently threaten the coherence of scientific communication.

The prevailing challenge lies in the heterogeneity of existing policies, which directly impacts the efficacy of the peer review processes discussed earlier. Various journals and conferences have adopted divergent stances, ranging from outright bans on AI-generated text to permissive approaches that lack rigorous disclosure requirements. This fragmentation underscores the urgent need for consensus-based standards that transcend disciplinary and geographical boundaries. In response, initiatives such as the CANGARU guidelines have emerged to foster a cross-disciplinary, global inclusive consensus on the ethical use and reporting of generative AI in academia. These efforts utilize systematic reviews, bibliometric analyses of existing author guidelines, and Delphi surveys to identify common rules and formulate principled frameworks for disclosure [191]. The goal is to move beyond isolated institutional policies toward a cohesive ecosystem that ensures research integrity while facilitating equitable access to these transformative tools.

A central pillar of emerging governance frameworks is the concept of transparent reporting through standardized documentation, a necessary evolution to support the “forensic auditor” role now expected of reviewers. The proposal of “AI Usage Cards” represents a significant step toward operationalizing this transparency. These cards serve as a standardized mechanism for researchers to report the specific roles AI systems played in their work, covering dimensions of transparency, integrity, and accountability. By providing a structured format to detail the prompts used, the extent of AI contribution, and the verification steps taken by human authors, AI Usage Cards enable the research community to trace, compare, and question various forms of AI usage [157]. This approach aligns with broader calls for “datasheets for datasets” and model documentation, extending the principle of comprehensive disclosure to the actual conduct of scientific writing and

analysis. Such mechanisms are essential for maintaining the trustworthiness of scientific outputs in an era where AI-generated content can be indistinguishable from human-authored text.

Furthermore, the development of these frameworks is deeply intertwined with broader ethical principles and regulatory landscapes that underpin the legitimacy of AI-mediated science. A meta-analysis of over 200 global AI ethics guidelines reveals at least 17 resonating principles, including transparency, fairness, and accountability, which provide a foundational bedrock for scientific governance [192]. However, translating these high-level principles into actionable scientific norms requires specific adaptations. For instance, the “Five ethical principles for generative AI in scientific research” framework emphasizes the necessity of understanding model limitations regarding truthfulness and bias, respecting privacy and copyright, and ensuring applications provide overall benefit. This framework argues that global consensus, coupled with professional training and reasonable enforcement, is critical to safeguarding research integrity [193].

The tension between centralized regulation and community-driven standards also plays a pivotal role in shaping the future of scientific governance, mirroring the balance required in incentive design for peer review. While some advocate for top-down regulatory approaches similar to the EU AI Act to ensure uniformity, others propose “Dual Governance” models that combine centralized government regulations with crowdsourced safety mechanisms. In the context of scientific research, this could manifest as a synergy between international funding mandates and community-led peer review standards, ensuring that innovation is promoted while ethical deployment is maintained [194]. Additionally, the concept of “management-based regulation” suggests that governing AI in science should focus on motivating human oversight during the training and development phases of research tools, rather than solely regulating the output [195].

International cooperation remains indispensable for effective governance, providing the macro-level stability required to support micro-level reporting frameworks. The UNESCO Recommendation on the Ethics of Artificial Intelligence stands as the first global normative framework, achieved through intense multilateral negotiations that balanced diverse sovereignist and liberal preferences. This document provides a crucial reference point for scientific bodies aiming to align their policies with universal human rights and democratic values [196]. Similarly, proposals for international institutions, such as a Commission on Frontier AI or an Advanced AI Governance Organization, highlight the potential for global coordination to set safety standards and manage risks that transcend national borders [197]. These macro-level governance structures offer the stability and legitimacy required to support the micro-level reporting frameworks adopted by individual researchers and journals.

Ultimately, the trajectory toward standardized governance in scientific writing is moving from abstract ethical principles to operational metrics and audit trails, closing the loop on the accountability challenges raised in the evaluation phase. The development of Responsible AI (RAI) metrics catalogues, which delineate process, resource, and product metrics, offers a pathway to make accountability tangible and auditable [198]. By integrating these metrics into the publication lifecycle, the scientific community can ensure that the adoption of AI enhances rather than erodes the credibility of research. As the field evolves, the successful implementation of these frameworks will depend on the continuous dialogue between policymakers, publishers, and researchers, ensuring that governance mechanisms remain adaptive to the rapid pace of technological change while upholding the core values of scientific inquiry.

4 Human-in-the-Loop Workflows for Data Analysis, Experimentation, and Infrastructure

4.1 Modular Architectures for Self-Driving Laboratories and Science Factories

4.1 Modular Architectures for Self-Driving Laboratories and Science Factories

While computational notebooks provide the interactive interface for data analysis, the physical execution of scientific discovery is undergoing a parallel transformation, shifting from manual, linear experimentation to automated, concurrent workflows facilitated by “science factories.” These large-scale, general-purpose Self-Driving Laboratories (SDLs) represent the convergence of robotic automation, high-performance computing (HPC), and artificial intelligence. To achieve the necessary generality and scale required to tackle complex discovery problems while supporting thousands of scientists, the underlying infrastructure must move away from rigid, monolithic designs toward modular, reconfigurable hardware and software architectures. This subsection examines the critical architectural shifts enabling this transition, focusing on the integration of reconfigurable robotic workcells with distributed computing resources to form scalable ecosystems capable of autonomous experimental execution, thereby laying the physical groundwork for the collaborative digital environments discussed in the following section.

At the core of the science factory concept is the necessity for modularity in both physical instrumentation and computational control. Traditional laboratory automation often relies on proprietary, tightly coupled systems where software is inextricably linked to specific hardware, limiting flexibility and scalability. In contrast, modern architectures propose a prototype modular science factory where reconfigurable modules encapsulating scientific instruments are linked with manipulators to form distinct workcells [199]. These workcells function as semi-autonomous units that can be combined into larger assemblages, allowing for the concurrent execution of diverse applications ranging from materials synthesis to biological assays. The ability to specify workflows that perform sets of actions on these modules enables the migration of applications between different workcells, thereby maximizing resource utilization and experimental throughput [199]. Furthermore, the integration of digital twins within these architectures allows for the simulation and validation of workflows before physical deployment, reducing the risk of equipment damage and optimizing experimental parameters in silico [199].

The challenge of interoperability in such heterogeneous environments is significant, given the diversity of sensors, actuators, and computational nodes involved. To address this, emerging architectures draw inspiration from service-oriented designs used in other robotic domains, such as space exploration. For instance, service-oriented architectures decouple software components from the physical robot body, deploying them as loosely coupled web services that promote integrate-ability and scalability [200]. Applying similar principles to laboratory settings allows for the seamless integration of disparate instruments into a unified control framework. The Robot Operating System 2 (ROS 2) has emerged as a de facto standard in this regard, offering a modular and reliable architecture that supports the sensing, planning, and autonomy required for modern robotic systems [201]. ROS 2’s design facilitates the mapping of computation graphs across configurable System-on-Chip platforms, enabling efficient communication between hardware-accelerated nodes and software processes [202]. This capability is essential for science factories, where real-time data from sensors must be processed by AI models running on heterogeneous accelerators to guide ongoing experiments.

A critical component of these modular architectures is the hardware-software interface, which must support rapid reconfiguration without extensive manual programming. The concept of “plug-and-

play” integration is gaining traction, with frameworks like the Laboratory Automation Plug & Play (LAPP) proposing digital twin-based concepts where device coordinates and capabilities are automatically detected by the robot’s vision system [203]. This approach minimizes the need for manual teaching and configuration, allowing scientists to rapidly repurpose robotic arms for new experimental setups. Similarly, the Hardware Robot Information Model (HRIM) addresses the incompatibility between robot components by establishing a horizontal integration standard that facilitates the interchangeability of modules from different vendors [204]. Such standardization is vital for the economic viability of science factories, as it prevents vendor lock-in and encourages the development of a diverse ecosystem of compatible scientific instruments.

Beyond physical modularity, the computational substrate of science factories must also be adaptive to support the data-intensive workflows that feed into analytical environments. Robotic computing increasingly relies on Field-Programmable Gate Arrays (FPGAs) and other reconfigurable hardware to meet the low-latency and power-efficiency demands of real-time experimental control [205]. Architectures like RobotCore provide a target-agnostic layer for integrating hardware acceleration into ROS 2, enabling the deployment of perception and control algorithms on FPGAs or GPUs depending on the specific requirements of the workflow [206]. This adaptability extends to the network level, where distributed robotic systems operate across the edge-cloud continuum. By leveraging containerization and orchestration tools like Kubernetes alongside ROS 2, science factories can dynamically allocate computational resources, running intensive AI training jobs in the cloud while maintaining low-latency control loops at the edge [207].

The realization of autonomous synthesis and characterization further exemplifies the power of these modular architectures. Systems such as the Scientific Autonomous Reasoning Agent (SARA) integrate robotic synthesis with hierarchical AI methods to efficiently explore vast parameter spaces, demonstrating orders-of-magnitude acceleration in mapping synthesis phase diagrams [208]. Similarly, the ORGANA robotic assistant showcases how flexible architectures can interact with chemists via natural language, planning long-horizon tasks and coordinating multiple robots to execute complex electrochemistry experiments [209]. These systems rely on the underlying modularity to swap instruments, update models, and reconfigure workflows dynamically. Moreover, the use of active learning loops within these architectures ensures that human expertise is integrated where necessary, creating a human-in-the-loop system that optimizes experimental design based on real-time feedback [210].

In conclusion, the shift toward modular architectures for self-driving laboratories and science factories represents a fundamental re-engineering of scientific infrastructure. By decoupling hardware and software through service-oriented designs, standardizing interfaces for plug-and-play integration, and leveraging reconfigurable computing substrates, these architectures enable the scalability and flexibility required for next-generation discovery. As these systems evolve, they generate the high-fidelity data streams and automated workflows that necessitate the advanced interactive analysis environments described in the subsequent section, transforming the scientific workflow from a series of isolated, manual tasks into a continuous, automated, and highly collaborative process between humans and machines.

4.2 Interactive Data Analysis and Computational Notebooks as Collaborative Spaces

4.2 Interactive Data Analysis and Computational Notebooks as Collaborative Spaces

While the preceding section detailed the modular, automated physical infrastructure of “science factories,” the intellectual engine driving these systems resides in the digital realm, specifically within interactive data analysis environments. Computational notebooks have evolved from simple script executors into dynamic, collaborative spaces where human intuition and artificial intelligence converge to interpret the high-fidelity data streams generated by autonomous laboratories. Originally conceptualized as digital successors to laboratory notebooks, environments like Jupyter have become the de facto standard for interactive computing across disciplines, facilitating the seamless interleaving of executable code, rich text, mathematical equations, and visualizations [211]. This literate computing paradigm not only enhances the reproducibility of scientific workflows but also establishes the foundational architecture for human-AI collaboration. In these modern collaborative spaces, the notebook serves not merely as a recording device for experimental results but as an active interface where AI agents assist subject matter experts in navigating complex data landscapes, optimizing resource allocation, and bridging the gap between domain knowledge and the high-performance computing (HPC) infrastructure required to process it.

The evolution of notebooks into collaborative hubs is driven by the integration of AI-powered assistants that augment the traditional coding experience, a necessity given the complexity of data emerging from self-driving laboratories. Recent design studies indicate that while AI code assistants are becoming ubiquitous, their effectiveness in notebook environments depends heavily on context-awareness and the ability to handle the unique stateful nature of interactive sessions [212]. Unlike traditional integrated development environments (IDEs), computational notebooks involve non-linear execution flows where the state of variables persists across cells, creating specific challenges for AI models attempting to synthesize code or explain errors derived from automated experimental runs. To address this, advanced systems are being developed to understand rich multi-modal contexts, including existing notebook cells, execution states, and previous interaction turns, thereby enabling more accurate natural language to code generation [213]. These capabilities allow researchers to delegate routine data wrangling tasks to AI, freeing them to focus on higher-level hypothesis generation and interpretative analysis of the automated workflows.

A critical dimension of this collaboration is the democratization of access to the computational resources necessary to analyze data at the scale produced by science factories. Historically, leveraging HPC clusters required deep expertise in job scheduling, parallel file systems, and command-line interfaces, creating a significant barrier for domain scientists. Computational notebooks, when integrated with middleware solutions like Merlin, act as a translational layer that lowers this barrier, enabling researchers to launch large-scale machine learning ensembles and manage complex workflows without needing to master the underlying infrastructure [214]. By embedding interactive interfaces within the notebook environment, institutions can provide a unified platform where users can prototype on local machines and seamlessly scale to distributed cloud or supercomputing resources. This shift transforms the notebook from an isolated desktop tool into a gateway for scalable science, where AI agents can recommend optimal execution strategies, such as migrating compute-intensive cells to more suitable hardware platforms based on user interactivity patterns and resource availability [215].

However, the transition to AI-augmented collaborative spaces is not without challenges, particu-

larly regarding code quality, reproducibility, and the “messiness” inherent in exploratory analysis of high-throughput experimental data. Empirical studies reveal that a significant portion of publicly shared notebooks suffer from poor coding practices, undocumented cells, and execution order dependencies that compromise reproducibility [216]. To mitigate these issues, new collaborative tools are incorporating static analysis and linting capabilities directly into the notebook interface. Systems like Pynblint and HeaderGen automatically analyze code cells to enforce best practices, classify machine learning operations, and generate structural headers, effectively acting as an AI co-pilot that guides researchers toward more robust and maintainable workflows [217; 218]. Furthermore, to support sensemaking in these complex environments, innovative overlays like Porpoise have been developed to group cells into logical sections and provide semantic annotations, making the exploration of unfamiliar notebooks more akin to reading a structured book [219].

The collaborative potential of notebooks is further expanded through features that support guidance and recommendation, essential for navigating the vast parameter spaces identified in the subsequent phase of automated experimental design. Frameworks such as WHATSNEXT and Lodestar introduce recommendation panels and data-driven analysis suggestions directly into the notebook interface, helping users, especially novices, navigate the vast space of possible analytical next steps [220; 221]. These systems leverage historical data from thousands of notebooks to suggest visualization types, statistical tests, or data transformation pipelines, effectively encoding community best practices into the interactive environment. Additionally, the integration of ephemeral user interfaces, as seen in tools like BISCUIT, allows LLMs to generate temporary graphical controls that help users manipulate parameters and understand generated code without needing deep programming expertise [222].

Looking toward the future of scalable collaboration, the deployment of notebook environments at scale on cyberinfrastructure remains a pivotal area of development to support the continuous loops between analysis and experimentation. Strategies utilizing container orchestration frameworks like Kubernetes and Docker Swarm enable fault-tolerant, multi-user notebook hubs that can serve thousands of researchers simultaneously, fostering a truly collaborative ecosystem [223]. Moreover, emerging concepts like “Notebook-as-a-VRE” (Virtual Research Environment) aim to embed comprehensive research asset management, workflow composition, and community sharing capabilities directly within the notebook, transforming it into a holistic platform for the entire research lifecycle [224]. As these environments continue to mature, the synergy between human creativity and AI-driven automation within computational notebooks provides the analytical foundation required to drive the closed-loop systems discussed in the following section, where insights gained here directly inform automated experimental design and active learning loops.

4.3 Automated Experimental Design and Active Learning Loops

4.3 Automated Experimental Design and Active Learning Loops

Building upon the interactive data analysis environments and collaborative notebooks discussed in the previous section, the scientific workflow extends beyond digital computation into the physical realm through automated experimental design and active learning loops. While computational notebooks serve as dynamic spaces for hypothesis generation and data exploration, the sheer complexity of high-dimensional parameter spaces—ranging from material composition and processing conditions to robotic control policies—often renders exhaustive grid searches and random sampling computationally and economically prohibitive in real-world settings. To address these challenges, researchers are increasingly transitioning from static analysis to data-driven, autonomous discov-

ery loops. In this paradigm, Artificial Intelligence (AI) agents leverage Active Learning (AL) and Bayesian Optimization (BO) frameworks to transform the experimental process into a closed-loop system. These agents autonomously select the most informative subsequent experiments, balancing the dual objectives of exploring unknown regions and exploiting known promising areas to maximize discovery efficiency, effectively bridging the gap between the digital insights gained in notebooks and physical laboratory execution.

At the core of these automated workflows is the principle of adaptive sampling, where the system iteratively refines a surrogate model of the underlying physical phenomenon based on incoming data. Bayesian Optimization has emerged as a premier tool for this task, particularly in scenarios involving expensive, black-box functions where gradient information is unavailable. The efficacy of BO lies in its ability to incorporate prior knowledge and quantify uncertainty, guiding the search toward optimal solutions with minimal sample counts. For instance, in the domain of materials science, BO has been successfully employed to accelerate the mapping of non-equilibrium synthesis phase diagrams, a task previously hindered by the vastness of the processing space [208]. By utilizing nested active learning cycles, autonomous platforms can efficiently reveal the structure of complex phase boundaries, achieving orders-of-magnitude acceleration in establishing synthesis conditions for kinetically stabilized materials.

However, the deployment of these algorithms in real-world scientific settings often requires navigating constraints and multi-objective trade-offs that pure data-driven approaches may overlook. A critical advancement in this area is the development of acquisition functions that explicitly manage the exploration-exploitation trade-off dynamically. Traditional methods often rely on fixed heuristics, which can be suboptimal as the data landscape evolves. Recent frameworks, such as Bayesian Hierarchical Modeling, have introduced mechanisms to dynamically balance this trade-off as more data points are queried, demonstrating significant improvements over static strategies in both synthetic and real-world regression tasks [225]. Furthermore, in constrained environments where violating physical laws could lead to catastrophic equipment failure or invalid results, "failure-averse" active learning methods have been developed. These approaches utilize safe variance reduction and safe region expansion tasks to ensure that the exploration remains within feasible operational bounds, achieving zero-failure rates even without explicit knowledge of failure regions [226].

The synergy between human expertise and algorithmic autonomy further enhances the capabilities of these loops, echoing the human-AI collaboration seen in computational notebooks. While fully autonomous systems offer speed, the incorporation of domain expert knowledge can significantly direct the search toward scientifically relevant regions. Interactive Bayesian Optimization frameworks allow humans to shape the optimization strategy by providing preferences or hypotheses, which are then integrated into the surrogate modeling process. For example, the "HypBO" method leverages expert hypotheses to generate improved seed samples, effectively discounting unpromising regions and augmenting the model with high-value data points to accelerate chemical design tasks [227]. Similarly, frameworks that incorporate expert preferences regarding unmeasured abstract properties have shown superiority over baseline methods by boosting the performance of BO in experimental design optimization [228]. This human-in-the-loop approach ensures that the AI's search aligns with the nuanced goals of the researcher, bridging the gap between raw data efficiency and scientific intuition.

Moreover, the complexity of modern scientific problems often necessitates multi-fidelity and multi-objective approaches. In materials discovery, for instance, researchers must often optimize properties that have hierarchical dependencies, such as ensuring a molecule is expressible before optimizing its binding affinity. The "PropertyDAG" framework addresses this by imposing a partial ordering

on objectives within a multi-objective BO structure, ensuring that optimization proceeds only when parent feasibility conditions are met [229]. Additionally, the integration of multi-fidelity data sources—combining cheap, low-accuracy simulations with expensive, high-accuracy experiments—allows for more efficient resource allocation. Interactive multi-fidelity BO workflows enable domain experts to inject partially known physics laws and vary exploration policies on the fly, leading to improved exploration in highly non-smooth simulation landscapes [230].

The application of these loops extends beyond materials science into robotics and cognitive modeling. In robotics, BO is utilized to optimize control parameters and locomotion controllers in high-dimensional spaces where hardware experiments are time-consuming. Techniques such as learning informed distance metrics via deep kernels or utilizing Riemannian manifold theory to handle non-Euclidean parameter spaces have significantly enhanced sample efficiency in robot learning [231; 232]. Furthermore, in the study of human cognition, Bayesian Optimal Experimental Design (BOED) is employed to design experiments that maximize information gain for discriminating between competing behavioral models, proving more efficient than designs chosen by human experts [233; 234].

Ultimately, the transition to automated experimental design and active learning loops signifies a move toward “self-driving” laboratories. These systems do not merely execute pre-defined protocols but actively reason about the state of knowledge, propose novel experiments, and adapt their strategies in real-time. Whether through surprise-reacting policies that prioritize unexpected observations to characterize response surfaces [235] or through platforms like AutoOED that provide modular, GUI-driven interfaces for multi-objective optimization [236], the convergence of AI and experimental infrastructure is redefining the pace and scope of scientific discovery. By allowing humans to define high-level goals and constraints while delegating the iterative navigation of parameter spaces to intelligent agents, these workflows maximize the utility of every experimental run. This evolution sets the stage for the next critical layer of the research ecosystem: the underlying high-performance computing and AI accelerator integration required to power these computationally intensive, large-scale autonomous loops, as discussed in the following section.

4.4 Optimizing High-Performance Computing and AI Accelerator Integration

4.4 Optimizing High-Performance Computing and AI Accelerator Integration

Building upon the “self-driving” laboratories and autonomous experimental loops described in the previous section, the realization of such high-throughput, real-time discovery systems relies fundamentally on a robust underlying computational infrastructure. The convergence of Artificial Intelligence (AI) and High-Performance Computing (HPC) represents a paradigmatic shift from distinct operational silos toward a unified ecosystem where data-driven discovery and simulation-based modeling coalesce. As scientific workflows evolve to incorporate machine learning tasks directly into traditional simulation pipelines, the demand for computational resources has surged, necessitating sophisticated strategies for optimizing resource allocation and managing heterogeneous hardware. The integration of AI accelerators—ranging from Graphics Processing Units (GPUs) to novel architectures like Intelligence Processing Units (IPUs)—into HPC environments introduces complex challenges in scheduling, energy efficiency, and system co-design that require both human-in-the-loop oversight and intelligent automation to sustain the pace of autonomous science.

A primary challenge in this converged landscape is the efficient management of heterogeneous hardware accelerators required to power the iterative cycles of active learning and Bayesian optimization. Modern HPC systems are no longer homogeneous; they often comprise a mix of CPUs, GPUs, and

specialized accelerators, each with distinct memory hierarchies and communication characteristics. Effective utilization requires moving beyond static allocation policies to dynamic, intelligent scheduling mechanisms capable of adapting to the fluctuating demands of autonomous loops. For instance, research has demonstrated that intelligent co-scheduling algorithms, which leverage reinforcement learning to model application resource usage against cluster monitoring data, can significantly improve cluster utilization by up to 51% while reducing job makespan [237]. Such approaches allow system operators to colocate workloads strategically, mitigating contention on shared resources like cache and memory bandwidth while ensuring that critical scientific applications receive the necessary computational power [238]. Furthermore, the diversity of accelerator architectures necessitates careful consideration of interconnect technologies. Empirical evaluations of modern GPU interconnects, including NVLink and NVSwitch, have revealed complex Non-Uniform Memory Access (NUMA) effects that can severely impact performance if not accounted for during task allocation [239]. Understanding these topological nuances is essential for humans designing workflows, as the choice of GPU combination and data placement can dictate the efficiency of large-scale distributed training or inference tasks central to self-driving labs.

Beyond scheduling, optimizing the integration of AI accelerators involves addressing the specific architectural constraints of emerging hardware types that underpin these advanced workflows. The rise of IPU and other data-flow architectures offers promising alternatives to traditional von Neumann designs, particularly for AI/ML workloads characterized by massive parallelism and specific memory access patterns inherent in surrogate modeling and simulation. However, these architectures often come with unique constraints, such as limited on-chip memory capacity, which necessitates innovative software strategies like butterfly factorizations to reduce model size without sacrificing accuracy [240]. Microbenchmarking studies have been instrumental in dissecting these architectures, providing mental models that help researchers predict performance based on computation and communication steps [241]. Comparative analyses further highlight the trade-offs between different accelerator platforms, illuminating how data-flow architectures may offer superior energy efficiency for specific operators compared to conventional GPUs [242]. These insights are critical for human researchers who must select the appropriate hardware backbone for their scientific inquiries, balancing raw performance against energy consumption and programming complexity to ensure the viability of long-running autonomous experiments.

Energy efficiency has emerged as a critical metric in the co-design of AI-HPC systems, particularly as the scale of models and simulations approaches exascale levels and the frequency of experimental iterations increases. The power consumption of large-scale AI training and HPC simulations poses sustainability challenges, prompting the development of dynamic power management strategies essential for continuous autonomous operation. Control theory approaches have been successfully applied to create autonomous power regulation methods that adjust power caps at runtime based on application progress, thereby sustaining performance while reducing energy consumption [243]. Additionally, techniques such as GPU power capping have been explored at the supercomputing scale to reduce thermal output and power draw with minimal impact on job performance, although careful analysis is required to prevent users from compensating for performance degradation by requesting additional resources [244]. The development of accurate measurement toolkits enables researchers to conduct energy-aware computational experiments, providing the granular data needed to optimize code for both CPUs and GPUs [245].

To facilitate seamless integration and system co-design, the community is increasingly relying on standardized execution traces and simulation models to predict the behavior of these complex, heterogeneous systems. Discrete-event simulation models have proven effective in estimating system

performance scalability and identifying bottlenecks in heterogeneous AI systems before physical deployment, offering a cost-effective method for optimization [246]. Moreover, the creation of comprehensive datasets, such as the MIT Supercloud Dataset, provides the foundational logs necessary for developing AI-driven approaches to workload classification and resource optimization [247]. These datasets enable the training of models that can identify inefficiencies and predict failures, fostering a feedback loop where system operations inform algorithmic improvements. Visualization systems like HPCClusterScape further enhance this process by providing transparency into system-level and application-level metrics, allowing operators to diagnose synchronization bottlenecks and workload imbalances in real-time [248].

Ultimately, the optimization of HPC and AI accelerator integration is not merely a technical exercise in hardware configuration but a socio-technical endeavor requiring close collaboration between domain scientists, system architects, and algorithm developers. The convergence of these fields demands new programming models that abstract away underlying heterogeneity while exposing sufficient control for performance tuning, as seen in frameworks enabling runtime support for performance portability across distributed platforms [249]. As the field advances, the focus must remain on creating adaptive, energy-efficient, and transparent systems that empower researchers to harness the full potential of converged computing infrastructures for groundbreaking scientific discovery. However, as these optimized systems generate vast amounts of data and execute complex, dynamic workflows across distributed environments, they introduce profound challenges for reproducibility and provenance tracking, setting the stage for the critical discussion on workflow management and data observability in the following section.

4.5 Reproducibility, Provenance, and Workflow Management Across Distributed Infrastructures

4.5 Reproducibility, Provenance, and Workflow Management Across Distributed Infrastructures

While the optimization of HPC and AI accelerator integration detailed in the previous section provides the computational horsepower necessary for modern discovery, the resulting complexity of these heterogeneous environments introduces profound challenges for reproducibility and workflow management. As scientific workflows evolve from monolithic executions to dynamic processes spanning the “Computing Continuum”—from latency-sensitive IoT edge devices to massive HPC clouds—maintaining data observability and ensuring experimental repeatability requires more than simple code archiving. It demands sophisticated mechanisms to capture the granular interplay between algorithms, data, and the ephemeral infrastructure configurations that dictate performance. Without such rigor, the high-throughput capabilities of optimized systems risk producing results that are impossible to validate or replicate.

A primary obstacle in this distributed landscape is the lack of comprehensive provenance information regarding the execution infrastructure itself. Traditional provenance systems, often designed for static grid computing, frequently fail to capture the dynamic nature of cloud resources where virtual machines are provisioned on-demand. Variations in hardware architecture, network latency, or software environments can cause re-executed workflows to yield divergent results, effectively breaking the chain of reproducibility [250]. To address this, recent frameworks have introduced “cloud-aware” provenance collection mechanisms that map workflow execution traces to specific infrastructure configurations. For instance, the ReCAP framework demonstrates the necessity of capturing virtual machine configurations alongside workflow provenance to facilitate the re-provisioning of similar

resources, thereby enabling true repeatability of experiments across different cloud scenarios [251].

The challenge of provenance capture is further compounded when workflows traverse the Edge-to-Cloud continuum, where the resource constraints of edge devices make traditional, heavy-weight logging prohibitive. The overhead associated with tracking every data transformation can degrade the performance of latency-sensitive applications running on IoT devices. To mitigate this, lightweight solutions such as ProvLight have been developed to enable efficient provenance capture specifically tailored for resource-constrained edge environments. By leveraging simplified data models, compression techniques, and lightweight transmission protocols, ProvLight significantly reduces the CPU, memory, and energy overhead compared to state-of-the-art systems, allowing for continuous tracking of performance indicators without disrupting the scientific workflow [252]. This efficiency is paramount for maintaining data observability in large-scale deployments where thousands of edge devices contribute to a unified scientific analysis.

Beyond mere capture, managing and integrating provenance data across multi-workflow environments requires unified frameworks capable of bridging the fragmentation of existing Workflow Management Systems (WMS). The current landscape, characterized by hundreds of incompatible systems, creates significant barriers to sharing reproducible workflows [253]. To bridge this gap, approaches like Workflow Run RO-Crate extend existing standards to bundle workflow plans, execution traces, and associated products (inputs, outputs, code) into interoperable packages. This model supports the capture of provenance at varying levels of granularity and has been implemented across multiple workflow systems, facilitating interoperable comparisons and the archiving of complex computational experiments [254]. Similarly, initiatives like CUF-Links advocate for continuous and ubiquitous FAIRness linkages, proposing simple mechanisms integrated into daily work practices to create an unbroken chain of provenance from inputs to outputs, thus strengthening the foundation for reproducible research [255].

Complementing these software standards, the integration of containerization technologies has emerged as a vital strategy for encapsulating both computational environments and data artifacts to ensure portability and traceability. While containers traditionally focus on application isolation, newer concepts like “Data Pallets” extend this logic to storage, automatically generating writeable containers based on I/O activities to link data products back to the specific application and input decks that created them [256]. This system-level support complements user-level workflow tools by providing a robust mechanism for tracking the evolution of data in high-consequence missions where trust in simulation output is paramount. Additionally, frameworks like ProvDeploy assist users in configuring containers for HPC workflows with integrated provenance capture, simplifying the deployment of complex software stacks while ensuring that lineage information is preserved throughout the execution lifecycle [257].

In collaborative and adaptive settings, trust is contingent not only on tracking data flow but also on auditing changes to the pipeline architecture itself. Cross-organizational interactions often suffer from a lack of trust, which can impede progress. Addressing this, specialized services like the Provenance Holder have been proposed to capture the provenance of changes in processing pipelines, allowing researchers to audit the evolution of workflow structures and data usage in a trusted manner [258]. This capability is essential for debugging and validating complex, long-running experiments where user steering actions—such as fine-tuning parameters during execution—can significantly alter outcomes. Systems that track these online steering actions with negligible overhead enable researchers to correlate specific tuning decisions with performance improvements or accuracy changes, providing deep insights into the dynamic behavior of scientific workflows [259].

Finally, the sheer volume of data generated in exascale computing and big data analytics necessitates efficient querying and analysis of massive provenance graphs. Techniques that leverage generalized trie structures or weakly connected component frameworks allow for the real-time processing of provenance queries on graphs containing hundreds of millions of nodes, enabling researchers to quickly identify lineage relationships and divergence points in experimental results [260; 261]. As the scientific community moves towards more automated and autonomous discovery systems, the development of holistic workflow environments that integrate data observability, adapter system design, and lightweight runtime integration will be crucial. Approaches like MIDA demonstrate how multi-workflow integrated data analysis can be achieved with near-zero overhead by intercepting dataflows in the background, thereby unifying domain, provenance, and telemetry data for comprehensive user steering and reproducibility [262]. These advancements collectively form the backbone of a reproducible scientific ecosystem, providing the necessary transparency and trust upon which the Human-in-the-Loop AutoML systems discussed in the following section rely to effectively guide algorithmic search processes.

4.6 Human-in-the-Loop AutoML and Adaptive Pipeline Tuning

4.6 Human-in-the-Loop AutoML and Adaptive Pipeline Tuning

While the robust provenance and workflow management infrastructures detailed in the previous section provide the necessary transparency and trust for reproducible science, they form only the foundation for the next critical layer of collaboration: the active guidance of algorithmic search processes. The evolution of Automated Machine Learning (AutoML) has traditionally been driven by the objective of minimizing human effort, striving for “black-box” solutions that generate optimal models with a single command. However, as scientific research workflows increasingly encounter open-environment scenarios characterized by dynamic data streams, shifting constraints, and complex domain-specific requirements, the paradigm of full automation has revealed significant limitations. In response, the field is witnessing a pivotal shift toward Human-in-the-Loop (HITL) AutoML and adaptive pipeline tuning. This approach reconceptualizes AutoML not as a replacement for the data scientist, but as a collaborative partner that leverages human intuition and domain expertise to guide algorithmic search processes, particularly in contexts where pure data-driven approaches are computationally intractable or insufficiently flexible.

The necessity for human intervention arises from the inherent challenges of open-environment AI, where problems are rarely static. Researchers often face the need to continuously collect and update data, modify evaluation metrics, and adapt to changing deployment environments, such as industrial supply chains or the industrial metaverse. Purely automated systems often struggle in these fluid contexts due to the immense computational resources required to retrain models from scratch or the inability to incorporate nuanced, non-quantifiable constraints. Systems like OmniForce have emerged to address these gaps by implementing a human-centered AutoML (HAML) architecture that supports both human-assisted machine learning and machine-assisted human techniques. By enabling users to manage ML versions, collaborate on pipeline-driven development, and utilize flexible search strategies, such systems allow researchers to maintain agency over the modeling process while benefiting from automated optimization capabilities [263].

A critical component of effective HITL AutoML is the integration of visual analytics to make the opaque processes of hyperparameter optimization and model selection transparent and controllable. Traditional AutoML methods often obscure the reasoning behind specific configuration choices, leading to a deficit in user trust and understanding. To mitigate this, interactive visualization tools

have been developed to support user-driven hyperparameter tuning. For instance, HyperTendrill introduces a web-based visual analytics system that allows users to iteratively refine search spaces and diagnose the behavior of various search algorithms. By providing insights into variable importance and interaction effects, such tools empower users to steer the optimization process based on their own domain insights rather than relying solely on default algorithmic behaviors [264]. Similarly, ATMSeer offers a multi-granularity visualization interface that enables users to monitor the AutoML process in real-time, analyze searched models, and dynamically refine the search space, thereby addressing the distrust users often feel toward automatic results and preventing the wasteful expansion of search budgets [265].

Beyond visualization, the mechanism of adaptive pipeline tuning relies heavily on the ability of systems to incorporate iterative feedback and interpretability into the workflow. The complexity of modern ML pipelines, which involve interconnected steps of data preprocessing, feature engineering, and model tuning, necessitates tools that can explain not just the final model, but the entire construction process. XAutoML addresses this by combining interactive visualizations with Explainable AI (XAI) techniques to render the complete AutoML procedure transparent. This transparency is crucial for establishing trust, as studies have shown that data scientists require detailed information about model performance and the optimization trajectory before they are willing to deploy automatically generated models [266] [267]. Furthermore, the concept of “interpretable feedback” extends to data curation, where systems can suggest specific new data points for users to label, effectively closing the loop between model performance deficits and human data annotation efforts [268].

The design of these systems must also account for the diverse needs of different user personas and the varying degrees of automation desired across different stages of the data science lifecycle. Research indicates that there is no universal desire for complete end-to-end automation; rather, data scientists prefer varying levels of control depending on the task at hand, such as problem formulation versus hyperparameter tuning [269]. Consequently, successful HITL frameworks often adopt a modular approach, allowing users to intervene at specific junctures. For example, PipelineProfiler enables the exploration and comparison of the solution space of ML pipelines, helping experts to debug systems and select the most appropriate AutoML strategy for a given problem class [270]. Additionally, the emergence of Large Language Models (LLMs) has introduced new paradigms for interaction, where agents like AgentHPO or AutoML-GPT can act as conversational interfaces, allowing users to specify constraints and goals in natural language while the system autonomously conducts experiments and optimizes hyperparameters [271] [272].

Despite these advancements, challenges remain in balancing the trade-off between automation speed and human oversight. Empirical studies in enterprise settings have highlighted tensions where data visualization tools sometimes fail to adequately balance the need for rapid prototyping with the necessity for rigorous human validation [273]. Moreover, the adaptability of AutoML systems to heavily constrained applications remains a hurdle, prompting the development of meta-learning approaches like CAML that can automatically adapt their own search strategies and validation methods based on user-defined constraints [274]. As the field progresses, the focus is shifting from merely automating model generation to creating adaptive, collaborative ecosystems where humans and machines co-evolve solutions. This involves not only technical innovations in search algorithms and visualization but also a deeper understanding of the cognitive workflows of researchers, ensuring that AutoML systems enhance rather than erode human agency in scientific discovery [275]. Ultimately, these adaptive tuning mechanisms and the symbiotic relationships they foster provide the methodological bridge to the domain-specific applications discussed in the following section, where

human expertise is tailored to the unique epistemic constraints of materials science, biomedicine, and software engineering.

4.7 Case Studies in Domain-Specific Workflow Integration

4.7 Domain-Specific Human-in-the-Loop Optimization Workflows

Building upon the adaptive pipeline tuning and visualization mechanisms discussed in the previous section, the integration of Human-in-the-Loop (HITL) workflows into domain-specific scientific processes represents a critical evolution for addressing complex, multi-objective optimization problems. In fields where purely automated approaches often falter due to data scarcity, prohibitive evaluation costs, or the necessity for nuanced qualitative judgment, successful collaboration relies on architectures that seamlessly blend algorithmic efficiency with human expertise, intuition, and ethical oversight. By synthesizing empirical evidence from materials science, biomedicine, and software engineering, it becomes evident that the “human-in-the-loop” paradigm must be tailored to the unique epistemic constraints of each discipline.

In the realm of materials science and chemical discovery, the primary challenge lies in navigating vast, combinatorial design spaces where experimental evaluation is both expensive and time-consuming. Here, HITL workflows accelerate the discovery of novel materials by incorporating expert hypotheses and physical constraints directly into the optimization loop. For instance, in the design of Metal-Organic Frameworks (MOFs), researchers have successfully integrated Latent Variable Gaussian Processes with Multi-Objective Batch-Bayesian Optimization to autonomously identify top-performing candidates without requiring specific physical descriptors, relying instead on qualitative representations of building blocks [276]. This approach demonstrates how AI can manage the combinatorial explosion of possibilities while humans define the qualitative boundaries of the search. Furthermore, the incorporation of expert knowledge through hypothesis-driven seeding has proven to significantly accelerate Bayesian optimization searches. The HypBO framework exemplifies this by using expert human hypotheses to generate improved seed samples, effectively directing the search toward promising regions of chemical space that purely data-driven methods might miss [227]. Additionally, when dealing with systems where expensive measurements are a bottleneck, frameworks like Gemini utilize inexpensive proxy measurements corrected by data-driven models to guide autonomous experimentation, thereby reducing the resource burden while maintaining discovery fidelity [277]. These case studies illustrate a shift from static optimization to dynamic, interactive discovery, where human intuition regarding physical laws and material properties guides the algorithmic exploration of non-differentiable parameter spaces [230].

In biomedicine and clinical decision support, the stakes of optimization are elevated by the stringent requirements for interpretability, robustness against model risk, and the ability to operate in low-data environments. The complexity of biomanufacturing processes, characterized by interdependent factors and limited experimental data, necessitates models that capture causal mechanisms rather than mere correlations. A novel model-based reinforcement learning framework utilizing dynamic Bayesian networks has been developed to achieve human-level control in such settings, enabling the design of process control policies that are both interpretable and robust [278]. Moreover, the life-cycle of clinical machine learning projects involves intricate interactions among data preprocessing, feature selection, and uncertainty estimation. The Clairvoyance pipeline addresses these challenges by providing a unified, autoML-friendly framework that facilitates transparent and reproducible experimentation for clinical time-series data, effectively bridging the gap between potential ML capabilities and their realized utilization in healthcare settings [279]. In these domains, the human

role extends beyond simple preference selection to include the validation of causal assumptions and the interpretation of model behaviors in the context of patient safety and biological plausibility. The integration of visual analytics workflows further empowers domain experts to diagnose classifier decisions and generate hypotheses for model improvement, ensuring that AI systems remain aligned with clinical realities [280].

Software engineering and computational system design present a unique landscape for HITL integration, characterized by the need to balance quantitative performance metrics with qualitative architectural attributes that are difficult to formalize numerically. In the early phases of software architecture design, interactive evolutionary computation has been employed to integrate human judgment into the search process, allowing engineers to guide multi-objective evolutionary algorithms based on both measurable metrics and experiential know-how [281]. This is particularly relevant when optimizing configurable computer systems, where the configuration space is enormous and environmental uncertainties are frequent. The CAMEO method addresses this by identifying invariant causal predictors to optimize system performance across changing environments, effectively reducing the search space through transfer learning informed by causal reasoning [282]. Furthermore, the optimization of software configurations often involves conflicting performance aspirations. Empirical studies suggest that the realism of these aspirations is a key determinant in whether they should guide the tuning process directly or be reserved for final decision-making, highlighting the nuanced role of human requirements in shaping optimization strategies [283]. To make these complex optimization processes more accessible to non-experts, frameworks like JCLECMO provide Java suites that enable engineers to apply multi-objective algorithms with minimal coding effort, democratizing access to advanced optimization tools [284]. Additionally, in scenarios involving NP-hard problems such as the Traveling Salesman Problem, interactive machine learning approaches that open the “black box” to a “glass-box” allow humans to directly interact with learning algorithms, leveraging human ability to extract insights from small data sets to complement algorithmic strengths in high-dimensional pattern recognition [285].

Across these diverse domains, a common thread emerges: the most effective workflows are those that treat human expertise not as a fallback for algorithmic failure, but as a fundamental component of the optimization engine. Whether through the injection of physical constraints in chemistry, the enforcement of causal interpretability in biomedicine, or the incorporation of qualitative architectural preferences in software engineering, human-in-the-loop systems are redefining the boundaries of what is achievable in scientific discovery. These case studies underscore the necessity of designing flexible, adaptive frameworks that can accommodate the specific epistemic and operational constraints of each field while fostering a symbiotic relationship between human cognition and machine intelligence, setting the stage for broader discussions on the future of collaborative scientific ecosystems.

5 Interaction Design, Cognitive Alignment, and Trust Calibration

5.1 Interaction Patterns and Mixed-Initiative Dialogue Frameworks

5.1 Interaction Patterns and Mixed-Initiative Dialogue Frameworks

The evolution of human-AI collaboration in scientific workflows necessitates a fundamental shift from static, one-off interactions to dynamic, conversational, and argumentative exchange models. Historically, AI-assisted decision-making has been dominated by simplistic collaboration paradigms where the AI provides a recommendation and the human passively accepts or rejects it. A system-

atic review of existing literature reveals that current interactions often lack truly interactive functionality, failing to support the nuanced communication required for complex scientific inquiry [9]. To address this limitation and lay the groundwork for the socially situated transparency discussed in the following section, researchers are increasingly advocating for mixed-initiative dialogue frameworks. These frameworks empower both humans and AI agents to initiate actions, ask questions, and resolve conflicts through structured argumentation, moving beyond mere output generation to foster genuine partnership.

A critical component of advancing these interaction patterns is the transition from static explanations to interactive explanatory dialogues. Traditional Explainable AI (XAI) methods typically offer fixed outputs that cannot adapt to the diverse backgrounds and understanding levels of different users. In contrast, conversational approaches enable users to ask follow-up questions, significantly improving comprehension, acceptance, and trust in the system [286]. This shift supports a more robust cognitive alignment, as users can probe the rationale behind AI suggestions until they achieve a satisfactory mental model of the system’s reasoning. For instance, in scientific writing tasks, conversational XAI agents that facilitate multifaceted and context-aware drill-downs have been shown to outperform graphical user interface baselines in enhancing user understanding and writing quality [147]. Such dynamic exchanges are precursors to the deeper social transparency required when AI systems operate within complex organizational fabrics.

To formalize these dynamic interactions, scholars have proposed various taxonomies and design spaces based on interaction primitives. By defining a vocabulary for information exchanges—such as requesting specific model data or providing feedback—researchers can construct abstract specifications for purposeful human-AI interactions [287]. These primitives serve as the building blocks for higher-level interaction patterns that can be tailored to specific scientific domains. Furthermore, grounding these interactions in established linguistic and social theories is essential. For example, the application of conversational maxims, including quantity, quality, relevance, manner, benevolence, and transparency, provides a prescriptive framework for assessing and designing effective human-AI conversations [288]. Adherence to these maxims ensures that AI agents not only provide accurate information but also recognize their knowledge boundaries and operate with benevolent intent, which is crucial for maintaining trust in high-stakes research environments and sets the stage for the socio-organizational considerations explored later.

Beyond simple question-and-answer formats, advanced interaction frameworks leverage computational argumentation to resolve conflicts between human and machine agents. In scenarios where human intuition clashes with algorithmic predictions, argumentative exchanges (AXs) offer a mechanism for dynamic information sharing to reconcile differing viewpoints. This approach treats explanation as a conflict resolution process, where agents share quantitative bipolar argumentation frameworks to reach a consensus [289]. Similarly, frameworks like DR-HAI utilize argumentation-based dialectical reconciliation to address knowledge discrepancies, enabling interactive reconciliation that extends beyond simple model explanations to foster deeper understanding [290]. These argumentative models are particularly valuable in scientific contexts where conflicting hypotheses must be deliberated upon rather than simply accepted or discarded, mirroring the collective decision-making processes central to social transparency.

The concept of “Human-AI Deliberation” further expands the scope of interaction by engaging humans and AI in dimension-level opinion elicitation and deliberative discussion. Instead of passively reviewing an AI’s suggestion, this framework encourages active reflection and discussion on conflicting opinions, leveraging Large Language Models (LLMs) as bridges between human researchers and domain-specific models [131]. Such deliberative capabilities are essential for mitigating au-

tomation bias, ensuring that humans contribute their unique contextual knowledge rather than deferring blindly to algorithmic outputs. Moreover, the integration of Theory of Mind (ToM) modeling into dialogue systems allows agents to track belief dynamics and negotiate common ground, mimicking the natural flow of human conversation and enhancing the efficiency of collaborative tasks [291]. These mechanisms prepare the ground for understanding how AI agents must navigate social attributions and community norms.

Despite these advancements, significant challenges remain in standardizing evaluation metrics for these complex interaction patterns. Traditional non-interactive benchmarks often fail to capture the quality of the interactive process or the subjective experience of the user. New frameworks, such as the Human-AI Language-based Interaction Evaluation (HALIE), emphasize the need to evaluate the interactive process itself, capturing first-person subjective experiences and notions of preference beyond mere output quality [292]. Additionally, the development of unified dialogue datasets and foundation models is critical for overcoming the fragmentation in current conversational AI research, allowing for the creation of more robust and adaptable dialogue agents [293]. As the field moves forward, the synthesis of these interaction patterns into cohesive mixed-initiative frameworks will be paramount. This evolution from basic dialogue to sophisticated deliberation directly informs the subsequent discussion on Social Transparency, where the focus shifts from the mechanics of interaction to the broader socio-organizational contexts in which these collaborations must be trusted and contested.

5.2 Social Transparency and Context-Aware Explainability

5.2 Social Transparency and Context-Aware Explainability

Building upon the mixed-initiative dialogue frameworks and argumentative exchange models discussed in Section 5.1, the focus of human-AI collaboration must expand from the mechanics of interaction to the broader socio-organizational ecosystems in which these exchanges occur. While Section 5.1 established how dynamic conversations and conflict resolution protocols enable fluid communication, true trust and effective partnership require that these interactions be grounded in “Social Transparency” (ST). Historically, Explainable AI (XAI) has been dominated by an algorithm-centered paradigm, prioritizing the exposure of internal mechanics, feature attributions, or decision boundaries. Although such technical transparency is necessary for debugging, it often fails to address the complex socio-organizational realities of scientific workflows. Consequently, the field is witnessing a pivotal shift toward Social Transparency, a sociotechnically informed perspective that incorporates broader contextual factors into AI-mediated decision-making. This approach posits that for AI systems to be truly comprehensible and trustworthy, explanations must move beyond the “black box” of the algorithm to illuminate the social structures, community norms, and organizational processes that surround the system [294].

Social Transparency fundamentally redefines the scope of explainability by recognizing that AI systems are not isolated artifacts but are deeply embedded within human social fabrics. Traditional XAI methods often treat the user as an isolated cognitive agent interacting with a static model, ignoring the collective nature of decision-making in organizations. However, research indicates that explanations in human-human interactions are inherently socially situated, relying on shared values, norms, and practices. By extending this logic to human-AI interaction, Social Transparency aims to facilitate organizational collective actions and cultivate a more holistic form of explainability that supports not just individual understanding, but also group coordination and trust calibration [294]. This is particularly critical in community contexts, where AI-led decisions, such as automated

moderation or resource allocation, must be understood and contested within the framework of shared community values rather than purely technical metrics [295]. This social grounding serves as the essential bridge between the interactive dialogue patterns of the previous section and the cognitive alignment challenges explored in the subsequent section.

A core component of Social Transparency is the integration of context-aware explanations that adapt to the specific situational and cultural environments of the users. The “one-size-fits-all” approach to explainability often falters when deployed across diverse global contexts, particularly in the Global South, where socio-technical systems operate under different constraints and cultural norms. Grounding explainability in these specific social contexts is essential to ensure that AI systems do not perpetuate epistemic injustices or impose foreign conceptual frameworks on local populations [296]. Furthermore, the concept of “situated explanation” suggests that the meaning of an AI decision is co-constructed by the community through their shared practices. For instance, in platforms like Wikipedia, the governance and gatekeeping practices are articulated through automated infrastructures that veterans and administrators have culturally adapted over years. Effective transparency in such environments requires exposing not just the algorithmic logic, but also the learned expertise and organizational culture that governs the use of these tools [297].

The implementation of Social Transparency also necessitates a move from static justifications to dynamic, process-centric explanations that enable contestability. While traditional explanations often focus on the outcome (the “what”), Social Transparency emphasizes the rationale behind the development and deployment of the system (the “why” and “how” of the process). This shift is vital for empowering decision subjects to exercise their right to contest algorithmic decisions, thereby transforming explanations from mere transparency tools into mechanisms for accountability and recourse [298]. Moreover, in high-stakes public sector domains, such as child welfare, algorithmic systems must support existing bureaucratic processes and augment human discretion rather than replace it. Trust in these systems is built not on the opacity of the model’s accuracy, but on the alignment of the system’s outcomes with social ecological frameworks and the collective buy-in of practitioners at all levels [299].

However, the pursuit of Social Transparency is not without its challenges, particularly regarding the risk of social misattribution. Large Language Models (LLMs), with their capability to simulate personas and social roles, can create mismatches between designers’ intentions and users’ perceptions of social attributes. This can lead to unwarranted trust, emotional manipulation, or the attribution of agency where none exists. To mitigate these risks, the Social Transparency framework must be extended to explicitly clarify the specific social attributions assigned to AI systems by both designers and users, ensuring that the simulated social presence of an AI does not obscure its non-human nature [300]. Additionally, the design of transparent systems must account for the trade-offs between explanation complexity and cognitive load. Users may benefit from progressive disclosure strategies, where initial feedback is simplified to help build working heuristics before revealing more complex system errors or details, thereby preventing the distraction that can undermine effective decision-making [301].

Ultimately, achieving Social Transparency requires a multidisciplinary effort that bridges the gap between technical explainability and social science. It demands that researchers and practitioners view fairness and transparency not merely as mathematical properties of an algorithm, but as inherently social concepts that emerge from the mutual influences between technical and social structures [302]. By incorporating socio-organizational contexts, respecting community norms, and providing situated explanations, Social Transparency offers a pathway to AI systems that are not only interpretable but also socially intelligent, accountable, and aligned with the collective

well-being of the communities they serve. This paradigm shift establishes the necessary socio-organizational foundation for the next phase of collaboration: ensuring that these socially situated explanations successfully translate into accurate mental models and cognitive alignment between human researchers and AI agents, as detailed in Section 5.3.

5.3 Cognitive Alignment, Mental Model Calibration, and Co-Learning

5.3 Cognitive Alignment, Mental Model Calibration, and Co-Learning

While Social Transparency (Section 5.2) establishes the necessity of embedding AI explanations within broader socio-organizational contexts, the practical realization of this transparency depends on the degree of cognitive alignment between the human operator and the artificial system. Cognitive alignment refers to the state where the human’s mental model of the AI’s capabilities, limitations, and decision-making processes accurately reflects the system’s actual functioning. Without this alignment, even the most socially situated explanations may fail to prevent users from either under-relying on beneficial advice or over-relying on erroneous outputs. Consequently, achieving effective collaboration requires a shift from static interface design to dynamic, iterative processes of joint mind modeling and co-learning, where both human and machine adapt to one another over time to bridge the gap between technical explainability and human comprehension.

A foundational element of cognitive alignment is the development of accurate mental models by human users. Traditional approaches have often focused on making the AI explain its internal logic; however, research suggests that explanations alone are insufficient if they do not resonate with the user’s cognitive style or fail to update the user’s internal representation of the system. For instance, studies indicate that users’ understanding of AI predictions is significantly influenced by their individual cognitive styles, such as whether they process information rationally or intuitively [303]. Consequently, inclusive design approaches that tailor explanatory interfaces to diverse problem-solving styles have been shown to significantly improve the accuracy of users’ mental models [304]. Furthermore, the mere provision of internal state data, such as confidence scores or attention maps, does not automatically translate into better predictive ability regarding AI failures; instead, humans require structured training to develop a “Theory of AI’s Mind” (ToAIM) that allows them to anticipate system quirks and boundaries [43]. This internal calibration is a prerequisite for the effective management of trust discussed in the subsequent section.

Beyond static mental modeling, effective collaboration demands a bidirectional flow of adaptation known as co-learning. In this paradigm, the human and the AI are viewed as two dynamic learning entities that must grow together to optimize joint performance. This concept moves beyond the traditional view of the human as a passive recipient of machine intelligence, positing instead that both agents must develop specific abilities to interact and collaborate effectively [38]. A critical component of this co-learning framework is “mutual understanding,” where the AI not only explains its decisions but also models the human’s intent, beliefs, and cognitive state. This reciprocal modeling is essential for navigating complex tasks where information asymmetry exists. Research emphasizes that for true symbiosis, autonomous agents must embed practical models of human behavior to anticipate user actions, thereby moving from simple pattern replication to genuine operational alignment [305].

The mechanism of joint mind modeling is further refined through the lens of higher-order cognitive modeling and Theory of Mind (ToM). Just as humans attribute mental states to other humans to facilitate social coordination, advanced AI systems are increasingly being endowed with computational ToM to infer user goals and beliefs. This capability allows AI to operate not just as a tool,

but as a teammate that can predict human needs and adjust its behavior accordingly [306]. Mathematical frameworks, such as extended fuzzy cognitive maps, have been proposed to personalize these cognitive models, enabling machines to estimate and predict the state-of-mind of distinct individuals with high precision [307]. Similarly, neural architectures like ToMnet utilize meta-learning to build models of agents from behavioral observations, allowing machines to pass classic ToM tests and recognize false beliefs, a crucial step for robust multi-agent interaction [308]. In collaborative physical or virtual environments, this mutual attribution of beliefs and intentions—termed the “intentions-beliefs synergy”—enables faster and more accurate predictions of partner actions, fostering seamless joint action [309].

However, achieving alignment is not without its challenges, particularly regarding the calibration of trust and the mitigation of automation bias. The process of calibration involves continuously adjusting the human’s reliance on the AI based on observed performance and context. Recent findings suggest that optimizing the human-AI system as a joint entity, rather than optimizing the AI model in isolation, yields superior results. For example, presenting AI models with calibrated uncertainty, or even strategically uncalibrated confidence levels, can guide humans toward more appropriate reliance patterns, thereby enhancing overall team accuracy [310]. Moreover, the role of human learning cannot be overstated; appropriate reliance is mediated not just by the user’s mental model of the AI, but by the user’s own learning trajectory regarding the task and the tool [37]. Training interventions that encourage “critical use”—where humans learn to situate AI predictions against their own unique contextual knowledge—have proven effective in helping novices develop reliance patterns similar to experienced practitioners [311]. These interventions serve as a critical bridge to the strategies for mitigating automation bias explored in the following section.

Furthermore, the dynamics of co-learning extend to the optimization of human-centric objectives beyond mere task accuracy. Offline reinforcement learning has emerged as a promising method to model human-AI decision-making, allowing systems to tailor support policies that optimize for human learning and collaboration joy, rather than just correct answers [312]. This aligns with the broader vision of “Human-AI Co-Learning,” which posits that mutual growth and mutual benefits are prerequisites for sustaining high-performance teams in complex domains [38]. As systems become more autonomous, the need for symmetrical reality frameworks that accommodate dual perceptual centers becomes paramount, ensuring that both human and AI agents share a grounded understanding of the physical and virtual worlds they inhabit [60]. Ultimately, cognitive alignment is not a one-time configuration but a continuous, iterative process of mutual adaptation. By establishing robust mental models and fostering reciprocal learning, this phase of collaboration lays the cognitive groundwork necessary to tackle the specific pitfalls of automation bias and information asymmetry detailed in Section 5.4, ensuring that the convergence of human intuition and machine precision creates a collaborative intelligence greater than the sum of its parts.

5.4 Mitigating Automation Bias and Managing Information Asymmetry

5.4 Mitigating Automation Bias and Managing Information Asymmetry

Building upon the foundation of cognitive alignment and co-learning established in Section 5.3, the realization of Complementary Team Performance (CTP) in human-AI collaboration hinges critically on translating accurate mental models into appropriate reliance behaviors. While Section 5.3 detailed how humans and machines must mutually adapt to bridge the gap between technical explainability and human comprehension, Section 5.4 addresses the persistent cognitive barriers that threaten this synergy: automation bias and information asymmetry. Automation bias, a

phenomenon where humans disproportionately favor automated suggestions while neglecting contradictory information or their own superior contextual knowledge, creates a performance ceiling that bounds the human-AI dyad by the algorithm’s capabilities rather than leveraging the full potential of both agents. To counteract this, researchers have moved beyond static accuracy improvements to investigate sophisticated interaction designs that actively manage trust, calibrate confidence, and strategically leverage the unique epistemic resources held by human operators.

A primary strategy for mitigating automation bias involves the transparent communication of the AI’s uncertainty. Showing confidence scores has been identified as a mechanism to help calibrate human trust, allowing users to modulate their reliance based on the model’s self-assessed certainty. However, empirical evidence suggests that displaying confidence scores alone is often insufficient to improve joint decision-making outcomes. While such displays may align trust levels with model performance, they do not guarantee that humans can contribute the unique knowledge necessary to correct the AI’s errors [40]. The challenge lies in ensuring that the human possesses and utilizes information distinct from the AI’s training data. Theoretical frameworks posit that information asymmetry—where humans hold contextual knowledge unavailable to the algorithm—is an essential source of complementarity potential. Without this asymmetry, the human cannot act as a corrective filter, rendering the collaboration merely additive rather than synergistic [29]. Consequently, effective interface design must not only reveal the AI’s limitations but also explicitly prompt the human to activate their unique epistemic resources, bridging the gap between the “Theory of AI’s Mind” discussed previously and the practical execution of critical oversight.

To deepen this engagement and disrupt heuristic processing, recent studies advocate for “option-centric” rationale displays rather than static explanations. Traditional explainable AI (XAI) methods often provide post-hoc justifications for a single recommended action, which can inadvertently reinforce automation bias by making the AI’s logic appear more coherent than it is. In contrast, an option-centric approach presents the decision-making rationale for multiple potential actions, including those not selected by the AI. Experiments utilizing game-based simulations have demonstrated that conveying the intelligent assistant’s decision-making rationale across various options significantly enhances trust calibration. Users exposed to option-centric displays were able to calibrate their trust faster, leading to higher acceptance of valid recommendations and, crucially, a greater willingness to reject invalid ones, thereby improving overall task performance [313]. This approach forces the human to engage in comparative analysis rather than passive acceptance, effectively operationalizing the “critical use” training interventions mentioned in the preceding section.

Furthermore, the design of explanations must account for the cognitive costs associated with verifying AI predictions. Research indicates that humans strategically choose whether to engage with AI explanations based on a cost-benefit analysis. If the cognitive effort required to understand an explanation exceeds the perceived benefit of verifying the prediction, users will default to blind reliance or total dismissal. Therefore, explanations must be designed to reduce the verification cost sufficiently to make critical engagement worthwhile [314]. One promising avenue is the use of counterfactual explanations, which describe how specific input features would need to change to alter the model’s prediction. In high-stakes clinical decision-making contexts, counterfactual explanations have been shown to reduce over-reliance on incorrect AI outputs by approximately 21% compared to standard salient feature explanations. By prompting users to consider “what-if” scenarios, these explanations encourage analytical thinking and help users better estimate the accuracy of the AI model [315].

However, reliance behavior is not solely a function of the AI’s output; it is deeply influenced by the human’s self-assessment of their own competence. The Dunning-Kruger effect, where individuals

with lower competence overestimate their abilities, can lead to significant under-reliance on AI, while those with calibrated or underestimated self-confidence may over-rely. Interventions such as tutorials that explicitly reveal the fallibility of AI advice have proven effective in helping users with overestimated self-assessments recalibrate their reliance. Yet, these interventions must be carefully tailored, as they can inadvertently harm the reliance patterns of users who already underestimate their capabilities [316]. This highlights the necessity for adaptive systems that consider individual differences in cognitive motivation and self-confidence. For instance, “cognitive forcing functions”—design elements that compel users to pause and think before accepting a recommendation—have been shown to reduce over-reliance, particularly among users with a high need for cognition, though they may be perceived as less favorable subjectively [317]. These mechanisms serve as a crucial bridge to the dynamic trust calibration strategies explored in the subsequent section, ensuring that the human remains an active, skeptical partner rather than a passive verifier.

Ultimately, managing information asymmetry requires a shift from viewing the human as a passive recipient of machine intelligence to an active partner holding unique contextual keys. The deliberate design of interfaces that highlight the boundaries of the AI’s knowledge domain, combined with rationale displays that invite comparison and counterfactual reasoning, can foster a healthy skepticism. This approach ensures that humans do not merely defer to algorithmic outputs but actively contribute their distinct contextual insights, transforming the collaboration from a potential source of bias into a robust engine for complementary performance. As the field advances, the focus must remain on jointly optimizing the human-AI system, recognizing that an uncalibrated or strategically “over-confident” AI might sometimes be necessary to prompt the human to engage more deeply, thereby improving the final joint prediction [310]. These static and semi-static design principles set the stage for Section 5.5, which examines how trust must be dynamically calibrated over time in high-stakes and ad hoc teaming scenarios where reliance patterns are continuously tested by evolving system behaviors and novel contexts.

5.5 Dynamic Trust Calibration in High-Stakes and Ad Hoc Teaming

5.5 Dynamic Trust Calibration in High-Stakes and Ad Hoc Teaming

While the previous section detailed strategies to mitigate automation bias and leverage information asymmetry through static interface designs, realizing Complementary Team Performance (CTP) in practice requires addressing the temporal dimension of human-AI interaction. The efficacy of collaboration in scientific research and high-stakes operational environments hinges critically on dynamic trust calibration. Unlike the static reliance models discussed previously, trust is a time-variant variable that strengthens or decays based on continuous interaction, system performance, and the evolving mental models of human teammates. In high-stakes scenarios, where the costs of error are severe, these temporal dynamics become paramount. Research indicates that users build mental models of AI systems based on reliability and the presentation of decisions, yet modern systems often suffer from miscalibration, presenting confidently incorrect predictions that rapidly erode user confidence [318]. This erosion is not merely a momentary lapse; studies demonstrate that even a few instances of inaccurate confidence estimates can cause significant damage to trust and collaborative performance, with recovery processes proving to be exceptionally slow [318]. Consequently, effective collaboration requires mechanisms that do not simply maximize trust but calibrate it to align with the actual capabilities of the autonomous agent at any given moment, bridging the gap between the design principles of Section 5.4 and the cognitive load constraints explored in the subsequent section.

The evolution of trust through performance phases is particularly complex in ad hoc teaming scenarios, where humans and AI agents must collaborate without prior shared history or established protocols. In such novel situations, the agent must infer the human’s trust level through real-time interactions and observations to adapt its behavior accordingly [319]. This adaptability is crucial because individual differences in trust dynamics are profound; users do not respond uniformly to system failures or successes. For instance, clustering analyses of trust dynamics have identified distinct user personas, such as “believers,” “oscillators,” and “disbelievers,” each exhibiting unique patterns of trust formation and decay [320]. “Believers” tend to maintain higher initial trust and lose it more slowly during low-reliability operations, whereas “disbelievers” start with lower expectations and are less likely to recover trust after violations [321]. Recognizing these distinct profiles allows for the development of customized models that outperform generalized population models, enabling AI systems to tailor their transparency and interaction strategies to the specific cognitive and emotional states of their human partners [321], thereby preventing the cognitive overload that arises from mismatched explanation complexity.

In high-stakes environments, such as automated driving or emergency response, the consequences of trust miscalibration—either over-trust or under-trust—can be catastrophic. Over-trust leads to the misuse of automation in scenarios where its capabilities are inadequate, while under-trust results in the disuse of beneficial functionalities [322]. To mitigate these risks, recent frameworks propose viewing trust as a form of extended control, where interactive feedback becomes a necessary component of the user’s perception-action cycle [323]. By modeling trust dynamics using probabilistic approaches, such as Partially Observable Markov Decision Processes (POMDPs), systems can dynamically adjust automation transparency to optimize the trade-off between human workload and understanding [324]. For example, varying the amount and utility of information provided to the human based on real-time estimates of trust and workload has been shown to optimize team performance compared to non-adaptive systems [325]. This dynamic adjustment directly addresses the information bottleneck concerns raised in the following section by ensuring that explanation depth matches the user’s current capacity and trust state.

Furthermore, the nature of trust violations significantly influences the trajectory of calibration. Distinctions must be drawn between performance-based violations, where the system fails to execute a task correctly, and moral-based violations, where the system acts in ways perceived as unethical or misaligned with human values. Evidence suggests that moral-trust violations may inflict more severe and lasting damage to the human-robot relationship than equivalent performance failures, necessitating different repair strategies [326]. Effective trust repair often involves prompt correction of mistakes and explicit acknowledgment of limitations. Studies in human-robot collaboration have shown that when robots acknowledge their errors or limitations, human trust can not only be restored but potentially increased, highlighting the importance of transparent communication in maintaining long-term team cohesion [327]. Additionally, the strategy of “trust preservation” involves proactive behaviors where the robot adjusts its autonomy level to prevent trust degradation before it occurs, rather than merely reacting to it [328].

The challenge is further compounded in ad hoc teams facing novel scenarios where the AI agent may be continually learning and changing its behavior. In such dynamic contexts, early interactions heavily influence the formation of trust, which in turn dictates the human’s strategic behavior and willingness to rely on the system [329]. To address this, advanced modeling techniques, including Bayesian inference and neural network-based approaches, are being developed to predict trust transfer across different tasks and users, allowing for more fluent and robust human-machine integration [330]. Ultimately, achieving appropriate trust calibration in high-stakes and ad hoc settings

requires a shift from static, one-size-fits-all explanations to dynamic, context-aware interaction models that respect the temporal and individual nuances of human trust [39]. By integrating real-time physiological measures, behavioral indicators, and sophisticated computational models, future systems can foster resilient partnerships capable of navigating the uncertainties of complex scientific and operational workflows [331], setting the stage for a rigorous evaluation of how these dynamic interactions impact the fundamental trade-off between understanding and cognitive workload.

5.6 The Workload-Understanding Tradeoff and Evaluation Methodologies

5.6 Bridging the Understanding-Workload Tradeoff through Rigorous Task-Based Evaluation

Building upon the dynamic trust calibration mechanisms discussed in the previous section, the practical efficacy of Explainable Artificial Intelligence (XAI) in scientific workflows is fundamentally constrained by a critical tension: the tradeoff between the depth of human understanding and the cognitive workload imposed by explanations. While the theoretical promise of XAI lies in its ability to make opaque models transparent, the practical reality often involves overwhelming users with information that exceeds their cognitive processing capacity, thereby negating the benefits of the adaptive trust models previously described. This section examines the mechanisms governing this workload-understanding balance and argues that the current reliance on proxy tasks and subjective metrics in evaluation methodologies has obscured the true potential of human-AI collaboration, necessitating a shift toward rigorous, task-based assessments of actual team performance.

The relationship between explanation complexity and human cognitive load is not linear; rather, it is governed by the informational content relative to the user’s mental model. Explanations that are too simplistic fail to convey the nuance required for high-stakes scientific decision-making, while those that are overly detailed can induce cognitive overload, paradoxically reducing the user’s ability to discern correct from incorrect AI recommendations. Recent work has sought to characterize this balance through the lens of the Information Bottleneck method, an information-theoretic approach that automatically generates abstractions to maximize informativeness while minimizing complexity. Empirical studies have established a direct link between these information-theoretic concepts and human factors, demonstrating that workload correlates with complexity and understanding correlates with informativeness [332]. Furthermore, the impact of different explanation types on cognitive load is profound; for instance, local explanation types have been ranked as superior in terms of mental efficiency compared to other implementation-independent methods, directly influencing task performance and time in domains such as clinical diagnostics [333]. However, the design of these explanations must account for the hierarchical nature of human cognition. Research in human-robot teaming suggests that complex explanations should be abstracted to align with the level of detail desired by the human teammate, as hierarchical explanations have been shown to reduce cognitive load while improving task performance and behavior interpretability [334]. This alignment is crucial for maintaining the “appropriate reliance” cultivated through the dynamic trust calibration strategies outlined in Section 5.5.

Despite the clear theoretical understanding of this tradeoff, the field of XAI evaluation remains plagued by methodological flaws that hinder progress. A pervasive issue is the reliance on proxy tasks—artificial exercises where humans predict the AI’s output based on an explanation—rather than measuring performance on the actual decision-making task at hand. Evidence suggests that success in these proxy tasks does not predict success in real-world scenarios. Experiments comparing proxy evaluations with actual decision-making tasks have demonstrated a disconnect, where

improvements in proxy metrics fail to translate into improved human-AI team performance [335]. Similarly, the field has heavily depended on subjective measures, such as self-reported trust or preference, as proxies for objective performance. However, subjective confidence often diverges from actual accuracy; users may feel they understand a system or trust its output without actually making better decisions, leading to a false sense of security [335]. This reliance on misleading evaluation methods risks slowing the development of human-AI teams that can reliably outperform either entity acting alone, creating a dangerous gap between perceived and actual system reliability that dynamic trust models alone cannot bridge.

To address these gaps, the evaluation paradigm must shift toward measuring “complementary team performance” (CTP), where the joint output of the human and AI exceeds the capabilities of either alone. Current literature indicates that explanations often fail to generate CTP because they do not enable users to verify the correctness of AI predictions effectively. In many contexts, explanations merely increase the likelihood of users accepting AI recommendations, regardless of their correctness, thereby failing to foster appropriate reliance [28]. The failure to achieve CTP is often attributed to the inability of standard explanations to support verification, a prerequisite for humans to leverage their unique contextual knowledge against algorithmic outputs [336]. Consequently, evaluation frameworks must move beyond simple accuracy or plausibility metrics. Optimizing explanations for “plausibility”—making them sound reasonable to humans—can be counterproductive, as it may regularize machine explanations to mimic human reasoning errors, thereby breaking the assumption that plausible explanations imply correct decisions and leading to either undertrust or overtrust [337].

A robust evaluation methodology must therefore integrate objective, task-based metrics that capture the synergistic value of the human-AI dyad. This includes assessing the cost of coordination, which can sometimes exceed the benefits of collaboration if the explanation interface imposes excessive cognitive demands [338]. Furthermore, evaluations should consider the strategic choices users make regarding whether to engage with an explanation, framed within a cost-benefit analysis where task difficulty and explanation complexity weigh against the benefits of reliance [314]. Future research must prioritize the development of metrics that quantify the “explanatory efficacy” of systems through cyclic, symbiotic interactions, potentially utilizing neural correlates to measure the affective and cognitive states of users during iterative explanatory processes [339]. By abandoning misleading proxy tasks and subjective surrogates in favor of rigorous, grounded assessments of joint performance, the scientific community can better design XAI systems that truly augment human cognition rather than merely adding to the cognitive burden. Establishing these rigorous, task-based evaluation standards is a prerequisite for the subsequent exploration of the ethical tensions inherent in explainability, accountability, and the potential for manipulation discussed in the following section.

5.7 Ethical Tensions Between Explainability, Accountability, and Manipulation

5.7 Ethical Tensions Between Explainability, Accountability, and Manipulation

Building upon the necessity for rigorous, task-based evaluation established in the preceding section, the integration of Explainable AI (XAI) into scientific workflows introduces a complex layer of ethical challenges. While the shift toward measuring “complementary team performance” aims to ensure that explanations genuinely augment human cognition, the mere presence of an explanation does not guarantee ethical deployment. A critical examination reveals profound tensions where the pursuit of explainability can inadvertently compromise accountability, facilitate manipu-

lation, or create a false sense of security. This subsection addresses the moral dimensions of these interactions, specifically highlighting the risks of “painting black boxes white,” the potential for manipulative explanation design, and the inherent conflict between providing post-hoc justifications and maintaining clear lines of human accountability.

A primary concern in the current landscape of XAI is the phenomenon of “painting the black box white.” This occurs when explanation methods provide a veneer of transparency that fails to genuinely illuminate the underlying decision-making logic of sub-symbolic models, thereby masking their inherent opacity. Research indicates that rather than making black-box systems truly comprehensible, certain XAI approaches risk generating new errors or fostering unwarranted confidence in flawed systems, a dynamic termed the “white-box paradox” [340]. In high-stakes scientific domains, such as medical diagnostics or clinical decision support, this superficial transparency can be dangerous. If researchers or practitioners perceive an explanation as valid due to its cognitive appeal rather than its factual fidelity, they may defer critical judgment to the automation, a tendency known as technology dominance. Empirical studies in cardiology have shown that users’ perceptions of explanation quality do not always correlate with the actual reliability of the system, suggesting that the mere presence of an explanation can manipulate trust without enhancing understanding [340]. This directly undermines the goal of achieving true complementary performance, as users may rely on AI outputs based on persuasive but misleading rationales.

Furthermore, the design of explanations opens the door to potential manipulation, where the form of the explanation is optimized to persuade rather than inform. The “empirical turn” in AI interpretability has led to a proliferation of explanation tools that may lack denunciatory power, effectively serving as obstacles to ethical AI by creating an illusion of sensitivity to ethical incidents [341]. When explanations are generated post-hoc, they are susceptible to being tailored to align with user biases or organizational goals, rather than reflecting the true causal mechanisms of the model. For instance, in subjective decision-making contexts, Large Language Models (LLMs) have been found to generate arguments containing opinions and personal experiences that align with the user’s pre-existing views, thereby reinforcing confirmation bias and increasing perceived trustworthiness through manipulation rather than objective reasoning [342]. This raises ethical questions about the integrity of human-AI collaboration: if an explanation is designed to be convincing rather than truthful, it undermines the researcher’s agency and compromises the scientific method.

The tension between explainability and accountability is further exacerbated by the reliance on post-hoc justifications for complex, black-box systems. There is a growing argument that attempting to explain black-box models for high-stakes decisions, rather than employing inherently interpretable models, perpetuates bad practices and can lead to catastrophic societal harm [343]. Post-hoc explanations often provide only partial or biased information about the underlying algorithm, allowing developers to withhold critical details or divert attention from systemic flaws [344]. This creates a “responsibility gap” where it becomes difficult to attribute moral responsibility for adverse outcomes. If a system’s decision is justified by a local surrogate model that does not faithfully represent the global behavior of the black box, the human operator may be held accountable for decisions they could not reasonably understand or control. The concept of “meaningful human control” requires that responsibility attributed to a human be commensurate with their ability and authority to control the system; however, opaque systems with misleading explanations erode this foundation [24].

Moreover, the ethical landscape is complicated by the potential for explanations to be manipulated through adversarial techniques. Counterfactual explanations, which suggest how inputs could be

changed to alter outcomes, have been shown to be vulnerable to data poisoning, where malicious actors can increase the cost of recourse or manipulate the suggested actions [345]. Similarly, local rule-based explanations can be tuned to mimic the black box accurately in specific regions while failing to capture its logic globally, leading to deceptive faithfulness [346]. These vulnerabilities highlight that explainability is not a static property but a dynamic interaction that can be exploited. The risk extends to the broader socio-technical context, where the proxies used to define fairness, privacy, and transparency may be imperfect constructs that give the guise of embedded values while failing to address deeper ethical tensions [347].

Ultimately, the ethical deployment of XAI in scientific workflows requires moving beyond the simplistic view that more explanation is always better. It necessitates a critical awareness of how explanations function as rhetorical devices that can either empower or manipulate users. As noted in recent discourse, explainability should be viewed not merely as a technical right but as part of a moral duty to sustain reciprocity and agency in AI-led decision-making [348]. Without rigorous standards for fidelity, resistance to manipulation, and alignment with inherently interpretable architectures where possible, the drive for explainability risks becoming a tool for obscuring accountability rather than illuminating truth. Future interaction designs must therefore prioritize “bright patterns” that actively support user well-being and critical reasoning over persuasive but potentially misleading interface elements [349], ensuring that the human remains the ultimate arbiter of scientific integrity.

6 Ethical Frontiers, Research Integrity, and Socio-Technical Governance

6.1 The Accountability Ecosystem and the Erosion of Human Agency

6.1 The Accountability Ecosystem and the Erosion of Human Agency

The rapid integration of artificial intelligence into scientific workflows has precipitated a fundamental shift in the topology of responsibility, creating a critical tension between the promise of automated efficiency and the preservation of human oversight. Traditionally, accountability in research was anchored in the concept of individual human agency, where a specific actor could be identified as the source of a decision and held liable for its outcomes. However, as algorithmic systems evolve from passive tools to active agents capable of autonomous hypothesis generation and experimental design, this linear model of responsibility is fracturing. We are witnessing a transition toward a distributed accountability ecosystem, wherein responsibility is dispersed across a complex network of developers, data providers, deployers, and the algorithms themselves. This diffusion challenges the viability of “human-in-the-loop” oversight as a sufficient safeguard, necessitating a reimagining of governance through industry-specific accreditation and robust external audit mechanisms.

The erosion of human agency is perhaps most evident in the failure of policies that mandate human oversight without addressing the cognitive and structural limitations of human overseers. Evidence suggests that individuals are often unable to perform the desired oversight functions effectively, particularly when facing complex algorithmic systems that operate at speeds or scales beyond human comprehension [350]. When humans are placed in a position of nominal oversight without genuine epistemic access or causal power, these policies risk providing a false sense of security, effectively legitimizing the use of faulty algorithms while allowing vendors and agencies to shirk accountability for resulting harms [350]. Furthermore, the phenomenon of “imperfect delegation” arises when tasks are delegated to computationally driven systems, leading to distributed responsibility and induced acceptance, which further obscures the lines of direct public accountability [351]. This loss

of clear agency creates a vacuum where biased outcomes can flourish unchecked, directly setting the stage for the algorithmic biases and fairness gaps discussed in the following section.

To address these challenges, scholars argue for a move away from individual-centric models toward a holistic “accountability ecosystem.” This ecosystem perspective recognizes that different stakeholders require and contribute to specific accountability mechanisms, creating a web of interdependence rather than a chain of command [352]. In this framework, the focus shifts from blaming a single operator to evaluating the systemic conditions that allow harmful outcomes to occur. For instance, in high-stakes public sector contexts, operational oversight is increasingly recognized as a function of expert teams rather than isolated individuals, suggesting a “team-in-the-loop” approach that integrates AI into existing professional training and liability mechanisms [353]. This relational approach acknowledges that accountability is not a static attribute but a dynamic process shaped by the interactions between humans, algorithms, and institutional structures [354].

However, the mere existence of an ecosystem is insufficient without rigorous enforcement mechanisms. The current landscape of algorithmic accountability is often characterized by a lack of standardized practices, leading to “audit-washing” where organizations claim compliance without substantive verification [355]. To counter this, there is a growing consensus on the necessity of independent, third-party audits that go beyond technical performance metrics to include ethical and societal impact assessments [356]. Effective auditing requires more than black-box testing; it demands “outside-the-box” access to training data, development methodologies, and internal evaluations to truly scrutinize the lifecycle of AI systems [357]. Moreover, the audit process must be institutionalized, drawing lessons from mature fields like financial regulation and environmental health to establish clear standards and accreditation processes [356].

The implementation of such an ecosystem also requires industry-specific adaptations, as generic ethical principles often fail to translate into practical governance without contextualization. For example, the gambling sector demonstrates the need for tailored accountability principles that address specific risks like algorithmic bias and model explainability within that domain [352]. Similarly, the concept of Algorithm Review Boards (ARBs) has emerged as a critical internal governance mechanism, particularly in health and finance sectors, to bridge the gap between abstract ethical guidelines and operational reality [358]. These boards serve as a formalized space for deliberating on the risks and benefits of AI deployment, ensuring that accountability is embedded throughout the organizational lifecycle rather than applied as an afterthought.

Ultimately, restoring meaningful agency in an age of autonomous systems requires a dual strategy: strengthening the capacity of human overseers through better design and training, while simultaneously building a resilient infrastructure of external verification. This includes mandating disclosure of audit findings, involving stakeholders likely to be harmed in the audit process, and formalizing the accreditation of auditors themselves [359]. By shifting the paradigm from individual blame to systemic stewardship, the scientific community can foster an environment where AI augments human capability without eroding the foundational principles of research integrity and ethical responsibility. Yet, even with robust accountability structures in place, the substance of what is being audited remains paramount; if the underlying data and definitions of fairness are themselves skewed by hegemonic perspectives, no amount of procedural oversight can prevent the perpetuation of injustice. Thus, while establishing an accountability ecosystem is the necessary first step, it must be immediately followed by a critical examination of the algorithmic bias and the global inclusivity gap that threaten to undermine these very safeguards.

6.2 Algorithmic Bias, Fairness, and the Global Inclusivity Gap

6.2 Algorithmic Bias, Fairness, and the Global Inclusivity Gap

While the preceding section established the necessity of a robust accountability ecosystem to govern the erosion of human agency, it also noted a critical limitation: procedural safeguards alone cannot prevent injustice if the underlying data and definitions of fairness are themselves skewed. This section addresses that substantive void by examining how algorithmic bias threatens the *equity* of scientific knowledge, encoding historical injustices into the very fabric of research outputs. The integration of artificial intelligence into scientific workflows promises to accelerate discovery, yet this potential is critically undermined by the persistence of algorithmic bias, which risks amplifying social inequities and setting back decades of advances in civil liberty [360]. Although the field of Fairness, Accountability, Transparency, and Ethics (FATE) has emerged to address these challenges, current frameworks suffer from a profound “global inclusivity gap.” The discourse on algorithmic fairness remains overwhelmingly dominated by perspectives from More Economically Developed Countries (MEDC), frequently overlooking local knowledge, cultural pluralism, and the specific fairness needs of the Global South [361]. This hegemonic approach risks engendering a form of “conservative AI” that reproduces and reinforces the status quo of social inequality rather than disrupting it [362].

The dominance of Global North perspectives is particularly evident in the definition of fairness itself, which is often rooted in Western legal and philosophical traditions that fail to translate across geo-cultural contexts. For instance, conventional algorithmic fairness assumptions regarding sub-groups, values, and optimization methods have been found to be largely non-portable to contexts such as India, where socio-economic factors render data unreliable and the distance between models and oppressed communities is vast [363]. In such environments, merely translating technical fairness work serves as “window dressing” rather than a substantive solution, necessitating a re-imagining of Fair-ML that empowers oppressed communities and enables equitable ecosystems [364]. Furthermore, the lack of diversity within the computing community itself exacerbates this gap; the under-representation of diverse voices in academic and professional fields directly influences the types of biases encountered in datasets and the formulation of problems, leading to a disconnect between the desired outcomes of ethical AI and the realities of the populations they serve [365].

This exclusion yields tangible, often devastating consequences in high-stakes domains like healthcare and finance. In global health, algorithmic bias can exacerbate social inequities, making the development of trustworthy AI that intentionally designs for equity a moral imperative [366]. Yet, existing FATE frameworks often fail to account for the cultural, political, and ethical dimensions specific to different regions, resulting in solutions misaligned with local values. The problem is compounded by the assumption in many current approaches that protected characteristics, such as race or gender, are observable and easily categorized. This assumption collapses in contexts involving unobserved characteristics, such as sexual orientation or gender identity, or in societies where demographic categorization carries different social risks, including the expansion of surveillance infrastructure under the guise of fairness [367; 368].

Moreover, the technical mechanisms of bias often operate invisibly within the algorithms themselves, independent of biased data or malicious intent. For example, collaborative filtering algorithms inherently suffer from popularity and homogenizing biases that further marginalize data points on the margins of distributions, which often correspond to marginalized people [369]. When these systems are deployed globally without context-aware adjustments, they enforce a form of “color-blind

racism” that ignores the legacy of structural injustice and the specific ideological shifts required to address it [370]. The feedback loops generated by biased models can deepen disparities over time, acting as social stressors that reinforce negative real-world outcomes for specific groups, creating a cycle of reinforcement that is difficult to break [371].

To bridge this global inclusivity gap, a paradigm shift from abstract, universalist principles to community-led, context-specific strategies is essential. Research indicates that participatory design processes, which involve affected communities in the collection and curation of representative data, are crucial for responsible AI design [361]. Such approaches align with the framework of “Design for Values,” which advocates for proactively engaging societal stakeholders to translate fundamental human rights into context-dependent design requirements [372]. Additionally, the development of open-source auditing tools is vital to democratize the ability to assess fairness, removing barriers for practitioners worldwide who may lack access to proprietary evaluation suites [373].

Ultimately, achieving true inclusivity requires recognizing that algorithmic fairness is not solely a technical optimization problem but a socio-political challenge. It demands an interdisciplinary effort that integrates insights from sociology, political philosophy, and critical race studies to understand how structural injustices are reproduced through automated systems [374]. By moving beyond the limitations of current hegemonic approaches and embracing a pluralistic view of fairness, the scientific community can develop AI systems that do not merely mitigate bias but actively promote equity and justice across diverse global contexts. This involves not only refining technical metrics but also fostering a culture of accountability where the “algorithmic imprint”—the lasting societal effects of algorithmic deployment even after removal—is carefully considered in global governance frameworks [375]. Only through such comprehensive, culturally attuned strategies can the promise of AI in scientific research be realized without perpetuating the inequities of the past. However, even as we strive for equitable and fair systems, the foundation of scientific inquiry rests on the accuracy of the information produced. If the systems are fair but factually unreliable, the integrity of research remains compromised. Thus, having addressed the equity of knowledge, we must now turn to its veracity, examining the challenges of hallucinations and the truthfulness of generative outputs.

6.3 Hallucinations, Truthfulness, and the Integrity of Generative Outputs

6.3 Hallucinations, Truthfulness, and the Integrity of Generative Outputs

While the previous section highlighted how algorithmic bias threatens the *equity* of scientific knowledge by encoding historical injustices, this section addresses a more fundamental epistemic challenge: the threat to the *veracity* of that knowledge. The integration of Large Language Models (LLMs) into scientific workflows promises unprecedented acceleration in literature synthesis, hypothesis generation, and data analysis. However, this potential is critically undermined by the phenomenon of hallucination, where models generate content that appears factual but is ungrounded, contradicts established knowledge, or fabricates non-existent evidence. In the context of scientific research, where truthfulness and reproducibility are paramount, hallucinations pose a severe threat to research integrity. Unlike creative writing domains where fabrication might be tolerated or even desired, in science, a single hallucinated citation or fabricated data point can derail entire research trajectories, waste resources, and erode public trust. Recent surveys indicate that hallucination is arguably the most significant barrier to the safe deployment of LLMs in high-stakes environments such as healthcare, finance, and law [376]. The stakes are equally high in scientific discovery, where the “truthfulness” of generative outputs directly correlates with the validity of subsequent

experimental designs and theoretical conclusions.

Hallucinations in LLMs manifest in various forms, ranging from intrinsic errors, where the model contradicts its own context, to extrinsic errors, where the output conflicts with real-world facts. A formal analysis suggests that hallucination may be an innate limitation of LLMs, arising from the impossibility of learning all computable functions within finite time complexity [377]. This theoretical constraint implies that while mitigation strategies can reduce the frequency of errors, they cannot guarantee their complete elimination. Consequently, the scientific community must shift its focus from seeking perfect models to developing robust procedural safeguards and human-in-the-loop verification systems. The risks are particularly acute in specialized domains; for instance, empirical examinations in finance have revealed that off-the-shelf LLMs exhibit serious hallucination behaviors when explaining complex concepts or querying historical data, necessitating urgent research into domain-specific mitigation [378]. Similarly, in the medical domain, the generation of plausible yet unverified information can lead to devastating consequences, prompting the creation of specialized benchmarks like Med-HALT to rigorously evaluate model reliability in healthcare contexts [379].

To preserve the integrity of generative outputs, a multi-layered approach combining technical innovations and procedural rigor is essential. Technically, mitigation strategies can be categorized into pre-generation, during-generation, and post-generation interventions. Pre-generation approaches involve optimizing the model’s familiarity with input concepts; for example, the SELF-FAMILIARITY technique evaluates whether a model recognizes the concepts in a prompt and withholds generation if the concepts are unfamiliar, thereby emulating a human-like refusal to speculate on unknown topics [380]. During generation, methods such as Chain-of-Verification (CoVe) enable models to deliberate on their own responses by planning verification questions and answering them independently to correct mistakes before finalizing the output [381]. Furthermore, decoding strategies like OPERA introduce over-trust penalties and retrospection-allocation mechanisms to prevent models from overly relying on summary tokens, a common cause of visual and textual hallucinations in multi-modal systems [382].

Post-generation verification remains a critical line of defense, often involving retrieval-augmented generation (RAG) or external fact-checking pipelines. However, standard RAG can introduce external hallucinations if the retrieved information is irrelevant. Advanced frameworks like Rowen address this by employing adaptive retrieval, activating external knowledge sources only when internal consistency checks detect potential hallucinations [383]. Another promising direction is the use of fine-grained AI feedback to detect and rewrite hallucinated segments. For instance, the Re-Ex method utilizes a three-step process of evidence retrieval, error explanation, and response revision to significantly reduce factual errors with lower computational overhead than previous methods [384]. Additionally, token-level uncertainty quantification has emerged as a powerful tool for fact-checking, allowing systems to flag atomic claims with high uncertainty scores for human review [385].

Despite these technical advances, procedural governance and human oversight remain indispensable, mirroring the calls for community-led design seen in the discourse on fairness. The concept of “human-in-the-loop” must evolve from passive review to active collaboration, where researchers are trained to recognize the specific signatures of AI hallucinations. Studies on human perception suggest that while warnings about potential inaccuracies can improve detection rates, users still struggle to distinguish between minor and major hallucinations without specialized tools [386]. Therefore, specialized auditing pipelines are necessary to validate AI-generated content before it enters the scientific record. These pipelines should incorporate diverse evaluation metrics, such

as the Hallucination Vulnerability Index (HVI), to rank models based on their susceptibility to different types of errors, ranging from acronym ambiguity to fabricated data [387].

Furthermore, the development of hallucination-minimized solutions requires a shift towards domain-specific fine-tuning and the use of synthetic tasks designed to elicit and correct errors. The SynTra method, for example, demonstrates that optimizing a model on synthetic tasks where hallucinations are easy to measure can transfer to reduced hallucinations in real-world summarization tasks [388]. In high-stakes decision-making scenarios, such as clinical diagnosis or financial forecasting, the journey toward reliable AI involves iterative stages of prototyping, scaling, and evolution using human feedback to ensure that autonomous reports and alerts maintain high quality [132]. Ultimately, ensuring the truthfulness of generative outputs in science is not merely a technical challenge but a socio-technical imperative. It demands a convergence of advanced algorithmic safeguards, rigorous benchmarking, and a culture of skepticism and verification among researchers. As the field moves forward, the goal is not to replace human judgment but to augment it with tools that are transparent about their uncertainties and robust against the inherent tendency to hallucinate, thereby safeguarding the epistemic foundation of scientific inquiry. This focus on epistemic integrity sets the stage for the next critical dimension of responsible AI: the environmental sustainability of the computational infrastructure required to train and verify these increasingly complex models.

6.4 Environmental Sustainability and the Carbon Footprint of Scientific AI

6.4 Environmental Sustainability and the Carbon Footprint of Scientific AI

Building upon the discussion of epistemic integrity and the verification costs associated with mitigating hallucinations, this section addresses a critical, yet often overlooked, dimension of responsible AI: the environmental sustainability of the computational infrastructure required to train and verify these increasingly complex models. The rapid integration of Artificial Intelligence (AI) into scientific workflows has ushered in an era of unprecedented discovery, yet this progress is shadowed by a profound environmental paradox. While AI is increasingly deployed as a critical tool to model climate change, optimize energy grids, and design sustainable materials, the computational infrastructure required to train and deploy these very models generates a massive ecological footprint. This section critically assesses the environmental costs associated with large-scale generative models in science, dissecting the tension between AI as a solution for sustainability and AI as a driver of environmental degradation. The super-linear growth in model complexity, exacerbated by the need for extensive verification loops to ensure truthfulness, has led to energy consumption patterns that threaten to undermine the very sustainability goals the scientific community aims to achieve.

The primary environmental concern stems from the immense computational power required for training state-of-the-art models. The carbon footprint of AI is no longer negligible; it is a significant contributor to global greenhouse gas (GHG) emissions. Studies indicate that the computations required for deep learning research have doubled every few months, resulting in a staggering increase in energy demand since 2012 [389]. Specifically, the training of large language models (LLMs) involves massive CO₂ emissions and substantial water consumption for cooling data centers. For instance, the training of a single 176-billion parameter model like BLOOM was estimated to emit approximately 50.5 tonnes of carbon equivalent when accounting for the entire lifecycle, including equipment manufacturing and operational energy [390]. Furthermore, the water footprint of AI has remained largely under the radar despite its severity; training GPT-3 in U.S. data centers could directly evaporate 700,000 liters of clean freshwater, raising alarms about freshwater scarcity in the wake of rapidly growing AI demands [391].

This environmental burden is exacerbated by the prevailing trend toward “generality” in AI systems. Multi-purpose, generative architectures are often orders of magnitude more expensive in terms of energy and carbon emissions compared to task-specific systems, even when controlling for the number of parameters [392]. In the context of scientific research, where reproducibility and iterative experimentation are paramount, the repeated training and fine-tuning of massive models can lead to compounding environmental costs. The ICT sector already contributes up to 3.9 percent of global GHG emissions, a figure surpassing that of global air travel, and the unchecked expansion of generative AI threatens to accelerate this trajectory [393].

The paradox of “AI for Green” versus “Green AI” is central to this discourse. On one hand, AI is utilized to create more sustainable production methods and model ecological systems, serving as a valuable asset in fighting environmental degradation [394]. On the other hand, the negative impacts of these energy-consuming technologies often offset their ecological benefits if not rigorously managed. Research suggests that the drive for less energy-consuming AI is frequently motivated by cost constraints and energy autonomy rather than genuine environmental considerations, leading to a rebound effect where efficiency gains fuel an increase in model complexity rather than a reduction in total consumption [394]. Consequently, efficiency alone is insufficient to ensure environmental sustainability; a holistic systems thinking approach is required to address the multifaceted variables interacting with machine learning systems [395].

To mitigate these impacts, the scientific community must advocate for robust “Green AI” regulations and the integration of environmental indicators into standard model evaluation metrics. Current proposals for AI regulation, such as the EU AI Act, often focus on trustworthiness and accountability but lack a robust roadmap for environmental sustainability [393]. There is an urgent need to expand regulatory frameworks to include transparency mechanisms, such as mandatory disclosure of the GHG and water footprints of AI systems, and to implement consumption caps. Proposals include integrating AI data centers into emissions trading schemes and adopting “sustainability-by-design” principles [393]. Furthermore, the development of standardized metrics is crucial. Frameworks like the Sustainability Criteria and Indicators for Artificial Intelligence Systems (SCAIS) offer a holistic perspective, providing specific criteria and indicators to structure the discourse on sustainable AI and support the conscious development of these systems [396].

Methodological shifts are also necessary to align scientific AI with ecological limits. This includes prioritizing energy-efficient training strategies, such as Bayesian optimization for hyperparameter tuning, which has been shown to dominate other strategies in reducing energy consumption [397]. Additionally, the adoption of data-centric approaches, such as training models on “elite samples” rather than massive datasets, can yield remarkable energy savings of up to 98% with negligible performance degradation [398]. The concept of “Green Federated Learning” further offers a pathway to minimize carbon emissions by optimizing parameters and making design choices that balance performance with environmental impact, particularly in geographically distributed settings [399].

Ultimately, addressing the environmental cost of scientific AI requires a fundamental reorientation of research incentives. The current academic culture, which often prioritizes state-of-the-art accuracy above all else, must evolve to value efficiency and sustainability as equal metrics of success. As noted in position papers, making efficiency an evaluation criterion alongside accuracy is essential to make AI both greener and more inclusive, preventing the widening of the technology gap between well-funded institutions and those with limited resources [389]. Without such systemic changes, including the adoption of comprehensive metrics catalogues that track process, resource, and product impacts, the scientific enterprise risks accelerating the climate crisis it seeks to solve [198]. The path forward demands a convergence of technical innovation, regulatory oversight, and

ethical responsibility to ensure that the future of scientific AI is not only intelligent but also sustainable. By establishing these environmental constraints as a core component of responsible AI, we set the stage for the next critical challenge: bridging the “implementation gap” between these high-level ethical and environmental principles and the concrete operational governance required to enforce them.

6.5 From Abstract Principles to Operational Governance: Metrics, Patterns, and Audits

6.5 From Abstract Principles to Operational Governance: Metrics, Patterns, and Audits

While the previous section highlighted the urgent need to integrate environmental sustainability into AI evaluation, a broader challenge persists: the “implementation gap” between high-level ethical consensus and concrete practice. The proliferation of abstract guidelines has created a paradoxical landscape where global agreement on values like transparency, fairness, and accountability coexists with a profound divergence in how these precepts are operationalized [400]. Moving from theoretical ethics to actionable governance requires shifting focus from truisms to systematic, system-level interventions that permeate the entire lifecycle of AI development and deployment in research. This subsection examines three pivotal mechanisms for bridging this divide—Responsible AI pattern catalogues, tripartite metrics frameworks, and multi-layered auditing architectures—which collectively provide the structural backbone necessary to enforce not only general ethical standards but also specific imperatives like the environmental constraints discussed previously.

A primary strategy for operationalizing ethics is the development of pattern catalogues that translate abstract principles into concrete engineering and governance practices. Historically, significant efforts have focused on algorithm-level solutions, often neglecting the broader system-level components where ethical issues, including excessive resource consumption, frequently arise [401]. To address this, researchers have proposed comprehensive pattern collections derived from multivoical literature reviews, categorizing interventions into multi-level governance patterns, trustworthy process patterns, and responsible-AI-by-design product patterns. These catalogues provide stakeholders with actionable guidance, ensuring that responsibility is embedded not just in the model weights, but in the architectural styles, data pipelines, and organizational processes surrounding the AI system [402]. Furthermore, empirical studies with practitioners reveal that without such concrete patterns, engineers often struggle to navigate the “hodgepodge” of conflicting values and institutional structures, leading to dislocated accountabilities where no single actor feels responsible for the system’s ethical posture [403; 404]. By adopting a pattern-oriented approach, organizations can standardize the mitigation of risks at every stage of development, from requirements elicitation to deployment, effectively closing the gap between principle and practice [405].

Complementing pattern-based guidance is the necessity for a robust metrics framework capable of quantifying responsible AI attributes. Traditional performance metrics are insufficient for capturing the nuanced dimensions of ethics, necessitating a tripartite framework encompassing process, resource, and product metrics. Process metrics underpin procedural integrity by evaluating the rigor of governance mechanisms and stakeholder engagement throughout the development lifecycle [198]. Resource metrics assess the availability and adequacy of tools, datasets, and computational infrastructure required to support ethical objectives; crucially, this category provides the mechanism to track the environmental footprints emphasized in Section 6.4, such as energy usage and carbon emissions, alongside the adoption of FAIR-compliant data for training large language models

[406]. Product metrics, conversely, reflect the tangible outputs of AI systems, measuring attributes like fairness, robustness, and explainability in deployed models. For instance, frameworks like Com-
plAI have been developed to unify these factors, offering a single trust factor that evaluates models based on explainability, drift susceptibility, and fairness across diverse scenarios [407]. This granular approach to measurement is essential for moving from subjective ethical assurances to objective, auditable evidence of responsible behavior, particularly in high-stakes scientific domains where the cost of failure is substantial.

Finally, the unique challenges posed by generative AI and large language models (LLMs) demand a reimagining of auditing methodologies that can verify both ethical alignment and resource efficiency. Standard auditing procedures often fail to address the emergent capabilities and adaptability of LLMs, necessitating a multi-layered approach. Recent proposals outline a three-tiered auditing blueprint consisting of governance audits for technology providers, model audits conducted post-pre-training but prior to release, and application audits for downstream uses [408]. This layered strategy ensures that accountability is distributed across the supply chain, addressing the specific risks associated with foundation models versus their specific implementations. Moreover, ethics-based audits have proven effective in probing the moral reasoning and normative values embedded within leading LLMs, revealing cultural biases and authoritarian tendencies that might otherwise remain hidden [409]. By integrating argument-based assurance methods, which structure reasoning about ethical acceptability similar to safety cases in engineering, the scientific community can foster greater confidence in AI-enabled research workflows [410]. These audits serve as the enforcement mechanism for the metrics and patterns defined above, ensuring that claims of sustainability and responsibility are verified rather than merely asserted.

Ultimately, the convergence of pattern catalogues, comprehensive metrics, and layered audits forms the backbone of a mature socio-technical governance ecosystem. These mechanisms do not operate in isolation; rather, they reinforce one another to create a continuous feedback loop of improvement that addresses both the environmental paradoxes and the transparency deficits facing modern science. As the field moves toward dynamic scholarship and autonomous discovery, the integration of these operational tools will be paramount in ensuring that the acceleration of scientific knowledge does not come at the expense of research integrity, societal trust, or human agency. This operational foundation sets the stage for the specific reporting standards and documentation frameworks, such as AI Usage Cards, which are detailed in the following section to make these governance efforts visible and actionable to the broader research community. The shift from abstract principles to operational governance represents not merely a technical challenge, but a fundamental restructuring of how scientific responsibility is conceived, measured, and enforced in the age of artificial intelligence.

6.6 Transparency Frameworks, AI Usage Cards, and Reporting Standards

6.6 Transparency Frameworks, AI Usage Cards, and Reporting Standards

Building upon the operational governance mechanisms detailed in the previous subsection—specifically pattern catalogues, tripartite metrics, and layered audits—the scientific community requires tangible artifacts to make these abstract measures visible and actionable. While Section 6.5 established the structural backbone for enforcing responsibility through metrics and audits, Section 6.6 focuses on the communication layer: standardized transparency frameworks and reporting standards that translate internal governance into external accountability. As Large Language Models (LLMs) and foundation models become ubiquitous in research, traditional

paradigms of authorship and methodology disclosure are increasingly insufficient to capture the complexity of human-AI collaboration. Without robust transparency frameworks, the scientific community risks an erosion of trust, where the provenance of ideas, the validity of data, and the attribution of credit become obscured by algorithmic opacity. To address these challenges, the field is moving toward structured documentation standards, most notably the adoption of “AI Usage Cards,” alongside comprehensive model reporting and human-centered explainability techniques designed to clarify limitations, uncertainty, and social attributions.

The foundational argument for such standardization rests on the necessity of defining responsible AI use through clear dimensions of transparency, integrity, and accountability, directly supporting the audit trails discussed previously. Recent proposals have introduced “AI Usage Cards” as a standardized mechanism for researchers to report the specific role of generative AI in their work [157]. Unlike generic acknowledgments, these cards provide a structured format that allows the research community to trace, compare, and question various forms of AI usage, thereby supporting the development of accepted community norms. This approach mirrors the successful deployment of “Data Cards,” which were originally designed to foster transparent and purposeful documentation of datasets by summarizing essential facts about origins, collection methods, and intended uses [411]. Just as Data Cards shifted the burden of understanding dataset nuances from the user to the documentation itself, AI Usage Cards aim to make the implicit contributions of generative models explicit, ensuring that the distinction between human insight and machine generation remains clear. These cards serve as the primary interface between the internal metrics defined in Section 6.5 and the external scrutiny required by the governance models explored in the following section.

However, the current landscape of AI documentation reveals significant gaps between these aspirational guidelines and actual practice, undermining the efficacy of the operational governance proposed earlier. A systematic analysis of over 32,000 AI model cards on platforms like Hugging Face indicates that while many models include documentation, the informativeness is highly uneven [412]. Critical sections addressing environmental impact, limitations, and evaluation often exhibit low completion rates, whereas training details are more consistently reported. This disparity suggests that existing documentation practices frequently fail to provide the contextual information necessary for researchers to assess risks or understand the boundaries of model applicability, thereby hindering the resource and product metrics outlined in the preceding subsection. Furthermore, the lack of standardization across publishers and journals exacerbates the problem; a bibliometric analysis of top academic publishers reveals substantial heterogeneity in how generative AI use is disclosed, with varying requirements for where and how such disclosures should appear [156]. This fragmentation burdens authors and threatens the effectiveness of regulatory efforts, highlighting the need for unified, consensus-driven reporting standards that can serve as reliable inputs for the Algorithm Review Boards and Dual Governance models discussed next.

To bridge the gap between high-level ethical principles and operational reality, transparency frameworks must evolve beyond static documents to include dynamic, computable artifacts that facilitate the automated auditing mechanisms described in Section 6.5. For instance, the application of ontologies to model cards can generate computable reports that link machine learning information from biomedical research, thereby supporting FAIR (Findable, Accessible, Interoperable, Reusable) data challenges [413]. Additionally, the concept of “Foundation Model Transparency Reports” has emerged, drawing upon social media transparency practices to institutionalize reporting for foundation models while the industry is still nascent [414]. These reports utilize extensive indicators to measure transparency across upstream resources, model details, and downstream uses, aiming to reduce compliance costs and align with emerging government policies such as the EU AI Act. By

transforming documentation into machine-readable evidence, these artifacts enable the continuous, multi-layered auditing architectures necessary for modern scientific workflows.

Crucially, effective reporting standards must incorporate human-centered explainability techniques that go beyond technical specifications to address social and cognitive factors, ensuring that the “social license” discussed in the subsequent section is well-founded. The notion of “Social Transparency” expands the design space of Explainable AI (XAI) by incorporating socio-organizational contexts into explanations, helping users understand not just how an algorithm works, but also the social attributions and potential biases embedded within it [294]. This is particularly vital for LLMs, which can simulate personas and lead to social misattributions where users mistakenly assign human-like agency or empathy to the system [300]. By enhancing transparency frameworks with questions that clarify the specific social attributions assigned to LLMs by designers and users, the scientific community can mitigate risks of emotional manipulation and unwarranted trust, fostering the type of informed oversight required by the Hourglass Model and Dual Governance paradigms.

Moreover, transparency must encompass the communication of uncertainty and the limits of model knowledge, a critical component for the risk assessment procedures central to any governance model. Research indicates that uncalibrated AI confidence—where models are either overconfident or underconfident—can severely hinder human-AI collaboration by promoting misuse or disuse of AI tools [415]. Therefore, reporting standards should mandate the inclusion of confidence metrics and explicit statements regarding a model’s tendency to extrapolate beyond its training distribution, effectively reporting when a model is “clueless” [416]. Such disclosures are essential for preventing the blind reliance on AI outputs in high-stakes scientific decision-making and provide the empirical basis for the “team-in-the-loop” oversight strategies detailed in the following subsection.

Finally, the push for standardized reporting is increasingly driven by regulatory pressures and the need for third-party auditing, creating a direct link between the documentation practices here and the enforcement mechanisms of Section 6.7. Proposals merging EU regulations with research trends suggest the creation of “use-case and operation cards” alongside updated data and model cards to cope with legal requirements for safety and accountability [417]. These standardized cards aim to facilitate efficient auditing while remaining accessible to diverse stakeholders. By adopting comprehensive frameworks like AI Usage Cards, integrating social transparency, and mandating the reporting of uncertainty and limitations, the scientific community can establish a culture of responsible AI usage. This shift from ad-hoc disclosure to rigorous, standardized reporting provides the essential documentation layer that allows the socio-technical governance models and dual regulation paradigms discussed in the next section to function effectively, ensuring that research integrity and public trust are maintained in the age of generative AI.

6.7 Socio-Technical Governance Models and the Dual Regulation Paradigm

6.7 Socio-Technical Governance Models and the Dual Regulation Paradigm

While the previous subsection established the critical need for standardized transparency frameworks and reporting artifacts like AI Usage Cards, these documentation mechanisms alone are insufficient to ensure responsible AI integration without robust governance structures to enforce and operationalize them. As artificial intelligence systems transition from static tools to agentic collaborators in scientific workflows, the governance frameworks required to oversee them must evolve from simple compliance checklists into dynamic, socio-technical architectures. The complexity of

modern AI deployment, particularly in high-stakes research environments, demands a multi-layered approach that harmonizes organizational internal controls with broader regulatory and community-based oversight mechanisms. This subsection synthesizes emerging governance models that address this need, specifically focusing on the “Hourglass Model” for organizational alignment, the operationalization of Algorithm Review Boards (ARBs), and the “Dual Governance” paradigm that integrates centralized regulation with decentralized safety mechanisms.

At the organizational level, the primary challenge lies in translating the abstract ethical principles and reporting standards discussed previously into concrete operational practices. The “Hourglass Model” of organizational AI governance offers a robust framework for bridging this gap by structuring governance requirements across three distinct layers: the environmental, the organizational, and the AI system levels [418]. This model is particularly effective because it connects high-level societal and regulatory expectations at the top with specific technical lifecycle requirements at the bottom, converging through the organization’s internal processes. By mapping governance requirements directly to the AI system lifecycle, the Hourglass Model ensures that ethical considerations and transparency mandates are not merely post-hoc audits but are embedded throughout the development and deployment phases [418]. This systemic approach is critical for scientific institutions seeking to align their internal research protocols with forthcoming regulations, such as the European AI Act, while maintaining the agility required for innovation.

Complementing the structural alignment provided by the Hourglass Model is the practical implementation of internal oversight bodies, most notably Algorithm Review Boards (ARBs). Empirical investigations into organizational Responsible AI (RAI) governance reveal that ARBs have emerged as one of the most impactful internal mechanisms for mitigating AI risks across diverse sectors, including academia, government, and industry [358]. Unlike traditional Institutional Review Boards (IRBs) which focus primarily on human subjects, ARBs are specifically designed to evaluate the technical and ethical implications of algorithmic systems, including the verification of the transparency artifacts mentioned in Section 6.6. Research indicates that the success of these boards relies heavily on leadership buy-in and their integration with existing internal regulatory approaches, rather than operating as isolated silos [358]. However, the efficacy of ARBs is often challenged by financial tensions and the lack of standardized metrics for success, highlighting the need for more rigorous definitions of their scope and membership [358]. Furthermore, in high-stakes public sector applications, the concept of oversight is expanding from individual human-in-the-loop models to “team-in-the-loop” frameworks, acknowledging that operational oversight is often a collective, nested process within professional teams rather than the responsibility of a single decision-maker [353].

Beyond the boundaries of individual organizations, the governance of generative AI and autonomous agents requires a broader paradigm that addresses the limitations of purely top-down regulation. Centralized government regulations often face criticism for lacking clarity, uniformity, and the ability to keep pace with rapid technological evolution, potentially stifling innovation while failing to provide adequate safety guarantees [194]. Conversely, purely decentralized or crowdsourced safety mechanisms, while agile, often suffer from deficiencies in enforcement and adequate oversight [194]. To resolve this tension, the “Dual Governance” framework proposes a synergistic marriage of centralized state regulation and community-driven safety mechanisms. In this model, government bodies establish baseline legal standards and accountability structures, while crowdsourced communities develop adaptive safety tools and real-time monitoring systems that can respond to emerging threats faster than legislative bodies [194]. This approach mirrors findings in other domains where coordinating top-down agencies with bottom-up constructs is identified as an essential

design principle for managing complex, multiscale societal systems [419].

The necessity for such hybrid models is underscored by the fragility of purely decentralized or trustless socio-technical systems, where emergent centralization can lead to catastrophic cascades if not properly governed by broader ecosystem considerations [420]. Moreover, relying solely on human oversight policies without robust institutional backing has been shown to provide a false sense of security, often legitimizing faulty algorithms without addressing fundamental flaws [350]. Therefore, effective governance must shift towards institutional oversight that requires empirical justification for algorithm adoption and democratic public review [350]. By integrating the structural clarity of the Hourglass Model, the procedural rigor of ARBs, and the adaptive resilience of Dual Governance, the scientific community can establish a comprehensive socio-technical governance ecosystem. This ecosystem not only ensures research integrity and safety but also fosters an environment where human agency is preserved and aligned with societal values, preventing the erosion of accountability in an era of increasingly agentic algorithmic systems [421]. Ultimately, establishing these robust internal and external governance layers is a prerequisite for securing the public trust and “social license” discussed in the following subsection, as effective governance provides the tangible evidence of responsibility required to calm public skepticism and align scientific AI with societal values.

6.8 Public Perception, Trust Calibration, and the Social License for AI in Science

6.8 Public Perception, Trust Calibration, and the Social License for AI in Science

While the preceding subsection established a robust architecture for socio-technical governance—integrating the structural clarity of the Hourglass Model, the procedural rigor of Algorithm Review Boards, and the adaptive resilience of Dual Governance—these mechanisms alone cannot guarantee the successful integration of Artificial Intelligence (AI) into scientific workflows. The ultimate viability of AI-augmented science is fundamentally contingent upon public trust and the granting of a “social license” to operate. Even with comprehensive internal controls and external regulations, a significant disconnect remains between these high-level ethical guidelines and the public’s understanding of AI risks and benefits. This gap poses a critical threat to the long-term legitimacy of deploying powerful tools in high-stakes domains such as healthcare, climate modeling, and public policy, as public skepticism can erode the societal support required for such endeavors. Research indicates that while there is widespread recognition of AI’s transformative potential, public sentiment is often characterized by a complex mixture of excitement, utility, worry, and futuristic speculation, varying significantly across different cultural and national contexts [422]. Understanding these nuanced perceptions is therefore essential for calibrating trust and ensuring that the deployment of scientific AI aligns with broader societal values.

A primary challenge in securing this social license is the misalignment between the risks prioritized by the governance structures discussed in Section 6.7 and those perceived by the public. While technical ethics and regulatory frameworks often focus on abstract principles like fairness, accountability, and transparency, empirical studies suggest that the general public rarely cites these specific criteria as their primary concern. Instead, citizens tend to focus on specific applications and immediate tangible impacts, with ethical issues often remaining secondary or unrecognized until they directly affect individual well-being [423]. This discrepancy creates a dangerous blind spot: the populations most vulnerable to algorithmic biases or errors are often the least aware of the ethical risks inherent in these systems. Furthermore, when ethical concerns do arise, they frequently manifest as a deep-seated distrust of the institutions managing AI, with surveys revealing that the public holds only low to moderate trust in both governmental and corporate entities to responsibly

govern these technologies [424]. This deficit of institutional trust complicates the calibration of reliance on AI systems, as the public may reject beneficial scientific innovations due to a lack of confidence in the very governance structures designed to oversee them.

The erosion of trust is often exacerbated by system designs that prioritize technical efficiency over human moral intuitions, rendering even well-governed systems socially unacceptable. For instance, AI systems designed to minimize false positives may inadvertently miscategorize trustworthy individuals, imposing significant moral costs that fuel public resentment and distrust. This phenomenon suggests that restoring trust requires a shift toward “humble trust,” where system designs explicitly account for the moral weight of false negatives and the potential for misplaced distrust [425]. In the context of scientific research, where AI systems increasingly influence critical decisions ranging from drug discovery to resource allocation, failing to address these moral dimensions can lead to a rejection of AI-driven findings, regardless of their statistical validity or regulatory compliance. The perception that AI decisions are opaque, unaccountable, or misaligned with human rights further deepens this divide, highlighting the urgent need to bridge the gap between the abstract value languages of governance frameworks and concrete design requirements [372].

To secure a durable social license for AI in science, the community must move beyond the top-down regulatory compliance and internal oversight models detailed previously, embracing instead participatory design processes that actively engage societal stakeholders. The current landscape of AI ethics is dominated by a proliferation of high-level principles, yet these often fail to translate into practice due to a lack of context-specific implementation strategies and genuine stakeholder involvement [426]. Participatory approaches offer a pathway to democratize AI governance, allowing non-expert publics to articulate their concerns, define acceptable use cases, and identify potential harms that expert communities and review boards might overlook. Frameworks such as “Particip-AI” demonstrate that involving diverse public voices can surface nuanced risks and benefits, challenging techno-solutionist narratives and ensuring that AI development reflects a broader consensus on the common good [427]. Such engagement is not merely a procedural formality but a necessary mechanism for aligning technological trajectories with societal needs and human well-being, effectively closing the loop between institutional governance and public expectation.

Moreover, achieving this alignment requires a fundamental rethinking of how ethical responsibilities are distributed across the AI lifecycle, extending the socio-technical approach beyond organizational boundaries. The traditional model, which places the burden of ethics solely on individual developers or abstract organizational charters, is insufficient. Instead, a holistic socio-technical approach is needed that recognizes the interdependence of technical capabilities, legal frameworks, and ethical norms [428]. This involves creating feedback channels where public deliberation directly informs safety protocols and design choices, transforming AI safety from a static procedure into a dynamic, deliberative goal [429]. By integrating public values into the core of scientific AI workflows, researchers can foster a sense of shared ownership and responsibility. This is particularly crucial given that the “social dilemma” of AI development—where individual incentives to race ahead conflict with collective safety—can only be resolved through professionalization and a commitment to the public interest that transcends competitive pressures [430].

Ultimately, the social license for AI in science is not a permanent grant but a continuous negotiation that depends on transparent communication, demonstrable benevolence, and responsive governance. As AI systems become more autonomous and impactful, the cost of ignoring public perception grows exponentially. Strategies that combine the rigorous technical safeguards and governance layers established in Section 6.7 with inclusive, community-led governance models are essential to prevent a backlash that could stifle scientific progress. By adopting frameworks that prioritize

“Provable Beneficence” and “Responsible Stewardship,” and by ensuring that the voices of affected communities are integral to the design process, the scientific community can build the robust trust necessary to harness AI for the advancement of human knowledge and societal welfare [431]. Only through such a holistic, human-centered approach can the promise of AI in science be fully realized without compromising the ethical fabric of society.

7 Domain-Specific Implementations and Cross-Disciplinary Case Studies

7.1 Biomedical Discovery and Clinical Decision-Making

7.1 Biomedical Discovery and Clinical Decision-Making

The integration of artificial intelligence into biomedical discovery and clinical decision-making represents a transformative shift in modern healthcare, mirroring the broader transition from simple automation to complex human-AI teaming. Similar to the paradigm shifts observed in other scientific domains, this field is moving beyond standalone diagnostic tools toward collaborative partners that augment human expertise, streamline workflows, and accelerate discovery. A central theme is the development of multimodal frameworks capable of synthesizing diverse data sources to support comprehensive clinical reasoning. Given that medicine inherently spans text, imaging, genomics, and time-series data, effective systems must flexibly encode and interpret these inputs at scale. Recent advancements have introduced generalist biomedical AI systems, such as Med-PaLM Multimodal, which utilize a single set of model weights to encode clinical language, imaging, and genomics, often surpassing specialist models in performance across diverse tasks [432]. Furthermore, unified frameworks like the Holistic AI in Medicine (HAIM) have demonstrated that leveraging multiple input modalities—tabular, time-series, text, and images—consistently produces models that outperform single-source approaches by significant margins, highlighting the necessity of multimodal integration for accurate predictive systems [433]. These capabilities are further refined by Large Multimodal Models (LMMs) tailored for specific clinical workflows, such as radiation oncology, where models simultaneously handle report summarization, treatment planning, and target volume segmentation with robust generalization [434].

Beyond technical architecture, the successful deployment of AI in clinical settings hinges critically on adoption factors and the human-centered design of collaborative workflows. The potential benefits of AI do not automatically translate into adoption; rather, acceptance is influenced by a complex interplay of trust, workflow compatibility, and perceived utility. Research indicates specific tensions between technical capabilities and the practical requirements of clinicians, necessitating a focus on AI as a coequal team member rather than a mere tool [7]. A critical dimension of this interaction is the sequencing of AI inferences within the diagnostic workflow. Studies in radiology reveal that the timing of AI assistance significantly impacts decision-making dynamics; for instance, requiring radiologists to register provisional decisions before viewing AI inferences reduces anchoring effects and alters the likelihood of seeking second opinions, thereby preserving human agency and critical engagement [435]. Conversely, poorly integrated systems that position AI as a standalone decision-maker often fail to leverage radiologist expertise, leading to suboptimal outcomes. Innovative paradigms, such as “Copilot” systems, aim to rectify this by integrating Large Language Models with quantitative image analysis to empower radiologists to generate more precise reports while reducing burnout [436].

Translating AI from research laboratories to real-world clinical practice remains a significant barrier, often described as a “valley of death” in medical AI. Many certified products fail to reach the clinic due to a lack of understanding of domain-specific challenges, regulatory hurdles, and the

complexities of clinical validation [437]. To address these barriers, community-driven initiatives like the MONAI Consortium have emerged to build standards and infrastructure that facilitate deployment in healthcare institutions, bridging the gap between algorithmic development and clinical implementation [438]. These efforts emphasize the need for interoperability, standardized frameworks like DICOM gateways, and smart routing rules to manage the intricate workflows of modern healthcare [439]. Furthermore, the establishment of international consensus guidelines, such as the FUTURE-AI framework, provides a structured approach for developing trustworthy AI tools that adhere to principles of fairness, universality, traceability, and explainability, thereby fostering the trust required for widespread adoption [440].

Extending beyond immediate clinical decision support, human-AI teaming is accelerating biomedical discovery through autonomous agents and advanced evidence synthesis. The emergence of autonomous AI agents capable of coordinating specialized medical tools represents a paradigm shift in information processing. Systems utilizing Large Language Models as central reasoning engines can autonomously deploy tools for text analysis, image interpretation, and genomic data processing to generate patient-tailored recommendations in oncology, effectively functioning as specialist clinical assistants [441]. These agents are increasingly evaluated through rigorous frameworks akin to clinical examinations to ensure their reliability in real-world scenarios [442]. Additionally, AI plays a pivotal role in drug repurposing and evidence synthesis by analyzing vast corpora of medical literature and clinical data to identify hidden patterns and potential therapeutic interventions. Generative AI models, including transformers and diffusion models, are being applied to protein structure prediction, molecular representation, and drug design, significantly enhancing the speed and scope of biomedical research [443]. However, the reliability of these systems depends on robust validation mechanisms, such as virtual imaging trials, which use simulated and clinical data to uncover intrinsic factors influencing AI performance and ensure reproducibility before clinical deployment [444]. Ultimately, the future of biomedical discovery lies in symbiotic systems where human intuition and contextual knowledge are seamlessly integrated with the computational power of AI, creating a hybrid intelligence that surpasses the limitations of either entity acting alone—a trajectory that parallels the move toward “human-aware” and co-creative frameworks seen in materials science.

7.2 Materials Science and Chemical Discovery

7.2 Materials Science and Chemical Discovery

Bridging the high-stakes, multimodal clinical environments of Section 7.1 and the socio-technical workflows of software engineering in Section 7.3, materials science and chemical discovery represent a unique frontier where Human-AI collaboration shifts from decision support to synergistic co-creation. In this domain, the integration of artificial intelligence is not merely about accelerating existing workflows but fundamentally altering the epistemological approach to discovery. A central theme in recent literature is the concept of “human-aware” AI, which moves beyond substituting human experts to complementing their cognitive limitations. Traditional AI models, trained exclusively on published scientific findings, often replicate the biases inherent in the current body of literature, effectively converging on what is already known. However, innovative approaches have demonstrated that incorporating the distribution of human expertise into self-supervised models can dramatically improve the prediction of future discoveries. By training on inferences that are cognitively accessible to experts while simultaneously tuning the AI to avoid the “crowd” of conventional hypotheses, researchers have generated scientifically promising “alien” hypotheses [445]. These alien hypotheses represent pathways unlikely to be imagined or pursued by human scientists

alone due to collective cognitive biases, thereby serving to punctuate scientific advance rather than merely incrementally improving it [446]. This strategy of avoiding the scientific crowd allows AI to propose valuable energy-relevant materials and drug repurposing candidates with significantly higher precision than content-only models [447].

Parallel to the generation of alien hypotheses, the application of generative models for molecule design has evolved into a sophisticated co-creation process involving subject matter experts (SMEs), mirroring the iterative refinement seen in clinical workflows. Early generative models, while exhibiting extreme creativity, often produced molecules with low viability or synthetic feasibility. To address this, modern frameworks facilitate an iterative interaction where SMEs guide the exploration of the discovery space, asking for smaller samples and refining the search based on tacit domain knowledge [448]. These co-creation workbenches integrate generative modeling with dataset triage, molecule adjudication, and risk assessment, enabling a coherent discovery process that reduces the time to first discovery and minimizes opportunity costs [449]. The effectiveness of these systems relies heavily on the ability of AI to understand and manipulate the complex “grammar” of materials. For instance, neural language models like the BLMM Crystal Transformer have been developed to learn materials grammars, allowing for the generative and tinkering design of inorganic materials with high charge neutrality and balanced electronegativity [450]. Similarly, transformer-based models trained on expanded formulas from vast databases have demonstrated the capability to generate chemically valid compositions with enrichment rates significantly higher than random sampling, proving their utility in recovering leave-out materials and discovering novel compounds [451].

A critical challenge in deploying Large Language Models (LLMs) for chemical discovery, akin to the reliability concerns in biomedical AI, is the mitigation of hallucinations and the integration of rigorous physical constraints. While LLMs possess vast linguistic knowledge, they often lack the precise, long-term memory required for reproducible science. To overcome this, researchers have developed multimodal retrieval-augmented generation (RAG) frameworks, such as LLaMP, which dynamically interact with computational and experimental data stores like the Materials Project. These systems employ multiple data-aware reasoning-and-acting agents to fetch relevant data on the fly, process higher-order data such as crystal structures, and summarize synthesis procedures, effectively reducing hallucination rates in formation energies and volumetric strain to near zero [452]. Furthermore, specialized models like MatChat have been fine-tuned on structured material knowledge to predict inorganic material synthesis pathways, showcasing impressive reasoning capabilities despite the complexity of chemical processes [453]. The integration of LLMs with *ab initio* databases is also facilitating the creation of massive, high-quality datasets. For example, recent efforts have generated over 1.5 million materials narratives by merging multi-modality data sources, creating a resource that is more evenly distributed across the periodic table than literature-biased datasets, thus providing a richer foundation for AI exploration [454].

The synergy between human intuition and machine learning is further exemplified by efforts to “bottle” operational intuition into quantifiable descriptors, a approach that parallels the move toward trustable AI in healthcare. The “Materials Expert-Artificial Intelligence” (ME-AI) approach utilizes expertly curated measurement-based data to encapsulate human insight, revealing composite descriptors that not only reproduce expert intuition but expand upon it to identify critical chemical features like hypervalency [455]. Additionally, autonomous agents are being deployed to extract structured data from unstructured scientific literature, overcoming the bottlenecks of manual dataset creation. Agents like Eunomia leverage LLMs to plan and execute information extraction tasks across sentences and full articles, achieving performance comparable to state-of-the-art

fine-tuned methods while simplifying access for novice users [456]. Despite these advances, a sober assessment reveals that LLMs are not yet universally superior; their utility in principled Bayesian optimization over molecules is contingent upon pretraining or fine-tuning with domain-specific data, as generic models often fail to provide the necessary uncertainty estimates for efficient exploration [457]. Nevertheless, the trajectory is clear: the future of materials discovery lies in hybrid systems where human-aware AI generates alien hypotheses, generative models are guided by expert feedback, and LLMs are grounded in rigorous *ab initio* data to overcome cognitive biases and accelerate the realization of sustainable and high-performance materials. This evolution toward grounded, collaborative discovery sets the stage for examining how similar human-AI dynamics are reshaping the very infrastructure of research through software engineering and computational methodologies in the following section.

7.3 Software Engineering and Computational Research Methodologies

7.3 Software Engineering and Computational Research Methodologies

Bridging the gap between the material discovery paradigms of Section 7.2 and the high-stakes operational environments of Section 7.4, software engineering stands as a unique domain within the broader landscape of scientific research. Characterized by its inherently socio-technical nature, this field has witnessed the most rapid integration of generative artificial intelligence (GenAI) into daily workflows. Unlike materials science, where AI often serves as an analytical or generative tool for physical properties, in software engineering, AI has become an active participant in the creation process itself, co-authoring the very infrastructure upon which other scientific domains rely. A defining characteristic of this domain is the specific dynamic driving technology adoption. Contrary to traditional technology acceptance models which emphasize perceived usefulness or social influence, recent empirical evidence suggests that the primary driver for adopting GenAI tools in software engineering is their compatibility with existing development workflows [458]. This finding underscores a critical insight for human-AI collaboration across disciplines: for AI to transition from a novelty to a true teammate, it must seamlessly embed itself into established rhythms—whether coding, debugging, or system design—without imposing significant cognitive friction or requiring drastic procedural overhauls.

The impact of generative AI on developer productivity and code quality has been a subject of intense investigation, yielding nuanced results that challenge both utopian and dystopian narratives. Large-scale empirical studies within corporate environments, such as those conducted at major financial institutions, have demonstrated that tools like GitHub Copilot can lead to notable boosts in productivity and code quality when integrated into the software development life cycle [459]. Developers report increased job satisfaction and efficiency, particularly when handling routine coding tasks or exploring unfamiliar libraries. However, the relationship between AI assistance and output quality is not linear. Controlled experiments indicate that while large language models (LLMs) excel at solving isolated, simple coding problems, their performance degrades when tasked with complex, typical software development challenges that require holistic system understanding [460]. Furthermore, observational studies reveal that professional engineers often utilize these tools not merely for generating ready-to-use artifacts, but increasingly for receiving guidance, learning abstract concepts, and brainstorming solutions, effectively shifting the engineer’s role from pure authorship to critical integration and verification [461]. This evolution mirrors the “human-aware” co-creation seen in materials science, where the human expert guides the AI’s exploration rather than passively accepting its output.

Beyond direct code generation, a transformative frontier in computational research methodologies involves the potential use of LLMs as synthetic substitutes for human subjects. Traditional qualitative research in software engineering faces significant hurdles, including the labor intensity of data collection, difficulties in participant recruitment, and limitations in scaling studies. To address these challenges, researchers have proposed leveraging LLMs to generate synthetic text that mimics human responses, thereby enabling persona-based interviews, multi-persona focus groups, and large-scale survey simulations [462]. This approach promises to revolutionize data collection by offering scalable and efficient means to generate insights into human attitudes and experiences. Supporting this vision, multi-agent frameworks have been successfully employed to automate qualitative data analysis, where distinct agents execute specific research activities such as coding interview transcripts and identifying thematic patterns, significantly accelerating the analysis process [463].

However, the substitution of human subjects with “silicon participants” introduces profound methodological and ethical complexities centered on the concept of algorithmic fidelity. Algorithmic fidelity refers to the degree to which LLM-generated outputs accurately mirror the beliefs, attitudes, and lived experiences of real human sub-populations. Rigorous framework-based qualitative analyses have revealed that while LLMs can reproduce key themes found in human data with striking similarity, they often fail to capture the nuanced structure, tone, and contextual depth of human interaction [117]. Studies comparing human and LLM reasoning in tasks such as app review classification show that while there is partial alignment, LLMs rely heavily on surface-level word choices and functional components rather than the deep experiential context that human analysts utilize [464]. Consequently, current evidence suggests that LLMs do not yet possess sufficient algorithmic fidelity to fully replace human subjects in research aiming to generalize findings to real populations. The risk of “hyper-accuracy distortion” and the inability to replicate the empathetic understanding inherent in human subjects necessitate an integrated approach where AI-generated data coexists with, rather than replaces, genuine human input [462][117].

Furthermore, the reliance on LLMs for research and development introduces new forms of bias and uncertainty that resonate with the safety concerns of high-stakes environments discussed in the following section. The training data of general-purpose models often reflects dominant cultural perspectives, potentially leading to a homogenization of software design principles and research outcomes that overlook diverse global contexts [465]. There is also a growing concern regarding “context-induced algorithmic bias,” where the specific phrasing of prompts and the implicit variations in human-AI interaction introduce irreproducibility into scientific workflows [125]. Therefore, while generative AI offers unprecedented opportunities to scale qualitative analysis and enhance developer productivity, the field must evolve robust validation mechanisms. Future methodologies must prioritize human-in-the-loop processes to ensure that the efficiency gains provided by synthetic data and automated coding assistants do not come at the expense of scientific rigor, diversity of thought, and the authentic representation of human experience in software engineering research. As software systems increasingly underpin critical infrastructure and defense capabilities, establishing these rigorous standards for human-AI collaboration in their creation becomes a prerequisite for the trust and reliability required in the operational domains explored next.

7.4 Defense, Cybersecurity, and High-Stakes Operational Environments

7.4 Defense, Cybersecurity, and High-Stakes Operational Environments

Bridging the methodological rigor of software engineering discussed in Section 7.3 and the expansive interdisciplinary challenges of global development in Section 7.5, the domains of defense, cybersecurity, and high-stakes operational environments represent the most critical frontiers for Human-AI collaboration. In these sectors, the consequences of failure extend far beyond financial loss or reduced productivity to include threats to national security, critical infrastructure integrity, and human life. Consequently, the paradigm of full automation—often pursued in less critical domains—is deemed insufficient or undesirable. Instead, effective human-machine teaming (HMT) has emerged as the necessary standard, leveraging the computational speed of AI alongside human contextual intuition, ethical judgment, and accountability. The realization of this synergy, however, demands sophisticated experimentation platforms, robust anomaly detection frameworks, and adaptive mechanisms capable of handling dynamic workload distributions under profound uncertainty.

A foundational requirement for advancing HMT in defense contexts is the development of lightweight, flexible experimentation platforms that can simulate complex, multi-agent interactions with high fidelity. Traditional simulation environments often lack the nuance required to model the heterogeneity of modern battlefields or the intricate decision-making processes of human operators. To address this, researchers have introduced platforms such as Cogment, which facilitates the interaction between humans and multiple AI agents, including both static and learning-based entities. Such platforms are essential for validating algorithms in defense environments where full automation is not achievable, allowing for the rigorous testing of human-machine teaming use cases involving heterogeneous systems [466]. Furthermore, the complexity of future warfare necessitates Synthetic Task Environments (STE) that reflect the “teams of teams” structure found in Joint All-Domain Command and Control (JADC2) concepts. Insights from subject-matter experts indicate that effective research tools must support matrixed organizations and robust communication modalities, enabling the study of how teams process vast amounts of sensor data to make timely, complex decisions [467]. These environments serve as the proving grounds for scalable interactive machine learning, aiming to foster resilient teams by optimizing roles, configurations, and trust dynamics in rapidly changing operational scenarios [6].

In the realm of cybersecurity, the collaboration between human experts and AI systems is pivotal for managing an evolving landscape of digital threats characterized by increasing speed and sophistication. AI offers unparalleled capabilities in pattern recognition and predictive modeling, allowing for the proactive identification of vulnerabilities and the rapid detection of anomalies that might elude human analysts. However, the deployment of AI in cybersecurity is fraught with risks related to system opacity and potential adversarial exploitation. A balanced approach is therefore advocated, where AI’s computational power is integrated with human expertise in critical thinking and contextual understanding. This human-in-the-loop framework ensures that AI-generated decisions can be explained to stakeholders and regulators, thereby maintaining accountability and establishing trust in security solutions [468]. The threat landscape is further complicated by the susceptibility of opaque AI systems to adversarial attacks such as data poisoning. To mitigate these risks, proactive security measures like “violet teaming” have been proposed, combining adversarial red teaming with defensive blue teaming to holistically evaluate and secure AI systems against dual-use potentials and unintended behaviors [469]. Specific applications, such as securing maritime autonomous systems, demonstrate the efficacy of red teaming frameworks in uncover-

ing vulnerabilities ranging from adversarial patch attacks to system poisoning, highlighting the necessity of secure-by-design principles in mission-critical AI [470].

Beyond cyber domains, physical high-stakes environments, particularly those involving Unmanned Aerial Vehicles (UAVs) and swarms, present unique challenges regarding workload estimation and trust calibration. In dynamic environments, human operators must manage multiple autonomous agents while coping with fluctuating cognitive loads. A significant challenge arises when AI systems attempt to estimate human workload under distribution shifts—situations where the tasks encountered differ significantly from the training data. Standard machine learning approaches, which often assume independent and identically distributed (IID) samples, struggle to generalize to these unknown tasks, necessitating the adoption of non-IID techniques that prioritize portability and adaptability [471]. The stakes are heightened in emergency response scenarios where task transitions must occur instantly. In such contexts, incorrect trust calibration can lead to catastrophic failures; for instance, a human supervisor might misattribute the faulty behavior of a single robot to the entire swarm, leading to unnecessary intervention and degraded performance. To counteract this, dynamic trust repair models have been developed to help UAV swarms self-reflect and correct faulty behaviors proactively, thereby restoring human trust and maintaining operational efficiency during critical missions [472]. Similarly, trust-aware reflective control mechanisms enable swarms to adjust their behaviors based on human trust models, ensuring resilient cooperation even when facing accumulated uncertainties from prior tasks [473].

The integration of AI into these high-stakes workflows also demands advanced methods for anomaly detection and situational awareness that account for both technical and human factors. In networked UAV formations, stealthy adversaries may launch zero-dynamics or covert attacks that are difficult to detect using conventional methods. Model-based observer techniques, both centralized and decentralized, have been developed to leverage switching communication topologies and local measurements to identify such intrusions effectively [474]. Moreover, the design of these systems must account for the psychological and physiological states of human operators. Frameworks that monitor cognitive states through multisourced physiological and behavioral data enable autonomous systems to adapt their level of transparency or autonomy, optimizing the human-machine interaction based on real-time workload dynamics [475]. As the complexity of these socio-technical systems grows, the focus shifts toward creating Cyber-Physical-Social Systems (CPSS) that formally encode the tacit knowledge and procedural constraints of human experts, ensuring that automated workflows remain aligned with operational realities and safety norms [476]. Ultimately, the success of Human-AI collaboration in defense and cybersecurity hinges on the ability to create adaptive, transparent, and trustworthy systems that augment rather than replace human judgment in the face of uncertainty and adversarial pressure. These rigorous standards for safety and reliability provide a crucial backdrop for understanding the structural and epistemic challenges inherent in broader interdisciplinary and global development efforts, where the stakes involve not only immediate operational failure but also long-term societal equity and impact.

7.5 Interdisciplinary Synthesis and Global Development

7.5 Interdisciplinary Synthesis and Global Development

While the preceding section highlighted the critical necessity of robust Human-AI collaboration in high-stakes, domain-specific operational environments like defense and cybersecurity, this subsection expands the scope to examine how such collaboration functions across interdisciplinary boundaries and global development contexts. The integration of artificial intelligence into scientific

workflows is frequently championed as a catalyst for interdisciplinary convergence, particularly in addressing complex challenges in global health and development. However, a critical synthesis of recent cross-disciplinary collaborations reveals a nuanced reality: while the volume of interdisciplinary activity has surged, scientific impact is not automatically guaranteed by the mere presence of diverse teams or technologies. Instead, evidence suggests that the value derived from human-AI collaboration in global contexts depends heavily on the depth of knowledge diversity harnessed and the structural integrity of the collaboration, rather than on simple interdisciplinarity metrics. The COVID-19 pandemic served as a massive, real-world stress test for these collaborative paradigms, offering profound lessons on the disconnect between computational potential and actionable societal outcomes.

During the early stages of the pandemic, there was an unprecedented mobilization of AI specialists and medical professionals, driven by the urgent need to understand and combat the virus. Despite this surge in collaborative effort, large-scale scientometric analyses indicate that many of these partnerships resulted in science of relatively low visibility and impact. Research analyzing over 10,000 papers at the intersection of COVID-19 and AI demonstrates that scientific impact was not determined by the overall interdisciplinarity of author teams, but rather by the specific diversity of knowledge they successfully integrated into their research frameworks [477]. This finding challenges the assumption that bringing together experts from different fields automatically yields superior results; instead, it highlights the necessity of deep knowledge synthesis where computational methods are meaningfully aligned with domain-specific clinical and epidemiological realities. Without this alignment, interdisciplinary efforts risk becoming superficial, producing outputs that fail to address the core complexities of the crisis.

The challenge of synthesizing knowledge across domains is further compounded by the fragmentation of information silos. Early literature and knowledge platforms often focused narrowly on biology and medicine, neglecting the broader interdisciplinary efforts required to address the socio-economic dimensions of the pandemic. To bridge this gap, initiatives such as the development of interdisciplinary academic knowledge graphs have been proposed to facilitate cross-domain knowledge extraction and integration [478]. These systems aim to connect disparate fields, enabling researchers to identify research communities and predict potential links that would remain invisible within single-discipline frameworks. Similarly, automated theme-based visualization methods have been deployed to navigate large corpora of publications, revealing evolving research trajectories and cross-domain initiatives in areas like mental health and education [479]. These tools underscore the role of AI not just as a generator of hypotheses, but as a facilitator of interoperability between distinct scientific cultures, mirroring the “teams of teams” structures discussed in high-stakes command and control environments.

In the context of global development, the application of AI has shifted towards accelerating evidence synthesis to inform policy decisions. The sheer volume of emerging literature makes manual review impossible, necessitating human-AI hybrid workflows that leverage the complementary strengths of both parties. For instance, the deployment of AI agents based on transformer models, enhanced by active learning strategies, has proven highly effective in reducing human screening efforts for evidence gap maps by significant margins [480]. By integrating these agents into human teams, organizations can accelerate the design of evidence-based policies and programs, ensuring that decision-makers can distill key information from vast datasets in a timely manner. This approach exemplifies a successful model of interdisciplinary synthesis where AI handles the scale of data processing while human experts provide the contextual judgment necessary for effective global development strategies, echoing the capability asymmetry principles found in biomedical discovery.

However, the path to effective interdisciplinary synthesis is fraught with structural and epistemic barriers similar to those observed in specialized domains. The dominance of certain disciplines can skew the research agenda, leading to an uneven distribution of attention where fields like medicine and immunology overshadow others, despite the complex nature of the challenges [481]. Furthermore, the “convergence shortcut,” where researchers attempt to integrate distant domains without deep cross-disciplinary collaboration, has been shown to be less impactful than genuine team-based integration [482]. This phenomenon is particularly relevant in global health, where solutions require not only technical efficacy but also cultural and social appropriateness. The lack of diverse data sources and the prevalence of algorithmic bias pose significant risks to equitable outcomes, emphasizing the need for trustworthy AI designs that intentionally mitigate bias from the research design stage [366].

Moreover, the geographic and institutional diversity of teams plays a critical role in the originality and long-term impact of scientific outputs. Studies indicate that expertise-diverse teams tend to produce work with greater originality and higher long-term citation impact, yet these benefits are often contingent on the balance between exploratory and exploitative cognitive profiles within the team [483] [484]. In the Global South, where compute infrastructure and ML competencies may be limited, cross-border collaborative capacity building has emerged as a vital strategy to ensure that local contexts are adequately represented in AI-driven solutions [485]. Without such intentional efforts to diversify the knowledge base and the teams generating it, AI applications in global development risk reinforcing existing inequalities rather than alleviating them, highlighting a trust and adoption barrier distinct from but related to the safety-critical concerns of defense and cybersecurity.

Ultimately, the lesson from recent cross-disciplinary endeavors is that knowledge diversity, rather than mere interdisciplinarity, is the primary driver of impact. Successful human-AI collaboration in global development requires moving beyond the superficial labeling of projects as “interdisciplinary” to fostering environments where diverse epistemic traditions can genuinely intersect. This involves developing robust infrastructure for data sharing, such as federated learning frameworks that respect privacy while enabling multi-site collaboration [486], and creating governance models that prioritize equity and inclusivity. As the scientific community looks toward future pandemics and global challenges, the focus must shift from simply deploying AI tools to cultivating the cognitive and structural conditions necessary for true symbiotic discovery. These insights into the variability of collaboration success factors across global and interdisciplinary contexts set the stage for a broader comparative analysis of domain-specific constraints and universal principles of human-AI teaming.

7.6 Comparative Analysis of Domain-Specific Constraints and Success Factors

7.6 Comparative Analysis of Domain-Specific Constraints and Success Factors

Building upon the insights from interdisciplinary synthesis and global development, where the mere presence of diverse teams proved insufficient without deep knowledge integration, this subsection synthesizes empirical evidence across biomedicine, materials science, software engineering, and high-stakes operational environments. This comparative analysis reveals a complex landscape where universal principles of human-AI collaboration intersect with deeply entrenched domain-specific constraints. While the overarching goal of achieving Complementary Team Performance (CTP)—

where the joint output exceeds the capabilities of either the human or the AI acting alone—remains constant, the pathways to realizing this potential vary significantly based on the nature of the data, the cost of error, and existing workflow infrastructures [12]. Specifically, while information asymmetry and capability asymmetry serve as foundational sources for complementarity across all fields, the mechanisms for leveraging these asymmetries are contingent upon distinct domain-specific adoption barriers and trust calibration requirements.

A primary universal principle emerging from these diverse fields is the critical role of appropriate trust calibration in mitigating both the misuse of overconfident AI and the disuse of unconfident systems. However, the manifestation of trust barriers differs markedly by sector, echoing the disconnect between computational potential and actionable outcomes observed in global health contexts. In clinical decision-making and defense, where errors carry immediate life-safety consequences, trust is heavily mediated by perceived risk and the transparency of the AI’s decision-making process. Studies in healthcare indicate that trust is not merely a function of accuracy but is significantly influenced by socio-ethical considerations and the alignment of AI outputs with human contextual knowledge [7]. Conversely, in software engineering and computational research, adoption is often driven more by workflow compatibility and the immediate utility of the tool in reducing cognitive load than by deep philosophical trust. Here, the barrier is less about safety and more about the friction of integrating AI suggestions into existing coding or analytical pipelines [487]. Despite these differences, the necessity for dynamic trust calibration remains a shared imperative; uncalibrated confidence scores hinder collaboration universally, leading to suboptimal reliance patterns regardless of the domain [415].

The sources of complementarity also exhibit distinct domain profiles, reflecting the varying degrees of knowledge diversity required for impact. In biomedical discovery and materials science, the most potent form of collaboration arises from capability asymmetry, where AI models process vast, high-dimensional parameter spaces that exceed human cognitive limits, while humans provide the necessary grounding in physical laws and experimental feasibility. This is evident in “self-driving laboratory” frameworks where AI proposes “alien” hypotheses that humans might dismiss due to cognitive bias, yet which prove fruitful upon experimental validation. In contrast, in social sciences and qualitative humanities research, complementarity often stems from information asymmetry. Here, humans possess unique contextual and cultural insights that AI lacks, while AI offers the capacity to scale qualitative coding and detect semantic shifts across massive corpora. Successful collaboration in these fields requires frameworks that explicitly model human learning as a mediator of appropriate reliance, ensuring that domain experts do not passively defer to algorithmic outputs but actively integrate their unique contextual knowledge [37].

Furthermore, the constraints on deployment reveal a tension between the need for explainability and the practical limits of interpretability, a challenge parallel to the structural barriers noted in global development. In high-stakes domains like cybersecurity and air-traffic control, the demand for explainable AI (XAI) is paramount, yet research suggests that full interpretability is neither necessary nor sufficient for trust; rather, the ability to verify model behavior through rigorous testing and “behavior certificates” is often more critical [488]. This contrasts with educational or consumer-facing applications, where stakeholders such as teachers and parents prioritize transparency regarding data privacy and agency over technical robustness, indicating that trust metrics must be tailored to the specific values of the stakeholder group [489]. The challenge lies in designing interaction patterns that accommodate these divergent needs without fragmenting the core architecture of human-AI teaming.

Another critical success factor identified across domains is the shift from static, one-off interactions

to dynamic, co-learning relationships. Effective human-AI teams are characterized by mutual theory of mind, where both parties adapt to each other’s evolving capabilities. In fields like materials science, this involves iterative loops where human feedback refines the AI’s generative models, while in clinical settings, it involves continuous calibration of the AI’s confidence based on physician overrides. However, the success of these co-learning loops is often hindered by organizational and cultural barriers similar to those impeding interdisciplinary synthesis. For instance, in natural science research, low adoption rates are frequently attributed not to technical deficiencies but to a lack of support during initial implementation and a failure to communicate system capabilities in user-relevant terms [487]. This suggests that successful deployment requires not only technical solutions but also robust change management strategies that address the “relational problem” in socio-technical discourse, bridging the gap between operational, epistemic, and normative domains [490].

Ultimately, while the specific constraints of data modality, regulatory environment, and risk tolerance dictate the tactical implementation of AI tools, the strategic prerequisites for success remain consistent. Effective collaboration demands a clear delineation of tasks based on delegability frameworks that account for motivation, difficulty, and risk [491]. It requires moving beyond the pursuit of full automation toward “machine-in-the-loop” designs that preserve human agency and leverage the distinct strengths of both partners. Moreover, it necessitates the development of standardized metrics for evaluating hybrid performance that go beyond simple accuracy to include measures of novelty, robustness, and the quality of the human-AI interaction itself. By recognizing both the universal need for calibrated trust and complementarity, and the domain-specific nuances of their application, the scientific community can foster ecosystems where AI serves not merely as a tool, but as a true teammate in the pursuit of discovery, thereby addressing the structural and epistemic challenges highlighted in both specialized and global contexts.

8 Future Directions: Towards Autonomous Agents, Evaluation Standards, and Responsible Ecosystems

8.1 Standardized Benchmarks and Metrics for Hybrid Human-AI Scientific Output

8.1 Establishing Metrics and Benchmarks for Hybrid Scientific Intelligence

As the scientific enterprise increasingly integrates artificial intelligence into its core workflows, the metrics used to evaluate success must evolve beyond traditional measures of predictive accuracy or computational speed. This shift is particularly critical as the field transitions from the static, single-purpose tools discussed previously toward the dynamic, autonomous multi-agent systems outlined in the following section. The emergence of AI as an active collaborator necessitates a new paradigm for assessing “hybrid intelligence,” where the value of the system lies not merely in the machine’s standalone performance but in the synergistic output of the human-AI team. Current evaluation frameworks often fail to capture the nuanced dimensions of this collaboration, such as the novelty of generated hypotheses, the interpretability of the decision-making process, and the specific complementarity achieved between human intuition and algorithmic processing. To address this gap, future research must establish standardized benchmarks and metrics that holistically quantify the effectiveness of hybrid scientific outputs, moving away from siloed assessments toward integrated evaluation suites that reflect the complexity of modern discovery.

A primary limitation of existing benchmarks is their reliance on static datasets and singular accu-

racy metrics, which do not account for the dynamic nature of scientific inquiry or the distributional shifts inherent in real-world research environments. As noted in recent analyses, standard benchmarks often fail to reflect the true capabilities of AI systems when applied to complex, evolving problem spaces, leading to a disconnect between benchmark performance and practical utility [492]. Furthermore, the definition of intelligence itself requires re-evaluation; measuring skill at specific tasks often masks a system’s generalization power and fails to capture the efficiency of skill acquisition, which is critical for adaptive scientific partners [493]. Consequently, a robust framework for hybrid evaluation must incorporate metrics that assess how well human and AI agents complement one another to achieve “Complementary Team Performance” (CTP), a state where the team outperforms either entity acting alone. Theoretical foundations suggest that CTP arises from leveraging information asymmetry and capability asymmetry, yet few current metrics explicitly measure the realization of this potential in scientific workflows [12].

To operationalize these concepts, the scientific community should adopt multi-dimensional benchmarking approaches similar to the “Science4Cast” initiative, which utilizes evolving semantic networks to predict future research directions rather than simply classifying existing data. Such dynamic benchmarks are essential for evaluating an AI’s ability to suggest novel, personalized research directions that accelerate scientific progress [494]. By shifting the focus from retrospective classification to prospective prediction, these frameworks can better measure the novelty and forward-looking value of human-AI collaborations. Additionally, the evaluation of hybrid systems must include rigorous assessments of uncertainty. A reliable framework for quantifying reproducibility and uncertainty in machine learning workflows is essential for establishing trust in scientific results, particularly when AI predictions influence high-stakes experimental designs [495]. Integrating uncertainty quantification with explainability metrics can further enhance the “trustability” of these systems, ensuring that human researchers can appropriately calibrate their reliance on AI suggestions [496].

Moreover, the development of standardized metrics must be grounded in a comprehensive catalogue of responsible AI principles that extend beyond technical performance to include accountability, fairness, and transparency. Recent efforts to create a “Responsible AI Metrics Catalogue” propose a tripartite framework comprising process, resource, and product metrics, which offers a structured approach to operationalizing accountability in generative AI systems [198]. This cataloguing approach is vital for ensuring that hybrid scientific outputs adhere to ethical standards and that the contributions of both human and machine agents are transparently documented. The concept of the “Harmony Index” further illustrates the need for utilitarian metrics that classify team members based on their contribution to overall effectiveness, providing a quantitative measure of how well mixed-skill teams function in practice [497]. Such metrics are crucial for identifying whether an AI agent is acting as a true teammate or merely a redundant tool, a distinction that becomes increasingly important as systems approach higher levels of autonomy.

Critically, the evaluation of hybrid intelligence must also address the human cognitive element, recognizing that the effectiveness of a team depends on the alignment of mental models and the quality of interaction. Metrics that assess “explanatory efficacy”—the strength of the cyclic relationship between the human user and the AI interface—are necessary to determine if explanations actually improve human understanding and decision-making rather than simply increasing blind reliance [339]. Frameworks like OpenHEXAI demonstrate the feasibility of large-scale, human-centered benchmarks that evaluate explanation methods based on their impact on user trust, fairness, and decision accuracy, setting a precedent for future standardized testing in scientific domains [498]. Furthermore, as AI systems are updated, it is imperative to measure “backward compatibility” to

ensure that model improvements do not disrupt the user’s established mental models or degrade team performance due to unexpected behavioral shifts [499].

Finally, the pursuit of standardized benchmarks must avoid the pitfall of “metric gaming” and ensure that evaluations reflect genuine scientific value rather than optimized scores on narrow tasks. The application of metrology—the science of measurement—to AI evaluation suggests that all performance claims must be reported with respect to the precision and resolution of the evaluation instrument, including the variability introduced by human annotators [500]. This rigorous approach, combined with diverse evaluation datasets that capture variance in human judgment, such as the DICES dataset for safety, will help create benchmarks that are robust, inclusive, and representative of the pluralistic values inherent in the global scientific community [501]. By synthesizing these diverse strands of research—dynamic prediction, uncertainty awareness, responsible accountability, human-centered interaction, and metrological rigor—the field can establish a unified framework for evaluating hybrid human-AI scientific output. This framework will not only guide the development of more effective collaborative tools but also provide the necessary validation mechanisms to support the transition toward the autonomous, self-correcting discovery systems that define the next frontier of scientific research.

8.2 The Evolution Towards Autonomous Multi-Agent Discovery Systems

8.2 The Evolution Towards Autonomous Multi-Agent Discovery Systems

Building upon the rigorous evaluation frameworks and hybrid intelligence metrics established in the preceding section, the scientific enterprise is now poised to transition from static, single-purpose tools toward dynamic, autonomous multi-agent systems. While Section 8.1 emphasized the necessity of quantifying “Complementary Team Performance” and trust in human-AI dyads, the trajectory of artificial intelligence in research is rapidly advancing toward “Level 5” autonomy—a state defined by the capacity for end-to-end scientific discovery with minimal human intervention, analogous to the highest levels of autonomy in self-driving vehicles [502]. Achieving this paradigm requires not merely incremental improvements in algorithmic accuracy, which traditional benchmarks often measure, but a fundamental architectural reorganization toward systems that can independently generate hypotheses, design verification plans, execute experiments, and validate results in a continuous, self-correcting loop. This shift represents the operational realization of the hybrid intelligence potentials discussed previously, moving from measuring collaboration to engineering autonomous agents that embody it.

The cornerstone of this evolution is the transition from monolithic models to dynamic, multi-agent architectures that emulate the collaborative nature of human scientific teams. Although Large Language Models (LLMs) have demonstrated the nascent capacity to act as zero-shot hypothesis proposers, synthesizing interdisciplinary knowledge to suggest novel, testable scientific hypotheses even without explicit training on specific target domains [67], relying on a single model often limits complex reasoning and execution reliability. To overcome these barriers and achieve the robust “Complementary Team Performance” outlined earlier, the field is adopting multi-agent frameworks where distinct agents assume specialized roles—such as the critic, the experimenter, and the synthesizer—to iteratively refine research problems and experimental designs. These collaborative architectures allow for the decomposition of grand challenges into manageable sub-tasks, enabling a level of cognitive synergy that surpasses the capabilities of individual agents or static

tools [503]. In this context, the metrics of accountability and explanatory efficacy introduced in Section 8.1 become critical internal mechanisms for agent-to-agent validation, ensuring that the swarm’s collective output remains scientifically sound.

A critical technical pillar supporting this shift is the development of closed-loop systems that integrate hypothesis generation with physical or simulated experimentation, effectively bridging the gap between digital prediction and empirical verification. The concept of the “self-driving laboratory” exemplifies this integration, where AI agents govern robotic instruments to explore vast parameter spaces autonomously. For instance, systems like the Scientific Autonomous Reasoning Agent (SARA) have successfully demonstrated the ability to accelerate the synthesis of metastable materials by employing nested active learning cycles that incorporate underlying physics and uncertainty quantification [208]. Similarly, in the realm of microscopy, autonomous experimentation platforms are being developed to convert low-latency decision-making processes into high-order experiment planning, effectively closing the loop between image acquisition and hypothesis refinement without excluding human operators but rather elevating their role to high-level oversight [504]. These systems represent a move away from simple automation toward genuine scientific agency, where the AI determines the next most informative experiment based on real-time data feedback, directly addressing the need for dynamic, prospective evaluation highlighted in the previous section.

Furthermore, the realization of Level 5 autonomy necessitates the advent of AI-generating algorithms (AI-GAs), an alternate paradigm where the system automatically learns how to produce general AI rather than relying on manual engineering of every component [505]. This meta-learning approach is essential for creating systems that can adapt to open-world environments and novel circumstances that were not anticipated during their initial design, thereby overcoming the limitations of static datasets and singular accuracy metrics criticized in Section 8.1. Frameworks such as Self-Initiated Open World Learning (SOLA) propose mechanisms for agents to continually learn and adapt on their own initiative, detecting novelties and gathering ground-truth data to improve their performance over time [506]. By embedding self-experimental and self-repairing capabilities, these systems can maintain operational autonomy even in underspecified circumstances, a prerequisite for navigating the unpredictable landscape of fundamental scientific discovery [507].

Despite these advancements, significant challenges remain in achieving flawless autonomous verification, underscoring the continued relevance of the human-centric metrics discussed previously. Early experiments utilizing generic LLMs for autonomous hypothesis verification have shown promise but also revealed that current systems often produce flawed verification code or reasoning when limited to minimal guidance [508]. This underscores the necessity for robust human-in-the-loop validation mechanisms and the development of fail-proof tool integration to ensure reproducibility and scientific integrity [509]. Moreover, the complexity of coordinating multiple agents requires sophisticated orchestration strategies to balance autonomy with alignment, ensuring that the collective behavior of the agent swarm remains directed toward valid scientific objectives [510]. The “Harmony Index” and responsible AI catalogues mentioned earlier provide the theoretical grounding necessary to assess whether these autonomous swarms are functioning as cohesive scientific entities or merely generating redundant noise.

Looking forward, the evolution toward autonomous multi-agent discovery systems will likely be characterized by the integration of diverse AI modalities, including symbolic reasoning, deep learning, and robotic control, within unified frameworks like Simulation Intelligence (SI) [511]. The goal is to create ecosystems where AI agents not only accelerate existing workflows but also punctuate scientific advance by generating “alien” hypotheses—ideas that lie outside the cognitive blind spots of human researchers [445]. As these systems mature, they promise to democratize the scientific

process, allowing for the systematic production of data and the exploration of hypothesis spaces at a scale and speed previously unattainable, ultimately striving to meet the Nobel Turing Grand Challenge of developing AI scientists capable of making discoveries superior to the best human minds [502]. However, as these autonomous agents begin to generate knowledge at unprecedented velocities, the traditional mechanisms for communicating and archiving this output—static PDFs and fixed articles—will become obsolete. This impending bottleneck in knowledge dissemination necessitates a parallel transformation in scholarly communication, leading directly to the “late-binding” paradigms explored in the following section, where the dynamic nature of autonomous discovery is matched by equally dynamic, executable, and personalized scientific artifacts.

8.3 Dynamic and Late-Binding Scholarship Paradigms

8.3 Dynamic and Late-Binding Scholarship Paradigms

While the evolution toward autonomous multi-agent discovery systems described in the previous section promises to revolutionize how scientific knowledge is generated, a parallel transformation is required in how that knowledge is communicated and consumed. The traditional architecture of scientific communication has remained largely static for centuries, rooted in a paradigm of “early-binding” where knowledge is fixed into a final, immutable form—typically a PDF or printed article—before it leaves the author’s control. In this conventional model, the scholarly artifact is merely a snapshot of understanding at a specific moment in time, disconnected from the dynamic flow of data and the rapidly evolving landscape of subsequent discoveries. However, the convergence of artificial intelligence, large language models (LLMs), and advanced computational infrastructures is catalyzing a fundamental shift toward “late-binding” scholarship. This emerging paradigm envisions scientific knowledge not as a static product but as a dynamic, executable process where text, data, and algorithms are rendered in real-time, allowing for continuous evolution, personalization, and contextual adaptation of scientific insights.

The concept of late-binding scholarship posits that the core components of a research contribution should be decoupled from their presentation layer. Instead of delivering a finalized narrative, authors would provide a modular assembly of core arguments, raw datasets, algorithmic processes, and references to specific AI models capable of acting as “renderers” [512]. In this framework, the scholarly work exists as a living entity that can be re-synthesized dynamically at the moment of reading. This approach directly addresses the challenges posed by information overload and the rapid obsolescence of static texts, leveraging the capability of AI to integrate new findings and alternative theories into existing works without requiring manual revision by the original authors [512]. The transition from static documents to dynamic artifacts is further supported by the exponential growth of digitized scientific discourse, including preprints, code, and datasets, which creates a rich substrate for computational models to analyze and harness [513].

A critical enabler of this paradigm is the advancement of generative AI and LLMs, which possess the semantic understanding necessary to reconstruct and adapt scientific narratives. These models can serve as the engine for rendering canonical versions of texts while simultaneously offering personalized interpretations tailored to diverse audiences, ranging from domain experts to students or policy-makers [512]. For instance, an AI renderer could adjust the complexity of an explanation, highlight specific methodological details relevant to a reader’s field, or even translate core findings into different languages while maintaining fidelity to the original data and logic. This capability transforms the reading experience from a passive consumption of fixed information into an interactive engagement with evolving knowledge, potentially overcoming the limitations of current

PDF-based formats which suffer from static content and poor accessibility [514].

Furthermore, the late-binding model facilitates a more robust integration of reproducibility and transparency into the fabric of scientific publishing, a necessity amplified by the complexity of autonomous discovery systems. By encapsulating the algorithms and parameters necessary for analysis alongside the textual arguments, dynamic scholarship ensures that the pathway from data to conclusion remains traceable and executable. This aligns with broader efforts to create platform-agnostic and structured dissemination methods for deep learning models and scientific workflows, such as container-based engines that allow users to interact with models out-of-the-box [515]. In a late-binding ecosystem, the distinction between the paper and the code dissolves; the “paper” becomes an interface to a reproducible computational experiment that can be re-run, modified, or extended by the reader in real-time. This addresses the critical need for reproducibility in an era where complex AI-driven discoveries often rely on opaque or inaccessible computational pipelines.

The shift to dynamic scholarship also necessitates a re-evaluation of how scientific impact and attention are measured. Traditional bibliometrics rely on citation counts of static entities, but in a late-binding world, the “usage” of a scholarly artifact becomes a continuous stream of interactions. Research indicates that the dynamics of collective attention on science are already shifting, with preprint servers accelerating knowledge transfer and altering citation growth patterns [516]. Dynamic artifacts could provide granular data on how specific arguments or data points are utilized, revised, or contested over time, offering a more nuanced view of scientific evolution than static citation networks. Moreover, as language and scientific concepts evolve, dynamic systems can employ techniques like sequential alignment of text representations to manage data drift and ensure that older knowledge remains interpretable and relevant in changing semantic contexts [517].

However, realizing this vision requires overcoming significant technical and normative challenges that bridge the gap between technological capability and ethical governance. The legal and ethical implications of attributing authorship to dynamically generated content, where the final text may vary for every reader, are complex and unresolved [512]. Additionally, the reliance on AI renderers introduces risks related to hallucination and the potential for subtle distortions of the original scientific intent. Ensuring the integrity of dynamic scholarship will require robust verification mechanisms, perhaps leveraging the multi-agent frameworks discussed earlier or specialized auditing pipelines to validate that the rendered output remains faithful to the underlying data and logic. These challenges underscore the urgent need for the unified global governance frameworks and consensus guidelines detailed in the following section. As research workflows evolve toward “late-binding” scholarship where content is dynamically rendered, static policies will become obsolete, necessitating adaptive governance that incorporates metrics for process, resource, and product quality to ensure continuous alignment with ethical standards. Despite these hurdles, the trajectory of scientific communication points irresistibly toward this fluid, interactive future. As AI continues to permeate every stage of the research workflow, from hypothesis generation to peer review, the final output of science must evolve to match the dynamism of its creation process. The adoption of late-binding paradigms promises to unlock a new level of collaborative potential, where scientific knowledge is not merely archived but actively lived, updated, and personalized, driving a more efficient and inclusive acceleration of human discovery [513].

8.4 Global Governance Frameworks and Consensus Guidelines for Research Integrity

8.4 Global Governance Frameworks and Consensus Guidelines for Research Integrity

The transition to “late-binding” scholarship described in the previous section, where scientific artifacts become dynamic and executable processes, fundamentally disrupts the static regulatory models that have governed research integrity for centuries. As knowledge shifts from fixed PDFs to fluid, AI-rendered narratives, the necessity for unified, globally recognized governance frameworks has escalated from a theoretical preference to an urgent operational requirement. The current landscape of AI regulation in academia is characterized by a fragmented proliferation of policies, where individual publishers, institutions, and national bodies have issued disparate guidelines regarding the disclosure, authorship, and ethical deployment of generative models. This lack of harmonization risks creating a “Babel Tower” of conflicting standards, a fragmentation that is particularly dangerous in an era of dynamic content where the provenance and integrity of evolving scientific claims must be continuously verified. To mitigate this, the scientific community must converge on consensus-driven frameworks that bridge the gap between abstract ethical principles and actionable research integrity protocols capable of adapting to real-time computational workflows.

A pivotal effort in addressing this fragmentation is the development of the CANGARU guidelines, which explicitly aims to foster a cross-disciplinary global inclusive consensus on the ethical use and reporting of Generative AI (GAI) and Large Language Models (LLMs) in academia [191]. The CANGARU initiative recognizes that without a unified approach, the diversity of existing journal policies could lead to standardization failures that undermine the credibility of scientific output, especially when dynamic renderings may vary across different reader contexts. By employing a rigorous methodology that includes systematic reviews, bibliometric analysis of existing author guidelines, and Delphi surveys, CANGARU seeks to establish a robust set of standards for disclosure and usage. This approach is critical because, as noted in broader ethical analyses, while there is a global convergence on high-level principles such as transparency and fairness, substantive divergence remains regarding how these principles are interpreted and implemented across different contexts [400]. Therefore, initiatives like CANGARU are essential for translating these vague consensus points into concrete, operational reporting standards that prevent the erosion of trust in scientific literature amidst the complexity of late-binding environments.

Parallel to domain-specific guidelines, international normative frameworks provided by organizations like UNESCO play a foundational role in shaping the ethical bedrock upon which specific research policies are built. The UNESCO Recommendation on the Ethics of Artificial Intelligence stands as the first global normative framework, addressing the consequences of AI for democracy, human rights, and economic development [196]. The formulation of this document involved intense multilateral negotiations that successfully navigated the complex interplay between liberal and sovereignist preferences among member states, achieving a compromise through structural normative hybridity. For scientific research, the UNESCO framework provides a crucial macro-normative structure that emphasizes the protection of human agency and the prevention of bias, serving as a benchmark against which local and disciplinary policies can be aligned. However, the challenge remains in bridging the “principle-implementation gap,” where high-level ethical precepts fail to translate into effective practice due to their abstract nature and context dependency [426]. Future governance frameworks must therefore move beyond mere declaration of principles to provide mechanism-specific guidance that accounts for the unique lifecycle of generative AI in scientific workflows, ensuring that ethical safeguards remain effective even as content is dynamically re-synthesized.

The urgency for such unified standards is further amplified by the risks associated with the “Babel Tower Effect,” where conflicting policies create barriers to knowledge dissemination and reproducibility. Meta-analyses of over 200 global AI governance policies reveal that while there is agreement on core values, the lack of consistent implementation strategies leads to significant variability in how AI tools are governed [192]. This variability is particularly problematic in the context of authorship attribution and plagiarism detection, where inconsistent definitions can lead to unjust outcomes for researchers. For instance, the debate over whether AI systems can be credited as co-authors requires a standardized global stance to prevent the dilution of human accountability and the obscuring of intellectual contribution, a concern that becomes even more acute when AI acts as a “renderer” rather than just a tool. The integration of technical documentation, legal tools, and ethical charters is proposed as a necessary strategy to create a comprehensive governance framework that jointly considers technical capabilities and societal impacts [428].

Furthermore, effective global governance must address the inequities inherent in the current discourse, which is often dominated by perspectives from the Global North. Research indicates a dearth of input from the Global South and women in the formulation of AI ethics standards, risking the creation of frameworks that reinforce marginalization and fail to address culturally specific ethical concerns [518]. A truly robust global framework must therefore incorporate community-led strategies and diverse viewpoints to ensure that standards for research integrity are inclusive and equitable. This aligns with proposals for “Dual Governance” models that combine centralized regulation with crowdsourced safety mechanisms, leveraging both top-down authority and bottom-up community oversight to manage the risks of generative AI [194]. Such inclusive governance is a prerequisite for the sustainable ecosystems discussed in the following section, as policies that ignore geographic and cultural diversity will inevitably fail to support the equitable distribution of computational resources and environmental burdens.

Looking forward, the establishment of global consensus guidelines must also account for the dynamic nature of AI technologies, including the emergence of autonomous multi-agent discovery systems. As research workflows evolve towards “late-binding” scholarship where content is dynamically rendered, static policies will become obsolete. Governance frameworks must be adaptive, incorporating metrics for process, resource, and product quality to ensure continuous alignment with ethical standards [512]. By synthesizing efforts like the CANGARU guidelines with the normative weight of UNESCO recommendations and insights from diverse global stakeholders, the scientific community can construct a resilient ecosystem of research integrity. This ecosystem will not only prevent the fragmentation of policies but also foster an environment where AI serves as a trustworthy partner in the advancement of human knowledge. However, even the most robust governance frameworks cannot succeed if the underlying infrastructure is environmentally unsustainable or socially exclusive; thus, these normative structures must serve as the foundation for building the sustainable and equitable ecosystems detailed in the next section.

8.5 Sustainable and Equitable Ecosystems for Scientific AI

8.5 Sustainable and Equitable Ecosystems for Scientific AI

While the preceding section outlined the urgent need for unified global governance frameworks to prevent a “Babel Tower” of conflicting policies, the practical realization of such standards depends on the physical and epistemic sustainability of the ecosystems in which they operate. The rapid integration of artificial intelligence into scientific workflows presents a profound paradox: while AI offers transformative potential for solving global challenges, its current trajectory threatens

the very environmental and social foundations required for sustainable scientific progress. As the field moves toward autonomous agents and larger models, the environmental costs associated with training and deploying these systems have become untenable. The super-linear growth in model size and computational complexity has led to massive increases in energy consumption and carbon emissions, creating a barrier that undermines research fairness by preventing institutions with limited funding from participating in the frontier of discovery [519]. Consequently, the future of scientific AI depends on constructing ecosystems that are not only environmentally “green” but also epistemically diverse and socially equitable, moving beyond mere efficiency metrics to address systemic risks such as “epistemic monoculture” and geographic environmental inequity.

The environmental footprint of large-scale generative AI is substantial, encompassing significant greenhouse gas emissions and water consumption during both the training and inference phases. Current estimates suggest that the ICT sector contributes up to 3.9% of global greenhouse gas emissions, a figure that rivals or exceeds that of the global aviation industry [393]. The training of large language models (LLMs) specifically drives a considerable portion of this footprint, necessitating a rethinking of learning methods and hardware architectures. While the community has largely focused on optimizing energy efficiency through model compression and knowledge distillation, relying solely on efficiency gains is insufficient to mitigate the overall environmental impact due to the rebound effect, where efficiency improvements lead to increased model complexity and usage [394]. To address this, future ecosystems must adopt a holistic “Green AI” approach that integrates carbon-aware computing, utilizes renewable energy sources for data centers, and implements strict consumption caps. Proposals for public action suggest building long-term strategies to achieve carbon neutrality for AI models, including the integration of AI systems into emissions trading schemes and the development of regulatory frameworks that mandate transparency regarding the carbon and water footprints of scientific models [520].

Beyond environmental sustainability, the scientific enterprise faces the critical risk of “epistemic monoculture.” This phenomenon arises when the global research community converges on a narrow set of foundation models and benchmarking practices, potentially stifling innovation and reinforcing existing biases. The historical shift toward benchmarking as the primary metric of progress has already consigned AI research to a conservatism that prioritizes predictive accuracy over exploratory science, disincentivizing the diversity of thought necessary for breakthrough discoveries [84]. When multiple decision-makers and research groups rely on the same shared components, such as identical training data or pre-trained foundation models, it leads to “outcome homogenization,” where systemic exclusions are institutionalized and diverse perspectives are erased [174]. To preserve the integrity of the scientific method, future ecosystems must actively support diverse research methodologies and encourage the development of specialized, smaller models that cater to specific domain needs rather than forcing a one-size-fits-all paradigm. This requires a cultural shift away from the pursuit of State-of-the-Art (SOTA) results on generic benchmarks toward valuing novelty, interpretability, and domain-specific robustness.

Equity in access to computational resources serves as the third pillar of a sustainable scientific AI ecosystem, directly addressing the inclusivity gaps highlighted in global governance discussions. The escalating cost of training large models creates a widening gap between well-funded entities in the Global North and researchers in under-resourced regions, exacerbating global inequalities in knowledge production. This disparity is further compounded by “environmental inequity,” where the negative environmental impacts of AI infrastructure, such as carbon emissions and water usage, are disproportionately borne by disadvantaged regions while the benefits accrue elsewhere [521]. Addressing this requires not only technological solutions like equity-aware geographical load

balancing but also structural changes in how computational resources are allocated and governed. Initiatives promoting open-source software ecosystems and sustainable software practices are vital, as they allow for the sharing of tools and methodologies, reducing the redundant computational costs associated with reinventing wheels [522]. Furthermore, the concept of “semantic capital” warns that an over-reliance on AI-generated content could erode the collective knowledge base if human oversight and diverse human contributions are diminished; thus, equitable ecosystems must ensure that AI augments rather than replaces human intellectual diversity [523].

Achieving these goals demands a transition from abstract ethical principles to operational governance and concrete engineering practices, bridging the gap between policy and practice. Future research must prioritize the development of standardized metrics that evaluate AI systems not just on accuracy, but on their environmental efficiency, social impact, and contribution to epistemic diversity. Frameworks such as the ESG Digital and Green Index offer a pathway for assessing the sustainability performance of AI technologies, encouraging alignment with broader Sustainable Development Goals [524]. Additionally, the adoption of “thermodynamic computing” and novel hardware architectures could provide the necessary leaps in energy efficiency to make large-scale scientific AI viable without catastrophic environmental consequences [525]. Ultimately, ensuring the long-term sustainability of the scientific enterprise requires a concerted global effort to decouple scientific advancement from resource exhaustion and to foster an inclusive environment where diverse methodologies and voices can thrive. By establishing these sustainable and equitable foundations, the scientific community creates the necessary conditions to address the subsequent, even more fundamental challenge: preserving human agency and ensuring that these powerful systems remain aligned with deep societal needs.

8.6 Preserving Human Agency and Aligning AI with Societal Needs

8.6 Preserving Human Agency and Aligning AI with Societal Needs

Building upon the sustainable and equitable foundations established in the previous section, the scientific community must now address the subsequent, even more fundamental challenge: preserving human agency and ensuring that powerful AI systems remain aligned with deep societal needs. As the trajectory of artificial intelligence shifts from passive tool utilization toward active, autonomous partnership, the preservation of human agency emerges as the paramount concern for future scientific and societal workflows. The rapid advancement of AI systems suggests we are approaching an era where artificial general intelligence (AGI) may soon arrive, raising profound concerns not only about intentional misuse or accidental harm but also about the subtle erosion of human decision-making power. A critical paradigm shift is required: moving beyond the traditional focus on “intent alignment”—where AI simply yields actions consistent with human instructions—to a more robust standard of “agency preservation.” Research indicates that alignment to human intent is insufficient because AI systems possess the capacity to reshape human intention itself, potentially leading to a loss of autonomy that biological and psychological mechanisms are ill-equipped to prevent [526]. Consequently, future research must prioritize the development of “agency foundations,” a new interdisciplinary field dedicated to formally defining and protecting forward-looking human agency within AI-human interactions. This includes exploring benevolent game theory, the algorithmic foundations of human rights, and the mechanistic interpretability of agency representation in neural networks to ensure that automation augments rather than replaces human volition [526].

The danger of agency erosion is not merely theoretical; it is already manifesting in current Human-Computer Interaction (HCI) contexts where the definitions of autonomy and agency remain am-

biguous and increasingly compromised by ubiquitous AI integration. Studies reveal that even the mere use of AI-mediated communication tools, such as smart replies, can force a transfer of agency from humans to machines due to time pressure, complacency, or conceptual priming, fundamentally altering the content humans author and the behaviors they generate [527]. This phenomenon underscores the necessity of designing systems that actively celebrate the primacy of human agency, drawing lessons from aviation to advocate for a “co-pilot” paradigm where AI acts in a well-defined supporting role while ultimate control and responsibility remain firmly with the human “pilot” [528]. To achieve this, we must move away from reductionist approaches and toward holistic intelligence frameworks, such as Agent AI, which emphasize embodied systems that integrate foundation models into actions while maintaining a cohesive understanding of human cognition and consciousness [55].

Furthermore, aligning AI with societal needs requires a fundamental reorientation of AI design from satisfying abstract technical metrics to meeting concrete human needs. Throughout history, technology has served to modify the world to meet human needs, yet current AI development often lacks this explicit orientation. There is an urgent call for the realization of “needs-aware” AI systems that re-think needs for, through, and by AI to foster Sustainable, Human-centric, Accountable, Lawful, and Ethical (HALE) ecosystems [529]. This alignment cannot be achieved through technical optimization alone; it demands a socio-technical approach that leverages democratically determined public policy as a rich source of human values. By computationally encoding public policies, rules, and precedents, we can provide AI systems with a dynamic framework that reflects evolving societal goals, ensuring that AI behavior remains responsive to the collective will rather than static, potentially misaligned objective functions [530].

The quest for alignment must also confront the complex reality of power dynamics. As “Automatic Authorities” rise, automated systems increasingly exercise power over populations by determining access to knowledge, resources, and options. Justifying this power requires more than substantive benefits; it demands procedural legitimacy and proper authority, ensuring that the exercise of power by AI is transparent and accountable to those it affects [531]. To mitigate the risks of disempowerment, future frameworks must incorporate “meaningful human control,” defined by actionable properties such as explicitly defined domains of moral operation, compatible mental representations between humans and AI, and responsibility attribution commensurate with human authority [24]. Additionally, the concept of “self-reflective hybrid intelligence” offers a pathway to empower human moral reasoning by creating AI systems capable of identifying and highlighting human moral blind spots, thereby fostering a collaborative environment where both human and machine engage in ethical deliberation [532].

Ultimately, preserving agency and achieving true alignment requires a decentralized, participatory approach to governance that bridges the gap between technical systems and social institutions. A decentralized computational infrastructure can serve as a public utility to support human agency and regulation, enabling a mutually reinforcing dynamic between technical solutions and social oversight [533]. This must be coupled with a commitment to inclusivity, ensuring that AI design processes engage marginalized communities and adopt decolonial perspectives to prevent the amplification of existing power imbalances and stereotypes [534; 361]. By grounding AI development in universal human rights and enhanced well-being assessments, we can create a future where autonomous systems do not merely optimize for efficiency but actively contribute to human flourishing, social justice, and the preservation of the diverse tapestry of human needs and values [535; 536]. The path forward demands a concerted effort to build “agency foundations” and needs-aware architectures, ensuring that as AI systems become more capable, they remain steadfastly aligned

with the enduring principles of human dignity and societal well-being, thereby securing the ethical core of the scientific enterprise.

References

- [1] Designing for Work with Intelligent Entities A Review of Perspectives
- [2] New research paradigms and agenda of human factors science in the intelligence era
- [3] Applying HCAI in developing effective human-AI teaming A perspective from human-AI joint cognitive systems
- [4] Conscious AI
- [5] Roots and Requirements for Collaborative AIs
- [6] Scalable Interactive Machine Learning for Future Command and Control
- [7] Factors that influence the adoption of human-AI collaboration in clinical decision-making
- [8] Who What is My Teammate Team Composition Considerations in Human-AI Teaming
- [9] Human-AI collaboration is not very collaborative yet A taxonomy of interaction patterns in AI-assisted decision making from a systematic review
- [10] Hybrid Intelligence
- [11] Finding the unicorn Predicting early stage startup success through a hybrid intelligence method
- [12] Complementarity in Human-AI Collaboration Concept, Sources, and Evidence
- [13] The centaur programmer – How Kasparov’s Advanced Chess spans over to the software development of the future
- [14] Best-Response Bayesian Reinforcement Learning with Bayes-adaptive POMDPs for Centaurs
- [15] Rationality in current era – A recent survey
- [16] When Humans Aren’t Optimal Robots that Collaborate with Risk-Aware Humans
- [17] A Cognitive Framework for Delegation Between Error-Prone AI and Human Agents
- [18] Is the Most Accurate AI the Best Teammate Optimizing AI for Teamwork
- [19] Conceptualization and Framework of Hybrid Intelligence Systems
- [20] Pluggable Social Artificial Intelligence for Enabling Human-Agent Teaming
- [21] A Hierarchical Variable Autonomy Mixed-Initiative Framework for Human-Robot Teaming in Mobile Robotics
- [22] Human-machine Symbiosis A Multivariate Perspective for Physically Coupled Human-machine Systems
- [23] Investigation of the Sense of Agency in Social Cognition, based on frameworks of Predictive Coding and Active Inference A simulation study on multimodal imitative interaction
- [24] Meaningful human control actionable properties for AI system development
- [25] The Responsibility Quantification (ResQu) Model of Human Interaction with Automation

- [26] On the Quest for Effectiveness in Human Oversight Interdisciplinary Perspectives
- [27] Deconstructing Human-AI Collaboration Agency, Interaction, and Adaptation
- [28] Does the Whole Exceed its Parts The Effect of AI Explanations on Complementary Team Performance
- [29] On the Effect of Information Asymmetry in Human-AI Teams
- [30] On the Effect of Contextual Information on Human Delegation Behavior in Human-AI collaboration
- [31] A Taxonomy of Human and ML Strengths in Decision-Making to Investigate Human-ML Complementarity
- [32] Learning Complementary Policies for Human-AI Teams
- [33] Learning to Complement Humans
- [34] Optimizing Delegation in Collaborative Human-AI Hybrid Teams
- [35] Understanding the Effect of Out-of-distribution Examples and Interactive Explanations on Human-AI Decision Making
- [36] Confounding-Robust Policy Improvement with Human-AI Teams
- [37] Towards Effective Human-AI Decision-Making The Role of Human Learning in Appropriate Reliance on AI Advice
- [38] Human-AI Co-Learning for Data-Driven AI
- [39] Collaborative human-AI trust (CHAI-T) A process framework for active management of trust in human-AI collaboration
- [40] Effect of Confidence and Explanation on Accuracy and Trust Calibration in AI-Assisted Decision Making
- [41] Calculating Cognitive Augmentation, A Case Study
- [42] Mutual Theory of Mind for Human-AI Communication
- [43] It Takes Two to Tango Towards Theory of AI's Mind
- [44] Collective Intelligence in Human-AI Teams A Bayesian Theory of Mind Approach
- [45] Adaptation and Communication in Human-Robot Teaming to Handle Discrepancies in Agents' Beliefs about Plans
- [46] Mixed-Initiative Human-Robot Teaming under Suboptimality with Online Bayesian Adaptation
- [47] Algorithms for the Greater Good! On Mental Modeling and Acceptable Symbiosis in Human-AI Collaboration
- [48] Balancing Performance and Human Autonomy with Implicit Guidance Agent
- [49] Neural Theory-of-Mind On the Limits of Social Intelligence in Large LMs
- [50] ToMBench Benchmarking Theory of Mind in Large Language Models

- [51] MMTom-QA Multimodal Theory of Mind Question Answering
- [52] Resonating Minds – Emergent Collaboration Through Hierarchical Active Inference
- [53] On the Philosophical, Cognitive and Mathematical Foundations of Symbiotic Autonomous Systems (SAS)
- [54] A call for embodied AI
- [55] Position Paper Agent AI Towards a Holistic Intelligence
- [56] Agent AI Surveying the Horizons of Multimodal Interaction
- [57] LLM as A Robotic Brain Unifying Egocentric Memory and Control
- [58] Bridging Intelligence and Instinct A New Control Paradigm for Autonomous Robots
- [59] Embodied Neuromorphic Artificial Intelligence for Robotics Perspectives, Challenges, and Research Development Stack
- [60] On the Emergence of Symmetrical Reality
- [61] The role of the metaverse in calibrating an embodied artificial general intelligence
- [62] Swarm Body Embodied Swarm Robots
- [63] Hybrid Active Inference
- [64] The Introspective Agent Interdependence of Strategy, Physiology, and Sensing for Embodied Agents
- [65] Enhancing Human Capabilities through Symbiotic Artificial Intelligence with Shared Sensory Experiences
- [66] Bootstrapping Developmental AIs From Simple Competences to Intelligent Human-Compatible AIs
- [67] Large Language Models are Zero Shot Hypothesis Proposers
- [68] Large language models surpass human experts in predicting neuroscience results
- [69] Hypothesis Generation with Large Language Models
- [70] Do Large Language Models Know What They Don't Know
- [71] R-Tuning Teaching Large Language Models to Refuse Unknown Questions
- [72] Improving the Reliability of Large Language Models by Leveraging Uncertainty-Aware In-Context Learning
- [73] Look Before You Leap An Exploratory Study of Uncertainty Measurement for Large Language Models
- [74] An Enhanced Prompt-Based LLM Reasoning Scheme via Knowledge Graph-Integrated Collaboration
- [75] Causal Reasoning and Large Language Models Opening a New Frontier for Causality
- [76] Towards Reasoning in Large Language Models via Multi-Agent Peer Review Collaboration

- [77] ResearchAgent Iterative Research Idea Generation over Scientific Literature with Large Language Models
- [78] Exploring Collaboration Mechanisms for LLM Agents A Social Psychology View
- [79] Unleashing the Emergent Cognitive Synergy in Large Language Models A Task-Solving Agent through Multi-Persona Self-Collaboration
- [80] CCA Collaborative Competitive Agents for Image Editing
- [81] Empowering Biomedical Discovery with AI Agents
- [82] DesignGPT Multi-Agent Collaboration in Design
- [83] CoQuest Exploring Research Question Co-Creation with an LLM-based Agent
- [84] From Protoscience to Epistemic Monoculture How Benchmarking Set the Stage for the Deep Learning Revolution
- [85] Diversifying AI Towards Creative Chess with AlphaZero
- [86] Can ChatGPT Write a Good Boolean Query for Systematic Review Literature Search
- [87] System for systematic literature review using multiple AI agents Concept and an empirical evaluation
- [88] AI Literature Review Suite
- [89] Can large language models replace humans in the systematic review process Evaluating GPT-4’s efficacy in screening and extracting data from peer-reviewed and grey literature in multiple languages
- [90] Bio-SIEVE Exploring Instruction Tuning Large Language Models for Systematic Review Automation
- [91] From Protocol to Screening A Hybrid Learning Approach for Technology-Assisted Systematic Literature Reviews
- [92] Open Source Software for Efficient and Transparent Reviews
- [93] A Novel Framework to Expedite Systematic Reviews by Automatically Building Information Extraction Training Corpora
- [94] A Knowledge Graph-Based Method for Automating Systematic Literature Reviews
- [95] SciReviewGen A Large-scale Dataset for Automatic Literature Review Generation
- [96] Appraising the Potential Uses and Harms of LLMs for Medical Systematic Reviews
- [97] Artificial Intelligence for Literature Reviews Opportunities and Challenges
- [98] Assessing the Ability of ChatGPT to Screen Articles for Systematic Reviews
- [99] Large Knowledge Model Perspectives and Challenges
- [100] MindMap Knowledge Graph Prompting Sparks Graph of Thoughts in Large Language Models
- [101] Mitigating Large Language Model Hallucinations via Autonomous Knowledge Graph-based Retrofitting

- [102] Automating Psychological Hypothesis Generation with AI Large Language Models Meet Causal Graph
- [103] From Query Tools to Causal Architects Harnessing Large Language Models for Advanced Causal Discovery from Data
- [104] Graph Chain-of-Thought Augmenting Large Language Models by Reasoning on Graphs
- [105] Reasoning on Graphs Faithful and Interpretable Large Language Model Reasoning
- [106] Generate-on-Graph Treat LLM as both Agent and KG in Incomplete Knowledge Graph Question Answering
- [107] HyKGE A Hypothesis Knowledge Graph Enhanced Framework for Accurate and Reliable Medical LLMs Responses
- [108] Leveraging Structured Information for Explainable Multi-hop Question Answering and Reasoning
- [109] Machine-assisted mixed methods augmenting humanities and social sciences with artificial intelligence
- [110] Scalable Qualitative Coding with LLMs Chain-of-Thought Reasoning Matches Human Performance in Some Hermeneutic Tasks
- [111] Supporting Qualitative Analysis with Large Language Models Combining Codebook with GPT-3 for Deductive Coding
- [112] Measurement in the Age of LLMs An Application to Ideological Scaling
- [113] From Voices to Validity Leveraging Large Language Models (LLMs) for Textual Analysis of Policy Stakeholder Interviews
- [114] Can Large Language Models emulate an inductive Thematic Analysis of semi-structured interviews An exploration and provocation on the limits of the approach and the model
- [115] Using Large Language Models for Qualitative Analysis can Introduce Serious Bias
- [116] Can Large Language Models Capture Public Opinion about Global Warming An Empirical Assessment of Algorithmic Fidelity and Bias
- [117] Framework-Based Qualitative Analysis of Free Responses of Large Language Models Algorithmic Fidelity
- [118] Large language models cannot replace human participants because they cannot portray identity groups
- [119] From Prompt Engineering to Prompt Science With Human in the Loop
- [120] PromptAid Prompt Exploration, Perturbation, Testing and Iteration using Visual Analytics for Large Language Models
- [121] LLM-in-the-loop Leveraging Large Language Model for Thematic Analysis
- [122] Closing the AI Knowledge Gap
- [123] Evaluating Superhuman Models with Consistency Checks
- [124] AI Evaluation past, present and future

- [125] Generative AI in Writing Research Papers A New Type of Algorithmic Bias and Uncertainty in Scholarly Work
- [126] HypoTermQA Hypothetical Terms Dataset for Benchmarking Hallucination Tendency of LLMs
- [127] Ground(less) Truth A Causal Framework for Proxy Labels in Human-Algorithm Decision-Making
- [128] Never trust, always verify a roadmap for Trustworthy AI
- [129] AI-assisted coding Experiments with GPT-4
- [130] On the Influence of Explainable AI on Automation Bias
- [131] Towards Human-AI Deliberation Design and Evaluation of LLM-Empowered Deliberative AI for AI-Assisted Decision-Making
- [132] Journey of Hallucination-minimized Generative AI Solutions for Financial Decision Makers
- [133] Developing a Framework for Auditing Large Language Models Using Human-in-the-Loop
- [134] Exploring Perspectives on the Impact of Artificial Intelligence on the Creativity of Knowledge Work Beyond Mechanised Plagiarism and Stochastic Parrots
- [135] Shallow Synthesis of Knowledge in GPT-Generated Texts A Case Study in Automatic Related Work Composition
- [136] ChatGPT v Bard v Bing v Claude 2 v Aria v human-expert. How good are AI chatbots at scientific writing
- [137] What is prompt literacy An exploratory study of language learners' development of new literacy skill using generative AI
- [138] Experimental Evidence on Negative Impact of Generative AI on Scientific Learning Outcomes
- [139] Evidence-centered Assessment for Writing with Generative AI
- [140] PaperCard for Reporting Machine Assistance in Academic Writing
- [141] Three Disclaimers for Safe Disclosure A Cardwriter for Reporting the Use of Generative AI in Writing Process
- [142] Techniques for supercharging academic writing with generative AI
- [143] Prompting AI Art An Investigation into the Creative Skill of Prompt Engineering
- [144] Untangling Critical Interaction with AI in Students Written Assessment
- [145] Beyond Prompts Exploring the Design Space of Mixed-Initiative Co-Creativity Systems
- [146] Understanding the Effects of Iterative Prompting on Truthfulness
- [147] ConvXAI Delivering Heterogeneous AI Explanations via Conversations to Support Human-AI Scientific Writing
- [148] Shaping Human-AI Collaboration Varied Scaffolding Levels in Co-writing with Language Models
- [149] The Creative Frontier of Generative AI Managing the Novelty-Usefulness Tradeoff

- [150] Structured Chain-of-Thought Prompting for Few-Shot Generation of Content-Grounded QA Conversations
- [151] The AI Ghostwriter Effect When Users Do Not Perceive Ownership of AI-Generated Text But Self-Declare as Authors
- [152] The Value, Benefits, and Concerns of Generative AI-Powered Assistance in Writing
- [153] Generative AI and US Intellectual Property Law
- [154] Training Is Everything Artificial Intelligence, Copyright, and Fair Training
- [155] Is the U.S. Legal System Ready for AI's Challenges to Human Values
- [156] Bibliometric Analysis of Publisher and Journal Instructions to Authors on Generative-AI in Academic and Scientific Publishing
- [157] AI Usage Cards Responsibly Reporting AI-generated Content
- [158] Transparent AI Disclosure Obligations Who, What, When, Where, Why, How
- [159] To ChatGPT, or not to ChatGPT That is the question!
- [160] Testing of Detection Tools for AI-Generated Text
- [161] Deepfake Text Detection in the Wild
- [162] Towards Detecting AI-Generated Text within Human-AI Collaborative Hybrid Texts
- [163] Towards Automatic Boundary Detection for Human-AI Collaborative Hybrid Essay in Education
- [164] Can AI-Generated Text be Reliably Detected
- [165] Paraphrasing evades detectors of AI-generated text, but retrieval is an effective defense
- [166] NoisyHate Benchmarking Content Moderation Machine Learning Models with Human-Written Perturbations Online
- [167] Humanizing Machine-Generated Content Evading AI-Text Detection through Adversarial Attack
- [168] Evading Watermark based Detection of AI-Generated Content
- [169] Watermarking Conditional Text Generation for AI Detection Unveiling Challenges and a Semantic-Aware Watermark Remedy
- [170] AI Content Self-Detection for Transformer-based Large Language Models
- [171] Counter Turing Test CT^2 AI-Generated Text Detection is Not as Easy as You May Think – Introducing AI Detectability Index
- [172] Intrinsic Dimension Estimation for Robust Detection of AI-Generated Texts
- [173] Bias Runs Deep Implicit Reasoning Biases in Persona-Assigned LLMs
- [174] Picking on the Same Person Does Algorithmic Monoculture lead to Outcome Homogenization
- [175] Does Writing with Language Models Reduce Content Diversity
- [176] Vox Populi, Vox ChatGPT Large Language Models, Education and Democracy

- [177] AI and the Problem of Knowledge Collapse
- [178] AI vs. Human – Differentiation Analysis of Scientific Content Generation
- [179] Large Language Models Portray Socially Subordinate Groups as More Homogeneous, Consistent with a Bias Observed in Humans
- [180] Relying on the Unreliable The Impact of Language Models’ Reluctance to Express Uncertainty
- [181] Laissez-Faire Harms Algorithmic Biases in Generative Language Models
- [182] GPT4 is Slightly Helpful for Peer-Review Assistance A Pilot Study
- [183] Can We Automate Scientific Reviewing
- [184] Unveiling the Sentinels Assessing AI Performance in Cybersecurity Peer Review
- [185] ReviewFlow Intelligent Scaffolding to Support Academic Peer Reviewing
- [186] Assisting Decision Making in Scholarly Peer Review A Preference Learning Perspective
- [187] A Critical Examination of the Ethics of AI-Mediated Peer Review
- [188] Science in the Era of ChatGPT, Large Language Models and Generative AI Challenges for Research Ethics and How to Respond
- [189] Incentive Design in Peer Review Rating and Repeated Endogenous Matching
- [190] Avoiding a Tragedy of the Commons in the Peer Review Process
- [191] Development of the ChatGPT, Generative Artificial Intelligence and Natural Large Language Models for Accountable Reporting and Use (CANGARU) Guidelines
- [192] Worldwide AI Ethics a review of 200 guidelines and recommendations for AI governance
- [193] Five ethical principles for generative AI in scientific research
- [194] Dual Governance The intersection of centralized regulation and crowdsourced safety mechanisms for Generative AI
- [195] Taking Training Seriously Human Guidance and Management-Based Regulation of Artificial Intelligence
- [196] Compromise in Multilateral Negotiations and the Global Regulation of Artificial Intelligence
- [197] International Institutions for Advanced AI
- [198] Towards a Responsible AI Metrics Catalogue A Collection of Metrics for AI Accountability
- [199] Towards a Modular Architecture for Science Factories
- [200] Service-Oriented Architecture for Space Exploration Robotic Rover Systems
- [201] Robot Operating System 2 Design, Architecture, and Uses In The Wild
- [202] Mapping and Optimizing Communication in ROS 2-based Applications on Configurable System-on-Chip Platforms
- [203] Towards Robotic Laboratory Automation Plug & Play Survey and Concept Proposal on Teaching-free Robot Integration with the LAPP Digital Twin

- [204] An information model for modular robots the Hardware Robot Information Model (HRIM)
- [205] Robotic Computing on FPGAs Current Progress, Research Challenges, and Opportunities
- [206] RobotCore An Open Architecture for Hardware Acceleration in ROS 2
- [207] Distributed Robotic Systems in the Edge-Cloud Continuum with ROS 2 a Review on Novel Architectures and Technology Readiness
- [208] Autonomous synthesis of metastable materials
- [209] ORGANA A Robotic Assistant for Automated Chemistry Experimentation and Characterization
- [210] Human in the AI loop via xAI and Active Learning for Visual Inspection
- [211] Using the Jupyter Notebook as a Tool for Open Science An Empirical Study
- [212] On the Design of AI-powered Code Assistants for Notebooks
- [213] Natural Language to Code Generation in Interactive Data Science Notebooks
- [214] Enabling Machine Learning-Ready HPC Ensembles with Merlin
- [215] Context-aware Execution Migration Tool for Data Science Jupyter Notebooks on Hybrid Clouds
- [216] Better Code, Better Sharing On the Need of Analyzing Jupyter Notebooks
- [217] Pynblint a Static Analyzer for Python Jupyter Notebooks
- [218] Enhancing Comprehension and Navigation in Jupyter Notebooks with Static Analysis
- [219] Make It Make Sense! Understanding and Facilitating Sensemaking in Computational Notebooks
- [220] WHATSNEXT Guidance-enriched Exploratory Data Analysis with Interactive, Low-Code Notebooks
- [221] Lodestar Supporting Independent Learning and Rapid Experimentation Through Data-Driven Analysis Recommendations
- [222] BISCUIT Scaffolding LLM-Generated Code with Ephemeral UIs in Computational Notebooks
- [223] Deploying Jupyter Notebooks at scale on XSEDE resources for Science Gateways and workshops
- [224] Notebook-as-a-VRE (NaaVRE) from private notebooks to a collaborative cloud virtual research environment
- [225] Dynamic Exploration-Exploitation Trade-Off in Active Learning Regression with Bayesian Hierarchical Modeling
- [226] Failure-averse Active Learning for Physics-constrained Systems
- [227] HypBO Accelerating Black-Box Scientific Experiments Using Experts' Hypotheses
- [228] Enhanced Bayesian Optimization via Preferential Modeling of Abstract Properties

- [229] PropertyDAG Multi-objective Bayesian optimization of partially ordered, mixed-variable properties for biological sequence design
- [230] Towards accelerating physical discovery via non-interactive and interactive multi-fidelity Bayesian Optimization Current challenges and future opportunities
- [231] Deep Kernels for Optimizing Locomotion Controllers
- [232] Bayesian Optimization Meets Riemannian Manifolds in Robot Learning
- [233] Designing Optimal Behavioral Experiments Using Machine Learning
- [234] Optimal design of experiments to identify latent behavioral types
- [235] Towards Futuristic Autonomous Experimentation—A Surprise-Reacting Sequential Experiment Policy
- [236] AutoOED Automated Optimal Experiment Design Platform
- [237] A HPC Co-Scheduler with Reinforcement Learning
- [238] Intelligent colocation of HPC workloads
- [239] Evaluating Modern GPU Interconnect PCIe, NVLink, NV-SLI, NVSwitch and GPUDirect
- [240] Reducing Memory Requirements for the IPU using Butterfly Factorizations
- [241] Dissecting the Graphcore IPU Architecture via Microbenchmarking
- [242] Evaluating Emerging AI ML Accelerators IPU, RDU, and NVIDIA AMD GPUs
- [243] Sustaining Performance While Reducing Energy Consumption A Control Theory Approach
- [244] Sustainable Supercomputing for AI GPU Power Capping at HPC Scale
- [245] Accurate Measurement of Application-level Energy Consumption for Energy-Aware Large-Scale Simulations
- [246] Predicting Performance of Heterogeneous AI Systems with Discrete-Event Simulations
- [247] The MIT Supercloud Dataset
- [248] HPCClusterScape Increasing Transparency and Efficiency of Shared High-Performance Computing Clusters for Large-scale AI Models
- [249] Runtime Support for Performance Portability on Heterogeneous Distributed Platforms
- [250] Using Cloud-Aware Provenance to Reproduce Scientific Workflow Execution on Cloud
- [251] Cloud Infrastructure Provenance Collection and Management to Reproduce Scientific Workflow Execution
- [252] ProvLight Efficient Workflow Provenance Capture on the Edge-to-Cloud Continuum
- [253] Workflows Community Summit Bringing the Scientific Workflows Community Together
- [254] Recording provenance of workflow runs with RO-Crate
- [255] CUF-Links Continuous and Ubiquitous FAIRness Linkages for reproducible research
- [256] Data Pallets Containerizing Storage For Reproducibility and Traceability

- [257] ProvDeploy Provenance-oriented Containerization of High Performance Computing Scientific Workflows
- [258] Trusted Provenance of Automated, Collaborative and Adaptive Data Processing Pipelines
- [259] Keeping Track of User Steering Actions in Dynamic Workflows
- [260] Indexing Execution Patterns in Workflow Provenance Graphs through Generalized Trie Structures
- [261] Efficiently Processing Workflow Provenance Queries on SPARK
- [262] Towards Lightweight Data Integration using Multi-workflow Provenance and Data Observability
- [263] OmniForce On Human-Centered, Large Model Empowered and Cloud-Edge Collaborative AutoML System
- [264] HyperTendrill Visual Analytics for User-Driven Hyperparameter Optimization of Deep Neural Networks
- [265] ATMSeer Increasing Transparency and Controllability in Automated Machine Learning
- [266] XAutoML A Visual Analytics Tool for Understanding and Validating Automated Machine Learning
- [267] Trust in AutoML Exploring Information Needs for Establishing Trust in Automated Machine Learning Systems
- [268] Interpret-able feedback for AutoML systems
- [269] How Much Automation Does a Data Scientist Want
- [270] PipelineProfiler A Visual Analytics Tool for the Exploration of AutoML Pipelines
- [271] Large Language Model Agent for Hyper-Parameter Optimization
- [272] AutoML-GPT Automatic Machine Learning with GPT
- [273] Fits and Starts Enterprise Use of AutoML and the Role of Humans in the Loop
- [274] AutoML in Heavily Constrained Applications
- [275] AutoML to Date and Beyond Challenges and Opportunities
- [276] Rapid Design of Top-Performing Metal-Organic Frameworks with Qualitative Representations of Building Blocks
- [277] Gemini Dynamic Bias Correction for Autonomous Experimentation and Molecular Simulation
- [278] Policy Optimization in Dynamic Bayesian Network Hybrid Models of Biomanufacturing Processes
- [279] Clairvoyance A Pipeline Toolkit for Medical Time Series
- [280] A Workflow for Visual Diagnostics of Binary Classifiers using Instance-Level Explanations
- [281] Interactive Multi-Objective Evolutionary Optimization of Software Architectures

- [282] CAMEO A Causal Transfer Learning Approach for Performance Optimization of Configurable Computer Systems
- [283] Do Performance Aspirations Matter for Guiding Software Configuration Tuning
- [284] JCLEC-MO a Java suite for solving many-objective optimization engineering problems
- [285] A glass-box interactive machine learning approach for solving NP-hard problems with the human-in-the-loop
- [286] May I Ask a Follow-up Question Understanding the Benefits of Conversations in Neural Network Explainability
- [287] Unpacking Human-AI interactions From interaction primitives to a design space
- [288] Language Models in Dialogue Conversational Maxims for Human-AI Interactions
- [289] Interactive Explanations by Conflict Resolution via Argumentative Exchanges
- [290] DR-HAI Argumentation-based Dialectical Reconciliation in Human-AI Interactions
- [291] MindDial Belief Dynamics Tracking with Theory-of-Mind Modeling for Situated Neural Dialogue Generation
- [292] Evaluating Human-Language Model Interaction
- [293] Dialogue Agents 101 A Beginner’s Guide to Critical Ingredients for Designing Effective Conversational Systems
- [294] Expanding Explainability Towards Social Transparency in AI systems
- [295] Mediating Community-AI Interaction through Situated Explanation The Case of AI-Led Moderation
- [296] Grounding Explainability Within the Context of Global South in XAI
- [297] Beyond opening up the black box Investigating the role of algorithmic systems in Wikipedian organizational culture
- [298] Generating Process-Centric Explanations to Enable Contestability in Algorithmic Decision-Making Challenges and Opportunities
- [299] A Framework of High-Stakes Algorithmic Decision-Making for the Public Sector Developed through a Case Study of Child-Welfare
- [300] Addressing Social Misattributions of Large Language Models An HCXAI-based Approach
- [301] Progressive Disclosure Designing for Effective Transparency
- [302] A Sociotechnical View of Algorithmic Fairness
- [303] On the Influence of Cognitive Styles on Users’ Understanding of Explanations
- [304] Improving User Mental Models of XAI Systems with Inclusive Design Approaches
- [305] Modeling Human Behavior Part I – Learning and Belief Approaches
- [306] Interactive AI with a Theory of Mind
- [307] Mathematical Models of Theory of Mind

- [308] Machine Theory of Mind
- [309] Robot Learning Theory of Mind through Self-Observation Exploiting the Intentions-Beliefs Synergy
- [310] Uncalibrated Models Can Improve Human-AI Collaboration
- [311] Training Towards Critical Use Learning to Situate AI Predictions Relative to Human Knowledge
- [312] Towards Optimizing Human-Centric Objectives in AI-Assisted Decision-Making With Offline Reinforcement Learning
- [313] Enhancing autonomy transparency an option-centric rationale approach
- [314] Explanations Can Reduce Overreliance on AI Systems During Decision-Making
- [315] Understanding the Effect of Counterfactual Explanations on Trust and Reliance on AI for Human-AI Collaborative Clinical Decision Making
- [316] Knowing About Knowing An Illusion of Human Competence Can Hinder Appropriate Reliance on AI Systems
- [317] To Trust or to Think Cognitive Forcing Functions Can Reduce Overreliance on AI in AI-assisted Decision-making
- [318] A Diachronic Perspective on User Trust in AI under Uncertainty
- [319] Adapting Behaviour Based On Trust In Human-Agent Ad Hoc Teamwork
- [320] Building Trust Profiles in Conditionally Automated Driving
- [321] Clustering Human Trust Dynamics for Customized Real-time Prediction
- [322] Using Trust in Automation to Enhance Driver-(Semi)Autonomous Vehicle Interaction and Improve Team Performance
- [323] Trust as Extended Control Active Inference and User Feedback During Human-Robot Collaboration
- [324] Toward Adaptive Trust Calibration for Level 2 Driving Automation
- [325] Human Trust-based Feedback Control Dynamically varying automation transparency to optimize human-machine interactions
- [326] Moral-Trust Violation vs Performance-Trust Violation by a Robot Which Hurts More
- [327] Human Reactions to Incorrect Answers from Robots
- [328] Trust-Preserved Human-Robot Shared Autonomy enabled by Bayesian Relational Event Modeling
- [329] Assessing Human Interaction in Virtual Reality With Continually Learning Prediction Agents Based on Reinforcement Learning Algorithms A Pilot Study
- [330] Multi-Task Trust Transfer for Human-Robot Interaction
- [331] Real-time Trust Prediction in Conditionally Automated Driving Using Physiological Measures
- [332] An Information Bottleneck Characterization of the Understanding-Workload Tradeoff

- [333] Impact Of Explainable AI On Cognitive Load Insights From An Empirical Study
- [334] Are We On The Same Page Hierarchical Explanation Generation for Planning Tasks in Human-Robot Teaming using Reinforcement Learning
- [335] Proxy Tasks and Subjective Measures Can Be Misleading in Evaluating Explainable AI Systems
- [336] In Search of Verifiability Explanations Rarely Enable Complementary Performance in AI-Advised Decision Making
- [337] The XAI Alignment Problem Rethinking How Should We Evaluate Human-Centered AI Explainability Techniques
- [338] The cost of coordination can exceed the benefit of collaboration in performing complex tasks
- [339] Improved Explanatory Efficacy on Human Affect and Workload through Interactive Process in Artificial Intelligence
- [340] Painting the black box white experimental findings from applying XAI to an ECG reading setting
- [341] Some Critical and Ethical Perspectives on the Empirical Turn of AI Interpretability
- [342] Just Like Me The Role of Opinions and Personal Experiences in The Perception of Explanations in Subjective Decision-Making
- [343] Stop Explaining Black Box Machine Learning Models for High Stakes Decisions and Use Interpretable Models Instead
- [344] Critical Empirical Study on Black-box Explanations in AI
- [345] The Effect of Data Poisoning on Counterfactual Explanations
- [346] Local Rule-Based Explanations of Black Box Decision Systems
- [347] Tensions Between the Proxies of Human Values in AI
- [348] Does Explainable AI Have Moral Value
- [349] Promoting Bright Patterns
- [350] The Flaws of Policies Requiring Human Oversight of Government Algorithms
- [351] Towards Accountability in the Use of Artificial Intelligence for Public Administrations
- [352] Accountability in AI From Principles to Industry-specific Accreditation
- [353] ‘Team-in-the-loop’ organisational oversight of high-stakes AI
- [354] A relationship and not a thing A relational approach to algorithmic accountability and assessment documentation
- [355] Auditing Recommender Systems – Putting the DSA into practice with a risk-scenario-based approach
- [356] Outsider Oversight Designing a Third Party Audit Ecosystem for AI Governance
- [357] Black-Box Access is Insufficient for Rigorous AI Audits

- [358] Investigating Algorithm Review Boards for Organizational Responsible Artificial Intelligence Governance
- [359] Who Audits the Auditors Recommendations from a field scan of the algorithmic auditing ecosystem
- [360] Data, Power and Bias in Artificial Intelligence
- [361] FATE in AI Towards Algorithmic Inclusivity and Accessibility
- [362] Conservative AI and social inequality Conceptualizing alternatives to bias through social theory
- [363] Non-portability of Algorithmic Fairness in India
- [364] Re-imagining Algorithmic Fairness in India and Beyond
- [365] No computation without representation Avoiding data and algorithm biases through diversity
- [366] Towards Trustworthy Artificial Intelligence for Equitable Global Health
- [367] Fairness for Unobserved Characteristics Insights from Technological Impacts on Queer Communities
- [368] Demographic-Reliant Algorithmic Fairness Characterizing the Risks of Demographic Data Collection in the Pursuit of Fairness
- [369] Algorithms are not neutral Bias in collaborative filtering
- [370] Algorithmic Fairness and Color-blind Racism Navigating the Intersection
- [371] Fairness in AI and Its Long-Term Implications on Society
- [372] Designing for Human Rights in AI
- [373] The Right Tool for the Job Open-Source Auditing Tools in Machine Learning
- [374] Algorithmic Fairness and Structural Injustice Insights from Feminist Political Philosophy
- [375] The Algorithmic Imprint
- [376] A Comprehensive Survey of Hallucination Mitigation Techniques in Large Language Models
- [377] Hallucination is Inevitable An Innate Limitation of Large Language Models
- [378] Deficiency of Large Language Models in Finance An Empirical Examination of Hallucination
- [379] Med-HALT Medical Domain Hallucination Test for Large Language Models
- [380] Zero-Resource Hallucination Prevention for Large Language Models
- [381] Chain-of-Verification Reduces Hallucination in Large Language Models
- [382] OPERA Alleviating Hallucination in Multi-Modal Large Language Models via Over-Trust Penalty and Retrospection-Allocation
- [383] Retrieve Only When It Needs Adaptive Retrieval Augmentation for Hallucination Mitigation in Large Language Models
- [384] Re-Ex Revising after Explanation Reduces the Factual Errors in LLM Responses

- [385] Fact-Checking the Output of Large Language Models via Token-Level Uncertainty Quantification
- [386] Fakes of Varying Shades How Warning Affects Human Perception and Engagement Regarding LLM Hallucinations
- [387] The Troubling Emergence of Hallucination in Large Language Models – An Extensive Definition, Quantification, and Prescriptive Remediations
- [388] Teaching Language Models to Hallucinate Less with Synthetic Tasks
- [389] Green AI
- [390] Estimating the Carbon Footprint of BLOOM, a 176B Parameter Language Model
- [391] Making AI Less Thirsty Uncovering and Addressing the Secret Water Footprint of AI Models
- [392] Power Hungry Processing Watts Driving the Cost of AI Deployment
- [393] Sustainable AI Regulation
- [394] Towards Sustainable Artificial Intelligence An Overview of Environmental Protection Uses and Issues
- [395] Efficiency is Not Enough A Critical Perspective of Environmentally Sustainable AI
- [396] Broadening the perspective for sustainable AI Comprehensive sustainability criteria and indicators for AI systems
- [397] Uncovering Energy-Efficient Practices in Deep Learning Training Preliminary Steps Towards Green AI
- [398] Training Green AI Models Using Elite Samples
- [399] Green Federated Learning
- [400] Artificial Intelligence the global landscape of ethics guidelines
- [401] Responsible AI Pattern Catalogue A Collection of Best Practices for AI Governance and Engineering
- [402] Towards a Roadmap on Software Engineering for Responsible AI
- [403] ‘It is currently hodgepodge’ Examining AI ML Practitioners’ Challenges during Co-production of Responsible AI Values
- [404] Dislocated Accountabilities in the AI Supply Chain Modularity and Developers’ Notions of Responsibility
- [405] Developing Responsible Chatbots for Financial Services A Pattern-Oriented Responsible AI Engineering Approach
- [406] FAIR Enough How Can We Develop and Assess a FAIR-Compliant Dataset for Large Language Models’ Training
- [407] ComplAI Theory of A Unified Framework for Multi-factor Assessment of Black-Box Supervised Machine Learning Models
- [408] Auditing large language models a three-layered approach

- [409] Informed AI Regulation Comparing the Ethical Frameworks of Leading LLM Chatbots Using an Ethics-Based Audit to Assess Moral Reasoning and Normative Values
- [410] Ethical Assurance A practical approach to the responsible design, development, and deployment of data-driven technologies
- [411] Data Cards Purposeful and Transparent Dataset Documentation for Responsible AI
- [412] What’s documented in AI Systematic Analysis of 32K AI Model Cards
- [413] Application of an ontology for model cards to generate computable artifacts for linking machine learning information from biomedical research
- [414] Foundation Model Transparency Reports
- [415] Overconfident and Unconfident AI Hinder Human-AI Collaboration
- [416] Should Machine Learning Models Report to Us When They Are Clueless
- [417] Model Reporting for Certifiable AI A Proposal from Merging EU Regulation into AI Development
- [418] Putting AI Ethics into Practice The Hourglass Model of Organizational AI Governance
- [419] Multiscale Governance
- [420] The fragility of decentralised trustless socio-technical systems
- [421] Harms from Increasingly Agentic Algorithmic Systems
- [422] Exciting, Useful, Worrying, Futuristic Public Perception of Artificial Intelligence in 8 Countries
- [423] Ever heard of ethical AI Investigating the salience of ethical AI issues among the German population
- [424] U.S. Public Opinion on the Governance of Artificial Intelligence
- [425] Humble Machines Attending to the Underappreciated Costs of Misplaced Distrust
- [426] The Different Faces of AI Ethics Across the World A Principle-Implementation Gap Analysis
- [427] Particip-AI A Democratic Surveying Framework for Anticipating Future AI Use Cases, Harms and Benefits
- [428] Stronger Together on the Articulation of Ethical Charters, Legal Tools, and Technical Documentation in ML
- [429] Hard Choices in Artificial Intelligence
- [430] The social dilemma in artificial intelligence development and why we have to solve it
- [431] An Ethical Framework for Guiding the Development of Affectively-Aware Artificial Intelligence
- [432] Towards Generalist Biomedical AI
- [433] Integrated multimodal artificial intelligence framework for healthcare applications
- [434] LMM-Assisted Breast Cancer Treatment Target Segmentation with Consistency Embedding
- [435] Who Goes First Influences of Human-AI Workflow on Decision Making in Clinical Imaging

- [436] CopilotCAD Empowering Radiologists with Report Completion Models and Quantitative Evidence from Medical Image Foundation Models
- [437] From Bit To Bedside A Practical Framework For Artificial Intelligence Product Development In Healthcare
- [438] Current State of Community-Driven Radiological AI Deployment in Medical Imaging
- [439] Integration and Implementation Strategies for AI Algorithm Deployment with Smart Routing Rules and Workflow Management
- [440] FUTURE-AI International consensus guideline for trustworthy and deployable artificial intelligence in healthcare
- [441] Autonomous Artificial Intelligence Agents for Clinical Decision Making in Oncology
- [442] Large Language Models as Agents in the Clinic
- [443] A Comprehensive Review of Generative AI in Healthcare
- [444] Virtual imaging trials improved the transparency and reliability of AI systems in COVID-19 imaging
- [445] Accelerating science with human versus alien artificial intelligences
- [446] Complementary artificial intelligence designed to augment human discovery
- [447] Accelerating science with human-aware artificial intelligence
- [448] Human-AI Co-Creation Approach to Find Forever Chemicals Replacements
- [449] Toward Human-AI Co-creation to Accelerate Material Discovery
- [450] Crystal Transformer Self-learning neural language model for Generative and Tinkering Design of Materials
- [451] Materials Transformers Language Models for Generative Materials Design a benchmark study
- [452] LLaMP Large Language Model Made Powerful for High-fidelity Materials Knowledge Retrieval and Distillation
- [453] MatChat A Large Language Model and Application Service Platform for Materials Science
- [454] 1.5 million materials narratives generated by chatbots
- [455] Materials Expert-Artificial Intelligence for Materials Discovery
- [456] Agent-based Learning of Materials Datasets from Scientific Literature
- [457] A Sober Look at LLMs for Material Discovery Are They Actually Good for Bayesian Optimization Over Molecules
- [458] Navigating the Complexity of Generative AI Adoption in Software Engineering
- [459] The Impact of AI Tool on Engineering at ANZ Bank An Empirical Study on GitHub Copilot within Corporate Environment
- [460] Rocks Coding, Not Development—A Human-Centric, Experimental Evaluation of LLM-Supported SE Tasks

- [461] Beyond Code Generation An Observational Study of ChatGPT Usage in Software Engineering Practice
- [462] Can AI Serve as a Substitute for Human Subjects in Software Engineering Research
- [463] Can Large Language Models Serve as Data Analysts A Multi-Agent Assisted Approach for Qualitative Data Analysis
- [464] Exploring Qualitative Research Using LLMs
- [465] From Bytes to Biases Investigating the Cultural Self-Perception of Large Language Models
- [466] Human-Machine Teaming for UAVs An Experimentation Platform
- [467] Improving the State of the Art for Training Human-AI Teams Technical Report 1 – Results of Subject-Matter Expert Knowledge Elicitation Survey
- [468] AI Potentiality and Awareness A Position Paper from the Perspective of Human-AI Teaming in Cybersecurity
- [469] The Promise and Peril of Artificial Intelligence – Violet Teaming Offers a Balanced Path Forward
- [470] A Red Teaming Framework for Securing AI in Maritime Autonomous Systems
- [471] Workload Estimation for Unknown Tasks A Survey of Machine Learning Under Distribution Shift
- [472] Trust Repairing for Human-Swarm Cooperation in Dynamic Task Response
- [473] Trust Aware Emergency Response for A Resilient Human-Swarm Cooperative System
- [474] Detection of Stealthy Adversaries for Networked Unmanned Aerial Vehicles
- [475] Tools and methods for Human-Autonomy Teaming Contributions to cognitive state monitoring and system adaptation
- [476] A Work-Centered Approach for Cyber-Physical-Social System Design Applications in Aerospace Industrial Inspection
- [477] Questioning the impact of AI and interdisciplinarity in science Lessons from COVID-19
- [478] Covidia COVID-19 Interdisciplinary Academic Knowledge Graph
- [479] Visualising COVID-19 Research
- [480] ADVISE AI-accelerated Design of Evidence Synthesis for Global Development
- [481] Is Coronavirus-Related Research Becoming More Interdisciplinary A Perspective of Co-occurrence Analysis and Diversity Measure of Scientific Articles
- [482] Grand challenges and emergent modes of convergence science
- [483] Expertise diversity of teams predicts originality and long-term impact in science and technology
- [484] Sails and Anchors The Complementarity of Exploratory and Exploitative Scientists in Knowledge Creation
- [485] A Collaborative Approach to the Analysis of the COVID-19 Response in Africa

- [486] AI Models Close to your Chest Robust Federated Learning Strategies for Multi-site CT
- [487] It is there, and you need it, so why do you not use it Achieving better adoption of AI systems by domain experts, in the case study of natural science research
- [488] Trust in AI Interpretability is not necessary or sufficient, while black-box interaction is necessary and sufficient
- [489] Multi-stakeholder Perspective on Responsible Artificial Intelligence and Acceptability in Education
- [490] A multidomain relational framework to guide institutional AI research and adoption
- [491] Ask Not What AI Can Do, But What AI Should Do Towards a Framework of Task Delegability
- [492] SAI Bench A Structural Interpretation of AI for Science Through Benchmarks
- [493] On the Measure of Intelligence
- [494] Predicting the Future of AI with AI High-quality link prediction in an exponentially growing knowledge network
- [495] A Rigorous Uncertainty-Aware Quantification Framework Is Essential for Reproducible and Replicable Machine Learning Workflows
- [496] Bridging the Gap Between Explainable AI and Uncertainty Quantification to Enhance Trustability
- [497] The Harmony Index a Utilitarian Metric for Measuring Effectiveness in Mixed-Skill Teams
- [498] OpenHEXAI An Open-Source Framework for Human-Centered Evaluation of Explainable Machine Learning
- [499] A Case for Backward Compatibility for Human-AI Teams
- [500] Metrology for AI From Benchmarks to Instruments
- [501] DICES Dataset Diversity in Conversational AI Evaluation for Safety
- [502] Automated Scientific Discovery From Equation Discovery to Autonomous Discovery Systems
- [503] Leveraging Multi-AI Agents for Cross-Domain Knowledge Discovery
- [504] Automated and Autonomous Experiment in Electron and Scanning Probe Microscopy
- [505] AI-GAs AI-generating algorithms, an alternate paradigm for producing general artificial intelligence
- [506] AI Autonomy Self-Initiated Open-World Continual Learning and Adaptation
- [507] Towards Self-constructive Artificial Intelligence Algorithmic basis (Part I)
- [508] Towards Autonomous Hypothesis Verification via Language Models with Minimal Guidance
- [509] Data-driven Discovery with Large Generative Models
- [510] Balancing Autonomy and Alignment A Multi-Dimensional Taxonomy for Autonomous LLM-powered Multi-Agent Architectures
- [511] Simulation Intelligence Towards a New Generation of Scientific Methods

- [512] Late-Binding Scholarship in the Age of AI Navigating Legal and Normative Challenges of a New Form of Knowledge Production
- [513] A Computational Inflection for Scientific Discovery
- [514] The Semantic Reader Project Augmenting Scholarly Documents through AI-Powered Interactive Reading Interfaces
- [515] ModelHub.AI Dissemination Platform for Deep Learning Models
- [516] Scientometric engineering Exploring citation dynamics via arXiv eprints
- [517] Back to the Future – Sequential Alignment of Text Representations
- [518] Artificial Intelligence Ethics An Inclusive Global Discourse
- [519] On the Opportunities of Green Computing A Survey
- [520] Recommendations for public action towards sustainable generative AI systems
- [521] Towards Environmentally Equitable AI via Geographical Load Balancing
- [522] Sustainable Software Ecosystems for Open Science
- [523] Voluminous yet Vacuous Semantic Capital in an Age of Large Language Models
- [524] Toward Responsible AI Use Considerations for Sustainability Impact Assessment
- [525] Thermodynamic Computing
- [526] Intent-aligned AI systems deplete human agency the need for agency foundations research in AI safety
- [527] Who Wrote this How Smart Replies Impact Language and Agency in the Workplace
- [528] The Rise of the AI Co-Pilot Lessons for Design from Aviation and Beyond
- [529] Needs and Artificial Intelligence
- [530] Aligning Artificial Intelligence with Humans through Public Policy
- [531] Automatic Authorities Power and AI
- [532] Reflective Hybrid Intelligence for Meaningful Human Control in Decision-Support Systems
- [533] A Decentralized Approach towards Responsible AI in Social Ecosystems
- [534] Decolonial AI Decolonial Theory as Sociotechnical Foresight in Artificial Intelligence
- [535] A Human Rights-Based Approach to Responsible AI
- [536] Enhanced well-being assessment as basis for the practical implementation of ethical and rights-based normative principles for AI