

# Retrieval-Augmented Generation for Large Language Models: A Survey

Retrieval-augmented generation (RAG) connects large language models to external evidence by retrieving documents, passages, memories, database records, or multimodal items and conditioning generation on the selected context. This survey reviews RAG for LLMs as a full pipeline rather than a single retrieval call. We synthesize an arXiv-derived corpus of 300 papers collected during this run and organize the literature around corpus construction, retrievers, query rewriting, reranking, evidence fusion, generator conditioning, attribution, evaluation, adaptation, domain-specific systems, and safety. The survey argues that RAG is best understood as a controllable interface between parametric knowledge and non-parametric evidence. Its strengths are updateability, provenance, and domain grounding; its weaknesses are retrieval misses, distractor sensitivity, stale corpora, context overload, and difficult end-to-end evaluation. We close with prospects for adaptive, auditable, and privacy-aware RAG systems.

## Keywords:

large language models; retrieval-augmented generation; dense retrieval; grounding; factuality; evaluation

## 1 Introduction

Retrieval-augmented generation (RAG) has become one of the central patterns for making large language models more grounded, updateable, and domain-aware. Instead of relying only on parametric knowledge learned during pretraining, a RAG system retrieves external evidence and conditions the model’s response on that evidence. This design is attractive because the corpus can be updated, inspected, filtered, cited, and specialized for a task without retraining the entire model [32, 90, 92, 102, 131, 267, 293, 299]. At the same time, RAG introduces a new systems problem: the final answer depends on the quality of the corpus, retriever, ranker, context assembler, generator, and verifier.

RAG is often described as retrieve-then-generate, but modern systems are more varied. Some rewrite the user’s query, perform iterative or multi-hop search, rerank passages, compress evidence, ask the generator to cite sources, or use feedback to decide whether more retrieval is needed [32, 81, 92, 102, 131, 154, 267, 291]. Others blend retrieval with long-context mod-

eling, agent memory, tool use, structured databases, knowledge graphs, or multimodal stores. The design space is therefore much broader than a vector database attached to a chatbot.

This survey treats RAG as an evidence interface between language models and external knowledge. A useful RAG system must answer four questions: what evidence is available, what evidence is relevant, how evidence should be represented in the prompt or model state, and how the generated answer should be checked against the evidence. These questions connect information retrieval, natural-language generation, databases, human-computer interaction, evaluation, and safety.

## 2 Background

### 2.1 Parametric and Non-Parametric Knowledge

LLMs store a large amount of knowledge in parameters, but that knowledge is difficult to update, inspect,

or attribute. RAG adds non-parametric knowledge in the form of documents, indexes, tables, memories, graphs, or other external stores. The result is a hybrid system: the model contributes linguistic generalization and reasoning, while retrieval contributes fresh evidence and domain specificity. The balance is delicate. If retrieval misses key evidence, generation may hallucinate; if retrieval returns distractors, generation may overfit to irrelevant context; if too much evidence is inserted, the model may fail to use the right passages.

## 2.2 The RAG Pipeline

A typical RAG pipeline includes ingestion, chunking, indexing, query construction, retrieval, reranking, prompt assembly, generation, citation, and evaluation. Each stage encodes assumptions. Chunking determines what semantic units can be retrieved; embeddings determine similarity; query rewriting determines recall; reranking determines precision; and prompt assembly determines what the model can attend to. End-to-end quality is bounded by the weakest stage, which is why many recent studies evaluate full pipelines rather than isolated retrievers [15, 32, 93, 102, 154, 244, 285, 299].

## 2.3 Grounding and Attribution

Grounding requires that generated claims be supported by retrieved evidence. Attribution goes further by exposing which source supports which claim. These goals are important because RAG is often deployed in settings where users expect traceability: science, medicine, law, finance, enterprise search, and technical support. Yet citations can be misleading when the answer cites a retrieved passage that is topically related but does not actually entail the claim. Grounded RAG therefore needs claim-level verification, evidence selection, and calibration [34, 92, 93, 102, 131, 267, 288, 299].

# 3 Methods

## 3.1 Search Protocol and Corpus Construction

The corpus for this survey was generated during this run from the arXiv structured API. Queries combined the topic with retrieval-augmented generation, RAG, dense retrieval, document retrieval, question answering, grounding, hallucination, query rewriting, reranking, multi-hop retrieval, evaluation, benchmarks, long context, memory, multimodal re-

trieval, and domain-specific RAG terms. Records were retained when title or abstract evidence connected LLMs or language models with retrieval, external evidence, generation, grounding, evaluation, training, or safety. Duplicates were removed by arXiv identifier, and the final bibliography cites 300 unique scholarly preprints.

## 3.2 Taxonomy

We organize RAG methods along six axes. The first is *corpus design*: document sources, freshness, chunking, metadata, and access control. The second is *retrieval*: sparse, dense, hybrid, graph-based, structured, or multimodal search. The third is *evidence selection*: query rewriting, decomposition, reranking, filtering, and compression. The fourth is *generation*: how retrieved evidence is fused, cited, and reconciled with model knowledge. The fifth is *learning*: whether retrievers, rerankers, generators, or verifiers are trained jointly or separately. The sixth is *governance*: safety, privacy, provenance, and auditability.

## 3.3 Retrievers, Indexes, and Corpus Construction

Retrieval quality starts before any query is issued. The corpus must be cleaned, segmented, enriched with metadata, and indexed in a way that matches expected questions. Dense retrievers improve semantic recall, sparse retrievers preserve lexical precision, and hybrid systems exploit both. Corpus construction also shapes failure modes: overly small chunks lose context, overly large chunks reduce precision, stale indexes return outdated information, and missing metadata makes filtering difficult. These issues motivate work on adaptive chunking, hierarchical retrieval, hybrid search, and domain-specific indexing [15, 32, 93, 102, 154, 244, 285, 299].

## 3.4 Query Rewriting, Multi-Hop Retrieval, and Reranking

User questions are rarely optimal retrieval queries. Query rewriting expands, decomposes, or clarifies the information need before retrieval. Multi-hop RAG retrieves evidence iteratively, where earlier passages guide later queries. Rerankers then estimate which candidate passages are most useful for generation. These components are crucial because the generator is sensitive to both missing evidence and distractors. The best RAG systems treat evidence selection as an active process rather than a static top- $k$  lookup [32, 81, 92, 102, 131, 154, 267, 291].

### 3.5 Generation, Fusion, and Answer Control

After retrieval, the model must fuse evidence with its own prior knowledge. Generation strategies vary from simple prompting to structured answer plans, citation-aware decoding, chain-of-thought over evidence, and verifier-guided revision. A central tension is that the model may know facts not present in the retrieved passages; using such knowledge can improve completeness but weaken attribution. RAG systems therefore need policies for abstention, conflict handling, source preference, and answer calibration [32, 90, 92, 102, 131, 267, 293, 299].

### 3.6 Training and Optimization

RAG components can be optimized separately or jointly. Retriever training improves recall, reranker training improves precision, generator tuning improves evidence use, and preference or reinforcement signals can reward faithful answers. End-to-end optimization is difficult because retrieval decisions are discrete, corpora evolve, and answer quality depends on both evidence and generation. Recent work explores synthetic data, self-training, contrastive learning, feedback loops, and verifier-based objectives [32, 34, 76, 93, 102, 131, 154, 299].

### 3.7 Evaluation and Diagnostics

RAG evaluation should distinguish retrieval quality from generation quality. A system may retrieve the right passage and still hallucinate, or retrieve weak evidence and answer correctly from parametric memory. Useful metrics include recall, answer correctness, citation faithfulness, context precision, robustness to distractors, abstention quality, latency, and cost. Benchmarks increasingly include multi-hop questions, long-context settings, domain corpora, and adversarial retrieval conditions [32, 76, 90, 102, 131, 267, 293, 299].

### 3.8 Domain-Specific and Multimodal RAG

Domain RAG systems adapt the pipeline to specialized evidence. Medical, legal, scientific, financial, code, table, graph, and multimodal settings differ in terminology, provenance requirements, latency, and risk. For example, biomedical RAG may prioritize controlled vocabularies and evidence hierarchy, while code RAG needs repository structure and executable verification. Multimodal RAG adds image, video, audio, and layout-aware retrieval, expanding the evidence interface beyond plain text.

### 3.9 Representative Literature Map

The following map cites the full retained corpus by theme, making the survey coverage explicit.

*Retrievers, indexes, and corpus construction* The retrievers, indexes, and corpus construction cluster contains work that treats LLMs as controllers of typed resources, stateful environments, or evaluation tasks [32, 68, 74, 106, 190, 191, 203, 204, 210]. A second strand in this cluster studies how such systems represent intermediate state, recover from failed calls, and compare tool-mediated behavior across datasets [61, 147, 171, 180, 213, 264].

*Augmentation architectures and generation control* The augmentation architectures and generation control cluster contains work that treats LLMs as controllers of typed resources, stateful environments, or evaluation tasks [76, 90, 92, 102, 131, 154, 267, 293, 299]. A second strand in this cluster studies how such systems represent intermediate state, recover from failed calls, and compare tool-mediated behavior across datasets [11, 34, 93, 165, 197, 244, 252, 254, 288]. Additional retained studies broaden the cluster with variants of interfaces, domains, training signals, and deployment assumptions [81, 107, 152, 159, 209, 221, 247, 285, 291]. Additional retained studies broaden the cluster with variants of interfaces, domains, training signals, and deployment assumptions [4, 12, 15, 49, 114, 115, 199, 217, 263]. Additional retained studies broaden the cluster with variants of interfaces, domains, training signals, and deployment assumptions [2, 30, 31, 60, 158, 200, 212, 266, 277]. Additional retained studies broaden the cluster with variants of interfaces, domains, training signals, and deployment assumptions [9, 62, 127, 142, 153, 214, 259, 298, 300]. Additional retained studies broaden the cluster with variants of interfaces, domains, training signals, and deployment assumptions [1, 13, 94, 129, 161, 219, 222, 240, 283]. Additional retained studies broaden the cluster with variants of interfaces, domains, training signals, and deployment assumptions [14, 85, 118, 120, 135, 193, 265, 279, 287]. Additional retained studies broaden the cluster with variants of interfaces, domains, training signals, and deployment assumptions [26, 113, 136, 145, 223, 224, 230, 237, 272]. Additional retained studies broaden the cluster with variants of interfaces, domains, training signals, and deployment assumptions [28, 66, 108, 116, 133, 179, 211, 248, 261]. Additional retained studies broaden the cluster with variants of interfaces, domains, training signals, and deployment assumptions [27, 37, 56, 70, 72, 103, 218, 243, 289]. Additional retained studies broaden the cluster with

variants of interfaces, domains, training signals, and deployment assumptions [20, 21, 58, 77, 155, 163, 189, 257, 297]. Additional retained studies broaden the cluster with variants of interfaces, domains, training signals, and deployment assumptions [3, 7, 59, 96, 101, 104, 141, 229, 286]. Additional retained studies broaden the cluster with variants of interfaces, domains, training signals, and deployment assumptions [6, 42, 69, 97, 140, 183, 228, 238, 250]. Additional retained studies broaden the cluster with variants of interfaces, domains, training signals, and deployment assumptions [54, 64, 117, 121, 168, 170, 182, 227, 258]. Additional retained studies broaden the cluster with variants of interfaces, domains, training signals, and deployment assumptions [65, 79, 91, 119, 125, 172, 231, 235, 274]. Additional retained studies broaden the cluster with variants of interfaces, domains, training signals, and deployment assumptions [48, 73, 84, 148, 162, 192, 256, 276, 280]. Additional retained studies broaden the cluster with variants of interfaces, domains, training signals, and deployment assumptions [110, 128, 130, 143, 201, 215, 225, 251, 273]. Additional retained studies broaden the cluster with variants of interfaces, domains, training signals, and deployment assumptions [35, 38, 41, 55, 150, 176, 239, 292, 295]. Additional retained studies broaden the cluster with variants of interfaces, domains, training signals, and deployment assumptions [8, 53, 86, 100, 126, 177, 198, 226, 241]. Additional retained studies broaden the cluster with variants of interfaces, domains, training signals, and deployment assumptions [18, 33, 44, 109, 124, 164, 169, 234, 246]. Additional retained studies broaden the cluster with variants of interfaces, domains, training signals, and deployment assumptions [22, 43, 47, 87, 95, 98, 146, 262, 296]. Additional retained studies broaden the cluster with variants of interfaces, domains, training signals, and deployment assumptions [24, 29, 112, 134, 137, 157, 195, 206, 290]. Additional retained studies broaden the cluster with variants of interfaces, domains, training signals, and deployment assumptions [19, 67, 82, 89, 166, 173, 220, 233, 275]. Additional retained studies broaden the cluster with variants of interfaces, domains, training signals, and deployment assumptions [17, 23, 39, 57, 149, 156, 186, 196, 282]. Additional retained studies broaden the cluster with variants of interfaces, domains, training signals, and deployment assumptions [40, 63, 99, 122, 151, 178, 232, 236, 260]. Additional retained studies broaden the cluster with variants of interfaces, domains, training signals, and deployment assumptions [25, 71, 105, 144, 185, 187, 205, 216, 270]. Additional retained studies broaden the cluster with variants of interfaces, domains, training signals, and deployment assumptions [5, 10, 36, 46, 50, 83, 139, 184, 242]. Ad-

ditional retained studies broaden the cluster with variants of interfaces, domains, training signals, and deployment assumptions [78, 194, 202].

#### *Query rewriting, reranking, and evidence selection*

The query rewriting, reranking, and evidence selection cluster contains work that treats LLMs as controllers of typed resources, stateful environments, or evaluation tasks [52, 160].

#### *Training, adaptation, and optimization*

The training, adaptation, and optimization cluster contains work that treats LLMs as controllers of typed resources, stateful environments, or evaluation tasks [188, 208].

#### *Benchmarks, evaluation, and diagnostics*

The benchmarks, evaluation, and diagnostics cluster contains work that treats LLMs as controllers of typed resources, stateful environments, or evaluation tasks [16, 80, 167, 174, 175, 253, 268, 281, 284]. A second strand in this cluster studies how such systems represent intermediate state, recover from failed calls, and compare tool-mediated behavior across datasets [45, 181, 207, 245, 249, 255, 269, 271, 278]. Additional retained studies broaden the cluster with variants of interfaces, domains, training signals, and deployment assumptions [51, 75, 88, 111, 123, 132, 294].

#### *Domain-specific and multimodal rag*

The domain-specific and multimodal rag cluster contains work that treats LLMs as controllers of typed resources, stateful environments, or evaluation tasks [138].

## 4 Challenges and Prospects

### 4.1 Retrieval Misses and Distractor Sensitivity

RAG systems can fail silently when they retrieve plausible but insufficient evidence. Increasing  $k$  improves recall but also increases distractors and context cost. Future systems need adaptive retrieval policies that decide when to search again, when to ask for clarification, and when to abstain.

### 4.2 Faithfulness and Citation Quality

Citation is not the same as support. A generated claim may cite a passage that contains related words but does not entail the claim. Better RAG systems will need claim-level attribution, entailment checks, and interfaces that expose uncertainty rather than



presenting citations as decorative references [34, 92, 93, 102, 131, 267, 288, 299].

### 4.3 Freshness, Privacy, and Security

External corpora create governance obligations. Documents may become stale, confidential, poisoned, or inappropriate for a user’s permission scope. RAG systems therefore need corpus versioning, access control, privacy filters, and monitoring for prompt injection and retrieval poisoning [4, 30, 32, 62, 212, 244, 266, 281].

### 4.4 Long-Context RAG and Memory

Long-context models change RAG but do not remove it. They can ingest more evidence, but retrieval remains valuable for freshness, filtering, provenance, and cost control. The future likely belongs to hybrid systems that combine long windows, hierarchical retrieval, memory stores, and evidence compression.

## 5 Conclusion

RAG has evolved from a simple retrieval-plus-generation recipe into a broad architecture for grounding LLMs in external evidence. This survey reviewed corpus construction, retrievers, rerankers, evidence fusion, training, evaluation, domain-specific systems, and safety. The central lesson is that RAG quality is pipeline quality: reliable answers require good evidence selection, faithful generation, transparent attribution, and governed corpora. Future RAG systems should be adaptive, auditable, privacy-aware, and explicit about uncertainty.

## References

- [1] Ali Abedi, Charlene H. Chu, and Shehroz S. Khan. Retrieval-Augmented Large Language Models for Evidence-Informed Guidance on Cannabidiol Use in Older Adults. *arXiv preprint arXiv:2604.09548*, 2026. URL <https://arxiv.org/abs/2604.09548>.
- [2] Parker Addison, Minh-Tuan H. Nguyen, Tomislav Medan, Jinali Shah, Mohammad T. Manzari, Brendan McElrone, Laksh Lalwani, Aboli More, Smita Sharma, Holger R. Roth, Isaac Yang, Chester Chen, Daguang Xu, Yan Cheng, Andrew Feng, and Ziyue Xu. C-FedRAG: A Confidential Federated Retrieval-Augmented Generation System. *arXiv preprint arXiv:2412.13163*, 2024. URL <https://arxiv.org/abs/2412.13163>.
- [3] Anum Afzal, Juraj Vladika, Gentrit Fazlija, Andrei Staradubets, and Florian Matthes. Towards Optimizing a Retrieval Augmented Generation using Large Language Model on Academic Data. *arXiv preprint arXiv:2411.08438*, 2024. URL <https://arxiv.org/abs/2411.08438>.
- [4] Khandakar Ashrafi Akbar, Md Nahiyen Uddin, Latifur Khan, Trayce Hockstad, Mizanur Rahman, Mashrur Chowdhury, and Bhavani Thuraisingham. Retrieval Augmented Generation-based Large Language Models for Bridging Transportation Cybersecurity Legal Knowledge Gaps. *arXiv preprint arXiv:2505.18426*, 2025. URL <https://arxiv.org/abs/2505.18426>.
- [5] Nurshat Fateh Ali, Md. Mahdi Mohtasim, Shakil Mosharrof, and T. Gopi Krishna. Automated Literature Review Using NLP Techniques and LLM-Based Retrieval-Augmented Generation. *arXiv preprint arXiv:2411.18583*, 2024. URL <https://arxiv.org/abs/2411.18583>.
- [6] Bang An, Shiyue Zhang, and Mark Dredze. RAG LLMs are Not Safer: A Safety Analysis of Retrieval-Augmented Generation for Large Language Models. *arXiv preprint arXiv:2504.18041*, 2025. URL <https://arxiv.org/abs/2504.18041>.
- [7] Raviteja Anantha, Tharun Bethi, Danil Vodianik, and Srinivas Chappidi. Context Tuning for Retrieval Augmented Generation. *arXiv preprint arXiv:2312.05708*, 2023. URL <https://arxiv.org/abs/2312.05708>.
- [8] Mohammad Aghajani Asl and Behrooz Minaei Bidgoli. FARSIQA: Faithful and Advanced RAG System for Islamic Question Answering. *arXiv preprint arXiv:2510.25621*, 2025. URL <https://arxiv.org/abs/2510.25621>.
- [9] Mohammad Aghajani Asl, Majid Asgari-Bidhendi, and Behrooz Minaei-Bidgoli. FAIR-RAG: Faithful Adaptive Iterative Refinement for Retrieval-Augmented Generation. *arXiv preprint arXiv:2510.22344*, 2025. URL <https://arxiv.org/abs/2510.22344>.
- [10] Byeonghun Bang, Jongsuk Yoon, Dong-Jin Chang, Seho Park, and Yong Oh Lee. Retrieval Augmented Large Language Model System for Comprehensive Drug Contraindications. *arXiv preprint arXiv:2508.06145*, 2025. URL <https://arxiv.org/abs/2508.06145>.
- [11] Prakhar Bansal and Shivangi Agarwal. Beyond the Parameters: A Technical Survey of Contextual Enrichment in Large Language Models: From In-Context Prompting to Causal Retrieval-Augmented Generation. *arXiv preprint arXiv:2604.03174*, 2026. URL <https://arxiv.org/abs/2604.03174>.
- [12] Andreas Baumann and Peter Eberhard. Experiments with Large Language Models on Retrieval-Augmented Generation for Closed-Source Simula-

- tion Software. *arXiv preprint arXiv:2502.03916*, 2025. URL <https://arxiv.org/abs/2502.03916>.
- [13] David Beauchemin and Richard Khoury. Benchmarking Large Language Models for Quebec Insurance: From Closed-Book to Retrieval-Augmented Generation. *arXiv preprint arXiv:2603.07825*, 2026. URL <https://arxiv.org/abs/2603.07825>.
  - [14] Patrice Bechard and Orlando Marquez Ayala. Reducing hallucination in structured outputs via Retrieval-Augmented Generation. *arXiv preprint arXiv:2404.08189*, 2024. URL <https://arxiv.org/abs/2404.08189>.
  - [15] Ruben Belo, Marta Guimaraes, and Claudia Soares. A Systematic Evaluation of Retrieval-Augmented Generation and Language Models for Space Operations. *arXiv preprint arXiv:2605.27444*, 2026. URL <https://arxiv.org/abs/2605.27444>.
  - [16] Gagan Bhatia, Hamdy Mubarak, Mustafa Jarrar, George Mikros, Fadi Zaraket, Mahmoud Alhirthani, Mutaz Al-Khatib, Logan Cochrane, Kareem Darwish, Rashid Yahiaoui, and Firoj Alam. From RAG to Agentic RAG for Faithful Islamic Question Answering. *arXiv preprint arXiv:2601.07528*, 2026. URL <https://arxiv.org/abs/2601.07528>.
  - [17] Kushagra Bhushan, Yatin Nandwani, Dinesh Khandelwal, Sonam Gupta, Gaurav Pandey, Dinesh Raghu, and Sachindra Joshi. Systematic Knowledge Injection into Large Language Models via Diverse Augmentation for Domain-Specific RAG. *arXiv preprint arXiv:2502.08356*, 2025. URL <https://arxiv.org/abs/2502.08356>.
  - [18] Yunus Serhat Bicakci, Joseph Shingleton, and Anahid Basiri. Street-Level Geolocalization Using Multimodal Large Language Models and Retrieval-Augmented Generation. *arXiv preprint arXiv:2509.01341*, 2025. URL <https://arxiv.org/abs/2509.01341>.
  - [19] Simon Binz, Heydar Soudani, and Faegheh Hasiibi. Uncertainty Quantification for Multimodal Retrieval Augmented Generation. *arXiv preprint arXiv:2605.29956*, 2026. URL <https://arxiv.org/abs/2605.29956>.
  - [20] Nitin Aravind Birur, Tanay Baswa, Divyanshu Kumar, Jatan Loya, Sahil Agarwal, and Prashanth Harshangi. VERA: Validation and Enhancement for Retrieval Augmented systems. *arXiv preprint arXiv:2409.15364*, 2024. URL <https://arxiv.org/abs/2409.15364>.
  - [21] Arnabh Borah, Md Tanvirul Alam, and Nidhi Rastogi. Adapting Large Language Models to Emerging Cybersecurity using Retrieval Augmented Generation. *arXiv preprint arXiv:2510.27080*, 2025. URL <https://arxiv.org/abs/2510.27080>.
  - [22] Andrei-Laurentiu Bornea, Fadhel Ayed, Antonio De Domenico, Nicola Piovesan, and Ali Maatouk. Telco-RAG: Navigating the Challenges of Retrieval-Augmented Language Models for Telecommunications. *arXiv preprint arXiv:2404.15939*, 2024. URL <https://arxiv.org/abs/2404.15939>.
  - [23] Roel Boumans, Lisa Cramer, Sascha van de Poll, and Henria Vermeulen. A Feasibility Study on Usability and Trust among Population Groups of a Medical Avatar Supported by Large Language Models with Retrieval Augmented Generation. *arXiv preprint arXiv:2510.15531*, 2025. URL <https://arxiv.org/abs/2510.15531>.
  - [24] Tilmann Bruckhaus. RAG Does Not Work for Enterprises. *arXiv preprint arXiv:2406.04369*, 2024. URL <https://arxiv.org/abs/2406.04369>.
  - [25] Tommaso Mario Buonocore and Enea Parimbelli. RAR: Setting Knowledge Tripwires for Retrieval Augmented Rejection. *arXiv preprint arXiv:2505.13581*, 2025. URL <https://arxiv.org/abs/2505.13581>.
  - [26] Yuzheng Cai, Zhenyue Guo, Yiwen Pei, Wanrui Bian, and Weiguo Zheng. SimGRAG: Leveraging Similar Subgraphs for Knowledge Graphs Driven Retrieval-Augmented Generation. *arXiv preprint arXiv:2412.15272*, 2024. URL <https://arxiv.org/abs/2412.15272>.
  - [27] Lang Cao, Qingyu Chen, and Yue Guo. EHR-RAG: Bridging Long-Horizon Structured Electronic Health Records and Large Language Models via Enhanced Retrieval-Augmented Generation. *arXiv preprint arXiv:2601.21340*, 2026. URL <https://arxiv.org/abs/2601.21340>.
  - [28] Mario Ceresa, Lorenzo Bertolini, Valentin Comte, Nicholas Spadaro, Barbara Raffael, Brigitte Tous-saint, Sergio Consoli, Amalia Munoz Pineiro, Alex Patak, Maddalena Querci, and Tobias Wiesenthal. Retrieval Augmented Generation Evaluation for Health Documents. *arXiv preprint arXiv:2505.04680*, 2025. URL <https://arxiv.org/abs/2505.04680>.
  - [29] Diego Ezequiel Cervera. Expert Mind: A Retrieval-Augmented Architecture for Expert Knowledge Preservation in the Energy Sector. *arXiv preprint arXiv:2603.14541*, 2026. URL <https://arxiv.org/abs/2603.14541>.
  - [30] Abhijit Chakraborty, Chahana Dahal, and Vivek Gupta. Federated Retrieval-Augmented Generation: A Systematic Mapping Study. *arXiv preprint arXiv:2505.18906*, 2025. URL <https://arxiv.org/abs/2505.18906>.
  - [31] Chi-Min Chan, Chunpu Xu, Ruibin Yuan, Hongyin Luo, Wei Xue, Yike Guo, and Jie Fu. RQ-RAG: Learning to Refine Queries for Retrieval Augmented Generation. *arXiv preprint arXiv:2404.00610*, 2024. URL <https://arxiv.org/abs/2404.00610>.

- [32] Ryan Sze-Yin Chan, Federico Nanni, Tomas Lazauskas, Rosie Wood, Penelope Yong, Lionel Tarassenko, Mark Girolami, James Geddes, and Andrew Duncan. Retrieval-augmented reasoning with lean language models. *arXiv preprint arXiv:2508.11386*, 2025. URL <https://arxiv.org/abs/2508.11386>.
- [33] Achuth Chandrasekhar, Omid Barati Farimani, Olabode T. Ajenifujah, Janghoon Ock, and Amir Barati Farimani. NANOGPT: A Query-Driven Large Language Model Retrieval-Augmented Generation System for Nanotechnology Research. *arXiv preprint arXiv:2502.20541*, 2025. URL <https://arxiv.org/abs/2502.20541>.
- [34] Chia-Yuan Chang, Zhimeng Jiang, Vineeth Rakesh, Menghai Pan, Chin-Chia Michael Yeh, Guanchu Wang, Mingzhi Hu, Zhichao Xu, Yan Zheng, Mahashweta Das, and Na Zou. MAIN-RAG: Multi-Agent Filtering Retrieval-Augmented Generation. *arXiv preprint arXiv:2501.00332*, 2024. URL <https://arxiv.org/abs/2501.00332>.
- [35] Jiawei Chen, Hongyu Lin, Xianpei Han, and Le Sun. Benchmarking Large Language Models in Retrieval-Augmented Generation. *arXiv preprint arXiv:2309.01431*, 2023. URL <https://arxiv.org/abs/2309.01431>.
- [36] Sibe Chen, Ju Fan, Bin Wu, Nan Tang, Chao Deng, Pengyi Wang, Ye Li, Jian Tan, Feifei Li, Jingren Zhou, and Xiaoyong Du. Automatic Database Configuration Debugging using Retrieval-Augmented Language Models. *arXiv preprint arXiv:2412.07548*, 2024. URL <https://arxiv.org/abs/2412.07548>.
- [37] Xin Chen, Saili Uday Gadgil, and Jiarong Qiu. Coordinated Semantic Alignment and Evidence Constraints for Retrieval-Augmented Generation with Large Language Models. *arXiv preprint arXiv:2603.04647*, 2026. URL <https://arxiv.org/abs/2603.04647>.
- [38] Yifu Chen, Shengpeng Ji, Haoxiao Wang, Ziqing Wang, Siyu Chen, Jinzheng He, Jin Xu, and Zhou Zhao. WavRAG: Audio-Integrated Retrieval Augmented Generation for Spoken Dialogue Models. *arXiv preprint arXiv:2502.14727*, 2025. URL <https://arxiv.org/abs/2502.14727>.
- [39] Zhe Chen, Yusheng Liao, Shuyang Jiang, Pingjie Wang, Yiqiu Guo, Yanfeng Wang, and Yu Wang. Towards Omni-RAG: Comprehensive Retrieval-Augmented Generation for Large Language Models in Medical Applications. *arXiv preprint arXiv:2501.02460*, 2025. URL <https://arxiv.org/abs/2501.02460>.
- [40] Zhuo Chen, Jiawei Liu, Haotan Liu, Qikai Cheng, Fan Zhang, Wei Lu, and Xiaozhong Liu. Black-Box Opinion Manipulation Attacks to Retrieval-Augmented Generation of Large Language Models. *arXiv preprint arXiv:2407.13757*, 2024. URL <https://arxiv.org/abs/2407.13757>.
- [41] Pengzhou Cheng, Yidong Ding, Tianjie Ju, Zongru Wu, Wei Du, Ping Yi, Zhuosheng Zhang, and Gongshen Liu. TrojanRAG: Retrieval-Augmented Generation Can Be Backdoor Driver in Large Language Models. *arXiv preprint arXiv:2405.13401*, 2024. URL <https://arxiv.org/abs/2405.13401>.
- [42] Eunseong Choi, June Park, Hyeri Lee, and Jongwuk Lee. Conflict-Aware Soft Prompting for Retrieval-Augmented Generation. *arXiv preprint arXiv:2508.15253*, 2025. URL <https://arxiv.org/abs/2508.15253>.
- [43] Cody Clop and Yannick Teglia. Backdoored Retrievers for Prompt Injection Attacks on Retrieval Augmented Generation of Large Language Models. *arXiv preprint arXiv:2410.14479*, 2024. URL <https://arxiv.org/abs/2410.14479>.
- [44] Xinbang Dai, Huikang Hu, Yuncheng Hua, Jiaqi Li, Yongrui Chen, Rihui Jin, Nan Hu, and Guilin Qi. After Retrieval, Before Generation: Enhancing the Trustworthiness of Large Language Models in Retrieval-Augmented Generation. *arXiv preprint arXiv:2505.17118*, 2025. URL <https://arxiv.org/abs/2505.17118>.
- [45] Souvik Das, Sallam Abualhaija, and Domenico Bianculli. Fine-grained Claim-level RAG Benchmark for Law. *arXiv preprint arXiv:2605.21071*, 2026. URL <https://arxiv.org/abs/2605.21071>.
- [46] Gelei Deng, Yi Liu, Kailong Wang, Yuekang Li, Tianwei Zhang, and Yang Liu. Pandora: Jailbreak GPTs by Retrieval Augmented Generation Poisoning. *arXiv preprint arXiv:2402.08416*, 2024. URL <https://arxiv.org/abs/2402.08416>.
- [47] Jingcheng Deng, Liang Pang, Huawei Shen, and Xueqi Cheng. RegaVAE: A Retrieval-Augmented Gaussian Mixture Variational Auto-Encoder for Language Modeling. *arXiv preprint arXiv:2310.10567*, 2023. URL <https://arxiv.org/abs/2310.10567>.
- [48] Peter Devine. ALoFTRAG: Automatic Local Fine Tuning for Retrieval Augmented Generation. *arXiv preprint arXiv:2501.11929*, 2025. URL <https://arxiv.org/abs/2501.11929>.
- [49] Jialin Dong, Bahare Fatemi, Bryan Perozzi, Lin F. Yang, and Anton Tsitsulin. Don't Forget to Connect! Improving RAG with Graph-based Reranking. *arXiv preprint arXiv:2405.18414*, 2024. URL <https://arxiv.org/abs/2405.18414>.
- [50] Ehsan Doostmohammadi and Marco Kuhlmann. Studying the Role of Input-Neighbor Overlap in Retrieval-Augmented Language Models Training Efficiency. *arXiv preprint arXiv:2505.14309*, 2025. URL <https://arxiv.org/abs/2505.14309>.

- [51] Sauhard Dubey. SciNets: Graph-Constrained Multi-Hop Reasoning for Scientific Literature Synthesis. *arXiv preprint arXiv:2601.09727*, 2025. URL <https://arxiv.org/abs/2601.09727>.
- [52] Matous Eibich, Shivay Nagpal, and Alexander Fred-Ojala. ARAGOG: Advanced RAG Output Grading. *arXiv preprint arXiv:2404.01037*, 2024. URL <https://arxiv.org/abs/2404.01037>.
- [53] Wenqi Fan, Yujuan Ding, Liangbo Ning, Shijie Wang, Hengyun Li, Dawei Yin, Tat-Seng Chua, and Qing Li. A Survey on RAG Meeting LLMs: Towards Retrieval-Augmented Large Language Models. *arXiv preprint arXiv:2405.06211*, 2024. URL <https://arxiv.org/abs/2405.06211>.
- [54] Feiteng Fang, Yuelin Bai, Shiwen Ni, Min Yang, Xiaojun Chen, and Ruifeng Xu. Enhancing Noise Robustness of Retrieval-Augmented Language Models with Adaptive Adversarial Training. *arXiv preprint arXiv:2405.20978*, 2024. URL <https://arxiv.org/abs/2405.20978>.
- [55] Mehrdad Farahani and Richard Johansson. Deciphering the Interplay of Parametric and Non-parametric Memory in Retrieval-augmented Language Models. *arXiv preprint arXiv:2410.05162*, 2024. URL <https://arxiv.org/abs/2410.05162>.
- [56] Masoomali Fatehikia, Ji Kim Lucas, and Sanjay Chawla. T-RAG: Lessons from the LLM Trenches. *arXiv preprint arXiv:2402.07483*, 2024. URL <https://arxiv.org/abs/2402.07483>.
- [57] Reza Fayyazi, Rozhina Taghdimi, and Shanchieh Jay Yang. Advancing TTP Analysis: Harnessing the Power of Large Language Models with Retrieval Augmented Generation. *arXiv preprint arXiv:2401.00280*, 2023. URL <https://arxiv.org/abs/2401.00280>.
- [58] Pietro Ferrazzi, Milica Cvjeticanin, Alessio Piracini, and Davide Giannuzzi. Is Agentic RAG worth it? An experimental comparison of RAG approaches. *arXiv preprint arXiv:2601.07711*, 2026. URL <https://arxiv.org/abs/2601.07711>.
- [59] Zoe Fingleton, Nazanin Siavash, and Armin Moin. Enhancing Large Language Models with Retrieval Augmented Generation for Software Testing and Inspection Automation. *arXiv preprint arXiv:2604.15270*, 2026. URL <https://arxiv.org/abs/2604.15270>.
- [60] Daniel Fleischer, Moshe Berchansky, Moshe Wasserblat, and Peter Izsak. RAG Foundry: A Framework for Enhancing LLMs for Retrieval Augmented Generation. *arXiv preprint arXiv:2408.02545*, 2024. URL <https://arxiv.org/abs/2408.02545>.
- [61] Anoushka Gade and Jorjeta Jetcheva. It’s About Time: Incorporating Temporality in Retrieval Augmented Language Models. *arXiv preprint arXiv:2401.13222*, 2024. URL <https://arxiv.org/abs/2401.13222>.
- [62] Aoran Gan, Hao Yu, Kai Zhang, Qi Liu, Wenyu Yan, Zhenya Huang, Shiwei Tong, and Guoping Hu. Retrieval Augmented Generation Evaluation in the Era of Large Language Models: A Comprehensive Survey. *arXiv preprint arXiv:2504.14891*, 2025. URL <https://arxiv.org/abs/2504.14891>.
- [63] Chunjing Gan, Dan Yang, Binbin Hu, Hanxiao Zhang, Siyuan Li, Ziqi Liu, Yue Shen, Lin Ju, Zhiqiang Zhang, Jinjie Gu, Lei Liang, and Jun Zhou. Similarity is Not All You Need: Endowing Retrieval Augmented Generation with Multi Layered Thoughts. *arXiv preprint arXiv:2405.19893*, 2024. URL <https://arxiv.org/abs/2405.19893>.
- [64] Yunfan Gao, Yun Xiong, Xinyu Gao, Kangxiang Jia, Jinliu Pan, Yuxi Bi, Yi Dai, Jiawei Sun, Meng Wang, and Haofen Wang. Retrieval-Augmented Generation for Large Language Models: A Survey. *arXiv preprint arXiv:2312.10997*, 2023. URL <https://arxiv.org/abs/2312.10997>.
- [65] Yao Ge, Sudeshna Das, Yuting Guo, and Abeed Sarker. Retrieval augmented generation based dynamic prompting for few-shot biomedical named entity recognition using large language models. *arXiv preprint arXiv:2508.06504*, 2025. URL <https://arxiv.org/abs/2508.06504>.
- [66] Zhuohan Ge, Darian Li, Yubo Wang, Nicole Hu, Xinyi Zhu, Haoyang Li, Xin Zhang, Mingtao Zhang, Shihao Qi, Yuming Xu, Han Shi, Chen Jason Zhang, and Qing Li. Survey and Experiments on Mental Disorder Detection via Social Media: From Large Language Models and RAG to Agents. *arXiv preprint arXiv:2504.02800*, 2025. URL <https://arxiv.org/abs/2504.02800>.
- [67] Komal Gilani, Marlo Verket, Christof Peters, Michel Dumontier, Hans-Peter Brunner-La Rocca, and Visara Urovi. CDE-Mapper: Using Retrieval-Augmented Language Models for Linking Clinical Data Elements to Controlled Vocabularies. *arXiv preprint arXiv:2505.04365*, 2025. URL <https://arxiv.org/abs/2505.04365>.
- [68] Larissa Guder, Joao Pedro de Moura, Arthur Accorsi, Gustavo Losch do Amaral, Mauricio Cecilio Magnaguagno, Felipe Meneguzzi, Marcio Sorraglia Pinho, and Dalvan Griebler. RAGE: A Retrieval-Augmented Generation Evaluation Framework. *arXiv preprint arXiv:2605.27445*, 2026. URL <https://arxiv.org/abs/2605.27445>.
- [69] Gauthier Guinet, Behrooz Omidvar-Tehrani, Anoop Deoras, and Laurent Callot. Automated Evaluation of Retrieval-Augmented Language Models with Task-Specific Exam Generation. *arXiv*



- preprint *arXiv:2405.13622*, 2024. URL <https://arxiv.org/abs/2405.13622>.
- [70] Junfeng Guo, Yiming Li, Ruibo Chen, Yihan Wu, Chenxi Liu, Yanshuo Chen, and Heng Huang. Towards Copyright Protection for Knowledge Bases of Retrieval-augmented Language Models via Reasoning. *arXiv preprint arXiv:2502.10440*, 2025. URL <https://arxiv.org/abs/2502.10440>.
  - [71] Shuyu Guo, Shuo Zhang, and Zhaochun Ren. Enhancing RAG Efficiency with Adaptive Context Compression. *arXiv preprint arXiv:2507.22931*, 2025. URL <https://arxiv.org/abs/2507.22931>.
  - [72] Zirui Guo, Xubin Ren, Lingrui Xu, Jiahao Zhang, and Chao Huang. RAG-Anything: All-in-One RAG Framework. *arXiv preprint arXiv:2510.12323*, 2025. URL <https://arxiv.org/abs/2510.12323>.
  - [73] Shailja Gupta, Rajesh Ranjan, and Surya Narayan Singh. A Comprehensive Survey of Retrieval-Augmented Generation (RAG): Evolution, Current Landscape and Future Directions. *arXiv preprint arXiv:2410.12837*, 2024. URL <https://arxiv.org/abs/2410.12837>.
  - [74] Bernal Jimenez Gutierrez, Yiheng Shu, Weijian Qi, Sizhe Zhou, and Yu Su. From RAG to Memory: Non-Parametric Continual Learning for Large Language Models. *arXiv preprint arXiv:2502.14802*, 2025. URL <https://arxiv.org/abs/2502.14802>.
  - [75] Kelvin Guu, Kenton Lee, Zora Tung, Panupong Pasupat, and Ming-Wei Chang. REALM: Retrieval-Augmented Language Model Pre-Training. *arXiv preprint arXiv:2002.08909*, 2020. URL <https://arxiv.org/abs/2002.08909>.
  - [76] Kaiqiao Han, LuAn Tang, Renliang Sun, Peng Yuan, Wei Cheng, Haoyu Wang, Wei Wang, Yizhou Sun, and Haifeng Chen. Structure-Aware RAG: Structured Retrieval Augmented Generation from Noisy Data for Conversational Agents. *arXiv preprint arXiv:2605.24366*, 2026. URL <https://arxiv.org/abs/2605.24366>.
  - [77] Tasnimul Hassan, Md Faisal Karim, Haziq Jee-lani, Elham Behnam, Robert Green, and Fayeque Jee-lani Syed. Optimizing Medical Question-Answering Systems: A Comparative Study of Fine-Tuned and Zero-Shot Large Language Models with RAG Framework. *arXiv preprint arXiv:2512.05863*, 2025. URL <https://arxiv.org/abs/2512.05863>.
  - [78] Teruaki Hayashi, Hiroki Sakaji, Jiayi Dai, and Randy Goebel. Metadata-based Data Exploration with Retrieval-Augmented Generation for Large Language Models. *arXiv preprint arXiv:2410.04231*, 2024. URL <https://arxiv.org/abs/2410.04231>.
  - [79] Mohammad Hassan Heydari, Arshia Hemmat, Erfan Naman, and Afsaneh Fatemi. Context Awareness Gate For Retrieval Augmented Generation. *arXiv preprint arXiv:2411.16133*, 2024. URL <https://arxiv.org/abs/2411.16133>.
  - [80] Samuel Hildebrand, Curtis Taylor, Sean Oesch, James M Ghawaly, Amir Sadovnik, Ryan Shivers, Brandon Schreiber, and Kevin Kurian. FATHOMS-RAG: A Framework for the Assessment of Thinking and Observation in Multimodal Systems that use Retrieval Augmented Generation. *arXiv preprint arXiv:2510.08945*, 2025. URL <https://arxiv.org/abs/2510.08945>.
  - [81] Yubin Hong, Chaofan Li, Jingyi Zhang, and Yingxia Shao. FG-RAG: Enhancing Query-Focused Summarization with Context-Aware Fine-Grained Graph RAG. *arXiv preprint arXiv:2504.07103*, 2025. URL <https://arxiv.org/abs/2504.07103>.
  - [82] Daniel Hopp. Free to play: UN Trade and Development’s experience with developing its own open-source Retrieval Augmented Generation Large Language Model application. *arXiv preprint arXiv:2407.16896*, 2024. URL <https://arxiv.org/abs/2407.16896>.
  - [83] Alexis Horde-Vo, Matt Duckham, Estrid He, and Rafe Benli. Identifying Origins of Place Names via Retrieval Augmented Generation. *arXiv preprint arXiv:2509.01030*, 2025. URL <https://arxiv.org/abs/2509.01030>.
  - [84] Yucheng Hu and Yuxing Lu. RAG and RAU: A Survey on Retrieval-Augmented Language Model in Natural Language Processing. *arXiv preprint arXiv:2404.19543*, 2024. URL <https://arxiv.org/abs/2404.19543>.
  - [85] Yuntong Hu, Zhihan Lei, Zheng Zhang, Bo Pan, Chen Ling, and Liang Zhao. GRAG: Graph Retrieval-Augmented Generation. *arXiv preprint arXiv:2405.16506*, 2024. URL <https://arxiv.org/abs/2405.16506>.
  - [86] Zhibo Hu, Chen Wang, Yanfeng Shu, Helen, Paik, and Liming Zhu. Prompt Perturbation in Retrieval-Augmented Generation based Large Language Models. *arXiv preprint arXiv:2402.07179*, 2024. URL <https://arxiv.org/abs/2402.07179>.
  - [87] Chengkai Huang, Yu Xia, Rui Wang, Kaige Xie, Tong Yu, Julian McAuley, and Lina Yao. Embedding-Informed Adaptive Retrieval-Augmented Generation of Large Language Models. *arXiv preprint arXiv:2404.03514*, 2024. URL <https://arxiv.org/abs/2404.03514>.
  - [88] Tianyi Huang, Caden Yang, Emily Yin, Eric Wang, and Michael Zhang. PAVE: Premise-Aware Validation and Editing for Retrieval-Augmented LLMs. *arXiv preprint arXiv:2603.20673*, 2026. URL <https://arxiv.org/abs/2603.20673>.
  - [89] Yizheng Huang and Jimmy Huang. A Survey on Retrieval-Augmented Text Generation for Large

- Language Models. *arXiv preprint arXiv:2404.10981*, 2024. URL <https://arxiv.org/abs/2404.10981>.
- [90] Andrei Cristian Ionescu, Sergey Titov, and Maliheh Izadi. A Multi-agent Onboarding Assistant based on Large Language Models, Retrieval Augmented Generation, and Chain-of-Thought. *arXiv preprint arXiv:2503.23421*, 2025. URL <https://arxiv.org/abs/2503.23421>.
- [91] Paulina Toro Isaza, Michael Nidd, Noah Zheutlin, Jae wook Ahn, Chidansh Amitkumar Bhatt, Yu Deng, Ruchi Mahindru, Martin Franz, Hans Florian, and Salim Roukos. Retrieval Augmented Generation-Based Incident Resolution Recommendation System for IT Support. *arXiv preprint arXiv:2409.13707*, 2024. URL <https://arxiv.org/abs/2409.13707>.
- [92] Shayekh Bin Islam, Md Asib Rahman, K S M Tozammel Hossain, Enamul Hoque, Shafiq Joty, and Md Rizwan Parvez. Open-RAG: Enhanced Retrieval-Augmented Reasoning with Open-Source Large Language Models. *arXiv preprint arXiv:2410.01782*, 2024. URL <https://arxiv.org/abs/2410.01782>.
- [93] Koki Itai, Shunichi Hasegawa, Yuta Yamamoto, Gouki Minegishi, and Masaki Otsuki. LIT-RAGBench: Benchmarking Generator Capabilities of Large Language Models in Retrieval-Augmented Generation. *arXiv preprint arXiv:2603.06198*, 2026. URL <https://arxiv.org/abs/2603.06198>.
- [94] Gabriel Iturra-Bocaz and Petra Galuscakova. A Reproducibility Study of Metacognitive Retrieval-Augmented Generation. *arXiv preprint arXiv:2604.19899*, 2026. URL <https://arxiv.org/abs/2604.19899>.
- [95] Mohamed Sobhi Jabal, Pranav Warman, Jikai Zhang, Kartikeye Gupta, Ayush Jain, Maciej Mazurowski, Walter Wiggins, Kirti Magudia, and Evan Calabrese. Language Models and Retrieval Augmented Generation for Automated Structured Data Extraction from Diagnostic Reports. *arXiv preprint arXiv:2409.10576*, 2024. URL <https://arxiv.org/abs/2409.10576>.
- [96] Chaeyun Jang, Deukhwan Cho, Seanie Lee, Hyungi Lee, and Juho Lee. Reliable Decision Making via Calibration Oriented Retrieval Augmented Generation. *arXiv preprint arXiv:2411.08891*, 2024. URL <https://arxiv.org/abs/2411.08891>.
- [97] Jinhao Jiang, Jiayi Chen, Junyi Li, Ruiyang Ren, Shijie Wang, Wayne Xin Zhao, Yang Song, and Tao Zhang. RAG-Star: Enhancing Deliberative Reasoning with Retrieval Augmented Verification and Refinement. *arXiv preprint arXiv:2412.12881*, 2024. URL <https://arxiv.org/abs/2412.12881>.
- [98] Wenqi Jiang, Suvinay Subramanian, Cat Graves, Gustavo Alonso, Amir Yazdanbakhsh, and Vidushi Dadu. RAGO: Systematic Performance Optimization for Retrieval-Augmented Generation Serving. *arXiv preprint arXiv:2503.14649*, 2025. URL <https://arxiv.org/abs/2503.14649>.
- [99] Zhengbao Jiang, Frank F. Xu, Luyu Gao, Zhiqing Sun, Qian Liu, Jane Dwivedi-Yu, Yiming Yang, Jamie Callan, and Graham Neubig. Active Retrieval Augmented Generation. *arXiv preprint arXiv:2305.06983*, 2023. URL <https://arxiv.org/abs/2305.06983>.
- [100] Dian Jiao, Li Cai, Jingsheng Huang, Wenqiao Zhang, Siliang Tang, and Yueting Zhuang. DueRAG: Collaborative Retrieval-Augmented Generation. *arXiv preprint arXiv:2405.13002*, 2024. URL <https://arxiv.org/abs/2405.13002>.
- [101] YiHan Jiao, ZheHao Tan, Dan Yang, DuoLin Sun, Jie Feng, Yue Shen, Jian Wang, and Peng Wei. HIRAG: Hierarchical-Thought Instruction-Tuning Retrieval-Augmented Generation. *arXiv preprint arXiv:2507.05714*, 2025. URL <https://arxiv.org/abs/2507.05714>.
- [102] Zhuoran Jin, Hongbang Yuan, Tianyi Men, Pengfei Cao, Yubo Chen, Kang Liu, and Jun Zhao. RAG-RewardBench: Benchmarking Reward Models in Retrieval Augmented Generation for Preference Alignment. *arXiv preprint arXiv:2412.13746*, 2024. URL <https://arxiv.org/abs/2412.13746>.
- [103] Mintong Kang, Nezihe Merve Gurel, Ning Yu, Dawn Song, and Bo Li. C-RAG: Certified Generation Risks for Retrieval-Augmented Language Models. *arXiv preprint arXiv:2402.03181*, 2024. URL <https://arxiv.org/abs/2402.03181>.
- [104] Yong-Bin Kang and Anthony McCosker. TA-RAG: Tone-Aware Retrieval-Augmented Generation for Peer-Support Health Communication. *arXiv preprint arXiv:2606.06794*, 2026. URL <https://arxiv.org/abs/2606.06794>.
- [105] A H M Rezaul Karim and Ozlem Uzuner. Multimodal Retrieval-Augmented Generation with Large Language Models for Medical VQA. *arXiv preprint arXiv:2510.13856*, 2025. URL <https://arxiv.org/abs/2510.13856>.
- [106] YuHe Ke, Liyuan Jin, Kabilan Elangovan, Hairil Rizal Abdullah, Nan Liu, Alex Tiong Heng Sia, Chai Rick Soh, Joshua Yi Min Tung, Jasmine Chiat Ling Ong, and Daniel Shu Wei Ting. Development and Testing of Retrieval Augmented Generation in Large Language Models – A Case Study Report. *arXiv preprint arXiv:2402.01733*, 2024. URL <https://arxiv.org/abs/2402.01733>.
- [107] Tristan Kenneweg, Philip Kenneweg, and Barbara Hammer. Retrieval Augmented Generation Systems: Automatic Dataset Creation, Evaluation and Boolean Agent Setup. *arXiv preprint*

- arXiv:2403.00820*, 2024. URL <https://arxiv.org/abs/2403.00820>.
- [108] Harshad Khadilkar and Abhay Gupta. Causal-Counterfactual RAG: The Integration of Causal-Counterfactual Reasoning into RAG. *arXiv preprint arXiv:2509.14435*, 2025. URL <https://arxiv.org/abs/2509.14435>.
- [109] Zahra Khalila, Arbi Haza Nasution, Winda Monika, Aytug Onan, Yohei Murakami, Yasir Bin Ismail Radi, and Noor Mohammad Osmani. Investigating Retrieval-Augmented Generation in Quranic Studies: A Study of 13 Open-Source Large Language Models. *arXiv preprint arXiv:2503.16581*, 2025. URL <https://arxiv.org/abs/2503.16581>.
- [110] M. Abdul Khaliq, P. Chang, M. Ma, B. Pflugfelder, and F. Miletic. RAGAR, Your Falsehood Radar: RAG-Augmented Reasoning for Political Fact-Checking using Multimodal Large Language Models. *arXiv preprint arXiv:2404.12065*, 2024. URL <https://arxiv.org/abs/2404.12065>.
- [111] Jaemin Kim, Jae O Lee, Sumyeong Ahn, and Seo Yeon Park. Neuro-RIT: Neuron-Guided Instruction Tuning for Robust Retrieval-Augmented Language Model. *arXiv preprint arXiv:2604.02194*, 2026. URL <https://arxiv.org/abs/2604.02194>.
- [112] Seonok Kim. MedBioLM: Optimizing Medical and Biological QA with Fine-Tuned Large Language Models and Retrieval-Augmented Generation. *arXiv preprint arXiv:2502.03004*, 2025. URL <https://arxiv.org/abs/2502.03004>.
- [113] Seonok Kim. MedBioRAG: Semantic Search and Retrieval-Augmented Generation with Large Language Models for Medical and Biological QA. *arXiv preprint arXiv:2512.10996*, 2025. URL <https://arxiv.org/abs/2512.10996>.
- [114] Singon Kim. Noise-Robust Abstractive Compression in Retrieval-Augmented Language Models. *arXiv preprint arXiv:2512.08943*, 2025. URL <https://arxiv.org/abs/2512.08943>.
- [115] Singon Kim, Gunho Jung, and Seong-Wan Lee. ACoRN: Noise-Robust Abstractive Compression in Retrieval-Augmented Language Models. *arXiv preprint arXiv:2504.12673*, 2025. URL <https://arxiv.org/abs/2504.12673>.
- [116] Antonios Minas Krasakis, Andrew Yates, and Evangelos Kanoulas. Corpus-informed Retrieval Augmented Generation of Clarifying Questions. *arXiv preprint arXiv:2409.18575*, 2024. URL <https://arxiv.org/abs/2409.18575>.
- [117] Valentin J. J. Kreileder, Johannes Reisinger, and Andreas Fischer. Evaluating Chunking Strategies for Retrieval-Augmented Generation on Academic Texts. *arXiv preprint arXiv:2607.01852*, 2026. URL <https://arxiv.org/abs/2607.01852>.
- [118] Tzu-Lin Kuo, Feng-Ting Liao, Mu-Wei Hsieh, Fu-Chieh Chang, Po-Chun Hsu, and Da-Shan Shiu. RAD-Bench: Evaluating Large Language Models Capabilities in Retrieval Augmented Dialogues. *arXiv preprint arXiv:2409.12558*, 2024. URL <https://arxiv.org/abs/2409.12558>.
- [119] Daeyong Kwon, SeungHeon Doh, and Juhan Nam. MUST-RAG: MUSical Text Question Answering with Retrieval Augmented Generation. *arXiv preprint arXiv:2507.23334*, 2025. URL <https://arxiv.org/abs/2507.23334>.
- [120] Hamza Landolsi, Kais Letaief, Nizar Taghouti, and Ines Abdeljaoued-Tej. CAPRAG: A Large Language Model Solution for Customer Service and Automatic Reporting using Vector and Graph Retrieval-Augmented Generation. *arXiv preprint arXiv:2501.13993*, 2025. URL <https://arxiv.org/abs/2501.13993>.
- [121] Hung-Shin Lee, Chen-Chi Chang, Ching-Yuan Chen, and Yun-Hsiang Hsu. Evaluating Cultural Knowledge Processing in Large Language Models: A Cognitive Benchmarking Framework Integrating Retrieval-Augmented Generation. *arXiv preprint arXiv:2511.01649*, 2025. URL <https://arxiv.org/abs/2511.01649>.
- [122] Juntae Lee, Jihwan Bang, Seunghan Yang, Kyuhong Shim, and Simyung Chang. Chain-of-Rank: Enhancing Large Language Models for Domain-Specific RAG in Edge Device. *arXiv preprint arXiv:2502.15134*, 2025. URL <https://arxiv.org/abs/2502.15134>.
- [123] Kyumin Lee, Minjin Jeon, Sanghwan Jang, and Hwanjo Yu. STEPER: Step-wise Knowledge Distillation for Enhancing Reasoning Ability in Multi-Step Retrieval-Augmented Language Models. *arXiv preprint arXiv:2510.07923*, 2025. URL <https://arxiv.org/abs/2510.07923>.
- [124] Quinn Leng, Jacob Portes, Sam Havens, Matei Zaharia, and Michael Carbin. Long Context RAG Performance of Large Language Models. *arXiv preprint arXiv:2411.03538*, 2024. URL <https://arxiv.org/abs/2411.03538>.
- [125] Alexandria Leto, Cecilia Aguerrebere, Ishwar Bhati, Ted Willke, Mariano Tepper, and Vy Ai Vo. Toward Optimal Search and Retrieval for RAG. *arXiv preprint arXiv:2411.07396*, 2024. URL <https://arxiv.org/abs/2411.07396>.
- [126] Dongyang Li, Junbing Yan, Taolin Zhang, Chengyu Wang, Xiaofeng He, Longtao Huang, Hui Xue, and Jun Huang. On the Role of Long-tail Knowledge in Retrieval Augmented Large Language Models. *arXiv preprint arXiv:2406.16367*, 2024. URL <https://arxiv.org/abs/2406.16367>.
- [127] Feiyang Li, Peng Fang, Zhan Shi, Arijit Khan, Fang Wang, Weihao Wang, Xin Zhang, and Yongjian

- Cui. CoT-RAG: Integrating Chain of Thought and Retrieval-Augmented Generation to Enhance Reasoning in Large Language Models. *arXiv preprint arXiv:2504.13534*, 2025. URL <https://arxiv.org/abs/2504.13534>.
- [128] Haoqing Li, Jun Shi, Xianmeng Chen, Qiwei Jia, Rui Wang, Wei Wei, Hong An, and Xiaowen Hu. Large Language Model Aided Birt-Hogg-Dube Syndrome Diagnosis with Multimodal Retrieval-Augmented Generation. *arXiv preprint arXiv:2511.19834*, 2025. URL <https://arxiv.org/abs/2511.19834>.
- [129] Kevin Chenhao Li, Vahid Zolfaghari, Nenad Petrovic, Fengjunjie Pan, and Alois Knoll. Optimizing Retrieval Augmented Generation for Object Constraint Language. *arXiv preprint arXiv:2505.13129*, 2025. URL <https://arxiv.org/abs/2505.13129>.
- [130] Mingchen Li, Halil Kilicoglu, Hua Xu, and Rui Zhang. BiomedRAG: A Retrieval Augmented Large Language Model for Biomedicine. *arXiv preprint arXiv:2405.00465*, 2024. URL <https://arxiv.org/abs/2405.00465>.
- [131] Mufei Li, Siqi Miao, and Pan Li. Simple Is Effective: The Roles of Graphs and Large Language Models in Knowledge-Graph-Based Retrieval-Augmented Generation. *arXiv preprint arXiv:2410.20724*, 2024. URL <https://arxiv.org/abs/2410.20724>.
- [132] Xiaodi Li, Yang Xiao, Munhwan Lee, Konstantinos Leventakos, Young J. Juhn, David Jones, Terence T. Sio, Wei Liu, Maria Vassilaki, and Nansu Zong. Lightweight Retrieval-Augmented Generation and Large Language Model-Based Modeling for Scalable Patient-Trial Matching. *arXiv preprint arXiv:2604.22061*, 2026. URL <https://arxiv.org/abs/2604.22061>.
- [133] Xinze Li, Hanbin Wang, Zhenghao Liu, Shi Yu, Shuo Wang, Yukun Yan, Yukai Fu, Yu Gu, and Ge Yu. Building A Coding Assistant via the Retrieval-Augmented Language Model. *arXiv preprint arXiv:2410.16229*, 2024. URL <https://arxiv.org/abs/2410.16229>.
- [134] Xinzhe Li, Ming Liu, and Shang Gao. GRAMMAR: Grounded and Modular Methodology for Assessment of Closed-Domain Retrieval-Augmented Language Model. *arXiv preprint arXiv:2404.19232*, 2024. URL <https://arxiv.org/abs/2404.19232>.
- [135] Yangning Li, Weizhi Zhang, Yuyao Yang, Wei-Chieh Huang, Yaozu Wu, Junyu Luo, Yuanchen Bei, Henry Peng Zou, Xiao Luo, Yusheng Zhao, Chunkit Chan, Yankai Chen, Zhongfen Deng, Yinghui Li, Hai-Tao Zheng, Dongyuan Li, Renhe Jiang, Ming Zhang, Yangqiu Song, and Philip S. Yu. Towards Agentic RAG with Deep Reasoning: A Survey of RAG-Reasoning Systems in LLMs. *arXiv preprint arXiv:2507.09477*, 2025. URL <https://arxiv.org/abs/2507.09477>.
- [136] Yihan Li, Xiyuan Fu, Ghanshyam Verma, Paul Buitelaar, and Mingming Liu. Mitigating Hallucination in Large Language Models (LLMs): An Application-Oriented Survey on RAG, Reasoning, and Agentic Systems. *arXiv preprint arXiv:2510.24476*, 2025. URL <https://arxiv.org/abs/2510.24476>.
- [137] Yongjian Li, HaoCheng Chu, Yukun Yan, Zhenghao Liu, Shi Yu, Zheni Zeng, Ruobing Wang, Sen Song, Zhiyuan Liu, and Maosong Sun. KARE-RAG: Knowledge-Aware Refinement and Enhancement for RAG. *arXiv preprint arXiv:2506.02503*, 2025. URL <https://arxiv.org/abs/2506.02503>.
- [138] Yuhan Li, Xinni Zhang, Linhao Luo, Heng Chang, Yuxiang Ren, Irwin King, and Jia Li. G-Refer: Graph Retrieval-Augmented Large Language Model for Explainable Recommendation. *arXiv preprint arXiv:2502.12586*, 2025. URL <https://arxiv.org/abs/2502.12586>.
- [139] Zizhong Li, Haopeng Zhang, and Jiawei Zhang. Unveiling the Magic: Investigating Attention Distillation in Retrieval-augmented Generation. *arXiv preprint arXiv:2402.11794*, 2024. URL <https://arxiv.org/abs/2402.11794>.
- [140] Xun Liang, Simin Niu, Zhiyu Li, Sensen Zhang, Hanyu Wang, Feiyu Xiong, Jason Zhaoxin Fan, Bo Tang, Shichao Song, Mengwei Wang, and Jiawei Yang. SafeRAG: Benchmarking Security in Retrieval-Augmented Generation of Large Language Model. *arXiv preprint arXiv:2501.18636*, 2025. URL <https://arxiv.org/abs/2501.18636>.
- [141] Shiyin Lin. Abductive Inference in Retrieval-Augmented Language Models: Generating and Validating Missing Premises. *arXiv preprint arXiv:2511.04020*, 2025. URL <https://arxiv.org/abs/2511.04020>.
- [142] Shuhang Lin, Zhenan Peng, Lingyao Li, Xiao Lin, Xi Zhu, and Yongfeng Zhang. Cache Mechanism for Agent RAG Systems. *arXiv preprint arXiv:2511.02919*, 2025. URL <https://arxiv.org/abs/2511.02919>.
- [143] Jasper Linders and Jakub M. Tomczak. Knowledge Graph-extended Retrieval Augmented Generation for Question Answering. *arXiv preprint arXiv:2504.08893*, 2025. URL <https://arxiv.org/abs/2504.08893>.
- [144] Hongjun Liu, Yilun Zhao, Arman Cohan, and Chen Zhao. SUCEA: Reasoning-Intensive Retrieval for Adversarial Fact-checking through Claim Decomposition and Editing. *arXiv preprint arXiv:2506.04583*, 2025. URL <https://arxiv.org/abs/2506.04583>.



- [145] Jiayu Liu, Rui Wang, Qing Zong, Yumeng Wang, Cheng Qian, Qingcheng Zeng, Tianshi Zheng, Haochen Shi, Dadi Guo, Baixuan Xu, Chunyang Li, and Yangqiu Song. NOVA: NOise-aware Verbal Confidence CALibration for Robust Large Language Models in RAG Systems. *arXiv preprint arXiv:2601.11004*, 2026. URL <https://arxiv.org/abs/2601.11004>.
- [146] Jingyu Liu, Jiaen Lin, and Yong Liu. How Much Can RAG Help the Reasoning of LLM? *arXiv preprint arXiv:2410.02338*, 2024. URL <https://arxiv.org/abs/2410.02338>.
- [147] Jiongnan Liu, Jiajie Jin, Zihan Wang, Jiehan Cheng, Zhicheng Dou, and Ji-Rong Wen. RETA-LLM: A Retrieval-Augmented Large Language Model Toolkit. *arXiv preprint arXiv:2306.05212*, 2023. URL <https://arxiv.org/abs/2306.05212>.
- [148] Junyi Liu, Liangzhi Li, Tong Xiang, Bowen Wang, and Yiming Qian. TCRA-LLM: Token Compression Retrieval Augmented Large Language Model for Inference Cost Reduction. *arXiv preprint arXiv:2310.15556*, 2023. URL <https://arxiv.org/abs/2310.15556>.
- [149] Runheng Liu, Xingchen Xiao, Heyan Huang, Zewen Chi, and Zhijing Wu. FlashBack:Efficient Retrieval-Augmented Language Modeling for Long Context Inference. *arXiv preprint arXiv:2405.04065*, 2024. URL <https://arxiv.org/abs/2405.04065>.
- [150] Shuliang Liu, Xinze Li, Zhenghao Liu, Yukun Yan, Cheng Yang, Zheni Zeng, Zhiyuan Liu, Maosong Sun, and Ge Yu. Judge as A Judge: Improving the Evaluation of Retrieval-Augmented Generation through the Judge-Consistency of Large Language Models. *arXiv preprint arXiv:2502.18817*, 2025. URL <https://arxiv.org/abs/2502.18817>.
- [151] Wanlong Liu, Junying Chen, Ke Ji, Li Zhou, Wenyu Chen, and Benyou Wang. RAG-Instruct: Boosting LLMs with Diverse Retrieval-Augmented Instructions. *arXiv preprint arXiv:2501.00353*, 2024. URL <https://arxiv.org/abs/2501.00353>.
- [152] Fatma Yasmine Loumachi, Mohamed Chahine Ghanem, and Mohamed Amine Ferrag. GenD-FIR: Advancing Cyber Incident Timeline Analysis Through Retrieval Augmented Generation and Large Language Models. *arXiv preprint arXiv:2409.02572*, 2024. URL <https://arxiv.org/abs/2409.02572>.
- [153] Elias Lumer, Alex Cardenas, Matt Melich, Myles Mason, Sara Dieter, Vamse Kumar Subbiah, Pradeep Honaganahalli Basavaraju, and Roberto Hernandez. Comparison of Text-Based and Image-Based Retrieval in Multimodal Retrieval Augmented Generation Large Language Model Systems. *arXiv preprint arXiv:2511.16654*, 2025. URL <https://arxiv.org/abs/2511.16654>.
- [154] Linhao Luo, Zicheng Zhao, Gholamreza Haffari, Dinh Phung, Chen Gong, and Shirui Pan. GFM-RAG: Graph Foundation Model for Retrieval Augmented Generation. *arXiv preprint arXiv:2502.01113*, 2025. URL <https://arxiv.org/abs/2502.01113>.
- [155] Sichun Luo, Jian Xu, Xiaojie Zhang, Linrong Wang, Sicong Liu, Hanxu Hou, and Linqi Song. RALL-Rec+: Retrieval Augmented Large Language Model Recommendation with Reasoning. *arXiv preprint arXiv:2503.20430*, 2025. URL <https://arxiv.org/abs/2503.20430>.
- [156] Peizhuo Lv, Mengjie Sun, Hao Wang, Xiaofeng Wang, Shengzhi Zhang, Yuxuan Chen, Kai Chen, and Limin Sun. RAG-WM: An Efficient Black-Box Watermarking Approach for Retrieval-Augmented Generation of Large Language Models. *arXiv preprint arXiv:2501.05249*, 2025. URL <https://arxiv.org/abs/2501.05249>.
- [157] Xingda Lyu, Gongfu Lyu, Zitai Yan, and Yuxin Jiang. P-RAG: Prompt-Enhanced Parametric RAG with LoRA and Selective CoT for Biomedical and Multi-Hop QA. *arXiv preprint arXiv:2602.15874*, 2026. URL <https://arxiv.org/abs/2602.15874>.
- [158] Yuanjie Lyu, Zhiyu Li, Simin Niu, Feiyu Xiong, Bo Tang, Wenjin Wang, Hao Wu, Huanyong Liu, Tong Xu, and Enhong Chen. CRUD-RAG: A Comprehensive Chinese Benchmark for Retrieval-Augmented Generation of Large Language Models. *arXiv preprint arXiv:2401.17043*, 2024. URL <https://arxiv.org/abs/2401.17043>.
- [159] Shengjie Ma, Chengjin Xu, Xuhui Jiang, Muzhi Li, Huaren Qu, Cehao Yang, Jiaxin Mao, and Jian Guo. Think-on-Graph 2.0: Deep and Faithful Large Language Model Reasoning with Knowledge-guided Retrieval Augmented Generation. *arXiv preprint arXiv:2407.10805*, 2024. URL <https://arxiv.org/abs/2407.10805>.
- [160] Xinbei Ma, Yeyun Gong, Pengcheng He, Hai Zhao, and Nan Duan. Query Rewriting for Retrieval-Augmented Large Language Models. *arXiv preprint arXiv:2305.14283*, 2023. URL <https://arxiv.org/abs/2305.14283>.
- [161] Xueguang Ma, Shengyao Zhuang, Bevan Koopman, Guido Zuccon, Wenhui Chen, and Jimmy Lin. VISA: Retrieval Augmented Generation with Visual Source Attribution. *arXiv preprint arXiv:2412.14457*, 2024. URL <https://arxiv.org/abs/2412.14457>.
- [162] Harish Tayyar Madabushi. FS-RAG: A Frame Semantics Based Approach for Improved Factual Accuracy in Large Language Models. *arXiv preprint arXiv:2406.16167*, 2024. URL <https://arxiv.org/abs/2406.16167>.

- [163] Ahmadreza Majlesara, Ali Majlesi, Ali Mamaghani, Alireza Shokrani, and Babak Hossein Khalaj. 5G Network Automation Using Local Large Language Models and Retrieval-Augmented Generation. *arXiv preprint arXiv:2511.21084*, 2025. URL <https://arxiv.org/abs/2511.21084>.
- [164] Yuren Mao, Xuemei Dong, Wenyi Xu, Yunjun Gao, Bin Wei, and Ying Zhang. FIT-RAG: Black-Box RAG with Factual Information and Token Reduction. *arXiv preprint arXiv:2403.14374*, 2024. URL <https://arxiv.org/abs/2403.14374>.
- [165] Reza Yousefi Maragheh, Pratheek Vadla, Priyank Gupta, Kai Zhao, Aysenur Inan, Kehui Yao, Jianpeng Xu, Praveen Kanumala, Jason Cho, and Sushant Kumar. ARAG: Agentic Retrieval Augmented Generation for Personalized Recommendation. *arXiv preprint arXiv:2506.21931*, 2025. URL <https://arxiv.org/abs/2506.21931>.
- [166] Julia Ann Mathew and Suining He. Lessons from A Large Language Model-based Outdoor Trail Recommendation Chatbot with Retrieval Augmented Generation. *arXiv preprint arXiv:2508.05652*, 2025. URL <https://arxiv.org/abs/2508.05652>.
- [167] Costas Mavromatis and George Karypis. GNN-RAG: Graph Neural Retrieval for Large Language Model Reasoning. *arXiv preprint arXiv:2405.20139*, 2024. URL <https://arxiv.org/abs/2405.20139>.
- [168] Lang Mei, Siyu Mo, Zhihan Yang, and Chong Chen. A Survey of Multimodal Retrieval-Augmented Generation. *arXiv preprint arXiv:2504.08748*, 2025. URL <https://arxiv.org/abs/2504.08748>.
- [169] Karan Mirhosseini, Arya Aftab, and Alireza Sheikh. RATE: An LLM-Powered Retrieval Augmented Generation Technology-Extraction Pipeline. *arXiv preprint arXiv:2507.21125*, 2025. URL <https://arxiv.org/abs/2507.21125>.
- [170] Fengran Mo, Zhan Su, Yuchen Hui, Jinghan Zhang, Jia Ao Sun, Zheyuan Liu, Chao Zhang, Tetsuya Sakai, and Jian-Yun Nie. OpenDecoder: Open Large Language Model Decoding to Incorporate Document Quality in RAG. *arXiv preprint arXiv:2601.09028*, 2026. URL <https://arxiv.org/abs/2601.09028>.
- [171] Sai Munikoti, Anurag Acharya, Sridevi Wagle, and Sameera Horawalavithana. ATLANTIC: Structure-Aware Retrieval-Augmented Language Model for Interdisciplinary Science. *arXiv preprint arXiv:2311.12289*, 2023. URL <https://arxiv.org/abs/2311.12289>.
- [172] Thushara Manjari Naduvilakandy, Hyeju Jang, and Mohammad Al Hasan. Retrieval Augmented Generation based Large Language Models for Causality Mining. *arXiv preprint arXiv:2505.23944*, 2025. URL <https://arxiv.org/abs/2505.23944>.
- [173] Ali Naseh, Yuefeng Peng, Anshuman Suri, Harsh Chaudhari, Alina Oprea, and Amir Houmansadr. Riddle Me This! Stealthy Membership Inference for Retrieval-Augmented Generation. *arXiv preprint arXiv:2502.00306*, 2025. URL <https://arxiv.org/abs/2502.00306>.
- [174] Nghia Trung Ngo, Chien Van Nguyen, Franck Deroncourt, and Thien Huu Nguyen. Comprehensive and Practical Evaluation of Retrieval-Augmented Generation Systems for Medical Question Answering. *arXiv preprint arXiv:2411.09213*, 2024. URL <https://arxiv.org/abs/2411.09213>.
- [175] Hang Ni, Fan Liu, Xinyu Ma, Lixin Su, Shuaiqiang Wang, Dawei Yin, Hui Xiong, and Hao Liu. TP-RAG: Benchmarking Retrieval-Augmented Large Language Model Agents for Spatiotemporal-Aware Travel Planning. *arXiv preprint arXiv:2504.08694*, 2025. URL <https://arxiv.org/abs/2504.08694>.
- [176] Cheng Niu, Yuanhao Wu, Juno Zhu, Siliang Xu, Kashun Shum, Randy Zhong, Juntong Song, and Tong Zhang. RAGTruth: A Hallucination Corpus for Developing Trustworthy Retrieval-Augmented Language Models. *arXiv preprint arXiv:2401.00396*, 2023. URL <https://arxiv.org/abs/2401.00396>.
- [177] Qinjiang Niu and Lu Yan. AOP-Smart: A RAG-Enhanced Large Language Model Framework for Adverse Outcome Pathway Analysis. *arXiv preprint arXiv:2604.10874*, 2026. URL <https://arxiv.org/abs/2604.10874>.
- [178] Jonas Oppenlaender. DangerMaps: Personalized Safety Advice for Travel in Urban Environments using a Retrieval-Augmented Language Model. *arXiv preprint arXiv:2503.14103*, 2025. URL <https://arxiv.org/abs/2503.14103>.
- [179] Jaewan Park, Solbee Cho, and Jay-Yoon Lee. Stop-RAG: Value-Based Retrieval Control for Iterative RAG. *arXiv preprint arXiv:2510.14337*, 2025. URL <https://arxiv.org/abs/2510.14337>.
- [180] Seong-Il Park and Jay-Yoon Lee. Toward Robust RALMs: Revealing the Impact of Imperfect Retrieval on Retrieval-Augmented Language Models. *arXiv preprint arXiv:2410.15107*, 2024. URL <https://arxiv.org/abs/2410.15107>.
- [181] Udit Patel, Rutu Mulkar, Jay Roberts, Cibi Chakravarthy Senthilkumar, Sujay Gandhi, Xiaofei Zheng, Naumaan Nayyar, Parul Kalra, and Rafael Castrillo. THELMA: Task Based Holistic Evaluation of Large Language Model Applications-RAG Question Answering. *arXiv preprint arXiv:2505.11626*, 2025. URL <https://arxiv.org/abs/2505.11626>.
- [182] Alexander R. Pelletier, Joseph Ramirez, Irsyad Adam, Simha Sankar, Yu Yan, Ding Wang, Dylan Steinecke, Wei Wang, and Peipei Ping. Explainable Biomedical Hypothesis Generation via Retrieval

- Augmented Generation enabled Large Language Models. *arXiv preprint arXiv:2407.12888*, 2024. URL <https://arxiv.org/abs/2407.12888>.
- [183] Mykhailo Poliakov and Nadiya Shvai. Multi-Meta-RAG: Improving RAG for Multi-Hop Queries using Database Filtering with LLM-Extracted Metadata. *arXiv preprint arXiv:2406.13213*, 2024. URL <https://arxiv.org/abs/2406.13213>.
- [184] Sonal Prabhune and Donald J. Berndt. Deploying Large Language Models With Retrieval Augmented Generation. *arXiv preprint arXiv:2411.11895*, 2024. URL <https://arxiv.org/abs/2411.11895>.
- [185] Ronak Pradeep, Nandan Thakur, Shivani Upadhyay, Daniel Campos, Nick Craswell, and Jimmy Lin. The Great Nugget Recall: Automating Fact Extraction and RAG Evaluation with Large Language Models. *arXiv preprint arXiv:2504.15068*, 2025. URL <https://arxiv.org/abs/2504.15068>.
- [186] Deeksha Prahlad, Chanhee Lee, Dongha Kim, and Hokeun Kim. Personalizing Large Language Models using Retrieval Augmented Generation and Knowledge Graph. *arXiv preprint arXiv:2505.09945*, 2025. URL <https://arxiv.org/abs/2505.09945>.
- [187] Gustavo Correa Publio and Jose Emilio Labra Gayo. xpSHACL: Explainable SHACL Validation using Retrieval-Augmented Generation and Large Language Models. *arXiv preprint arXiv:2507.08432*, 2025. URL <https://arxiv.org/abs/2507.08432>.
- [188] Zile Qiao, Wei Ye, Yong Jiang, Tong Mo, Pengjun Xie, Weiping Li, Fei Huang, and Shikun Zhang. Supportiveness-based Knowledge Rewriting for Retrieval-augmented Language Modeling. *arXiv preprint arXiv:2406.08116*, 2024. URL <https://arxiv.org/abs/2406.08116>.
- [189] Derrick Quinn, Mohammad Nouri, Neel Patel, John Salihu, Alireza Salemi, Sukhan Lee, Hamed Zamani, and Mohammad Alian. Accelerating Retrieval-Augmented Generation. *arXiv preprint arXiv:2412.15246*, 2024. URL <https://arxiv.org/abs/2412.15246>.
- [190] Aryan Raj, Astitva Veer Garg, and Anitha D. Structured Relevance Assessment for Robust Retrieval-Augmented Language Models. *arXiv preprint arXiv:2507.21287*, 2025. URL <https://arxiv.org/abs/2507.21287>.
- [191] Ori Ram, Yoav Levine, Itay Dalmedigos, Dor Muhlgay, Amnon Shashua, Kevin Leyton-Brown, and Yoav Shoham. In-Context Retrieval-Augmented Language Models. *arXiv preprint arXiv:2302.00083*, 2023. URL <https://arxiv.org/abs/2302.00083>.
- [192] Leonardo Ranaldi, Marco Valentino, and Andre Freitas. Eliciting Critical Reasoning in Retrieval-Augmented Language Models via Contrastive Explanations. *arXiv preprint arXiv:2410.22874*, 2024. URL <https://arxiv.org/abs/2410.22874>.
- [193] Leonardo Ranaldi, Federico Ranaldi, Fabio Massimo Zanzotto, Barry Haddow, and Alexandra Birch. Improving Multilingual Retrieval-Augmented Language Models through Dialectic Reasoning Argumentations. *arXiv preprint arXiv:2504.04771*, 2025. URL <https://arxiv.org/abs/2504.04771>.
- [194] Jiarui Rao and Jionghao Lin. RAMO: Retrieval-Augmented Generation for Enhancing MOOCs Recommendations. *arXiv preprint arXiv:2407.04925*, 2024. URL <https://arxiv.org/abs/2407.04925>.
- [195] Varun Nagaraj Rao, Siddharth Choudhary, Aditya Deshpande, Ravi Kumar Satzoda, and Srikar Appalaraju. RAVEN: Multitask Retrieval Augmented Vision-Language Learning. *arXiv preprint arXiv:2406.19150*, 2024. URL <https://arxiv.org/abs/2406.19150>.
- [196] Siddhant Ray, Rui Pan, Zhuohan Gu, Kuntai Du, Shaoting Feng, Ganesh Ananthanarayanan, Ravi Netravali, and Junchen Jiang. METIS: Fast Quality-Aware RAG Systems with Configuration Adaptation. *arXiv preprint arXiv:2412.10543*, 2024. URL <https://arxiv.org/abs/2412.10543>.
- [197] Xubin Ren, Lingrui Xu, Long Xia, Shuaiqiang Wang, Dawei Yin, and Chao Huang. VideoRAG: Retrieval-Augmented Generation with Extreme Long-Context Videos. *arXiv preprint arXiv:2502.01549*, 2025. URL <https://arxiv.org/abs/2502.01549>.
- [198] Mattia Rengo, Senad Beadini, Domenico Alfano, and Roberto Abbruzzese. A System for Comprehensive Assessment of RAG Frameworks. *arXiv preprint arXiv:2504.07803*, 2025. URL <https://arxiv.org/abs/2504.07803>.
- [199] Mohammad Reza Rezaei and Adji Bousso Dieng. Vendi-RAG: Adaptively Trading-Off Diversity And Quality Significantly Improves Retrieval Augmented Generation With LLMs. *arXiv preprint arXiv:2502.11228*, 2025. URL <https://arxiv.org/abs/2502.11228>.
- [200] Giorgio Roffo. Exploring Advanced Large Language Models with LLMsuite. *arXiv preprint arXiv:2407.12036*, 2024. URL <https://arxiv.org/abs/2407.12036>.
- [201] Kevin Dela Rosa. Video Enriched Retrieval Augmented Generation Using Aligned Video Captions. *arXiv preprint arXiv:2405.17706*, 2024. URL <https://arxiv.org/abs/2405.17706>.
- [202] Pouria Rouzrokh, Shahriar Faghani, Cooper U. Gamble, Moein Shariatnia, and Bradley J. Erickson. CONFLARE: CONFormal LArge language model

- REtrieval. *arXiv preprint arXiv:2404.04287*, 2024. URL <https://arxiv.org/abs/2404.04287>.
- [203] Soham Roy, Mitul Goswami, Nisharg Nargund, Suneeta Mohanty, and Prasant Kumar Pattnaik. Conversational Text Extraction with Large Language Models Using Retrieval-Augmented Systems. *arXiv preprint arXiv:2501.09801*, 2025. URL <https://arxiv.org/abs/2501.09801>.
- [204] Ohad Rubin and Jonathan Berant. Retrieval-Pretrained Transformer: Long-range Language Modeling with Self-retrieval. *arXiv preprint arXiv:2306.13421*, 2023. URL <https://arxiv.org/abs/2306.13421>.
- [205] Timothy Rupperecht, Enfu Nan, Arash Akbari, Arman Akbari, Lei Lu, Priyanka Maan, Sean Duffy, Pu Zhao, Yumei He, David Kaeli, and Yanzhi Wang. RAGs to Riches: RAG-like Few-shot Learning for Large Language Model Role-playing. *arXiv preprint arXiv:2509.12168*, 2025. URL <https://arxiv.org/abs/2509.12168>.
- [206] Michael J. Ryan, Danmei Xu, Chris Nivera, and Daniel Campos. EnronQA: Towards Personalized RAG over Private Documents. *arXiv preprint arXiv:2505.00263*, 2025. URL <https://arxiv.org/abs/2505.00263>.
- [207] Alireza Salemi and Hamed Zamani. Evaluating Retrieval Quality in Retrieval-Augmented Generation. *arXiv preprint arXiv:2404.13781*, 2024. URL <https://arxiv.org/abs/2404.13781>.
- [208] Alireza Salemi, Surya Kallumadi, and Hamed Zamani. Optimization Methods for Personalizing Large Language Models through Retrieval Augmentation. *arXiv preprint arXiv:2404.05970*, 2024. URL <https://arxiv.org/abs/2404.05970>.
- [209] Aniruddha Salve, Saba Attar, Mahesh Deshmukh, Sayali Shivpuje, and Arnab Mitra Utsab. A Collaborative Multi-Agent Approach to Retrieval-Augmented Generation Across Diverse Data. *arXiv preprint arXiv:2412.05838*, 2024. URL <https://arxiv.org/abs/2412.05838>.
- [210] Kunal Sawarkar, Abhilasha Mangal, and Shivam Raj Solanki. Blended RAG: Improving RAG (Retriever-Augmented Generation) Accuracy with Semantic Search and Hybrid Query-Based Retrievers. *arXiv preprint arXiv:2404.07220*, 2024. URL <https://arxiv.org/abs/2404.07220>.
- [211] Bhargav Shandilya and Alexis Palmer. Boosting the Capabilities of Compact Models in Low-Data Contexts with Large Language Models and Retrieval-Augmented Generation. *arXiv preprint arXiv:2410.00387*, 2024. URL <https://arxiv.org/abs/2410.00387>.
- [212] Chaitanya Sharma. Retrieval-Augmented Generation: A Comprehensive Survey of Architectures, Enhancements, and Robustness Frontiers. *arXiv preprint arXiv:2506.00054*, 2025. URL <https://arxiv.org/abs/2506.00054>.
- [213] Weijia Shi, Sewon Min, Michihiro Yasunaga, Minjoon Seo, Rich James, Mike Lewis, Luke Zettlemoyer, and Wen tau Yih. REPLUG: Retrieval-Augmented Black-Box Language Models. *arXiv preprint arXiv:2301.12652*, 2023. URL <https://arxiv.org/abs/2301.12652>.
- [214] Yunxiao Shi, Xing Zi, Zijing Shi, Haimin Zhang, Qiang Wu, and Min Xu. ERAGent: Enhancing Retrieval-Augmented Language Models with Improved Accuracy, Efficiency, and Personalization. *arXiv preprint arXiv:2405.06683*, 2024. URL <https://arxiv.org/abs/2405.06683>.
- [215] L. Siddharth and Jianxi Luo. Retrieval Augmented Generation using Engineering Design Knowledge. *arXiv preprint arXiv:2307.069850*, 2023. URL <https://arxiv.org/abs/2307.069850>.
- [216] Xiaoyu Song, William Han, Tony Chen, Chaojing Duan, Michael A. Rosenberg, Emerson Liu, and Ding Zhao. Retrieval-Augmented Generation for Electrocardiogram-Language Models. *arXiv preprint arXiv:2510.00261*, 2025. URL <https://arxiv.org/abs/2510.00261>.
- [217] Francesco Sovrano and Alberto Bacchelli. Illocutionary Explanation Planning for Source-Faithful Explanations in Retrieval-Augmented Language Models. *arXiv preprint arXiv:2604.06211*, 2026. URL <https://arxiv.org/abs/2604.06211>.
- [218] Weihang Su, Yichen Tang, Qingyao Ai, Zhijing Wu, and Yiqun Liu. DRAGIN: Dynamic Retrieval Augmented Generation based on the Information Needs of Large Language Models. *arXiv preprint arXiv:2403.10081*, 2024. URL <https://arxiv.org/abs/2403.10081>.
- [219] Weihang Su, Yichen Tang, Qingyao Ai, Junxi Yan, Changyue Wang, Hongning Wang, Ziyi Ye, Yujia Zhou, and Yiqun Liu. Parametric Retrieval Augmented Generation. *arXiv preprint arXiv:2501.15915*, 2025. URL <https://arxiv.org/abs/2501.15915>.
- [220] Chunyu Sun, Bingyu Liu, Zhichao Cui, Junhan Shi, Anbin Qi, Tian hao Zhang, Dinghao Zhou, and Lewei Lu. SEAL: Speech Embedding Alignment Learning for Speech Large Language Model with Retrieval-Augmented Generation. *arXiv preprint arXiv:2502.02603*, 2025. URL <https://arxiv.org/abs/2502.02603>.
- [221] Jiashuo Sun, Xianrui Zhong, Sizhe Zhou, and Jiawei Han. DynamicRAG: Leveraging Outputs of Large Language Model as Feedback for Dynamic Reranking in Retrieval-Augmented Generation. *arXiv preprint arXiv:2505.07233*, 2025. URL <https://arxiv.org/abs/2505.07233>.



- [222] Qixin Sun, Ziqin Wang, Hengyuan Zhao, Yilin Li, Kaiyou Song, Linjiang Huang, Xiaolin Hu, Qingpei Guo, and Si Liu. VaccineRAG: Boosting Multimodal Large Language Models’ Immunity to Harmful RAG Samples. *arXiv preprint arXiv:2509.04502*, 2025. URL <https://arxiv.org/abs/2509.04502>.
- [223] Yiyu Sun, Junjie Hu, Wei Cheng, and Haifeng Chen. DFA-RAG: Conversational Semantic Router for Large Language Model with Definite Finite Automaton. *arXiv preprint arXiv:2402.04411*, 2024. URL <https://arxiv.org/abs/2402.04411>.
- [224] Yixuan Tang and Yi Yang. MultiHop-RAG: Benchmarking Retrieval-Augmented Generation for Multi-Hop Queries. *arXiv preprint arXiv:2401.15391*, 2024. URL <https://arxiv.org/abs/2401.15391>.
- [225] Yannis Tevissen, Khalil Guetari, and Frederic Petitpont. Towards Retrieval Augmented Generation over Large Video Libraries. *arXiv preprint arXiv:2406.14938*, 2024. URL <https://arxiv.org/abs/2406.14938>.
- [226] Ryota Tozuka, Hisashi Johno, Akitomo Amakawa, Junichi Sato, Mizuki Muto, Shoichiro Seki, Atsushi Komaba, and Hiroshi Onishi. Application of NotebookLM, a Large Language Model with Retrieval-Augmented Generation, for Lung Cancer Staging. *arXiv preprint arXiv:2410.10869*, 2024. URL <https://arxiv.org/abs/2410.10869>.
- [227] Sourav Verma. Contextual Compression in Retrieval-Augmented Generation for Large Language Models: A Survey. *arXiv preprint arXiv:2409.13385*, 2024. URL <https://arxiv.org/abs/2409.13385>.
- [228] Sriram Veturi, Saurabh Vaichal, Reshma Lal Jagadheesh, Nafis Irtiza Tripto, and Nian Yan. RAG based Question-Answering for Contextual Response Prediction System. *arXiv preprint arXiv:2409.03708*, 2024. URL <https://arxiv.org/abs/2409.03708>.
- [229] Sergii Voloshyn. L-RAG: Balancing Context and Retrieval with Entropy-Based Lazy Loading. *arXiv preprint arXiv:2601.06551*, 2026. URL <https://arxiv.org/abs/2601.06551>.
- [230] Sridevi Wagle, Sai Munikoti, Anurag Acharya, Sara Smith, and Sameera Horawalavithana. Empirical evaluation of Uncertainty Quantification in Retrieval-Augmented Language Models for Science. *arXiv preprint arXiv:2311.09358*, 2023. URL <https://arxiv.org/abs/2311.09358>.
- [231] Nicholas Thomas Walker, Stefan Ultes, and Pierre Lison. Retrieval-Augmented Neural Response Generation Using Logical Reasoning and Relevance Scoring. *arXiv preprint arXiv:2310.13566*, 2023. URL <https://arxiv.org/abs/2310.13566>.
- [232] Fei Wang, Xingchen Wan, Ruoxi Sun, Jiefeng Chen, and Sercan O. Ark. Astute RAG: Overcoming Imperfect Retrieval Augmentation and Knowledge Conflicts for Large Language Models. *arXiv preprint arXiv:2410.07176*, 2024. URL <https://arxiv.org/abs/2410.07176>.
- [233] Haijin Wang, Mianrong Zhang, Zheng Chen, Nan Shang, Shangheng Yao, Fushuan Wen, and Junhua Zhao. Carbon Footprint Accounting Driven by Large Language Models and Retrieval-augmented Generation. *arXiv preprint arXiv:2408.09713*, 2024. URL <https://arxiv.org/abs/2408.09713>.
- [234] Haoran Wang, Xiong Xiao Xu, Baixiang Huang, and Kai Shu. Privacy-Aware Decoding: Mitigating Privacy Leakage of Large Language Models in Retrieval-Augmented Generation. *arXiv preprint arXiv:2508.03098*, 2025. URL <https://arxiv.org/abs/2508.03098>.
- [235] Zheng Wang, Shu Xian Teo, Jieer Ouyang, Yongjun Xu, and Wei Shi. M-RAG: Reinforcing Large Language Model Performance through Retrieval-Augmented Generation with Multiple Partitions. *arXiv preprint arXiv:2405.16420*, 2024. URL <https://arxiv.org/abs/2405.16420>.
- [236] Zhengchao Wang, Yitao Hu, Jianing Ye, Zhuxuan Chang, Jiazheng Yu, Youpeng Deng, and Keqiu Li. RAGPulse: An Open-Source RAG Workload Trace to Optimize RAG Serving Systems. *arXiv preprint arXiv:2511.12979*, 2025. URL <https://arxiv.org/abs/2511.12979>.
- [237] Zilong Wang, Zifeng Wang, Long Le, Huaixiu Steven Zheng, Swaroop Mishra, Vincent Perot, Yuwei Zhang, Anush Mattapalli, Ankur Taly, Jingbo Shang, Chen-Yu Lee, and Tomas Pfister. Speculative RAG: Enhancing Retrieval Augmented Generation through Drafting. *arXiv preprint arXiv:2407.08223*, 2024. URL <https://arxiv.org/abs/2407.08223>.
- [238] Ziyu Wang, Elahe Khatibi, and Amir M. Rahmani. MedCoT-RAG: Causal Chain-of-Thought RAG for Medical Question Answering. *arXiv preprint arXiv:2508.15849*, 2025. URL <https://arxiv.org/abs/2508.15849>.
- [239] Aniket Wattamwar and Mrunal Kakirwar. \$\$-RAG: Oblivious Retrieval via Semantic Quantization and Transcendental Addressing for Large Language Models. *arXiv preprint arXiv:2606.22153*, 2026. URL <https://arxiv.org/abs/2606.22153>.
- [240] Qikai Wei, Huansheng Ning, Chunlong Han, and Jianguo Ding. A Query-Aware Multi-Path Knowledge Graph Fusion Approach for Enhancing Retrieval-Augmented Generation in Large Language Models. *arXiv preprint arXiv:2507.16826*, 2025. URL <https://arxiv.org/abs/2507.16826>.

- [241] Di Wu, Jia-Chen Gu, Fan Yin, Nanyun Peng, and Kai-Wei Chang. Synchronous Faithfulness Monitoring for Trustworthy Retrieval-Augmented Generation. *arXiv preprint arXiv:2406.13692*, 2024. URL <https://arxiv.org/abs/2406.13692>.
- [242] Jiayi Wu, Hengyi Cai, Lingyong Yan, Hao Sun, Xiang Li, Shuaiqiang Wang, Dawei Yin, and Ming Gao. PA-RAG: RAG Alignment via Multi-Perspective Preference Optimization. *arXiv preprint arXiv:2412.14510*, 2024. URL <https://arxiv.org/abs/2412.14510>.
- [243] Jinyang Wu, Shuai Zhang, Feihu Che, Mingkuan Feng, Chuyuan Zhang, Pengpeng Shao, and Jianhua Tao. Pandora’s Box or Aladdin’s Lamp: A Comprehensive Analysis Revealing the Role of RAG Noise in Large Language Models. *arXiv preprint arXiv:2408.13533*, 2024. URL <https://arxiv.org/abs/2408.13533>.
- [244] Junde Wu, Jiayuan Zhu, Yunli Qi, Jingkun Chen, Min Xu, Filippo Menolascina, and Vicente Grau. Medical Graph RAG: Towards Safe Medical Large Language Model via Graph Retrieval-Augmented Generation. *arXiv preprint arXiv:2408.04187*, 2024. URL <https://arxiv.org/abs/2408.04187>.
- [245] Peilin Wu, Xinlu Zhang, Wenhao Yu, Xingyu Liu, Xinya Du, and Zhiyu Zoey Chen. Do Retrieval-Augmented Language Models Adapt to Varying User Needs? *arXiv preprint arXiv:2502.19779*, 2025. URL <https://arxiv.org/abs/2502.19779>.
- [246] Pingyu Wu, Daiheng Gao, Jing Tang, Huimin Chen, Wenbo Zhou, Weiming Zhang, and Nenghai Yu. MES-RAG: Bringing Multi-modal, Entity-Storage, and Secure Enhancements to RAG. *arXiv preprint arXiv:2503.13563*, 2025. URL <https://arxiv.org/abs/2503.13563>.
- [247] Wenlong Wu, Haofen Wang, Bohan Li, Peixuan Huang, Xinzhe Zhao, and Lei Liang. MultiRAG: A Knowledge-guided Framework for Mitigating Hallucination in Multi-source Retrieval Augmented Generation. *arXiv preprint arXiv:2508.03553*, 2025. URL <https://arxiv.org/abs/2508.03553>.
- [248] Yaxiong Wu, Jianyuan Bo, Yongyue Zhang, Sheng Liang, and Yong Liu. Query-Centric Graph Retrieval Augmented Generation. *arXiv preprint arXiv:2509.21237*, 2025. URL <https://arxiv.org/abs/2509.21237>.
- [249] Haizhou Xia. Diagnosing Evidence Utilization in Long-Context and Retrieval-Augmented Language Models under Matched Evidence Conditions. *arXiv preprint arXiv:2606.06758*, 2026. URL <https://arxiv.org/abs/2606.06758>.
- [250] Yikuan Xia, Jiazun Chen, Yirui Zhan, Suifeng Zhao, Weipeng Jiang, Chaorui Zhang, Wei Han, Bo Bai, and Jun Gao. ER-RAG: Enhance RAG with ER-Based Unified Modeling of Heterogeneous Data Sources. *arXiv preprint arXiv:2504.06271*, 2025. URL <https://arxiv.org/abs/2504.06271>.
- [251] Yuan Xia, Jingbo Zhou, Zhenhui Shi, Jun Chen, and Haifeng Huang. Improving Retrieval Augmented Language Model with Self-Reasoning. *arXiv preprint arXiv:2407.19813*, 2024. URL <https://arxiv.org/abs/2407.19813>.
- [252] Xingchen Xiao, Heyan Huang, Runheng Liu, and Jincheng Xie. MASS-RAG: Multi-Agent Synthesis Retrieval-Augmented Generation. *arXiv preprint arXiv:2604.18509*, 2026. URL <https://arxiv.org/abs/2604.18509>.
- [253] Miao Xie, Xiao Zhang, Yi Li, and Chunli Lv. Structure Guided Retrieval-Augmented Generation for Factual Queries. *arXiv preprint arXiv:2604.22843*, 2026. URL <https://arxiv.org/abs/2604.22843>.
- [254] Pengzhen Xie and Huizhi Liang. Graph Your Way to Inspiration: Integrating Co-Author Graphs with Retrieval-Augmented Generation for Large Language Model Based Scientific Idea Generation. *arXiv preprint arXiv:2602.22215*, 2025. URL <https://arxiv.org/abs/2602.22215>.
- [255] Guangzhi Xiong, Qiao Jin, Zhiyong Lu, and Aidong Zhang. Benchmarking Retrieval-Augmented Generation for Medicine. *arXiv preprint arXiv:2402.13178*, 2024. URL <https://arxiv.org/abs/2402.13178>.
- [256] Can Xu, Lingyong Yan, Jiayi Wu, Haosen Wang, Shuaiqiang Wang, Yuchen Li, Jizhou Huang, Dawei Yin, and Xiang Li. Adversarial Yet Cooperative: Multi-Perspective Reasoning in Retrieved-Augmented Language Models. *arXiv preprint arXiv:2601.04651*, 2026. URL <https://arxiv.org/abs/2601.04651>.
- [257] Chenhao Xu, Longxiang Gao, Yuan Miao, and Xi Zheng. Distributed Retrieval-Augmented Generation. *arXiv preprint arXiv:2505.00443*, 2025. URL <https://arxiv.org/abs/2505.00443>.
- [258] Chuankai Xu, Dongming Zhao, Bo Wang, and Hanwen Xing. Enhancing Retrieval-Augmented LMs with a Two-stage Consistency Learning Compressor. *arXiv preprint arXiv:2406.02266*, 2024. URL <https://arxiv.org/abs/2406.02266>.
- [259] Jian Xu, Sichun Luo, Xiangyu Chen, Haoming Huang, Hanxu Hou, and Linqi Song. RALL-Rec: Improving Retrieval Augmented Large Language Model Recommendation with Representation Learning. *arXiv preprint arXiv:2502.06101*, 2025. URL <https://arxiv.org/abs/2502.06101>.
- [260] Ran Xu, Hui Liu, Sreyashi Nag, Zhenwei Dai, Yaochen Xie, Xianfeng Tang, Chen Luo, Yang Li, Joyce C. Ho, Carl Yang, and Qi He. SimRAG: Self-Improving Retrieval-Augmented Generation for Adapting Large Language Models to Specialized

- Domains. *arXiv preprint arXiv:2410.17952*, 2024. URL <https://arxiv.org/abs/2410.17952>.
- [261] Ran Xu, Wenqi Shi, Yuchen Zhuang, Yue Yu, Joyce C. Ho, Haoyu Wang, and Carl Yang. Collab-RAG: Boosting Retrieval-Augmented Generation for Complex Question Answering via White-Box and Black-Box LLM Collaboration. *arXiv preprint arXiv:2504.04915*, 2025. URL <https://arxiv.org/abs/2504.04915>.
- [262] Shicheng Xu, Liang Pang, Mo Yu, Fandong Meng, Huawei Shen, Xueqi Cheng, and Jie Zhou. Unsupervised Information Refinement Training of Large Language Models for Retrieval-Augmented Generation. *arXiv preprint arXiv:2402.18150*, 2024. URL <https://arxiv.org/abs/2402.18150>.
- [263] Sun Xu, Tongkai Xu, Baiheng Xie, Li Huang, Qiang Gao, and Kunpeng Zhang. M-RAG: Making RAG Faster, Stronger, and More Efficient. *arXiv preprint arXiv:2603.26667*, 2026. URL <https://arxiv.org/abs/2603.26667>.
- [264] Yilong Xu, Jinhua Gao, Xiaoming Yu, Yuanhai Xue, Baolong Bi, Huawei Shen, and Xueqi Cheng. Training a Utility-based Retriever Through Shared Context Attribution for Retrieval-Augmented Language Models. *arXiv preprint arXiv:2504.00573*, 2025. URL <https://arxiv.org/abs/2504.00573>.
- [265] Jiaqi Xue, Mengxin Zheng, Yebowen Hu, Fei Liu, Xun Chen, and Qian Lou. BadRAG: Identifying Vulnerabilities in Retrieval Augmented Generation of Large Language Models. *arXiv preprint arXiv:2406.00083*, 2024. URL <https://arxiv.org/abs/2406.00083>.
- [266] Guohang Yan, Yue Zhang, Pinlong Cai, Ding Wang, Song Mao, Hongwei Zhang, Yaoze Zhang, Hairong Zhang, Xinyu Cai, and Botian Shi. Het-aRAG: Hybrid Deep Retrieval-Augmented Generation across Heterogeneous Data Stores. *arXiv preprint arXiv:2509.21336*, 2025. URL <https://arxiv.org/abs/2509.21336>.
- [267] Shi-Qi Yan, Jia-Chen Gu, Yun Zhu, and Zhen-Hua Ling. Corrective Retrieval Augmented Generation. *arXiv preprint arXiv:2401.15884*, 2024. URL <https://arxiv.org/abs/2401.15884>.
- [268] Kaiyu Yang, Aidan M. Swope, Alex Gu, Rahul Chalamala, Peiyang Song, Shixing Yu, Saad Godil, Ryan Prenger, and Anima Anandkumar. LeanDojo: Theorem Proving with Retrieval-Augmented Language Models. *arXiv preprint arXiv:2306.15626*, 2023. URL <https://arxiv.org/abs/2306.15626>.
- [269] Shiping Yang, Jie Wu, Wenbiao Ding, Ning Wu, Shining Liang, Ming Gong, Hongzhi Li, Hengyuan Zhang, Angel X. Chang, and Dongmei Zhang. Quantifying and Improving the Robustness of Retrieval-Augmented Language Models Against Spurious Features in Grounding Data. *arXiv preprint arXiv:2503.05587*, 2025. URL <https://arxiv.org/abs/2503.05587>.
- [270] Tianbo Yang, Mingqi Yan, Hongyi Zhao, and Tianshuo Yang. LemmaHead: RAG Assisted Proof Generation Using Large Language Models. *arXiv preprint arXiv:2501.15797*, 2025. URL <https://arxiv.org/abs/2501.15797>.
- [271] Xiao Yang, Kai Sun, Hao Xin, Yushi Sun, Nikita Bhalla, Xiangsen Chen, Sajal Choudhary, Rongze Daniel Gui, Ziran Will Jiang, Ziyu Jiang, Lingkun Kong, Brian Moran, Jiaqi Wang, Yifan Ethan Xu, An Yan, Chenyu Yang, Eting Yuan, Hanwen Zha, Nan Tang, Lei Chen, et al. CRAG – Comprehensive RAG Benchmark. *arXiv preprint arXiv:2406.04744*, 2024. URL <https://arxiv.org/abs/2406.04744>.
- [272] Jiayi Yao, Hanchen Li, Yuhua Liu, Siddhant Ray, Yihua Cheng, Qizheng Zhang, Kuntai Du, Shan Lu, and Junchen Jiang. CacheBlend: Fast Large Language Model Serving for RAG with Cached Knowledge Fusion. *arXiv preprint arXiv:2405.16444*, 2024. URL <https://arxiv.org/abs/2405.16444>.
- [273] Fuda Ye, Shuangyin Li, Yongqi Zhang, and Lei Chen. R<sup>2</sup>AG: Incorporating Retrieval Information into Retrieval Augmented Generation. *arXiv preprint arXiv:2406.13249*, 2024. URL <https://arxiv.org/abs/2406.13249>.
- [274] Hailong Yin, Bin Zhu, Jingjing Chen, and Chong-Wah Ngo. Efficient Test-Time Retrieval Augmented Generation. *arXiv preprint arXiv:2511.01059*, 2025. URL <https://arxiv.org/abs/2511.01059>.
- [275] Ori Yoran, Tomer Wolfson, Ori Ram, and Jonathan Berant. Making Retrieval-Augmented Language Models Robust to Irrelevant Context. *arXiv preprint arXiv:2310.01558*, 2023. URL <https://arxiv.org/abs/2310.01558>.
- [276] Chuanyue Yu, Jiahui Wang, Yuhua Li, Heng Chang, Ge Lan, Qingyun Sun, Jia Li, Jianxin Li, and Ziwei Zhang. Unlocking the Potentials of Retrieval-Augmented Generation for Diffusion Language Models. *arXiv preprint arXiv:2601.11342*, 2026. URL <https://arxiv.org/abs/2601.11342>.
- [277] Tian Yu, Shaolei Zhang, and Yang Feng. Auto-RAG: Autonomous Retrieval-Augmented Generation for Large Language Models. *arXiv preprint arXiv:2411.19443*, 2024. URL <https://arxiv.org/abs/2411.19443>.
- [278] Wenhao Yu, Hongming Zhang, Xiaoman Pan, Kaixin Ma, Hongwei Wang, and Dong Yu. Chain-of-Note: Enhancing Robustness in Retrieval-Augmented Language Models. *arXiv preprint arXiv:2311.09210*, 2023. URL <https://arxiv.org/abs/2311.09210>.

- [279] Xiaohan Yu, Zhihan Yang, and Chong Chen. Cog-Planner: Unveiling the Potential of Agentic Multimodal Retrieval Augmented Generation with Planning. *arXiv preprint arXiv:2501.15470*, 2025. URL <https://arxiv.org/abs/2501.15470>.
- [280] Jianhao Yuan, Shuyang Sun, Daniel Omeiza, Bo Zhao, Paul Newman, Lars Kunze, and Matthew Gadd. RAG-Driver: Generalisable Driving Explanations with Retrieval-Augmented In-Context Learning in Multi-Modal Large Language Model. *arXiv preprint arXiv:2402.10828*, 2024. URL <https://arxiv.org/abs/2402.10828>.
- [281] Cyril Zakka, Akash Chaurasia, Rohan Shad, Alex R. Dalal, Jennifer L. Kim, Michael Moor, Kevin Alexander, Euan Ashley, Jack Boyd, Kathleen Boyd, Karen Hirsch, Curt Langlotz, Joanna Nelson, and William Hiesinger. Almanac: Retrieval-Augmented Language Models for Clinical Medicine. *arXiv preprint arXiv:2303.01229*, 2023. URL <https://arxiv.org/abs/2303.01229>.
- [282] Huimin Zeng, Zhenrui Yue, Qian Jiang, and Dong Wang. Federated Recommendation via Hybrid Retrieval Augmented Generation. *arXiv preprint arXiv:2403.04256*, 2024. URL <https://arxiv.org/abs/2403.04256>.
- [283] Saber Zerhoubi and Michael Granitzer. Person-aRAG: Enhancing Retrieval-Augmented Generation Systems with User-Centric Agents. *arXiv preprint arXiv:2407.09394*, 2024. URL <https://arxiv.org/abs/2407.09394>.
- [284] Chenghao Zhang, Guanting Dong, Xinyu Yang, and Zhicheng Dou. Towards Mixed-Modal Retrieval for Universal Retrieval-Augmented Generation. *arXiv preprint arXiv:2510.17354*, 2025. URL <https://arxiv.org/abs/2510.17354>.
- [285] Chi Zhang, Qiyang Chen, and Mengqi Zhang. Mixture-of-RAG: Integrating Text and Tables with Large Language Models. *arXiv preprint arXiv:2504.09554*, 2025. URL <https://arxiv.org/abs/2504.09554>.
- [286] Gongbo Zhang, Zihan Xu, Qiao Jin, Fangyi Chen, Yilu Fang, Yi Liu, Justin F. Rousseau, Ziyang Xu, Zhiyong Lu, Chunhua Weng, and Yifan Peng. A MapReduce Approach to Effectively Utilize Long Context Information in Retrieval Augmented Language Models. *arXiv preprint arXiv:2412.15271*, 2024. URL <https://arxiv.org/abs/2412.15271>.
- [287] Gongbo Zhang, Yifan Peng, and Chunhua Weng. Improving Retrieval-Augmented Generation without Taxonomy-based Error Categorization. *arXiv preprint arXiv:2605.18772*, 2026. URL <https://arxiv.org/abs/2605.18772>.
- [288] Hairong Zhang, Jiaheng Si, Guohang Yan, Boyuan Qi, Pinlong Cai, Song Mao, Ding Wang, and Botian Shi. RAKG: Document-level Retrieval Augmented Knowledge Graph Construction. *arXiv preprint arXiv:2504.09823*, 2025. URL <https://arxiv.org/abs/2504.09823>.
- [289] Jiarui Zhang, Xiangyu Liu, Yong Hu, Chaoyue Niu, Fan Wu, and Guihai Chen. RAGRouter: Learning to Route Queries to Multiple Retrieval-Augmented Language Models. *arXiv preprint arXiv:2505.23052*, 2025. URL <https://arxiv.org/abs/2505.23052>.
- [290] Maosen Zhang, Jianshuo Dong, Boting Lu, Wenyue Li, Xiaoping Zhang, Tianwei Zhang, and Han Qiu. LeakDojo: Decoding the Leakage Threats of RAG Systems. *arXiv preprint arXiv:2605.05818*, 2026. URL <https://arxiv.org/abs/2605.05818>.
- [291] Qinggang Zhang, Shengyuan Chen, Yuanchen Bei, Zheng Yuan, Huachi Zhou, Zijin Hong, Hao Chen, Yilin Xiao, Chuang Zhou, Junnan Dong, Yi Chang, and Xiao Huang. A Survey of Graph Retrieval-Augmented Generation for Customized Large Language Models. *arXiv preprint arXiv:2501.13958*, 2025. URL <https://arxiv.org/abs/2501.13958>.
- [292] Ruizhe Zhang, Yongxin Xu, Yuzhen Xiao, Runchuan Zhu, Xinke Jiang, Xu Chu, Junfeng Zhao, and Yasha Wang. KnowPO: Knowledge-aware Preference Optimization for Controllable Knowledge Selection in Retrieval-Augmented Language Models. *arXiv preprint arXiv:2408.03297*, 2024. URL <https://arxiv.org/abs/2408.03297>.
- [293] Xin Zhang, Lissette Iturburu, Juan Nicolas Vilamizar, Xiaoyu Liu, Manuel Salmeron, Shirley J. Dyke, and Julio Ramirez. Large Language Model Agent for Structural Drawing Generation Using ReAct Prompt Engineering and Retrieval Augmented Generation. *arXiv preprint arXiv:2507.19771*, 2025. URL <https://arxiv.org/abs/2507.19771>.
- [294] Zhihao Zhang, Alan Zhu, Lijie Yang, Yihua Xu, Lanting Li, Phitchaya Mangpo Phothilimthana, and Zhihao Jia. Accelerating Retrieval-Augmented Language Model Serving with Speculation. *arXiv preprint arXiv:2401.14021*, 2024. URL <https://arxiv.org/abs/2401.14021>.
- [295] Zhuocheng Zhang, Yang Feng, and Min Zhang. FlexRAG: A Flexible and Comprehensive Framework for Retrieval-Augmented Generation. *arXiv preprint arXiv:2506.12494*, 2025. URL <https://arxiv.org/abs/2506.12494>.
- [296] Zongmeng Zhang, Yufeng Shi, Jinhua Zhu, Wengang Zhou, Xiang Qi, Peng Zhang, and Houqiang Li. Trustworthy Alignment of Retrieval-Augmented Large Language Models via Reinforcement Learning. *arXiv preprint arXiv:2410.16843*, 2024. URL <https://arxiv.org/abs/2410.16843>.
- [297] Shengwei Zhao, Jingwen Yao, Sitong Wei, Linhai Xu, Yuying Liu, Dong Zhang, Zhiqiang Tian, and Shaoyi Du. MMRAG-RFT: Two-stage Reinforcement Fine-tuning for Explainable Multi-modal



- Retrieval-augmented Generation. *arXiv preprint arXiv:2512.17194*, 2025. URL <https://arxiv.org/abs/2512.17194>.
- [298] Xiangrong Zhu, Yuexiang Xie, Yi Liu, Yaliang Li, and Wei Hu. Knowledge Graph-Guided Retrieval Augmented Generation. *arXiv preprint arXiv:2502.06864*, 2025. URL <https://arxiv.org/abs/2502.06864>.
- [299] Deyu Zou, Yongqiang Chen, Mufei Li, Siqu Miao, Chenxi Liu, Bo Han, James Cheng, and Pan Li. Weak-to-Strong GraphRAG: Aligning Weak Retrievers with Large Language Models for Graph-based Retrieval Augmented Generation. *arXiv preprint arXiv:2506.22518*, 2025. URL <https://arxiv.org/abs/2506.22518>.
- [300] Jiaru Zou, Dongqi Fu, Sirui Chen, Xinrui He, Zihao Li, Yada Zhu, Jiawei Han, and Jingrui He. RAG over Tables: Hierarchical Memory Index, Multi-Stage Retrieval, and Benchmarking. *arXiv preprint arXiv:2504.01346*, 2025. URL <https://arxiv.org/abs/2504.01346>.