

Multi-Sensor Fusion for Autonomous Driving Perception: A Comprehensive Survey

Abstract

Multi-modal sensor fusion has become the cornerstone of robust autonomous driving systems, enabling perception models to integrate complementary cues from heterogeneous sensors, including visual cameras, Light Detection and Ranging (LiDAR), and millimeter-wave (MMW) radar. As autonomous driving scales toward Level 4 and Level 5 automation, environmental perception must transition from isolated, modality-specific processing to unified, holistic scene understanding. This survey provides an exhaustive, structured overview of recent advancements in deep learning-based multi-sensor fusion paradigms. We systematically categorize fusion architectures by their theoretical foundations, spanning early (data-level), mid (feature-level), late (decision-level), and emerging dynamic fusion strategies. Particular emphasis is placed on dominant architectural trends, notably unified Bird's-Eye-View (BEV) representations, sequential point-decoration techniques, and query-based cross-modal attention mechanisms. Furthermore, we explore the adaptation of these architectures across diverse downstream perception tasks, including 3D object detection, semantic occupancy prediction, and trajectory forecasting. Moving beyond conventional single-agent paradigms, this paper deeply analyzes cutting-edge frontiers such as 4D imaging radar integration, Vehicle-to-Everything (V2X) cooperative perception, and robustness against adverse weather and sensor spoofing. Finally, we discuss the transformative potential of vision-language foundation models (VLMs) and end-to-end learning in unlocking the next generation of scalable, interpretable, and trustworthy autonomous perception systems.

1. Introduction

Autonomous driving represents a highly disruptive technological frontier poised to redefine modern transportation by mitigating traffic accidents, optimizing fuel efficiency, and enhancing global mobility¹. At the core of any Automated Driving System (ADS) lies the environmental perception module, which serves as the foundational precursor to downstream planning, decision-making, and vehicle control⁴. Because real-world Operational Design Domains (ODDs) are fraught with dynamic complexities—ranging from dense urban traffic and unpredictable pedestrian behavior to severe weather conditions—relying on a single sensing modality is fundamentally inadequate for ensuring certifiable safety³. To realize the stringent safety requirements of Level 4 and Level 5 autonomy defined by the SAE J3016 standard, systems must build a comprehensive, panoramic understanding of the ego-vehicle's surroundings³.

Consequently, modern autonomous vehicles are equipped with diverse sensor suites designed

to capitalize on the complementary physical properties of different sensing modalities¹. Visual cameras, the most mature and ubiquitous sensors, provide dense, high-resolution photometric data capturing rich semantic, textural, and color information. This is critical for recognizing traffic signs, lane markings, and pedestrian boundaries⁷. However, as passive sensors, cameras suffer from inherent depth ambiguity, limited dynamic range, and severe performance degradation in low-light, high-glare, or adverse weather conditions¹. Conversely, LiDAR operates as an active time-of-flight sensor, emitting laser pulses to generate precise, sparse 3D point clouds. While LiDAR provides unmatched geometric fidelity and direct depth measurements independent of ambient illumination, it is historically cost-prohibitive, computationally heavy, and highly susceptible to atmospheric scattering caused by rain, snow, and fog¹. Traditional millimeter-wave (MMW) radar offers robust, all-weather performance with long-range detection capabilities and direct radial velocity measurement via the Doppler effect. Yet, standard 3D radars suffer from extreme sparsity, lack of elevation data, and heavy clutter noise, limiting their standalone utility for complex semantic classification⁷. To overcome these isolated limitations, multi-sensor fusion has evolved into a critical research domain¹³. The primary objective of fusion is to synergize heterogeneous data streams, mathematically compensating for the deficiencies of individual sensors while maximizing overall perceptual redundancy and fault tolerance¹. Recent years have witnessed a massive paradigm shift from traditional rule-based filtering (e.g., Extended Kalman Filters) toward end-to-end deep representation learning. Innovations such as Bird's-Eye-View (BEV) transformations, Transformer-based cross-attention, and contrastive alignment have redefined how high-dimensional sensory data is aggregated and processed¹.

Sensor Modality	Primary Data Format	Key Strengths	Inherent Weaknesses
Camera	2D RGB/RGB-D Images	High semantic richness, dense texture, color, low cost.	Depth ambiguity, vulnerable to lighting and weather.
LiDAR	3D Point Clouds	High geometric precision, accurate depth, illumination invariant.	Sparse at long range, high cost, degraded by rain/snow/fog.
3D Radar	Range-Azimuth Heatmaps	All-weather robustness, Doppler velocity, high penetration.	Extreme sparsity, lack of elevation, noisy clutter, poor semantics.

4D Radar	4D Point Clouds	Adds elevation to radar, dense returns, reliable Doppler tracking.	Still sparser than LiDAR, requires complex noise filtering.
Ultrasonic	1D Distance	Ultra-low cost, excellent for near-field obstacle detection.	Very short range, low resolution, highly directional.

This survey provides a comprehensive and in-depth review of the state-of-the-art in multi-sensor fusion for autonomous driving perception. Section 2 details the architectural methodologies and theoretical frameworks governing modality integration. Section 3 reviews the application of these methods across core downstream perception tasks. Section 4 highlights emerging frontiers, including 4D radar, cooperative V2X perception, and physical robustness. Finally, Section 5 outlines current challenges and future prospects, particularly the integration of large foundation models, before Section 6 concludes the manuscript.

2. Methods: Theoretical Foundations and Fusion Architectures

The architecture of multi-sensor fusion networks dictates how and when information from divergent sources is synthesized. Historically categorized by the abstraction level of the merged data, modern deep learning approaches have increasingly blurred these lines, opting for unified spatial representations and dynamic routing.

2.1 Theoretical Frameworks and the Evolution of Fusion Stages

The theoretical foundation of multimodal sensor fusion builds upon the confluence of estimation theory, information theory, and deep representation learning¹. Fundamentally, the fusion process aims to compute a posterior distribution of the environmental state given multiple heterogeneous sensor observations, accounting for cross-modal redundancy and temporal complementarity. Early theoretical frameworks relied heavily on Bayesian estimation and probabilistic filtering. Under this classical framework, each sensor contributes a likelihood term, and the posterior is updated through recursive filtering. Common implementations include the Extended Kalman Filter (EKF) for nonlinear systems and Particle Filters (PF) for nonparametric estimation, which are still extensively used for sensor fusion in localization and tracking¹.

In deep learning-based multi-view learning, fusion strategies are broadly categorized into specific stages based on when the representations are aggregated:

Early Fusion (Data-Level): Modalities are combined at the raw data level before feature extraction. A common methodology involves projecting LiDAR point clouds onto an image

plane to create dense RGB-Depth inputs. While early fusion preserves maximal raw information and provides strong spatial correlation, it is highly sensitive to spatiotemporal misalignment, intrinsic calibration errors, and sensor noise¹.

Mid-Level Fusion (Feature-Level): Currently the most dominant paradigm, mid-level fusion utilizes modality-specific encoders to extract high-dimensional latent representations, which are subsequently aggregated. By leveraging conditional correlation among sensors, neural networks can learn rich spatial and semantic alignments, effectively filtering out modality-specific noise before the final decision layers¹. Feature-level fusion architectures benefit heavily from cross-attention mechanisms and unified Cartesian mapping.

Late Fusion (Decision-Level): Each sensor stream independently produces predictions, such as bounding boxes or semantic segmentation masks, which are then fused using heuristics like Non-Maximum Suppression (NMS) or Bayesian updating. Late fusion offers high fault tolerance and modularity, isolating the failure of one sensor from the rest of the network. However, it fails to exploit cross-modal geometric correlations during the feature extraction phase, often resulting in suboptimal performance in highly occluded scenarios where partial observations could have been combined¹.

Dynamic and Multi-Stage Fusion Architectures: Recent advancements recognize that static fusion strategies are often suboptimal under varying environmental conditions. Dynamic frameworks, such as the Feature Disagreement Score Network (FDSNet), conditionally adapt the fusion stage based on real-time semantic consistency metrics¹⁹. By calculating the spatial variance—or disagreement—across modality-specific BEV feature maps, FDSNet quantifies the extent to which one modality deviates from the average representation. The network applies mid-level fusion for confident, consistent scenes, taking advantage of deep feature aggregation, but shifts to late-stage decision fusion in regions with high ambiguity, such as localized weather degradation or sensor occlusion, thereby preventing the propagation of cross-modal noise¹⁹. Similarly, multi-stage frameworks like MS-Occ employ hierarchical fusion pipelines. In the middle stage, dense geometric priors derived from LiDAR are injected into camera features, while a late-stage voxel fusion module utilizes self-attention to resolve residual semantic ambiguities and dynamically balance geometric and semantic characteristics¹⁸.

2.2 Bird's-Eye-View (BEV) Centric Fusion

A transformative trend in autonomous perception is the transition to Bird's-Eye-View (BEV) representations. The BEV paradigm elegantly solves the spatial misalignment between 2D perspective images and 3D point clouds by mapping all heterogeneous inputs into a shared, top-down Cartesian coordinate system¹. This mapping establishes a unified physical reference frame, providing consistent geometry for downstream tasks such as object detection, motion forecasting, and path planning, which inherently occur in Euclidean space¹.

To fuse camera and LiDAR data in BEV space, visual features must be geometrically "lifted" from the 2D image plane to 3D space. The widely adopted Lift-Splat-Shoot (LSS) paradigm predicts categorical depth distributions for each image pixel, computing the outer product of image features and depth probabilities to generate a 3D frustum. These frustums are then "splatted"

into a BEV grid via voxel pooling²¹. In parallel, LiDAR points are processed via voxelization, such as in VoxelNet, or pillarization, as seen in PointPillars, and scattered into the same BEV grid²⁵. Architectures like BEVFusion have decisively demonstrated that performing fusion in this shared BEV space preserves both the visual semantic density of cameras and the geometric structure of LiDAR. By maintaining this structure, BEV-centric fusion significantly outperforms prior mid-level methods that relied on globally pooled one-dimensional features, which inherently destroyed spatial grounding¹.

Despite its perceptual superiority, dense BEV fusion is computationally expensive. Constructing a high-resolution BEV grid involves massive memory bandwidth, particularly during voxel pooling operations¹. To mitigate this bottleneck, researchers have introduced content-aware pruning mechanisms and adaptive scheduling. By utilizing the BEV map as a unified anchor representation, networks can dynamically predict and prune computationally expensive yet task-irrelevant regions, such as the sky, building facades, and distant static objects, before they enter the heavy backbone processors. Empirical results indicate that adaptive joint input pruning can reduce computational overhead by up to 50% without compromising the detection of safety-critical dynamic objects²². Furthermore, real-time scheduling frameworks like FLEX adaptively allocate limited machine attention among different spatial views partitioned by camera facing directions. This elastic fusion strategy dynamically adjusts the fusion configuration based on the available computation time budget and immediate context²⁷.

2.3 Query-Based and Transformer Fusion

While BEV grids provide explicit geometric alignment, Transformer-based architectures offer an implicit, learning-driven alternative by treating multi-sensor fusion as a sequence-to-sequence or set prediction problem¹. Leveraging cross-modal attention mechanisms, these architectures dynamically route information across sensors based on learned spatial and semantic correlations, eliminating the need for rigid coordinate transformations.

Frameworks such as FUTR3D bypass explicit view transformations and heuristic late fusion altogether. FUTR3D employs a modality-agnostic feature sampler (MAFS) combined with a DETR-like Transformer decoder. Instead of building a dense 3D grid, the network defines a sparse set of learnable 3D object queries. Each query projects its 3D reference point onto the encoded feature maps of every available sensor, including cameras, low-resolution LiDAR, high-resolution LiDAR, and radar. The Transformer decoder then aggregates multi-modal features via cross-attention. This mechanism naturally handles arbitrary sensor configurations in a modality-agnostic way and demonstrates extreme robustness against individual sensor failure. For instance, FUTR3D utilizing only a low-cost 4-beam LiDAR alongside cameras achieves parity with state-of-the-art models relying on expensive 32-beam LiDARs³¹.

The Cross-Modal Transformer (CMT) further simplifies the pipeline by directly encoding 3D spatial coordinates into the multi-modal tokens (image patches and point cloud clusters). Object queries subsequently interact with these position-encoded, multi-modal tokens simultaneously through attention mechanisms. By implicitly embedding 3D geometry into the feature space through a Coordinates Encoding Module (CEM), CMT avoids the quantization

errors and computational bottlenecks associated with grid sampling and voxel-pooling. CMT achieves state-of-the-art inference speeds while maintaining robust performance, acting nearly seamlessly as a vision-only detector when the LiDAR signal is abruptly lost¹⁶. Similarly, the TransFuser architecture integrates image and LiDAR representations using self-attention at multiple resolutions. Rather than relying on simple geometric projections, TransFuser incorporates global context and pairwise interactions into the feature extraction layers. Tested on the CARLA simulator's Longest6 benchmark, TransFuser demonstrated that imitation learning policies based on transformer fusion drastically reduce collisions in complex urban scenarios characterized by dense, dynamic traffic, outperforming geometry-based fusion baselines by massive margins³⁰.

2.4 Sequential and Point-Decoration Fusion

An alternative paradigm to unifying modalities in a shared latent space is to enhance the raw data of one modality with the processed semantics of another. This family of methods, often referred to as point-decoration or sequential fusion, primarily relies on augmenting sparse LiDAR geometry with rich 2D image semantics.

The seminal PointPainting architecture operates by first processing 2D images through an off-the-shelf semantic segmentation network. The LiDAR point cloud is then projected into the camera view, and the resulting pixel-wise class scores are concatenated directly onto the LiDAR points, appending semantic probabilities to the standard Cartesian coordinates and intensity channels. The resulting "painted" point cloud is subsequently fed into standard LiDAR-only 3D detectors like PointPillars or VoxelNet. This sequential approach is highly modular, avoids the feature-blurring effects associated with projecting sparse 3D features into dense 2D space, and significantly improves the detection of small or hard-to-see classes such as pedestrians and cyclists³⁸.

While PointPainting is strictly unidirectional, transferring features from image to point, advanced architectures like EPNet and EPNet++ introduce bi-directional cascade fusion. These models establish point-to-pixel correspondences that allow dense image features to enrich point clouds, while simultaneously enabling sparse, geometrically accurate point features to flow inversely back into the 2D image branch. This cascade bi-directional fusion ensures that the gradients back-propagated from auxiliary 2D tasks, such as normalized local coordinates map estimation, explicitly boost the representation capacity of the 3D backbone. By explicitly guaranteeing consistency between predicted scores from two modalities through a Multi-Modal Consistency (MC) loss, EPNet++ generates highly discriminative feature representations, proving particularly effective in environments with highly sparse point cloud returns⁴².

3. Methods: Downstream Perception Tasks

The architectural paradigms of multi-sensor fusion are fundamentally tailored to support specific downstream perception tasks, converting abstract representations into actionable scene understanding for the autonomous agent.

3.1 3D Object Detection

3D object detection is the quintessential perception task, requiring the precise localization and classification of entities in 3D space, typically represented by oriented bounding boxes characterized by dimensions, location, and heading angle⁷. Fusion architectures have catalyzed the transition from anchor-based detectors, which rely on predefined bounding box priors, to anchor-free and center-based models. In the context of LiDAR-camera fusion, methods frequently adopt a center-point paradigm in the BEV space, where the network predicts a heatmap of object centers and regresses box attributes directly from the center features²³. A persistent challenge in 3D detection is handling small or heavily occluded objects. Geometric guidance and semantic-aware sampling prove vital here. Frustum-based networks, such as FrustumFusionNets, leverage 2D bounding boxes generated by a mature image detector to extrude a 3D frustum into the point cloud. This effectively narrows the search space, allowing the network to focus computational resources solely on regions of interest identified by the high-resolution camera⁴³. Furthermore, addressing Out-Of-Distribution (OOD) scenarios in monocular and fused 3D detection has driven research into Monocular Test-Time Adaptation (MonoTTA). By implementing reliability-driven adaptation, models dynamically optimize high-score detections during deployment to enhance overall confidence and mitigate the severe omissions that occur when test environments diverge from training distributions due to weather or camera corruptions⁴⁷.

3.2 3D Semantic Occupancy Prediction

While 3D bounding boxes discretize the world into specific foreground objects, autonomous vehicles must navigate complex, irregular environments that defy standard bounding box parameterization, such as construction zones, overhanging debris, and uniquely shaped vehicles. 3D semantic occupancy prediction addresses this by voxelizing the entire 3D space and assigning a semantic class and occupancy state to every voxel, providing a continuous volumetric representation of the scene¹⁸.

Multi-modal fusion is imperative for occupancy prediction because vision-centric methods suffer from severe depth ambiguities that lead to geometric inaccuracies and "floating" artifacts, while LiDAR-only methods yield precise geometries but semantically impoverished voxels¹⁸. Frameworks like MS-Occ perform multi-stage fusion specifically tailored for this task. By leveraging Gaussian kernel rendering on sparse LiDAR depth maps, they inject dense geometric priors into the 2D camera features. Subsequently, deformable attention transfers the enriched image semantics directly into the 3D LiDAR voxels, drastically improving intersection-over-union (IoU) scores for safety-critical, irregular geometries¹⁸.

To mitigate the massive computational burden of processing dense 3D grids, researchers are exploring alternative representations. Dual Range-Voxel Representations (DRVR) project point clouds into a spherical range-view image to extract compact, dense contextual features using standard 2D convolutions, and then seamlessly fuse these with sparse 3D voxel geometries. This facilitates highly accurate semantic occupancy prediction from a single-sweep point cloud without relying on computationally heavy multi-sweep temporal aggregation⁴⁹. Additionally, approaches like Semantic and Depth-Guided BEV Transformation (SDG-OCC) and RenderFusion introduce 3D Gaussian Splatting as an auxiliary supervision mechanism,

rendering semantic occupancy into 2D spaces to align image and LiDAR features efficiently, allowing for real-time processing of unstructured scenes²⁴.

3.3 Trajectory Prediction and Multi-Object Tracking

Beyond static scene understanding, autonomous driving requires accurate spatiotemporal modeling. Fusion models must maintain temporal consistency across frames to track objects and forecast future trajectories. Standard approaches utilize 3D Kalman Filters paired with data association algorithms, such as Hungarian matching, on the output of frame-by-frame 3D detectors⁵¹. However, multi-sensor systems inherently suffer from sensor asynchrony; cameras and LiDARs operate at different frequencies and capture data at disparate timestamps⁵³. Architectures like AsyncBEV and MTC-BEV incorporate specialized temporal fusion modules to address asynchrony. By employing explicit ego-motion compensation and temporal Transformers, these networks accurately align past BEV features with current observations, showing remarkable robustness against significant time offsets⁵³. Furthermore, sequence modeling techniques are transitioning from standard attention mechanisms to state-space models like Mamba. Frameworks such as PIE utilize Bidirectional Mamba fusion alongside a Mixture-of-Experts (MoE) decoder. This allows for efficient, long-horizon temporal feature aggregation without the quadratic complexity of traditional attention, ensuring highly accurate trajectory forecasting and dynamic interactions modeling for complex traffic scenarios⁵⁵. Furthermore, multi-task systems like UniMT combine a DETR-like detection decoder with a Reference Trajectory-based Multi-Class Transformer (RTMCT) to predict flexible-length future trajectories for diverse object classes concurrently with detection⁵⁶.

4. Emerging Frontiers and Specialized Domains

As foundational fusion methodologies mature, the research focus is expanding to integrate novel hardware, coordinate across multiple agents, and ensure absolute safety under worst-case environmental conditions.

4.1 4D Millimeter-Wave Radar Integration

Traditional 3D MMW radar is prized for its low cost, extreme weather resilience, and Doppler velocity measurement, but it is heavily criticized for extreme point cloud sparsity and a complete lack of vertical resolution. The advent of 4D imaging radar—which introduces elevation measurement (pitch angle resolution) and higher point density alongside range, azimuth, and Doppler velocity—has catalyzed a new wave of radar-camera fusion research¹¹. Despite these improvements, 4D radar point clouds remain significantly sparser than LiDAR, averaging hundreds of points per scan versus tens of thousands⁷. Direct application of LiDAR-centric fusion models to 4D radar often results in catastrophic feature sparsity and latency. To address this, specialized 4D radar-camera fusion frameworks, such as RPGFusion and RCGDet3D, have been developed to exploit the unique physical priors of radar. RPGFusion leverages radar Doppler velocity and Radar Cross Section (RCS) reflectivity to generate spatial confidence maps. These maps guide the image-to-BEV transformation, preventing uneven feature distribution and densifying sparse radar pillars prior to cross-modal attention⁵⁷.

Furthermore, emerging representations like 3D Gaussian Splatting are being adapted specifically for 4D radar. RCGDet3D predicts ray-centric Gaussian primitives aligned with the radar coordinate system, allowing the model to accurately reconstruct dense geometric features. By decoupling coordinate transformations from semantic representation learning and utilizing Semantic Injection (SI) from RGB images, this method achieves high geometric fidelity and real-time inference speed without relying on computationally heavy dense feature maps²⁴. Other approaches, such as the Doppler-Aware Radar Encoder (DARE), enhance motion-sensitive features via velocity-guided attention, systematically exploiting temporal information to model dynamic object motion accurately¹¹.

4.2 Cooperative Perception in V2X Systems

Single-agent perception is fundamentally constrained by line-of-sight physical occlusions and limited sensor ranges. Cooperative Perception (CP) leverages Vehicle-to-Everything (V2X) communication, allowing connected autonomous vehicles (CAVs) and intelligent roadside infrastructure (RSUs) to share sensory data. This establishes a holistic, beyond-line-of-sight situational awareness that profoundly expands the operational design domain of autonomous fleets⁶.

Collaboration Strategies and Communication Constraints: CP strategies parallel the stages of sensor fusion:

1. *Early Collaboration:* Agents share raw point clouds or images. While this maximizes spatial accuracy, it demands prohibitive communication bandwidth, rendering it largely infeasible for real-world wireless networks⁶.
2. *Intermediate Collaboration:* Agents process raw data into compressed latent features or BEV maps before transmission. This strategy offers an optimal trade-off between bandwidth and accuracy, making it the dominant paradigm in recent research (e.g., V2VNet, CoBEVT)⁵².
3. *Late Collaboration:* Agents transmit only final bounding boxes or tracked trajectories. While highly bandwidth-efficient, late collaboration struggles to detect objects that are only partially visible to multiple agents and cannot exploit feature-level synergies⁶. To bridge this gap, algorithms like DMSTrack employ differentiable multi-sensor Kalman filters to improve tracking accuracy while maintaining the minimal communication costs associated with late fusion⁵².

Heterogeneous Agent Fusion Architectures: Fusing data from multiple moving agents introduces severe challenges regarding asynchronous transmission delays, localization (pose) errors, and vastly varying sensor viewpoints. The V2X-ViT (Vehicle-to-Everything Vision Transformer) addresses this by alternating heterogeneous multi-agent self-attention layers with multi-scale window self-attention⁶⁰. V2X-ViT incorporates a delay-aware positional encoding mechanism to mathematically compensate for the time-lag of received packets, robustly fusing features even amidst GPS noise and high communication latency⁶⁰. Moreover, adaptive collaboration frameworks like HydraCollab dynamically adjust the communication payload based on spatial confidence maps. Rather than transmitting uniform feature grids across the network, agents selectively transmit only the most informative,

high-uncertainty sensor features, toggling between intermediate and late collaboration strategies on the fly to maximize perception gains while minimizing communication overhead⁶¹. For highly occluded intersections, Infrastructure-to-Vehicle (I2V) cooperation has proven particularly powerful. Roadside sensors provide an elevated, unobstructed view that seamlessly complements the ego-vehicle's ground-level perspective. Class-adaptive fusion frameworks selectively amplify these viewpoint advantages, yielding massive improvements in the detection of vulnerable road users like pedestrians and cyclists compared to uniform fusion strategies⁵⁹.

4.3 Robustness in Adverse Weather and Sensor Malfunction

The ultimate test of an automated driving system is its fail-operational capability during unexpected hardware malfunction or extreme weather conditions³. Standard LiDAR sensors suffer catastrophic signal degradation in rain, snow, and fog due to beam attenuation and backscatter, while cameras simultaneously lose contrast, color accuracy, and edge fidelity¹.

Weather-Aware Denoising and Fusion: Recent algorithms specifically target weather-induced corruption to maintain environmental awareness. Fog-Aware Feature Denoising Modules (FADM) utilize cross-modal attention to explicitly suppress modality inconsistencies. In heavy rain, where LiDAR returns are decimated and the frustum support for geometry-based checks is severely thinned, these modules dynamically increase the network's reliance on radar Doppler cues and unoccluded camera semantics to maintain spatial tracking¹².

Adversarial and Spoofing Defenses: Beyond natural phenomena, fusion systems must be hardened against malicious adversarial attacks. LiDAR spoofing can inject fake points to create "ghost objects" or actively cancel out returning beams to execute hiding attacks¹². Empirical research indicates that adverse weather actually exacerbates these vulnerabilities. Rain naturally thins point clouds, creating voids that mimic occlusion shadows. This allows attackers to execute successful hiding attacks with a significantly reduced "attack budget" (i.e., injecting fewer points), completely bypassing physical-invariant defenses and temporal-consistency checks¹². Developing fusion architectures that rigorously cross-verify semantic boundaries in the camera stream against geometric returns in the LiDAR and velocity vectors in the radar is a critical ongoing endeavor for securing autonomous perception against interference. Fault injection testing methodologies, such as FADE, are increasingly utilized to systematically assess the system-level behaviors of perception modules under synthetic sensor faults⁶⁸.

5. Challenges and Prospects

While the progress in multi-sensor fusion has been exponential, the deployment of Level 4 and 5 autonomous vehicles requires overcoming several lingering theoretical and practical bottlenecks related to reasoning, efficiency, and real-world domain adaptation.

5.1 Integration of Large Foundation Models (LLMs/VLMs)

The next major frontier in autonomous driving perception is the integration of multi-modal foundation models, particularly Large Language Models (LLMs) and Vision-Language Models

(VLMs)⁶⁹. Traditional perception models are constrained by closed-set vocabularies and rigidly defined bounding boxes, causing them to struggle immensely when interpreting rare, long-tail corner cases—such as a pedestrian carrying a large, irregularly shaped object, or an overturned truck blocking multiple lanes⁷¹.

VLMs introduce open-vocabulary understanding, zero-shot adaptation, and deep semantic reasoning. By aligning LiDAR and camera features within the latent space of a pre-trained VLM (such as CLIP or LLaVA), networks can leverage vast, internet-scale knowledge to interpret complex traffic scenarios⁷¹. Frameworks like SAM4D and DriveTransformer translate this linguistic context directly into spatial segmentation and planning operations. SAM4D introduces Unified Multi-modal Positional Encoding (UMPE) to align camera and LiDAR features, enabling seamless cross-modal prompting and interaction via text⁷².

Furthermore, Foundation Models enable "chain-of-thought" reasoning. Moving beyond simple regression outputs, models like DriveVLM and datasets like nuReasoning provide natural language explanations for driving decisions, bridging the gap between raw multi-modal perception and human-like contextual reasoning. These end-to-end (E2E) autonomous driving models integrate perception, prediction, and planning into a single differentiable framework, utilizing language-augmented maps and predictive intent to proactively respond to dynamic environments⁷⁰.

5.2 Computational Efficiency and Edge Deployment

Despite the elegance of advanced Transformer architectures and dense 3D BEV grids, their staggering computational overhead poses a severe challenge for onboard vehicle edge computing. Deploying massive models within the strict power budgets (often <100W) and thermal constraints of real-world electric vehicles remains a significant practical hurdle¹.

Future research must prioritize computationally efficient fusion paradigms. Innovations in hardware-aware neural architecture search, aggressive quantization, and dynamic token sparsification are paramount. Algorithms like FLEX demonstrate the potential of real-time, context-aware scheduling, adaptively allocating limited machine attention among different spatial views based on the immediate threat level of the environment²⁷. The integration of state-space models offers another promising pathway; utilizing architectures like Mamba achieves the global receptive field and sequence modeling capabilities of Transformers but with linear time complexity, substantially reducing GPU memory utilization during inference⁵⁵. Furthermore, completely parallel end-to-end architectures like PARA-Drive significantly enhance runtime speeds by simplifying complex query interactions into unified self-attention blocks⁷⁴.

5.3 Standardization, Datasets, and the Sim-to-Real Gap

The rapid iteration of perception algorithms relies heavily on comprehensive datasets. However, the lack of standardized evaluation metrics and the persistent simulation-to-real (sim-to-real) gap hinder the deployment of robust fusion models. To contextualize the current landscape, Table 1 outlines several highly influential datasets driving recent innovations.

Dataset Name	Key Modalities	Focus Area	Unique Contributions
nuScenes [cite: 21, 50, 81]	Camera, LiDAR, Radar	Urban Driving	Benchmark for 3D detection and 3D semantic occupancy.
Waymo Open [cite: 3, 82]	High-Res LiDAR, Camera	Diverse Environments	Large-scale, high-density point clouds; long-range perception.
TJ4D / VoD [cite: 11, 26, 57]	4D Radar, Camera, LiDAR	Radar Integration	Benchmarks for 4D imaging radar, capturing Doppler and elevation.
DAIR-V2X [cite: 66, 83]	Camera, LiDAR (V2X)	Cooperative Perception	Large-scale Vehicle-to-Infrastructure (V2I) collaboration.
V2V4Real [cite: 52, 83]	Camera, LiDAR (V2V)	Real-World V2V	Multi-object tracking and tracking ID annotations across vehicles.
UrbanIng-V2X [cite: 84]	Thermal, Camera, LiDAR	Urban Intersections	Dense multi-agent infrastructure tracking with thermal cameras.
MAN TruckScenes [cite: 82]	4D Radar, Camera, LiDAR	Heavy Vehicles	360-degree 4D radar, long-range annotations for trucks.

nuReasoning [cite: 76]	VLM annotations, V2X	Foundation Models	247K reasoning annotations for spatial/decision reasoning.
----------------------------------	-------------------------	-------------------	---

Models trained on pristine, well-calibrated data often fail when subjected to real-world sensor drift, dynamic misalignment, and hardware degradation over time³. Developing unsupervised domain adaptation techniques—such as query-based adversarial learning frameworks that exploit the complementary nature of image and BEV features to continuously self-calibrate and adapt to new environments without human intervention—is critical for achieving truly robust, generalizable autonomy⁸⁸.

6. Conclusion

Multi-sensor fusion has fundamentally redefined the landscape of autonomous driving perception. The evolution from disjointed, late-fusion heuristics to mathematically elegant, end-to-end representations—anchored by Bird's-Eye-View transformations, sequential point-decorations, and cross-modal attention mechanisms—has enabled unprecedented levels of geometric and semantic scene understanding. The capability to seamlessly integrate dense visual context with precise 3D geometry has pushed the boundaries of object detection, semantic occupancy prediction, and complex trajectory forecasting. As the field rapidly expands to incorporate 4D imaging radar, V2X cooperative infrastructure, and the vast reasoning capabilities of multi-modal foundation models, perception systems are becoming increasingly resilient to occlusion, adverse weather, and long-tail anomalies. However, realizing the ultimate goal of ubiquitous, Level 5 autonomy will require sustained interdisciplinary efforts to optimize the computational efficiency of these fusion pipelines for edge deployment, bridge the persistent sim-to-real gap, and mathematically guarantee fail-operational safety under all conceivable, dynamic driving conditions.

7. References

1. O. Natan and J. Miura, "Towards Compact Autonomous Driving Perception with Balanced Learning and Multi-sensor Fusion," *arXiv preprint arXiv:2606.02979*, 2026.
2. Keleti Károly Faculty of Business and Management, et al., "Multimodal Sensor Fusion in Autonomous Vehicles: Technologies," *Sensors*, vol. 25, no. 19, p. 6033.
3. Anonymous, "Table of Contents - Tasks and Open Competitions," *arXiv preprint arXiv:2202.02703v3*, 2022.
4. S. Ruan, et al., "A Survey of Multi-sensor Fusion Perception for Embodied AI: Background, Methods, Challenges and Prospects," *arXiv preprint arXiv:2506.19769*, 2025.
5. "A Survey of Multi-sensor Fusion Perception for Embodied AI: Background, Methods, Challenges and Prospects," *arXiv preprint arXiv:2506.19769v1*, 2025.
6. C. Wei, et al., "Integrating Multi-Modal Sensors: A Review of Fusion Techniques for Intelligent Vehicles," *arXiv preprint arXiv:2506.21885*, 2025.

7. Anonymous, "Multi-Sensor Fusion in Autonomous Driving," Atlantis Press, Article 126016749.
8. Anonymous, "Multi-sensor fusion object detection," *PMC*, PMC12074113.
9. T. Lan, "Research on Perception Optimization of Multi-Sensor Fusion in Unmanned Vehicle-Road Coordination," *Journal of Big Data and Computing*, vol. 3, no. 3, 2025.
10. Data Insights Market, "Autonomous Vehicle Lidar Sensors Market," Report 834691.
11. Mordor Intelligence, "Sensor Fusion Market In Autonomous Vehicles."
12. Anonymous, "Multi-sensor fusion object detection algorithm," *Sensors*, vol. 25, no. 9, p. 2794.
13. Solution1, "Comprehensive Review of the Autonomous Vehicle Industry in 2025."
14. Basic AI, "15 Notable AI Datasets for Autonomous Driving in 2024-2025."
15. Anonymous, "Unified BEV Framework for Beam Prediction," *arXiv preprint arXiv:2604.05668v1*, 2026.
16. Anonymous, "BEVCar: Joint BEV Object and Map Segmentation," *arXiv preprint arXiv:2403.11761v2*, 2024.
17. A. Singh, "Vision-RADAR fusion for Robotics BEV Detections: A Survey," *arXiv preprint arXiv:2302.06643*, 2023.
18. Anonymous, "RaLiBEV: Radar and LiDAR BEV Fusion," *arXiv preprint arXiv:2211.06108v5*, 2022.
19. D. Wu, F. Yang, and B. Liu, "A Survey of Deep Learning Based Radar and Vision Fusion for 3D Object Detection in Autonomous Driving," *arXiv preprint arXiv:2406.00714v1*, 2024.
20. Anonymous, "Radar and Camera Fusion for Object Detection and Tracking: A Comprehensive Survey," *arXiv preprint arXiv:2410.19872v1*, 2024.
21. L. Kong, et al., "Generative world models for autonomous driving."
22. Anonymous, "Rain-aware threat model focusing on LiDAR spoofing attacks," *ASTAR OAR**.
23. Z. Wei, et al., "MS-Occ: Multi-Stage LiDAR-Camera Fusion for 3D Semantic Occupancy Prediction," *arXiv preprint arXiv:2504.15888v1*, 2025.
24. Y. Song and A. L. Wang, "BiCo-Fusion: Bidirectional Complementary LiDAR-Camera Fusion."
25. Anonymous, "PIE: Bidirectional Mamba Fusion," *arXiv preprint arXiv:2509.18609v1*, 2025.
26. Robotics Research Centre (RRC), "Autonomous Vehicle and UAV Robotics," *NTU*.
27. O. Natan and J. Miura, "We present a novel compact deep multi-task learning model," *arXiv preprint arXiv:2606.02979v1*, 2026.
28. Anonymous, "Systematic analysis of the machine learning and deep learning models underpinning vehicle autonomy," *Preprints.org*, 202605.0908.
29. Anonymous, "Multimodal Sensor Fusion in Autonomous Vehicles: Technologies," *Sensors*, vol. 26, no. 11, p. 3528.
30. Anonymous, "Self-driving vehicles equipped with sensors," *PMC*, PMC7436174.
31. Anonymous, "Testing the fault-tolerance of multi-sensor fusion perception in autonomous driving systems," *ACM*, 2025.
32. Y. X. Li, "Efficient perception methods for autonomous driving via bird's-eye-view representation," *NTU Doctoral Thesis*, 2025.

33. Anonymous, "MTC-BEV: A novel multi-modal 3D object detection framework," *Sensors*, vol. 16, no. 9, p. 493.
34. Anonymous, "Bird's-Eye-View representation serving as a transformative framework," *arXiv preprint arXiv:2410.07268v1*, 2024.
35. Anonymous, "BEV-CLIP: The first multimodal Bird's-Eye View retrieval methodology," *arXiv preprint arXiv:2401.01065v1*, 2024.
36. H. X. Ruan, "BEV-CasDM: a radar BEV enhancement module," *NTU Master's Thesis*, 2025.
37. Anonymous, "FDSNet (Feature Disagreement Score Network)," *PMC*, PMC12717259.
38. Anonymous, "Point-block Carving (PC) for point cloud completion," *TMM*.
39. Anonymous, "Monocular Test-Time Adaptation (MonoTTA)," *arXiv preprint arXiv:2405.19682v1*, 2024.
40. Z. Luo, et al., "PTTR: Point Tracking TRansformer," *TPAMI*, 2024.
41. L. Kong, et al., "DS-Net: Dynamic Shifting Network," *TPAMI*, 2024.
42. Anonymous, "Radar and Camera Fusion for Object Detection and Tracking: A Comprehensive Survey," *ResearchGate*.
43. Z. Luo, et al., "Exploring point-BEV fusion for 3D point cloud object tracking with transformer."
44. Anonymous, "UnsOcc: Multi-modal fusion framework for 3D semantic occupancy prediction," *arXiv preprint arXiv:2606.03581v1*, 2026.
45. Anonymous, "Dual Range-Voxel Representation (DRVR)," *arXiv preprint arXiv:2606.31688v1*, 2026.
46. Anonymous, "SDGOCC: Semantic and Depth-Guided Bird's-Eye View Transformation," *CVPR*, 2025.
47. Anonymous, "Sparse voxel fusion approach for efficient and accurate 3D scene understanding," *LRS4Fusion*, 2025.
48. H. Xu, "3D Occupancy Perception for Autonomous Driving," *GitHub*.
49. Z. Wei, et al., "MS-Occ: Multi-Stage LiDAR-Camera Fusion for 3D Semantic Occupancy Prediction," *arXiv preprint arXiv:2504.15888v1*, 2025.
50. Anonymous, "Motion-aware framework to address key challenges in 4D millimeter-wave radar detection," *PMC*, PMC12900112.
51. Anonymous, "R4Det: 4D radar-camera sensing configuration," *arXiv preprint arXiv:2603.11566v1*, 2026.
52. Anonymous, "RCGDet3D: Ray-centric Gaussian representation," *arXiv preprint arXiv:2605.21112*, 2026.
53. Anonymous, "Doppler-Aware Radar Encoder (DARE)," *Sensors*, vol. 15, no. 6, p. 1284.
54. Anonymous, "3D Object Detection in Adverse Weather Paper List," *GitHub*.
55. Anonymous, "RPGFusion: Radar Prior Guided Fusion," *CVPR*, 2026.
56. Anonymous, "The emergence of multi-modal foundation models," *arXiv preprint arXiv:2512.00021v2*, 2025.
57. DriveX Workshop, "4th edition of the DriveX Workshop," *CVPR*, 2026.
58. Anonymous, "Large Language Models (LLMs) showcasing remarkable proficiency," *EAAI*, 2026.
59. Anonymous, "Multimodal Sensor Fusion in Autonomous Vehicles: Technologies," *Sensors*,

vol. 26, no. 11, p. 3528.

60. Anonymous, "Multi-modal sensor fusion has become a cornerstone," *Sensors*, vol. 25, no. 19, p. 6033.
61. Anonymous, "Recent advances in foundation models (FMs)," *arXiv preprint arXiv:2504.00911v2*, 2025.
62. Anonymous, "Class-Adaptive Cooperative Perception for Multi-Class LiDAR-based 3D Object Detection in V2X Systems," *arXiv preprint arXiv:2604.10305*, 2026.
63. Anonymous, "Class-Adaptive Cooperative Perception," *arXiv preprint arXiv:2604.10305v1*, 2026.
64. Anonymous, "Collaborative perception has attracted growing interest," *arXiv preprint arXiv:2504.12696*, 2025.
65. Anonymous, "UrbanIng-V2X: Large-scale, multi-modal dataset supporting cooperative perception," *arXiv preprint arXiv:2510.23478v1*, 2025.
66. Anonymous, "Vehicle-to-everything cooperative perception," *arXiv preprint arXiv:2310.03525v5*, 2023.
67. Anonymous, "Collaboration-Based Perception Datasets," *arXiv preprint arXiv:2504.12696v2*, 2025.
68. Anonymous, "EPNet++: Cascade Bi-Directional Fusion for Multi-Modal 3D Object Detection," *ResearchGate*.
69. Anonymous, "Multi-modal 3D intelligence," *arXiv preprint arXiv:2310.15676v2*, 2023.
70. Anonymous, "FrustumFusionNets," *arXiv preprint arXiv:2503.13951*, 2025.
71. Anonymous, "Slight bias on calibration can reduce the performance seriously," *arXiv preprint arXiv:2405.16848v1*, 2024.
72. Anonymous, "Feature-level fusion for cross-modal 3D object detection," *OpenReview*.
73. Anonymous, "Semantic-aware Multi-branch Sampling (SMS) module," *arXiv preprint arXiv:2407.05769v3*, 2024.
74. Anonymous, "MIC-BEV: Multi-Infrastructure Camera Bird's-Eye-View Transformer," *ResearchGate*.
75. Anonymous, "Egocentric hand-object interaction (Ego-HOI)," *arXiv preprint arXiv:2403.05856v1*, 2024.
76. Anonymous, "MonoTAKD: Cross-modal distillation for Mono3D," *arXiv preprint arXiv:2404.04910v2*, 2024.
77. Anonymous, "DA-BEV: Domain adaptive camera-only BEV framework," *ECCV*, 2024.
78. Anonymous, "Content-aware multi-modal joint input pruning technique," *arXiv preprint arXiv:2410.07268*, 2024.
79. Anonymous, "Multi-modal predictive interaction," *NeurIPS*, 2022.
80. Z. Wei, et al., "MS-Occ: Multi-Stage LiDAR-Camera Fusion for 3D Semantic Occupancy Prediction," *arXiv*, 2025.
81. Anonymous, "OccFusion, a novel sensor fusion framework," *Semantic Scholar*.
82. Anonymous, "Semantic Fusion: Integrating Multimodal Data," *EmergentMind*.
83. Anonymous, "Geometric Fusion," *HuggingFace*.
84. Z. Wei, et al., "Accurate 3D semantic occupancy perception," *arXiv preprint arXiv:2504.15888v1*, 2025.

85. Z. Wei, et al., "MS-Occ: Multi-Stage LiDAR-Camera Fusion," *HuggingFace*, 2025.
86. A. Mohammed, "FDSNet: dynamic multimodal fusion stage selection for autonomous driving via feature disagreement scoring," *ResearchGate*, 2025.
87. Anonymous, "Optimizing RetinaNet anchors using differential evolution," *ResearchGate*, 2025.
88. D. Wu, et al., "A Survey of Deep Learning Based Radar and Vision Fusion for 3D Object Detection," *Semantic Scholar*.
89. Anonymous, "Advancing low-light object detection with you only look once," *ResearchGate*, 2024.
90. Anonymous, "Decision-Making Policy for Autonomous Vehicles on Highways," *ResearchGate*, 2024.
91. G. Raut and A. Patole, "End-to-End 3D Object Detection Using LiDAR Point Cloud," *Semantic Scholar*, 2023.
92. Anonymous, "UniMT: Unified modality detector with Mamba and Transformer," *arXiv preprint arXiv:2504.13647v2*, 2025.
93. Anonymous, "DAL: Detecting As Labeling," *arXiv preprint arXiv:2311.07152*, 2023.
94. Anonymous, "Lift-Attend-Splat," *CVPRW*, 2024.
95. Anonymous, "AsyncBEV," *OpenReview*.
96. Anonymous, "Instance feature Enhancement Module (IE-Module)," *PMC*, PMC12473358.
97. Z. Yan, et al., "Cross Modal Transformer: Towards Fast and Robust 3D Object Detection," *ICCV*, 2023.
98. X. Chen, et al., "FUTR3D: A Unified Sensor Fusion Framework for 3D Detection," *ResearchGate*.
99. X. Chen, et al., "FUTR3D: A Unified Sensor Fusion Framework for 3D Detection," *arXiv preprint arXiv:2203.10642*, 2022.
100. X. Chen, et al., "FUTR3D: A Unified Sensor Fusion Framework for 3D Detection," *CVPRW*, 2023.
101. X. Chen, et al., "FUTR3D Webpage."
102. X. Chen, et al., "FUTR3D: A Unified Sensor Fusion Framework for 3D Detection," *CVPR Workshops*, 2023.
103. Anonymous, "Sensor Fusion Technique," *arXiv preprint arXiv:2411.10535v1*, 2024.
104. Anonymous, "BEVFusion: A Simple and Robust LiDAR-Camera Fusion Framework," *arXiv preprint arXiv:2205.13790*, 2022.
105. Anonymous, "GaussianFusion," *HuggingFace*.
106. Anonymous, "FLEX: Real-time scheduling framework," *OpenReview*.
107. Anonymous, "MSMDFusion," *CVPR Slides*, 2023.
108. Anonymous, "SAM4D: Promptable multi-modal segmentation," *arXiv preprint arXiv:2506.21547v1*, 2025.
109. Anonymous, "Physical AI," *Preprints.org*, 202604.0549.
110. R. Xu, et al., "V2X-ViT: Vehicle-to-Everything Cooperative Perception with Vision Transformer," *ECCV*, 2022.
111. S. Chen, "CollaborativePerception_paper," *GitHub*.
112. R. Xu, et al., "V2X-ViT: Vehicle-to-Everything Cooperative Perception with Vision

- Transformer," *arXiv preprint arXiv:2203.10638*, 2022.
113. Anonymous, "Vehicle-to-Everything (V2X) cooperative perception," *arXiv preprint arXiv:2509.24927*, 2025.
 114. R. Xu, et al., "V2X-ViT: Vehicle-to-Everything Cooperative Perception with Vision Transformer," *ResearchGate*.
 115. Y. Li, et al., "V2U4Real: A Real-world Large-scale Dataset for Vehicle-to-UAV Cooperative Perception," *CVPR*, 2026.
 116. Anonymous, "HydraCollab: Adaptive collaborative-perception framework," *arXiv preprint arXiv:2607.00191v1*, 2026.
 117. Anonymous, "Vehicle-infrastructure cooperative perception (VICP)," *MDPI*.
 118. Anonymous, "Cooperative Perception (CP)," *arXiv preprint arXiv:2310.03525v5*, 2023.
 119. Anonymous, "DMSTrack: Differentiable multi-sensor Kalman Filter," *arXiv preprint arXiv:2309.14655v2*, 2023.
 120. K. Chitta, et al., "TransFuser: Imitation with Transformer-Based Sensor Fusion for Autonomous Driving," *PAMI*, 2022.
 121. K. Chitta, et al., "TransFuser," *Computer.org*.
 122. Anonymous, "PARA-Drive: A fully parallel end-to-end AV architecture," *CVPR*, 2024.
 123. Anonymous, "DriveTransformer," *arXiv preprint arXiv:2503.07656v2*, 2025.
 124. K. Chitta, et al., "TransFuser Code," *GitHub*.
 125. Anonymous, "TransFuser v6 (TFv6)," *arXiv preprint arXiv:2512.20563v1*, 2025.
 126. S. Vora, et al., "PointPainting: Sequential Fusion for 3D Object Detection," *CVPR*, 2020.
 127. S. Vora, et al., "PointPainting: Sequential Fusion for 3D Object Detection," *OpenAIRE*.
 128. S. Vora, et al., "PointPainting: Sequential Fusion for 3D Object Detection," *CVPR Open Access*.
 129. Anonymous, "PeP module," *arXiv preprint arXiv:2310.07591*, 2023.
 130. S. Vora, et al., "PointPainting," *DBLP*.
 131. "Autoware pointpainting_fusion," *Autoware*.

引用的著作

1. A Review of Multi-Sensor Fusion in Autonomous Driving - MDPI, <https://www.mdpi.com/1424-8220/25/19/6033>
2. Research on Perception Optimization of Multi-Sensor Fusion in Unmanned Vehicle-Road Coordination_Vol. 3 No. 3 (JBDC 2025)_Journal of Big Data and Computing (ISSN, <http://www.stemmpress.com/jbdc/jbdc20253/3107.html>)
3. Multimodal Sensor Fusion in Autonomous Vehicles: Technologies, Architectures, and Open Challenges - MDPI, <https://www.mdpi.com/1424-8220/26/11/3528>
4. Towards Compact Autonomous Driving Perception with Balanced Learning and Multi-sensor Fusion - arXiv, <https://arxiv.org/html/2606.02979v1>
5. Sensor Fusion and Perception for Autonomous Driving: A Critical Review of Modalities, AI Models, Algorithms, and Industry Configurations - Preprints.org, <https://www.preprints.org/manuscript/202605.0908>
6. Vehicle-to-Everything Cooperative Perception for Autonomous Driving - arXiv, <https://arxiv.org/html/2310.03525v5>

7. A Survey of Deep Learning Based Radar and Vision Fusion for 3D Object Detection in Autonomous Driving This work was supported in part by the National Natural Science Foundation of China (No. 61374159), Shaanxi Natural Fund (No. 2018MJ6048), Space Science and Technology Fund, the Foundation of CETC Key Laboratory of Data Link Technology (- arXiv, <https://arxiv.org/html/2406.00714v1>
8. Multi-Sensor Fusion in Autonomous Driving - Atlantis Press, <https://www.atlantis-press.com/article/126016749.pdf>
9. Deep Learning Sensor Fusion for Autonomous Vehicle Perception and Localization: A Review - PMC - NIH, <https://pmc.ncbi.nlm.nih.gov/articles/PMC7436174/>
10. A Survey of the Multi-Sensor Fusion Object Detection Task in Autonomous Driving - MDPI, <https://www.mdpi.com/1424-8220/25/9/2794>
11. Dynamic Feature Fusion for Sparse Radar Detection: Motion-Centric BEV Learning with Adaptive Task Balancing - PMC, <https://pmc.ncbi.nlm.nih.gov/articles/PMC12900112/>
12. Leveraging Adverse Weather for Enhanced LiDAR Spoofing in Autonomous Driving - A*OAR, <https://oar.a-star.edu.sg/storage/2/2dp6qdypj5/mag-submission-camera-ready.pdf>
13. Radar and Camera Fusion for Object Detection and Tracking: A Comprehensive Survey - arXiv, <https://arxiv.org/html/2410.19872v1>
14. A Survey of the Multi-Sensor Fusion Object Detection Task in Autonomous Driving - PMC, <https://pmc.ncbi.nlm.nih.gov/articles/PMC12074113/>
15. A BEV-Fusion Based Framework for Sequential Multi-Modal Beam Prediction in mmWave Systems - arXiv, <https://arxiv.org/html/2604.05668v1>
16. Cross Modal Transformer: Towards Fast and Robust 3D Object Detection - CVF Open Access, https://openaccess.thecvf.com/content/ICCV2023/papers/Yan_Cross_Modal_Transformer_Towards_Fast_and_Robust_3D_Object_Detection_ICCV_2023_paper.pdf
17. [2302.06643] Vision-RADAR fusion for Robotics BEV Detections: A Survey - arXiv, <https://arxiv.org/abs/2302.06643>
18. MS-Occ: Multi-Stage LiDAR-Camera Fusion for 3D Semantic Occupancy Prediction - arXiv, <https://arxiv.org/html/2504.15888v1>
19. (PDF) FDSNet: dynamic multimodal fusion stage selection for autonomous driving via feature disagreement scoring - ResearchGate, https://www.researchgate.net/publication/398873144_FDSNet_dynamic_multimodal_fusion_stage_selection_for_autonomous_driving_via_feature_disagreement_scoring
20. FDSNet: dynamic multimodal fusion stage selection for autonomous driving via feature disagreement scoring - PMC, <https://pmc.ncbi.nlm.nih.gov/articles/PMC12717259/>
21. BEVCar: Camera-Radar Fusion for BEV Map and Object Segmentation - arXiv, <https://arxiv.org/html/2403.11761v2>
22. Learning Content-Aware Multi-Modal Joint Input Pruning via Birds'-Eye-View Representation - arXiv, <https://arxiv.org/html/2410.07268v1>

23. [2205.13790] BEVFusion: A Simple and Robust LiDAR-Camera Fusion Framework - arXiv, <https://arxiv.org/html/2205.13790>
24. CVPR Poster SDGOCC: Semantic and Depth-Guided Bird's-Eye View Transformation for 3D Multimodal Occupancy Prediction, <https://cvpr.thecvf.com/virtual/2025/poster/32524>
25. RaLiBEV: Radar and LiDAR BEV Fusion Learning for Anchor Box Free Object Detection Systems - arXiv, <https://arxiv.org/html/2211.06108v5>
26. RCGDet3D: Rethinking 4D Radar-Camera Fusion-based 3D Object Detection with Enhanced Radar Feature Encoding - arXiv, <https://arxiv.org/pdf/2605.21112>
27. Adaptive Task Batch Scheduling with Elastic Fusion in Multi-Modal Multi-View Machine Perception - OpenReview, <https://openreview.net/pdf?id=SUuH5QcERL>
28. Efficient perception methods for autonomous driving via bird's-eye-view representation - DR-NTU, <https://dr.ntu.edu.sg/entities/publication/22e28fbf-fff7-4e9d-8d34-c9e737a311bd>
29. Learning Content-Aware Multi-Modal Joint Input Pruning via Birds'-Eye-View Representation - arXiv, <https://arxiv.org/pdf/2410.07268>
30. TransFuser: Imitation with Transformer-Based Sensor Fusion for Autonomous Driving - Andreas Geiger, <https://www.cvlbs.net/publications/Chitta2022PAMI.pdf>
31. (PDF) FUTR3D: A Unified Sensor Fusion Framework for 3D Detection - ResearchGate, https://www.researchgate.net/publication/359390091_FUTR3D_A_Unified_Sensor_Fusion_Framework_for_3D_Detection
32. [2203.10642] FUTR3D: A Unified Sensor Fusion Framework for 3D Detection - arXiv, <https://arxiv.org/abs/2203.10642>
33. FUTR3D: A Unified Sensor Fusion Framework for 3D Detection - CVF Open Access, https://openaccess.thecvf.com/content/CVPR2023W/WAD/papers/Chen_FUTR3D_A_Unified_Sensor_Fusion_Framework_for_3D_Detection_CVPRW_2023_paper.pdf
34. FUTR3D: A Unified Sensor Fusion Framework for 3D Detection, <https://tsinghua-mars-lab.github.io/futr3d/>
35. IFE-CMT: Instance-Aware Fine-Grained Feature Enhancement Cross Modal Transformer for 3D Object Detection - PMC, <https://pmc.ncbi.nlm.nih.gov/articles/PMC12473358/>
36. TransFuser: Imitation With Transformer-Based Sensor Fusion for Autonomous Driving - IEEE Computer Society, <https://www.computer.org/csdl/journal/tp/2023/11/09863660/1FXWSbjkLm>
37. LEAD: Minimizing Learner-Expert Asymmetry in End-to-End Driving - arXiv, <https://arxiv.org/html/2512.20563v1>
38. PointPainting: Sequential Fusion for 3D Object Detection - CVF Open Access, https://openaccess.thecvf.com/content_CVPR_2020/papers/Vora_PointPainting_Sequential_Fusion_for_3D_Object_Detection_CVPR_2020_paper.pdf
39. PointPainting: Sequential Fusion for 3D Object Detection - National Open Access Monitor, Ireland, <https://oamonitor.ireland.openaire.eu/rpo/rcsi/search/publication?pid=10.1109%2Fcvpr42600.2020.00466>

40. Sourabh Vora, Alex H. Lang, Bassam Helou, Oscar Beijbom; Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2020, pp. 4604-4612,
https://openaccess.thecvf.com/content_CVPR_2020/html/Vora_PointPainting_Sequential_Fusion_for_3D_Object_Detection_CVPR_2020_paper.html
41. PeP: a Point enhanced Painting method for unified point cloud tasks - arXiv, <https://arxiv.org/pdf/2310.07591>
42. EPNNet++: Cascade Bi-Directional Fusion for Multi-Modal 3D Object Detection, https://www.researchgate.net/publication/366264339_EPNNet_Cascade_Bi-Directional_Fusion_for_Multi-Modal_3D_Object_Detection
43. FrustumFusionNets: A Three-Dimensional Object Detection Network Based on Tractor Road Scene - arXiv, <https://arxiv.org/pdf/2503.13951>
44. Bidirectional Propagation for Cross-Modal 3D Object Detection - OpenReview, https://openreview.net/forum?id=gYs_cRuK7V
45. Boosting 3D Object Detection with Semantic-Aware Multi-Branch Framework - arXiv, <https://arxiv.org/html/2407.05769v3>
46. Multi-modal Sensor Fusion for Auto Driving Perception: A Survey - arXiv, <https://arxiv.org/html/2202.02703v3>
47. Fully Test-Time Adaptation for Monocular 3D Object Detection - arXiv, <https://arxiv.org/html/2405.19682v1>
48. UnsOcc: 3D Semantic Occupancy Prediction in Unstructured Scene via Rendering Fusion, <https://arxiv.org/html/2606.03581v1>
49. Semantic Occupancy Prediction with Dual Range-Voxel Representation - arXiv, <https://arxiv.org/html/2606.31688v1>
50. MS-Occ: Multi-Stage LiDAR-Camera Fusion for 3D Semantic Occupancy Prediction - arXiv, <https://arxiv.org/abs/2504.15888>
51. Exploring Point-BEV Fusion for 3D Point Cloud Object Tracking With Transformer - IEEE Computer Society, <https://www.computer.org/csdl/journal/tp/2024/09/10460168/1V4sjD0wojC>
52. Probabilistic 3D Multi-Object Cooperative Tracking for Autonomous Driving via Differentiable Multi-Sensor Kalman Filter - arXiv, <https://arxiv.org/html/2309.14655v2>
53. ASYNCBEV: CROSS-MODAL FLOW ALIGNMENT IN ASYNCHRONOUS 3D OBJECT DETECTION - OpenReview, <https://openreview.net/pdf/f9c436252a52383576e468a88aa03ea5c4981afb.pdf>
54. MTC-BEV: Semantic-Guided Temporal and Cross-Modal BEV Feature Fusion for 3D Object Detection - MDPI, <https://www.mdpi.com/2032-6653/16/9/493>
55. PIE: Perception and Interaction Enhanced End-to-End Motion Planning for Autonomous Driving - arXiv, <https://arxiv.org/html/2509.18609v1>
56. An Efficient LiDAR-Camera Fusion Network for Multi-Class 3D Dynamic Object Detection and Trajectory Prediction - arXiv, <https://arxiv.org/html/2504.13647v2>
57. RPGFusion: 4D Radar Prior-Guided Multi-Modal Fusion for 3D Detection, https://openaccess.thecvf.com/content/CVPR2026/papers/Qiu_RPGFusion_4D_Radar_Prior-Guided_Multi-Modal_Fusion_for_3D_Detection_CVPR_2026_paper.pdf
58. Robust BEV Perception via Dual 4D Radar-Camera Fusion Under Adverse

- Conditions with Fog-Aware Enhancement - MDPI,
<https://www.mdpi.com/2079-9292/15/6/1284>
59. Class-Adaptive Cooperative Perception for Multi-Class LiDAR-based 3D Object Detection in V2X Systems - arXiv, <https://arxiv.org/html/2604.10305v1>
 60. V2X-ViT: Vehicle-to-Everything Cooperative Perception with Vision Transformer - ECVA,
https://www.ecva.net/papers/eccv_2022/papers_ECCV/papers/136990106.pdf
 61. HydraCollab: Adaptive Collaborative-Perception for Distributed Autonomous Systems - arXiv, <https://arxiv.org/html/2607.00191v1>
 62. Siheng-Chen/CollaborativePerception_paper: A paper list about collaborative perception., https://github.com/Siheng-Chen/CollaborativePerception_paper
 63. (PDF) V2X-ViT: Vehicle-to-Everything Cooperative Perception with Vision Transformer,
https://www.researchgate.net/publication/359389780_V2X-ViT_Vehicle-to-Everything_Cooperative_Perception_with_Vision_Transformer
 64. V2X-ViT: Vehicle-to-Everything Cooperative Perception with Vision Transformer - arXiv, <https://arxiv.org/abs/2203.10638>
 65. [2604.10305] Class-Adaptive Cooperative Perception for Multi-Class LiDAR-based 3D Object Detection in V2X Systems - arXiv,
<https://arxiv.org/abs/2604.10305>
 66. Feature-Level Vehicle-Infrastructure Cooperative Perception with Adaptive Fusion for 3D Object Detection - MDPI, <https://www.mdpi.com/2624-6511/8/5/171>
 67. ylwhxht/Awesome-3D_Object_Detection_in_Adverse_Weather: Paper List of 3D Object Detection in Adverse Weather - GitHub,
https://github.com/ylwhxht/3D_Object_Detection_in_Adverse_Weather_Paper_List
 68. Testing the fault-tolerance of multi-sensor fusion perception in autonomous driving systems,
<https://dr.ntu.edu.sg/entities/publication/d5856a7d-ec76-4086-a4f2-d896d01e63cf>
 69. Foundation Models for Trajectory Planning in Autonomous Driving: A Review of Progress and Open Challenges - arXiv, <https://arxiv.org/html/2512.00021v2>
 70. Foundation models for autonomous driving: A comprehensive survey - ÉTS,
<https://profs.etsmtl.ca/rlangar/Papers/EAAI2026.pdf>
 71. Foundation Models for Autonomous Driving Systems: An Initial Roadmap - arXiv,
<https://arxiv.org/html/2504.00911v2>
 72. SAM4D: Segment Anything in Camera and LiDAR Streams - arXiv,
<https://arxiv.org/html/2506.21547v1>
 73. [PDF] OccFusion: Multi-Sensor Fusion Framework for 3D Semantic Occupancy Prediction,
<https://www.semanticscholar.org/paper/3c0106ce657e71b0803e9521c49713e0a9ec94e4>
 74. DriveTransformer: Unified Transformer for Scalable End-to-End Autonomous Driving - arXiv, <https://arxiv.org/html/2503.07656v2>
 75. Lingdong Kong Doctor of Philosophy PhD Student at National University of Singapore - ResearchGate,

- <https://www.researchgate.net/profile/Lingdong-Kong-3>
76. Foundation Models for Autonomous Driving - DriveX Workshop, <https://drivex-workshop.github.io/cvpr2026/>
 77. Sensor Fusion Market In Autonomous Vehicles Report 2031 - Mordor Intelligence, <https://www.mordorintelligence.com/industry-reports/sensor-fusion-market-in-a-utonomous-vehicle>
 78. A Survey of Multi-sensor Fusion Perception for Embodied AI: Background, Methods, Challenges and Prospects - arXiv, <https://arxiv.org/html/2506.19769v1>
 79. We research 3D Occupancy Perception for Autonomous Driving - GitHub, <https://github.com/HuaiyuanXu/3D-Occupancy-Perception>
 80. CVPR Poster PARA-Drive: Parallelized Architecture for Real-time Autonomous Driving, <https://cvpr.thecvf.com/virtual/2024/poster/29664>
 81. Sourabh Vora - DBLP, <https://dblp.org/pid/203/9954>
 82. 15 Notable AI Datasets for Autonomous Driving in 2024-2025 | BasicAI's Blog, <https://www.basic.ai/blog-post/15-new-autonomous-driving-datasets-in-2024-2025>
 83. Collaborative Perception Datasets for Autonomous Driving: A Review - arXiv, <https://arxiv.org/html/2504.12696v2>
 84. UrbanIng-V2X: A Large-Scale Multi-Vehicle, Multi-Infrastructure Dataset Across Multiple Intersections for Cooperative Perception - arXiv, <https://arxiv.org/html/2510.23478v1>
 85. Collaborative Perception Datasets for Autonomous Driving: A Review - arXiv, <https://arxiv.org/pdf/2504.12696>
 86. A re-calibration method for object detection with multi-modal alignment bias in autonomous driving - arXiv, <https://arxiv.org/html/2405.16848v1>
 87. When Autonomous Vehicle Meets V2X Cooperative Perception: How Far Are We? - arXiv, <https://arxiv.org/pdf/2509.24927?>
 88. DA-BEV: Unsupervised Domain Adaptation for Bird's Eye View Perception - ECVA | European Computer Vision Association, https://www.ecva.net/papers/eccv_2024/papers_ECCV/papers/10792.pdf