

Deep Learning for Financial Risk Modeling under Uncertainty

Abstract

The integration of deep learning (DL) into financial risk modeling has precipitated a paradigm shift in quantitative finance, offering unprecedented capabilities for capturing non-linear dependencies, high-dimensional interactions, and complex temporal dynamics. However, the deployment of standard deterministic neural networks in high-stakes financial environments is fundamentally constrained by their inability to quantify uncertainty, often yielding overconfident predictions in out-of-distribution scenarios such as market crashes or structural regime shifts. This comprehensive survey systematically reviews the state-of-the-art in deep learning methodologies engineered for financial risk modeling under uncertainty. The review examines foundational advancements in Uncertainty Quantification (UQ), including Bayesian Deep Learning and Evidential Deep Learning, detailing their capacity to disentangle aleatoric and epistemic uncertainty. Furthermore, it explores the convergence of deep learning with stochastic calculus and econometric theory, detailing innovations in Deep Backward Stochastic Differential Equations (BSDEs), Physics-Informed Neural Networks (PINNs), and hybrid GARCH-recurrent architectures. The survey addresses the transition from risk-neutral pricing to real-world risk management through Deep Hedging and Distributional Reinforcement Learning. Finally, the role of Generative AI—encompassing Normalizing Flows, Diffusion Models, and Generative Adversarial Networks—in heavy-tailed density estimation and scenario generation is analyzed, alongside the application of Graph Neural Networks for modeling systemic risk and financial contagion. By synthesizing these subfields, this paper identifies critical challenges, including non-stationarity, regime fragility, and calibration inconsistency, outlining the future research trajectories necessary for developing trustworthy, uncertainty-aware financial AI systems.

Introduction

Global financial markets are governed by persistent, multi-faceted uncertainty, characterized by stochastic volatility, heavy-tailed asset returns, and complex cross-asset dependencies¹. Classical financial mathematics and econometrics, built upon frameworks such as the Black-Scholes-Merton option pricing model, Markowitz portfolio optimization, and Generalized Autoregressive Conditional Heteroskedasticity (GARCH) models, rely heavily on rigid parametric assumptions⁴. While these models provide analytical tractability, they frequently fail to capture the full spectrum of market dynamics, particularly during periods of macroeconomic stress, structural breaks, and liquidity crises⁷. The assumption of log-normal price diffusions with constant volatility, for instance, systematically underestimates the probability of extreme tail events, exposing financial institutions to catastrophic risks¹⁰. In response, the financial sector has increasingly adopted deep learning algorithms—ranging

from Long Short-Term Memory (LSTM) networks to advanced Transformer architectures—to extract latent representations from massive, unstructured, and high-dimensional datasets¹². Deep learning has demonstrated superior performance in diverse applications, including algorithmic trading, credit scoring, algorithmic market making, and high-frequency volatility forecasting¹⁴. However, conventional deep neural networks act as deterministic function approximators, producing singular point estimates without a rigorous measure of predictive confidence¹⁶. In financial risk management, where the cost of error is asymmetric and potentially ruinous, the inability to distinguish between high-confidence predictions and arbitrary extrapolations represents a critical vulnerability¹⁹.

Consequently, the intersection of deep learning and Uncertainty Quantification (UQ) has emerged as a paramount research frontier. Quantifying uncertainty enables financial models to communicate the limits of their knowledge, facilitating risk-averse decision-making, selective prediction, and robust capital allocation²¹. This survey provides an exhaustive examination of how deep learning architectures are being explicitly engineered or augmented to model, quantify, and optimize for uncertainty in financial risk contexts. The subsequent sections categorize this rapidly evolving domain into distinct methodological families, exploring their theoretical foundations, mathematical formalisms, and empirical applications.

Methods

Uncertainty Quantification in Deep Learning

The theoretical foundation of uncertainty-aware deep learning relies on the decomposition of predictive uncertainty into two distinct components: aleatoric uncertainty and epistemic uncertainty²⁰. Aleatoric uncertainty, or data uncertainty, arises from the inherent stochasticity and irreducible noise in the observation process, such as bid-ask bounce or market microstructure noise²². It cannot be eliminated by increasing the size of the dataset. Epistemic uncertainty, or model uncertainty, stems from a lack of knowledge about the true underlying data-generating mechanism²⁰. This uncertainty peaks in regions of the feature space where training data is sparse—such as rare market crashes or novel macroeconomic regimes—and can theoretically be reduced by observing additional data¹⁶.

Bayesian Deep Learning and Ensemble Methods

Bayesian Neural Networks (BNNs) offer a principled probabilistic framework for capturing epistemic uncertainty by replacing deterministic network weights with probability

distributions⁷. Given a dataset $\mathcal{D} = \{X, Y\}$, Bayesian inference seeks to compute the

posterior distribution over the network parameters W , expressed via Bayes' theorem as

$$P(W|\mathcal{D}) = \frac{P(\mathcal{D}|W)P(W)}{P(\mathcal{D})}$$
²⁴. Because the marginal likelihood $P(\mathcal{D})$ requires integrating over a high-dimensional parameter space and is thus intractable for deep networks, approximate inference techniques are strictly required¹⁶.

Common approximations in financial modeling include Variational Inference (VI), where the true

posterior is approximated by a tractable parameterized distribution by minimizing the Kullback-Leibler (KL) divergence, and Monte Carlo (MC) Dropout²⁴. MC Dropout, interpreted as a variational approximation, enables scalable uncertainty estimation by stochastically dropping neurons during both training and inference, thereby generating a distribution of predictions whose variance serves as a proxy for epistemic uncertainty²⁴.

Deep Ensembles have also proven highly effective in financial forecasting. By training multiple identical architectures with varying weight initializations and random seeds, the variance across the ensemble provides robust, well-calibrated uncertainty estimates that often outperform explicit Bayesian methods in out-of-distribution (OOD) detection¹⁶. In volatility forecasting, Bayesian LSTMs and Deep Ensembles have demonstrated the ability to deliver conservative, risk-conscious predictive intervals during highly volatile regimes, such as the COVID-19 pandemic, where traditional point-forecasting networks suffered from severe degradation⁷. However, these sampling-based techniques suffer from high computational overhead during inference, limiting their viability for latency-sensitive applications like high-frequency trading²⁹.

Evidential Deep Learning

Evidential Deep Learning (EDL) circumvents the computational bottleneck of Bayesian sampling by formulating uncertainty quantification as a deterministic process requiring only a single forward pass¹⁷. Grounded in Dempster-Shafer Theory (DST) and Subjective Logic (SL), EDL treats the neural network as an evidence collector²⁰.

For classification tasks, such as credit default prediction, standard networks output a softmax probability vector. In contrast, EDL replaces the softmax layer with an activation function (e.g.,

Softplus) that outputs a non-negative evidence vector $\mathbf{e} = [e_1, \dots, e_K]$ for K classes³⁰.

This evidence parameterizes a Dirichlet distribution, $\text{Dir}(\mathbf{p}|\boldsymbol{\alpha})$, where the concentration

parameters are defined as $\boldsymbol{\alpha} = \mathbf{e} + \mathbf{1}$ ³⁰. The predictive probability is derived as the mean of the Dirichlet distribution, while the epistemic uncertainty is inversely proportional to the total

accumulated evidence, mathematically defined as $u = \frac{K}{\sum \alpha_k}$ ²⁰.

Despite its computational elegance, standard EDL exhibits significant vulnerabilities. The original formulation includes a KL-divergence regularization term designed to shrink evidence for incorrect classes to zero. However, this often induces optimization shortcuts where the network blindly inflates feature magnitudes to produce excessively high target evidence, resulting in severe overconfidence³⁰. Recent theoretical advancements have addressed this by reformulating EDL through the lens of variational inference, deriving a rigorous Evidence Lower Bound (ELBO) that structurally penalizes magnitude-induced overconfidence³⁰. Other extensions, such as Information Bottleneck-regularized EDL (IB-EDL), actively suppress spurious information in the evidence generation process, ensuring that only truly predictive features influence the uncertainty estimates¹⁷. In financial applications like fraud detection and algorithmic trading, Evidential Probing Networks (EPNs) and plug-in loss variants have been successfully deployed to efficiently quantify uncertainty without requiring fundamental

changes to underlying pretrained architectures³¹.

UQ Methodology	Core Mechanism	Computational Cost	Primary Financial Applications
Bayesian Neural Networks	Places prior distributions over weights; approximates posterior via VI ²⁴ .	Very High (Training)	Long-term macroeconomic forecasting; uncertainty-aware credit scoring ⁷ .
MC Dropout	Stochastically masks neurons during forward passes to approximate inference ²⁷ .	Medium (Inference)	Realized volatility prediction; robust time-series forecasting ²⁷ .
Deep Ensembles	Aggregates predictions across multiple independently trained networks ²² .	High (Training & Inference)	Out-of-distribution market detection; ensemble deep hedging ²² .
Evidential Deep Learning	Predicts Dirichlet concentration parameters directly in a single forward pass ²⁵ .	Low	High-frequency trading; real-time fraud detection; anomaly detection ²⁵ .

Deep Solvers for Stochastic Differential Equations

Pricing complex financial derivatives and managing high-dimensional risk portfolios inherently requires solving non-linear Partial Differential Equations (PDEs), most notably the Hamilton-Jacobi-Bellman (HJB) equations. Traditional numerical methods, such as finite difference schemes and binomial trees, suffer catastrophically from the "curse of dimensionality," rendering them computationally intractable for portfolios containing multiple correlated underlying assets³⁶.

The Deep BSDE Solver

To overcome this fundamental limitation, deep learning leverages the non-linear Feynman-Kac theorem, which establishes a formal mathematical correspondence between parabolic PDEs and Backward Stochastic Differential Equations (BSDEs)³⁷. The pioneering Deep BSDE method,

introduced by E, Han, and Jentzen, reformulates the PDE as a stochastic optimal control problem³⁷. Given a forward Stochastic Differential Equation (SDE) describing the asset price dynamics X_t , the solution to the PDE $u(t, X_t)$ and its spatial gradient $\nabla_x u(t, X_t)$ correspond precisely to the processes Y_t and Z_t in the BSDE framework³⁷.

The Deep BSDE algorithm parameterizes the unknown gradient function Z_t (representing the optimal delta hedging strategy) at each discrete time step using a sequence of feed-forward neural networks³⁷. The networks take the simulated asset paths and Brownian increments as inputs and are trained to minimize the mean squared error between the forward-simulated terminal value Y_T and the known derivative payoff function $g(X_T)$ ³⁷. By simulating millions of stochastic paths, the model accesses virtually infinite training data, allowing Stochastic Gradient Descent (SGD) to discover the optimal hedging strategy directly³⁷. To accelerate convergence and reduce the loss function in highly complex domains, researchers frequently inject prior knowledge via asymptotic expansion approximations, utilizing the neural network strictly to learn the residual error term⁴⁰.

Physics-Informed Neural Networks (PINNs)

Physics-Informed Neural Networks (PINNs) provide a mesh-free, unsupervised framework for option pricing by embedding the Black-Scholes or Heston PDE directly into the neural network's loss function⁴². The PINN is trained to minimize a composite loss comprising the PDE residual (ensuring adherence to the financial mathematical model), the initial condition loss, and the boundary condition loss⁴⁵.

Because PINNs utilize automatic differentiation to compute the exact spatial and temporal derivatives of the neural network with respect to stock price and time, they do not require computationally expensive finite difference grids⁴⁴. For complex instruments such as American options, PINNs tackle the free-boundary problem via obstacle-style relaxations, retaining the PDE residual exclusively within the continuation region⁴². Recent developments in Uncertainty-Aware PINNs (AT-PINN) utilize anchored-ensemble fine-tuning to provide epistemic uncertainty bands alongside point estimates. This dual capability circumvents the catastrophic error accumulation common in time-marching schemes and aligns the surrogate model strictly with no-arbitrage constraints⁴².

Deep Hedging and Rough Volatility Calibration

The restrictive assumptions of continuous trading, infinite liquidity, and zero transaction costs inherent in classical Black-Scholes replication are heavily violated in real-world markets. Deep Hedging addresses these market frictions by entirely bypassing the risk-neutral valuation framework, operating directly under the physical probability measure⁴⁸.

The Deep Hedging Framework

Introduced as an end-to-end framework, Deep Hedging formulates dynamic risk management as a discrete-time stochastic optimization problem⁵⁰. A neural network, acting as the trading

policy, maps market states (e.g., spot prices, implied volatility surfaces, elapsed time) to optimal hedge ratios⁵². The network is trained using reinforcement learning (RL) techniques, or direct stochastic optimization, to minimize a convex risk measure—such as Conditional Value-at-Risk (CVaR) or Entropic Risk—applied to the terminal Profit & Loss (P&L) of the hedged portfolio³⁵. Crucially, the objective function explicitly incorporates transaction costs, liquidity constraints, and permanent market impact⁴⁸. Empirical studies on S&P 500 options reveal that Deep Hedging agents learn a systematic "delta haircut"—underhedging relative to the Black-Scholes delta in order to minimize trading turnover and transaction costs⁵². However, these learned strategies exhibit "regime fragility"; an agent optimized on a specific historical trajectory or a highly stylized simulation may suffer severe losses in adversarial or previously unseen market regimes⁵². To mitigate this, uncertainty-aware Deep Hedging blends classical delta hedges with neural hedges, utilizing ensemble disagreement to dynamically weight the strategy. When ensemble variance is high (indicating significant epistemic uncertainty), the system heavily weights the classical analytical model, reducing CVaR by up to 250 basis points over optimal theoretical strategies³⁵.

Deep Calibration of Rough Volatility

In classical quantitative modeling, calibrating stochastic volatility models (e.g., the Heston model) to market implied volatility (IV) surfaces is an ill-posed, computationally intensive inverse problem⁵⁷. This computational burden is vastly exacerbated in Rough Volatility models (such as Rough Heston or rBergomi), where the volatility process is driven by fractional

Brownian motion characterized by a Hurst parameter $H < 1/2$ ⁵⁹. Rough models brilliantly capture the explosive short-term skew of IV surfaces and the "Zumbach effect" (the asymmetry between past returns and future volatility), but lack closed-form pricing formulas, requiring exceptionally slow Monte Carlo simulations⁶¹.

Deep Calibration maps the non-linear relationship between model parameters and the resulting IV surface using supervised deep learning⁵⁸. An extensive offline dataset is generated via Monte Carlo simulation, and a neural network is trained to approximate the pricing function. In production, the network performs a virtually instantaneous forward pass, entirely bypassing the simulation bottleneck⁵⁷. For rough volatility, deep neural networks enable the rapid identification of the Hurst parameter and forward variance curve, integrating deep calibration with recurrent Deep Hedging networks to provide model-free replication strategies in non-Markovian environments⁵⁹.

Generative AI and Heavy-Tailed Density Estimation

A fundamental challenge in data-driven financial risk modeling is the extreme scarcity of historical data representing tail events or "flash crashes"³. Statistical models rely on generative architectures to synthesize realistic counterfactual scenarios, augment datasets, and accurately estimate complex, multi-modal probability densities³.

Normalizing Flows and Lévy-Flows

Normalizing Flows (NFs) constitute a class of generative models that provide exact, tractable

density estimation by learning a sequence of invertible and differentiable transformations (diffeomorphisms) mapping a simple base distribution to a complex target distribution¹⁰. Given a base distribution $p_Z(z)$ and an invertible mapping f_θ , the target likelihood is precisely evaluated via the change-of-variables formula:

$$\log p_X(x) = \log p_Z(f_\theta^{-1}(x)) + \log \left| \det \frac{\partial f_\theta^{-1}}{\partial x} \right|$$

¹¹.

However, standard NFs rely on Gaussian base distributions. In finance, where asset returns exhibit extreme leptokurtosis (fat tails), a Gaussian base systematically underestimates extreme events, forcing the flow transformation to aggressively stretch the tails—a process that is numerically unstable and statistically inefficient¹⁰. To rectify this, researchers introduced Lévy-Flows, which substitute the Gaussian base with heavy-tailed distributions derived from Lévy processes, specifically the Variance Gamma (VG) and Normal-Inverse Gaussian (NIG) distributions¹⁰. The VG distribution captures semi-heavy tails via Gamma-subordinated Brownian motion, while the NIG distribution offers precise skewness and kurtosis control via inverse Gaussian subordination¹⁰.

Theoretical analyses utilizing extreme value theory demonstrate that Lévy-Flows preserve the tail index of the base distribution under asymptotically linear transformations⁶⁹. Empirically, VG-based flows have been shown to reduce test negative log-likelihood by up to 69% relative to Gaussian flows on S&P 500 returns, achieving exact 95% VaR calibration and vastly superior Expected Shortfall (ES) estimation⁶⁹. Furthermore, Normalising Flow GARCH (NF-GARCH) models combine parametric volatility recursions with nonparametric flow-based residuals, resolving structural non-stationarity while averting overfitting via likelihood regularization⁷¹.

GANs and Diffusion Models

Generative Adversarial Networks (GANs) and Denoising Diffusion Probabilistic Models (DDPMs) are increasingly utilized to simulate synthetic multivariate financial time series⁶⁶. Frameworks like QuantGAN leverage Temporal Convolutional Networks (TCNs) within the generator and discriminator to faithfully reproduce volatility clustering and long-memory effects⁶⁵. Recent innovations such as CoMeTS-GAN (Correlated Multivariate Time Series GAN) model the joint dynamics of mid-prices and trading volumes across correlated assets⁶⁶. CoMeTS-GAN employs a novel cross-correlation distance metric to evaluate the realism of generated correlation structures and utilizes a Critic-Guided diffusion process to steer the synthesis of high-fidelity synthetic market environments for institutional stress testing⁶⁶.

Diffusion models, renowned for superior mode coverage and training stability, have also been adapted for financial time series³. By transforming financial sequences (prices, spreads, volumes) into full-color spectrogram images via continuous wavelet transforms, DDPMs learn intricate spectral characteristics before utilizing inverse wavelet transforms to recover the synthetic series³. Advanced architectures incorporate Geometric Brownian Motion (GBM)

directly into the forward noising process, injecting noise proportionally to asset prices to mathematically mirror the heteroskedasticity innate to financial markets⁷⁴.

Generative Methodology	Core Mechanism	Financial Application & Advantages
Lévy-Flows	Invertible transformations on Variance Gamma / NIG bases ¹⁰ .	Exact log-likelihood evaluation; superior heavy-tail VaR and ES estimation ⁶⁹ .
Quant/Time-GANs	Adversarial training over temporal sequences via TCNs/LSTMs ⁶⁵ .	Replicates volatility clustering and dependencies; allows fast scenario generation ⁶⁵ .
Wavelet-DDPMs	Image-based diffusion applied to wavelet spectrograms of returns ³ .	Highly stable training; accurately reproduces multi-scale stylized facts ³ .
GBM-Diffusion	Forward noising process scaled by Geometric Brownian Motion ⁷⁴ .	Aligns with stochastic calculus; natively handles financial heteroskedasticity ⁷⁴ .

Hybrid Econometric and Temporal Fusion Architectures

Financial time series are characterized by specific, pervasive stylized facts: volatility clustering, long memory, leverage effects (asymmetric responses to positive versus negative shocks), and mean-reverting behavior³. While standard recurrent neural networks (RNNs) and LSTMs excel at capturing arbitrary non-linear temporal dependencies, they do not inherently respect these econometric properties, often resulting in uninterpretable outputs⁵.

GARCH-Informed Neural Networks

To bridge the gap between statistical rigor and deep learning flexibility, researchers have developed hybrid architectures that integrate Generalized Autoregressive Conditional Heteroskedasticity (GARCH) dynamics directly into neural network cells⁵. The classical GARCH model explicitly formulates the conditional variance as a function of past squared residuals and past variances, algorithmically enforcing volatility clustering⁷⁵.

In architectures such as the GARCH-LSTM, the recurrent cell is modified so that econometric volatility formulations dictate the hidden state transitions⁵. Instead of static concatenation of GARCH outputs to neural network inputs, the internal gating mechanisms of the LSTM are modulated by a variance process⁷⁵. For example, the HAR-LSTM-GARCH model applies a

Heterogeneous Autoregressive (HAR) component to capture multi-scale, long-memory persistence in realized volatility, utilizes an LSTM to learn non-linear residual patterns, and applies a GARCH framework to model the time-varying volatility of the residuals themselves⁶. These hybrid models significantly outperform standalone deep learning models and classical GARCH variants in predicting crude oil futures volatility and Value-at-Risk (VaR), offering a unified framework that is accurate, economically interpretable, and mathematically robust⁶.

Temporal Fusion Transformers (TFT)

The Temporal Fusion Transformer (TFT) has emerged as a state-of-the-art attention-based architecture explicitly designed for multi-horizon financial forecasting under uncertainty⁷⁷. Unlike standard Transformers utilized in natural language processing, the TFT integrates several domain-specific components tailored for heterogeneous time series:

1. **Variable Selection Networks (VSNs):** These networks dynamically learn the importance of different input features at each time step, providing critical interpretability for financial analysts seeking to understand the drivers of market risk⁷⁷.
2. **Gated Residual Networks (GRNs):** GRNs allow the model to selectively route information, enabling the network to learn complex spatial representations while explicitly filtering out pervasive market noise⁷⁷.
3. **Multi-Head Attention:** This mechanism captures long-term dependencies and identifies complex attention patterns across historical windows⁷⁷.

The TFT inherently outputs probabilistic forecasts by utilizing quantile regression loss, explicitly bounding future uncertainty⁸⁰. Extensions such as the Hierarchical Temporal Fusion (HTF) model introduce coherence-aware loss functions and structured embeddings to ensure that volatility or demand forecasts at lower aggregation levels perfectly reconcile with macro-level forecasts, a strict mathematical requirement in institutional portfolio risk aggregation⁸⁰.

Graph Neural Networks for Systemic Risk and Credit Contagion

Financial networks are deeply interconnected; localized distress in a single institution can rapidly cascade through counterparty exposures, leading to systemic collapse⁸. Traditional systemic risk metrics evaluate institutions in isolation, ignoring topological dependencies and the non-linear amplification of shocks⁸. Graph Neural Networks (GNNs) overcome this limitation by performing message passing across financial networks, encoding both entity-level features and structural connectivity into dense, actionable representations⁸¹.

In interbank contagion modeling, GNNs extend classical Eisenberg-Noe clearing models by predicting the minimal bailout capital required to secure the network prior to systemic default⁸³. Dynamic spatio-temporal architectures, such as GAT-GRU (Graph Attention Networks coupled with Gated Recurrent Units), evaluate continuously evolving liability networks⁸¹. The GAT mechanism dynamically assigns attention weights to contagion sources (e.g., identifying heavily leveraged regional banks prior to insolvency), while the GRU memorizes the trajectory of structural deterioration⁸. Recent empirical applications on US banking data demonstrated that GNNs successfully identified accumulating vulnerabilities in institutions like Silicon Valley Bank several quarters prior to failure, during periods when

traditional regulatory stress tests indicated sufficient capitalization⁸.

Furthermore, GNNs are increasingly applied to massive-scale heterogeneous graphs for consumer credit risk, linking tens of millions of nodes representing borrower attributes, installment payments, and point-of-sale behaviors⁸⁶. While standalone GNNs offer marginal lift over Gradient Boosted Trees (GBMs) in tabular credit scoring, hybrid models that ingest GNN-derived topological embeddings as auxiliary features achieve state-of-the-art Area Under the Precision-Recall Curve (PR-AUC) metrics⁸⁶. Multi-market studies further utilize GNNs to model cross-market spillovers, highlighting that cryptocurrency network integration acts as a significant amplifier of systemic volatility during macroeconomic stress regimes, confirming the necessity of time-varying network analysis⁸².

Distributional and Risk-Sensitive Reinforcement Learning

Standard Reinforcement Learning (RL) optimizes the expected cumulative reward, an objective fundamentally unsuited for risk-sensitive financial applications⁹⁰. In financial portfolio management and automated hedging, relying solely on expected returns ignores the variance, skewness, and severe tail risks of the portfolio⁹¹. Distributional Reinforcement Learning (DRL) reformulates the traditional Bellman equation to learn the entire probability distribution of future returns rather than a simple scalar expectation⁹².

In risk-sensitive portfolio optimization, Distributional RL utilizes Quantile Regression or Radial Basis Function (RBF) feature expansions to map the value distribution⁹². This enables the RL agent to explicitly optimize for coherent risk measures, such as Spectral Risk Measures (SRM) or Conditional Value-at-Risk (CVaR)⁹¹. By integrating maximum drawdown constraints and CVaR penalties directly into the proximal policy objective, DRL agents learn dynamic asset allocation strategies that aggressively reduce tail exposure during turbulent regimes while maintaining competitive Sharpe ratios during bull markets⁹¹. To facilitate the learning of highly generalizable allocation strategies, circular block bootstrap methods and simulated synthetic data environments are frequently employed during training, vastly improving the agent's capacity to extrapolate robust policies into out-of-distribution market conditions⁹⁵.

Challenges and Prospects

Despite rapid theoretical advancements, the deployment of uncertainty-aware deep learning in financial markets faces substantial operational, computational, and mathematical hurdles¹³.

1. **Non-Stationarity and Regime Shifts:** Financial data is profoundly non-stationary. Models trained during prolonged periods of low volatility often fail catastrophically during black-swan events¹⁸. Ensuring that uncertainty estimates (both epistemic and aleatoric) dynamically scale in response to structural market breaks remains a prevailing challenge¹⁸. The "regime fragility" of deep hedging agents highlights that outperformance in one economic cycle does not guarantee robustness in another⁵².
2. **Calibration Inconsistency:** Recent extensive surveys indicate severe discrepancies across standard UQ evaluation metrics (e.g., Expected Calibration Error vs. Brier Score). Metrics frequently disagree on the efficacy of recalibration techniques, permitting the

cherry-picking of empirical results and obscuring the true reliability of the model¹⁹. Developing unified, mathematically rigorous calibration standards is paramount for institutional trust.

3. **Data Sparsity and Simulation Bias:** Advanced methods like Deep Hedging, Deep Calibration, and Distributional RL rely heavily on synthetic data generated via Monte Carlo simulators. If the underlying stochastic simulator is misspecified, the neural network optimizes a flawed representation of reality—a vulnerability vastly exacerbated by the lack of historical high-frequency data for extreme tail events³⁵.
4. **Interpretability and Regulatory Compliance:** Financial regulators mandate high transparency (e.g., the SR 11-7 model risk management guidelines in the US). While models like the Temporal Fusion Transformer and PINNs offer built-in interpretability via attention weights and physical equation residuals, massive multi-layered architectures generally remain "black boxes." Bridging the gap between empirical predictive performance and auditable, Explainable AI (XAI) is critical for widespread institutional adoption⁸.

Conclusion

The application of deep learning to financial risk modeling has evolved from naive pattern recognition and deterministic forecasting to highly sophisticated, uncertainty-aware computational paradigms. By merging deep learning with stochastic calculus (Deep BSDEs), economic structural constraints (PINNs and GARCH-LSTMs), and rigorous probabilistic reasoning (Evidential and Bayesian DL), modern frameworks can effectively navigate the frictions, non-linearities, and heavy-tailed realities of global financial markets. Advancements in Generative AI and Graph Neural Networks further enable the modeling of complex systemic topologies, cross-market contagion, and counterfactual risk scenarios. However, realizing the full potential of these architectures requires overcoming persistent, structural challenges related to model fragility, metric consistency, and interpretability. As research shifts toward robust out-of-distribution generalization, causal inference, and distributional optimization, uncertainty-aware deep learning stands poised to fundamentally redefine quantitative risk management, derivative pricing, and regulatory surveillance in the financial sector.

References

[1] Deep Learning in Finance: A Survey of Applications and Techniques. [2] Prediction of financial risk behavior based on deep learning and big data algorithms. arXiv:2410.19394. [3] A Survey on Uncertainty Quantification Methods for Deep Learning. arXiv:2302.13425. [5] A Structured Review of Literature on Uncertainty in Machine Learning & Deep Learning. arXiv:2406.00332. [8] Comparative analysis of Bayesian LSTM and GPR for financial risk. [10] Bayesian neural networks for multi-step-ahead stock price forecasting before and during COVID-19. [11] Forecasting financial volatility under structural breaks. [12] What Uncertainties Do We Need in Bayesian Deep Learning for Computer Vision. [16] Option pricing using Deep BSDE solvers. [17] Deep learning algorithms for high-dimensional PDEs. [18] Deep BSDE method for high-dimensional nonlinear PDEs. arXiv:2505.17032. [19] CoMeTS-GAN: Correlated

Multivariate Time Series GAN. arXiv:2605.27113. [20] Generation of synthetic financial time series by diffusion models. arXiv:2410.18897. [21] Denoising diffusion probabilistic models for synthetic financial time series. [23] A diffusion-based generative model for financial time series via geometric Brownian motion. arXiv:2507.19003. [24] TimeGAN and QuantGAN for synthetic financial data. [28] MiT-WGAN: A hybrid generative adversarial framework for financial time-series. [31] Deep Hedging. arXiv:1802.03042. [32] Deep Hedging: Learning Risk-Neutral Implied Volatility Dynamics. [34] Reinforcement learning algorithms to learn adaptive hedging strategies. [36] Deep Hedging to Manage Tail Risk. arXiv:2506.22611. [38] Risk-aware portfolio optimization based on conditional value-at-risk. [39] Beyond CVaR: Leveraging Static Spectral Risk Measures for Enhanced Decision-Making in Distributional Reinforcement Learning. arXiv:2501.02087. [41] Distributional Reinforcement Learning for Option Pricing. arXiv:2507.12657. [42] Deep reinforcement learning framework for dynamic portfolio optimization. arXiv:2511.11481. [43] Graph Neural Networks for Systemic Financial Risk Forecasting. [44] Computing systemic risk measures with graph neural networks. arXiv:2410.07222. [46] Explainable Graph Neural Networks for Interbank Contagion Surveillance. arXiv:2604.14232. [49] Lévy-Flow Models: Heavy-Tail-Aware Normalizing Flows for Financial Risk Management. arXiv:2604.00195. [50] Variance Gamma and Normal-Inverse Gaussian distribution for Lévy-Flows. [51] Lévy-Flow Models for Financial Risk Management. [52] Normalising Flow GARCH (NF-GARCH). [53] Theoretical analysis of tail-shape preservation in Normalizing Flows. [54] Normalizing Flows Policy for visuomotor and sequential learning. [56] Evidential deep learning for sampling-free uncertainty estimation. [57] Evidential Probing Network (EPN) for Graph Neural Networks. arXiv:2503.08097. [58] Evidential deep learning (EDL) for deep neural network models. arXiv:2501.05656. [59] Regularizing EDL by incorporating an information bottleneck (IB-EDL). arXiv:2502.06351. [61] Temporal Fusion Transformers for time series forecasting. [67] Neural Stochastic Differential Equations for option hedging. [74] Risk-sensitive reinforcement learning framework for risk-aware portfolio optimization. [75] Future-looking reward function and circular block bootstrap method in DRL. arXiv:2502.02619. [76] Sharpe Ratio-Based Objective Function in Deep Reinforcement Learning. [78] Multiple risk aversion actor-critic algorithms. [79] An uncertainty-aware physics-informed neural network solution for the Black-Scholes equation. arXiv:2511.05519. [80] Physics-Informed Neural Networks for Option Pricing. [81] Coupled physics-informed neural networks for American options. arXiv:2605.06688. [82] Physics-Informed Neural Networks for Black-Scholes equation. arXiv:2312.06711. [83] Automatic differentiation in PINNs for PDE residual. [84] Uncertainty-Aware PINNs (AT-PINN) for free-boundary problems. [86] Asymptotic expansion as prior knowledge in the deep BSDE solver. [88] General convergence rate of the deep BSDE loss functional. arXiv:2508.14215. [89] Acceleration technique for deep learning methods for high-dimensional BSDEs. arXiv:1710.07030. [90] Deep BSDE solver's effectiveness in complex financial settings. [91] Modeling Cross-Market Contagion Between Banking Systems and Cryptocurrency Markets. [92] Heterogeneous graph neural networks for credit default prediction. arXiv:2601.14633. [95] Massive-scale heterogeneous graph for credit risk. [96] AI4Risk interbank dataset and GNN-based credit models. [98] Generative AI and mechanical nondeterminism in financial AI. arXiv:2605.23955. [101] Deep learning methods applied to credit risk prediction. [103] Uncertainty Quantification with Machine Learning and Reinforcement

Learning for Business Decision-Making. [104] Probabilistic models for time series forecasting in finance. [105] Inconsistencies in UQ metrics and recalibration evaluation. arXiv:2508.17761. [106] Critical survey on the consistency of uncertainty quantification used in deep learning. [107] Monte Carlo Dropout and Deep Ensembles for DNN predictions. arXiv:2503.17385. [108] Advances in Uncertainty Quantification Techniques. arXiv:2501.03282. [110] Foundation models for tabular data in credit risk modeling. arXiv:2605.18147. [111] Sequential Deep Learning for Credit Risk Monitoring with Tabular Financial Data. arXiv:2012.15330. [112] FinLangNet: Deep learning risk prediction framework. arXiv:2404.13004. [113] What Does Deep Hedging Actually Learn? Delta Corrections, Regime Fragility. arXiv:2605.21696. [114] Ensemble's disagreement in deep hedging and CVaR optimization. arXiv:2603.10137. [115] Regime Fragility and Symbolic Distillation in Deep Hedging. [116] Applications of deep learning in hedging and calibration techniques. [117] Deep Hedging and Deep Calibration. [118] Deep Hedging with Reinforcement Learning: A Practical Framework. arXiv:2512.12420. [119] Hierarchical Temporal Fusion (HTF). arXiv:2606.28553. [120] Quantum Temporal Fusion Transformer. arXiv:2508.04048. [121] Variable Selection Networks and Gated Residual Networks in TFT. [123] Temporal Fusion Transformers for Interpretable Multi-horizon Time Series Forecasting. [125] Normal-Inverse Gaussian (NIG) distributions for Normalizing Flows. [126] Extreme value theory in tail-shape preservation. [127] Lévy-Flow Models: Heavy-Tail-Aware Normalizing Flows. [130] Density estimation with Normalizing flows. arXiv:2507.22429. [131] Evidential Deep Learning as parameterized generalized Bayesian inference. arXiv:2605.25599. [132] Evidence Lower Bound (ELBO) in Evidential Deep Learning. arXiv:2605.26477. [133] Aleatoric and Epistemic uncertainty in Deep Learning. [134] Plug-in Losses for Evidential Deep Learning. arXiv:2605.22746. [135] Re-EDL: Simplified Evidential Deep Learning. arXiv:2410.00393. [136] Evidential Transformation Network (ETN). [137] Deep Hedging under Rough Volatility. arXiv:2102.01962. [139] rBergomi model and Hurst parameter analysis in deep hedging. [141] Pricing and hedging under rough volatility models. arXiv:2105.04073. [143] Hybrid deep learning framework for Heston calibration. [144] Deep Calibration of stochastic volatility models. arXiv:2510.24074. [145] Quadratic rough Heston model calibration. arXiv:2107.01611. [146] Multi-factor approximation for rough volatility paradigms. [148] Deep calibration of rough stochastic volatility models. [149] Equivalence Relation Between GARCH and Neural Networks. arXiv:2402.06642. [150] HAR-LSTM-GARCH model for realized volatility forecasting. [151] Embedding GARCH dynamics within recurrent neural networks. arXiv:2504.09380. [154] GARCH-Informed Neural Networks (GINN). arXiv:2410.00288. [155] Predictive models in credit risk management. [161] Dynamic graph neural network (GNN) framework for systemic risk. [163] Systemic risk measures for stochastic financial networks. [164] Explainable Graph Neural Networks for Interbank Contagion Surveillance. [165] Market clearing algorithms in systemic risk networks.

Works cited

1. Analysis of Financial Risk Behavior Prediction Using Deep Learning and Big Data Algorithms - arXiv, <https://arxiv.org/pdf/2410.19394>
2. Forecasting Financial Volatility Under Structural Breaks: A Comparative Study of GARCH Models and Deep Learning Techniques - MDPI,

- <https://www.mdpi.com/1911-8074/18/9/494>
3. Generation of synthetic financial time series by diffusion models - arXiv, <https://arxiv.org/pdf/2410.18897>
 4. NEURAL STOCHASTIC DIFFERENTIAL EQUATIONS FOR MODEL-FREE OPTION HEDGING: CONVERGENCE, CALIBRATION, AND RISK BOUNDS - OpenReview, <https://openreview.net/pdf?id=faCnKp5qgQ>
 5. From GARCH to Neural Network for Volatility Forecast - arXiv, <https://arxiv.org/html/2402.06642v1>
 6. A Hybrid HAR-LSTM-GARCH Model for Forecasting Volatility in Energy Markets - MDPI, <https://www.mdpi.com/1911-8074/19/1/77>
 7. Bayesian neural networks for stock price forecasting before and during COVID-19 pandemic, <https://pmc.ncbi.nlm.nih.gov/articles/PMC8248663/>
 8. Explainable Graph Neural Networks for Interbank Contagion Surveillance: A Regulatory-Aligned Framework for the U.S. Banking Sector - arXiv, <https://arxiv.org/pdf/2604.14232>
 9. Sequential Deep Learning for Credit Risk Monitoring with Tabular Financial Data - arXiv, <https://arxiv.org/abs/2012.15330>
 10. Heavy-Tail-Aware Normalizing Flows for Financial Risk Management - arXiv, <https://arxiv.org/pdf/2604.00195>
 11. Lévy-Flow Models: Heavy-Tail-Aware Normalizing Flows for Financial Risk Management, <https://arxiv.org/html/2604.00195v1>
 12. Deep Learning in Finance: A Survey of Applications and Techniques - ResearchGate, https://www.researchgate.net/publication/385337069_Deep_Learning_in_Finance_A_Survey_of_Applications_and_Techniques
 13. From Accuracy to Auditability: A Survey of Determinism in Financial AI Systems - arXiv, <https://arxiv.org/html/2605.23955v1>
 14. Deep Learning for Credit Risk Prediction: A Survey of Methods, Applications, and Challenges - MDPI, <https://www.mdpi.com/2078-2489/17/4/395>
 15. FinLangNet: A Novel Deep Learning Framework for Credit Risk Prediction Using Linguistic Analogy in Financial Data - arXiv, <https://arxiv.org/html/2404.13004v2>
 16. [2302.13425] A Survey on Uncertainty Quantification Methods for Deep Learning - arXiv, <https://arxiv.org/abs/2302.13425>
 17. Calibrating LLMs with Information-Theoretic Evidential Deep Learning - arXiv, <https://arxiv.org/html/2502.06351v2>
 18. Uncertainty Quantification with Machine Learning and Reinforcement Learning for Business Decision-Making - ORBi, <https://orbi.uliege.be/bitstream/2268/343707/1/DoctoralThesis-SINGH-Akash.pdf>
 19. Uncertainty quantification for deep learning | Environmental Data Science | Cambridge Core, <https://www.cambridge.org/core/journals/environmental-data-science/article/uncertainty-quantification-for-deep-learning/84DE5004FA69851F6770B448B8FDDACQ>
 20. A Comprehensive Survey on Evidential Deep Learning and its Applications - IEEE Computer Society,

- <https://www.computer.org/csdl/journal/tp/2026/03/11217233/2b4UDI712yk>
21. Uncertainty Quantification | IBM,
<https://www.ibm.com/think/topics/uncertainty-quantification>
 22. From Aleatoric to Epistemic: Exploring Uncertainty Quantification Techniques in Artificial Intelligence - arXiv, <https://arxiv.org/html/2501.03282v1>
 23. [2406.00332] A Structured Review of Literature on Uncertainty in Machine Learning & Deep Learning - arXiv, <https://arxiv.org/abs/2406.00332>
 24. What Uncertainties Do We Need in Bayesian Deep Learning for Computer Vision? - NIPS,
<https://papers.neurips.cc/paper/7141-what-uncertainties-do-we-need-in-bayesian-deep-learning-for-computer-vision.pdf>
 25. Generalized Evidential Deep Learning: From a Bayesian Perspective - arXiv,
<https://arxiv.org/html/2605.25599v1>
 26. Evidential Deep Learning for Guided Molecular Property Prediction and Discovery - PMC, <https://pmc.ncbi.nlm.nih.gov/articles/PMC8393200/>
 27. Uncertainty Quantification for Data-Driven Machine Learning Models in Nuclear Engineering Applications: Where We Are and What Do We Need? - arXiv,
<https://arxiv.org/html/2503.17385v1>
 28. Bayesian Deep Learning and Probabilistic Forecasting of Stock Prices - MDPI,
<https://www.mdpi.com/1999-4893/19/5/391>
 29. Evidential Deep Learning for Uncertainty Quantification and Out-of-Distribution Detection in Jet Identification using Deep Neural Networks - arXiv,
<https://arxiv.org/html/2501.05656v1>
 30. Variational Inference for Evidential Deep Learning - arXiv,
<https://arxiv.org/html/2605.26477v1>
 31. Evidential Transformation Network: Turning Pretrained Models into Evidential Models for Post-hoc Uncertainty Estimation - CVF Open Access,
https://openaccess.thecvf.com/content/CVPR2026/papers/Chun_Evidential_Transformation_Network_Turning_Pretrained_Models_into_Evidential_Models_for_CVPR_2026_paper.pdf
 32. [2410.00393] Revisiting Essential and Nonessential Settings of Evidential Deep Learning - arXiv, <https://arxiv.org/abs/2410.00393>
 33. [2503.08097] Evidential Uncertainty Probes for Graph Neural Networks - arXiv,
<https://arxiv.org/abs/2503.08097>
 34. [2605.22746] Plug-in Losses for Evidential Deep Learning: A Simplified Framework for Uncertainty Estimation that Includes the Softmax Classifier - arXiv,
<https://arxiv.org/abs/2605.22746>
 35. Uncertainty-Aware Deep Hedging - arXiv, <https://arxiv.org/pdf/2603.10137>
 36. Numerical Investigation of the Deep BSDE Solver for Pricing European Options by Carl Genlöv Brouwer - DiVA portal,
<https://www.diva-portal.org/smash/get/diva2:1979264/FULLTEXT01.pdf>
 37. A brief review of the Deep BSDE method for solving high-dimensional partial differential equations - arXiv, <https://arxiv.org/pdf/2505.17032>
 38. A brief review of the Deep BSDE method for solving high-dimensional partial differential equations - arXiv, <https://arxiv.org/html/2505.17032v1>

39. Deep BSDE Solver on Bounded Domains Part I: General Loss Rate - arXiv, <https://arxiv.org/pdf/2508.14215>
40. Asymptotic Expansion as Prior Knowledge in Deep Learning Method for high dimensional BSDEs - arXiv, <https://arxiv.org/pdf/1710.07030>
41. Asymptotic Expansion as Prior Knowledge in Deep Learning Method for high dimensional BSDEs, <https://www.carf.e.u-tokyo.ac.jp/pdf/workingpaper/fseries/F423.pdf>
42. An uncertainty-aware physics-informed neural network solution for the Black-Scholes equation: a novel framework for option pricing - IDEAS/RePEc, <https://ideas.repec.org/p/arx/papers/2511.05519.html>
43. Physics Informed Neural Network for Option Pricing - arXiv, <https://arxiv.org/html/2312.06711v1>
44. An Uncertainty-Aware Physics-Informed Neural Network Solution for the Black-Scholes Equation: A Novel Framework for Option Pricing - arXiv, <https://arxiv.org/html/2511.05519v1>
45. matlab-deep-learning/PINNsOptionPricing: Apply Physics-informed neural networks in solving Black-Scholes Equations - GitHub, <https://github.com/matlab-deep-learning/PINNsOptionPricing>
46. Physics-Informed Neural Networks (PINNs) for Option Pricing - MATLAB Central Blogs, <https://blogs.mathworks.com/finance/2025/01/07/physics-informed-neural-networks-pinns-for-option-pricing/>
47. American Options Pricing under Heston Model via Curriculum Learning in Coupled PINNs, <https://arxiv.org/html/2605.06688v1>
48. DEEP LEARNING IN FINANCE: A REVIEW OF DEEP HEDGING AND DEEP CALIBRATION TECHNIQUES | International Journal of Theoretical and Applied Finance - World Scientific Publishing, <https://www.worldscientific.com/doi/10.1142/S021902492530001X>
49. A Review of New Developments in Finance with Deep Learning: Deep Hedging and Deep Calibration, <https://www.imes.boj.or.jp/research/papers/english/24-E-02.pdf>
50. [1802.03042] Deep Hedging - arXiv, <https://arxiv.org/abs/1802.03042>
51. Deep Hedging - IDEAS/RePEc, <https://ideas.repec.org/p/arx/papers/1802.03042.html>
52. What Does Deep Hedging Actually Learn? Delta Corrections, Regime Fragility, and Symbolic Distillation - arXiv, <https://arxiv.org/pdf/2605.21696>
53. Deep Hedging with Reinforcement Learning: A Practical Framework for Option Risk Management - arXiv, <https://arxiv.org/pdf/2512.12420>
54. [2506.22611] Deep Hedging to Manage Tail Risk - arXiv, <https://arxiv.org/abs/2506.22611>
55. Deep Hedging Under Market Frictions: A Comparison of DRL Models for Options Hedging with Impact and Transaction Costs - MDPI, <https://www.mdpi.com/1911-8074/18/9/497>
56. [2605.21696] What Does Deep Hedging Actually Learn? Delta Corrections, Regime Fragility, and Symbolic Distillation - arXiv, <https://arxiv.org/abs/2605.21696>

57. Deep Learning-Enhanced Calibration of the Heston Model: A Unified Framework - MDPI, <https://www.mdpi.com/2227-7390/14/13/2322>
58. Deep Learning-Enhanced Calibration of the Heston Model: A Unified Framework - arXiv, <https://arxiv.org/html/2510.24074v1>
59. Deep Hedging under Rough Volatility arXiv:2102.01962v1 [q-fin.CP] 3 Feb 2021, <https://arxiv.org/pdf/2102.01962>
60. Deep Hedging under Rough Volatility - MDPI, <https://www.mdpi.com/2227-9091/9/7/138>
61. Hedging under rough volatility arXiv:2105.04073v1 [q-fin.MF] 10 May 2021, <https://arxiv.org/pdf/2105.04073>
62. Deep calibration of the quadratic rough Heston model - IDEAS/RePEc, <https://ideas.repec.org/p/arx/papers/2107.01611.html>
63. [2107.01611] Deep calibration of the quadratic rough Heston model - arXiv, <https://arxiv.org/abs/2107.01611>
64. Deep calibration of rough stochastic volatility models - Semantic Scholar, <https://www.semanticscholar.org/paper/Deep-calibration-of-rough-stochastic-volatility-Bayer-Stemper/1063ed4f49eacf7a82d4b0ef5639c38531df1bd1>
65. Generation of synthetic financial time series by diffusion models - arXiv, <https://arxiv.org/html/2410.18897v1>
66. High-Quality Synthetic Financial Time-Series using a GAN-Diffusion Framework - arXiv, <https://arxiv.org/html/2605.27113v1>
67. Normalizing Flows are Capable Visuomotor Policy Learning Models - arXiv, <https://arxiv.org/html/2509.21073v1>
68. [2507.22429] Comparing Normalizing Flows with Kernel Density Estimation in Estimating Risk of Automated Driving Systems - arXiv, <https://arxiv.org/abs/2507.22429>
69. (PDF) Lévy-Flow Models: Heavy-Tail-Aware Normalizing Flows for Financial Risk Management - ResearchGate, https://www.researchgate.net/publication/403429666_Levy-Flow_Models_Heavy-Tail-Aware_Normalizing_Flows_for_Financial_Risk_Management
70. [2604.00195] Lévy-Flow Models: Heavy-Tail-Aware Normalizing Flows for Financial Risk Management - arXiv, <https://arxiv.org/abs/2604.00195>
71. Normalising Flow Enhanced GARCH Models: A Two-Stage Framework for Flexible Innovation Modelling in Financial Time Series - MDPI, <https://www.mdpi.com/2227-9091/14/5/100>
72. Generation of synthetic financial time series by diffusion models - IDEAS/RePEc, <https://ideas.repec.org/p/arx/papers/2410.18897.html>
73. MiT-WGAN: Financial Time-Series Generation GAN Based on Multi-Convolution Dynamic Fusion and iTransformer - MDPI, <https://www.mdpi.com/2073-8994/17/10/1740>
74. [2507.19003] A diffusion-based generative model for financial time series via geometric Brownian motion - arXiv, <https://arxiv.org/abs/2507.19003>
75. Unified GARCH-Recurrent Neural Networks in Financial Volatility Forecasting - arXiv, <https://arxiv.org/html/2504.09380v2>
76. GARCH-Informed Neural Networks for Volatility Prediction in Financial Markets -

- arXiv, <https://arxiv.org/abs/2410.00288>
77. Master the State-of-the-Art: coding a Temporal Fusion Transformer by Hand - Medium, <https://medium.com/@marcelboersma/temporal-fusion-transformers-from-scratch-a6c10866ccd6>
 78. Quantum Temporal Fusion Transformer - arXiv, <https://arxiv.org/pdf/2508.04048>
 79. Temporal Fusion Transformers for Interpretable Multi-horizon Time Series Forecasting, <https://www.semanticscholar.org/paper/Temporal-Fusion-Transformers-for-Interpretable-Time-Lim-Arik/6a9d69fb35414b8461573df333dba800f254519f>
 80. Improving Coherence in Hierarchical Time Series Forecasting using Structured Temporal Fusion - arXiv, <https://arxiv.org/html/2606.28553v1>
 81. Explainable Graph Neural Networks for Interbank Contagion Surveillance: A Regulatory-Aligned Framework for the U.S. Banking Sector - ResearchGate, https://www.researchgate.net/publication/403905282_Explainable_Graph_Neural_Networks_for_Interbank_Contagion_Surveillance_A_Regulatory-Aligned_Framework_for_the_US_Banking_Sector
 82. Graph Neural Networks for Systemic Financial Risk Forecasting: Modeling Cross-Market Contagion Between Banking Systems and Cryptocurrency Markets | International Journal of Accounting and Economics Studies - Science Publishing Corporation, <https://www.sciencepubco.com/index.php/IJAES/article/view/37895>
 83. Computing Systemic Risk Measures with Graph Neural Networks - IDEAS/RePEc, <https://ideas.repec.org/p/arx/papers/2410.07222.html>
 84. Computing Systemic Risk Measures with Graph Neural Networks - arXiv, <https://arxiv.org/pdf/2410.07222>
 85. Computing Systemic Risk Measures with Graph Neural Networks - arXiv, <https://arxiv.org/html/2410.07222v2>
 86. Relational Graph Modeling for Credit Default Prediction: Heterogeneous GNNs and Hybrid Ensemble Learning - arXiv, <https://arxiv.org/pdf/2601.14633>
 87. Relational Graph Modeling for Credit Default Prediction: Heterogeneous GNNs and Hybrid Ensemble Learning - arXiv, <https://arxiv.org/html/2601.14633v1>
 88. Foundation Models for Credit Risk Prediction: A Game Changer? - arXiv, <https://arxiv.org/html/2605.18147v1>
 89. (PDF) Graph Neural Networks for Systemic Financial Risk Forecasting: Modeling Cross-Market Contagion Between Banking Systems and Cryptocurrency Markets - ResearchGate, https://www.researchgate.net/publication/401392167_Graph_Neural_Networks_for_Systemic_Financial_Risk_Forecasting_Modeling_Cross-Market_Contagion_Between_Banking_Systems_and_Cryptocurrency_Markets
 90. Continuous Reinforcement Learning for Hedging of General Portfolios across Multiple Risk Aversions - arXiv, <https://arxiv.org/pdf/2207.07467>
 91. Risk-Sensitive Reinforcement Learning for Portfolio Optimization Under Stochastic Market Dynamics - MDPI, <https://www.mdpi.com/2227-7390/14/8/1334>
 92. Distributional Reinforcement Learning on Path-dependent Options - arXiv, <https://arxiv.org/html/2507.12657v1>

93. [2501.02087] Beyond CVaR: Leveraging Static Spectral Risk Measures for Enhanced Decision-Making in Distributional Reinforcement Learning - arXiv, <https://arxiv.org/abs/2501.02087>
94. Risk-Aware Deep Reinforcement Learning for Dynamic Portfolio Optimization - arXiv, <https://arxiv.org/html/2511.11481v1>
95. Regret-Optimized Portfolio Enhancement through Deep Reinforcement Learning and Future Looking Rewards - arXiv, <https://arxiv.org/html/2502.02619v1>
96. Evaluating the Quality of the Quantified Uncertainty for (Re)Calibration of Data-Driven Regression Models - arXiv, <https://arxiv.org/html/2508.17761v3>