ChatGPT

Retrieval-Augmented Generation for Large Language Models

Abstract

Retrieval-augmented generation (RAG) has become a central systems paradigm for large language models (LLMs) because it addresses three structural weaknesses of purely parametric models: incomplete coverage of long-tail knowledge, temporal staleness, and weak provenance. Rather than asking a model to internalize all relevant facts during pretraining, RAG augments inference with external memory, typically implemented as textual, graph-structured, multimodal, or task-specific repositories queried by learned or hybrid retrievers. This design turns knowledge access into a dynamic process and recasts generation as conditional reasoning over retrieved evidence. Since the introduction of REALM and RAG, the area has broadened from open-domain question answering into a full research program including retrieval-aware pretraining, query reformulation, reranking, context compression, iterative retrieval-generation loops, graph and memory-augmented retrieval, multimodal RAG, domain-specific adaptation, attribution, and benchmark design [3-15]. 1

This survey synthesizes the modern RAG literature from the perspective of computer science systems design, information retrieval, and language-model alignment. The main argument is that RAG should no longer be viewed as a single prompting trick or a narrow open-book QA recipe. Instead, it is best understood as a family of semi-parametric knowledge systems whose performance depends on coordinated advances in corpus construction, indexing, retrieval, evidence selection, generator control, training, and evaluation. Across the literature, the most effective systems are not those that merely retrieve more documents; they are those that retrieve the right granularity of evidence, expose it to the LLM in an LLM-compatible form, decide when retrieval is necessary, and evaluate correctness jointly with grounding and attribution [16-90]. 2

Introduction

The motivation for RAG follows directly from empirical studies of LLM knowledge. Large pretrained models remain strong pattern learners, but they do not uniformly memorize rare facts, recent events, or domain-specific information; scaling helps, yet long-tail factual coverage remains uneven, as shown by work on knowledge boundaries, temporal freshness, and popularity-sensitive QA [1,8,138-160]. In knowledge-intensive settings, a purely parametric model must either answer from incomplete memory or hallucinate. RAG offers a different contract: the model may consult a mutable external store at inference time, condition its reasoning on retrieved evidence, and, in stronger formulations, expose the provenance supporting its answer [1,3-8,138-171]. 3

Historically, the field emerged from open-domain question answering and retrieval-augmented language modeling. REALM introduced latent retrieval during pretraining, RAG coupled dense retrieval with sequence generation, and kNN-LM and RETRO showed that non-parametric memory could improve perplexity, factuality, and updateability without re-pretraining the entire model [3-6]. Subsequent work diversified along at least four axes. First, retrievers improved from sparse lexical baselines to dense, late-interaction, hybrid, and generative retrieval. Second, generator coupling moved from simple concatenation to fusion-in-decoder, retrieval-aware prompting, compression, reranking, and instruction tuning. Third, reasoning-time control evolved from one-shot retrieve-then-read pipelines to iterative, self-reflective, and agentic loops. Fourth, the knowledge substrate expanded from passage corpora to training sets, graphs, tables, source planners, code repositories, audio, and multimodal memories [3-6,16-24,49-58,83-120]. 4

A useful conceptual distinction is between retrieval as access, retrieval as supervision, and retrieval as reasoning. Retrieval as access uses an external store to provide evidence at inference time. Retrieval as supervision uses retrieved examples, training instances, or queries to improve representation learning, instruction tuning, or domain adaptation. Retrieval as reasoning treats search itself as part of the computational graph of problem solving, allowing the LLM to decompose, refine, sequence, or critique retrieval actions. Much of the recent literature can be interpreted as moving from the first regime to the third [25-48,59-98]. 5

Methods

Foundational formulations and retrieval-side design. At the retrieval layer, the core design choice is how to represent both query and corpus. Sparse lexical methods such as BM25 remain competitive, especially out of domain and in terminology-heavy settings, while dense dual encoders such as DPR and ANCE reduce vocabulary mismatch and scale well with approximate nearest-neighbor search [16-18]. Late-interaction models such as ColBERT and ColBERTv2 retain token-level similarity and therefore often improve first-stage retrieval quality without paying full cross-encoder cost [19-20]. Unsupervised or weakly supervised dense retrieval, represented by Contriever, HyDE-style zero-shot retrieval, Promptagator, and domain-adaptive dense retrievers, broadens applicability when labeled relevance data are scarce [21,24-27,46]. Learned sparse retrievers such as SPLADE and SPLADE v2 further blur the sparse/dense distinction, often delivering strong zero-shot robustness on BEIR-like evaluations [22-24]. At the same time, the literature shows that retrieval granularity matters: passage size, hierarchical retrieval, dense-x retrieval, and structure-aware retrieval all alter the evidence frontier seen by the generator [29,39,42]. 6

Query construction and pre-retrieval optimization. A large fraction of RAG performance variance comes from pre-retrieval decisions rather than the generator itself. Work on query rewriting, step-back abstraction, question clarification, decomposition, and refinement argues that the raw user query is often not an optimal retrieval object [30-33,83-90]. Query Rewriting for Retrieval-Augmented Large Language Models, RQ-RAG, Tree of Clarifications, and related lines show that multi-hop or ambiguous requests benefit from LLM-mediated reformulation before dense retrieval is invoked [30,32-33,88]. In parallel, prompt retrieval and in-context example retrieval—UPRISE, in-context example retrieval, CorpusBrain, unified generative retrievers, and language models as semantic indexers—extend “retrieval” beyond passages to demonstrations, labels, and training exemplars [34-40]. This broadens RAG from document retrieval toward general non-parametric augmentation of the entire inference state. 7

Generator-side coupling and post-retrieval optimization. Early RAG systems mostly concatenated retrieved passages into prompts; subsequent work shows that this is often suboptimal because LLMs are sensitive to order, redundancy, distraction, and positional effects [49-69,180]. Fusion-in-Decoder (FiD) established a powerful architecture for aggregating many passages with stronger cross-passage reasoning than early RAG-style encoder-decoder coupling [49]. RePlug, In-Context RALM, Atlas, Raven, and Zemi extended this direction to black-box LLMs, few-shot learning, and retrieval-augmented in-context learning [50-54]. A major theme since 2023 has been context selection: RECOMP, token elimination, filter learning, LongLLMLingua, RankRAG, and R2AG all attempt to map retrieved material into a representation the generator can use more reliably and more efficiently [55-60,48,47]. These methods are motivated by empirical findings that long contexts are not fully utilized and that models are easily distracted by irrelevant documents [61-67,180]. The emerging consensus is that good RAG is not simply retrieve-more; it is retrieve-then-shape evidence into a generator-compatible context. 8

Joint training, alignment, and domain adaptation. A second major method family treats RAG as a training problem rather than a frozen pipeline. Retrieval-aware pretraining, instruction tuning, dual instruction tuning, RAFT, and domain adaptation studies all show that models trained to use retrieved evidence behave differently from models merely prompted with it [70-82]. The distinction is especially important when distractors are present, when citation is required, or when the retriever and LLM prefer different kinds of evidence. InstructRetro, RA-DIT, RAFT, context tuning, preference-gap bridging, and domain-adaptive RAG collectively suggest that alignment between retriever and generator is a first-order systems issue [43,47,48,71-75]. More broadly, this literature reframes RAG as semi-parametric co-adaptation: the retriever should not only find relevant text for humans, but also compose an evidence package matched to the LLM’ s inductive biases and instruction-following behavior [46-48,70-79]. 9

Adaptive, iterative, and agentic RAG. One-shot retrieval is often insufficient for multi-hop and compositional queries. This has produced a rich line of work in which the model interleaves search and reasoning. Iter-RetGen, FLARE-style active retrieval, IRCoT, Self-RAG, Demonstrate-Search-Predict, Auto-RAG, IM-RAG, Open-RAG, AirRAG, MCTS-RAG, and IterDRAG together define a transition from static pipelines to controlled search policies [83-98]. Their common intuition is that retrieval should be decided during reasoning, not only before reasoning. Self-RAG adds self-reflection tokens controlling when to retrieve and critique evidence [85]. CRAG, TRACE, RuleRAG, and DynamicRAG further show that retrieval quality estimation, chain construction, symbolic rule hints, and reinforcement learning over reranking can substantially improve robustness under imperfect retrieval [89-90,97-98]. The strongest versions of these approaches begin to resemble specialized agents whose primary external tool is a retriever and whose internal action space includes rewrite, continue-search, stop-search, rerank, critique, and abstain [88-98]. 10

Structured memory, graphs, and long-context hybrids. A parallel direction argues that flat passage retrieval is not enough when evidence is relational, global, or temporally accumulated. Graph-based methods—Chain of Knowledge, knowledge-graph prompting, G-Retriever, SubgraphRAG, MedGraphRAG, and the broader GraphRAG literature—attempt to preserve entity relations, traversal structure, and community-level summaries that simple chunk retrieval discards [100-115]. Memory-inspired methods such as HippoRAG, HippoRAG 2, RAPTOR, MemoRAG, MemGPT, and gist-memory readers shift emphasis from isolated relevance matching toward associative memory, hierarchical abstraction, and corpus summarization [104-110]. These methods are especially important for long-form reasoning, multi-hop QA, literature synthesis, and document collections whose useful structure spans many passages. They also sharpen the ongoing debate between long-context LLMs and retrieval. Studies on long-context behavior, especially “lost in the middle,” show that larger windows do not automatically eliminate the need for retrieval; rather, they create new optimization problems around context order, compression, and salience [80,82,104-110,176-181]. 11

Beyond text: multimodal, speech, code, dialogue, and domain-specific RAG. The text-only formulation is now clearly too narrow. MuRAG, RA-CM3, and REVEAL retrieve from multimodal memories containing images, text, question-answer pairs, and knowledge sources; this broadens RAG into a general framework for open-world multimodal grounding [113-115]. InfoSeek and related visual information-seeking work show that retrieval remains crucial even for strong vision-language models because visual recognition alone rarely supplies sufficient world knowledge [116]. Audio and speech variants such as WavRAG and retrieval-augmented spoken dialog models bring the same logic to raw audio and hybrid speech-text knowledge stores [119-120]. Code and scientific RAG introduce repository retrieval, example retrieval, or paper-grounded generation for software engineering and research assistance [121-124].

Dialogue, recommendation, event extraction, legal reasoning, and medical QA further show that domain-specific RAG often depends on source planning, evidence structuring, or controlled knowledge bases rather than general-purpose retrieval alone [123-137]. 12

Challenges and Prospects

The first enduring challenge is **retrieval quality under realistic uncertainty**. Standard benchmarks often assume well-formed queries, stable corpora, and a clean separation between relevant and irrelevant evidence. Real deployments face misspellings, conversational ellipsis, evolving corpora, source conflict, ambiguous intents, incomplete corpora, and adversarial or low-quality documents. QE-RAG, CRUD-RAG, DomainRAG, MultiHop-RAG, HAGRID, and expert-curated attributed benchmarks make progress, but the field still lacks a universally accepted evaluation protocol that jointly measures retrieval recall, faithfulness, attribution, abstention, robustness to counterfactual evidence, and end-to-end utility [161-174]. Moreover, because retrieval, reranking, and generation can fail independently or compensate for one another, accurate diagnosis requires component-level as well as pipeline-level evaluation [12-14,161-174]. 13

The second challenge is **coordination between retriever, generator, and controller**. Much recent work can be read as evidence that naïvely combining a strong retriever with a strong LLM is insufficient. Preference gaps, context-format mismatches, long-context positional biases, and distractor sensitivity all imply that the components of a RAG system have incompatible optimization objectives unless explicitly aligned [47-48,55-67,176-181]. This suggests several promising directions. One is end-to-end or partially end-to-end optimization with retrieval-aware supervision. A second is adaptive routing that decides whether to use retrieval, long context, or parametric memory for a particular query. A third is calibrated abstention, in which the model can state that the evidence base is insufficient or conflicting instead of forcing a low-confidence answer. The best future systems are likely to be mixture-of-memory architectures rather than pure “retrieve everything” or pure “put everything into context window” systems [8,47-48,73,80,83-98,176-181]. 14

The third challenge is **trustworthiness, provenance, and lifecycle management**. In many applications, the core question is not merely whether RAG improves accuracy, but whether the system is auditable, updateable, privacy-preserving, and safe. Domain deployments in medicine, law, enterprise search, and scientific assistance reveal tensions between freshness and source control, between personalization and privacy, and between graph-construction quality and hallucination propagation [110-137]. Multimodal RAG introduces additional problems: cross-modal misalignment, private-document leakage, sparse gold labels, and evaluation difficulty when evidence spans text, tables, charts, or audio [113-120]. Looking forward, several prospects appear especially important: retrieval over structured and semi-structured enterprise data; persistent memory with governance and forgetting policies; citation-aware decoding and answer editing; retrieval over model-generated reasoning traces; and benchmark ecosystems that better reflect multi-source, multi-turn, and long-horizon use rather than static one-shot QA alone [14-15,104-120,161-181]. 15

Conclusion

RAG has evolved from a pragmatic answer-grounding technique into a general architecture for semi-parametric intelligence. The mature view that emerges from the literature is that successful RAG is not defined by the presence of a retriever, but by the quality of coordination between knowledge organization, evidence access, inference-time control, and evaluation. Seminal systems established that external retrieval can extend or refresh model knowledge [3-6]; later work showed that the real frontier lies in evidence shaping, adaptive retrieval, graph and memory structures, joint tuning, and robust evaluation [47-98,104-181]. The most promising future research will likely unify these strands into systems that can decide when to retrieve, what to retrieve, how to represent what was retrieved, how to reason over it transparently, and when not to answer at all. In that sense, RAG is not merely an accessory to LLMs. It is one of the clearest pathways toward language systems that are more updateable, more interpretable, and more scientifically controllable than static pretrained models alone. [16]

References

[1] N. Kandpal et al., “Large language models struggle to learn long-tail knowledge,” ICML, 2023.

[2] Y. Zhang et al., “Siren’ s song in the AI ocean: A survey on hallucination in large language models,” arXiv preprint, 2023, arXiv:2309.01219.

[3] K. Guu et al., “REALM: Retrieval-augmented language model pre-training,” ICML, 2020, arXiv: 2002.08909.

[4] P. Lewis et al., “Retrieval-augmented generation for knowledge-intensive NLP tasks,” NeurIPS, 2020, arXiv:2005.11401.

[5] U. Khandelwal et al., “Generalization through memorization: Nearest neighbor language models,” ICLR, 2020, arXiv:1911.00172.

[6] S. Borgeaud et al., “Improving language models by retrieving from trillions of tokens,” ICML, 2022.

[7] L. Ouyang et al., “Training language models to follow instructions with human feedback,” NeurIPS, 2022.

[8] A. Mallen et al., “When not to trust language models: Investigating effectiveness of parametric and non-parametric memories,” ACL Findings, 2023, arXiv:2212.10511.

[9] Y. Gao et al., “Retrieval-augmented generation for large language models: A survey,” arXiv preprint, 2023, arXiv:2312.10997.

[10] Y. Huang and J. Huang, “A survey on retrieval-augmented text generation for large language models,” ACM Computing Surveys, 2026, arXiv:2404.10981.

[11] W. Fan et al., “RAG and RAU: A survey on retrieval-augmented language models,” arXiv preprint, 2024, arXiv:2404.19543.

[12] H. Yu et al., “Evaluation of retrieval-augmented generation: A survey,” arXiv preprint, 2024, arXiv: 2405.07437.

[13] A. Gan et al., “Retrieval augmented generation evaluation in the era of large language models: A comprehensive survey,” arXiv preprint, 2025, arXiv:2504.14891.

[14] B. Ni et al., “Towards trustworthy retrieval-augmented generation for large language models: A survey,” arXiv preprint, 2025, arXiv:2502.06872.

[15] A. Singh et al., “Agentic retrieval-augmented generation: A survey on agentic RAG,” arXiv preprint, 2025, arXiv:2501.09136.

[16] S. Robertson and H. Zaragoza, “The probabilistic relevance framework: BM25 and beyond,” Foundations and Trends in Information Retrieval, 2009.

[17] V. Karpukhin et al., “Dense passage retrieval for open-domain question answering,” EMNLP, 2020, arXiv:2004.04906.

[18] L. Xiong et al., “Approximate nearest neighbor negative contrastive learning for dense text retrieval,” ICLR, 2021, arXiv:2007.00808.

[19] O. Khattab and M. Zaharia, “ColBERT: Efficient and effective passage search via contextualized late interaction over BERT,” SIGIR, 2020, arXiv:2004.12832.

[20] K. Santhanam et al., “ColBERTv2: Effective and efficient retrieval via lightweight late interaction,” NAACL, 2022, arXiv:2112.01488.

[21] G. Izacard et al., “Unsupervised dense information retrieval with contrastive learning,” TMLR, 2022, arXiv:2112.09118.

[22] T. Formal et al., “SPLADE: Sparse lexical and expansion model for first stage ranking,” SIGIR, 2021, arXiv:2107.05720.

[23] T. Formal et al., “SPLADE v2: Sparse lexical and expansion model for information retrieval,” arXiv preprint, 2021, arXiv:2109.10086.

[24] N. Thakur et al., “BEIR: A heterogeneous benchmark for zero-shot evaluation of information retrieval models,” NeurIPS Datasets and Benchmarks, 2021, arXiv:2104.08663.

[25] J. Du et al., “General purpose text embeddings from pre-trained language models for scalable inference,” Findings of EMNLP, 2020, arXiv:2004.14287.

[26] L. Gao et al., “Precise zero-shot dense retrieval without relevance labels,” ACL, 2023, arXiv: 2212.10496.

[27] Y. Qu et al., “RocketQA: An optimized training approach to dense passage retrieval for open-domain question answering,” NAACL, 2021, arXiv:2010.08191.

[28] Y. Li et al., “More robust dense retrieval with contrastive dual learning,” CIKM, 2021, arXiv: 2107.07773.

[29] Y. Liu et al., “Dense hierarchical retrieval for open-domain question answering,” ACL Findings, 2022, arXiv:2110.15439.

[30] X. Ma et al., “Query rewriting for retrieval-augmented large language models,” EMNLP Findings, 2023, arXiv:2305.14283.

[31] H. S. Zheng et al., “Take a step back: Evoking reasoning via abstraction in large language models,” ICLR, 2024, arXiv:2310.06117.

[32] C.-M. Chan et al., “RQ-RAG: Learning to refine queries for retrieval-augmented generation,” arXiv preprint, 2024, arXiv:2404.00610.

[33] G. Kim et al., “Tree of clarifications: Answering ambiguous questions with retrieval-augmented large language models,” arXiv preprint, 2023, arXiv:2310.14696.

[34] D. Cheng et al., “UPRISE: Universal prompt retrieval for improving zero-shot evaluation,” EMNLP, 2023, arXiv:2303.08518.

[35] L. Wang et al., “Learning to retrieve in-context examples for large language models,” EACL, 2024, arXiv:2307.07164.

[36] J. Chen et al., “A unified generative retriever for knowledge-intensive language tasks,” ACL Findings, 2023, arXiv:2304.14856.

[37] J. Chen et al., “CorpusBrain: Pre-train a generative retrieval model for knowledge-intensive language tasks,” NeurIPS, 2022, arXiv:2208.07652.

[38] J. Guo et al., “CorpusBrain++: A continual generative pre-training framework for knowledge-intensive language tasks,” arXiv preprint, 2024, arXiv:2402.16767.

[39] X. Li et al., “Structure-aware language model pretraining improves dense retrieval on structured data,” arXiv preprint, 2023, arXiv:2305.19912.

[40] B. Jin et al., “Language models as semantic indexers,” arXiv preprint, 2023, arXiv:2310.07815.

[41] P. Zhang et al., “Retrieve anything to augment large language models,” arXiv preprint, 2023, arXiv: 2310.07554.

[42] T. Chen et al., “Dense x retrieval: What retrieval granularity should we use?” arXiv preprint, 2023, arXiv:2312.06648.

[43] R. Anantha et al., “Context tuning for retrieval augmented generation,” arXiv preprint, 2023, arXiv: 2312.05708.

[44] H. Wang et al., “UniMS-RAG: A unified multi-source retrieval-augmented generation for personalized dialogue systems,” arXiv preprint, 2024, arXiv:2401.13256.

[45] H. Wang et al., “Large language models as source planner for personalized knowledge-grounded dialogue,” arXiv preprint, 2023, arXiv:2310.08840.

[46] Z. Yu et al., “Augmentation-adapted retriever improves generalization of language models as generic plug-in,” arXiv preprint, 2023, arXiv:2305.17331.

[47] Z. Ke et al., “Bridging the preference gap between retrievers and LLMs,” ACL, 2024, arXiv: 2401.06954.

[48] F. Ye et al., “R^2AG: Incorporating retrieval information into retrieval-augmented generation,” arXiv preprint, 2024, arXiv:2406.13249.

[49] G. Izacard and E. Grave, “Leveraging passage retrieval with generative models for open domain question answering,” EACL, 2021, arXiv:2007.01282.

[50] W. Shi et al., “REPLUG: Retrieval-augmented black-box language models,” NAACL, 2024, arXiv: 2301.12652.

[51] O. Ram et al., “In-context retrieval-augmented language models,” TACL, 2023, arXiv:2302.00083.

[52] G. Izacard et al., “Atlas: Few-shot learning with retrieval augmented language models,” JMLR, 2023, arXiv:2208.03299.

[53] J. Huang et al., “Raven: In-context learning with retrieval-augmented encoder-decoder language models,” arXiv preprint, 2023, arXiv:2308.07922.

[54] Z. Wang et al., “Zemi: Learning zero-shot semi-parametric language models from multiple tasks,” ICLR, 2023, arXiv:2210.00185.

[55] F. Xu et al., “RECOMP: Improving retrieval-augmented LMs with compression and selective augmentation,” ICLR, 2024, arXiv:2310.04408.

[56] M. Berchansky et al., “Optimizing retrieval-augmented reader models via token elimination,” arXiv preprint, 2023, arXiv:2310.13682.

[57] Z. Wang et al., “Learning to filter context for retrieval-augmented generation,” arXiv preprint, 2023, arXiv:2311.08377.

[58] H. Jiang et al., “LongLLMLingua: Accelerating and enhancing LLMs in long context scenarios via prompt compression,” ACL, 2024, arXiv:2310.06839.

[59] Y. Yu et al., “RankRAG: Unifying context ranking with retrieval-augmented generation in LLMs,” NeurIPS, 2024, arXiv:2407.02485.

[60] Z. Wei et al., “InstructRAG: Instructing retrieval-augmented generation via self-synthesized rationales,” ICLR, 2025, arXiv:2406.13629.

[61] W. Yu et al., “Chain-of-Note: Enhancing robustness in retrieval-augmented language models,” arXiv preprint, 2023, arXiv:2311.09210.

[62] O. Yoran et al., “Making retrieval-augmented language models robust to irrelevant context,” ICLR, 2024, arXiv:2310.01558.

[63] F. Shi et al., “Large language models can be easily distracted by irrelevant context,” ICML, 2023.

[64] F. Cuconasu et al., “The power of noise: Redefining retrieval for RAG systems,” arXiv preprint, 2024, arXiv:2401.14887.

[65] N. Thakur et al., “NoMIRACL: Knowing when you don’t know for robust multilingual retrieval-augmented generation,” arXiv preprint, 2023, arXiv:2312.11361.

[66] Y. Liu et al., “ReCall: A benchmark for LLMs’ robustness against external counterfactual knowledge,” arXiv preprint, 2023, arXiv:2311.08147.

[67] J. Baek et al., “Knowledge-augmented language model verification,” arXiv preprint, 2023, arXiv: 2310.12836.

[68] S. Xu et al., “Search-in-the-chain: Towards accurate, credible and traceable large language models for knowledge-intensive tasks,” arXiv preprint, 2023, arXiv:2304.14732.

[69] P. Jain et al., “1-Pager: One pass answer generation and evidence retrieval,” arXiv preprint, 2023, arXiv:2310.16568.

[70] B. Wang et al., “Shall we pretrain autoregressive language models with retrieval? A comprehensive study,” arXiv preprint, 2023, arXiv:2304.06762.

[71] B. Wang et al., “InstructRetro: Instruction tuning post retrieval-augmented pretraining,” arXiv preprint, 2023, arXiv:2310.07713.

[72] X. V. Lin et al., “RA-DIT: Retrieval-augmented dual instruction tuning,” ICLR, 2024, arXiv: 2310.01352.

[73] T. Zhang et al., “RAFT: Adapting language model to domain specific RAG,” arXiv preprint, 2024, arXiv:2403.10131.

[74] O. Ovadia et al., “Fine-tuning or retrieval? Comparing knowledge injection in LLMs,” arXiv preprint, 2023, arXiv:2312.05934.

[75] S. Siriwardhana et al., “Improving the domain adaptation of retrieval augmented generation models for open domain question answering,” TACL, 2023.

[76] Z. Feng et al., “Retrieval-generation synergy augmented large language models,” arXiv preprint, 2023, arXiv:2310.05149.

[77] Y. Wang et al., “Self-knowledge guided retrieval augmentation for large language models,” arXiv preprint, 2023, arXiv:2310.05002.

[78] Z. Luo et al., “Augmented large language models with parametric knowledge guiding,” arXiv preprint, 2023, arXiv:2305.04757.

[79] X. Wang et al., “KnowledgeGPT: Enhancing large language models with retrieval and storage access on knowledge bases,” arXiv preprint, 2023, arXiv:2308.11761.

[80] P. Xu et al., “Retrieval meets long context large language models,” arXiv preprint, 2023, arXiv: 2310.03025.

[81] Y. Liu et al., “Exploring the integration strategies of retriever and large language models,” arXiv preprint, 2023, arXiv:2308.12574.

[82] H.-T. Chen et al., “Understanding retrieval augmentation for long-form question answering,” arXiv preprint, 2023, arXiv:2310.12150.

[83] Z. Shao et al., “Enhancing retrieval-augmented large language models with iterative retrieval-generation synergy,” EMNLP Findings, 2023, arXiv:2305.15294.

[84] Z. Jiang et al., “Active retrieval augmented generation,” EMNLP Findings, 2023, arXiv:2305.06983.

[85] A. Asai et al., “Self-RAG: Learning to retrieve, generate, and critique through self-reflection,” ICLR, 2024, arXiv:2310.11511.

[86] H. Trivedi et al., “Interleaving retrieval with chain-of-thought reasoning for knowledge-intensive multi-step questions,” ACL, 2023, arXiv:2212.10509.

[87] O. Khattab et al., “Demonstrate-Search-Predict: Composing retrieval and language models for knowledge-intensive NLP,” TMLR, 2024, arXiv:2212.14024.

[88] T. Yu et al., “Auto-RAG: Autonomous retrieval-augmented generation for large language models,” arXiv preprint, 2024, arXiv:2411.19443.

[89] S.-Q. Yan et al., “Corrective retrieval augmented generation,” arXiv preprint, 2024, arXiv:2401.15884.

[90] J. Fang et al., “TRACE the Evidence: Constructing knowledge-grounded reasoning chains for retrieval-augmented generation,” SIGIR, 2024, arXiv:2406.11460.

[91] Z. Zhang et al., “IAG: Induction-augmented generation framework for answering reasoning questions,” EMNLP, 2023.

[92] S. B. Islam et al., “Open-RAG: Enhanced retrieval-augmented reasoning with open-source large language models,” arXiv preprint, 2024, arXiv:2410.01782.

[93] D. Yang et al., “IM-RAG: Multi-round retrieval-augmented generation through learning inner monologues,” SIGIR, 2024, arXiv:2405.13021.

[94] W. Feng et al., “AirRAG: Activating intrinsic reasoning for retrieval augmented generation using tree-based search,” arXiv preprint, 2025, arXiv:2501.10053.

[95] Y. Hu et al., “MCTS-RAG: Enhancing retrieval-augmented generation with Monte Carlo tree search,” arXiv preprint, 2025, arXiv:2503.20757.

[96] Z. Yue et al., “Inference scaling for long-context retrieval augmented generation,” arXiv preprint, 2024, arXiv:2410.04343.

[97] Z. Chen et al., “RuleRAG: Rule-guided retrieval-augmented generation with language models for question answering,” arXiv preprint, 2024, arXiv:2410.22353.

[98] J. Sun et al., “DynamicRAG: Leveraging outputs of large language model as feedback for dynamic reranking in retrieval-augmented generation,” arXiv preprint, 2025, arXiv:2505.07233.

[99] L. Luo et al., “Reasoning on graphs: Faithful and interpretable large language model reasoning,” arXiv preprint, 2023, arXiv:2310.01061.

[100] Y. Wang et al., “Knowledge graph prompting for multi-document question answering,” arXiv preprint, 2023, arXiv:2308.11730.

[101] X. Li et al., “Chain of knowledge: A framework for grounding large language models with structured knowledge bases,” arXiv preprint, 2023, arXiv:2305.13269.

[102] X. He et al., “G-Retriever: Retrieval-augmented generation for textual graph understanding and question answering,” arXiv preprint, 2024, arXiv:2402.07630.

[103] M. Li et al., “Simple is effective: The roles of graphs and large language models in knowledge-graph-based retrieval-augmented generation,” arXiv preprint, 2024, arXiv:2410.20724.

[104] B. J. Gutiérrez et al., “HippoRAG: Neurobiologically inspired long-term memory for large language models,” arXiv preprint, 2024, arXiv:2405.14831.

[105] B. J. Gutiérrez et al., “From RAG to memory: Non-parametric continual learning for large language models,” arXiv preprint, 2025, arXiv:2502.14802.

[106] P. Sarthi et al., “RAPTOR: Recursive abstractive processing for tree-organized retrieval,” arXiv preprint, 2024, arXiv:2401.18059.

[107] H. Qian et al., “MemoRAG: Moving towards next-gen RAG via memory-inspired knowledge discovery,” arXiv preprint, 2024, arXiv:2409.05591.

[108] C. Packer et al., “MemGPT: Towards LLMs as operating systems,” arXiv preprint, 2023, arXiv: 2310.08560.

[109] K.-H. Lee et al., “A human-inspired reading agent with gist memory of very long contexts,” arXiv preprint, 2024, arXiv:2402.09727.

[110] J. Wu et al., “Medical Graph RAG: Towards safe medical large language model via graph retrieval-augmented generation,” arXiv preprint, 2024, arXiv:2408.04187.

[111] X. Zhao et al., “MedRAG: Enhancing retrieval-augmented generation with knowledge-graph-elicited reasoning for healthcare copilot,” arXiv preprint, 2025, arXiv:2502.04413.

[112] X. Jiang et al., “Think and retrieval: A hypothesis knowledge graph enhanced medical large language model,” arXiv preprint, 2023, arXiv:2312.15883.

[113] W. Chen et al., “MuRAG: Multimodal retrieval-augmented generator for open question answering over images and text,” EMNLP, 2022, arXiv:2210.02928.

[114] M. Yasunaga et al., "Retrieval-augmented multimodal language modeling," ICML, 2023, arXiv: 2211.12561.

[115] Z. Hu et al., "REVEAL: Retrieval-augmented visual-language pre-training with multi-source multimodal knowledge memory," CVPR, 2023, arXiv:2212.05221.

[116] Y. Chen et al., "Can pre-trained vision and language models answer visual information-seeking questions?" EMNLP, 2023, arXiv:2302.11713.

[117] J. Li et al., "BLIP-2: Bootstrapping language-image pre-training with frozen image encoders and large language models," ICML, 2023, arXiv:2301.12597.

[118] A. Yang et al., "Vid2Seq: Large-scale pretraining of a visual language model for dense video captioning," CVPR, 2023.

[119] Y. Chen et al., "WavRAG: Audio-integrated retrieval augmented generation for spoken dialogue models," arXiv preprint, 2025, arXiv:2502.14727.

[120] Y. Chen et al., "Retrieval-augmented end-to-end spoken dialog models," arXiv preprint, 2024, arXiv:2402.01828.

[121] N. Nashid et al., "Retrieval-based prompt selection for code-related few-shot learning," ICSE, 2023.

[122] H. Husain et al., "CodeSearchNet challenge: Evaluating the state of semantic code search," arXiv preprint, 2019, arXiv:1909.09436.

[123] J. Cui et al., "ChatLaw: Open-source legal large language model with integrated external knowledge bases," arXiv preprint, 2023, arXiv:2306.16092.

[124] J. Lala et al., "PaperQA: Retrieval-augmented generative agent for scientific research," arXiv preprint, 2023, arXiv:2312.07559.

[125] P. Ranade and A. Joshi, "Fabula: Intelligence report generation using retrieval-augmented narrative construction," arXiv preprint, 2023, arXiv:2310.13848.

[126] Y. Gao et al., "Chat-REC: Towards interactive and explainable LLMs-augmented recommender system," arXiv preprint, 2023, arXiv:2303.14524.

[127] S. Rajput et al., "Recommender systems with generative retrieval," arXiv preprint, 2023, arXiv: 2305.05065.

[128] M. Seo et al., "Retrieval-augmented data augmentation for low-resource domain tasks," arXiv preprint, 2024, arXiv:2402.13482.

[129] Z. Guo et al., "Prompt-guided retrieval augmentation for non-knowledge-intensive tasks," arXiv preprint, 2023, arXiv:2305.17653.

[130] F. Luo and M. Surdeanu, “Divide & conquer for entailment-aware multi-hop evidence retrieval,” arXiv preprint, 2023, arXiv:2311.02616.

[131] Y. Ren et al., “Retrieve-and-sample: Document-level event argument extraction via hybrid retrieval augmentation,” ACL, 2023.

[132] X. Du and H. Ji, “Retrieval-augmented generative question answering for event argument extraction,” arXiv preprint, 2022, arXiv:2211.07067.

[133] M. Kang et al., “Knowledge graph-augmented language models for knowledge-grounded dialogue generation,” arXiv preprint, 2023, arXiv:2305.18846.

[134] W. Shen et al., “Retrieval-generation alignment for end-to-end task-oriented dialogue system,” arXiv preprint, 2023, arXiv:2310.08877.

[135] T. Shi et al., “Dual-feedback knowledge retrieval for task-oriented dialogue systems,” arXiv preprint, 2023, arXiv:2310.14528.

[136] X. Xu et al., “Long time no see! Open-domain conversation with long-term persona memory,” ACL Findings, 2022, arXiv:2203.05797.

[137] E. Dinan et al., “Wizard of Wikipedia: Knowledge-powered conversational agents,” ICLR, 2019, arXiv:1811.01241.

[138] T. Kwiatkowski et al., “Natural Questions: A benchmark for question answering research,” TACL, 2019.

[139] M. Joshi et al., “TriviaQA: A large scale distantly supervised challenge dataset for reading comprehension,” ACL, 2017, arXiv:1705.03551.

[140] P. Rajpurkar et al., “SQuAD: 100,000+ questions for machine comprehension of text,” EMNLP, 2016, arXiv:1606.05250.

[141] Z. Yang et al., “HotpotQA: A dataset for diverse, explainable multi-hop question answering,” EMNLP, 2018, arXiv:1809.09600.

[142] H. Trivedi et al., “MuSiQue: Multihop questions via single-hop question composition,” TACL, 2022.

[143] X. Ho et al., “Constructing a multi-hop QA dataset for comprehensive evaluation of reasoning steps,” COLING, 2020, arXiv:2011.01060.

[144] A. Fan et al., “ELI5: Long form question answering,” ACL, 2019, arXiv:1907.09190.

[145] T. Kočiský et al., “The NarrativeQA reading comprehension challenge,” TACL, 2018.

[146] I. Stelmakh et al., “ASQA: Factoid questions meet long-form answers,” EMNLP, 2022, arXiv: 2204.06092.

[147] M. Zhong et al., “QMSum: A new benchmark for query-based multi-domain meeting summarization,” NAACL, 2021, arXiv:2104.05938.

[148] P. Dasigi et al., “A dataset of information-seeking questions and answers anchored in research papers,” NAACL, 2021, arXiv:2105.03011.

[149] T. Möller et al., “COVID-QA: A question answering dataset for COVID-19,” ACL NLP-COVID Workshop, 2020.

[150] R. Y. Pang et al., “QuALITY: Question answering with long input texts, yes!,” NAACL, 2022, arXiv: 2112.08608.

[151] P. Clark et al., “Think you have solved question answering? Try ARC, the AI2 reasoning challenge,” arXiv preprint, 2018, arXiv:1803.05457.

[152] A. Talmor et al., “CommonsenseQA: A question answering challenge targeting commonsense knowledge,” NAACL, 2019, arXiv:1811.00937.

[153] R. Zellers et al., “HellaSwag: Can a machine really finish your sentence?,” ACL, 2019, arXiv: 1905.07830.

[154] D. Hendrycks et al., “Measuring massive multitask language understanding,” ICLR, 2021, arXiv: 2009.03300.

[155] M. Geva et al., “Did Aristotle use a laptop? A question answering benchmark with implicit reasoning strategies,” TACL, 2021.

[156] J. Thorne et al., “FEVER: A large-scale dataset for fact extraction and verification,” NAACL, 2018, arXiv:1803.05355.

[157] F. Petroni et al., “KILT: A benchmark for knowledge intensive language tasks,” NAACL, 2021, arXiv: 2009.02252.

[158] T. Vu et al., “FreshLLMs: Refreshing large language models with search engine augmentation,” arXiv preprint, 2023, arXiv:2310.03214.

[159] J. Kasai et al., “RealTime QA: What’s the answer right now?,” NeurIPS Datasets and Benchmarks, 2022, arXiv:2207.13332.

[160] S. J. Amouyal et al., “QAMPARI: An open-domain question answering benchmark for questions with many answers from multiple paragraphs,” INLG Workshop, 2023, arXiv:2205.12665.

[161] S. Es et al., “RAGAs: Automated evaluation of retrieval augmented generation,” EACL Demo/ Industry Track, 2024, arXiv:2309.15217.

[162] J. Saad-Falcon et al., “ARES: An automated evaluation framework for retrieval-augmented generation systems,” NAACL, 2024, arXiv:2311.09476.

[163] R. Friel et al., “RAGBench: Explainable benchmark for retrieval-augmented generation systems,” arXiv preprint, 2024, arXiv:2407.11005.

[164] Y. Lyu et al., “CRUD-RAG: A comprehensive Chinese benchmark for retrieval-augmented generation of large language models,” arXiv preprint, 2024, arXiv:2401.17043.

[165] J. Chen et al., “Benchmarking large language models in retrieval-augmented generation,” arXiv preprint, 2023, arXiv:2309.01431.

[166] Y. Hoshi et al., “RALLE: A framework for developing and evaluating retrieval-augmented large language models,” arXiv preprint, 2023, arXiv:2308.10633.

[167] K. Zhang et al., “QE-RAG: A robust retrieval-augmented generation benchmark for query entry errors,” arXiv preprint, 2025, arXiv:2504.04062.

[168] S. Wang et al., “DomainRAG: A Chinese benchmark for evaluating domain-specific retrieval-augmented generation,” arXiv preprint, 2024, arXiv:2406.05654.

[169] Y. Tang and Y. Yang, “MultiHop-RAG: Benchmarking retrieval-augmented generation for multi-hop queries,” arXiv preprint, 2024, arXiv:2401.15391.

[170] S. Rosenthal et al., “CLAPNQ: Cohesive long-form answers from passages in Natural Questions for RAG systems,” TACL, 2025, arXiv:2404.02103.

[171] E. Kamalloo et al., “HAGRID: A human-LLM collaborative dataset for generative information-seeking with attribution,” arXiv preprint, 2023, arXiv:2307.16883.

[172] C. Malaviya et al., “ExpertQA: Expert-curated questions and attributed answers,” arXiv preprint, 2023, arXiv:2309.07852.

[173] M. Zhong et al., “FActScore: Fine-grained atomic evaluation of factual precision in long-form text generation,” EMNLP, 2023, arXiv:2305.14251.

[174] N. F. Liu et al., “Lost in the middle: How language models use long contexts,” TACL, 2024, arXiv:2307.03172.

[175] G. Xiao et al., “Efficient streaming language models with attention sinks,” arXiv preprint, 2023, arXiv:2309.17453.

[176] Y. Bai et al., “LongBench: A bilingual, multitask benchmark for long context understanding,” ACL, 2024, arXiv:2308.14508.

[177] Y. Bai et al., “LongBench v2: Towards deeper understanding and reasoning on realistic long-context multitasks,” arXiv preprint, 2024, arXiv:2412.15204.

[178] Z. Chen et al., “LongBench Pro: A more realistic and comprehensive bilingual long-context evaluation benchmark,” arXiv preprint, 2026, arXiv:2601.02872.

[179] J. Kaplan et al., “Scaling laws for neural language models,” arXiv preprint, 2020, arXiv:2001.08361.

[180] D. Zhou et al., "Least-to-most prompting enables complex reasoning in large language models," ICLR, 2023, arXiv:2205.10625.

[181] S. Dhuliawala et al., "Chain-of-verification reduces hallucination in large language models," arXiv preprint, 2023, arXiv:2309.11495.

[182] S. Kim et al., "The CoT collection: Improving zero-shot and few-shot learning of language models via chain-of-thought fine-tuning," EMNLP, 2023, arXiv:2305.14045.

[183] K. Cobbe et al., "Training verifiers to solve math word problems," arXiv preprint, 2021, arXiv: 2110.14168.

[184] R. Nakano et al., "WebGPT: Browser-assisted question-answering with human feedback," arXiv preprint, 2021, arXiv:2112.09332.

[185] T. Schick et al., "Toolformer: Language models can teach themselves to use tools," NeurIPS, 2023, arXiv:2302.04761.

---

1 4 16 REALM: Retrieval-Augmented Language Model Pre-Training
https://arxiv.org/abs/2002.08909?utm_source=chatgpt.com

2 6 Dense Passage Retrieval for Open-Domain Question Answering
https://arxiv.org/abs/2004.04906?utm_source=chatgpt.com

3 Generalization through Memorization: Nearest Neighbor Language Models
https://arxiv.org/abs/1911.00172?utm_source=chatgpt.com

5 RQ-RAG: Learning to Refine Queries for Retrieval Augmented Generation
https://arxiv.org/abs/2404.00610?utm_source=chatgpt.com

7 MultiHop-RAG: Benchmarking Retrieval-Augmented Generation for Multi-Hop Queries
https://arxiv.org/abs/2401.15391?utm_source=chatgpt.com

8 In-Context Retrieval-Augmented Language Models
https://arxiv.org/abs/2302.00083?utm_source=chatgpt.com

9 RAFT: Adapting Language Model to Domain Specific RAG
https://arxiv.org/abs/2403.10131?utm_source=chatgpt.com

10 Corrective Retrieval Augmented Generation
https://arxiv.org/abs/2401.15884?utm_source=chatgpt.com

11 G-Retriever: Retrieval-Augmented Generation for Textual Graph Understanding and Question Answering
https://arxiv.org/abs/2402.07630?utm_source=chatgpt.com

12 MuRAG: Multimodal Retrieval-Augmented Generator for Open Question Answering over Images and Text
https://arxiv.org/abs/2210.02928?utm_source=chatgpt.com

13 RAGBench: Explainable Benchmark for Retrieval-Augmented Generation Systems
https://arxiv.org/abs/2407.11005?utm_source=chatgpt.com

14 Bridging the Preference Gap between Retrievers and LLMs
https://arxiv.org/abs/2401.06954?utm_source=chatgpt.com

15 Towards Trustworthy Retrieval Augmented Generation for ...
https://arxiv.org/abs/2502.06872?utm_source=chatgpt.com