# Deep Learning for Financial Risk Modeling under Uncertainty

## Abstract

Financial risk modeling is fundamentally an exercise in decision-making under uncertainty: the target quantities of practical interest are not merely conditional means, but tail losses, default intensities, stress scenarios, confidence regions, and policy decisions that remain reliable under regime shifts, sparse extremes, and structural breaks. Recent work has moved the field from narrow point forecasting toward representation-rich, probabilistic, and decision-aware deep learning systems that target value-at-risk, expected shortfall, loss-given-default, mortgage transition risk, systemic contagion, portfolio downside risk, and hedging loss distributions. At the same time, this literature has imported a broad uncertainty-quantification toolkit from machine learning, including Bayesian neural networks, deep ensembles, evidential methods, calibrated quantile regression, conformal prediction, graph neural networks, generative scenario models, and risk-sensitive reinforcement learning. This survey synthesizes that literature with a deliberate focus on the interaction between deep architectures and uncertainty handling. It argues that the central scientific question is no longer whether deep learning can improve predictive accuracy in finance, but when uncertainty-aware deep models provide better calibrated, tail-sensitive, and decision-relevant risk estimates than classical econometric or scorecard baselines. The survey organizes the area around method families, clarifies the relationship between predictive uncertainty and financial risk measures, reviews empirical evidence across market, credit, mortgage, systemic, and portfolio settings, and identifies unresolved challenges in calibration, temporal distribution shift, scarcity of extreme events, explainability, governance, and evaluation. The resulting picture is mixed but encouraging: deep learning is most compelling when it is used not as a black-box forecaster of average outcomes, but as a probabilistic modeling layer for tails, transitions, scenarios, and robust controls under uncertainty. [1]

## Introduction

Financial risk modeling sits at the intersection of statistical forecasting, optimization, and institutional decision-making. Classical frameworks established the language of variance, conditional heteroskedasticity, coherent risk measures, quantiles, expected shortfall, and formal backtesting [1–11]. Yet these tools were primarily developed for low-dimensional settings and relatively rigid distributional assumptions. As financial datasets became higher-dimensional and more heterogeneous—combining prices, order books, macro variables, text, borrower trajectories, collateral information, and networks of obligations—deep learning began to attract attention as a representation-learning alternative to hand-crafted feature engineering [12–35]. Reviews of financial machine learning and deep learning now consistently identify risk management, credit scoring, bankruptcy and distress prediction, market forecasting, and portfolio allocation as core application areas, with sequence models, convolutional architectures, transformers, graph neural networks, generative models, and reinforcement learning forming the main methodological backbone [25–31, 99–109]. [2]

What makes this area scientifically distinctive is that "uncertainty" appears at several levels simultaneously. There is **aleatoric uncertainty** arising from noisy and partially observed market or borrower outcomes; **epistemic uncertainty** arising from finite samples, model misspecification, and limited tail support; **distributional uncertainty** caused by temporal drift, crises, and policy changes; and **decision uncertainty** produced when a learned predictor is embedded inside a portfolio, capital, or hedging workflow [60–83]. In finance, these layers are inseparable. A default probability that is slightly miscalibrated can distort pricing and capital allocation. A volatility model that is sharp on average but underestimates tail risk can fail regulatory backtests. A portfolio learner that maximizes expected return without uncertainty controls can produce unstable allocations. For this reason, recent work increasingly treats uncertainty quantification not as an auxiliary diagnostic, but as part of the primary modeling target [23, 24, 41–46, 58–83]. 3

A second reason for renewed interest is that modern risk tasks are often **non-Euclidean, sequential, and counterfactual**. Mortgage and consumer finance involve longitudinal transition behavior rather than one-shot classification [16–18, 87–91]. Systemic risk depends on network topology and contagion channels rather than only node-level covariates [84–91]. Internal market-risk models require synthetic scenarios that preserve dependence structure and generate plausible tail events, raising questions of calibration, memorization, and validation [37–45]. Hedging and allocation problems are embedded control tasks in which the relevant loss is path-dependent and risk-sensitive [46–59]. Deep learning has become important in these domains not only because of its expressive power, but because it can unify representation learning, probabilistic forecasting, and downstream optimization inside one trainable pipeline. 4

The literature nevertheless remains fragmented. Market-risk papers often emphasize VaR breaches and ES scoring, while credit-risk papers focus on AUC, KS, or Brier-type metrics. Uncertainty papers emphasize coverage, calibration, or out-of-distribution detection. Portfolio and hedging papers emphasize realized utility, drawdown, or CVaR. The central aim of this survey is therefore integrative: to connect these strands and make explicit how deep learning architectures, uncertainty quantification techniques, and financial risk objectives interact. Throughout, the survey adopts a deliberately conservative view of relevance. It prioritizes publications that directly concern risk modeling, tail-sensitive decision-making, or uncertainty-aware deep learning in finance, rather than mechanically inflating the bibliography with weakly related trading papers.

## Methods

**Risk targets, supervision, and loss design.**
A useful organizing principle is to begin with the target quantity. Some models estimate **distributional features of asset returns**, such as volatility, quantiles, VaR, or ES [2, 3, 5, 8–11, 41–45]. Others estimate **borrower or instrument transition probabilities**, such as probability of default, delinquency trajectories, prepayment, foreclosure, or recovery rates [12–24]. A third group estimates **system-level quantities**, including systemic importance scores, contagion exposure, and bailout allocations on financial networks [84–91]. A fourth group learns **decision policies** directly, targeting portfolio utility, downside risk, or hedging losses rather than a single predictive statistic [46–59, 110]. This distinction matters because the right learning objective changes with the target. For VaR or prediction intervals, quantile losses and conformal calibration are natural [8, 76–83]. For ES or CVaR, coherent-risk objectives and joint elicitation become important [5, 10, 11, 52–59]. For borrower trajectories, multi-task or survival-style supervision can better reflect temporal risk horizons [16–18, 87–91]. For systemic allocation or hedging, the training loss should mirror the downstream capital or risk measure rather than a purely statistical proxy [46–59, 84–89]. Recent finance papers increasingly make this connection explicit, for example by jointly modeling VaR and

ES with recurrent networks, selecting portfolios based on conformal intervals, or blending deep hedging policies with uncertainty-sensitive risk control [45, 46, 59]. 5

**Sequence, multimodal, and representation-learning architectures.**
Deep learning entered financial risk primarily through feed-forward and recurrent models that improved nonlinear prediction on tabular or temporal data [12–18, 25, 32]. LSTMs and related recurrent architectures remain influential when risk unfolds through time, as in market states, mortgage behavior, delinquency progression, and VaR/ES forecasting [16, 17, 32, 41, 45]. CNN-style architectures became important when financial inputs could be arranged as structured tensors, such as limit order books, local temporal windows, or spatialized feature maps [17, 33, 34]. Attention-based models and transformers generalized this trend by enabling variable selection, long-range temporal dependency modeling, and hybrid static–dynamic conditioning [35, 36, 90]. In borrower-level credit applications, the most promising recent pattern is not any single architecture, but modality matching: tabular MLPs or NODE-like learners for application features, sequence models for repayment histories and transaction streams, transformers for mixed heterogeneous covariates, and graph neural networks when the data generating process is relational—for example, in SME ecosystems, supply-chain finance, or interbank connectivity [17, 18, 31, 88–91, 99–107]. Newer credit-risk frameworks such as FinLangNet and related sequence-aware systems underscore that temporal financial behavior can often be modeled more effectively as structured trajectories than as static snapshots [90]. 6

**Probabilistic deep learning and uncertainty quantification.**
The most important methodological shift in this area is the move from deterministic prediction to explicit uncertainty decomposition. Bayesian neural networks and approximate posterior methods provide one family of solutions, represented by variational Bayes, dropout-as-approximate-inference, and related posterior approximation methods [61–64]. Ensemble methods provide another, often empirically stronger, route by using model disagreement as a proxy for epistemic uncertainty [65, 66, 98]. A third family treats uncertainty as a learnable output distribution, including heteroskedastic mean–variance networks, evidential classifiers and regressors, prior networks, and single-model quantile approaches [69–73]. The finance literature has used all three strategies, but the adoption pattern is domain-specific. In credit risk, uncertainty-aware deep learning has been most explicit in LGD estimation and default scoring, where Bayesian neural networks, deep evidential regression, and uncertainty diagnostics are attractive because they support capital planning, risk ranking, and model-risk governance [22–24]. In market and hedging applications, ensembles and quantile-style methods have been more common because they scale well and align with tail-sensitive decision criteria [41–46, 59]. Across domains, the central empirical lesson is that high predictive accuracy does not imply calibrated uncertainty. Calibration, sharpness, and coverage must therefore be evaluated separately, and sometimes optimized jointly [67, 68, 72, 73, 79–83]. 7

**Conformal, quantile, and coverage-oriented methods.**
Distribution-free uncertainty quantification has become particularly attractive for financial risk because finite-sample coverage statements are conceptually aligned with capital protection and interval-based decision rules [78–83]. Conformalized quantile regression and jackknife-style methods allow researchers to place uncertainty bands around forecasts even when the underlying deep model is flexible and black-box [79, 80]. In finance, this has motivated conformal portfolio selection, interval-informed return ranking, and recent efforts to adapt conformal methods to temporal dependence and shift [46, 78–82]. The appeal is clear: unlike purely Bayesian approximations, conformal methods provide validity guarantees under relatively weak assumptions, and unlike plain quantile training, they explicitly target coverage. The difficulty is equally clear: financial data violate exchangeability, experience regime shifts, and exhibit heavy tails, so coverage guarantees may become approximate or delayed under temporal adaptation. As a result, the

most promising direction is hybridization—deep feature extractors or quantile predictors calibrated by conformal layers and monitored online for drift [45, 46, 79–83]. 8

**Tail-aware market risk and scenario generation.**
A large share of finance-specific deep risk modeling focuses on market-risk estimation, where the key question is not simply the conditional mean return but the lower tail of the predictive distribution. Here the literature includes direct VaR estimation, joint VaR–ES forecasting, and generative scenario models that simulate stress paths while preserving cross-sectional and temporal dependence [41–45]. Recurrent models have recently shown competitive or superior performance relative to rolling-window, GARCH-type, and GAS benchmarks in joint VaR/ES forecasting [45]. VAE- and GAN-based approaches have been used to construct alternative scenario generators that learn dependency structure without strong parametric assumptions [37–44]. In the insurance and internal-model literature, this line of work has progressed furthest because one-year solvency capital calculations naturally require scenario generation rather than only single-step forecasting [43, 44]. The methodological bottleneck is validation: a realistic-looking scenario generator may still memorize historical data, distort dependence, or underrepresent extreme losses. The proposal of memorization-ratio diagnostics and dependence checks is therefore an important contribution, because it reframes model validation around the actual failure modes of generative risk systems [43, 44]. 9

**Credit, bankruptcy, and borrower-level risk under uncertainty.**
Credit risk remains the most mature domain for deep learning under uncertainty because the supervision is often clearer than in market prediction, and the decisions are closer to operational lending workflows [12–24, 31, 99–107]. Early machine-learning work demonstrated that nonlinear models could improve consumer credit scoring and default discrimination relative to classical scorecards [12–15]. Deep learning then extended this agenda in several directions: mortgage risk with massive loan-level panel data [16], convolutional models for mortgage default [17], text-enhanced bankruptcy prediction from disclosures [19, 20], and tabular or sequence learners for retail, peer-to-peer, and online-lending risk [18, 21, 90, 99–107]. The literature is also notable for its sobriety. Some large empirical studies conclude that deep learning is not universally superior to boosted trees in conventional scorecard settings, especially on medium-sized tabular datasets [18]. That negative result is scientifically valuable: it indicates that the benefit of deep learning is concentrated in settings with temporal depth, heterogeneous modalities, or large-scale nonlinear interactions, not in every credit task. Uncertainty-aware credit models add another layer by producing predictive intervals, posterior variability, or evidential confidence around PD or LGD estimates [22–24]. In regulated environments, these uncertainty signals are often as important as the point predictions themselves, because they can guide human review, policy overrides, conservative pricing, or abstention mechanisms for borderline cases. 10

**Relational and systemic risk with graph representation learning.**
Systemic risk is intrinsically relational: shocks propagate through bilateral exposures, common asset holdings, governance ties, clearing mechanisms, and correlated funding channels. This makes graph neural networks and permutation-equivariant architectures natural candidates when the learning target depends on network structure [84–91]. Recent work shows two distinct but complementary directions. One is **predictive graph learning**, where GNNs are trained to classify or regress institutional systemic importance based on graph features and simulated or partially observed labels [88]. The other is **structural graph approximation**, where permutation-equivariant neural networks are used to approximate systemic risk measures derived from clearing-network models and bailout optimization problems [89]. Together, these papers mark an important transition: systemic-risk learning is moving from graph-inspired feature engineering toward architectures that preserve the mathematical symmetries of the underlying financial

network. The challenge, however, is that ground truth remains scarce; many labels are simulated, aggregated, or proxy-based. This raises a familiar uncertainty problem: graph models may appear strong in-sample while remaining fragile to network misspecification, partial observability, or changes in market microstructure. Explainable GNNs and uncertainty-aware graph learning are therefore likely to be crucial next steps [88–91]. 11

**Decision-aware learning, reinforcement learning, and hedging.**
A final family of methods departs from "predict-then-optimize" altogether and learns policies directly under a risk-sensitive objective. In portfolio management, this includes deep reinforcement learning methods that map state representations to allocations while optimizing a reward linked to wealth growth, Sharpe-like criteria, or downside control [49–58]. In derivative risk management, deep hedging learns trading strategies to minimize convex risk measures under frictions, high dimensionality, and model uncertainty [47, 48, 59]. The conceptual attraction of decision-aware learning is that it aligns training with what practitioners actually care about: tail losses, transaction costs, risk limits, or hedging error. But it also intensifies the uncertainty problem because the learner must now extrapolate counterfactually over state–action trajectories that may not be well represented in the data. This is why risk-sensitive RL, distributional RL, robust CVaR optimization, and uncertainty-aware ensembling are particularly relevant in finance [52–59]. The recent introduction of uncertainty-aware deep hedging is especially illustrative: ensemble disagreement can be converted into a confidence-weighted blend between a learned hedge and a classical model-based hedge, thereby turning uncertainty estimation into a live control variable rather than a mere post hoc diagnostic [59]. 12

## Challenges and Prospects

The first grand challenge is **temporal distribution shift**. Financial risk models fail less often because they cannot interpolate, and more often because the data-generating process changes. Crises, monetary tightening, macroprudential interventions, borrower-selection shifts, and changes in market microstructure all induce nonstationarity. This is precisely where uncertainty quantification should matter most, yet many finance papers still report only static train/test metrics. Future work should treat temporal adaptation, online recalibration, prequential evaluation, and stress-period transfer as first-class research objects rather than robustness appendices [18, 23, 27, 44, 60–83, 96–98]. In practice, the most credible systems will likely combine multiple uncertainty signals: ensemble disagreement, interval coverage drift, density or embedding-shift diagnostics, and explicit scenario stress tests. 13

The second challenge is **tail scarcity**. Finance cares most about rare events, but these are exactly the observations that standard deep learning sees least often. Oversampling alone is not enough; it can distort dependence or encourage memorization. The literature on scenario generation and graph risk already points toward more structured solutions: generative models constrained by stress properties, dependence-aware validation, hybrid parametric–neural tail models, and risk-aware objectives that emphasize rare but consequential outcomes [37–45, 84–89]. A particularly promising direction is to merge extreme-value ideas and coherent-risk scoring with representation learning, so that the network learns not just from typical states, but from how losses concentrate in the tail.

The third challenge is **governance, interpretability, and responsible deployment**. In many financial institutions, the strongest barrier to adoption is not computational cost but model-risk policy. A model that is marginally more accurate yet poorly calibrated, opaque, or unstable under group shifts may be inferior in practice to a simpler scorecard or econometric benchmark [18, 23, 31, 92–98]. The literature increasingly recognizes this. Explainability tools such as LIME and SHAP can be useful for local or global auditing [92, 93], but the finance setting often requires more than post hoc explanation. It requires monotonicity constraints, stable variable effects, conservative uncertainty bands, documented validation workflows, and escalation rules for abstention or human override. This is why interpretable-by-design architectures, monotone networks, sparse attention, and uncertainty-aware decision thresholds deserve more attention than generic feature-attribution plots. Recent credit surveys rightly frame “trustworthy deep learning” as the next frontier rather than a secondary concern [31, 99–107]. 14

The fourth challenge is **evaluation mismatch**. AUC, accuracy, and even likelihood can be poor proxies for financial usefulness. A market-risk model should be judged by calibration, backtesting performance, and economic capital implications. A credit model should be judged by discrimination, calibration, reject-option quality, and robustness under portfolio drift. A portfolio learner should be judged by out-of-sample utility, turnover, downside risk, and failure behavior under crises. A scenario generator should be judged by dependence fidelity, novelty, tail realism, and its effect on downstream risk measures. The best prospect for the field is therefore methodological pluralism: no single metric, architecture, or uncertainty formalism will dominate. Instead, successful systems will be those that align architecture, uncertainty method, training objective, and evaluation protocol with the financial decision they support.

Finally, there is a major opportunity in **foundation-style and multimodal risk models**. Credit risk increasingly involves documents, statements, transactions, collateral images, customer support interactions, and network context; market risk increasingly involves text, options surfaces, order books, and cross-asset states. A natural next step is the integration of deep tabular models, sequence encoders, text transformers, graph modules, and calibrated decision layers into a unified risk stack. But the lesson of the current literature is clear: multimodality without calibrated uncertainty is not enough. The future of deep financial risk modeling will belong to systems that couple representation richness with credible uncertainty estimates and decision-aware validation.

## Conclusion

Deep learning has changed financial risk modeling most convincingly where the problem itself is rich in structure: high-dimensional borrower panels, long temporal trajectories, synthetic scenario generation, networked systemic risk, and control problems such as hedging or dynamic allocation. In these settings, deep models offer something classical methods struggle to match: the ability to learn nonlinear representations across modalities and to couple prediction with downstream risk-sensitive decision-making [16–18, 32–47, 84–91]. 15

At the same time, the survey supports a more cautious conclusion than simple “deep learning beats classical baselines” narratives. The most reliable progress comes when uncertainty is modeled explicitly, calibration is evaluated directly, and losses are aligned with financial objectives. When that alignment is absent, deep learning can easily become an overfit point forecaster in a domain where the tails matter most. The scientific frontier is therefore not merely larger neural networks for finance, but trustworthy deep risk systems that preserve calibration under nonstationarity, produce usable uncertainty estimates, and remain robust when embedded in real decisions.

## References

[1] Markowitz, H. Portfolio Selection. *The Journal of Finance* 1952, 7(1), 77–91.

[2] Engle, R.F. Autoregressive Conditional Heteroscedasticity with Estimates of the Variance of United Kingdom Inflation. *Econometrica* 1982, 50(4), 987–1007.

[3] Bollerslev, T. Generalized Autoregressive Conditional Heteroskedasticity. *Journal of Econometrics* 1986, 31(3), 307–327.

[4] Artzner, P.; Delbaen, F.; Eber, J.-M.; Heath, D. Coherent Measures of Risk. *Mathematical Finance* 1999, 9(3), 203–228.

[5] Rockafellar, R.T.; Uryasev, S. Optimization of Conditional Value-at-Risk. *Journal of Risk* 2000, 2(3), 21–41.

[6] Kupiec, P.H. Techniques for Verifying the Accuracy of Risk Measurement Models. *The Journal of Derivatives* 1995, 3(2), 73–84.

[7] Christoffersen, P.F. Evaluating Interval Forecasts. *International Economic Review* 1998, 39(4), 841–862.

[8] Koenker, R.; Bassett, G. Regression Quantiles. *Econometrica* 1978, 46(1), 33–50.

[9] Gneiting, T. Making and Evaluating Point Forecasts. *Journal of the American Statistical Association* 2011, 106(494), 746–762.

[10] Fissler, T.; Ziegel, J.F. Higher Order Elicitability and Osband's Principle. *The Annals of Statistics* 2016, 44(4), 1680–1707.

[11] Acerbi, C.; Székely, B. Back-Testing Expected Shortfall. *Risk* 2014, 27(11), 76–81.

[12] Khandani, A.E.; Kim, A.J.; Lo, A.W. Consumer Credit-Risk Models via Machine-Learning Algorithms. *Journal of Banking & Finance* 2010, 34(11), 2767–2787.

[13] Lessmann, S.; Baesens, B.; Seow, H.-V.; Thomas, L.C. Benchmarking State-of-the-Art Classification Algorithms for Credit Scoring: An Update of Research. *European Journal of Operational Research* 2015, 247(1), 124–136.

[14] Brown, I.; Mues, C. An Experimental Comparison of Classification Algorithms for Imbalanced Credit Scoring Data Sets. *Expert Systems with Applications* 2012, 39(3), 3446–3453.

[15] Baesens, B.; Setiono, R.; Mues, C.; Vanthienen, J. Using Neural Network Rule Extraction and Decision Tables for Credit-Risk Evaluation. *Management Science* 2003, 49(3), 312–329.

[16] Sadhwani, A.; Sirignano, J.; Giesecke, K. Deep Learning for Mortgage Risk. *Journal of Financial Econometrics* 2021, 19(2), 313–368. arXiv:1607.02470.

[17] Kvamme, H.; Sellereite, N.; Aas, K.; Sjursen, S. Predicting Mortgage Default Using Convolutional Neural Networks. *Expert Systems with Applications* 2018, 102, 207–217.

[18] Gunnarsson, B.R.; vanden Broucke, S.; Baesens, B.; Óskarsdóttir, M.; Lemahieu, Y. Deep Learning for Credit Scoring: Do or Don't? *European Journal of Operational Research* 2021, 295(1), 292–305.

[19] Mai, F.; Tian, S.; Lee, C.; Ma, L. Deep Learning Models for Bankruptcy Prediction Using Textual Disclosures. *European Journal of Operational Research* 2019, 274(2), 743–758.

[20] Ahmadi, Z.; Martens, P.; Koch, C.; Kremer, H.; Teuteberg, F. Towards Bankruptcy Prediction: Deep Sentiment Mining to Detect Financial Distress from Business Management Reports. In *Proceedings of the IEEE International Conference on Data Science and Advanced Analytics*; 2018; pp. 293–302.

[21] Ma, X.; Lv, S. Financial Credit Risk Prediction in Internet Finance Driven by Machine Learning. *Neural Computing and Applications* 2019, 31, 8359–8367.

[22] Mbuvha, R.; Boulkaibet, I.; Marwala, T. Automatic Relevance Determination Bayesian Neural Networks for Credit Card Default Modelling. *arXiv* 2019, arXiv:1906.06382.

[23] Nagl, M.; Nagl, M.; Rösch, D. Quantifying Uncertainty of Machine Learning Methods for Loss Given Default. *Frontiers in Applied Mathematics and Statistics* 2022, 8, 1076083.

[24] Dhiman, A. UQ for Credit Risk Management: A Deep Evidence Regression Approach. *arXiv* 2023, arXiv: 2305.04967.

[25] Heaton, J.B.; Polson, N.G.; Witte, J.H. Deep Learning in Finance. *Annual Review of Financial Economics* 2017, 11, 1–22.

[26] Mashrur, A.; Luo, W.; Zaidi, N.A.; Robles-Kelly, A. Machine Learning for Financial Risk Management: A Survey. *IEEE Access* 2020, 8, 203203–203223.

[27] Sezer, O.B.; Gudelek, M.U.; Ozbayoglu, A.M. Financial Time Series Forecasting with Deep Learning: A Systematic Literature Review. *Applied Soft Computing* 2020, 90, 106181.

[28] Ozbayoglu, A.M.; Gudelek, M.U.; Sezer, O.B. Deep Learning for Financial Applications: A Survey. *Applied Soft Computing* 2020, 93, 106384.

[29] Peng, K.; Yan, G. A Survey on Deep Learning for Financial Risk Prediction. *Quantitative Finance and Economics* 2021, 5(4), 716–737.

[30] Mienye, E.; Jere, N.; Obaido, G.; Mienye, I.D.; Aruleba, K. Deep Learning in Finance: A Survey of Applications and Techniques. *AI* 2024, 5, 2066–2091.

[31] Mienye, I.D.; Esenogho, E.; Modisane, C. Deep Learning for Credit Risk Prediction: A Survey of Methods, Applications, and Challenges. *Information* 2026, 17(4), 395.

[32] Fischer, T.; Krauss, C. Deep Learning with Long Short-Term Memory Networks for Financial Market Predictions. *European Journal of Operational Research* 2018, 270(2), 654–669.

[33] Sirignano, J.; Cont, R. Universal Features of Price Formation in Financial Markets: Perspectives from Deep Learning. *Quantitative Finance* 2019, 19(9), 1449–1459.

[34] Zhang, Z.; Zohren, S.; Roberts, S. DeepLOB: Deep Convolutional Neural Networks for Limit Order Books. *IEEE Transactions on Signal Processing* 2020, 68, 3001–3012. arXiv:1808.03668.

[35] Lim, B.; Arik, S.O.; Loeff, N.; Pfister, T. Temporal Fusion Transformers for Interpretable Multi-Horizon Time Series Forecasting. *International Journal of Forecasting* 2021, 37(4), 1748–1764. arXiv:1912.09363.

[36] Vaswani, A.; Shazeer, N.; Parmar, N.; Uszkoreit, J.; Jones, L.; Gomez, A.N.; Kaiser, Ł.; Polosukhin, I. Attention Is All You Need. In *Advances in Neural Information Processing Systems*; 2017; Vol. 30.

[37] Goodfellow, I.; Pouget-Abadie, J.; Mirza, M.; Xu, B.; Warde-Farley, D.; Ozair, S.; Courville, A.; Bengio, Y. Generative Adversarial Nets. In *Advances in Neural Information Processing Systems*; 2014; Vol. 27.

[38] Kingma, D.P.; Welling, M. Auto-Encoding Variational Bayes. In *International Conference on Learning Representations*; 2014. arXiv:1312.6114.

[39] Yoon, J.; Jarrett, D.; van der Schaar, M. Time-Series Generative Adversarial Networks. In *Advances in Neural Information Processing Systems*; 2019; Vol. 32.

[40] Wiese, M.; Knobloch, R.; Korn, R.; Kretschmer, P. Quant GANs: Deep Generation of Financial Time Series. *Quantitative Finance* 2020, 20(9), 1419–1440. arXiv:1907.06673.

[41] Arimond, A.; Borth, D.; Hoepner, A.; Klawunn, M.; Weisheit, S. Neural Networks and Value at Risk. *arXiv* 2020, arXiv:2005.01686.

[42] Arian, H.; Moghimi, M.; Tabatabaei, E.; Zamani, S. Encoded Value-at-Risk: A Predictive Machine for Financial Risk Management. *arXiv* 2020, arXiv:2011.06742.

[43] Flaig, S.; Junike, G. Scenario Generation for Market Risk Models Using Generative Neural Networks. *arXiv* 2022, arXiv:2109.10072.

[44] Junike, G.; Flaig, S.; Werner, R. Validation of Machine Learning Based Scenario Generators. *Journal of Risk and Insurance* 2026, 93(2), 508–533. arXiv:2301.12719.

[45] Qiu, Z.; Lazar, E.; Nakata, K. VaR and ES Forecasting via Recurrent Neural Network-Based Stateful Models. *International Review of Financial Analysis* 2024, 92, 103102.

[46] Kato, M. Conformal Predictive Portfolio Selection. *arXiv* 2025, arXiv:2410.16333.

[47] Bühler, H.; Gonon, L.; Teichmann, J.; Wood, B. Deep Hedging. *arXiv* 2018, arXiv:1802.03042.

[48] Halperin, I. QLBS: Q-Learner in the Black-Scholes(-Merton) Worlds. *arXiv* 2017, arXiv:1712.04609.

[49] Jiang, Z.; Xu, D.; Liang, J. A Deep Reinforcement Learning Framework for the Financial Portfolio Management Problem. *arXiv* 2017, arXiv:1706.10059.

[50] Liang, Z.; Chen, H.; Zhu, J.; Jiang, K.; Li, Y. Adversarial Deep Reinforcement Learning in Portfolio Management. *arXiv* 2018, arXiv:1808.09940.

[51] Moody, J.; Saffell, M. Learning to Trade via Direct Reinforcement. *IEEE Transactions on Neural Networks* 2001, 12(4), 875–889.

[52] Tamar, A.; Glassner, Y.; Mannor, S. Policy Gradients Beyond Expectations: Conditional Value-at-Risk in Markov Decision Processes. In *AAAI Conference on Artificial Intelligence*; 2015.

[53] Tamar, A.; Chow, Y.; Ghavamzadeh, M.; Mannor, S. Policy Gradient for Coherent Risk Measures. In *Advances in Neural Information Processing Systems*; 2015; Vol. 28.

[54] Chow, Y.; Ghavamzadeh, M. Algorithms for CVaR Optimization in MDPs. In *Advances in Neural Information Processing Systems*; 2014; Vol. 27.

[55] Chow, Y.; Tamar, A.; Mannor, S.; Pavone, M. Risk-Sensitive and Robust Decision-Making: A CVaR Optimization Approach. In *Advances in Neural Information Processing Systems*; 2015; Vol. 28.

[56] Bellemare, M.G.; Dabney, W.; Munos, R. A Distributional Perspective on Reinforcement Learning. In *International Conference on Machine Learning*; 2017; pp. 449–458.

[57] Dabney, W.; Rowland, M.; Bellemare, M.G.; Munos, R. Distributional Reinforcement Learning with Quantile Regression. In *AAAI Conference on Artificial Intelligence*; 2018.

[58] Ni, X.; Lai, L. Robust Risk-Sensitive Reinforcement Learning with Conditional Value-at-Risk. *arXiv* 2024, arXiv:2405.01718.

[59] Poddar, M. Uncertainty-Aware Deep Hedging. *arXiv* 2026, arXiv:2603.10137.

[60] Hüllermeier, E.; Waegeman, W. Aleatoric and Epistemic Uncertainty in Machine Learning: An Introduction to Concepts and Methods. *Machine Learning* 2021, 110, 457–506.

[61] Kendall, A.; Gal, Y. What Uncertainties Do We Need in Bayesian Deep Learning for Computer Vision? In *Advances in Neural Information Processing Systems*; 2017; Vol. 30. arXiv:1703.04977.

[62] Gal, Y.; Ghahramani, Z. Dropout as a Bayesian Approximation: Representing Model Uncertainty in Deep Learning. In *International Conference on Machine Learning*; 2016; pp. 1050–1059. arXiv:1506.02142.

[63] Blundell, C.; Cornebise, J.; Kavukcuoglu, K.; Wierstra, D. Weight Uncertainty in Neural Networks. In *International Conference on Machine Learning*; 2015; pp. 1613–1622. arXiv:1505.05424.

[64] Hernández-Lobato, J.M.; Adams, R. Probabilistic Backpropagation for Scalable Learning of Bayesian Neural Networks. In *International Conference on Machine Learning*; 2015; pp. 1861–1869. arXiv:1502.05336.

[65] Lakshminarayanan, B.; Pritzel, A.; Blundell, C. Simple and Scalable Predictive Uncertainty Estimation Using Deep Ensembles. In *Advances in Neural Information Processing Systems*; 2017; Vol. 30. arXiv:1612.01474.

[66] Ovadia, Y.; Fertig, E.; Ren, J.; Nado, Z.; Sculley, D.; Nowozin, S.; Dillon, J.; Lakshminarayanan, B.; Snoek, J. Can You Trust Your Model's Uncertainty? Evaluating Predictive Uncertainty under Dataset Shift. In *Advances in Neural Information Processing Systems*; 2019; Vol. 32. arXiv:1906.02530.

[67] Guo, C.; Pleiss, G.; Sun, Y.; Weinberger, K.Q. On Calibration of Modern Neural Networks. In *International Conference on Machine Learning*; 2017; pp. 1321–1330. arXiv:1706.04599.

[68] Kuleshov, V.; Fenner, N.; Ermon, S. Accurate Uncertainties for Deep Learning Using Calibrated Regression. In *International Conference on Machine Learning*; 2018; pp. 2796–2804. arXiv:1807.00263.

[69] Sensoy, M.; Kaplan, L.; Kandemir, M. Evidential Deep Learning to Quantify Classification Uncertainty. In *Advances in Neural Information Processing Systems*; 2018; Vol. 31. arXiv:1806.01768.

[70] Malinin, A.; Gales, M. Predictive Uncertainty Estimation via Prior Networks. In *Advances in Neural Information Processing Systems*; 2018; Vol. 31. arXiv:1802.10501.

[71] Amini, A.; Schwarting, W.; Soleimany, A.; Rus, D. Deep Evidential Regression. In *Advances in Neural Information Processing Systems*; 2020; Vol. 33. arXiv:1910.02600.

[72] Tagasovska, N.; Lopez-Paz, D. Single-Model Uncertainties for Deep Learning. In *Advances in Neural Information Processing Systems*; 2019; Vol. 32. arXiv:1811.00908.

[73] Pearce, T.; Zaki, M.; Brintrup, A.; Neely, A. High-Quality Prediction Intervals for Deep Learning: A Distribution-Free, Ensembled Approach. In *International Conference on Machine Learning*; 2018 Workshop. arXiv:1802.07167.

[74] Abdar, M.; Samami, M.; Dehghani, M.; Khosravi, A.; Nahavandi, S.; et al. A Review of Uncertainty Quantification in Deep Learning: Techniques, Applications and Challenges. *Information Fusion* 2021, 76, 243–297.

[75] Gawlikowski, J.; Tassi, C.; Ali, M.; Lee, J.; Humt, M.; Feng, J.; Kruspe, A.; Triebel, R.; Jung, P.; Roscher, R.; et al. A Survey of Uncertainty in Deep Neural Networks. *Artificial Intelligence Review* 2023, 56, 1513–1589.

[76] Yao, J.; Pan, W.; Ghosh, S.; Doshi-Velez, F. Quality of Uncertainty Quantification for Bayesian Neural Network Inference. In *International Conference on Machine Learning*; 2019; pp. 6993–7002. arXiv:1906.09686.

[77] Caprio, M.; Dutta, S.; Müller, S.; Simkus, T.; Pfreundt, F.-J.; Schütze, H. Credal Bayesian Deep Learning. In *Uncertainty in Artificial Intelligence*; 2023. arXiv:2302.09656.

[78] Shafer, G.; Vovk, V. A Tutorial on Conformal Prediction. *Journal of Machine Learning Research* 2008, 9, 371–421.

[79] Romano, Y.; Patterson, E.; Candès, E. Conformalized Quantile Regression. In *Advances in Neural Information Processing Systems*; 2019; Vol. 32. arXiv:1905.03222.

[80] Barber, R.F.; Candès, E.J.; Ramdas, A.; Tibshirani, R.J. Predictive Inference with the Jackknife+. *The Annals of Statistics* 2021, 49(1), 486–507.

[81] Gibbs, I.; Candès, E. Adaptive Conformal Inference under Distribution Shift. In *Advances in Neural Information Processing Systems*; 2021; Vol. 34. arXiv:2106.00170.

[82] Angelopoulos, A.N.; Bates, S. Conformal Prediction: A Gentle Introduction. *Foundations and Trends in Machine Learning* 2023, 16(4), 494–591. arXiv:2107.07511.

[83] Chung, Y.; Neiswanger, W.; Char, I.; Schneider, J.; Póczos, B. Beyond Pinball Loss: Quantile Methods for Calibrated Uncertainty Quantification. In *Advances in Neural Information Processing Systems*; 2021; Vol. 34. arXiv:2011.09588.

[84] Eisenberg, L.; Noe, T.H. Systemic Risk in Financial Systems. *Management Science* 2001, 47(2), 236–249.

[85] Battiston, S.; Puliga, M.; Kaushik, R.; Tasca, P.; Caldarelli, G. DebtRank: Too Central to Fail? Financial Networks, the FED and Systemic Risk. *Scientific Reports* 2012, 2, 541.

[86] Glasserman, P.; Young, H.P. Contagion in Financial Networks. *Journal of Economic Literature* 2016, 54(3), 779–831.

[87] Biagini, F.; Fouque, J.-P.; Fritelli, M.; Meyer-Brandis, T. A Unifying Framework for Systemic Risk Measures. *SIAM Journal on Financial Mathematics* 2019, 10(2), 435–464.

[88] Balmaseda, V.; Coronado, M.; de Cadenas-Santiago, G. Predicting Systemic Risk in Financial Systems Using Deep Graph Learning. *Intelligent Systems with Applications* 2023, 19, 200240.

[89] Gonon, L.; Meyer-Brandis, T.; Weber, N. Computing Systemic Risk Measures with Graph Neural Networks. *arXiv* 2024, arXiv:2410.07222.

[90] Liu, B.; et al. Unveiling the Potential of Graph Neural Networks in SME Credit Risk Prediction. *arXiv* 2024, arXiv:2409.17909.

[91] Sun, M.; Sun, W.; Sun, Y.; Liu, S.; Jiang, M.; Xu, Z. Applying Hybrid Graph Neural Networks to Strengthen Credit Risk Analysis. *arXiv* 2024, arXiv:2410.04283.

[92] Ribeiro, M.T.; Singh, S.; Guestrin, C. "Why Should I Trust You?" Explaining the Predictions of Any Classifier. In *ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*; 2016; pp. 1135–1144. arXiv:1602.04938.

[93] Lundberg, S.M.; Lee, S.-I. A Unified Approach to Interpreting Model Predictions. In *Advances in Neural Information Processing Systems*; 2017; Vol. 30. arXiv:1705.07874.

[94] Doshi-Velez, F.; Kim, B. Towards a Rigorous Science of Interpretable Machine Learning. *arXiv* 2017, arXiv: 1702.08608.

[95] Rudin, C. Stop Explaining Black Box Machine Learning Models for High Stakes Decisions and Use Interpretable Models Instead. *Nature Machine Intelligence* 2019, 1, 206–215.

[96] Sagawa, S.; Koh, P.W.; Hashimoto, T.B.; Liang, P. Distributionally Robust Neural Networks for Group Shifts: On the Importance of Regularization for Worst-Case Generalization. In *International Conference on Learning Representations*; 2020. arXiv:1911.08731.

[97] Koh, P.W.; Sagawa, S.; Marklund, H.; Xie, S.M.; Zhang, M.; Balsubramani, A.; Hu, W.; Yasunaga, M.; Liang, P. WILDS: A Benchmark of in-the-Wild Distribution Shifts. In *International Conference on Machine Learning*; 2021; pp. 5637–5664. arXiv:2012.07421.

[98] Wilson, A.G.; Izmailov, P. Bayesian Deep Learning and a Probabilistic Perspective of Generalization. In *Advances in Neural Information Processing Systems*; 2020; Vol. 33. arXiv:2002.08791.

[99] Hayashi, Y. Emerging Trends in Deep Learning for Credit Scoring: A Review. *Electronics* 2022, 11, 3181.

[100] Shi, S.; Tse, R.; Luo, W.; D'Addona, S.; Pau, G. Machine Learning-Driven Credit Risk: A Systemic Review. *Neural Computing and Applications* 2022, 34, 14327–14339.

[101] Noriega, J.P.; Rivera, L.A.; Herrera, J.A. Machine Learning for Credit Risk Prediction: A Systematic Literature Review. *Data* 2023, 8, 169.

[102] Montevechi, A.A.; de Carvalho Miranda, R.; Medeiros, A.L.; Montevechi, J.A.B. Advancing Credit Risk Modelling with Machine Learning: A Comprehensive Review of the State of the Art. *Engineering Applications of Artificial Intelligence* 2024, 137, 109082.

[103] Bhatore, S.; Mohan, L.; Reddy, Y.R. Machine Learning Techniques for Credit Risk Evaluation: A Systematic Literature Review. *Journal of Banking and Financial Technology* 2020, 4, 111–138.

[104] Çallı, B.A.; Coşkun, E. A Longitudinal Systematic Review of Credit Risk Assessment and Credit Default Predictors. *SAGE Open* 2021, 11, 21582440211061333.

[105] Mhlanga, D. Financial Inclusion in Emerging Economies: The Application of Machine Learning and Artificial Intelligence in Credit Risk Assessment. *International Journal of Financial Studies* 2021, 9, 39.

[106] Hoyos Gutiérrez, S.P.; Santos López, F.M. Credit Risk Assessment System Based on Deep Learning: A Systematic Literature Review. In *International Conference on Computer Science, Electronics and Industrial Engineering*; Springer, 2023; pp. 395–413.

[107] Demma Wube, H.; Zekarias Esubalew, S.; Fayiso Weldesellasie, F.; Girma Debelee, T. Deep Learning and Machine Learning Techniques for Credit Scoring: A Review. In *Pan African Conference on Artificial Intelligence*; Springer, 2024; pp. 30–61.

[108] Sonkavde, G.; Dharrao, D.S.; Bongale, A.M.; Deokate, S.T.; Doreswamy, D.; Bhat, S.K. Forecasting Stock Market Prices Using Machine Learning and Deep Learning Models: A Systematic Review, Performance Analysis and Discussion of Implications. *International Journal of Financial Studies* 2023, 11, 94.

[109] Pricope, T.V. Deep Reinforcement Learning in Quantitative Algorithmic Trading: A Review. *arXiv* 2021, arXiv:2106.00123.

[110] Zhou, Q.; et al. Portfolio Optimization with Robust Covariance and Conditional Value-at-Risk Constraints. *arXiv* 2024, arXiv:2406.00610.

---

[1] [2] From Deep Learning to LLMs: A survey of AI in Quantitative Investment
https://arxiv.org/html/2503.21422v1

[3] [7] [2305.04967] UQ for Credit Risk Management: A deep evidence regression approach
https://arxiv.org/abs/2305.04967

[4] [2109.10072] Scenario generation for market risk models using generative neural networks
https://arxiv.org/abs/2109.10072

[5] VaR and ES forecasting via recurrent neural network-based stateful models - ScienceDirect
https://www.sciencedirect.com/science/article/pii/S1057521924000346

[6] [10] [14] Deep Learning for Credit Risk Prediction: A Survey of Methods, Applications, and Challenges | MDPI
https://www.mdpi.com/2078-2489/17/4/395

[8] Conformal Predictive Portfolio Selection
https://arxiv.org/html/2410.16333v2

[9] [2005.01686] Neural Networks and Value at Risk
https://arxiv.org/abs/2005.01686

[11] Predicting systemic risk in financial systems using Deep Graph Learning - ScienceDirect
https://www.sciencedirect.com/science/article/pii/S2667305323000650

[12] Deep Hedging
https://arxiv.org/abs/1802.03042?utm_source=chatgpt.com

[13] Validation of machine learning based scenario generators
https://arxiv.org/html/2301.12719v4

[15] Deep Learning for Mortgage Risk
https://arxiv.org/abs/1607.02470?utm_source=chatgpt.com