

A Survey of AI Software Supply-Chain Security and Vulnerability Detection

InteractiveSurvey

1 Introduction

The rapid integration of artificial intelligence into software development lifecycles has fundamentally transformed how applications are built, deployed, and maintained. While these advancements accelerate innovation, they have simultaneously expanded the attack surface, introducing complex vulnerabilities within the software supply chain [1]. Traditional static analysis tools, often reliant on rigid rule-based systems, struggle to detect subtle semantic flaws or zero-day exploits hidden within vast, interconnected codebases. As software ecosystems grow increasingly dependent on third-party libraries and automated generation tools, the need for intelligent, adaptive security mechanisms has become paramount. Consequently, the intersection of AI and cybersecurity has emerged as a critical frontier, demanding sophisticated approaches to identify and mitigate risks before they compromise global digital infrastructure.

This survey paper focuses on the evolving landscape of AI-driven code vulnerability detection, specifically examining how modern machine learning architectures address the challenges of software supply-chain security [2]. Unlike previous reviews that broadly cover software testing, this work provides a targeted analysis of deep learning models designed to understand code semantics, structural dependencies, and contextual nuances. The primary objective is to synthesize recent advancements in neural network architectures, learning paradigms, and large language model applications that enhance the precision and robustness of vulnerability identification [3]. By concentrating on techniques that overcome data scarcity, class imbalance, and adversarial evasion, this study aims to clarify the current state-of-the-art and identify gaps hindering the practical deployment of these technologies in real-world de-

velopment environments.

The core of this review begins with a comprehensive examination of deep learning architectures tailored for semantic code representation. We analyze Graph Neural Networks for their unique capacity to model non-local structural dependencies through code property graphs, contrasting their topological awareness with the limitations of sequence-based methods. Simultaneously, the discussion explores Transformer-based encoders, which leverage self-attention mechanisms to capture long-range contextual relationships within source code, outperforming traditional embedding techniques. Furthermore, the enduring relevance of Recurrent Neural Networks, particularly LSTM and GRU variants, is evaluated for their effectiveness in recognizing sequential patterns and temporal dynamics in code streams, providing a foundational understanding of how different architectures interpret program logic.

Following the architectural overview, the survey delves into advanced learning paradigms and data optimization strategies essential for training effective detection models. Special attention is given to self-supervised and contrastive learning techniques that exploit vast unlabeled corpora to learn robust representations without expensive manual annotation. The review critically assesses methods for addressing the severe class imbalance inherent in vulnerability datasets, including resampling strategies, synthetic data generation, and specialized evaluation metrics that offer more accurate performance insights than standard accuracy scores. Additionally, the role of adversarial training is explored, highlighting how injecting perturbed inputs and hard negative samples fortifies models against evasive attacks and ensures generalization across diverse coding styles and obfuscation techniques [4].

The final technical section investigates the transformative impact of Large Language Models on




Figure 1: The outline of our survey: A Survey of AI Software Supply-Chain Security and Vulnerability Detection






Figure 2: Chart from 'Data and Context Matter: Towards Generalizing AI-based Software Vulnerability Detection' (arXiv:2508.16625)



((a)) Confusion matrix for Clean Data (QNN)



((b)) Confusion matrix for Perturbed Data (QNN)

Figure 3: Chart from 'Exploring the Vulnerabilities of Machine Learning and Quantum Machine Learning to Adversarial Attacks using a Malware Dataset: A Comparative Analysis' (arXiv:2305.19593)

vulnerability identification, shifting the paradigm from extensive fine-tuning to efficient prompt engineering and in-context learning [5]. This part analyzes how carefully crafted prompts and few-shot examples enable pre-trained models to detect security flaws with minimal computational overhead, leveraging their inherent understanding of code syntax and semantics [3]. By synthesizing insights from architectural designs, data optimization, and generative AI applications, the survey constructs a holistic view of the current technological landscape, setting the stage for a discussion on future research directions and the practical challenges of integrating these AI-driven solutions into continuous integration pipelines.




Figure 4: Chart from 'SHERLOCK: A Deep Learning Approach To Detect Software Vulnerabilities' (arXiv:2512.12593)

The primary contribution of this survey lies in its systematic categorization and critical analysis of AI techniques specifically applied to software supply-chain security and vulnerability detection. By bridging the gap between theoretical deep learning advancements and practical cybersecu-

rity applications, this work offers a unified framework for understanding how structural, sequential, and contextual code features are leveraged for threat mitigation. It provides researchers and practitioners with a consolidated reference on overcoming data scarcity and adversarial robustness, while also highlighting the emerging potential of large language models in automating security audits. Ultimately, this paper serves as a foundational guide for developing next-generation static analysis tools capable of securing modern, complex software ecosystems against evolving threats.

2 AI-Driven Code Vulnerability Detection

2.1 Deep Learning Architectures for Semantic Code Representation

2.1.1 Graph Neural Networks for Structural Dependency Modeling

Graph Neural Networks (GNNs) have emerged as a pivotal architecture for modeling structural dependencies within software systems, addressing the limitations of sequence-based models in capturing non-local code relationships. By representing source code as rich graph structures—such as Abstract Syntax Trees, Control Flow Graphs, or Code Property Graphs—GNNs effectively encode the intricate topological connections between functions, classes, and variables [6]. This representation allows the model to propagate information across nodes via message-passing mechanisms, thereby learning comprehensive program semantics that are critical for tasks like vulnerability detection and dependency alert classification.

The efficacy of GNNs in this domain stems from their ability to aggregate features from neighboring nodes while preserving the hierarchical and relational nature of code. Unlike traditional token-based approaches that often overlook long-range dependencies due to fixed context windows, graph-based models explicitly map direct and transitive dependencies, such as package reachability and function call chains. Variants including Graph Convolutional Networks, Graph Attention Networks, and Heterogeneous GNNs have been employed to handle diverse node types and edge attributes, enabling the precise identification of structural anomalies and silent dependency alerts within complex software ecosystems.

Despite their superior performance in capturing structural nuances, the deployment of GNNs



Figure 5: Chart from 'AI-Based Software Vulnerability Detection: A Systematic Literature Review' (arXiv:2506.10280)

presents challenges related to computational overhead and graph construction complexity. Generating accurate code property graphs requires sophisticated parsing and can be resource-intensive, potentially hindering integration into continuous integration pipelines or resource-constrained environments. Consequently, while GNNs currently represent the state-of-the-art for structural dependency modeling by outperforming recurrent and convolutional baselines, ongoing research focuses on optimizing graph construction efficiency and exploring lightweight architectures that retain semantic fidelity without prohibitive training costs.

2.1.2 Transformer-Based Encoders for Contextual Token Analysis

Transformer-based encoders have emerged as the dominant paradigm for contextual token analysis in source code, surpassing static embedding techniques like Word2Vec and n-gram features. Unlike earlier methods that treat code as simple token sequences or rely on sparse binary representations, these models leverage self-attention mechanisms to capture long-range dependencies and complex syntactic structures inherent in programming languages. By mapping tokens into dense, continuous vector spaces, encoder-only architectures effectively preserve both semantic meaning and structural relationships, enabling a nuanced under-

standing of code fragments that traditional bag-of-words or shallow neural network approaches fail to achieve.

Recent advancements prioritize pre-trained models specifically adapted for software engineering tasks, such as CodeBERT and UniXCoder, which are trained on vast corpora of source code to learn robust code representations. These encoder-based models significantly outperform decoder-only or encoder-decoder variants in vulnerability detection and code understanding tasks, as they are optimized for bidirectional context aggregation rather than autoregressive generation. Empirical studies indicate that architectures capable of handling extended context windows, such as those supporting 1024-token sequences, yield superior performance metrics, including higher F1 scores and recall rates, by comprehensively analyzing the global context surrounding potential security flaws.

The efficacy of these encoders is further enhanced when integrated with specialized attention layers that calculate the importance of specific tokens relative to the entire code structure [7]. This capability allows the model to distinguish between critical variables, operators, and keywords while ignoring irrelevant noise, a distinct advantage over methods that average token embeddings and consequently blur semantic details. As the field evolves, the focus remains on refining these transformer encoders to manage diverse programming languages and larger codebases, establishing them as essential components for next-generation static analysis tools and automated vulnerability mitigation systems.

2.1.3 Recurrent Neural Networks for Sequential Pattern Recognition

Recurrent Neural Networks (RNNs) constitute a fundamental architecture for sequential pattern recognition, specifically designed to capture long-term dependencies within temporal data streams. Among these, Long Short-Term Memory (LSTM) networks have emerged as the predominant variant, utilizing a sophisticated gating mechanism comprising input, forget, and output gates to regulate information flow [7]. This structural design effectively mitigates the vanishing gradient problem inherent in traditional RNNs, enabling the model to retain critical context over extended sequences. Consequently, LSTMs have become ubiquitous in domains requiring deep contextual understanding,




Figure 6: Chart from 'Silent Vulnerable Dependency Alert Prediction with Vulnerability Key Aspect Explanation' (arXiv:2302.07445)

such as acoustic modeling, handwritten recognition, and machine translation, where maintaining historical state is paramount for accurate prediction.

To further enhance contextual acquisition, Bidirectional LSTM (BiLSTM) architectures process input sequences in both forward and backward directions through separate hidden layers [8]. This dual-pass approach allows the network to integrate past and future context simultaneously, significantly improving performance in tasks like source code vulnerability detection and natural language understanding. Complementing this, the Gated Recurrent Unit (GRU) offers a computationally efficient alternative by merging the forget and input gates into a single update gate while retaining a reset gate [9]. Although simplified, GRUs often achieve comparable accuracy to LSTMs with reduced parameter counts, making them advantageous for resource-constrained environments or scenarios involving massive datasets where training efficiency is critical.

Empirical studies across software engineering and cybersecurity demonstrate that these recurrent models, particularly when augmented with attention mechanisms and static word embeddings like Word2Vec, yield optimal results for identifying complex sequential patterns. Research indicates that BiLSTM and attention-enhanced GRU con-




Figure 7: Chart from 'SafePyScript: A Web-Based Solution for Machine Learning-Driven Vulnerability Detection in Python' (arXiv:2411.00636)

figurations frequently outperform standard unidirectional variants in detecting security vulnerabilities within source code. Despite the recent ascendancy of transformer-based models, recurrent architectures remain highly relevant for specific sequential tasks due to their inherent ability to model temporal dynamics explicitly. Their continued application underscores their robustness in handling variable-length inputs and capturing intricate sequential dependencies that define modern pattern recognition challenges.

2.2 Advanced Learning Paradigms and Data Optimization

2.2.1 Self-Supervised and Contrastive Learning Strategies

Self-supervised learning has emerged as a pivotal paradigm for leveraging vast amounts of unlabeled data, particularly in domains where manual annotation is prohibitively expensive. By formulating pretext tasks such as masked token prediction or sequence reconstruction, models learn robust semantic representations without explicit class labels. This approach allows architectures, including transformers and graph neural networks, to capture intrinsic structural and contextual patterns within the data, establishing a strong foundation for downstream tasks like vulnerability detection or code classification.

Complementing these generative objectives, contrastive learning strategies refine feature spaces by maximizing the agreement between differently augmented views of the same sample while pushing apart representations of distinct instances. Through metric learning objectives, these methods enforce a geometry where semantically similar in-

puts cluster tightly, effectively separating classes even in the absence of labeled boundaries. This mechanism is particularly effective for handling noisy labels and improving generalization across diverse datasets, as it relies on relative similarity rather than absolute categorical assignments.

The integration of self-supervised pre-training with contrastive fine-tuning has demonstrated superior efficacy compared to traditional supervised baselines, especially when labeled data is scarce. By initializing models with weights derived from large-scale unlabeled corpora, the subsequent learning process converges faster and achieves higher accuracy on specialized tasks. Furthermore, these strategies enhance model robustness against adversarial perturbations and distribution shifts, making them essential components in modern machine learning pipelines designed for complex, real-world applications [4].

2.2.2 Techniques for Addressing Class Imbalance and Data Scarcity

The inherent scarcity of labeled vulnerable commits compared to the abundance of non-vulnerable code creates severe class imbalance, often leading to models that are biased toward the majority class and yield inflated accuracy metrics. To mitigate this, researchers frequently employ resampling strategies, including random down-sampling of non-vulnerable data to match the volume of vulnerable instances or stratified sampling to ensure representative distributions across training and testing sets. While these methods facilitate binary classifier training, aggressive down-sampling risks discarding valuable contextual information from the majority class, potentially limiting the model's ability to generalize to real-world repository conditions where imbalance is extreme.

Beyond resampling, advanced algorithmic adjustments and synthetic data generation have emerged as critical techniques for enhancing detection capabilities in data-scarce environments. Approaches such as sample weighting during the pre-processing phase assign higher penalties to misclassifications of the minority class, effectively forcing the model to focus on vulnerable patterns without altering the dataset size. Furthermore, deep learning frameworks integrated with fuzzy oversampling or text generation mechanisms, such as those leveraging self-admitted technical debt patterns, synthesize realistic vulnerable samples to augment training data. These generative meth-

ods aim to capture complex control and data dependencies that traditional static analysis often misses, thereby improving recall without significantly compromising precision.

Evaluation methodologies must also evolve to accurately reflect performance under imbalanced conditions, as standard accuracy metrics are often misleading. Consequently, the field has shifted toward using Area Under the Precision-Recall Curve (PR-AUC) and Recall at specific triage budgets, which provide more informative assessments than ROC-AUC in highly skewed distributions. Despite these advancements, challenges remain regarding label noise, data duplication, and the replicability of results on diverse, real-world datasets. Future directions emphasize the need for calibrated decision thresholds and robust validation on held-out target domains to ensure that vulnerability detection models maintain efficacy when deployed in production environments with naturally occurring data scarcity.

2.2.3 Adversarial Training and Robustness Enhancement

Adversarial training has emerged as a pivotal strategy for fortifying software vulnerability detection models against evasive attacks and distributional shifts [4]. By intentionally injecting perturbed inputs, such as obfuscated code or semantically equivalent variants with renamed identifiers, during the training phase, models learn to recognize invariant vulnerability patterns rather than superficial syntactic features. This process mitigates the risk of overfitting to specific coding styles or dataset artifacts, thereby enhancing the classifier's ability to generalize across diverse, unseen codebases where traditional static analysis tools often fail due to context limitations.

A critical component of robustness enhancement involves the strategic incorporation of hard negative samples into the training curriculum. These samples consist of secure code snippets that are semantically similar to vulnerable instances, forcing the model to discern fine-grained distinctions rather than relying on coarse-grained heuristics. Furthermore, addressing severe class imbalance through adversarial generation ensures that the decision boundary is not biased toward the majority class of secure code. Techniques such as generative adversarial networks tailored for sequential data have been employed to synthesize realistic vulnerability patches, providing a richer

ground truth for learning complex repair strategies and improving detection accuracy in imbalanced scenarios [10].




Figure 8: Chart from 'Vulnerability Detection with Fine-Grained Interpretations' (arXiv:2106.10478)

Despite these advancements, significant challenges remain regarding the stability of token-based features under adversarial transformations. While deep neural networks demonstrate superior resilience compared to shallow models by autonomously learning high-level representations, they remain susceptible to sophisticated evasion techniques that exploit semantic ambiguities. Future directions emphasize the integration of explainable AI to interpret model decisions and the development of cross-platform detection mechanisms that maintain robustness against rapidly evolving malicious behaviors. Ultimately, achieving true robustness requires a holistic approach that combines data augmentation, transfer learning, and continuous fine-tuning on diverse, real-world software repositories to ensure reliable vulnerability identification in dynamic development environments.

2.3 Large Language Models in Vulnerability Identification

2.3.1 Prompt Engineering and In-Context Learning Approaches

Prompt engineering and in-context learning have emerged as pivotal strategies for leveraging large language models in vulnerability detection without extensive fine-tuning [11]. By carefully crafting input prompts, researchers can guide pre-trained models to identify security flaws through few-shot examples embedded directly within the context window. This approach mitigates the computational intensity associated with training massive architectures like BERT or CodeBERT from scratch, offering a resource-efficient alternative that capitalizes on the model's inherent semantic understanding of code syntax and structure.

However, the efficacy of these methods is intrinsically bound by the limitations of the input context window, which often forces the truncation of critical long-range dependencies essential for accurate vulnerability identification. When the code




Figure 9: Chart from 'VulEval: Towards Repository-Level Evaluation of Software Vulnerability Detection' (arXiv:2404.15596)

snippet exceeds the token limit, vital contextual information regarding data flow or control dependencies may be omitted, leading to false negatives or superficial analysis. Consequently, optimizing prompt design involves a delicate compromise between including sufficient surrounding context for granular detection and adhering to strict length constraints, necessitating advanced techniques for context selection and compression.

To address these granularity challenges, recent advancements propose iterative, reflection-based frameworks where the model dynamically refines its understanding through multiple interaction rounds rather than relying on single-step queries. These systems employ distinct roles to expand context, infer code logic, and validate findings, effectively simulating a deeper analytical process that mimics human expert reasoning. Such iterative in-context learning approaches significantly enhance the model's ability to discern complex, context-sensitive vulnerabilities, marking a shift towards more robust and explainable automated security analysis tools.

2.3.2 Fine-Tuning Strategies for Domain-Specific Security Tasks

Most existing vulnerability detection models operate primarily at the function level, creating a significant gap in fine-grained analysis required for precise remediation. While recent efforts have begun to explore statement or line-level identification, the majority of approaches struggle to transfer vulnerability signatures effectively across diverse software domains [8]. This limitation stems from the reliance on heavy representations like code graphs that, although powerful for C/C++ analysis, often lack the adaptability needed for heterogeneous environments where security contexts vary significantly between projects and programming languages.

To address these challenges, advanced fine-

tuning strategies leverage transformer-based architectures pre-trained on source code to capture domain-specific semantic nuances. These methods move beyond surface-level token matching by incorporating fine-grained interpretations that pinpoint exact vulnerable statements, thereby reducing false positives and aiding developer triage. By adapting pre-trained models through specialized objective functions focused on security patterns, researchers have demonstrated improved capabilities in distinguishing subtle exploit vectors from benign code logic, outperforming traditional static analyzers and generic deep learning baselines.

Furthermore, effective fine-tuning requires careful handling of domain shifts, particularly when applying models to third-party libraries or emerging threat landscapes. Strategies involving lightweight feature extraction, such as sparse token n-grams combined with security metrics, serve as transparent baselines that complement complex neural architectures. Future directions emphasize normalizing code literals to prevent overfitting and integrating compiler-assisted specialization to enhance detection robustness. Ultimately, successful domain adaptation hinges on continuous model refinement using curated security datasets that reflect evolving attack methodologies and coding practices.

2.3.3 Hybrid Neuro-Symbolic Reasoning Frameworks

Hybrid neuro-symbolic reasoning frameworks emerge as a critical paradigm to bridge the gap between the statistical proficiency of deep learning and the rigorous logical guarantees of symbolic systems. While neural architectures like Graph Neural Networks effectively encode structural representations such as Abstract Syntax Trees and Control Flow Graphs, they often struggle with deep semantic reasoning and explicit constraint satisfaction. By integrating symbolic logic, including axiomatic foundations and formal preconditions, these hybrid models aim to overcome the limitations of purely data-driven approaches that frequently fail to capture complex program semantics or rely on superficial syntactic heuristics.

The architectural synthesis typically involves coupling neural encoders with symbolic reasoners to enhance both detection accuracy and interpretability. In this configuration, neural components process high-dimensional code features, such as variable dependencies and token se-

quences, to generate probabilistic embeddings, while symbolic modules enforce logical constraints and validate reasoning paths against formal specifications. This bidirectional interaction allows the system to leverage the generalization capabilities of pre-trained language models for code while utilizing symbolic rules to guide the reasoning process, ensuring that predictions align with established software engineering principles and security properties.

Furthermore, these frameworks significantly advance explainability by generating human-understandable justifications for their outputs, a capability often absent in black-box neural networks. By mapping neural attention mechanisms to explicit logical derivations or counterfactual explanations, hybrid systems provide transparent insights into vulnerability detection and program analysis. This integration not only improves the reliability of automated auditing tools but also facilitates the verification of logic-based flaws that require multi-step deduction, ultimately fostering more trustworthy and robust artificial intelligence systems for software security and maintenance.

2.4 Evaluation Methodologies and Benchmarking

2.4.1 Dataset Construction and Temporal Validation Protocols

Constructing robust datasets for vulnerability detection necessitates rigorous multi-source aggregation and manual verification to ensure ground truth fidelity. Researchers typically synthesize data from repositories like NVD, IBM X-Force, and GitHub, followed by a critical manual validation phase where experts assess extracted key aspects to eliminate false positives. Inter-annotator agreement is quantitatively measured using metrics such as Cohen's Kappa to guarantee consistency. Furthermore, advanced pipelines address historical limitations by removing mislabeled duplicates, integrating hard negative samples, and enriching entries with comprehensive metadata and Common Weakness Enumeration mappings to support diverse machine learning architectures.

A paramount challenge in dataset construction involves maintaining external validity across heterogeneous programming languages and domains. Due to the scarcity of language-specific patch data, many studies employ hybrid datasets combining code-based, repository-based, and cross-

domain sources like OWASP, Juliet, and Android applications. While this diversity enhances model generalizability, it introduces complexities in feature representation, requiring distinct preprocessing techniques for syntax and semantic information. Consequently, systematic characterization of these datasets is essential to document their distributions, limitations, and suitability for specific AI-driven detection tasks, ensuring that models are not biased toward specific repository structures or vulnerability types.

Temporal validation protocols are increasingly adopted to mitigate data leakage and evaluate model performance under realistic chronological conditions. Unlike random splitting, which risks training on future vulnerabilities, temporal validation strictly separates data based on disclosure dates, often spanning over a decade to capture evolving attack patterns. This approach tests a model's ability to generalize to unseen, recent vulnerabilities rather than memorizing historical patterns. By extending temporal coverage and utilizing dedicated test sets from the most recent periods, researchers can more accurately assess the longevity and adaptability of detection systems against emerging threats in dynamic software ecosystems.

2.4.2 Performance Metrics and Generalization Assessment

The evaluation of software engineering models predominantly relies on standard classification metrics, including accuracy, precision, recall, and the F1-score, alongside area under the receiver operating characteristic curve (AUC) for binary tasks. For generative applications, such as commit message synthesis, lexical overlap measures like ROUGE are frequently employed to assess output quality. While these metrics provide a foundational understanding of model efficacy on held-out test sets, they often fail to capture the nuances of real-world deployment scenarios where class imbalance and varying error costs significantly impact utility.

A critical limitation identified in current literature is the substantial degradation of performance when models are applied to cross-project or out-of-distribution datasets. State-of-the-art systems frequently demonstrate high efficacy on curated benchmarks but suffer severe drops in F1-scores when evaluated on more realistic, diverse corpora, indicating poor generalization capabilities. This

discrepancy suggests that many models overfit to specific project characteristics or benchmark artifacts rather than learning intrinsic semantic patterns, thereby limiting their practical applicability in heterogeneous development environments.

To address these generalization gaps, recent studies advocate for the integration of architectural and execution complexity metrics alongside traditional token-based representations. Comparative analyses reveal that combining structural features with deep learning embeddings yields statistically significant improvements in robustness across different domains. Furthermore, rigorous assessment now necessitates evaluating not only predictive accuracy but also runtime efficiency and reproducibility, ensuring that proposed solutions maintain performance stability under varying computational constraints and data distributions.

2.4.3 Interpretability and Explainability in Detection Models

Interpretability in deep learning-based vulnerability detection addresses the critical "black-box" nature of modern classifiers, which often hinders adoption by security practitioners [8]. While high-performance models utilizing graph neural networks or transformers achieve superior accuracy, their decision logic remains opaque without explicit explanation mechanisms. This lack of transparency creates a significant trust deficit, as developers cannot verify whether a flagged vulnerability stems from genuine semantic flaws or spurious correlations within the training data, such as specific coding styles or comment patterns unrelated to actual security risks.

To mitigate these challenges, recent approaches integrate intrinsic explainability features directly into model architectures. Techniques such as hierarchical attention mechanisms operate at both line and token levels to highlight specific code segments contributing to a prediction, offering fine-grained insights into the model's reasoning process. Similarly, post-hoc methods adapted from general explainable AI, including graph-based explainer frameworks, are employed to isolate subgraphs or syntactic structures that drive classification decisions. These methods aim to transform abstract vector representations into human-understandable rationales, distinguishing between true vulnerabilities and false positives caused by obfuscation or dataset bias.

Despite these advancements, substantial gaps

remain in evaluating the fidelity and utility of generated explanations. Current studies indicate that while visual or verbal explanations can enhance user confidence, incorrect or misleading rationales may inadvertently cause experts to overlook genuine threats or waste resources on benign code. Furthermore, the complexity of large-scale software projects often overwhelms simple explanation interfaces, necessitating more sophisticated interaction designs that align with developer workflows. Future research must prioritize rigorous user studies to validate whether these interpretability tools genuinely improve triage efficiency and reduce the cognitive load on security analysts in real-world deployment scenarios.

3 Malicious Package and Supply Chain Attack Analysis

3.1 Taxonomy of Supply Chain Attack Vectors

3.1.1 Dependency Confusion and Namespace Squatting Mechanisms

Dependency confusion exploits the precedence logic inherent in modern package managers, where public repositories are often queried before private registries during resolution. Attackers register malicious packages on public indices using names identical to an organization's internal, private dependencies [12]. When a build system attempts to resolve these dependencies, it inadvertently retrieves the public malicious artifact instead of the intended private library, effectively hijacking the supply chain. This mechanism is particularly potent because it bypasses traditional authentication barriers, relying solely on the ambiguity of namespace resolution rather than compromising developer credentials or infiltrating private infrastructure directly.




Figure 10: Chart from 'An Analysis of Malicious Packages in Open-Source Software in the Wild' (arXiv:2404.04991)

Namespace squatting extends this threat model by manipulating hierarchical naming conventions

to create deceptive identities that mimic legitimate libraries. Unlike simple typosquatting which relies on character-level errors, namespace squatting leverages compound structures and sub-domain similarities to impersonate trusted authors or organizations. In ecosystems supporting hierarchical namespaces, attackers can register packages with slight variations in the namespace path, such as altering a single character in the organization prefix while maintaining a familiar package suffix. This technique exploits the cognitive bias of developers who visually verify package origins, allowing malicious code to infiltrate projects under the guise of authorized updates or common utilities [13].




Figure 11: Chart from 'Detecting Malicious Source Code in PyPI Packages with LLMs: Does RAG Come in Handy?' (arXiv:2504.13769)

The convergence of these mechanisms creates a complex attack surface where static analysis often fails to distinguish between legitimate and malicious artifacts based solely on metadata. Because the malicious packages adhere to standard structural formats and often contain valid dependency graphs, they appear benign until execution [12]. The risk is compounded in environments with transitive dependencies, where a compromised indirect library can trigger the confusion or squatting exploit without direct interaction from the downstream developer. Consequently, detection requires moving beyond simple name matching to analyze resolution behaviors, repository trust levels, and the semantic consistency of namespace hierarchies across the entire dependency tree.

3.1.2 Account Compromise and Maintainer Identity Theft

Account compromise and maintainer identity theft represent critical vectors in software supply chain attacks, often serving as the initial phase for widespread malware propagation [14]. Attackers frequently employ credential stuffing, phishing, or social engineering to hijack legitimate developer accounts, thereby bypassing registry security fil-

ters that trust established identities. Once access is gained, adversaries can publish malicious updates to popular packages, leveraging the implicit trust of downstream consumers to distribute payloads such as credential stealers or reverse shells without raising immediate suspicion [15].




Figure 12: Chart from 'Towards Measuring Supply Chain Attacks on Package Managers for Interpreted Languages' (arXiv:2002.01139)

The severity of these incidents is exemplified by high-profile breaches where compromised maintainer credentials led to the injection of exfiltration code into widely used libraries. In such scenarios, the malicious code executes during the installation process, harvesting sensitive data like API keys, authentication tokens, and local configuration files from victim environments. Because the attack originates from a verified account within a trusted package, traditional static analysis tools often fail to detect the anomaly, allowing the threat to persist across thousands of repositories before discovery and remediation efforts can contain the spread [16].

$$SFR = \frac{|F_{sliced} \cap F_{original}|}{|F_{original}|} \times 100\%$$

Figure 13: Chart from 'Taint-Based Code Slicing for LLMs-based Malicious NPM Package Detection' (arXiv:2512.12313)

To mitigate these risks, the security community increasingly advocates for mandatory multi-factor authentication (MFA) for maintainers of critical or highly downloaded packages [17]. Registry operators are urged to implement automated scoring systems that enforce strict identity verification protocols once a package reaches a specific popularity threshold. Furthermore, enhancing detection mechanisms to identify abnormal login patterns and enforcing policies around ownership transfers are essential steps. These measures aim to reduce the attack surface by ensuring that account compromise does not automatically equate to supply chain contamination, thereby preserving the in-

egrity of the open-source ecosystem.

3.1.3 Malicious Code Injection via Build and Install Scripts

Malicious code injection via build and install scripts represents a critical vector in software supply chain attacks, exploiting the privileged execution context granted during package installation [18]. Attackers frequently embed payloads within standard build artifacts, such as `setup.py` in Python or `install.js` in Node.js, which are automatically executed by package managers. These scripts often leverage obfuscation techniques, including base64 encoding and dynamic command construction, to download and execute remote binaries or exfiltrate sensitive data before the legitimate package functionality is ever invoked.

The sophistication of these attacks has evolved from simple payload delivery to complex, multi-stage intrusion patterns that evade traditional static analysis. By distributing malicious logic across multiple dependencies or hiding it within seemingly benign configuration files, adversaries can bypass signature-based detectors that rely on known malicious patterns. Furthermore, attackers increasingly utilize stealth strategies, introducing malicious code only after a package has established a trusted user base, thereby complicating the distinction between legitimate build processes and compromised installation routines [19].




Figure 14: Chart from 'An Empirical Study of Malicious Code In PyPI Ecosystem' (arXiv:2309.11021)

Current defense mechanisms face significant challenges in addressing these threats due to the inherent necessity of script execution during the installation phase. While dynamic analysis tools offer visibility into runtime behaviors, they often struggle with scalability and the automatic differentiation of malicious intent from legitimate system interactions. Emerging approaches propose integrating lightweight monitoring systems and semantic analysis to inspect installation scripts for anomalous network requests or file system mod-

ifications, yet a comprehensive solution that balances security with the flexibility required by diverse build ecosystems remains an open research problem.

3.2 Detection Mechanisms for Compromised Artifacts

3.2.1 Static Analysis and Behavioral Pattern Matching

Static analysis remains a foundational technique for identifying malicious patterns within software packages without execution, primarily leveraging Abstract Syntax Tree (AST) parsing and control flow graph construction [20]. Traditional approaches extract syntactic features such as Portable Executable headers, string literals, and byte sequences to classify binaries using machine learning models like Naive Bayes. While computationally efficient and scalable for large ecosystems, these methods often struggle with dynamic language features, reflection, and sophisticated obfuscation techniques that alter code structure at runtime, leading to significant false negatives in complex supply chain attacks.

$$is_malicious(p) = \bigvee_{i=1}^k match(r_i, p)$$

Figure 15: Chart from 'Cutting the Gordian Knot: Detecting Malicious PyPI Packages via a Knowledge-Mining Framework' (arXiv:2601.16463)

To mitigate the limitations of purely syntactic inspection, behavioral pattern matching integrates static signatures with inferred runtime characteristics to detect anomalies. This hybrid strategy involves mapping API usage sequences and data flow paths to known malicious behaviors, such as unauthorized network connections or file system modifications, directly from the source or binary code. By correlating static artifacts with behavioral heuristics, scanners can identify suspicious logic chains indicative of malware, although reliance on predefined rules often restricts adaptability against novel attack vectors or polymorphic code variations.

Recent advancements address these rigidity issues by incorporating semantic analysis and large language models to enhance pattern recognition capabilities beyond simple signature matching. These evolved static analyzers aim to understand

the intent behind code segments, distinguishing between benign utility functions and malicious payloads even when obfuscated. Despite improvements in detecting higher-order regularities, static analysis alone frequently fails to capture context-dependent vulnerabilities introduced during installation or interaction with the runtime environment, necessitating a complementary dynamic verification layer for comprehensive security assurance.

3.2.2 Dynamic Execution Tracing and Sandbox Monitoring

Dynamic execution tracing and sandbox monitoring have emerged as critical methodologies for detecting malicious behaviors that evade static analysis. By executing software artifacts within isolated environments, researchers can observe runtime interactions such as system calls, network traffic, and file system modifications. Tools like Sysdig, Strace, and DTrace are frequently employed to capture low-level kernel and user-space activities, generating comprehensive behavioral traces. These traces are often embedded into vector spaces or analyzed via rule-based matching to distinguish between benign operations and malicious payloads, including privilege escalation, anti-debugging mechanisms, and unauthorized data exfiltration.

To mitigate the risks associated with executing untrusted code, modern frameworks leverage containerization technologies, such as Docker, to create secure, reproducible sandboxes. This approach ensures that potential sabotage, abuse, or environment-adaptive evasion techniques do not compromise the host infrastructure. Advanced systems integrate dynamic tracing with machine learning models to automate the classification of suspicious activities, analyzing patterns like repeated loops, payload triggering, and anomalous dependency fetching. Furthermore, some methodologies combine static pre-filtering with dynamic verification to optimize resource usage, addressing the computational costs inherent in full-scale runtime monitoring across diverse operating systems.

Despite their efficacy, dynamic analysis techniques face significant challenges, including limited code coverage, high resource consumption, and sophisticated anti-analysis evasion strategies. Malware increasingly employs environment-aware logic to remain dormant when detecting sandbox signatures, leading to false negatives. Additionally, behavioral sequences often lack explicit

semantic intent, complicating the differentiation between legitimate administrative tasks and malicious actions. Current research focuses on enhancing sandbox diversity, implementing forced execution paths, and integrating explainable AI to interpret complex behavioral data, aiming to reduce false positives while maintaining scalability for large-scale software supply chain security.

3.2.3 Machine Learning Classifiers for Anomaly Detection

Machine learning classifiers have emerged as pivotal tools for identifying anomalous behaviors within software supply chains, particularly in detecting malicious packages across ecosystems like npm and PyPI [21]. Traditional supervised approaches leverage ensemble methods, such as Support Vector Machines and Random Forests, trained on static and dynamic features including executed commands, network destinations, and code metrics. These models demonstrate high efficacy in distinguishing benign from malicious artifacts, with recent studies reporting precision rates exceeding 0.98 by utilizing tree-based boosting techniques that effectively capture non-linear relationships in high-dimensional feature spaces.




Figure 16: Chart from 'A Machine Learning-Based Approach For Detecting Malicious PyPI Packages' (arXiv:2412.05259)

Despite their success, conventional classifiers face significant challenges regarding data imbalance and the evolving nature of attack vectors. The scarcity of labeled malicious samples often leads to elevated false negative rates, while reliance on handcrafted features limits the detection of novel obfuscation techniques or zero-day threats. Furthermore, static analysis alone frequently fails to uncover runtime anomalies, necessitating the integration of dynamic execution traces. To mitigate these limitations, researchers are increasingly adopting hybrid architectures that combine heuristic rules with deep learning representations, allowing for the automatic extraction of semantic patterns directly from source code without extensive

manual feature engineering.

Recent advancements have shifted towards leveraging Large Language Models and unsupervised anomaly detection to enhance adaptability against sophisticated adversarial strategies. By fine-tuning models on labeled domains or employing clustering algorithms to identify outliers, these systems can detect malware variants and previously unseen attack patterns with greater robustness. Although issues such as model poisoning and evasion attacks remain critical concerns, the integration of explainable AI techniques provides necessary transparency, enabling security analysts to validate predictions and iteratively refine detection rules. This evolution marks a transition from rigid, signature-based scanning to intelligent, behavior-driven security frameworks capable of real-time threat mitigation.

3.3 Large Language Models in Threat Intelligence

3.3.1 Automated Malware Analysis and Semantic Reasoning

Automated malware analysis has evolved from static signature matching to sophisticated dynamic behavioral profiling, yet significant challenges remain in interpreting semantic intent. Traditional dynamic analysis often suffers from weak semantic cohesion in API associations, where raw execution traces fail to distinguish between benign administrative tasks and malicious activities. Furthermore, the increasing prevalence of anti-analysis techniques, such as environment-adaptive execution, allows modern malware to evade sandboxed environments, rendering purely observational methods insufficient for comprehensive threat detection in complex software supply chains.

To address these limitations, recent research integrates semantic reasoning through knowledge graphs and large language models to infer malicious intent beyond superficial system calls. By constructing cybersecurity knowledge graphs from after-action reports, systems can map relationships between disparate indicators of compromise, enabling the identification of novel variants through clustering and logical inference. Advanced frameworks now employ explainable deep learning and chain-of-thought prompting to guide automated agents in dissecting computational graphs, effectively isolating core attack functions

within serialized models or heterogeneous configuration files that rule-based scanners frequently miss.

The convergence of semantic understanding with automated vetting pipelines significantly reduces false positives while enhancing the detection of subtle threats like dependency confusion and model injection. Unlike previous lexical approaches that struggle with the nuances of natural language and code obfuscation, semantic reasoning engines evaluate the purpose behind system interactions, distinguishing accidental errors from deliberate exploitation. This shift towards cognitively inspired inference not only accelerates the identification of malicious packages across diverse registries but also provides the explanatory power necessary for human analysts to triage alerts and refine security ontologies effectively.

3.3.2 Generative Approaches for Rule Synthesis and Explanation

Generative approaches for rule synthesis leverage large language models to automatically derive detection logic from raw code artifacts and behavioral traces. By analyzing abstract syntax trees and dynamic execution flows, these systems identify recurring patterns associated with malicious activities, such as indirect sink propagation or suspicious API invocations. The synthesis process involves transforming observed feature sequences into formalized rules, often utilizing iterative self-refinement loops where the model generates initial candidates, critiques their precision against held-out corpora, and refines lexical triggers and argument structures to minimize false positives while maintaining syntactic validity.

Beyond mere generation, these frameworks prioritize the explainability of synthesized rules by translating complex symbolic logic into natural language descriptions. This dual-output mechanism ensures that each generated rule is accompanied by a semantic justification, detailing the specific code constructs and data flows that triggered its creation. Techniques such as embedding rule metadata and utilizing retrieval-augmented generation allow analysts to query the rationale behind specific detections, effectively bridging the gap between opaque neural feature extraction and interpretable, human-readable security policies. This transparency is critical for validating that the generated rules capture genuine adversarial intent rather than spurious correlations.

The integration of generative synthesis with verification stages creates a robust pipeline for evolving defense mechanisms against novel attack vectors. Once candidate rules are generated and explained, they undergo rigorous entity verification where the model cross-references the proposed logic against original source texts to ensure consistency and completeness. Incorrect or ambiguous rules trigger a feedback loop, prompting the system to adjust prompts and incorporate specific counter-examples for correction. This continuous refinement cycle not only enhances the accuracy of the rule base but also facilitates the automated discovery of subtle obfuscation techniques, such as typosquatting or dynamic code generation, that traditional static analysis often misses.

3.3.3 Multi-Agent Frameworks for Collaborative Forensics

Multi-agent frameworks for collaborative forensics address the limitations of standalone Large Language Models by orchestrating specialized Worker Agents under a central supervisor. This architecture employs a Plan-and-Execute coordination model to decompose complex supply chain investigations into manageable, executable steps, thereby mitigating issues such as hallucination and context confusion. By assigning distinct roles—ranging from static metadata examination to dynamic sandbox monitoring—these systems enable autonomous, multi-step problem solving that surpasses the capabilities of single-turn interaction models in detecting stealthy malicious injections.

To enhance detection accuracy, these frameworks integrate deterministic external tools directly into the agent workflow, allowing for real-time interaction with the analysis environment. Specialized utilities, such as YARA rule compilers, serializable method extractors, and retrieval-augmented generation modules for API documentation, provide ground-truth evidence that pure reasoning models often miss. The supervisor agent maintains a global analysis state, aggregating intermediate reports from diverse sources like system call traces and network traffic logs, while employing voting mechanisms to resolve conflicts and ensure the consistency of aggregated intelligence across heterogeneous data streams.

Despite their efficacy, current implementations face significant challenges regarding computational feasibility, particularly the high memory footprint and inference latency associated with

continuous real-time scanning of production registries. Furthermore, the semantic complexity of diverse configuration files across different ecosystems complicates the development of unified detection heuristics, often leading to overfitting on known attack patterns. Future advancements aim to extend these agentic systems to support cross-platform sandboxing and kernel-level monitoring, ensuring robust generalizability against evolving evasion techniques while maintaining explainable and reproducible forensic outcomes.

3.4 Evasion Techniques and Adversarial Robustness

3.4.1 Code Obfuscation and Packing Strategies

Code obfuscation and packing represent primary evasion mechanisms in software supply chain attacks, designed to conceal malicious logic from both manual inspection and automated static analysis. Attackers employ diverse techniques, including randomization, encoding, and logic structure transformation, to distort code semantics and hinder decompilation efforts. By embedding malicious components within unreachable code sections or utilizing complex lambda chains to alias critical functions like `exec`, adversaries effectively reduce the centrality scores of malicious APIs in static call graphs, thereby bypassing signature-based detectors and heuristic filters [22].




Figure 17: Chart from 'CHASE: LLM Agents for Dissecting Malicious PyPI Packages' (arXiv:2601.06838)

Packing strategies further complicate detection by compressing or encrypting payload code, which remains dormant until runtime execution. This approach necessitates dynamic analysis or specialized unpacking tools to extract the hidden executable logic, as static methods often fail to resolve indirect calls or retrieve remote payloads embedded within compressed archives. The re-

silience of these packed binaries against binary diffing and disassembly creates a significant blind spot for traditional security scanners, allowing multi-stage payloads and logic bombs to persist within package ecosystems until they are activated in a victim's environment.

To counter these sophisticated evasion tactics, recent research advocates for hybrid detection frameworks that integrate taint-based code slicing with large language models. By isolating security-relevant code segments and filtering out non-reachable noise, these methods reduce context overhead while preserving critical semantic information necessary for identifying obfuscated threats. Furthermore, incorporating dynamic behavior monitoring within isolated sandbox environments provides a essential defense-in-depth layer, ensuring that even successfully obfuscated or packed malicious code is detected upon execution before it can compromise the host system or exfiltrate sensitive data.

3.4.2 Adversarial Perturbations Against Detection Systems

Adversarial perturbations against detection systems represent a critical vulnerability in machine learning-based malware identification, where attackers deliberately craft inputs to induce misclassification [23]. In the context of software package ecosystems, these attacks often manifest as evasion techniques where malicious code is augmented with benign components or structural modifications designed to bypass static and dynamic analyzers [13]. By exploiting the decision boundaries of classifiers, adversaries can inject payloads that retain their harmful functionality while appearing statistically similar to legitimate software, effectively reducing the recall of detection models without triggering alarm thresholds.

Specific methodologies include visualization-based attacks and graph neural network manipulations, where subtle changes in package metadata, dependency structures, or code representation fool deep learning architectures. For instance, attackers may alter control flow graphs or inject noise into feature vectors used by models trained on bytecode or abstract syntax trees, causing the system to label malicious packages as benign. These perturbations are particularly dangerous in supply chain scenarios, as they allow compromised updates to propagate undetected by leveraging the trust established in previous versions or by mim-




Figure 18: Chart from 'Enhancing the Insertion of NOP Instructions to Obfuscate Malware via Deep Reinforcement Learning' (arXiv:2111.09626)

icking the semantic patterns of popular libraries.

The resilience of current detection frameworks remains insufficient against such adaptive threats, highlighting a gap between theoretical robustness and practical deployment. While some studies propose adversarial training to harden models, the dynamic nature of package repositories and the low cost of generating new perturbations create an ongoing arms race. Consequently, ensuring model stability requires not only rigorous evaluation against known attack vectors but also the development of detection mechanisms capable of identifying semantic inconsistencies that persist despite surface-level obfuscation, thereby mitigating the risk of false negatives in high-stakes environments.

3.4.3 Resilience of Models to Feature Manipulation

The resilience of machine learning models against feature manipulation is a critical determinant of their security posture, particularly when facing semantic-preserving obfuscations. Traditional approaches relying on hand-engineered static features often lack robustness, as adversaries can easily alter code structure through instruction reordering or variable renaming without affecting functionality, thereby evading detection. Such manipulations degrade the performance of classifiers that depend on rigid syntactic patterns, creating a dangerous illusion of security where models fail to generalize against evolved threats or cross-ecosystem variations.

To mitigate these vulnerabilities, recent advancements have shifted toward deep learning architectures capable of learning invariant represen-

tations from raw code or dynamic behavioral data. Techniques utilizing graph neural networks and transformer-based models incorporate control flow and data-flow awareness, enabling the extraction of features resilient to compiler optimizations and superficial code transformations [24]. Furthermore, integrating dynamic analysis allows for the capture of runtime behaviors that static analyzers cannot detect, ensuring that the underlying malicious intent remains visible even when the source code is heavily obfuscated or transformed at execution time.

Despite these improvements, the stability of detection mechanisms remains contingent upon rigorous feature selection and the elimination of noise introduced by irrelevant metadata. Empirical studies indicate that specific feature combinations can inadvertently degrade detection capabilities if not carefully curated, necessitating the use of explainable AI techniques to validate decision boundaries. Ultimately, achieving true resilience requires a holistic strategy that combines multimodal feature representations with continuous empirical feedback loops, ensuring that models maintain high fidelity and transparency when confronted with sophisticated, evasive manipulation tactics.

4 Software Supply Chain Risk Assessment and Governance

4.1 Dependency Graph Analysis and Vulnerability Propagation

4.1.1 Transitive Dependency Resolution and Reachability Analysis

Transitive dependency resolution constitutes the foundational mechanism for constructing accurate Software Bill of Materials (SBOMs), yet it faces significant challenges due to non-deterministic resolution algorithms and ecosystem-specific constraints. Unlike direct dependencies explicitly declared in manifests, transitive chains are dynamically inferred by package managers, often resulting in divergent dependency graphs across different build environments. Recent empirical studies highlight that resolution logic varies substantially between ecosystems like Maven, npm, and Cargo, where version conflicts and overlapping constraints can lead to unpredictable artifact selection. Consequently, static analysis tools frequently struggle to reconstruct the exact runtime graph without executing the native build system,

creating a gap between declared specifications and actual consumed artifacts.

Reachability analysis addresses the critical limitation of traditional vulnerability scanning by distinguishing between present dependencies and those actively invoked during execution. Given that a vast majority of vulnerabilities reside in transitive layers, blindly flagging all present libraries generates excessive false positives and obscures genuine risks. Advanced methodologies now employ heterogeneous graph models and semantic lifting to map call paths from application entry points to vulnerable functions, effectively filtering out unreachable code. This approach shifts the security paradigm from component-level scoring to path-based risk assessment, identifying multi-step cascaded threats that propagate through complex dependency chains while ignoring dormant libraries that pose no immediate exploit vector.

The integration of reachability data with transitive resolution enables prioritized remediation strategies, focusing resources on high-severity vulnerabilities within active execution paths. While automated bots facilitate routine updates, they often lack the contextual awareness to resolve deep transitive conflicts or assess the behavioral impact of upgrades on reachable code. Emerging research suggests that combining static dependency graphs with dynamic tracing and large language model-driven semantic understanding can enhance the precision of vulnerability propagation analysis. However, current industrial implementations often lack reproducible frameworks for validating these autonomous remediation agents, necessitating further development of open benchmarks to ensure reliable supply chain security across diverse software architectures.

4.1.2 Knowledge Graph Construction for Risk Aggregation

Knowledge graph construction for risk aggregation begins by transforming raw Software Bill of Materials (SBOMs) and dependency trees into structured, semantic representations capable of capturing complex supply chain topologies. This process involves defining a surjective graph transformation that compresses million-node code property graphs into scalable, domain-specific subgraphs governed by a rigorous vulnerability relation ontology. By mapping disparate artifacts—such as packages, versions, and known Common Vulnerabilities and Exposures (CVEs)—onto unified

nodes, the framework establishes a heterogeneous network that preserves critical lineage information while eliminating redundant structural noise inherent in flat dependency lists.

The core of this construction lies in encoding multi-hop dependencies and contextual relationships that facilitate the propagation of risk scores across the network. Unlike static scoring systems that evaluate components in isolation, the generated graph embeds semantic relations reflecting upstream and downstream connectivity, allowing for the identification of bridge nodes and single points of failure. These structural features are essential for modeling cascading failures, as they enable the system to distinguish between locally dense clusters and globally sparse connections, thereby highlighting paths where a compromise in a foundational library could amplify risks throughout the entire ecosystem.

Finally, the constructed graph serves as the substrate for advanced reasoning engines that compute composite risk metrics through learned node embeddings. By integrating attention mechanisms modulated by semantic relation types, the model aggregates inherent vulnerability severity with contextually propagated risk from neighboring entities. This approach supports the detection of both known vulnerabilities and zero-day anomalous behaviors by analyzing long-range dependencies that traditional static analysis overlooks. The resulting composite scores provide a dynamic, quantitative basis for prioritizing mitigation efforts, effectively translating complex graph topology into actionable security intelligence for supply chain governance.

4.1.3 Temporal Modeling of Vulnerability Evolution

Temporal modeling of vulnerability evolution addresses the dynamic nature of software security, where vulnerability patterns shift due to emerging CWE categories and updated attack matrices. Traditional static analysis often fails to capture these temporal dynamics, leading to concept drift where predictive accuracy degrades over multi-year deployments without continual learning mechanisms. Effective frameworks must therefore integrate historical attack data with real-time policy rules to maintain validity across evolving software ecosystems, ensuring that detection models adapt to novel exploit classes rather than relying solely on syntactic pattern matching.

The vulnerability lifecycle is typically segmented into four distinct stages: introduction, discovery, fixing, and disclosure, each characterized by specific temporal metrics. Empirical studies utilizing Kaplan-Meier estimators reveal a sharp decline in survival probability shortly after introduction, indicating rapid disclosure for many issues, yet a persistent long-tail distribution exists where vulnerabilities remain undetected for years. In ecosystems like npm and RubyGems, the median time between introduction and fix identification can span five to seven years, highlighting a critical latency gap that temporal models aim to quantify and reduce through proactive monitoring of dependency trees and commit histories.

Advanced approaches leverage graph-based transformations and survival analysis to map vulnerability propagation paths and forecast future risks within software supply chains. By encoding affect relationships and transitivity granularities, these models compute the proportion of dependency updates introducing known vulnerabilities and assess the net impact of semantic version differences between vulnerable and patched states [25]. This granular temporal reasoning enables the prediction of adoption delays for security fixes, allowing developers to understand how package release frequencies and maintenance practices directly influence the speed of vulnerability remediation propagation throughout the ecosystem.

4.2 Software Bill of Materials and Provenance Verification

4.2.1 Automated SBOM Generation and Enrichment Techniques

Automated Software Bill of Materials (SBOM) generation has become a foundational control for mitigating supply chain risks, shifting from manual inventory to integrated build workflow enforcement [26]. Modern pipelines leverage tools like Syft, Trivy, and Microsoft's sbom-tool to parse package manifests and binary artifacts, producing standardized outputs in SPDX or CycloneDX formats. These tools facilitate the precise mapping of direct and transitive dependencies, addressing the critical gap where projects lack visibility into their component lineage. By embedding generation steps directly into continuous integration environments, organizations ensure that every build artifact is accompanied by a machine-readable inventory, effectively countering the exposure seen

in ecosystems lacking such automation.




Figure 19: Chart from 'An Empirical Study on Software Bill of Materials: Where We Stand and the Road Ahead' (arXiv:2301.05362)

Beyond mere enumeration, contemporary techniques focus on SBOM enrichment through real-time vulnerability correlation and reachability analysis [27]. Static scanners often flag outdated packages based on known CVE databases, but advanced workflows integrate dynamic analysis to determine if vulnerable code paths are actually executable within the specific application context. This differentiation between theoretical and exploitable risk reduces false positives and prioritizes remediation efforts. Furthermore, enrichment processes now incorporate metadata regarding license compliance and maintenance status, transforming the SBOM from a static list into a dynamic security asset that informs policy enforcement and update strategies without requiring exhaustive manual review.




Figure 20: Chart from 'A Reality Check on SBOM-based Vulnerability Management: An Empirical Study and A Path Forward' (arXiv:2511.20313)

Despite these advancements, significant challenges remain in handling complex, polyglot codebases and ensuring domain generalization across diverse hardware architectures. Current approaches often struggle with opaque binaries or proprietary firmware where source manifests are unavailable, necessitating static binary analysis to recover behavioral semantics. Future direc-

tions point toward agentic systems and large language models capable of autonomously reasoning about dependency graphs and generating platform-specific remediation configurations. However, the integration of such intelligent mechanisms must balance memory complexity and the need for continuous updates to detect novel attack vectors, ensuring that automated generation evolves alongside the sophistication of supply chain threats.

4.2.2 Binary Artifact Attestation and Build Tracing

Binary artifact attestation establishes a cryptographic link between build outputs and their provenance, addressing the limitations of post-production verification in modern supply chains. By integrating frameworks like SLSA and in-toto, organizations can generate tamper-evident metadata that records the complete build lifecycle, including source inputs, build environment specifications, and execution steps. This approach ensures that release artifacts are not merely signed but are accompanied by verifiable claims regarding their origin, effectively mitigating risks where compromised pipelines inject malicious code into seemingly legitimate, signed binaries.

Build tracing extends this security posture by enforcing reproducible builds, where identical source code and environments yield bit-identical outputs. Techniques such as OmniBOR automatically construct Artifact Dependency Graphs to trace transitive relationships without developer intervention, enabling the detection of injected vulnerabilities that traditional static analysis might miss in stripped binaries. Unlike dynamic analysis, which often faces hardware emulation barriers, static tracing of binary artifacts allows for high-coverage semantic recovery, ensuring that any deviation from the expected build path is immediately flagged as a potential integrity violation before deployment.

Despite these advancements, significant gaps remain in verifying the behavioral semantics of compound AI systems and trained models, where bit-for-bit reproducibility is often unattainable. Current mechanisms primarily focus on structural provenance rather than active vulnerability removal, necessitating a shift toward negotiated versioning contracts and cryptographically bound state snapshots. To achieve true defense-in-depth, attestation workflows must evolve beyond simple presence checks to include reachability analysis

and real-time policy enforcement, ensuring that vulnerable dependencies are identified and remediated within the build pipeline rather than relying solely on downstream scanning.

4.2.3 Dynamic Consistency Verification in Runtime Environments

Dynamic consistency verification in runtime environments addresses the critical gap left by static provenance frameworks like SLSA and in-toto, which primarily validate post-production artifacts rather than monitoring live execution. While traditional mechanisms ensure ledger consistency through Byzantine Fault Tolerant consensus and Merkle-tree hashing, they often fail to detect dynamic threats such as injection attacks, insecure deserialization, or broken access control that emerge only during operation. Consequently, relying solely on static analysis or artifact signing proves insufficient for preventing real-time pipeline compromises, necessitating a shift toward continuous, context-aware runtime monitoring that can autonomously respond to evolving security posture deviations.

The implementation of robust runtime verification faces significant architectural bottlenecks, particularly the "Environment Wall" where emulating proprietary hardware peripherals for high-coverage dynamic testing becomes prohibitively expensive. Although tools like fuzzers and runtime monitors aim to penetrate this opacity, they frequently struggle with scalability when applied to monolithic firmware or stripped binaries lacking source code. To mitigate these limitations, emerging approaches leverage static binary artifacts to recover high-level behavioral semantics without hardware emulation, while integrating rate limiting and staking-based penalties to deter denial-of-service attempts. This hybrid strategy enhances tamper resistance by correlating observed runtime behaviors against expected policy definitions, ensuring that data distribution and process interactions remain authentic and verified.

Furthermore, effective dynamic consistency requires bridging the disconnect between automated dependency management tools and actual runtime behavior. Although utilities like Dependabot automate updates, they often lack the contextual awareness to verify if a patched dependency introduces runtime inconsistencies or fails to resolve specific vulnerability chains under load. Future directions emphasize compositional verification, decompos-

ing pipelines into deterministic segments where observability ensures that changes in runtime behavior are attributable to specific upstream artifacts. By establishing clear policies for process-level assurance and integrating dynamic tracing with continuous integration checks, systems can achieve a resilient state where provenance, attestation, and real-time integrity verification operate synergistically to secure the software supply chain.

4.3 Trust Metrics and Reputation Systems

4.3.1 Quantitative Assessment of Project Health and Maintenance

Quantitative assessment of project health relies on a multidimensional analysis of maintenance indicators, including release frequency, the recency of the latest version, and the trajectory of contributor activity. These metrics extend beyond simple commit counts to evaluate contributor diversity, organizational backing, and the robustness of the underlying governance model. By synthesizing data on issue resolution times, response latencies, and recent engagement levels, researchers can derive a composite health score that reflects the project's capacity to sustain long-term development and respond effectively to emerging challenges.

System overhead and operational stability serve as critical proxies for maintenance efficiency, measured through continuous integration build durations, resource consumption per agent, and the frequency of pipeline interruptions. Empirical studies indicate that high-performing teams maintain build overheads below six percent, often executing mitigations without human intervention. This quantitative lens further encompasses dependency management strategies, where the ratio of dependencies to project size is analyzed against policy adherence to determine if increased complexity correlates with heightened security risks or delayed vulnerability remediation.

Advanced assessment frameworks increasingly integrate automated agents to scan commits, software bills of materials, and infrastructure-as-code configurations, providing granular visibility into the supply chain. These tools facilitate the correlation of vulnerability fix rates with project popularity metrics, such as star and fork counts, revealing that widespread adoption does not always guarantee timely patching [28]. Consequently, modern evaluation methodologies prioritize reproducible,

automated scoring systems that account for both static metadata and dynamic behavioral patterns to accurately predict a project’s resilience and maintenance sustainability.




Figure 21: Chart from 'Code Ownership in Open-Source AI Software Security' (arXiv:2312.10861)

4.3.2 Contributor Behavior Analysis and Identity Verification

Contributor behavior analysis leverages multidimensional network metrics, including in-degree, out-degree, and composite centrality, to systematically identify high-risk nodes within the software supply chain. By examining contribution graphs, researchers can distinguish between established maintainers and transient volunteers, assessing the implications of inconsistent peer review and variable commit frequencies on project integrity. This behavioral profiling is critical for understanding how security risks propagate through dependency trees, particularly when vulnerable dependencies persist due to delayed adoption rather than fix availability.

Identity verification serves as a foundational mechanism for ensuring code continuity and preventing impersonation in ecosystems where volunteer commit access is prevalent. Robust verification frameworks enforce account security through multi-factor authentication and mandate cryptographic signing for commits and releases to guarantee non-repudiation. Furthermore, the implementation of permissioned security ledgers records agent identities, observed signals, and mitigation actions, creating an immutable audit trail that links specific behavioral patterns to distinct, verified entities while preserving the integrity of the development lifecycle.

Despite these advances, significant challenges remain in machine-readably capturing qualitative aspects of contributor trust, such as the rigor of code reviews or response times to disclosed vulnerabilities. Current automated approaches often miss formal contributor declarations and explicit security testing policies, necessitating a hybrid

strategy that combines structured identity checks with context-aware behavioral analysis. Future methodologies must bridge this gap by integrating reproducible build artifacts with verified contributor identities, ensuring that artifact immutability aligns with the authenticated provenance of the human or agent actors modifying the codebase.

4.3.3 Community-Driven Risk Scoring and Alert Prioritization

Community-driven risk scoring leverages collective intelligence to refine vulnerability severity beyond static metrics like CVSS, addressing the gap between theoretical impact and real-world exploitability [29]. By aggregating data from diverse ecosystems such as npm and PyPI, these frameworks incorporate contextual signals including adoption delays, dependency centrality, and historical fix rates to generate dynamic risk profiles. This approach recognizes that a vulnerability’s true threat level is often dictated by its propagation potential within the software supply chain rather than isolated component characteristics [30].




Figure 22: Chart from 'A Survey on Automated Software Vulnerability Detection Using Machine Learning and Deep Learning' (arXiv:2306.11673)

Alert prioritization mechanisms increasingly utilize graph-based reasoning and machine learning to mitigate alert fatigue, a critical issue where high-severity findings are obscured by noise. Advanced systems employ heterogeneous graphs to model multi-hop dependencies, identifying bridge nodes that disproportionately amplify cascading risks across clusters. Furthermore, integrating Large Language Models enhances semantic understanding of vulnerability contexts, allowing for the adaptive filtering of false positives and the elevation of alerts tied to active exploitation patterns or critical business assets.

Despite these advancements, significant challenges remain in standardizing community inputs and ensuring the robustness of automated scoring against biased datasets. Current implementations often struggle with the latency between vulnerability disclosure and community consensus, leading to delayed mitigation for emerging threats. Future directions must focus on harmonizing qualitative community insights with quantitative loss exposure models to create resilient, real-time prioritization engines that effectively guide resource allocation in complex, interdependent software environments.

4.4 Governance Frameworks and Automated Mitigation

4.4.1 Policy Enforcement via Zero Trust Architectures

Policy enforcement within Zero Trust Architectures (ZTA) for software supply chains necessitates a shift from perimeter-based defenses to continuous, identity-centric verification of all dependencies and build artifacts. Unlike traditional models that assume implicit trust within internal networks, ZTA mandates strict validation of every transaction, treating direct and transitive dependencies as untrusted until proven otherwise. This approach integrates automated policy engines that evaluate artifacts against dynamic criteria, including license compliance, supplier approval status, and real-time vulnerability data, ensuring that only components satisfying explicit security constraints are admitted into the development lifecycle.

The implementation of these architectures relies heavily on immutable ledgers and cryptographic attestation to guarantee the integrity and non-repudiation of agent actions throughout the pipeline. By recording agent identities, reasoning summaries, and mitigation outcomes in a permissioned blockchain structure, organizations can establish an auditable trail that prevents tampering and facilitates rapid forensic analysis. This mechanism supports the enforcement of least-privilege principles, where repository-level agents are restricted to minimal permissions required for specific tasks, thereby mitigating risks associated with compromised credentials or malicious insider threats that could otherwise exfiltrate secrets or inject vulnerable code.

Furthermore, ZTA frameworks address the stochastic nature of modern AI-driven develop-

ment by enforcing deterministic policies on probabilistic outputs. As large language models introduce bounded non-determinism into code generation, policy engines must verify that generated artifacts adhere to predefined safety lattices and do not propagate latent vulnerabilities from training data. This involves continuous monitoring of dependency graphs to detect drift and automatically triggering remediation workflows when policy violations occur. Consequently, Zero Trust transforms security from a post-hoc verification step into an embedded, unending enforcement layer that adapts to evolving threat landscapes without relying on vendor cooperation or static signatures.

4.4.2 Agentic Workflows for Autonomous Remediation

Agentic workflows for autonomous remediation represent a paradigm shift from reactive verification to proactive, self-defending software supply chains. These systems employ structured, directed acyclic graphs (DAGs) orchestrated by modern frameworks to execute complex remediation tasks without human intervention. By integrating domain-specific prompts with reasoning capabilities, agents continuously monitor build environments, evaluate code provenance, and assess dependency behaviors. This architectural approach allows for the dynamic isolation of anomalies and the automatic curation of packages, effectively reducing the window of opportunity for attackers while minimizing mitigation latency compared to traditional rule-based baselines.

The core mechanism of these workflows relies on neuro-symbolic integration, combining large language models with formal verification tools to ensure both flexibility and correctness. Agents utilize reflexive prompting and reinforcement learning to navigate abstract security lattices, translating probabilistic risk detections into deterministic reachability proofs via constraint solvers. This hybrid methodology enables the system to not only identify high-risk paths in stripped binaries but also to generate and validate fixes that adhere to strict compliance policies. Consequently, the workflow bridges the semantic gap between high-level security intent and low-level binary constraints, offering verifiable guarantees rather than heuristic suggestions.

Despite promising industrial prototypes demonstrating capabilities in self-healing dependencies and automated patching, significant challenges re-

main regarding trust and reproducibility. Current implementations often lack peer-reviewed validation and robust datasets necessary to evaluate the frequency and security consequences of agent-driven dependency modifications. Future research must address the establishment of reliable provenance signals to distinguish AI-generated remediations from human interventions, ensuring that autonomous agents do not inadvertently introduce new vulnerabilities or bypass critical security controls. Ultimately, realizing fully autonomous remediation requires rigorous benchmarking of agent decision-making processes against real-world attack scenarios to foster confidence in their deployment within critical infrastructure.

4.4.3 Integration of Security Controls in CI/CD Pipelines

The integration of security controls within CI/CD pipelines has evolved from periodic manual audits to continuous, automated enforcement mechanisms essential for modern software supply chain integrity. Central to this paradigm is the implementation of cryptographically signed release artifacts and rigorous identity verification, which prevent tampering during the build and deployment phases. By embedding static and dynamic code analysis directly into pipeline workflows, organizations can detect vulnerabilities such as injection attacks, insecure deserialization, and broken access controls before artifacts reach production. This shift ensures that security properties are validated at every stage, transforming the pipeline into a trusted boundary that mitigates risks associated with compromised build environments or malicious code injections.

Advanced dependency management strategies further fortify these pipelines by automating the generation and analysis of Software Bill of Materials (SBOMs). Automated tools now facilitate real-time tracking of transitive dependencies, enforcing policies that identify outdated libraries and known vulnerabilities without halting development velocity. The adoption of autonomous agents and machine learning models enables semantic vulnerability analysis, reducing false positives and providing explainable decisions for complex security scenarios. These systems continuously learn from simulated and real-world attack vectors, including configuration misconfigurations and privilege escalation attempts, to adaptively refine security postures across diverse infrastructure settings like

GitHub Actions and Jenkins.

Despite these advancements, significant challenges remain in aligning technical controls with regulatory frameworks such as GDPR, HIPAA, and the EU AI Act. Many existing implementations lack comprehensive audit logging and model explainability features required for full compliance, creating gaps in accountability. Furthermore, the reliance on AI-generated configuration files introduces new attack surfaces where traditional human review processes are bypassed. Future architectures must therefore prioritize end-to-end traceability and immutable logging to satisfy both security mandates and legal obligations, ensuring that automated remediation processes do not inadvertently introduce biases or overlook critical context-specific risks in the software lifecycle.

5 Future Directions

Despite significant advancements in AI-driven vulnerability detection, several critical limitations hinder the widespread adoption of these technologies in real-world software supply chains. Current deep learning architectures, particularly Graph Neural Networks and Transformer-based encoders, often struggle with the computational overhead required to process massive, interconnected codebases, making their integration into continuous integration pipelines challenging. Furthermore, the field continues to grapple with severe data scarcity and class imbalance, where vulnerable code snippets are vastly outnumbered by benign samples, leading to models that suffer from high false-negative rates. Existing approaches also lack robustness against adversarial evasion techniques, where subtle code obfuscation can easily bypass detection mechanisms. Finally, while Large Language Models show promise, their reliance on prompt engineering often results in inconsistent performance across diverse programming languages and complex, multi-file dependency structures, limiting their reliability for automated security auditing.

Future research must prioritize the development of lightweight, scalable architectures that maintain high semantic fidelity without prohibitive resource costs. This includes exploring hybrid models that combine the structural awareness of graph-based methods with the contextual efficiency of attention mechanisms, optimized specifically for incremental analysis in dynamic development environ-

ments. Simultaneously, there is an urgent need for advanced data synthesis techniques, such as generative adversarial networks tailored for code, to create high-quality, diverse vulnerability datasets that address class imbalance and reduce reliance on scarce manual annotations. Investigating self-supervised learning paradigms that can effectively leverage unlabeled open-source repositories will further enhance model generalization across different coding styles and languages. Additionally, future work should focus on adversarial robustness by incorporating defensive training strategies that expose models to a wider range of obfuscation and evasion tactics, ensuring resilience against sophisticated attacks.

Addressing these challenges will fundamentally transform the landscape of software security, enabling the deployment of proactive, intelligent defense systems capable of securing modern digital infrastructure. By overcoming current limitations in scalability and data efficiency, next-generation tools will be able to perform real-time vulnerability scanning within complex supply chains, significantly reducing the window of exposure for zero-day exploits. The successful integration of robust, adaptive AI models into standard development workflows will not only automate tedious security audits but also empower developers to identify and mitigate risks earlier in the lifecycle. Ultimately, these advancements will foster a more resilient software ecosystem, minimizing the economic and societal impacts of cyberattacks while accelerating the pace of safe innovation in an increasingly interconnected world.

References

- [1] A Benchmark Comparison of Python Malware Detection Approaches
- [2] Machine Learning Techniques for Python Source Code Vulnerability Detection
- [3] Data and Context Matter Towards Generalizing AI-based Software Vulnerability Detection
- [4] Exploring the Vulnerabilities of Machine Learning and Quantum Machine Learning to Adversarial Attacks using a Malware Dataset A Comparative Analysis
- [5] SHERLOCK A Deep Learning Approach To Detect Software Vulnerabilities
- [6] AI-Based Software Vulnerability Detection A Systematic Literature Review
- [7] Silent Vulnerable Dependency Alert Prediction with Vulnerability Key Aspect Explanation
- [8] SafePyScript A Web-Based Solution for Machine Learning-Driven Vulnerability Detection in Python
- [9] Classifying Issues in Open-source GitHub Repositories
- [10] Vulnerability Detection with Fine-Grained Interpretations
- [11] VulEval Towards Repository-Level Evaluation of Software Vulnerability Detection
- [12] An Analysis of Malicious Packages in Open-Source Software in the Wild
- [13] Detecting Malicious Source Code in PyPI Packages with LLMs Does RAG Come in Handy
- [14] Towards Measuring Supply Chain Attacks on Package Managers for Interpreted Languages
- [15] SourceBroken A large-scale analysis on the (un)reliability of SourceRank in the PyPI ecosystem
- [16] Taint-Based Code Slicing for LLMs-based Malicious NPM Package Detection
- [17] ENHANCING SOFTWARE SUPPLY CHAIN RESILIENCE STRATEGY FOR MITIGATING SOFTWARE SUPPLY-CHAIN SECURITY RISKS AND ENSURING SECURITYCONTINUITY IN DEVELOPMENT LIFECYCLE
- [18] A Review of Attacks Against Language-Based Package Managers
- [19] An Empirical Study of Malicious Code In PyPI Ecosystem
- [20] Cutting the Gordian Knot Detecting Malicious PyPI Packages via a Knowledge-Mining Framework
- [21] A Machine Learning-Based Approach For Detecting Malicious PyPI Packages
- [22] CHASE LLM Agents for Dissecting Malicious PyPI Packages
- [23] Enhancing the Insertion of NOP Instructions to Obfuscate Malware via Deep Reinforcement Learning
- [24] Utilization of machine learning for the detection of self-admitted vulnerabilities
- [25] Towards a Benchmark for Dependency Decision-Making
- [26] An Empirical Study on Software Bill of Materials Where We Stand and the Road Ahead
- [27] A Reality Check on SBOM-based Vulnerability Management An Empirical Study and A Path Forward

- [28] Code Ownership in Open-Source AI Software Security
- [29] A Survey on Automated Software Vulnerability Detection Using Machine Learning and Deep Learning
- [30] Cascaded Vulnerability Attacks in Software Supply Chains