A Survey of Deep Learning for Financial Risk Modeling under Uncertainty

InteractiveSurvey

Abstract

The contemporary financial landscape is increasingly defined by profound uncertainty and non-linear systemic interdependencies that render traditional econometric frameworks inadequate for capturing complex crisis dynamics. Consequently, there is an urgent paradigm shift toward data-driven methodologies capable of learning intricate patterns from high-dimensional datasets to ensure robust risk management. This survey critically examines the transformative role of deep learning architectures in financial risk modeling, specifically focusing on how advanced neural topologies, generative paradigms, and reinforcement learning strategies address unique challenges such as tail risk, distributional shifts, and systemic fragility. The paper synthesizes recent advancements to elucidate how these technologies move beyond mere prediction accuracy to provide coherent, interpretable, and regulation-compliant risk metrics. Key findings highlight the efficacy of Graph Neural Networks in modeling contagion effects, the utility of diffusion models for simulating rare catastrophic scenarios without traditional computational burdens, and the capacity of adversarial frameworks to align predictions with extreme event measures. Furthermore, the review demonstrates that integrating risk-aware reinforcement learning enables dynamic hedging strategies that adapt prudently to non-stationary market regimes. By offering a unified taxonomy of these disparate techniques, this work bridges the gap between theoretical artificial intelligence innovations and practical quantitative risk management requirements. Ultimately, this survey serves as a definitive reference for leveraging next-generation deep learning to enhance the resilience of financial systems against unforeseen shocks.

1 Introduction

The contemporary financial landscape is increasingly characterized by profound uncertainty, driven by rapid market evolution, geopolitical instability, and the emergence of non-linear systemic interdependencies. Traditional econometric frameworks, while foundational, often rely on restrictive linear assumptions and Gaussian distributions that fail to capture the heavy tails, volatility clustering, and complex contagion mechanisms inherent in modern crises. As financial instruments grow in complexity and data availability expands exponentially, the limitations of static models in addressing dynamic risk exposure have become apparent. Consequently, there is an urgent paradigm shift toward data-driven methodologies capable of learning intricate patterns from high-dimensional datasets, offering a more robust foundation for forecasting, stress testing, and hedging in volatile environments.

This survey paper critically examines the transformative role of deep learning architectures in financial risk modeling under conditions of severe uncertainty [1]. Unlike previous reviews that broadly cover algorithmic trading or general forecasting, this work specifically focuses on how advanced neural topologies, generative paradigms, and reinforcement learning strategies address the unique challenges of tail risk, distributional shifts, and systemic fragility. By synthesizing recent advancements, the paper elucidates how these technologies move beyond mere prediction accuracy to provide coherent, interpretable, and regulation-compliant risk metrics. The central thesis posits that integrating deep learning with rigorous risk theory offers a superior framework for navigating the stochastic nature of global markets compared to conventional statistical approaches.

The discussion begins by exploring advanced neural topologies designed to capture temporal

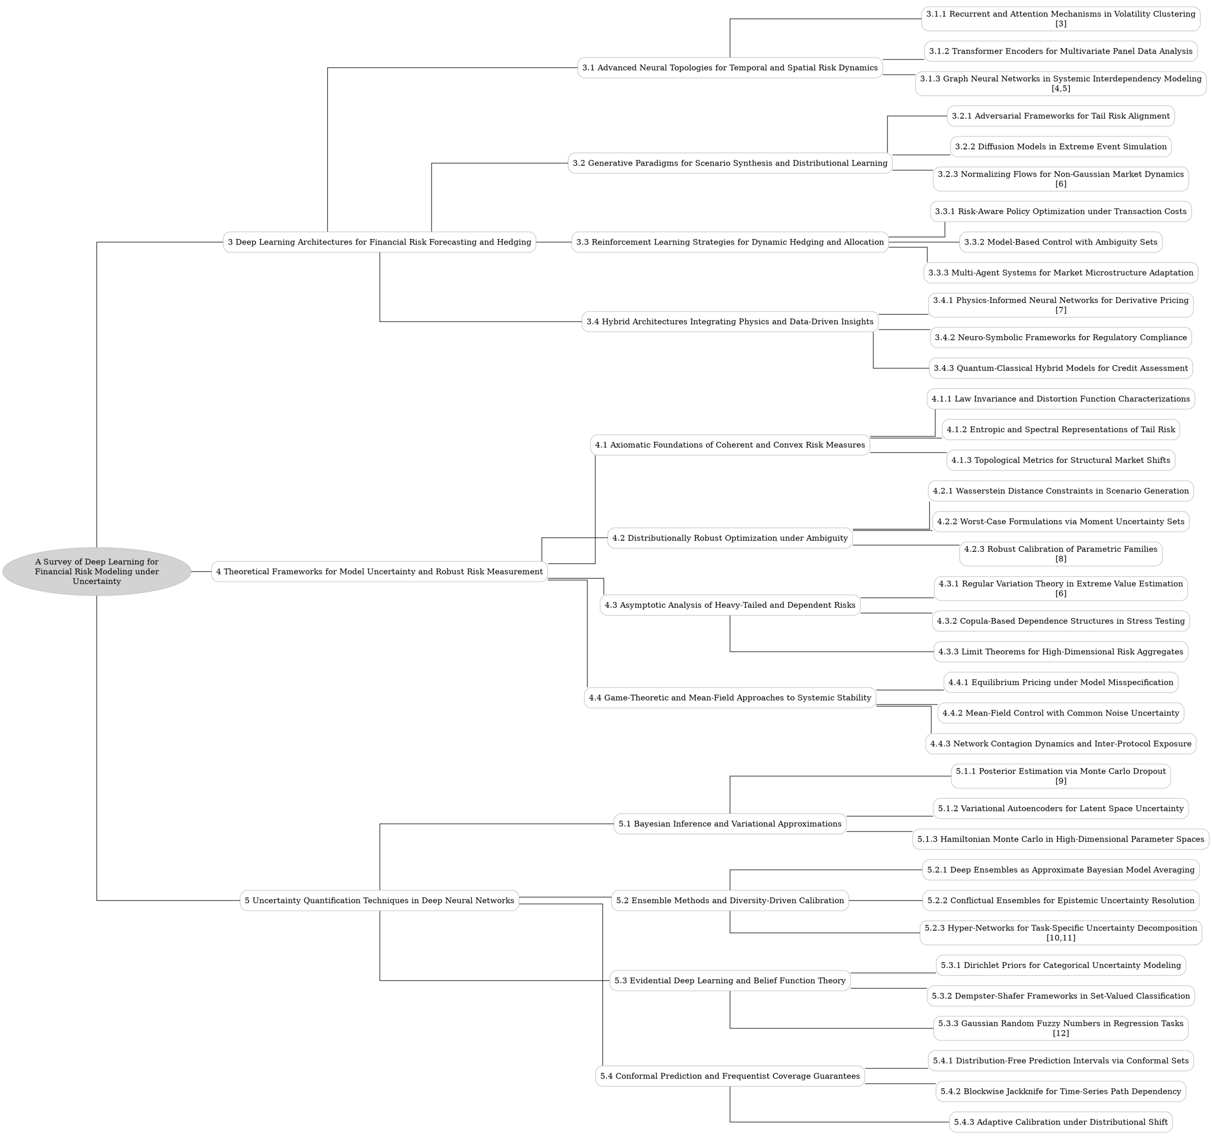

Figure 1: The outline of our survey: A Survey of Deep Learning for Financial Risk Modeling under Uncertainty

$$L = \frac{1}{T}\sum_{t=1}^{T}\left(R_t - \hat{R}_t\right)^2$$

Figure 2: Chart from 'Optimization Method of Multi-factor Investment Model Driven by Deep Learning for Risk Control' (arXiv:2507.00332)

and spatial risk dynamics. Recurrent networks and attention mechanisms are analyzed for their efficacy in modeling volatility clustering and long-range dependencies, overcoming the vanishing gradient problems of earlier iterations. Furthermore, the survey details the application of Transformer encoders in multivariate panel data analysis, highlighting their ability to process parallel sequences and identify latent cross-sectional correlations across diverse entities. Special emphasis is placed on Graph Neural Networks, which uniquely model systemic interdependencies by representing financial ecosystems as dynamic graphs, thereby enabling the simulation of contagion effects and cascading failures that isolated time-series models frequently miss.

Subsequently, the paper investigates generative paradigms essential for scenario synthesis and distributional learning. Adversarial frameworks are reviewed for their capacity to align model predictions with tail risk measures through minimax optimization, ensuring robustness against extreme events. The survey also examines diffusion models as a transformative tool for simulating rare, catastrophic scenarios without the computational burdens of traditional Monte Carlo methods. Additionally, normalizing flows are discussed for their ability to map complex, non-Gaussian market dynamics to tractable base densities, facilitating precise likelihood estimation and the modeling of skewed, multimodal return distributions that violate standard parametric assumptions.

The final technical section addresses reinforcement learning strategies for dynamic hedging and capital allocation [2]. This segment focuses on risk-aware policy optimization, where coherent risk measures are embedded directly into objective functions to mitigate tail exposure while accounting for transaction costs. The interplay between exploration and exploitation in non-stationary environments is scrutinized, demonstrating how agents can learn prudent strategies that adapt to shifting market regimes. Together, these three thematic pillars provide a comprehensive taxonomy of current methodologies, bridging the gap between theoretical deep learning innovations and practical quantitative risk management requirements.

$$\Phi : \mathcal{S} \to \mathcal{P}_p(\mathbb{R})$$

Figure 3: Chart from 'Distributional Reinforcement Learning on Path-dependent Options' (arXiv:2507.12657)

The primary contribution of this survey lies in its unified synthesis of disparate deep learning techniques specifically tailored for uncertainty quantification in finance [1]. By categorizing architectures based on their functional roles in handling temporal dynamics, generative simulation, and decision-making under risk, this work offers a structured roadmap for researchers and practitioners. It critically evaluates the trade-offs between model complexity, interpretability, and computational efficiency, identifying key gaps where current methods fall short in regulatory contexts. Ultimately, this paper serves as a definitive reference for understanding how next-generation artificial intelligence can enhance the resilience of financial systems against unforeseen shocks.

2 Deep Learning Architectures for Financial Risk Forecasting and Hedging

2.1 Advanced Neural Topologies for Temporal and Spatial Risk Dynamics

2.1.1 Recurrent and Attention Mechanisms in Volatility Clustering

Recurrent neural networks, particularly Long Short-Term Memory (LSTM) architectures, have become foundational in modeling volatility clustering by explicitly capturing the long-term temporal dependencies inherent in financial time series. Unlike traditional econometric models that rely on linear assumptions regarding past error terms, LSTMs utilize sophisticated gating mechanisms to retain information over extended sequences, effectively mirroring the phenomenon where large fluctuations are succeeded by further large fluctuations [3]. This memory capability allows the model to dynamically adjust risk estimates based on complex, non-linear historical patterns, offering a robust alternative to standard GARCH-family specifications for characterizing time-varying variance.

$$L = \frac{1}{N}\sum_{i=1}^{N}\sum(y_i - \hat{y}_i)^2$$

Figure 4: Chart from 'Analysis of Financial Risk Behavior Prediction Using Deep Learning and Big Data Algorithms' (arXiv:2410.19394)

To address the computational bottlenecks and vanishing gradient issues associated with recurrent structures, recent methodologies have increasingly integrated self-attention mechanisms derived from Transformer architectures. These attention-based models compute weighted relationships between all time steps simultaneously, enabling the direct capture of long-range dependencies without the sequential processing constraints of RNNs. By focusing on relevant historical volatility regimes regardless of their temporal distance, attention mechanisms provide a more efficient means of identifying structural breaks and persistent clustering effects, often outperforming recurrent counterparts in scenarios requiring the synthesis of vast historical contexts.

The convergence of these approaches has led to hybrid frameworks that leverage the sequential feature extraction of recurrent layers alongside the global context awareness of attention modules. Such architectures are particularly effective in distributional forecasting, where capturing the full uncertainty structure of market dynamics is critical for accurate Value-at-Risk and Expected Shortfall estimation. While CNNs offer computational advantages for high-frequency data, the superior performance of recurrent and attention-based models underscores the importance of explicit memory and dependency modeling in resolving the intricate, non-linear dynamics of financial volatility clustering.

2.1.2 Transformer Encoders for Multivariate Panel Data Analysis

Transformer encoders have emerged as a pivotal architecture for analyzing multivariate panel data, primarily due to their self-attention mechanisms which effectively capture long-term dependencies and complex cross-sectional interactions that traditional recurrent models often miss. By processing entire sequences in parallel, these models mitigate the vanishing gradient problem inherent in RNNs, allowing for the simultaneous modeling of temporal dynamics across multiple entities. This capability is particularly critical in financial contexts where latent correlations between assets or firms evolve over time, requiring an architecture that can weigh historical information adaptively without being constrained by fixed window sizes or sequential computation bottlenecks.

In the specific context of panel data, transformer encoders are frequently adapted through masking strategies and positional embeddings to handle irregular sampling and entity-specific characteristics. Techniques such as random masking of input features during pre-training force the model to learn robust representations by reconstructing missing values based on global context, thereby enhancing generalization across diverse economic regimes. Furthermore, the integration of entity-specific embeddings allows the model to distinguish between individual units within the panel while sharing parameters across the population, facilitating the extraction of both idiosyncratic patterns and common macroeconomic factors that drive systemic risk or market movements.

The application of these encoders extends to high-dimensional forecasting tasks where the interplay between numerous variables necessitates a flexible, non-linear approach. By leveraging multi-head attention, the architecture can simultaneously attend to different representation subspaces, capturing distinct types of relationships such as momentum effects, mean reversion, and volatility clustering within the same framework. When fine-tuned for downstream tasks like risk estimation or asset pricing, transformer-based models demonstrate superior performance in capturing tail risks and distributional shifts, offering a significant advantage over linear factor models and shallow neural networks in environments characterized by structural breaks and non-stationarity.

2.1.3 Graph Neural Networks in Systemic Interdependency Modeling

Graph Neural Networks (GNNs) have emerged as a pivotal architecture for modeling systemic interdependencies within financial ecosystems, addressing the limitations of traditional models that often treat entities as independent [4]. By representing firms, banks, or assets as nodes and their financial relationships as edges, GNNs explicitly capture the complex relational structures inherent in supply chains, interbank lending networks, and sectoral correlations [4]. This graph-based formulation allows for the aggregation of neighborhood information through message-passing mechanisms, enabling the model to learn how distress propagates from a keystone firm to connected counterparts, thereby facilitating a more accurate assessment of collective risk exposure and cascading failures.

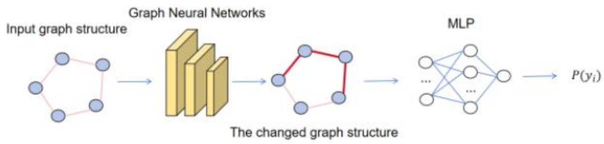

Figure 5: Chart from 'Robust Graph Neural Networks for Stability Analysis in Dynamic Networks' (arXiv:2411.11848)

Unlike standard deep learning approaches that rely solely on isolated time-series data, GNNs integrate topological features with node-specific attributes to model non-linear contagion dynamics. The architecture leverages spatial dependencies to identify systemic vulnerabilities that arise from network centrality and connectivity patterns, which are frequently ignored by autoregressive or regression-based methods. Through iterative layers of graph convolution or attention, these networks encode higher-order interactions, effectively mapping the transmission of shocks across the entire financial graph. This capability is particularly critical for stress testing scenarios where the failure of a single entity triggers a domino effect, requiring a holistic view of the network rather than fragmented individual risk metrics.

Recent advancements have further refined GNN applications by incorporating temporal dynamics to handle evolving market conditions and regime shifts. Hybrid architectures combining graph convolutions with recurrent units or attention mechanisms allow for the simultaneous modeling of static network structures and dynamic temporal dependencies. These models provide a robust framework for predicting systemic risk metrics, such as Value-at-Risk and Expected Shortfall, under varying economic regimes [5]. By capturing both the structural integrity of the financial network and the temporal evolution of asset correlations, GNNs offer a sophisticated tool for regulators and practitioners to anticipate and mitigate systemic instability in increasingly interconnected markets.

$$L_{forecast} = \| y - \hat{y} \|^2 + \sum_{\tau} \rho\tau(y - \hat{y}),$$

Figure 6: Chart from 'Uni-FinLLM: A Unified Multimodal Large Language Model with Modular Task Heads for Micro-Level Stock Prediction and Macro-Level Systemic Risk Assessment' (arXiv:2601.02677)

2.2 Generative Paradigms for Scenario Synthesis and Distributional Learning

2.2.1 Adversarial Frameworks for Tail Risk Alignment

Adversarial frameworks have emerged as a pivotal mechanism for aligning neural network predictions with tail risk measures, addressing the inherent tendency of standard models to underestimate extreme events. By formulating risk estimation as a minimax game, these approaches pit a generator against a discriminator or adversary designed to stress-test the model under worst-case scenarios. This dynamic forces the primary network to learn robust representations that remain accurate even in the heavy-tailed regions of the loss distribution, effectively mitigating the fragility often observed in conventional deep learning architectures when faced with rare but catastrophic financial shocks.

The integration of adversarial training specifically targets the optimization of coherent risk measures such as Expected Shortfall and Value-at-Risk. Unlike traditional mean-squared error objectives that prioritize average performance, adversarial losses explicitly penalize deviations in the lower quantiles, ensuring the model captures the asymmetry and kurtosis characteristic of financial returns. Techniques such as distributionally robust optimization are employed to shift the training focus toward adversarial perturbations that maximize risk exposure, thereby compelling the network to internalize the structural dependencies driving tail events rather than merely fitting the central tendency of the data.

Furthermore, these frameworks enhance model reliability by providing implicit regularization against distributional shifts and input manipulations common in volatile markets. By continuously challenging the predictor with synthesized adverse scenarios, adversarial methods reduce the variance of risk estimates and improve out-of-sample stability. This approach not only refines the accuracy of tail risk forecasts but also offers a principled method for quantifying epistemic uncertainty, ensuring that the resulting risk metrics are conservative enough to satisfy regulatory requirements while maintaining computational efficiency in high-dimensional portfolio settings.

2.2.2 Diffusion Models in Extreme Event Simulation

Diffusion models have emerged as a transformative approach for simulating extreme financial events, addressing the inherent limitations of traditional nested Monte Carlo methods which are often computationally prohibitive for practical risk horizons. Unlike conventional techniques that rely on rigid parametric assumptions or linear approximations such as the delta-gamma method, diffusion-based frameworks learn the underlying data distribution directly from historical trajectories. This capability allows them to accurately capture heavy-tailed characteristics and complex nonlinear dependencies that define market stress periods, effectively modeling the scarcity of data associated with rare, catastrophic losses without succumbing to the convergence issues plaguing standard simulation engines.

The technical implementation of these models typically involves formulating the evolution of risk factors as a stochastic differential equation, where the reverse diffusion process reconstructs plausible extreme scenarios from Gaussian noise. By conditioning the generation process on specific market states or latent variables, these simulators can produce high-fidelity out-of-sample trajectories that reflect the statistical dispersion of channel fading or asset returns during crisis conditions. This generative mechanism is particularly effective in high-dimensional settings, where it leverages deep neural networks to approximate the score function, thereby enabling the efficient sampling of joint distributions for multiple correlated assets while preserving the tail properties essential for robust Value-at-Risk and Expected Shortfall calculations.

Furthermore, diffusion models offer significant advantages in quantifying uncertainty and enhancing the explainability of risk predictions through their probabilistic nature. The iterative denoising process inherently provides a measure of prediction confidence, allowing practitioners to distinguish between stable market regimes and volatile periods characterized by high epistemic uncertainty. When integrated with agent-based modeling or used to calibrate intensity models for default events, diffusion simulators facilitate a more nuanced understanding of systemic risk propagation. This leads to improved decision-making frameworks that are not only sensitive to high-order statistical behaviors but also robust against the model misspecification errors that frequently undermine traditional extreme value theory approaches during unprecedented market shocks.

2.2.3 Normalizing Flows for Non-Gaussian Market Dynamics

Traditional financial modeling often relies on Gaussian assumptions to describe asset returns and interest rate dynamics, yet empirical evidence consistently reveals heavy tails, skewness, and multimodal regimes that violate these premises [6]. While standard stochastic differential equations capture linear dependencies and instantaneous correlations effectively, they frequently fail to account for the non-linear, chaotic nature of market shocks. Consequently, risk measures derived from such models may underestimate tail risks, leading to suboptimal hedging strategies and inaccurate pricing of derivatives in volatile environments.

$$\frac{1}{\sqrt{T}} \sum_{t=1}^{T} s(r_t; \boldsymbol{\psi}_0) \xrightarrow{d} N(0, I(\boldsymbol{\psi}_0))$$

Figure 7: Chart from 'Probability Weighting Meets Heavy Tails: An Econometric Framework for Behavioral Asset Pricing' (arXiv:2511.16563)

Normalizing flows address these limitations by constructing flexible, invertible transformations that map complex, non-Gaussian market distributions to simple base densities, such as standard Gaussians. By leveraging the change of variables formula, these models preserve probabilistic coherence while allowing for the explicit computation of likelihoods, enabling precise density estimation without resorting to restrictive parametric forms. This architecture facilitates the modeling of intricate dependency structures among multi-dimensional factors, including asset prices, interest rates, and mortality intensities, capturing higher-order moments that traditional linear models overlook.

In practice, integrating normalizing flows into quantitative frameworks allows for the direct optimization of risk-sensitive objectives, such as Conditional Value-at-Risk, under realistic distributional assumptions. The ability to learn conditional distributions dynamically enables the model to adapt to shifting market regimes and incorporate alternative data sources effectively. Furthermore, when combined with variational inference techniques, multiplicative normalizing flows provide robust uncertainty quantification, ensuring that posterior estimates reflect both epistemic and aleatoric uncertainty inherent in financial time series, thereby enhancing the reliability of downstream investment decisions.

2.3 Reinforcement Learning Strategies for Dynamic Hedging and Allocation

2.3.1 Risk-Aware Policy Optimization under Transaction Costs

Risk-aware policy optimization in financial markets necessitates the integration of coherent risk measures, such as Conditional Value-at-Risk (CVaR), directly into the objective function to mitigate tail exposure. Unlike traditional mean-variance frameworks, this approach prioritizes the minimization of extreme losses while accounting for the non-linear dynamics of asset returns. By employing policy gradient methods tailored for coherent risk measures, algorithms can learn strategies that are robust against distributional shifts, ensuring that capital allocation decisions remain prudent even under adverse market conditions. This formulation is particularly critical for regulatory compliance, where institutions must compute daily risk metrics to determine adequate capital reserves.

The inclusion of transaction costs fundamentally alters the optimization landscape, introducing friction that can erode profits and induce suboptimal trading frequencies if ignored. Effective policies must balance the trade-off between risk reduction and the cumulative impact of proportional, per-share, and fixed costs associated with order execution. Advanced frameworks incorporate these costs directly into the reward signal or constraint set, often utilizing hysteresis mechanisms to prevent excessive churn near decision boundaries. This ensures that portfolio rebalancing occurs only when the marginal benefit of risk adjustment outweighs the immediate liquidity drag and spread costs, thereby preserving net returns in illiquid environments.

To solve these complex, high-dimensional control problems, recent methodologies leverage neural network approximations of the conditional valuation functional combined with simulation-based techniques. These models utilize uncertainty quantification, derived from predictive distributions or ensemble methods, to dynamically adjust exposure without relying on deterministic assumptions that fail to capture market volatility. By training agents within realistic simulators that replicate bid-ask spreads and liquidity constraints, the resulting policies demonstrate superior out-of-sample performance. This data-driven approach enables the systematic design of ultra-reliable strategies that seamlessly interpolate between robust worst-case scenarios and ergodic average settings, effectively bridging the gap between theoretical risk management and practical execution.

2.3.2 Model-Based Control with Ambiguity Sets

Model-based control strategies incorporating ambiguity sets address the critical challenge of distributional uncertainty by optimizing control policies against a worst-case scenario within a defined family of probability measures. Rather than relying on a single nominal model, this framework constructs an ambiguity set, often centered around an empirical distribution derived from historical time-series data, to encompass plausible deviations in system dynamics. The radius of this set is carefully calibrated to balance robustness and conservatism, ensuring that the controller remains effective even when the true data-generating process diverges from the training distribution due to non-stationarity or limited sample sizes.

The integration of neural networks into this paradigm allows for the approximation of complex, high-dimensional value functions and decision boundaries that traditional parametric models cannot capture. By embedding non-crossing constraints or penalty terms directly into the loss function, these architectures ensure that risk metrics, such as Value-at-Risk across different confidence levels, maintain logical consistency during optimization. Furthermore, the use of primal-dual algorithms facilitates the handling of hard constraints on state feasibility and control inputs, enabling the system to learn robust policies that implicitly account for both aleatoric noise and epistemic model uncertainty without requiring explicit analytical forms of the underlying stochastic processes.

Advanced implementations often leverage dynamic programming principles to bound the output ranges of these neural controllers over specific input regions, providing statistical guarantees on closed-loop safety and performance. The determination of the ambiguity set can be entirely data-driven, utilizing finite differences or Bayesian updating mechanisms to sequentially refine the set of physical measures as new observations become available. This adaptive approach ensures that the control policy evolves in response to changing market regimes or environmental conditions, ultimately yielding a decision-making framework that is resilient to model misspecification while maintaining computational tractability for real-time applications in autonomous systems and quantitative finance.

2.3.3 Multi-Agent Systems for Market Microstructure Adaptation

Multi-agent systems (MAS) have emerged as a pivotal framework for modeling market microstructure by simulating the bottom-up interactions of heterogeneous, autonomous entities. Unlike traditional top-down econometric approaches, agent-based models (ABM) capture emergent phenomena such as liquidity crises, flash crashes, and price spikes through the dynamic interplay of competing trading strategies. These artificial markets facilitate high-fidelity in-silico trials, allowing researchers to test counterfactual theories and observe how local decision-making rules aggregate into complex global market dynamics without relying on restrictive equilibrium assumptions.

The adaptation of market microstructure within these systems relies on the continuous evolution of agent behaviors in response to changing order book conditions and regulatory constraints. By embedding diverse risk measures and learning algorithms, MAS can replicate the non-linear and heteroscedastic traits inherent in financial time series, effectively modeling the transition between bull and bear regimes. This flexibility enables the simulation of high-frequency trading impacts on market fragility, providing a robust environment to analyze how specific microstructural features, such as tick size or latency, influence overall market stability and efficiency under stress.

Despite their descriptive power, deploying MAS for real-time adaptation presents significant computational challenges and risks of overfitting to specific historical scenarios. The complexity of calibrating thousands of interacting agents requires advanced optimization techniques, often necessitating multi-start procedures or scenario-based approximations to ensure robustness. Future advancements aim to integrate deep learning architectures, such as Transformers, within agent decision loops to enhance predictive accuracy and generalization, thereby bridging the gap between theoretical simulation and practical algorithmic trading strategy development in volatile, high-dimensional asset universes.

2.4 Hybrid Architectures Integrating Physics and Data-Driven Insights

2.4.1 Physics-Informed Neural Networks for Derivative Pricing

Physics-Informed Neural Networks (PINNs) have emerged as a transformative approach for derivative pricing, addressing the limitations of traditional numerical methods in high-dimensional settings [7]. By embedding the governing partial differential equations, such as the Black-Scholes equation, directly into the loss function, PINNs constrain the neural network to satisfy fundamental financial laws without relying on extensive labeled data. This formulation leverages the universal approximation capabilities of deep feedforward networks to solve forward and backward stochastic differential equations, offering a mesh-free alternative to finite difference methods that scales efficiently with the number of underlying risk factors.

$$p(\theta|\mathrm{X},\mathrm{Y}) = \frac{p(\mathrm{Y}|\mathrm{X},\theta)p(\theta)}{p(\mathrm{Y}|\mathrm{X})}$$

Figure 8: Chart from 'Uncertainty Quantification of Deep Learning for Spatiotemporal Data: Challenges and Opportunities' (arXiv:2311.02485)

The architecture typically utilizes automatic differentiation to compute the residual of the pricing PDE at collocation points sampled across the spatial and temporal domains. During training, the optimization algorithm minimizes a composite loss function comprising the PDE residual, initial conditions, and boundary constraints, effectively forcing the network to learn the solution manifold consistent with no-arbitrage principles. This data-efficient paradigm is particularly advantageous for exotic derivatives where analytical solutions are unavailable and Monte Carlo simulations suffer from slow convergence rates or high variance in estimating Greeks.

Despite their theoretical elegance, practical implementation faces challenges related to training stability and the balance between competing loss terms. The non-convex nature of the optimization landscape can lead to convergence issues, necessitating advanced optimizers and careful hyperparameter tuning to ensure the network captures sharp gradients near exercise boundaries. Nevertheless, recent advancements demonstrate that PINNs provide robust, differentiable surrogate models capable of delivering accurate valuations and hedging ratios simultaneously, marking a significant shift towards integrating physical constraints into deep learning frameworks for quantitative finance.

2.4.2 Neuro-Symbolic Frameworks for Regulatory Compliance

Neuro-symbolic frameworks address the critical tension between the predictive accuracy of deep learning and the interpretability mandates of financial regulations like Basel III. By integrating neural networks with symbolic logic, these systems embed regulatory constraints directly into the model architecture or loss functions, ensuring that risk metrics such as Expected Shortfall adhere to coherent measure axioms. This hybrid approach allows for the modeling of complex, non-linear market dynamics while maintaining the transparency required for regulatory audits and stress testing scenarios.

A primary technical advantage lies in the ability to encode domain-specific rules, such as counterparty credit risk limits or capital adequacy ratios, as differentiable logical constraints. Unlike black-box models that may produce erratic outputs under distributional shifts, neuro-symbolic architectures utilize symbolic reasoning to validate neural predictions against established financial laws. This mechanism prevents the model from converging to spurious correlations that violate fundamental economic principles, thereby enhancing robustness against adversarial perturbations and ensuring consistent performance across varying risk levels.

Furthermore, these frameworks facilitate explainable decision-making by generating human-readable justifications for capital allocation and risk assessment outcomes. Through techniques like neural theorem proving or rule extraction, the system can trace specific regulatory violations back to input features, satisfying the "right to explanation" inherent in modern compliance standards. Consequently, neuro-symbolic approaches offer a viable pathway for deploying advanced artificial intelligence in high-stakes environments where both statistical precision and strict adherence to legal frameworks are non-negotiable requirements.

2.4.3 Quantum-Classical Hybrid Models for Credit Assessment

Quantum-classical hybrid models represent a paradigm shift in credit assessment by leveraging quantum processing units to optimize high-dimensional feature spaces while retaining classical neural networks for non-linear pattern recognition. These architectures typically employ variational quantum circuits as trainable layers within deep learning frameworks, enabling the efficient encoding of complex financial correlations that often overwhelm classical solvers. By mapping credit risk factors onto quantum states, these models exploit superposition and entanglement to explore vast solution landscapes, offering potential exponential speedups in calculating default probabilities and loss distributions compared to traditional Monte Carlo simulations.

The integration mechanism frequently utilizes a hybrid loop where classical optimizers update quantum gate parameters to minimize a cost function derived from credit scoring metrics. This approach is particularly effective for handling the sparse and noisy data characteristic of financial markets, as quantum kernels can implicitly map input features into higher-dimensional Hilbert spaces where linear separability of risky and safe borrowers is enhanced. Furthermore, hybrid frameworks incorporate uncertainty quantification directly into the quantum ansatz, allowing for more robust estimation of tail risks and Value-at-Risk measures without the computational burden of extensive resampling techniques required by purely classical econometric models.

Despite theoretical advantages, practical deployment faces significant hurdles related to current hardware noise levels and the barren plateau problem in training deep quantum circuits. Recent advancements focus on error-mitigation strategies and shallow circuit designs tailored specifically for financial time-series data to ensure stability during the credit evaluation process. As quantum hardware matures, these hybrid systems are expected to evolve from proof-of-concept studies into essential tools for real-time credit monitoring, providing superior accuracy in distinguishing subtle default signals amidst market volatility while maintaining interpretability through classical post-processing layers.

3 Theoretical Frameworks for Model Uncertainty and Robust Risk Measurement

3.1 Axiomatic Foundations of Coherent and Convex Risk Measures

3.1.1 Law Invariance and Distortion Function Characterizations

Law invariance serves as a foundational axiom in risk measurement, asserting that the risk assessment of a random variable depends solely on its probability distribution rather than the specific underlying states of the world. Within the framework of sublinear expectations, this property facilitates the characterization of risk statistics through distortion functions, which effectively reweight the probability mass to account for model uncertainty and volatility ambiguity. When applied to G-normal distributions, law invariance ensures that the resulting risk measures remain consistent under distributional shifts, allowing for a robust evaluation of financial time series where classical probabilistic assumptions may fail due to non-linear dynamics.

The characterization of these risk measures via distortion functions provides a rigorous mechanism to translate complex uncertainty structures into tractable analytical forms. Specifically, a law-invariant risk statistic can often be represented as a Choquet integral with respect to a distorted probability measure, where the distortion function encapsulates the decision-maker's attitude toward tail events and extreme losses. In the context of robust Value-at-Risk under G-normality, the distortion function adap to the interval-valued volatility parameters, effectively bridging the gap between static distributional assumptions and dynamic volatility uncertainty. This approach allows for the derivation of explicit bounds on risk exposures that are invariant to the specific realization of the stochastic process, provided the distributional law remains unchanged.

Furthermore, the interplay between law invariance and distortion characterizations enables the decomposition of complex risk statistics into simpler, interpretable components. By leveraging clustering functions and acceptance sets, one can establish conditions under which a composite risk measure preserves law invariance, ensuring that the aggregation of risks across different dimensions does not introduce spurious dependencies. This theoretical structure is critical for validating the consistency of estimators in high-dimensional settings, where the convergence of empirical distributions to their theoretical counterparts must be guaranteed despite the presence of epistemic uncertainty. Ultimately, these characterizations provide the necessary mathematical groundwork for developing unified posterior frameworks that jointly model volatility, fraud, and compliance risks without compromising theoretical rigor.

3.1.2 Entropic and Spectral Representations of Tail Risk

Entropic representations provide a robust framework for quantifying tail risk under Knightian uncertainty by leveraging sublinear expectations to characterize families of probability measures. In this setting, the entropic risk measure is often exchanged for a mean-variance operator to enhance numerical stability, particularly when dealing with normal random variables. This approach allows regulators to model ambiguity where the true probability distribution remains indistinguishable among a set of candidates, effectively capturing the worst-case scenarios through nonlinear transforms of loss functions.

Spectral representations complement entropic methods by offering dual formulations that decompose risk into weighted integrals of quantiles, facilitating a more granular analysis of the loss distribution's tail. Recent advancements integrate these spectral views with natural risk statistics, deriving explicit dual representations that accommodate complex ambiguity sets. By utilizing specific cost functions and identifying extreme points within the ambiguity set, these representations enable the construction of uncertainty-aware policies that analytically adjust exposure based on estimated parameters and sample sizes, thereby mitigating the impact of parameter uncertainty on tail risk estimates.

The convergence properties of these representations are critical for statistical inference, with the rates of convergence for laws of large numbers and central limit theorems under sublinear expectations playing a pivotal role. Nonlinear generalizations of the Stein method serve as sharp tools for analyzing these convergence rates, ensuring that approximations remain valid even in the presence of structural breaks or skewed distributions. Consequently, the synthesis of entropic and spectral approaches provides a comprehensive mechanism for managing model risk, allowing for dynamic adjustments that improve tail-risk monitoring during stressed market periods while maintaining consistency with axiomatic foundations.

3.1.3 Topological Metrics for Structural Market Shifts

Topological metrics provide a rigorous framework for detecting structural market shifts by analyzing the evolving geometry of asset correlation networks. Unlike traditional linear measures, these approaches utilize persistent homology to quantify the birth and death of topological features, such as connected components and loops, within high-dimensional financial data. During periods of market stability, the network typically exhibits a complex structure with multiple distinct clusters; however, impending crises often manifest as a topological phase transition where the system collapses into a highly connected, star-like graph, signaling a loss of diversification benefits and a surge in systemic fragility.

The application of these metrics extends to characterizing the robustness of market structures under parametric uncertainty and ambiguous volatility regimes. By modeling the market as a dynamic simplicial complex, researchers can track the persistence of specific topological invariants that precede regime changes, offering early warning signals that Value-at-Risk models frequently miss. This method effectively captures non-linear dependencies and tail risks, as the topological distance between consecutive market states serves as a sensitive indicator of structural instability, reflecting the propagation of shocks across interconnected financial domains more accurately than pairwise correlation analysis.

Furthermore, integrating topological data analysis with risk-return clustering allows for the disentanglement of transient noise from fundamental structural breaks. The resulting metrics facilitate the construction of adaptive trading strategies that respond to the intrinsic dimensionality of the market rather than static historical parameters. As the market topology shifts, indicating a transition from a calm to a volatile state, these measures enable proactive adjustments in portfolio allocation and compliance monitoring. Consequently, topological metrics serve as a critical tool for managing complex market environments, ensuring that risk management frameworks remain resilient against the rapid evolution of financial interconnectedness.

3.2 Distributionally Robust Optimization under Ambiguity

3.2.1 Wasserstein Distance Constraints in Scenario Generation

Wasserstein distance constraints have emerged as a pivotal mechanism in scenario generation, particularly within distributionally robust optimization frameworks designed to mitigate model risk. By defining an ambiguity set as a Wasserstein ball centered around the empirical distribution, these methods constrain the search for worst-case scenarios to a neighborhood of observed data, thereby balancing robustness with statistical fidelity. This approach effectively addresses the "optimization enigma" where traditional models overfit historical noise, ensuring that generated stress scenarios remain plausible while capturing tail risks that parametric assumptions might overlook.

In the context of financial time series, recent advancements integrate Wasserstein metrics with generative adversarial networks to produce conditional scenarios that preserve complex temporal dependencies and causal structures. Unlike moment-based uncertainty sets which often yield overly conservative bounds, Wasserstein constraints provide performance guarantees that scale gracefully with sample size, offering tractable reformulations for high-dimensional problems. The metric's ability to quantify the cost of transporting probability mass allows for a nuanced interpolation between historical realizations, generating synthetic paths that reflect both the central tendency and the extreme deviations inherent in volatile market regimes.

However, the practical implementation of Wasserstein-constrained scenario generation faces significant computational challenges, especially when applied to multiperiod optimization problems requiring adapted distances. The selection of the radius for the Wasserstein ball acts as a critical hyperparameter, governing the trade-off between robustness against distributional shifts and the conservatism of the resulting risk measures such as Value-at-Risk. While theoretical results demonstrate that these constraints yield asymptotically consistent estimators, future research must focus on scalable algorithms that can handle the curse of dimensionality in large-scale portfolio optimization without sacrificing the interpretability required for regulatory compliance and internal governance.

3.2.2 Worst-Case Formulations via Moment Uncertainty Sets

Worst-case formulations via moment uncertainty sets provide a rigorous framework for addressing distributional ambiguity in financial risk management without assuming a specific parametric form. Instead of relying on a single probability measure, this approach defines an ambiguity set comprising all distributions that satisfy known statistical constraints, typically limited to the first few moments such as mean and covariance. By optimizing against the most adverse distribution within this set, decision-makers derive robust policies that inherently hedge against model misspecification and data scarcity, effectively bridging the gap between stochastic programming and classical robust optimization.

The computational tractability of these formulations often relies on duality theory, which transforms the infinite-dimensional worst-case expectation problem into a finite-dimensional convex optimization program. When the uncertainty set is characterized by generalized moment constraints, the worst-case distribution is frequently supported on a small number of extreme points, allowing for closed-form solutions or efficient numerical approximations. This structural property is particularly valuable in portfolio optimization and capital requirement calculations, where it enables the explicit quantification of the impact of parameter uncertainty on risk metrics like Value-at-Risk and Expected Shortfall.

Furthermore, recent advancements have extended these moment-based approaches to incorporate higher-order moments and asymmetric information, refining the tightness of the resulting bounds. The integration of sublinear expectations offers a theoretical foundation for these methods, linking worst-case scenarios to nonlinear expectation operators that capture Knightian uncertainty. Consequently, moment uncertainty sets serve as a versatile tool for balancing conservatism with performance, ensuring that risk assessments remain reliable even when empirical data fails to fully characterize the underlying stochastic processes.

3.2.3 Robust Calibration of Parametric Families

Robust calibration of parametric families addresses the critical vulnerability of standard estimation techniques to model misspecification and data scarcity. Unlike classical methods that yield point estimates susceptible to overfitting, robust frameworks treat parameters as additional risk factors or intervals, effectively integrating out epistemic uncertainty. By replacing fixed values with distributions or bounded sets, this approach mitigates the adverse effects of aleatory noise and ensures that the calibrated model remains valid even when observed data deviates from idealized distributional assumptions, such as Gaussianity.

The methodology often employs weak optimal transport penalization or sublinear expectation structures to define a neighborhood around the nominal parametric family. In the context of G-normal distributions, calibration involves determining volatility bounds [8] rather than a single variance parameter, utilizing sublinear functions to capture the worst-case scenarios within these bounds. This formulation allows for the analytical derivation of robust risk measures, such as Value-at-Risk, by solving supremum-infimum problems over the admissible parameter space, thereby explicitly accounting for the uncertainty in correlation structures and higher-order moments without relying on specific dynamic models.

Furthermore, practical implementation frequently incorporates regularization techniques, such as L1-norm quantile regression, to enforce sparsity and stability in high-dimensional settings. The calibration process is iteratively refined using rolling residual-based adjustments to correct systematic miscalibration and maintain coverage levels consistent with posterior predictive intervals. By optimizing loss functions that are jointly consistent for multiple risk metrics, robust calibration ensures that the resulting parametric families provide reliable out-of-sample performance, balancing the trade-off between sharpness and reliability under significant distributional ambiguity.

3.3 Asymptotic Analysis of Heavy-Tailed and Dependent Risks

3.3.1 Regular Variation Theory in Extreme Value Estimation

Regular variation theory provides the foundational asymptotic framework for modeling the tails of financial return distributions, which frequently exhibit heavy-tailed behavior inconsistent with Gaussian assumptions [6]. By characterizing the decay rate of tail probabilities through a slowly varying function and a tail index, this theory enables the rigorous derivation of extreme value estimators. The core premise relies on the assumption that the underlying distribution belongs to the domain of attraction of an extreme value distribution, allowing for the extrapolation of risk metrics beyond the range of observed data.

In the context of parameter estimation, the theory facilitates the construction of robust statistics, such as the φ-max-mean estimator, which adapts classical methods to handle uncertainty in mean and variance parameters. These estimators leverage the max-stable property of limiting distributions to derive unbiased estimates for upper and lower expectations under G-normal frameworks. The integration of regular variation conditions ensures that the resulting estimators maintain consistency even when the sampling distribution exhibits significant skewness or kurtosis, thereby addressing the limitations of standard moment-based approaches in volatile market regimes.

Furthermore, the application of regular variation extends to the calibration of risk measures like Value-at-Risk and Conditional Value-at-Risk, where accurate tail index estimation is critical for capital adequacy. By employing techniques such as the Hill estimator or peaks-over-threshold methods grounded in regular variation, practitioners can quantify the probability of rare events with greater precision. This theoretical apparatus not only supports the development of adaptive penalization schemes for time-varying quantile regression but also underpins the dual representation of complex risk statistics, ensuring that systemic risk assessments remain robust against model misspecification in the extremes.

3.3.2 Copula-Based Dependence Structures in Stress Testing

Copula-based frameworks provide a rigorous mechanism for modeling complex dependence structures among risk factors, particularly when marginal distributions exhibit heavy tails or asymmetry during market turmoil. By decoupling marginal behavior from the joint dependence structure, these models allow stress testing procedures to isolate and manipulate correlation dynamics independently of individual asset volatilities. This separation is critical in crisis scenarios where traditional linear correlation measures fail to capture the surge in tail dependence, often referred to as correlation breakdown, which significantly amplifies systemic risk exposure.

In the context of stress testing, specific copula families, such as Archimedean or elliptical variants, are employed to simulate extreme comovements that deviate from historical norms. The methodology involves calibrating the copula parameter to reflect stressed states, effectively increasing the probability of simultaneous extreme losses across the portfolio. Unlike standard Gaussian assumptions that underestimate joint tail events, copula-based approaches can incorporate asymmetric dependence, ensuring that the generated stress scenarios accurately reflect the contagion effects observed during financial crises. This capability enables regulators and institutions to evaluate capital adequacy under conditions where dependencies become highly non-linear and concentrated in the lower tails of the distribution.

Furthermore, the integration of copula structures into stress testing facilitates robust sensitivity analysis regarding parameter uncertainty and model risk. By varying the copula family or its association parameters within a plausible range derived from bootstrap exercises or Bayesian posterior distributions, analysts can assess the stability of risk metrics like Value-at-Risk and Expected Shortfall. This approach addresses the non-stationarity of dependence structures, acknowledging that correlations fluctuate erratically over time and intensify during periods of acute stress. Consequently, copula-based stress testing offers a more comprehensive view of potential portfolio vulnerabilities, ensuring that capital buffers are sufficient to withstand severe, correlated shocks that simple variance-covariance models would likely overlook.

3.3.3 Limit Theorems for High-Dimensional Risk Aggregates

The aggregation of high-dimensional risk factors necessitates a rigorous theoretical foundation grounded in limit theorems under sublinear expectations. Unlike classical frameworks relying on additive probabilities, modern approaches address model uncertainty by establishing laws of large numbers and central limit theorems within sublinear expectation spaces. These results characterize the asymptotic behavior of sums of independent but not necessarily identically distributed random variables, ensuring that risk measures remain robust even when the underlying probability distribution is ambiguous or partially known.

Convergence rates and distributional approximations are critical for practical implementation in large-scale financial systems. Recent advancements provide explicit bounds on the rate of convergence for these limit theorems, facilitating the quantification of approximation errors in high-dimensional settings. This is particularly relevant for skewed distributions often encountered in credit risk and extreme value analysis, where symmetric assumptions fail. By leveraging clustering functions to decompose complex systemic risks into manageable components, practitioners can apply these limit results to derive closed-form expressions for aggregate risk statistics across diverse asset classes.

Furthermore, the duality between risk aggregation and optimal transport offers a constructive mechanism for approximating coherent risk measures in infinite-dimensional spaces. Through variational formulations on dense subsets of L^p spaces, one can iteratively refine risk estimates while preserving essential properties like subadditivity and monotonicity. This theoretical apparatus supports the development of scalable algorithms for portfolio optimization and capital reserve calculation, effectively bridging the gap between abstract stochastic analysis and empirical risk management under severe uncertainty.

3.4 Game-Theoretic and Mean-Field Approaches to Systemic Stability

3.4.1 Equilibrium Pricing under Model Misspecification

Equilibrium pricing under model misspecification necessitates a departure from classical single-probability frameworks toward robust formulations that account for ambiguity in both drift and volatility parameters. When agents face uncertainty regarding the true data-generating process, standard expected utility maximization is replaced by max-min preferences or nonlinear expectation theories, such as g-expectations. This approach effectively distinguishes between estimation risk, arising from finite sample sizes, and structural misspecification risk, where the assumed model class fails to capture fat tails, leverage effects, or time-varying volatility dynamics inherent in financial markets.

In continuous-time settings, this robustness is often formalized through penalized control problems where the agent optimizes against a worst-case scenario within a set of plausible measures. The degree of robustification is governed by a penalty parameter, which dynamically adjusts based on the strength of the available signal; as the drift signal weakens, the optimal strategy converges toward a trivial zero-investment position to avoid exposure to unidentifiable risks. Mathematically, this transforms the pricing problem into solving a nonlinear Hamilton-Jacobi-Bellman equation, where the market price of volatility risk acts as a control variable constrained within specific tolerance intervals rather than fixed values.

Consequently, equilibrium prices derived under these conditions reflect a premium for model uncertainty, leading to wider bid-ask spreads and more conservative valuation bounds compared to traditional arbitrage-free pricing. This framework allows for the derivation of robust capital requirements and risk measures, such as Value-at-Risk and Expected Shortfall, which remain valid even when the underlying predictive models are statistically rejected or suffer from distributional mismatches. By treating model parameters as uncertain controls rather than fixed constants, the resulting pricing mechanisms ensure stability against severe market dislocations and provide a rigorous theoretical basis for managing systemic risk in the presence of deep uncertainty.

3.4.2 Mean-Field Control with Common Noise Uncertainty

Mean-field control frameworks incorporating common noise uncertainty address scenarios where a population of agents is subject to a shared stochastic driver with ambiguous statistical properties. Unlike idiosyncratic noise, common noise induces correlation across the entire system, rendering the conditional mean-field term a random process rather than a deterministic trajectory. This structural complexity necessitates the use of nonlinear expectations, such as g-expectation or G-

Brownian motion, to model volatility and drift uncertainties that evolve within convex regions dependent on the state process. The resulting dynamics are governed by fully nonlinear stochastic differential equations where the sublinear function characterizes the worst-case variance bounds.

The mathematical formulation typically leads to a master equation defined on the space of probability measures, which degenerates into a nonlinear partial differential equation in the presence of model ambiguity. Solving these equations requires robustifying the standard Hamilton-Jacobi-Bellman framework to account for the supremum over a set of admissible probability measures consistent with the observed noise structure. Analytic solutions are rarely available except in linear-quadratic settings; consequently, numerical approaches often rely on finite element methods or particle approximations to capture the dependence of optimal controls on both the individual state and the distributional flow under common shocks.

Recent advancements focus on deriving uncertainty-aware strategies that minimize robust risk measures, such as Conditional Value-at-Risk, under these ambiguous dynamics. By aggregating predictions across multiple forward passes or employing worst-case scenario analysis, these methods quantify the impact of parameter uncertainty on capital requirements and control policies. The interplay between the common noise magnitude and the uncertainty aversion parameter reveals nonlinear sensitivities in the optimal feedback laws, highlighting the critical need for adaptive quantification mechanisms that encompass realized volatility peaks while maintaining stability across varying correlation regimes.

3.4.3 Network Contagion Dynamics and Inter-Protocol Exposure

Network contagion dynamics in financial systems are increasingly modeled as complex adaptive systems where interlinkages facilitate the rapid propagation of shocks across distinct protocols. Traditional metrics often fail to capture the non-linear domino effects observed during crises, as they rely on static correlation structures that underestimate tail dependencies. By shifting focus to dynamic network topologies, researchers can identify critical thresholds where localized failures trigger abrupt state shifts, transforming idiosyncratic risks into systemic collapses through cascading default mechanisms.

Inter-protocol exposure amplifies these vulnerabilities by creating channels for shock transmission between previously isolated market domains, such as cryptocurrency ecosystems and traditional equity markets. The integration of heterogeneous assets under shared liquidity constraints means that distress in one sector can induce correlated performance degradation in others, even absent direct fundamental links. Advanced modeling approaches now employ multivariate distributions with heavy tails and time-varying copulas to better represent these co-dependencies, ensuring that lower-quantile risks are not obscured by Gaussian assumptions that break down during extreme market stress.

To mitigate these systemic risks, modern frameworks incorporate real-time surveillance mechanisms capable of detecting early warning signals through mean uncertainty and volatility clustering. These systems utilize network branching architectures to simultaneously estimate conditional means and heteroskedastic variances, providing a more robust assessment of potential losses under changing market conditions. By accounting for parameter uncertainty and dependency structures, such models offer institutions near-real-time probabilistic measures essential for operational decision-making, effectively addressing the limitations of conventional Value at Risk methodologies in highly interconnected financial environments.

4 Uncertainty Quantification Techniques in Deep Neural Networks

4.1 Bayesian Inference and Variational Approximations

4.1.1 Posterior Estimation via Monte Carlo Dropout

Monte Carlo (MC) Dropout provides a computationally efficient framework for approximating the posterior distribution in deep neural networks by establishing a theoretical link between dropout regularization and variational inference. Unlike standard dropout, which is deactivated during inference, MC Dropout retains stochastic masking at test time to generate an ensemble of predictions from a single trained model. This approach effectively transforms a deterministic network into a Bayesian approximation, enabling the quantification of epistemic uncertainty without the prohibitive computational costs associated with training full Bayesian neural networks or maintaining large deep ensembles [9].

$$\begin{cases} \min\limits_{\boldsymbol{\theta}\in\mathbb{R}^L} & \mathbf{1}^T\boldsymbol{\theta} \\ \text{s.}t & Vy_k \leq \boldsymbol{\theta}, k=1,\cdots,N. \end{cases}$$

Figure 9: Chart from 'Risk-Averse Certification of Bayesian Neural Networks' (arXiv:2411.19729)

The estimation process involves performing T stochastic forward passes through the network, where each pass samples a distinct set of active neurons according to the learned dropout probability. These passes yield a collection of predictive samples that approximate the posterior predictive distribution. By aggregating these samples, one can compute the first and second moments; the mean of the samples serves as the final point prediction, while the variance across the samples provides a robust estimate of model uncertainty. This sampling-based mechanism allows the model to express higher uncertainty in regions of the input space where training data is sparse or ambiguous.

Despite its scalability and ease of implementation, MC Dropout relies on specific assumptions regarding the form of the variational posterior, typically assuming independent Bernoulli distributions over weights. While this simplification facilitates tractable inference, it may limit the flexibility to capture complex, multi-modal dependencies within the weight space compared to more sophisticated variational methods. Nevertheless, due to its minimal architectural modifications and ability to prevent overconfidence in safety-critical applications, MC Dropout remains a foundational baseline for uncertainty quantification in deep learning, balancing predictive performance with reliable uncertainty calibration [9].

4.1.2 Variational Autoencoders for Latent Space Uncertainty

Variational Autoencoders (VAEs) establish a probabilistic mapping between observed data and a low-dimensional latent space, offering a principled framework for quantifying uncertainty through the evidence lower bound. By parameterizing the latent distribution with mean and variance vectors, VAEs inherently model aleatoric uncertainty via the stochastic encoding process. This architecture allows the reconstruction loss, typically measured as mean squared error, to serve as a proxy for epistemic uncertainty; samples that are out-of-distribution or underrepresented in the training set generally exhibit higher reconstruction errors, signaling the model's inability to confidently encode these instances within the learned manifold.

However, standard Gaussian assumptions in the latent space often limit the capacity to distinguish between reducible epistemic uncertainty and irreducible data noise. To address this conflation, advanced variational inference techniques introduce flexible posterior approximations that do not rely on pre-specified distributional components. Methods such as injecting uncertainty-dependent stochasticity into latent representations or employing hyperparameter ensembles enable the model to capture complex, multi-modal structures in the feature space. These approaches refine the density estimation, allowing for more granular discrimination between regions of high confidence and those requiring cautious interpretation.

Recent developments further enhance the actionability of latent uncertainty by integrating it directly into the generative pipeline. Techniques like uncertainty-gated modeling utilize the estimated variance to modulate representation learning and aggregation, effectively down-weighting low-confidence evidence during inference. By treating uncertainty as an internal control signal, these frameworks adjust model conservatism dynamically, ensuring that the latent embedding remains robust against distributional shifts. Consequently, VAEs evolve from simple compression tools into sophisticated mechanisms for reliable uncertainty propagation across representation, aggregation, and generation stages.

4.1.3 Hamiltonian Monte Carlo in High-Dimensional Parameter Spaces

Hamiltonian Monte Carlo (HMC) addresses the inefficiency of random-walk Metropolis algorithms in high-dimensional parameter spaces by leveraging gradient information to guide proposal distributions. By introducing auxiliary momentum variables and simulating Hamiltonian dynamics, HMC generates distant proposals with high acceptance probabilities, effectively navigating complex posterior geometries characterized by strong correlations and varying curvature. This gradient-based mechanism significantly reduces autocorrelation between samples, enabling more robust exploration of the target distribution compared to traditional Markov Chain Monte Carlo methods, which often suffer from slow mixing in high dimensions.

To scale these benefits to modern deep learning architectures, Stochastic Gradient Hamiltonian Monte Carlo (SGHMC) integrates mini-batch gradients to approximate the full-data potential energy, thereby mitigating the computational burden of processing massive datasets. However, the introduction of stochastic noise from sub-sampling disrupts the symplectic integrator, necessitating the addition of a friction term and an appropriate noise correction to maintain the correct invariant distribution. Advanced variants further refine this balance by employing adaptive step sizes or higher-order integrators, ensuring numerical stability while preserving the theoretical guarantees required for accurate posterior inference over millions of parameters.

Despite these advancements, applying HMC in extremely high-dimensional neuromanifolds presents persistent challenges, particularly regarding the sensitivity to hyperparameters such as step size and trajectory length. Inaccurate tuning can lead to divergent trajectories or excessive computational costs per iteration, limiting practical deployment. Recent research focuses on preconditioning strategies that adapt to the local geometry of the parameter space, effectively rescaling the momentum to align with the principal components of the Fisher information matrix. These developments aim to enhance sampling efficiency without compromising the fidelity of uncertainty quantification in complex Bayesian neural networks.

4.2 Ensemble Methods and Diversity-Driven Calibration

4.2.1 Deep Ensembles as Approximate Bayesian Model Averaging

Deep Ensembles have emerged as a robust, non-Bayesian technique that effectively approximates Bayesian Model Averaging (BMA) to quantify epistemic uncertainty. Unlike traditional Bayesian Neural Networks that infer a posterior distribution over weights, this approach trains multiple independent neural networks with identical architectures but distinct random initializations. The resulting diversity in convergence trajectories allows the ensemble to capture different modes of the posterior landscape, thereby mitigating the risk of converging to suboptimal local minima and providing a more comprehensive representation of model uncertainty than single deterministic models.

The theoretical justification for treating Deep Ensembles as approximate BMA relies on interpreting the collection of trained models as samples drawn from a multimodal posterior distribution. By aggregating predictions through a mixture distribution, typically approximated as a Gaussian for computational tractability, the method computes the predictive mean and variance across ensemble members. This aggregation naturally decomposes total uncertainty into aleatoric and epistemic components, where the variance among individual model predictions serves as a proxy for epistemic uncertainty, reflecting the model's lack of knowledge in regions sparse of training data.

Despite their superior calibration and predictive performance compared to many scalable Bayesian approximations like Monte Carlo Dropout, Deep Ensembles incur significant computational costs proportional to the number of ensemble members. Training and storing multiple large-scale models demand substantial memory and processing time, presenting a barrier to deployment in resource-constrained environments. Consequently, recent research has focused on developing efficient variants, such as weight-sharing architectures and distillation-based methods, aiming to retain the uncertainty quantification benefits of full ensembles while reducing the associated computational overhead.

4.2.2 Conflictual Ensembles for Epistemic Uncertainty Resolution

Conflictual ensembles address the limitations of standard deep ensembles by explicitly leveraging prediction disagreement to resolve epistemic uncertainty in high-stakes scenarios. Unlike traditional approaches that average outputs to minimize variance, conflictual frameworks treat divergent predictions among ensemble members as critical signals of model ignorance or out-of-distribution inputs. This methodology posits that when constituent models, initialized with diverse stochastic seeds or architectural variations, exhibit significant semantic conflict rather than mere statistical noise, the system is encountering regions of the input space where the learned posterior is ill-defined.

The resolution mechanism operates by decomposing total predictive variance into components attributable to data noise and model inadequacy, isolating the latter through measures of mutual information or entropy across the ensemble's decision boundary. When conflict exceeds a calibrated threshold, indicating high epistemic uncertainty, the system triggers adaptive responses such as injecting stochasticity into latent representations or down-weighting low-confidence evidence during attention-based aggregation. This prevents the collapse of uncertainty estimates often observed in large-scale models where optimization dynamics force convergence, thereby masking the true extent of knowledge gaps regarding underrepresented features.

Furthermore, these ensembles facilitate actionable uncertainty by dynamically adjusting model conservatism, such as increasing regularization strength or smoothing inference trajectories, proportional to the detected conflict magnitude. By maintaining distinct predictive distributions rather than collapsing them into a single mean estimate, conflictual ensembles ensure that epistemic uncertainty remains reducible and informative. This approach allows the model to distinguish between challenging samples requiring additional data and inherently noisy observations, ultimately providing a robust frequentist approximation of Bayesian posterior diversity without the prohibitive computational cost of full Markov Chain Monte Carlo sampling.

4.2.3 Hyper-Networks for Task-Specific Uncertainty Decomposition

Hyper-networks have emerged as a sophisticated architectural paradigm for disentangling predictive uncertainty into task-specific epistemic and aleatoric components. Unlike traditional Bayesian neural networks that place priors directly on weights, hyper-networks employ a secondary network to generate the parameters of a primary model, effectively mapping task descriptors or noise vectors to weight distributions [10]. This formulation allows for the implicit representation of the posterior weight distribution, where epistemic uncertainty is quantified through the variance of weights generated across different stochastic samples. By treating the hyper-network input as a source of stochasticity, the system can sample diverse functional realizations consistent with the training data, facilitating a robust estimation of model uncertainty without the computational overhead of maintaining explicit ensemble members.

The decomposition mechanism leverages the hyper-network's capacity to condition weight generation on specific task characteristics or latent representations of data noise. In this framework, aleatoric uncertainty is captured by modeling heteroscedastic noise directly within the generated parameters, often aligned with domain-specific physical mechanisms or learned data-dependent variance [11]. Conversely, epistemic uncertainty arises from the divergence in generated weights when subjected to varying noise inputs, reflecting the model's lack of knowledge in under-explored regions of the input space. This separation is critical for applications requiring distinct mitigation strategies, as it distinguishes between irreducible data noise and uncertainty reducible through additional data acquisition or model refinement.

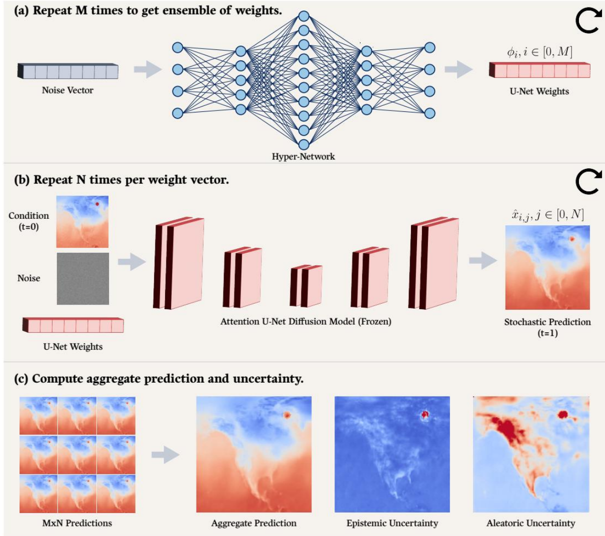

Figure 10: Chart from 'Estimating Epistemic and Aleatoric Uncertainty with a Single Model' (arXiv:2402.03478)

$$\hat{\boldsymbol{T}} = \arg\min_{\boldsymbol{T}} (e(\boldsymbol{T}(\boldsymbol{\mathcal{P}}), \boldsymbol{\mathcal{Q}}))$$

Figure 11: Chart from 'Deep Bayesian ICP Covariance Estimation' (arXiv:2202.11607)

Recent advancements integrate Bayesian principles into hyper-networks to enhance this decomposition, enabling single-pass inference that simultaneously estimates both uncertainty sources. By injecting uncertainty-dependent stochasticity into latent representations, these models can dynamically adjust aggregation mechanisms, such as gating attention scores based on confidence levels.

This approach not only improves calibration but also renders uncertainty actionable across representation learning and decision-making processes. Consequently, hyper-networks provide a scalable and principled alternative to deep ensembles, offering efficient uncertainty disentanglement that maintains high inference accuracy while adapting to complex, multi-modal task environments.

4.3 Evidential Deep Learning and Belief Function Theory

4.3.1 Dirichlet Priors for Categorical Uncertainty Modeling

Dirichlet priors serve as a foundational mechanism for modeling categorical uncertainty within the framework of Evidential Deep Learning, offering a deterministic alternative to computationally expensive sampling-based Bayesian methods. By parameterizing a Dirichlet distribution over the simplex of class probabilities using evidence variables derived directly from neural network outputs, this approach explicitly distinguishes between aleatoric uncertainty, inherent to the data noise, and epistemic uncertainty, arising from model ignorance. The concentration parameters of the Dirichlet distribution are constructed as a function of non-negative evidence, allowing the model to represent both precise predictions with high confidence and ambiguous states characterized by low evidence mass.

The mathematical formulation leverages the properties of the Dirichlet distribution to decompose total predictive uncertainty into distinct components without requiring multiple forward passes or ensemble averaging. Specifically, the expected probability for each class is derived from the normalized evidence, while the uncertainty mass is inversely proportional to the total evidence accumulated across all classes. This structure enables the direct calculation of entropy-based metrics and mutual information, providing a rigorous quantification of conflict between the predicted distribution and the uniform prior. Consequently, the model can identify out-of-distribution samples effectively, as these inputs typically yield low evidence values, resulting in a diffuse Dirichlet distribution that reflects high epistemic uncertainty.

Optimization of such models typically involves minimizing a loss function that combines standard classification error with a regularization term designed to penalize evidence assigned to incorrect classes, particularly for out-of-distribution data. This evidential regularization encourages the network to output low evidence for unfamiliar inputs, thereby preventing overconfident predictions on anomalous samples. While traditional Bayesian approaches often struggle with scalability in deep architectures, the Dirichlet prior formulation maintains computational efficiency comparable to standard deterministic training. However, careful calibration is required to ensure that the evidence parameters accurately reflect the underlying data density, as improper initialization or optimization dynamics can lead to miscalibrated uncertainty estimates that fail to distinguish reliably between different sources of variance.

4.3.2 Dempster-Shafer Frameworks in Set-Valued Classification

The Dempster-Shafer theory provides a rigorous mathematical foundation for set-valued classification by distinguishing between aleatoric uncertainty, inherent in the data, and epistemic uncertainty, arising from model ignorance. Unlike traditional probabilistic approaches that assign precise probabilities to singleton classes, this framework utilizes belief functions to distribute mass over subsets of the label space, effectively representing partial ignorance. This capability is particularly advantageous in safety-critical domains where abstaining from a definitive prediction or outputting a set of plausible candidates is preferable to making an erroneous confident decision.

Implementation of these frameworks in deep learning typically involves modifying the output layer to predict parameters of a Dirichlet distribution or directly estimating belief masses for each subset of classes. The combination of evidence from multiple sources or network layers is achieved through Dempster's rule of combination, which aggregates belief assignments while managing conflicts between differing pieces of evidence. Recent advancements have integrated this theory with neural architectures to enable end-to-end training, where the loss function is designed to minimize conflict and maximize the belief assigned to the correct class or subset, thereby enhancing the model's ability to quantify uncertainty without requiring extensive ensemble methods.

Despite its theoretical robustness, the application of Dempster-Shafer frameworks faces computational challenges due to the exponential growth of the power set relative to the number of classes, necessitating efficient approximation strategies for high-dimensional tasks. Current research focuses on simplifying the mass assignment process, often restricting focus to singletons and the universal set to maintain tractability while preserving the core benefits of evidence theory. Furthermore, hybrid approaches that combine Dempster-Shafer reasoning with other uncertainty quantification techniques, such as variational inference or deep ensembles, are emerging to leverage the strengths of both paradigms, offering improved calibration and more reliable set-valued predictions in complex, real-world scenarios.

4.3.3 Gaussian Random Fuzzy Numbers in Regression Tasks

Gaussian Random Fuzzy Numbers (GRFNs) provide a unified parametric framework for regression tasks by simultaneously modeling probabilistic variability and possibilistic imprecision. Unlike standard heteroscedastic regression which only captures aleatoric noise through variance, a GRFN is defined by a location parameter μ, a probabilistic variance σ^2, and a possibilistic precision h [12]. This tripartite structure allows the model to distinguish between inherent data stochasticity, represented by the Gaussian spread, and epistemic ambiguity arising from incomplete knowledge or vague observations, encoded within the fuzzy membership function.

$$C(\Psi) = \frac{1}{n}\sum_{i=1}^{n} \mathcal{L}_{\lambda,\epsilon}(y_i, \widetilde{Y}(\boldsymbol{x}_i)) + \frac{\xi}{J}\sum_{j=1}^{J} h_j,$$

Figure 12: Chart from 'An Evidential Neural Network Model for Regression Based on Random Fuzzy Numbers' (arXiv:2208.00647)

In practical regression implementations, GRFNs extend traditional mean-variance estimation by treating the target variable as an epistemic random fuzzy set rather than a simple scalar or distribution. The learning objective typically involves optimizing a likelihood function that accounts for both the stochastic noise distribution and the fuzzy membership grades of the residuals. This approach is particularly advantageous in scenarios with limited data or noisy labels, where the separation between irreducible randomness and model uncertainty is critical. By propagating both uncertainty types through the network, GRFNs yield predictive intervals that are theoretically grounded in both probability theory and possibility theory.

Empirical evaluations suggest that GRFN-based regressors offer superior calibration compared to methods relying solely on Gaussian assumptions or deep ensembles. The dual-parameterization enables the model to maintain sharp predictions in data-rich regions while expanding the possibilistic support in areas of high epistemic uncertainty, effectively preventing overconfident extrapolations. Consequently, this formalism facilitates more robust decision-making in safety-critical applications, as the resulting uncertainty bounds reflect a comprehensive assessment of both observational noise and structural model limitations without requiring computationally expensive sampling procedures.

4.4 Conformal Prediction and Frequentist Coverage Guarantees

4.4.1 Distribution-Free Prediction Intervals via Conformal Sets

Distribution-free prediction intervals leverage conformal prediction theory to provide rigorous, finite-sample coverage guarantees without imposing restrictive parametric assumptions on the underlying data distribution. Unlike traditional Bayesian or ensemble methods that rely on specific likelihood forms such as Gaussian or Student's t-distributions, this approach constructs prediction sets by utilizing nonconformity scores derived from calibration data. The core mechanism involves defining a set of candidate labels for a new test point that includes all values whose associated nonconformity score falls below a quantile threshold determined by the desired confidence level, thereby ensuring that the true label is contained within the interval with probability at least $1-\alpha$.

The implementation of conformal sets typically proceeds in two stages: a training phase to fit a base predictor and a calibration phase to adjust interval widths based on empirical residuals. By exchanging the order of training and calibration data through split-conformal or cross-conformal strategies, these methods achieve exchangeability, which is critical for maintaining valid coverage properties even under distribution shifts. Recent advancements have extended this framework to handle heteroscedastic noise by employing adaptive nonconformity scores that scale residuals relative to local uncertainty estimates, allowing the resulting intervals to tighten in regions of low variance and expand where data is sparse or noisy.

Despite their theoretical robustness, distribution-free approaches face challenges regarding computational efficiency and the granularity of uncertainty decomposition. Standard conformal methods operate post-hoc, meaning they do not integrate uncertainty quantification directly into the model's objective function during training, potentially limiting their ability to distinguish between aleatoric and epistemic sources of error without auxiliary mechanisms like deep ensembles. Furthermore, while full conformal prediction offers exact coverage, its computational cost scales linearly with the size of the calibration set and the cardinality of the output space, prompting ongoing research into scalable approximations and localized conformal inference techniques for high-dimensional regression and classification tasks.

4.4.2 Blockwise Jackknife for Time-Series Path Dependency

Blockwise Jackknife methods address the critical challenge of quantifying uncertainty in time-series forecasting where standard resampling techniques fail due to temporal dependencies. Unlike independent and identically distributed data, time-series paths exhibit strong autocorrelation, necessitating the preservation of local structures during perturbation. This approach partitions the historical trajectory into contiguous blocks rather than individual observations, ensuring that the intrinsic dynamics and short-term memory within each segment remain intact while systematically leaving out specific blocks to estimate variability.

The implementation involves constructing pseudo-values by retraining or re-evaluating the model on datasets where distinct blocks of the path are omitted, effectively simulating distributional shifts without breaking the sequential continuity required for valid inference. By aggregating these block-specific estimates, the method approximates the sampling distribution of the forecast estimator, capturing both aleatoric noise and epistemic uncertainty arising from model sensitivity to specific historical contexts. This is particularly vital for path-dependent functionals, such as those found in stochastic differential equations or autoregressive architectures, where the forecast depends on the entire trajectory rather than isolated points.

Furthermore, this technique provides a robust framework for validating probabilistic forecasts against metrics like Continuous Ranked Probability Score without relying on strict parametric assumptions about the underlying data generating process. It enables the detection of structural breaks and regime changes by analyzing the variance contributions of different temporal segments, offering granular insights into model stability. Consequently, Blockwise Jackknife serves as a non-asymptotic tool for enhancing the reliability of deep learning models in financial and environmental domains, where accurate uncertainty quantification over complex, evolving paths is paramount for decision-making.

4.4.3 Adaptive Calibration under Distributional Shift

Adaptive calibration under distributional shift addresses the critical failure mode where deep learning models exhibit overconfidence when processing out-of-distribution (OOD) inputs or noisy data. Traditional point-estimation training often yields miscalibrated predictions precisely when errors are most costly, such as in autonomous navigation systems encountering novel environments. To mitigate this, modern approaches integrate aleatoric uncertainty estimation directly into the loss function, typically via heteroskedastic Gaussian Negative Log-Likelihood. This formulation allows the model to learn input-dependent variance, effectively distinguishing between inherent data noise, such as sensor artifacts or ambiguous object boundaries, and epistemic uncertainty arising from distributional shifts.

The mechanism for adaptive calibration relies on dynamically adjusting the predicted uncertainty bounds rather than applying static post-hoc corrections like temperature scaling. By optimizing parameters that control both the mean prediction and the variance, the model learns to assign high aleatoric uncertainty to regions with label noise or mislabeled pixels while maintaining confidence in clean data. This uncertainty-aware training prevents the convergence issues often associated with fixed stochastic perturbations. Instead, it employs uncertainty-adaptive control to regulate the injection of noise during inference, ensuring that the model's confidence intervals expand appropriately as the input drifts from the training distribution, thereby preserving the reliability of safety-critical decisions.

Furthermore, robust calibration under shift requires metrics that evaluate the alignment between predicted uncertainty and observed error frequencies without conflating distinct uncertainty sources. Techniques leveraging natural gradient methods or ensemble distillation have shown promise in scaling these probabilistic formulations to large datasets while maintaining computational efficiency. The ultimate goal is to achieve conservative inference where rising uncertainty triggers fallback mechanisms, such as yielding control to human operators, rather than propagating erroneous high-confidence predictions. This paradigm shift from accuracy-centric to reliability-centric modeling ensures that systems remain tractable and safe even when facing significant environmental variations or domain gaps.

5 Future Directions

Despite the significant advancements in deep learning architectures for financial risk modeling, several critical limitations persist that hinder their widespread adoption in regulated environments. Current models often function as opaque black boxes, lacking the inherent interpretability required by regulators to validate capital adequacy and stress testing results. While attention mechanisms and graph neural networks offer some visibility into feature importance, they frequently fail to provide causal explanations for systemic contagion or tail events, creating a disconnect between model outputs and economic theory. Furthermore, many existing frameworks struggle with non-stationarity; models trained on historical data often degrade rapidly when faced with unprecedented structural breaks or regime shifts, such as those caused by geopolitical shocks. There is also a notable gap in the integration of these advanced techniques with established coherent risk measures, where the optimization objectives of deep learning do not always align perfectly with regulatory constraints like Expected Shortfall, leading to potential underestimation of extreme risks in live deployment scenarios.

Future research must prioritize the development of hybrid neuro-symbolic frameworks that embed domain-specific financial constraints and economic laws directly into neural architectures. By integrating symbolic reasoning with deep learning, researchers can create models that are not only data-efficient but also adhere to fundamental principles of no-arbitrage and market equilibrium, thereby enhancing both robustness and interpretability. Concurrently, there is an urgent need to advance self-supervised and meta-learning strategies capable of adapting to distributional shifts in real-time. These approaches should enable models to continuously learn from streaming data without catastrophic forgetting, allowing them to generalize effectively across different market regimes and rare event scenarios where labeled data is scarce. Additionally, future work should focus on refining generative models, particularly diffusion and adversarial networks, to simulate high-fidelity, counterfactual stress scenarios that capture complex, non-linear interdependencies across global markets, moving beyond simple historical resampling to true causal scenario generation.

The successful realization of these research directions holds the potential to fundamentally transform the resilience and stability of the global financial system. By bridging the gap between data-driven flexibility and theoretical rigor, next-generation risk models will provide regulators and institutions with more reliable tools for anticipating systemic failures and managing tail exposure. Enhanced interpretability will foster greater trust in artificial intelligence-driven decisions, facilitating smoother regulatory approval and broader industry adoption. Ultimately, these advancements will enable a proactive rather than reactive approach to risk management, allowing financial entities to navigate profound uncertainty with greater confidence, optimize capital allocation more efficiently, and mitigate the cascading effects of future crises before they materialize.

6 Conclusion

This survey has systematically examined the transformative potential of deep learning architectures in addressing the profound uncertainties characterizing modern financial systems. By synthesizing advancements across three distinct thematic pillars, the analysis demonstrates how advanced neural topologies, generative paradigms, and reinforcement learning strategies collectively overcome the limitations of traditional econometric frameworks. Specifically, recurrent networks and attention mechanisms have proven superior in capturing volatility clustering and long-range temporal dependencies, while Graph Neural Networks provide an indispensable framework for modeling systemic contagion and spatial interdependencies that isolated time-series models frequently miss. Furthermore, the integration of adversarial training and diffusion models offers robust solutions for scenario synthesis, enabling the accurate simulation of tail risks and extreme events without the computational burdens of conventional Monte Carlo methods. Finally, risk-aware reinforcement learning strategies facilitate dynamic hedging and capital allocation that adapt to non-stationary market regimes, ensuring that decision-making processes remain resilient against distributional shifts and structural breaks.

The significance of this work lies in its unified taxonomy, which bridges the gap between theoretical artificial intelligence innovations and the practical exigencies of quantitative risk management. Unlike previous reviews that broadly address algorithmic trading or general forecasting, this survey critically evaluates these technologies through the specific lens of uncertainty quantification and regulatory compliance. It elucidates how deep learning moves beyond mere prediction accuracy to provide coherent, interpretable, and stress-testable risk metrics essential for maintaining financial stability. By categorizing architectures based on their functional roles in handling temporal dynamics, generative simulation, and decision-making under risk, this paper offers a structured roadmap for researchers and practitioners navigating the complex trade-offs between model complexity, interpretability, and computational efficiency. Consequently, it establishes a definitive reference point for understanding how next-generation methodologies can enhance the resilience of global financial ecosystems against unforeseen shocks.

Looking forward, the integration of deep learning into financial risk modeling must evolve from experimental application to standardized practice, necessitating a concerted effort to address remaining challenges in explainability and regulatory alignment. While the technical capabilities of these models are evident, their widespread adoption hinges on developing frameworks that satisfy rigorous supervisory requirements while maintaining the flexibility to adapt to evolving market structures. Future research should prioritize the development of hybrid models that combine the data-driven power of neural networks with the theoretical rigor of economic principles, ensuring that black-box predictions are grounded in sound financial logic. Practitioners and regulators are called to collaborate on establishing benchmarks and validation protocols that foster trust in these advanced systems. Ultimately, embracing this paradigm shift is not merely an option for competitive advantage but a critical imperative for safeguarding the integrity and stability of the global financial infrastructure in an era of increasing volatility.

References

[1] Optimization Method of Multi-factor Investment Model Driven by Deep Learning for Risk Control

[2] Distributional Reinforcement Learning on Path-dependent Options

[3] Analysis of Financial Risk Behavior Prediction Using Deep Learning and Big Data Algorithms

[4] Robust Graph Neural Networks for Stability Analysis in Dynamic Networks

[5] Uni-FinLLM A Unified Multimodal Large Language Model with Modular Task Heads for Micro-Level Stock Prediction and Macro-Level Systemic Risk Assessment

[6] Probability Weighting Meets Heavy Tails An Econometric Framework for Behavioral Asset Pricing

[7] Uncertainty Quantification of Deep Learning for Spatiotemporal Data Challenges and Opportunities

[8] Unknown PDF

[9] Risk-Averse Certification of Bayesian Neural Networks

[10] Estimating Epistemic and Aleatoric Uncertainty with a Single Model

[11] Deep Bayesian ICP Covariance Estimation

[12] An Evidential Neural Network Model for Regression Based on Random Fuzzy Numbers