Bayesian Deep Learning for Uncertainty Quantification: A Comprehensive Survey

Abstract

The deployment of deep learning systems in safety-critical domains such as healthcare, autonomous driving, and scientific discovery necessitates not only high predictive accuracy but also reliable estimates of predictive uncertainty. Bayesian Deep Learning (BDL) has emerged as a principled framework for quantifying uncertainty by treating network weights as random variables rather than point estimates, thereby distinguishing between reducible epistemic uncertainty and irreducible aleatoric uncertainty. This survey provides a comprehensive review of the state-of-the-art in Bayesian Deep Learning, synthesizing insights from 300 recent academic papers. We present a detailed taxonomy of methodological families, ranging from approximate inference techniques like Variational Inference (VI) and Markov Chain Monte Carlo (MCMC) to scalable approximations such as Deep Ensembles, Monte Carlo Dropout, and Laplace approximations. Furthermore, we examine advanced architectures including Bayesian Neural Operators, Spiking Neural Networks, and Neurosymbolic frameworks. We critically analyze evaluation metrics, benchmarking results across diverse domains including medical imaging, climate modeling, and materials science, and discuss persistent challenges such as computational intractability, the "cold posterior" effect, and the failure of standard methods to capture uncertainty in overparameterized regimes. By integrating theoretical foundations with empirical evidence, this survey aims to guide future research toward robust, scalable, and trustworthy Bayesian deep learning systems.

1. Introduction

Deep neural networks (DNNs) have achieved unprecedented success in tasks ranging from image classification to natural language processing. However, standard deterministic training via maximum likelihood estimation often yields overconfident predictions, particularly when models encounter out-of-distribution (OOD) data or operate in data-scarce regimes [1], [2]. In high-stakes applications, the inability to quantify "what the model does not know" poses significant risks. Bayesian Deep Learning (BDL) addresses this limitation by placing probability distributions over network parameters, allowing the model to express uncertainty through the posterior predictive distribution [3], [4].

The core distinction in uncertainty quantification (UQ) lies between aleatoric uncertainty, which arises from inherent noise in the data generation process and is irreducible, and epistemic uncertainty, which stems from model ignorance due to limited data and is reducible [5], [6]. While early works focused on exact Bayesian inference, the high dimensionality of modern neural networks rendered exact posterior computation intractable, necessitating the development of approximate inference methods [7], [8]. These include Variational Inference (VI), which approximates the posterior with a tractable distribution [9], [10], and Markov Chain Monte Carlo (MCMC) methods, which sample from the posterior but suffer from slow convergence in high dimensions [3], [11].

Recent years have witnessed a proliferation of scalable approximations and hybrid architectures. Methods such as Monte Carlo (MC) Dropout [12], [13] and Deep Ensembles [1], [14] have gained traction for their computational efficiency, though they often trade off theoretical rigor for scalability. Concurrently, theoretical advances have linked finite-width networks to Gaussian Processes (GPs) in the infinite-width limit, providing analytical insights into uncertainty behavior [15], [16]. Despite these advances, significant challenges remain. Empirical studies reveal that popular methods like MC Dropout can fail to increase uncertainty in extrapolation regions [13], and large overparameterized models may exhibit an "uncertainty hole" where epistemic uncertainty collapses despite data scarcity [17].

This survey synthesizes findings from 300 relevant papers to provide a holistic view of the field. We organize the discussion into a structured taxonomy of methods, covering variational approaches, sampling-based techniques, ensemble methods, and post-hoc approximations. We further explore domain-specific applications, evaluation protocols, and open theoretical questions, offering a roadmap for developing next-generation uncertainty-aware AI systems.

2. Method

The methodological landscape of Bayesian Deep Learning is diverse, encompassing techniques that modify the training objective, alter the network architecture, or apply post-hoc corrections to pre-trained models. We categorize these approaches into several coherent families: Approximate Inference via Variational Methods, Sampling-Based Inference, Ensemble and Hybrid Techniques, Linearized and Post-Hoc Approximations, and Specialized Architectures for Specific Data Modalities.

2.1 Variational Inference and Posterior Approximation

Variational Inference (VI) remains a cornerstone of scalable Bayesian Deep Learning, framing posterior estimation as an optimization problem that minimizes the Kullback-Leibler (KL) divergence between a variational distribution $q(\theta)$ and the true posterior $p(\theta|D)$ [3], [4].

2.1.1 Mean-Field and Structured Variational Inference The most common implementation, Mean-Field Variational Inference (MFVI), assumes independence among weight parameters, typically modeling them as factorized Gaussians [3], [9]. While computationally efficient, MFVI often underestimates variance and ignores parameter correlations, leading to overconfident predictions [3], [18]. To mitigate this, recent works have introduced structured variational families. For instance, Multiplicative Normalizing Flows (MNF) enhance posterior flexibility by transforming simple base distributions through invertible mappings, enabling the capture of complex multimodal posteriors in cosmological parameter inference [19]. Similarly, Boosted Variational Inference constructs mixture densities iteratively, adding components to refine the posterior approximation and better model multimodal weight distributions [20].

Hybrid architectures have also emerged to balance efficiency and expressiveness. Hybrid Bayesian Neural Networks (Hybrid BNNs) partition the network into deterministic feature extraction layers and stochastic prediction layers, reducing computational redundancy while preserving uncertainty estimation at the output [9]. Another approach, Implicitly Bayesian Neural Networks (iBNN), shifts stochasticity from high-dimensional weights to low-dimensional layer-wise input nodes, reducing the dimensionality of the inference problem by orders of magnitude while maintaining calibration [21].

2.1.2 Function-Space and Output-Space Variational Inference Traditional VI operates in weight space, which can be inefficient for overparameterized networks. Function-Space Variational Inference (FSVI) addresses this by performing inference directly over the space of functions defined by the network [22]. FSVI employs local linearization to approximate the KL divergence between function distributions, allowing for the incorporation of prior knowledge via context distributions.

This approach has shown superior performance in OOD detection compared to weight-space baselines [22].

Extending this concept, Variational Inference in the Final-layer Output space (VIFO) learns the mean and variance of the final layer outputs directly, bypassing weight-space inference entirely [23]. By operating in the lower-dimensional output space, VIFO achieves efficient variational inference with competitive calibration, particularly in shifted distribution scenarios. Furthermore, Function-Space Regularization with Dirichlet Priors interprets stochastic network outputs as samples from a Dirichlet distribution, enabling explicit control over predictive uncertainty independent of parameterization [24].

2.1.3 Advanced Priors and Regularization The choice of prior significantly influences posterior geometry. Spike-and-slab priors enable automatic sparsification and feature selection by introducing latent binary indicators for weight inclusion, effectively performing structural learning alongside parameter inference [10], [25]. In the context of sparse deep ReLU networks, semi-parametric Bernstein-von Mises theorems have been extended to show that sparse posteriors approximate Gaussian limits for low-dimensional functionals under specific smoothness constraints [26].

Regularization techniques have also been adapted to improve uncertainty. Maximum Mean Discrepancy (MMD) regularization replaces the standard KL divergence in the Evidence Lower Bound (ELBO) with an MMD estimator, preventing distribution collapse and enhancing robustness to prior selection [27]. Additionally, Calibration-Aware Bayesian Neural Networks (CA-BNN) integrate a differentiable approximation of Expected Calibration Error (ECE) directly into the variational objective, mitigating misspecification effects without sacrificing uncertainty quality [28].

2.2 Sampling-Based Inference and MCMC

While VI offers scalability, Sampling-based methods like Markov Chain Monte Carlo (MCMC) provide asymptotically exact posterior estimates, serving as a gold standard for evaluation [3], [8].

2.2.1 Stochastic Gradient MCMC Stochastic Gradient MCMC (SG-MCMC) methods, such as Stochastic Gradient Langevin Dynamics (SGLD) and Stochastic Gradient Hamiltonian Monte Carlo (SGHMC), enable scaling to large datasets by using mini-batch gradients [29]. These methods inject noise into the gradient updates to simulate Langevin dynamics, allowing the sampler to explore the posterior landscape. However, standard SG-MCMC can suffer from slow mixing in high-dimensional spaces. Recent innovations include Condensed Stein Variational Gradient Descent (cSVGD), which integrates sparsifying priors into the Stein variational flow, driving particle ensembles toward sparse manifolds while maintaining uncertainty quantification [30].

In specific domains, MCMC remains indispensable. For seismic imaging, preconditioned SGLD (pSGLD) has been used to sample deep prior posteriors, propagating data-driven image uncertainty to downstream tasks like horizon tracking [31]. Similarly, in radio galaxy classification, Hamiltonian Monte Carlo (HMC) has demonstrated superior uncertainty calibration compared to cheaper approximations, albeit at a prohibitive computational cost for large-scale deployment [32], [33].

2.2.2 Microcanonical and Ensemble Sampling To improve sampling efficiency, Microcanonical Langevin Ensembles (MILE) propose a two-stage hybrid workflow: first optimizing independent ensemble members to locate diverse high-likelihood modes, then executing Microcanonical Langevin

Monte Carlo to explore local posterior structures [34]. This approach combines the mode-seeking capability of optimization with the local exploration of MCMC, yielding calibrated uncertainty estimates with strong predictive accuracy. Additionally, Multi-Level Monte Carlo (MLMC) methods have been adapted for Bayesian DNNs using Trace-Class Neural Network (TNN) priors, constructing telescoping sums to reduce the variance of posterior estimators and achieve optimal computational complexity [35].

2.3 Deep Ensembles and Hybrid Approximations

Deep Ensembles (DE), which average predictions from multiple independently trained networks, have consistently outperformed many Bayesian approximations in terms of calibration and accuracy, despite lacking a strict Bayesian interpretation in their standard form [1], [14], [2].

2.3.1 Theoretical Foundations and Empirical Bayes Recent theoretical work has bridged the gap between ensembles and Bayesian inference. It has been shown that Deep Ensembles implicitly perform Empirical Bayes by learning a prior supported entirely on maximum-likelihood estimators, effectively maximizing the Evidence Lower Bound with flexible priors [36]. Furthermore, in the infinite-width limit, the squared prediction errors of Random Network Distillation (RND) have been proven equivalent to Deep Ensemble predictive variance, linking these seemingly distinct approaches through Neural Tangent Kernel (NTK) theory [37].

2.3.2 Optimized and Efficient Ensembles To enhance ensemble diversity and efficiency, Bayesian Optimized Deep Ensembles (BODE) utilize Bayesian Optimization to tune hyperparameters for individual ensemble members, generating diverse architectures that better separate aleatoric and epistemic uncertainties [38]. Shallow Ensembles with shared weights (DPOSE) offer a computationally efficient alternative, splitting only the last layer into an ensemble of heads while sharing preceding weights, achieving comparable accuracy to deep ensembles at lower cost [39].

In the context of Large Language Models (LLMs), LoRA Ensembles construct Bayesian posterior approximations via ensembles of Low-Rank Adaptation modules. This approach provides accurate posterior approximations comparable to larger ensembles with minimal overhead, making it suitable for parameter-efficient fine-tuning [40]. Additionally, Uncertainty-Aware Perceivers integrate multiple ensemble strategies, including Stochastic Weight Averaging (SWA), into the Perceiver architecture, significantly enhancing generalization and calibration [41].

2.3.3 Cooperative and Disentangled Frameworks Cooperative Bayesian and Variance Networks (BNN-VE) sequentially train distinct components: a deterministic mean network, a variance network predicting squared residuals via Gamma likelihood, and a Bayesian network for epistemic uncertainty. This decoupled strategy stabilizes training and improves the disentanglement of uncertainty sources [42]. Similarly, Direct Aleatoric and Deep Ensemble-based Epistemic (DADEE) frameworks combine Anchored Ensembles for epistemic uncertainty with Direct Epistemic Uncertainty Prediction (DEUP) networks for aleatoric uncertainty, addressing the collapse of model variance-based algorithms in-domain [43].

2.4 Linearized and Post-Hoc Approximations

Post-hoc methods approximate the posterior of a pre-trained deterministic network, offering a practical route to uncertainty quantification without retraining.

2.4.1 Laplace Approximation and Variants The Laplace Approximation (LA) fits a Gaussian distribution around a Maximum A Posteriori (MAP) estimate using the Hessian of the loss function. Variational Laplace (VL) improves upon this by replacing the Hessian with the Fisher Information Matrix to avoid pathological optima and deriving a squared-gradient regularizer for stable optimization [44], [45]. Last-Layer Laplace Approximation focuses inference solely on the final layer, significantly reducing computational cost while maintaining competitive performance [16], [46].

Mixtures of Laplace Approximations (MoLA) extend this by applying last-layer LA to each member of a deep ensemble, constructing a Gaussian mixture posterior that mitigates overconfidence far from the training data [47]. Regularization Variation (RegVar) recovers Laplace variance via finite differences of prediction-regularized MAP estimates, avoiding explicit Hessian computation [48].

2.4.2 Neural Tangent Kernel and Linearized Sampling Leveraging the NTK regime, Neural Uncertainty Quantification by Linearized Sampling (NUQLS) linearizes a trained network and generates an ensemble by applying stochastic gradient descent to the linearized surrogate. This method explores the solution manifold efficiently without requiring artificial priors [49]. Similarly, LinearSampling compares full-network versus last-layer linearization, demonstrating via Random Matrix Theory that last-layer linearization offers no intrinsic Bayes Free Energy disadvantage in large-width limits, justifying its widespread use for efficiency [16].

2.4.3 Calibration and Correction Post-hoc calibration is critical for reliable uncertainty. Sigma Scaling applies a single learnable scalar to rescale the predictive standard deviation, jointly recalibrating aleatoric and epistemic uncertainty in regression tasks [50]. Conformalized Generative Bayesian Imaging (CGBI) integrates split conformal prediction with deep ensembling of generative models, ensuring valid coverage guarantees for computational imaging tasks [51]. Additionally, the Laplace Bridge (LB) provides an analytic mapping from Gaussian logits to Dirichlet softmax outputs, enabling fast predictive uncertainty estimation without sampling [52].

2.5 Specialized Architectures and Domains

Bayesian Deep Learning has been adapted to various specialized architectures and domains, addressing unique challenges in each.

2.5.1 Graph Neural Networks and Spatiotemporal Models For Graph Neural Networks (GNNs), Bayesian frameworks propagate aleatoric uncertainty via Assumed Density Filtering (ADF) by matching moments across layers, while treating links and features as probabilistic [53]. Variational Graph Neural Networks (VGNN) place variational layers exclusively in the decoder to balance efficiency and robustness in inverse problems [54]. In spatiotemporal forecasting, models like SYMHnet combine Graph Neural Networks with Bidirectional LSTMs and Bayesian dropout to capture both parameter correlations and temporal dynamics for geomagnetic index prediction [55]. Deep Ensembles have also been applied to Spatiotemporal GNNs for traffic forecasting, using Gaussian copulas to sample diverse hyperparameter configurations [56].

2.5.2 Physics-Informed and Scientific Machine Learning Physics-Informed Neural Networks (PINNs) have been extended to Bayesian frameworks (B-PINN) by integrating PDE residuals into the ELBO loss function, enabling principled uncertainty quantification in scientific discovery [57]. In turbulent flow modeling, Bayesian PINNs with tempered likelihoods and Repulsive Deep

Ensembles enforce diversity in function space, achieving consistent uncertainty calibration across latent variables [58]. For operator learning, Generative In-Context Operator Networks (GenICON) formulate in-context learning as implicit Bayesian inference on Banach spaces, extending deterministic operators to generative models for sampling posterior solutions [59].

2.5.3 Spiking and Neuromorphic Computing In Spiking Neural Networks (SNNs), Bayesian learning via Improved Variational Online Newton (IVON) replaces point estimates with weight distributions, smoothing angular loss landscapes and improving accuracy in speech recognition [60]. Average-Over-Time SNNs (AOT-SNN) approximate Monte Carlo dropout by averaging outputs across intrinsic time steps within a single forward pass, enabling efficient uncertainty estimation for regression tasks [61]. Deep Bayesian Convolutional SNNs integrate Monte Carlo Dropout into spiking architectures, demonstrating that rate-coding schemes yield lower uncertainty compared to temporal coding [62].

2.5.4 Medical Imaging and Healthcare Medical applications demand high reliability. Bayesian U-Nets with Flipout estimators and attention modules capture epistemic and aleatoric uncertainty in brain tumor segmentation [63]. Uncertainty-Aware BI-RADS prediction models map predictive entropy to severity levels, mirroring radiologist uncertainty in mammography [64]. In survival analysis, NeuralSurv employs two-stage data augmentation with Pólya-Gamma variables to transform intractable continuous-time likelihoods into conditionally Gaussian forms amenable to variational inference [65]. Furthermore, Bayesian frameworks for structural heart disease triage utilize selective prediction protocols based on credible intervals to identify inconclusive cases [66].

2.5.5 Climate, Weather, and Environmental Modeling In weather forecasting, hybrid Bayesian-Physics ensembles integrate variational inference with flow-dependent stochastic perturbations, achieving significant speedups over traditional ensemble methods while maintaining skill [67]. Heteroscedastic BNNs parameterize subgrid-scale tendencies in climate models, disentangling aleatoric (subgrid variability) and epistemic (parametric) uncertainties across timescales [68]. For wildfire danger forecasting, unified frameworks employ double Monte Carlo estimation to jointly model epistemic and aleatoric uncertainty, revealing divergent behaviors over temporal horizons [69]. Spatial uncertainty quantification in wildfire spread has identified coherent spatial buffers along fire perimeters, informing emergency planning [70].

2.6 Advanced Uncertainty Decomposition and Metrics

Accurate decomposition of uncertainty is vital for interpretability. Information-theoretic measures have been refined, with Expected Pairwise KL-Divergence proposed as a superior alternative to mutual information for quantifying epistemic uncertainty, particularly in OOD detection [71]. Axiomatic assessments have highlighted violations in standard entropy-based measures for regression, advocating for variance-based measures that satisfy non-negativity and location-shift invariance [72].

Disentanglement Error (DE) metrics have been introduced to quantify the orthogonality of aleatoric and epistemic estimates, revealing that Information Theoretic disentanglement often outperforms Gaussian Logits approaches [73]. In regression, Proper Scoring Rules (e.g., CRPS, Logarithmic) provide a unified framework for decomposing pointwise risk into Bayes risk (aleatoric) and excess risk (epistemic), offering theoretically grounded uncertainty measures [74], [75]. Additionally, the "Epistemic Uncertainty Hole" phenomenon has been identified, where large models with limited data exhibit collapsing epistemic uncertainty, necessitating careful model selection and regularization [17].

2.7 Toolkits and Infrastructure

The proliferation of methods has spurred the development of specialized toolkits. Fortuna offers a unified interface for conformal prediction, post-hoc calibration, and scalable Bayesian inference within the JAX/Flax ecosystem [76]. TyXe provides a modular pipeline for converting standard PyTorch models into BNNs using Pyro effect handlers, decoupling architecture from inference [77]. Borch serves as a universal probabilistic programming language with native PyTorch integration, enabling seamless conversion of modules to Bayesian layers [78]. Surveys of these toolkits highlight a divergence between comprehensive method coverage (e.g., UQ360) and modular framework integration, indicating a need for standardized benchmarks [79].

3. Challenges and Outlook

Despite significant progress, Bayesian Deep Learning faces several critical challenges that hinder its widespread adoption in real-world systems.

3.1 Computational Intractability and Scalability

Exact Bayesian inference remains computationally prohibitive for large-scale models [8], [80]. While approximations like MFVI and MC Dropout scale well, they often sacrifice accuracy and calibration [3], [13]. MCMC methods, though accurate, suffer from slow mixing and high memory requirements, limiting their use to small networks or specific scientific applications [32], [29]. Future research must focus on developing scalable inference algorithms that retain the theoretical guarantees of MCMC while approaching the efficiency of point-estimate training. Techniques like subspace inference [81], sparse variational methods [82], and linearized sampling [49] offer promising directions.

3.2 The Cold Posterior and Prior Misspecification

The “cold posterior” effect, where artificially sharpening the posterior (tempering) improves performance, challenges the validity of standard Bayesian priors in deep learning [83], [33]. This suggests that standard priors (e.g., isotropic Gaussians) may be misspecified for deep networks, or that data augmentation and other practical tricks violate Bayesian assumptions [83]. Understanding the interplay between priors, likelihoods, and data augmentation is crucial for developing principled Bayesian frameworks that do not require ad-hoc tempering.

3.3 Failure Modes in Overparameterized Regimes

Recent studies reveal that standard Bayesian approximations can fail in overparameterized regimes. The “Epistemic Uncertainty Hole” demonstrates that large models trained on small datasets may exhibit collapsing epistemic uncertainty, failing to signal OOD inputs [17]. Similarly, MC Dropout has been shown to yield constant variance across the input space, failing to increase uncertainty in interpolation gaps or extrapolation regions [13], [84]. Addressing these failures requires new architectural inductive biases or inference methods that correctly scale uncertainty with model capacity and data density.

3.4 Evaluation and Benchmarking Bottlenecks

A lack of standardized evaluation metrics and benchmarks hampers progress. Many studies rely on synthetic data or limited benchmarks, making it difficult to compare methods fairly [5], [2]. The absence of ground truth uncertainty in real-world datasets complicates validation, often forcing reliance on proxy metrics like OOD detection AUROC or calibration error, which may not fully capture uncertainty quality [6], [85]. Comprehensive benchmarks covering diverse domains, data regimes, and uncertainty types are needed to drive the field forward. Recent efforts like the UQ benchmark for prognostics [86] and regression benchmarks [87] are steps in this direction.

3.5 Safety, Robustness, and Deployment

In safety-critical applications, uncertainty estimates must be reliable under distributional shift and adversarial attacks. While Bayesian methods theoretically offer robustness, empirical results show mixed performance, with some methods failing to detect sophisticated adversarial examples [88]. Integrating uncertainty quantification with robustness certification and conformal prediction [76], [51] is essential for deploying trustworthy AI. Furthermore, the computational overhead of Bayesian inference poses challenges for real-time deployment in resource-constrained environments like edge devices or autonomous vehicles [89], [90]. Developing lightweight Bayesian models and efficient inference engines is a priority for practical adoption.

3.6 Future Directions

Future research should explore the integration of Bayesian Deep Learning with foundation models and large language models, where uncertainty quantification is critical for mitigating hallucinations and ensuring reliability [80], [40]. The development of hybrid neurosymbolic frameworks that combine Bayesian uncertainty with logical constraints offers a promising avenue for scientific discovery and reasoning [91]. Additionally, advancing theoretical understanding of uncertainty in infinite-width limits and non-identifiable manifolds [15], [92] will provide deeper insights into the behavior of practical finite-width networks. Finally, fostering interdisciplinary collaboration to apply BDL in emerging domains such as quantum computing, genomics, and climate science will further validate and refine these methodologies.

4. Conclusion

Bayesian Deep Learning has established itself as a vital framework for quantifying uncertainty in deep neural networks, offering a principled approach to distinguishing between data noise and model ignorance. Through a comprehensive survey of 300 papers, we have highlighted the diversity of methodological approaches, from variational inference and MCMC to deep ensembles and post-hoc linearizations. While significant strides have been made in scalability and application-specific adaptations, challenges related to computational cost, prior specification, and failure modes in overparameterized regimes remain. Addressing these challenges requires a concerted effort in developing more robust inference algorithms, standardized evaluation protocols, and theoretically grounded architectures. As the field matures, Bayesian Deep Learning promises to play a central role in building trustworthy, reliable, and safe AI systems capable of operating effectively in the complex and uncertain real world.

References

[1] Moloud Abdar, Farhad Pourpanah, Sadiq Hussain. (2021). A Review of Uncertainty Quantification in Deep Learning: Techniques, Applications and Challenges. arXiv:2011.06225.

[2] Jakob Gawlikowski, Cedrique Rovile Njieutcheu Tassi, Mohsin Ali. (2023). A Survey of Uncertainty in Deep Neural Networks. arXiv:2107.03342.

[3] Ethan Goan, Clinton Fookes. (2017). Bayesian Neural Networks: An Introduction and Survey. arXiv:2006.12024.

[4] Laurent Valentin Jospin, Hamid Laga, Farid Boussaid. (2022). Hands-on Bayesian Neural Networks - A Tutorial for Deep Learning Users. arXiv:2007.06823.

[5] Wenchong He, Zhe Jiang, Tingsong Xiao. (2023). A Survey on Uncertainty Quantification Methods for Deep Learning. arXiv:2302.13425.

[6] Fahimeh Fakour, Ali Mosleh, Ramin Ramezani. (2024). A Structured Review of Literature on Uncertainty in Machine Learning & Deep Learning. arXiv:2406.00332.

[7] Daniel T. Chang. (2024). Probabilistic Deep Learning with Probabilistic Neural Networks and Deep Probabilistic Models. arXiv:2106.00120.

[8] Tom Charnock, Laurence Perreault-Levasseur, François Lanusse. (2024). Bayesian Neural Networks. arXiv:2006.01490.

[9] Daniel T. Chang. (2024). Bayesian Neural Networks: Essentials. arXiv:2106.13594.

[10] Aliaksandr Hubin, Geir Storvik. (2024). Variational Inference for Bayesian Neural Networks under Model and Parameter Uncertainty. arXiv:2305.00934.

[11] Lukas Mosser, Ehsan Zabihi Naeini. (2021). Calibration and Uncertainty Quantification of Bayesian Convolutional Neural Networks for Geophysical Applications. arXiv:2105.12115.

[12] Ronald Seoh. (2019). Qualitative Analysis of Monte Carlo Dropout. arXiv:2007.01720.

[13] Aslak Djupskås, Signe Riemer-Sørensen, Alexander Johannes Stasik. (2024). Unreliable Monte Carlo Dropout Uncertainty Estimation. arXiv:2512.14851.

[14] Riccardo Farris, Emanuele Telari, Nongnuch Artrith. (2025). BAYESIAN NEURAL NETWORKS VERSUS DEEP ENSEMBLES FOR UNCERTAINTY QUANTIFICATION IN MACHINE LEARNING INTERATOMIC POTENTIALS. arXiv:2509.19180.

[15] Florian Juengermann, Maxime Laasri, Marius Merkle. (2024). An Overview of Uncertainty Quantification Methods for Infinite Neural Networks. arXiv:2201.04746.

[16] Joseph Wilson, Chris van der Heide, Liam Hodgkinson. (2025). Is the Last Layer Sufficient for Uncertainty Quantification?. arXiv:2605.30741.

[17] M. Fellaji, F. Pennerath. (2024). The Epistemic Uncertainty Hole: an issue of Bayesian Neural Networks. arXiv:2407.01985.

[18] Stephen Brown, William L. Rodi, Marco Seracini. (2023). Bayesian Neural Networks for Geothermal Resource Assessment: Prediction with Uncertainty. arXiv:2209.15543.

[19] Hector J. Hortua, Luz Angela Garcia, Leonardo Castaneda C.. (2024). Constraining cosmological parameters from N-body simulations with Variational Bayesian Neural Networks. arXiv:2301.03991.

[20] Pavia Bera, Sanjukta Bhanja. (2025). QUANTIFICATION OF UNCERTAINTIES IN PROBABILISTIC DEEP NEURAL NETWORK BY IMPLEMENTING BOOSTING OF VARIATIONAL INFERENCE. arXiv:2503.13909.

[21] Trung Trinh, Samuel Kaski, Markus Heinonen. (2024). Scalable Bayesian neural networks by layer-wise input augmentation. arXiv:2010.13498.

[22] Tim G. J. Rudner, Zonghao Chen, Yee Whye Teh. (2022). Tractable Function-Space Variational Inference in Bayesian Neural Networks. arXiv:2312.17199.

[23] Yadi Wei, Roni Khardon. (2024). Variational Inference on the Final-Layer Output of Neural Networks. arXiv:2302.02420.

[24] Jihao Andreas Lin, Joe Watson, Pascal Klink. (2021). Function-Space Regularization for Deep Bayesian Classification. arXiv:2307.06055.

[25] Hrushikesh Loya, Pranav Poduval, Deepak Anand. (2024). UNCERTAINTY ESTIMATION IN CANCER SURVIVAL PREDICTION. arXiv:2003.08573.

[26] Yuexi Wang, Veronika Ročková. (2024). Uncertainty Quantification for Sparse Deep Learning. arXiv:2002.11815.

[27] Jary Pomponi, Simone Scardapane, Aurelio Uncini. (2021). Bayesian Neural Networks With Maximum Mean Discrepancy Regularization. arXiv:2003.00952.

[28] Jiayi Huang, Sangwoo Park, Osvaldo Simeone. (2023). Calibration-Aware Bayesian Learning. arXiv:2305.07504.

[29] Wenlong Chen, Bolian Li, Ruqi Zhang. (2023). Bayesian Computation in Deep Learning. arXiv:2502.18300.

[30] Govinda Anantha Padmanabha, Cosmin Safta, Nikolaos Bouklas. (2024). Condensed Stein Variational Gradient Descent for Uncertainty Quantification of Neural Networks. arXiv:2412.16462.

[31] Ali Siahkoohi, Gabrio Rizzuti, Felix J. Herrmann. (2024). Uncertainty quantification in imaging and automatic horizon tracking—a Bayesian deep-prior based approach. arXiv:2004.00227.

[32] Devina Mohan, Anna M. M. Scaife. (2024). Evaluating Bayesian deep learning for radio galaxy classification. arXiv:2405.18351.

[33] Devina Mohan, Anna M. M. Scaife, Fiona Porter. (2024). Quantifying Uncertainty in Deep Learning Approaches to Radio Galaxy Classification. arXiv:2201.01203.

[34] Vyron Arvanitis, Angelos Aslanidis, Emanuel Sommer. (2026). bde: A Python Package for Bayesian Deep Ensembles via MILE. arXiv:2605.14146.

[35] Neil K. Chada, Ajay Jasra, Kody J. H. Law. (2024). Bayesian Deep Learning with Multilevel Trace-class Neural Networks. arXiv:2203.12961.

[36] Gabriel Loaiza-Ganem, Valentin Villecroze, Yixin Wang. (2024). Deep Ensembles Secretly Perform Empirical Bayes. arXiv:2501.17917.

[37] Moritz A. Zanger, Yijun Wu, Pascal R. Van der Vaart. (2024). On the Equivalence of Random Network Distillation, Deep Ensembles, and Bayesian Inference. arXiv:2602.19964.

[38] Zaid Abulawia, Rui Hu, Prasanna Balaprakash. (2024). Bayesian Optimized Deep Ensemble for Uncertainty Quantification of Deep Neural Networks: a System Safety Case Study on Sodium Fast Reactor Thermal Stratification Modeling. arXiv:2412.08776.

[39] Matthias Kellner, Michele Ceriotti. (2024). Uncertainty quantification by direct propagation of shallow ensembles. arXiv:2402.16621.

[40] Oleksandr Balabanov, Hampus Linander. (2024). Uncertainty Quantification in Fine-Tuned LLMs Using LoRA Ensembles. arXiv:2402.12264.

[41] EuiYul Song. (2024). Uncertainty-Aware Perceiver. arXiv:2402.02433.

[42] Jiaxiang Yi, Miguel A. Bessa. (2024). Cooperative Bayesian and variance networks disentangle aleatoric and epistemic uncertainties. arXiv:2505.02743.

[43] Masoud Ataei, Vikas Dhiman. (2024). DADEE: Well-calibrated uncertainty quantification in neural networks for barriers-based robot safety. arXiv:2407.00616.

[44] Ali Unlu, Laurence Aitchison. (2021). Variational Laplace for Bayesian neural networks. arXiv:2011.10443.

[45] Ali Unlu, Laurence Aitchison. (2021). Variational Laplace for Bayesian neural networks. arXiv:2103.00222.

[46] Felix Fiedler, Sergio Lucia. (2024). Improved uncertainty quantification for neural networks with Bayesian last layer. arXiv:2302.10975.

[47] Runa Eschenhagen, Erik Daxberger, Philipp Hennig. (2021). Mixtures of Laplace Approximations for Improved Post-Hoc Uncertainty in Deep Learning. arXiv:2111.03577.

[48] James McInerney, Nathan Kallus. (2024). Variation Due to Regularization Tractably Recovers Bayesian Deep Learning. arXiv:2403.10671.

[49] Joseph Wilson, Chris van der Heide, Liam Hodgkinson. (2024). Uncertainty Quantification with the Empirical Neural Tangent Kernel. arXiv:2502.02870.

[50] Max-Heinrich Laves, Sontje Ihler, Jacob F. Fast. (2021). Recalibration of Aleatoric and Epistemic Regression Uncertainty in Medical Imaging. arXiv:2104.12376.

[51] Canberk Ekmekci, Mujdat Cetin. (2024). Conformalized Generative Bayesian Imaging: An Uncertainty Quantification Framework for Computational Imaging. arXiv:2504.07696.

[52] Marius Hobbhahn, Agustinus Kristiadi, Philipp Hennig. (2020). Fast Predictive Uncertainty for Classification with Bayesian Deep Networks. arXiv:2003.01227.

[53] Sai Munikoti, Deepesh Agarwal, Laya Das. (2024). A General Framework for quantifying Aleatoric and Epistemic uncertainty in Graph Neural Networks. arXiv:2205.09968.

[54] David González, Alba Muixí, Beatriz Moya. (2026). Variational Graph Neural Networks for Uncertainty Quantification in Inverse Problems. arXiv:2603.29515.

[55] Yasser Abduallah, Khalid A. Alobaid, Jason T. L. Wang. (2024). Prediction of the SYM-H Index Using a Bayesian Deep Learning Method with Uncertainty Quantification. arXiv:2402.17196.

[56] Tanwi Mallick, Prasanna Balaprakash, Jane Macfarlane. (2022). Deep-Ensemble-Based Uncertainty Quantification in Spatiotemporal Graph Neural Networks for Traffic Forecasting. arXiv:2204.01618.

[57] Ibai Ramirez, Jokin Alcibar, Joel Pino. (2025). BAYESIAN PHYSICS INFORMED NEURAL NETWORKS FOR RELIABLE TRANSFORMER PROGNOSTICS. arXiv:2509.15933.

[58] Khemraj Shukla, Zongren Zou, Theo Kaeufer. (2024). Uncertainty Quantification in PINNs for Turbulent Flows: Bayesian Inference and Repulsive Ensembles. arXiv:2604.17156.

[59] Benjamin J. Zhang, Siting Liu, Stanley J. Osher. (2025). Probabilistic operator learning: GENERATIVE MODELING AND UNCERTAINTY QUANTIFICATION for foundation models of differential equations. arXiv:2509.05186.

[60] Yesmine Abdennadher, Philip N. Garner. (2024). Practical Bayesian Inference for Speech SNNs: Uncertainty and Loss-Landscape Smoothing. arXiv:2604.08624.

[61] Tao Sun, Sander Bohte. (2024). Average-Over-Time Spiking Neural Networks for Uncertainty Estimation in Regression. arXiv:2412.00278.

[62] Mohaddeseh Chegini, Ali Mahloojifar. (2024). A Deep Bayesian Convolutional Spiking Neural Network-based CAD system with Uncertainty Quantification for Medical Images Classification. arXiv:2504.17819.

[63] Lohith Konathala. (2024). Bayesian Neural Networks for 2D MRI Segmentation. arXiv:2311.14875.

[64] Mohaddeseh Cheginia, Ali Mahloojifara. (2024). BI-RADS prediction of mammographic masses using uncertainty information extracted from a Bayesian Deep Learning model. arXiv:2503.13999.

[65] Mélodie Monod, Alessandro Micheli, Samir Bhatt. (2024). NeuralSurv: Deep Survival Analysis with Bayesian Uncertainty Quantification. arXiv:2505.11054.

[66] Mitchel J. Colebank. (2024). Uncertainty-aware classification and triage of structural heart disease using electrocardiography and echocardiography metrics. arXiv:2605.22968.

[67] Xinlei Xiong, Wenbo Hu, Shuxun Zhou. (2024). BRIDGING THE GAP BETWEEN BAYESIAN DEEP LEARNING AND ENSEMBLE WEATHER FORECASTS. arXiv:2511.14218.

[68] Laura A. Mansfield, Hannah M. Christensen. (2025). Epistemic and Aleatoric Uncertainty Quantification in Weather and Climate Models. arXiv:2511.23448.

[69] Spyros Kondylatos, Gustau Camps-Valls, Ioannis Papoutsis. (2025). Uncertainty-Aware Deep Learning for Wildfire Danger Forecasting. arXiv:2509.25017.

[70] Aditya Chakravarty. (2025). Spatial Uncertainty Quantification in Wildfire Forecasting for Climate-Resilient Emergency Planning. arXiv:2510.09666.

[71] Kajetan Schweighofer, Lukas Aichberger, Mykyta Ielanskyi. (2024). Introducing an Improved Information-Theoretic Measure of Predictive Uncertainty. arXiv:2311.08309.

[72] Christopher Bülte, Yusuf Sale, Timo Löhr. (2025). An Axiomatic Assessment of Entropy- and Variance-based Uncertainty Quantification in Regression. arXiv:2504.18433.

[73] Ivo Pascal de Jong, Andreea Ioana Sburlea, Matthia Sabatelli. (2024). Measuring Uncertainty Disentanglement Error in Classification. arXiv:2408.12175.

[74] Alexander Fishkov, Nikita Kotelevskii, Kajetan Schweighofer. (2025). UNCERTAINTY QUANTIFICATION FOR REGRESSION USING PROPER SCORING RULES. arXiv:2509.26610.

[75] Christopher Bülte, Yusuf Sale, Gitta Kutyniok. (2025). Uncertainty Quantification for Regression: A Unified Framework Based on Kernel Scores. arXiv:2510.25599.

[76] Gianluca Detommaso, Alberto Gasparin, Michele Donini. (2023). Fortuna: A Library for Uncertainty Quantification in Deep Learning. arXiv:2302.04019.

[77] Hippolyt Ritter, Theofanis Karaletsos. (2021). TyXe: Pyro-based Bayesian neural nets for Pytorch. arXiv:2110.00276.

[78] Unknown Author 1, Unknown Author 2, Unknown Author 3. (2022). BORCH: A DEEP UNIVERSAL PROBABILISTIC PROGRAMMING LANGUAGE. arXiv:2209.06168.

[79] Maximilian Pintz, Joachim Sicking, Maximilian Poretschkin. (2022). A SURVEY ON UNCERTAINTY TOOLKITS FOR DEEP LEARNING. arXiv:2205.01040.

[80] Theodore Papamarkou, Maria Skoularidou, Konstantina Palla. (2024). Position: Bayesian Deep Learning is Needed in the Age of Large-Scale AI. arXiv:2402.00809.

[81] Sanket Jantre, Nathan M. Urban, Xiaoning Qian. (2024). LEARNING ACTIVE SUBSPACES FOR EFFECTIVE AND SCALABLE UNCERTAINTY QUANTIFICATION IN DEEP NEURAL NETWORKS. arXiv:2309.03061.

[82] Junbo Li, Zichen Miao, Qiang Qiu. (2024). TRAINING BAYESIAN NEURAL NETWORKS WITH SPARSE SUBSPACE VARIATIONAL INFERENCE. arXiv:2402.11025.

[83] Julyan Arbel, Konstantinos Pitas, Mariia Vladimirova. (2024). A Primer on Bayesian Neural Networks: Review and Debates. arXiv:2309.16314.

[84] Francesco Verdoja, Ville Kyrki. (2020). Notes on the Behavior of MC Dropout. arXiv:2008.02627.

[85] Daniel Ries, Joshua Michalenko, Tyler Ganter. (2022). Comparing the quality of neural network uncertainty estimates for classification problems. arXiv:2308.05903.

[86] Luis Basora, Arthur Viens, Manuel Arias Chao. (2024). A Benchmark on Uncertainty Quantification for Deep Learning Prognostics. arXiv:2302.04730.

[87] Brian Staber, Sebastien Da Veiga. (2024). Benchmarking Bayesian neural networks and evaluation metrics for regression tasks. arXiv:2206.06779.

[88] Carsten Hartmann, Lorenz Richter. (2023). Transgressing the boundaries: towards a rigorous understanding of deep learning and its (non-)robustness. arXiv:2307.02454.

[89] Hans Weytjens, Jochen De Weerdt. (2024). Learning Uncertainty with Artificial Neural Networks for Predictive Process Monitoring. arXiv:2206.06317.

[90] Richard J. Licata, Piyush M. Mehta. (2022). Uncertainty Quantification Techniques for Space Weather Modeling: Thermospheric Density Application. arXiv:2201.02067.

[91] Shahnawaz Alam, Mohammed Mudassir Uddin, Mohammed Kaif Pasha. (2024). Constraint-Aware Neurosymbolic Uncertainty Quantification with Bayesian Deep Learning for Scientific Discovery. arXiv:2601.12442.

[92] David Rugamer. (2024). On the Epistemic Uncertainty of Overparametrized Neural Networks. arXiv:2605.25234.