

Vision-Language Models for Embodied Reasoning: Architectures, Grounding, and Interactive Decision-Making

1 Introduction

Embodied reasoning represents a fundamental paradigm shift from static, passive vision-language understanding to dynamic agency, where intelligent systems must actively perceive, reason, and act within continuous physical environments. Unlike traditional multimodal models that process isolated image-text pairs, embodied agents operate under the constraints of temporal continuity, partial observability, and the necessity for real-time environmental feedback. This subsection establishes the conceptual foundation of this field, distinguishing it by the imperative of active physical interaction and the challenge of grounding abstract linguistic symbols in high-dimensional sensorimotor streams. The scope of this survey encompasses the evolution from modular, decoupled pipelines to unified foundation models, analyzing how recent architectures bridge the gap between semantic reasoning and low-level control.

Historically, embodied AI research relied on modular frameworks that separated perception, planning, and execution. Early approaches utilized probabilistic grammars and symbolic planners to map language to actions, often requiring extensive feature engineering and domain-specific knowledge [1; 2]. While these neuro-symbolic hybrids offered interpretability, they struggled with the scalability and robustness required for open-world interactions. The advent of deep reinforcement learning enabled end-to-end training from pixels to actions, demonstrating the emergence of grounded communication protocols without explicit supervision [3; 4]. However, these early deep models often lacked the broad semantic knowledge necessary for generalization beyond narrow task distributions.

The current landscape is defined by the integration of Large Language Models (LLMs) and Vision-Language Models (VLMs) into embodied loops, creating a new class of Vision-Language-Action (VLA) models. This transition addresses the “symbol grounding problem” in continuous state spaces, where discrete natural language concepts must be mapped to continuous proprioceptive signals and visual observations. Recent works such as [5] and [6] exemplify this shift by directly injecting continuous sensor modalities into language model latent spaces, enabling zero-shot generalization across diverse robotic embodiments. Furthermore, approaches like [7] highlight the critical need to constrain LLM-generated plans with physical affordances to ensure feasibility, moving beyond purely semantic reasoning to physically grounded decision-making.

Despite these advances, core challenges persist in aligning the discrete, symbolic nature of language with the continuous, stochastic dynamics of the physical world. The discrepancy between allocentric training data and egocentric robotic control remains a significant hurdle, necessitating novel architectural biases for spatial schema learning [8]. Additionally, long-horizon tasks require mechanisms for maintaining temporal consistency and memory, prompting the development of inner monologue frameworks that leverage environmental feedback for self-correction [9]. As the field progresses towards artificial general intelligence, the convergence of multimodal perception, generative world modeling, and interactive action becomes paramount. This survey critically examines these architectural paradigms, reasoning mechanisms, and learning strategies, aiming to provide a comprehensive taxonomy that guides future research in building trustworthy, generalizable, and truly embodied intelligent agents.

2 Architectural Paradigms for Multimodal Perception and Action

2.1 Unified Vision-Language-Action Transformers

The emergence of Unified Vision-Language-Action (VLA) Transformers represents a paradigmatic shift in embodied AI, moving from modular pipelines to monolithic architectures that treat robotic control as a next-token prediction problem. By tokenizing continuous sensorimotor streams alongside visual observations and linguistic instructions, these models seamlessly integrate perception, reasoning, and action into a single autoregressive sequence. This formulation allows VLAs to leverage the vast semantic knowledge encoded in internet-scale pre-training, enabling zero-shot generalization to novel objects and commands that were absent during robotic training [10]. The core mechanism involves mapping low-level motor commands, such as joint positions or end-effector poses, into discrete tokens compatible with large language model vocabularies, thereby unifying the objective functions of language modeling and policy learning [11].

Early iterations of this approach, such as [10], demonstrated that co-fine-tuning vision-language models on robotic trajectory data alongside web-scale corpora yields emergent capabilities, including the ability to interpret abstract commands like “pick up the thing you drink from” without explicit task-specific supervision. This transfer of semantic reasoning is further enhanced by architectures like [5], which incorporates continuous state estimation directly into the language model’s input stream, establishing a robust link between words and percepts across diverse embodiments. Similarly, [11] illustrates that open-source VLAs trained on diverse demonstration collections can outperform larger closed models, highlighting the critical role of data diversity and efficient fine-tuning strategies like Low-Rank Adaptation (LoRA) in scaling embodied intelligence.

A significant advancement within this domain is the integration of Chain-of-Thought (CoT) reasoning directly into the action generation process. While standard VLAs predict actions autoregressively, recent works argue that purely reactive policies lack the deliberative capacity required for complex, multi-stage manipulation. [12] introduces Embodied CoT (ECoT), training VLAs to generate intermediate reasoning steps regarding plans, sub-tasks, and visually grounded features before outputting motor commands. This approach significantly improves success rates on generalization tasks by forcing the model to ground its semantic reasoning in sensory observations. Complementary to this, [13] utilizes prefix tuning on large-scale embodied planning datasets to generate high-quality sub-goals, creating a closed loop between high-level planning and low-level control. Furthermore, [14] employs CoT to predict contact points and end-effector directions, enhancing stability in object-centric tasks through active impedance adaptation.

Despite their promise, unified transformers face challenges related to the discretization of continuous action spaces and the computational latency inherent in autoregressive generation. Methods like [15] address the alignment gap by predicting intermediate 2D “visual traces” to better synchronize vision and action spaces. Additionally, [16] explores state space models as an alternative to standard transformers to achieve linear inference complexity, addressing the latency constraints of real-time control. Nevertheless, the dominant trend remains the scaling of monolithic transformers, where the boundary between semantic understanding and physical actuation continues to blur. Future research must address the trade-off between the rich semantic generalization afforded by large-scale pre-training and the precise, high-frequency control loops required for dynamic physical interaction, potentially through hybrid architectures that combine the reasoning depth of VLAs with the efficiency of specialized control heads.

2.2 Neuro-Symbolic Hybrid Frameworks for Interpretable Planning

Bridging the gap between the monolithic flexibility of Unified Vision-Language-Action (VLA) Transformers and the rigorous geometric requirements of 3D agency, Neuro-symbolic hybrid frameworks represent a pivotal architectural paradigm that strategically decouples high-level semantic planning from low-level motor execution. While the preceding section highlighted how VLAs unify perception and action into a single autoregressive stream, neuro-symbolic systems address the resulting challenges in verifiable safety and interpretability by leveraging the semantic flexibility of Large Language Models (LLMs) alongside the rigorous guarantees of symbolic planners and formal logic. A dominant manifestation of this approach is “code-as-policies,” where LLMs synthesize executable Python scripts or domain-specific programs to orchestrate perception APIs and control primitives. For instance, [17] demonstrates how code-generation models can compose reactive policies capable of precise geometric reasoning and arithmetic calculations, effectively translating ambiguous natural language commands into deterministic robot behaviors. Similarly, [18] utilizes code generation to modularly compose vision-language subroutines, achieving state-of-the-art performance on complex visual tasks without task-specific training by explicitly separating reasoning steps from perceptual execution.

To ensure long-horizon consistency and safety, advanced frameworks map linguistic intents to formal logical representations, such as Linear Temporal Logic (LTL) or Planning Domain Definition Language (PDDL), prior to execution. [19] introduces a grounded semantic parser that hypothesizes LTL formulas to capture temporal constraints, ensuring that generated plans adhere to strict logical validity before being passed to a symbolic executor. This disentanglement allows for the verification of plan feasibility and the mitigation of hallucination risks inherent in purely neural approaches. Furthermore, [20] illustrates how discrete symbols and probabilistic rules can be distilled from continuous sensorimotor experiences, enabling off-the-shelf symbolic planners to operate on knowledge extracted directly from robot interaction.

The integration of neuro-symbolic reasoning extends to dynamic environment adaptation through structured scene representations, a necessity that anticipates the detailed discussion on 3D spatial encoding in the subsequent section. Systems like [21] construct semantic maps from visual observations which are then queried by LLMs for sub-goal planning, effectively bridging the gap between continuous spatial data and discrete logical inference. This modular design facilitates robust failure recovery; [22] employs a dual-role LLM to generate execution constraints and utilizes Vision-Language Models (VLMs) to detect violations, triggering immediate replanning when physical outcomes diverge from symbolic intentions. Additionally, [23] highlights the efficacy of using temporal logic formulae to reason about long-term event evolution, significantly outperforming end-to-end models in scenarios requiring memory and temporal coherence.

Despite their advantages, neuro-symbolic frameworks face significant challenges in the seamless grounding of abstract symbols to noisy, continuous sensory streams, particularly when lacking the explicit 3D priors discussed later. The “symbol grounding problem” persists when neural perception modules misidentify objects, leading to logically valid but physically impossible symbolic plans. Emerging trends address this by incorporating probabilistic abduction, as seen in [24], which reasons over probabilistic scene representations to handle perceptual uncertainty. Future research must focus on tightening the feedback loop between symbolic verification and neural perception, potentially through differentiable logic layers that allow gradient-based refinement of symbolic parameters based on execution feedback. This synthesis aims to unify the data efficiency and logical rigor of symbolic methods with the generalization power of deep learning, while relying on

the robust 3D representations that form the foundation of the following discourse.

2.3 Geometric Encoding and 3D Spatial Representations

The transition from passive 2D perception to active 3D agency necessitates architectural paradigms capable of encoding metric geometry and volumetric structure within language-model latent spaces. Standard vision-language models, predominantly trained on internet-scale image-text pairs, often exhibit profound deficiencies in spatial reasoning, struggling with basic directional concepts and metric relationships due to the lack of explicit 3D priors in their training corpora [25; 26]. To bridge this gap, contemporary architectures employ three primary strategies: direct integration of depth and point cloud features, the fusion of neural implicit fields with semantic embeddings, and the learning of egocentric spatial schemas through embodied interaction.

The most direct approach involves augmenting visual encoders with explicit geometric data. Methods such as [27] demonstrate that feeding both RGB and aligned depth images significantly enhances a model’s ability to answer metric queries, while [28] introduces flexible plugin modules to inject depth information into existing visual encoders, enabling precise reasoning about relative distances and directions. Beyond 2.5D depth maps, architectures like [29] and [30] process raw 3D point clouds directly, utilizing specialized 3D encoders (e.g., ReCon++) to extract geometric features that are projected into the language embedding space. This allows for open-vocabulary querying of 3D scenes, supporting tasks ranging from dense captioning to task decomposition [29; 30]. However, these discrete representations can be sparse and noisy; consequently, [31] proposes leveraging rich 2D semantics during training to assist 3D grounding, effectively distilling semantic knowledge from the 2D domain to overcome the inherent limitations of point cloud data.

A more recent innovation lies in leveraging neural implicit representations, such as Neural Radiance Fields (NeRFs) and voxel grids, to create continuous, searchable scene models. Frameworks like [32] fuse pretrained visual-language features with 3D reconstructions to generate spatial maps that support natural language indexing without additional labeled data. Similarly, [33] utilizes implicit fields to identify objects in 3D scenes, employing large language models to decompose complex queries and evaluate spatial relations in a zero-shot manner. These approaches address the “amodal” nature of occlusion, allowing agents to reason about object permanence and hidden structures more effectively than passive 2D systems [34]. Furthermore, [6] extends this by integrating generative world models, predicting future goal images and point clouds to link 3D perception with action planning, thereby capturing the dynamics of the physical world often neglected in static representations.

Despite these advances, a critical challenge remains in aligning allocentric 3D data with the egocentric perspective required for robotic control. Models trained on static scans often fail to generalize to the dynamic viewpoint changes inherent in embodiment. To mitigate this, [35] integrates both scene-level and ego-centric 3D features, utilizing ego-centric frames for feature alignment to enhance localization of small objects. Additionally, [36] and [37] argue that explicit spatial representations, such as top-down semantic maps, provide a stronger grounding signal than raw egocentric observations alone, facilitating more robust long-horizon navigation. Nevertheless, empirical evaluations indicate that even state-of-the-art models lag significantly behind human performance in top-view and complex spatial reasoning tasks [38], suggesting that current architectural interventions are necessary but insufficient. Future directions must likely move beyond simple feature fusion toward architectures that inherently model spatial causality and physical constraints, potentially through multi-sensory integration as explored in [39], to achieve true geometric fluency in embodied agents.

2.4 Memory-Augmented Architectures for Long-Horizon Context

Building upon the geometric foundations established in 3D perception, long-horizon embodied reasoning necessitates architectures capable of maintaining temporal consistency and spatial fidelity across partially observable environments, where standard transformer context windows often prove insufficient. Memory-augmented paradigms address this limitation by decoupling the retention of historical state from immediate processing, enabling agents to construct persistent cognitive maps and retrieve relevant precedents via attention mechanisms. A primary architectural strategy involves the integration of external vector memory banks, which store compressed representations of past visual-language experiences. Systems like [40] exemplify this by maintaining a topological map that encodes navigation history, viewpoints, and object relationships, serving as a global action space that enhances interpretability and strategic diversity. Similarly, [41] unifies egocentric memory and control through closed-loop multi-round dialogues, allowing the agent to recall prior observations for active exploration and embodied question answering. These retrieval-augmented approaches mitigate the “forgetting” phenomenon inherent in recurrent models, allowing for dynamic re-weighting of historical evidence based on current query relevance.

Beyond static retrieval, advanced frameworks employ generative mechanisms to construct and update internal cognitive maps that fuse temporal, spatial, and semantic elements, effectively grounding the abstract memory structures in the physical world. The [42] pioneered the use of explicit attention hops to select relevant spatial regions from memory for spatial inference, a concept evolved in modern systems like [32], which fuses pretrained visual-language features directly with 3D reconstructions. This integration allows for natural language indexing of the map without additional labeled data, enabling precise localization of open-vocabulary goals. Furthermore, [36] posits that language-vision association is strongest within explicit spatial representations, utilizing cross-modal attention over predicted top-down semantics to guide waypoint prediction. Such cognitive maps facilitate not only localization but also counterfactual simulation; for instance, [43] prompts large language models to act as world models, imagining next observations to validate navigational decisions before execution, thereby bridging the domain gap between offline training and online deployment.

Crucially, to transition from passive mapping to active agency, memory architectures must support closed-loop feedback mechanisms that refine plans based on execution outcomes, directly addressing the alignment challenges noted in 3D perception. The concept of an “inner monologue,” as formalized in [9], leverages environmental feedback—such as success detection and scene description—to form a recurring reasoning loop. This allows agents to self-correct and replan by processing the discrepancy between intended and actual states. Extending this, [22] introduces a dual-role framework where language models generate constraints to detect misalignments between high-level plans and low-level execution, triggering immediate recovery protocols. Similarly, [44] utilizes the physical grounding of vision-language models to adapt actions dynamically when initial plans fail due to unforeseen obstacles. These feedback loops transform memory from a passive store into an active component of the reasoning process, enabling robust adaptation in dynamic settings.

Despite these advances, challenges remain in balancing the granularity of memory representation with computational efficiency, particularly when integrating the dense geometric data discussed previously. While external memory offers scalability, the latency of retrieval and the potential for noise accumulation in long sequences pose significant hurdles. Future directions likely involve hybrid architectures that combine the precise geometric reasoning of neural implicit fields with the semantic flexibility of large language models, potentially utilizing hierarchical memory structures

to manage abstraction levels. As agents move toward lifelong learning, the development of memory consolidation mechanisms that prevent catastrophic forgetting while integrating new experiences will be paramount for achieving truly autonomous, generalizable embodied intelligence.

3 Reasoning Mechanisms and Cognitive Capabilities

3.1 Hierarchical Task Decomposition and Structured Planning

Hierarchical task decomposition serves as the cognitive bridge translating abstract natural language directives into executable sequences of motor primitives, addressing the fundamental disconnect between high-level semantic goals and low-level physical constraints. Early approaches often relied on monolithic policy generation, which struggled with long-horizon dependencies and compositional generalization. In contrast, contemporary frameworks increasingly adopt structured planning mechanisms that leverage Chain-of-Thought (CoT) reasoning to explicitly generate intermediate sub-goals before action execution. For instance, [45] demonstrated that contextualized language models could generate multi-step plans from high-level instructions alone, while [13] advanced this by integrating egocentric visual contexts to refine sub-goal generation, significantly improving success rates in complex manipulation tasks. Furthermore, [12] highlights that naive CoT prompting is insufficient for robotics; instead, policies must be trained to reason about visually grounded features, such as object bounding boxes and end-effector positions, to ensure semantic plans remain physically realizable.

Parallel to sequential reasoning, program synthesis has emerged as a robust paradigm for encoding complex logic and geometric constraints directly into executable code. By treating robot control as a code generation problem, frameworks like [17] enable agents to synthesize Python scripts that invoke perception APIs and control primitives, facilitating precise arithmetic and spatial reasoning unattainable through standard token prediction. This neuro-symbolic approach is further enhanced by [46], which integrates foundation models like SAM and CLIP into the generated code loop, allowing for dynamic perception-planning-action cycles. Similarly, [47] utilizes LLMs to compose 3D value maps from free-form instructions, grounding semantic knowledge into the agent’s observation space to synthesize robust, closed-loop trajectories zero-shot. These methods collectively shift the burden from learning implicit policies to explicitly constructing verifiable behavioral logic.

To balance computational efficiency with solution optimality, recent research explores dual-process cognitive architectures inspired by human fast and slow thinking. [48] proposes a framework where an instruction discriminator routes tasks to either a reactive “fast” system for immediate responses or a deliberative “slow” system for complex reasoning, optimizing resource allocation. This deliberative capacity is often augmented by search algorithms; [49] introduces a framework where the LLM acts as both a world model and a reasoning agent, employing Monte Carlo Tree Search to explore alternative reasoning paths and anticipate future states, thereby overcoming the limitations of greedy decoding. Moreover, [50] presents a neuro-symbolic hybrid where an LLM collaborates with a traditional planner to handle noisy observations and uncertainty, achieving greater efficiency than pure neural baselines.

Despite these advances, challenges persist in maintaining consistency between high-level plans and dynamic environmental feedback. Methods like [9] address this by incorporating environmental feedback into a recurring language loop, allowing agents to self-correct and replan based on execution outcomes. Future directions must focus on unifying these disparate strategies into cohesive architectures that seamlessly integrate probabilistic inference with symbolic rigor, ensuring that hierarchical decomposition remains robust against the stochasticity of real-world physical interac-

tions.

3.2 Internal World Modeling and Counterfactual Simulation

Building upon the explicit behavioral logic and hierarchical decomposition strategies discussed previously, the construction of internal world models represents a pivotal shift from generating discrete action sequences to maintaining continuous, dynamic representations of the physical environment. This capacity enables agents to anticipate consequences and optimize actions through mental simulation, effectively reducing the reliance on costly physical trial-and-error. Unlike the reactive or code-synthesis approaches that dominate earlier paradigms, internal world modeling requires the agent to sustain a state representation that supports counterfactual inference and deliberative planning. Recent frameworks have repurposed Large Language Models (LLMs) to serve as both reasoning agents and implicit world models; for instance, the Reasoning via Planning (RAP) framework utilizes LLMs to incrementally build reasoning trees, simulating future states and rewards to balance exploration and exploitation in complex spaces [49]. Similarly, the “Dream” phase in the DECKARD agent employs LLMs to hypothesize Abstract World Models (AWMs) regarding subgoal transitions, which are subsequently verified through environmental interaction, thereby blending noisy semantic priors with grounded dynamics [51].

A critical challenge in evolving from symbolic planning to continuous simulation is bridging the gap between abstract linguistic reasoning and 3D spatial dynamics, as traditional language models often lack the geometric fidelity required for precise manipulation. To address this, generative world models like 3D-VLA integrate 3D perception with action by employing diffusion models to predict future goal images and point clouds, effectively allowing the agent to “imagine” the physical outcome of an action before execution [6]. Complementary approaches, such as Mental Imagery for Robotic Affordances (MIRA), leverage novel-view synthesis to construct consistent 3D scene representations, enabling agents to optimize actions over imagined orthographic views for 6-DoF manipulation tasks [52]. Furthermore, the NavCoT strategy enhances Vision-and-Language Navigation by prompting LLMs to act as world models that forecast next observations based on instructions, selecting actions that best align with these imagined futures [53].

Extending these predictive capabilities, counterfactual reasoning allows agents to evaluate “what-if” scenarios, a necessity for robust decision-making under uncertainty and for establishing causal understanding beyond statistical correlation. Benchmarks like CLEVRER have highlighted the deficiency of standard models in handling counterfactual queries, necessitating architectures that explicitly model causal structures [54]. Innovative frameworks such as CounterCurate utilize semantic counterfactuals generated by high-performing models to fine-tune visio-linguistic systems, significantly improving compositional reasoning and physical grounding [55]. Moreover, the CFMM benchmark reveals that current Multimodal Large Language Models (MLLMs) often struggle to override visual evidence with counterfactual presuppositions, indicating a significant frontier for improving hypothetical reasoning [56].

To further ground these simulations and mitigate the limitations of purely internal computation, recent works introduce visual intermediates as cognitive scaffolds. The Visual Sketchpad framework enables models to draw auxiliary lines and marks, externalizing working memory and facilitating spatial reasoning akin to human diagrammatic thought [57]. This aligns with the “Mind’s Eye” paradigm, where computational physics engines simulate outcomes to provide grounded context for language models, drastically improving performance on physical reasoning tasks [58]. Despite these advances, limitations persist in the fidelity of simulated dynamics and the computational cost of

iterative look-ahead processes. Consequently, the field is moving toward integrating these simulation capabilities with explicit physical commonsense and causal inference, ensuring that internal models not only predict visual outcomes but also adhere to the underlying laws of physics, a transition explored in the subsequent section.

3.3 Physical Commonsense and Causal Inference

The integration of physical commonsense and causal inference represents a pivotal transition in embodied reasoning, moving agents from passive pattern recognition to active understanding of dynamic physical laws. While large language models encode vast semantic knowledge, they often lack the grounded experience necessary to reason about unstated physical constraints, leading to hallucinations in manipulation tasks [59; 60]. To bridge this gap, recent architectures explicitly model object affordances and material properties, enabling agents to deduce functionality beyond visual appearance. For instance, frameworks leveraging pre-trained vision-language models have demonstrated the ability to generalize to unseen object categories by reasoning about usage rather than relying on fixed class labels [61; 62]. Furthermore, specialized datasets and fine-tuning strategies, such as those incorporating tactile feedback or physical concept annotations, have proven essential for embedding tacit knowledge of fragility, mass, and friction into model latent spaces [63; 64; 65].

Beyond static properties, genuine embodied intelligence requires robust causal reasoning to predict the temporal evolution of scenes and the consequences of actions. Traditional video understanding models often struggle with counterfactual and explanatory queries, frequently failing to distinguish between correlation and causation [54; 66]. Advanced approaches address this by integrating differentiable physics engines or neural dynamics models that allow agents to simulate future states and validate plans before execution [67; 68]. By treating physical interaction as a generative process, systems like PIGLeT and Mind’s Eye utilize internal simulation to forecast outcomes of linguistic commands, significantly improving accuracy in “what-if” scenarios compared to purely text-based reasoning [69; 58]. Moreover, causal intervention techniques have been employed to disentangle spurious cross-modal correlations, ensuring that agents base decisions on invariant causal structures rather than dataset biases [70; 71].

Despite these advances, significant challenges remain in scaling these capabilities to open-world environments. Current models often exhibit brittleness when faced with novel physical interactions or complex multi-object dynamics, as evidenced by performance gaps in benchmarks designed to test compositional physical reasoning [72; 73]. A promising direction involves neuro-symbolic hybrids that combine the flexibility of neural perception with the rigor of symbolic planners, allowing for verifiable reasoning about physical constraints [74; 21]. Additionally, the emergence of embodied chain-of-thought reasoning suggests that iterating through physically grounded intermediate steps—such as predicting contact points or force vectors—can enhance policy robustness [12]. Future research must prioritize the development of unified world models that seamlessly integrate intuitive physics with semantic planning, enabling agents to not only perceive the world but to truly understand the causal mechanisms governing it.

3.4 Active Perception and Uncertainty-Driven Exploration

Building upon the unified world models and causal mechanisms established in the previous section, active perception in embodied reasoning represents the operationalization of these internal simulations to resolve external uncertainty. While prior discussions focused on understanding physical

laws and predicting outcomes, active perception transcends passive observation by requiring agents to execute epistemic actions that actively reduce ambiguity in partially observable environments. This paradigm shift transforms the agent from a static interpreter into an active investigator, where decision-making is driven by the maximization of information gain rather than immediate task reward. Early frameworks grounded in Bayesian inference established the theoretical basis for ranking hypotheses based on sensor models and geometric relations [75], evolving into multi-world approaches that represent uncertainty about perceived scenes to answer complex queries under ambiguous inputs [76]. In navigation contexts, this manifests as agents selecting sensing actions—such as changing viewpoints or manipulating occluders—to minimize entropy in their belief states regarding goal locations or object properties.

Recent advancements leverage Large Language Models (LLMs) and Vision-Language Models (VLMs) to orchestrate these exploratory behaviors through semantic priors, effectively bridging the gap between abstract causal reasoning and concrete sensorimotor interaction. For instance, language-guided exploration algorithms utilize LLM commonsense to hypothesize Abstract World Models (AWMs), decomposing tasks into subgoals that drive efficient search in unseen environments [51]. Similarly, frameworks like LGX employ LLMs to make sequential navigational decisions for zero-shot object goal navigation, significantly outperforming baselines by integrating semantic reasoning with visual grounding [77]. To address the challenge of miscalibrated confidence which can lead to premature termination of exploration, conformal prediction techniques have been integrated with VLMs to calibrate stopping criteria, ensuring agents explore until statistically confident in their answers [78]. Furthermore, the integration of spatial mapping with language, as seen in LLM-Brain, unifies egocentric memory and control, enabling agents to maintain context over long horizons while actively querying the environment [41].

Beyond autonomous exploration, uncertainty-driven mechanisms increasingly incorporate adaptive help-seeking behaviors, serving as a precursor to the introspective self-correction loops discussed in the subsequent section. When internal reasoning reaches an impasse due to insufficient information, agents can initiate dialogue with human operators or external knowledge bases. The VNLA task exemplifies this by enabling agents to sense when they are lost and query an advisor for language subgoals, modeled through Imitation Learning with Indirect Intervention [79]. This proactive engagement is further refined in frameworks like MACAROON, which train models to shift from passive answering to proactive clarification requests when faced with ambiguous or unanswerable questions, thereby reducing hallucination rates [80]. The Matcha agent extends this by instructing epistemic actions across multiple modalities—vision, sound, and haptics—reasoning over interactively acquired sensations to plan task execution in partially observable states [81].

Despite these advances, significant challenges remain in balancing the computational cost of look-ahead planning with real-time control constraints. While methods like NavCoT introduce navigational chains of thought to imagine future observations and select actions accordingly [53], the integration of such deliberative processes into low-level control loops remains non-trivial. Future research must focus on developing unified world models that seamlessly blend active inference with reactive control, allowing agents to dynamically switch between exploration and exploitation based on precise uncertainty quantification. Additionally, extending these capabilities to dynamic, non-stationary environments where the “ground truth” shifts during interaction, as highlighted in dynamic disaster scenarios [82], will be critical for robust real-world deployment. These limitations in external uncertainty management naturally motivate the transition to internal cognitive control mechanisms, where agents must not only perceive the world actively but also critique and refine their own reasoning processes, as explored in the following section on self-reflection.

3.5 Self-Reflection and Critical Evaluation Mechanisms

The deployment of embodied agents in unstructured physical environments necessitates mechanisms that transcend single-pass planning, requiring systems capable of introspective critique and iterative refinement. Self-reflection in Vision-Language Models (VLMs) functions as a cognitive control loop, enabling agents to detect hallucinations, validate logical consistency, and recover from execution failures without external intervention. A primary paradigm involves the utilization of an “inner monologue,” where agents generate natural language descriptions of their execution history to identify discrepancies between intended goals and observed states. For instance, the [9] framework demonstrates that closing the loop with environmental feedback—such as success detection and scene descriptions—significantly enhances high-level instruction completion by allowing the model to replan dynamically upon detecting failure. Similarly, [22] advances this by employing VLMs as continuous monitors that detect constraint violations during low-level execution, triggering immediate recovery protocols when plan-execution misalignments occur.

Beyond internal monologues, adversarial and multi-agent frameworks have emerged to mitigate the inherent biases and hallucination tendencies of single-model architectures. By instantiating multiple model instances with distinct roles, such as proposers and critics, systems can leverage debate dynamics to filter out infeasible actions. The [83] study highlights how decentralized agents can achieve emergent verification through communication, while [84] introduces reinforced advantage feedback to streamline this self-refinement process, reducing the computational overhead of excessive querying. Furthermore, specialized verifier modules act as external critics to assess the physical feasibility of generated reasoning chains. Research in [85] establishes that VLMs can serve as scalable behavior critics, identifying undesirable policies that satisfy goal conditions but violate implicit safety or style constraints, a capability crucial for safe real-world deployment.

A critical distinction exists between semantic self-correction and physically grounded validation. While language-based reflection addresses logical coherence, it often lacks metric precision. To bridge this gap, [86] integrates VLMs directly into the control loop to verify task progress against visual evidence, enabling adaptive planning in open-ended scenarios. Additionally, [87] proposes an in-context self-critic mechanism that utilizes vision-specific metrics to select optimal responses, thereby improving modality alignment without relying on external data. However, these approaches face trade-offs between computational latency and reasoning depth; iterative refinement cycles inherently increase inference time, which can be prohibitive for real-time control. Future research must therefore focus on developing lightweight, hierarchical reflection mechanisms that balance the rigor of neuro-symbolic verification with the efficiency required for dynamic physical interaction, ensuring that self-correction remains responsive to the continuous stream of sensorimotor data.

4 Learning Paradigms and Data Ecosystems

4.1 Scalable Multimodal Instruction Tuning and Data Curation

The transition from generic web-scale pre-training to embodied reasoning necessitates a paradigm shift in instruction tuning, moving beyond static image-text alignment toward dynamic sensorimotor integration. Scalable multimodal instruction tuning for embodied agents requires unifying diverse robotic tasks into a coherent training objective that bridges the modality gap between high-dimensional visual observations and continuous action spaces. A foundational strategy involves the tokenization of actions, where continuous robot joint positions and end-effector poses are discretized into text-like tokens, enabling end-to-end policy learning through next-token prediction [10]. This approach allows models like [11] to leverage internet-scale vision-language data alongside

diverse robot demonstrations, achieving robust generalization across multiple embodiments with significantly fewer parameters than closed-source counterparts.

However, the efficacy of such models is contingent upon the quality and diversity of the training corpus. The scarcity of high-quality, diverse instruction-response pairs poses a significant bottleneck for generalization [88]. To address this, recent works advocate for scaling up visual instruction tuning by constructing massive datasets that include complex reasoning and referring expressions, often generated automatically via large language models to ensure diversity and balance [89]. Furthermore, specialized datasets like [88] and [90] emphasize the importance of multi-modal in-context information and conversation-style formulations to empower agents with perception, reasoning, and planning capabilities. The curation of such data is critical; naive scaling can lead to redundancy, necessitating algorithmic approaches for data value estimation and redundancy elimination to focus training on high-value examples that maximize generalization.

Beyond mere volume, the semantic richness of the data must be enhanced to support fine-grained grounding. Models such as [91] and [92] introduce mechanisms for region-level understanding, representing refer expressions as links to bounding boxes or decomposing images into region tokens. This fine-grained alignment is crucial for manipulation tasks where precise spatial reasoning is required. Additionally, incorporating physical commonsense and affordance knowledge into the tuning process is essential for realistic interaction. Frameworks like [93] and [63] demonstrate that fine-tuning on object-centric datasets containing physical concept annotations significantly improves an agent’s ability to reason about material properties and tool usage, thereby reducing hallucinations in physical planning.

Emerging trends also highlight the potential of synthetic data and automated pipelines to bridge the sim-to-real gap and scale data collection. Systems like [94] leverage foundation models to propose diverse instructions and guide data collection in unseen scenarios, effectively scaling up real-world robot episodes with minimal human supervision. Similarly, synthetic datasets such as [95] are utilized to teach visual language concepts beyond nouns, enhancing compositional reasoning without compromising zero-shot capabilities. Despite these advances, challenges remain in balancing the trade-offs between data diversity, computational efficiency, and the fidelity of physical grounding. Future directions must focus on developing adaptive curriculum learning strategies and self-improvement mechanisms, such as those proposed in [87], to continuously refine instruction tuning datasets and align embodied agents more closely with human intent and physical constraints.

4.2 Vision-Language Models as Reward Functions and Behavior Critics

Building upon the robust, grounded policies established through advanced instruction tuning, the next critical evolution involves transitioning VLMs from direct policy generators to sophisticated critics and reward models that refine agent behavior through learned feedback signals. This shift from manually engineered reward functions to dense, learned signals derived from pre-trained VLMs represents a pivotal advancement in embodied reinforcement learning. Traditional approaches often struggle with the sparsity of environmental feedback and the prohibitive cost of human annotation, limiting the scalability of policy optimization. VLMs address these bottlenecks by serving as universal, zero-shot reward models that compute continuous alignment scores between visual observations and natural language goals. By leveraging the semantic richness embedded in internet-scale pre-training, these models eliminate the need for task-specific reward engineering, enabling agents to interpret complex instructions such as “kneel” or “sit in a lotus position” through simple textual prompts [96]. This paradigm shift allows for the direct translation of high-level human intent into

differentiable reward landscapes, facilitating the training of agents in sparse-reward environments where binary success signals are insufficient for efficient exploration.

Beyond static goal specification, VLMs function as sophisticated behavior critics capable of evaluating temporal consistency and procedural quality, thereby complementing the foundational policies with nuanced behavioral alignment. While early success detectors often relied on binary classification of final states, recent frameworks utilize video-language models to critique entire trajectories, identifying undesirable behaviors even when the primary goal is achieved. For instance, a robot may successfully grasp an object but do so with excessive force or awkward kinematics; VLM-based critics can detect these nuances by analyzing video sequences against natural language constraints regarding style and safety [85]. This capability is further enhanced by approaches that integrate temporal dynamics, allowing the critic to penalize inefficient motion patterns or safety violations that static frame evaluators would miss. The integration of such critics into iterative refinement loops enables agents to self-correct and align their policies with complex, multi-faceted human preferences that are difficult to codify mathematically.

The efficacy of VLMs as reward functions is further amplified by techniques that refine the reward signal through contrastive learning and preference alignment, creating a more granular feedback mechanism for policy optimization. Methods such as Language-Image Value (LIV) learning establish a unified objective that trains multimodal representations to implicitly encode universal value functions from action-free videos, assigning dense rewards to frames based on their semantic proximity to language or image goals [97]. Furthermore, the incorporation of natural language feedback allows for preference alignment where humans provide qualitative critiques rather than scalar rewards, enabling agents to learn nuanced aspects of behavior such as social appropriateness or operational smoothness [98]. Emerging self-improvement frameworks also leverage VLMs to generate and evaluate their own candidate responses, employing fine-grained reward modeling to reduce hallucinations and enhance visual grounding without extensive human supervision [99; 100].

Despite these advances, challenges remain regarding the robustness of VLM-derived rewards against distribution shifts and the computational overhead of querying large models during training, issues that directly impact the reliability of agents when deployed in physical environments. Current VLMs may exhibit biases or fail in scenarios requiring precise spatial reasoning, potentially leading to Reward hacking where agents optimize for semantic plausibility rather than physical correctness [96]. Future research must focus on developing hybrid critic architectures that combine the semantic flexibility of VLMs with the geometric rigor of symbolic verifiers, ensuring that learned reward functions remain faithful to physical constraints while scaling to open-world complexity. Addressing these limitations is essential, as the ultimate efficacy of such reward mechanisms and interactive feedback loops hinges on the agent’s ability to operate reliably amidst the visual discrepancies and dynamic mismatches inherent in the transition from simulation to reality.

4.3 Simulation-to-Reality Transfer and Domain Randomization Strategies

The transition from simulation to reality (Sim2Real) remains a pivotal bottleneck in deploying embodied vision-language agents, primarily due to the “reality gap” encompassing visual discrepancies, dynamic mismatches, and sensor noise. While photorealistic simulators offer scalable ground-truth annotations [101], policies trained exclusively in synthetic domains often fail when confronted with the unstructured variability of physical environments. To mitigate this, advanced domain randomization (DR) has emerged as a foundational strategy, systematically perturbing visual textures, lighting conditions, and physical parameters such as friction and mass during training [101]. By

forcing policies to rely on invariant semantic features rather than superficial visual cues, DR enhances robustness against real-world distribution shifts. Recent innovations extend this paradigm by leveraging natural language as a domain-invariant bridge; [102] demonstrates that training image encoders to predict linguistic descriptions of scenes creates a shared semantic space, allowing policies to generalize from simulation to reality with significantly fewer real-world demonstrations compared to purely vision-based baselines.

Beyond static randomization, generative world models are reshaping Sim2Real strategies by enabling agents to practice in “imagined” trajectories that closely approximate real-world physics. Frameworks like [6] integrate 3D perception with generative diffusion models to predict future states, allowing agents to refine policies through counterfactual simulation before physical execution. Similarly, [103] utilizes multimodal large language models to synthesize reliable instruction instances for fine-tuning, effectively expanding the training distribution to cover edge cases rarely encountered in standard simulators. These approaches shift the focus from merely masking the sim-real divide to actively modeling the dynamics of the target environment, thereby reducing the sample complexity of real-world adaptation.

Curriculum learning further optimizes this transfer by automating the synthesis of progressively complex training scenarios. Systems such as [13] employ chain-of-thought reasoning to generate structured sub-goals, creating a scaffolded learning environment that prepares agents for long-horizon tasks. This is complemented by methods like [51], which hypothesize abstract world models via large language models and verify them through interaction, blending internet-scale knowledge with grounded environmental dynamics. However, despite these advances, challenges persist in handling high-frequency sensor noise and non-linear contact dynamics that are computationally prohibitive to simulate accurately. Future directions must therefore explore hybrid paradigms that combine the scalability of generative world models with real-time adaptive mechanisms, ensuring that embodied agents can not only survive but thrive in the inherent unpredictability of physical deployment.

4.4 Interactive Learning and Human-in-the-Loop Refinement

Building upon the hybrid paradigms and real-time adaptive mechanisms established for Sim2Real transfer, interactive learning and human-in-the-loop (HITL) refinement represent the critical next step: transitioning from static, offline training to dynamic, lifelong adaptation in embodied agents. While large-scale pre-training and generative world models provide a robust foundation, the inherent ambiguity of unstructured physical environments necessitates mechanisms for continuous correction and guidance post-deployment. This subsection examines frameworks where agents leverage natural language dialogue, kinesthetic demonstration, and multimodal feedback to iteratively refine policies, effectively closing the gap between abstract intent and physical execution.

A primary mechanism for this refinement is the integration of natural language correction and demonstration. Rather than relying solely on the sparse reward signals or simulated trajectories discussed previously, agents can ingest explicit human feedback to rectify failures in real-time. The [9] framework exemplifies this by enabling agents to form a closed-loop reasoning process where environmental feedback and human interaction are verbalized, allowing the model to self-correct high-level plans without additional gradient updates. Similarly, [7] grounds abstract language models by constraining them with pretrained skills, yet it relies on human intervention to resolve ambiguities where semantic knowledge diverges from physical affordances. Extending this, [79] formalizes a setting where agents actively query human advisors for subgoals when lost, treating

uncertainty resolution as a learnable skill within the imitation learning objective. These approaches demonstrate that treating language as an interactive control signal significantly enhances robustness compared to passive instruction following.

Beyond corrective dialogue, learning from demonstration remains a potent tool for policy adaptation, particularly when combined with active exploration. The [3] study illustrates how agents can emerge with grounded communication protocols through cooperative games, suggesting that interactive contexts accelerate the alignment of symbolic and perceptual representations. In more complex manipulation tasks, [104] shows that leveraging self-supervised visuo-lingual affordance models allows agents to learn from unstructured, reset-free data, requiring language annotations for only a fraction of trajectories. This data efficiency is further amplified by systems like [81], where agents perform epistemic actions to gather multimodal information (vision, sound, haptics) before executing tasks, effectively asking the environment questions to reduce planning uncertainty. Such active perception strategies transform the agent from a passive learner into a proactive investigator, optimizing the information gain from human interventions.

However, integrating human feedback introduces challenges regarding scalability and the potential for conflicting signals. The [105] approach addresses misalignment and hallucination by adapting Reinforcement Learning from Human Feedback (RLHF) to the vision-language domain, using factual augmentation to prevent reward hacking. Nevertheless, purely human-driven refinement can be bottlenecked by operator availability. Consequently, emerging trends point toward hybrid systems where VLMs act as autonomous critics, mirroring the self-correction capabilities of generative world models. For instance, [85] proposes using VLMs as behavior critics to detect undesirable but successful behaviors, creating a scalable proxy for human preference that can guide iterative policy refinement. Furthermore, [106] demonstrates that fine-tuning VLMs with RL using chain-of-thought reasoning significantly improves decision-making, suggesting that internal reasoning traces can serve as a dense feedback signal complementary to external human input.

Future directions must address the latency of interactive loops and the synthesis of diverse feedback modalities into unified optimization objectives. As agents increasingly operate in safety-critical domains, the development of calibrated confidence estimation mechanisms, as explored in [78], will be essential to determine precisely when human intervention is required versus when autonomous exploration suffices. Ultimately, the convergence of interactive dialogue, demonstrative learning, and automated self-critique will define the next generation of embodied agents capable of safe, adaptable, and lifelong learning in dynamic human environments.

5 Evaluation Frameworks, Benchmarks, and Simulation Environments

5.1 Comprehensive Benchmarks for Language-Guided Embodied Tasks

The evaluation of language-guided embodied agents has undergone a paradigmatic shift from static, single-image perception tasks to dynamic, long-horizon scenarios that demand tight coupling between sequential reasoning, spatial grounding, and physical action. Early benchmarks primarily focused on isolated visual question answering or simple navigation, yet the frontier of embodied reasoning now necessitates environments where agents must interpret complex instructions within continuous sensorimotor streams. foundational work in Vision-and-Language Navigation (VLN), such as the Room-to-Room (R2R) dataset established within the Matterport3D simulator, first formalized the challenge of following natural language instructions through real-world 3D environments [107]. Subsequent advancements introduced Speaker-Follower models to address data scarcity and pragmatic reasoning, significantly improving success rates by synthesizing instructions

and evaluating action sequences against linguistic intent [108]. However, these navigation-centric benchmarks often abstract away the complexities of object manipulation and state change.

To bridge the gap between navigation and interaction, benchmarks like ALFRED introduced grounded instructions for everyday household tasks, requiring agents to execute long, compositional sequences involving non-reversible state changes [109]. Similarly, CALVIN extended this scope to long-horizon policy learning, emphasizing the composition of diverse skills from unconstrained language specifications in simulated environments [110]. These datasets highlight a critical limitation in prior evaluations: the inability of models to maintain temporal consistency over extended horizons. Recent efforts have further diversified the evaluation landscape to include urban spatial reasoning, as seen in the Touchdown dataset, which challenges agents to navigate street environments and identify specific locations based on rich spatial descriptions [111]. Moreover, the emergence of situated question answering in 3D scenes, exemplified by SQA3D, forces agents to reason about their own position and orientation relative to the environment, testing a spectrum of capabilities from spatial relation comprehension to multi-hop reasoning [112].

A distinct trend in contemporary benchmarking is the move towards multi-image and dynamic context understanding. Traditional static evaluations fail to capture the temporal dynamics essential for real-world agency. Benchmarks like Mementos explicitly assess reasoning over image sequences, revealing that current multimodal large language models (MLLMs) struggle with behavioral hallucinations and temporal consistency when extrapolating from static frames to dynamic events [113]. Furthermore, the STAR benchmark evaluates situated reasoning in real-world videos by representing situations as hyper-graphs, demanding logical grounding of actions and interactions over time [114]. In the domain of manipulation, newer suites evaluate not just task completion but the quality of spatial grounding and physical commonsense. For instance, SpatialRGPT-Bench provides ground-truth 3D annotations to rigorously test 3D spatial cognition, while benchmarks like AGQA focus on compositional spatio-temporal reasoning to uncover biases in video understanding [115; 28].

Despite these advances, significant challenges remain. Many benchmarks still rely on binary success metrics that obscure failures in intermediate reasoning steps or physical feasibility. The integration of physical properties, such as fragility and material constraints, is often underrepresented, limiting the assessment of an agent’s physical commonsense [63]. Future evaluation frameworks must evolve to include fine-grained metrics for reasoning transparency, hallucination detection, and robustness to distribution shifts, ensuring that progress in benchmark performance translates to reliable deployment in unstructured physical environments. The trajectory suggests a convergence towards holistic evaluation suites that simultaneously stress-test perception, cognition, and control within unified, photorealistic, and dynamically changing worlds.

5.2 Metrics for Reasoning Transparency, Consistency, and Hallucination Detection

Building on the call for holistic evaluation suites that stress-test perception, cognition, and control, this section dissects the intermediate reasoning processes that govern action generation, moving beyond the binary success metrics that often obscure flawed logic. A critical focus is the quantification of “object hallucination,” where models generate descriptions of entities absent from the visual scene—a phenomenon exacerbated by text inertia where language priors override visual evidence [116; 117; 118]. To mitigate evaluation biases inherent in generation styles, polling-based methods like POPE provide stable, flexible quantification of hallucination rates, while logical closed-loop frameworks such as LogicCheckGPT probe consistency by verifying if attribute-object relationships

form valid deductive cycles [116; 119].

Beyond detecting hallucinations, assessing the transparency of Chain-of-Thought (CoT) reasoning requires metrics that validate the alignment between intermediate textual rationales and visual features. Standard CoT prompting often fails in embodied settings due to a disconnect between semantic reasoning and sensorimotor grounding. Consequently, novel approaches like Embodied Chain-of-Thought (ECoT) and Navigational Chain-of-Thought (NavCoT) advocate for training models to generate visually grounded intermediate steps, such as object bounding boxes or imagined future observations, prior to action execution [12; 53]. The efficacy of these methods is evaluated by verifying whether the generated rationale actually influences the final policy, addressing the limitations of static models that struggle with multi-image relational reasoning [120]. The Prism framework further advances this diagnostic capability by disentangling perception from reasoning, allowing for isolated assessment of whether errors stem from visual misinterpretation or logical fallacies, thereby offering a more granular tool than end-to-end accuracy [121].

Consistency metrics also extend to counterfactual and temporal reasoning, which are critical for agents operating in the dynamic environments highlighted in prior benchmarks. Suites like CFMM and Mementos evaluate an agent’s ability to reason about “what-if” scenarios and sequential image changes, revealing that current models often ignore counterfactual presuppositions in favor of static visual priors [56; 113]. To scale these assessments, the use of Large Language Models as automated judges (LLM-as-a-Judge) and behavior critics has emerged as an alternative to human annotation, capable of detecting undesirable behaviors that satisfy goal conditions but violate safety or efficiency constraints [85; 122]. However, reliance on automated critics introduces risks of alignment bias, necessitating factually augmented reward models to prevent hallucinated feedback [105]. Future evaluation paradigms must therefore integrate fine-grained reward modeling, such as ViGoR, to continuously align generated text with visual ground truth, ensuring that embodied agents possess not only high success rates but also robust, interpretable, and faithful reasoning mechanisms [99]. These rigorous diagnostic tools lay the groundwork for the scalable simulation platforms and synthetic data strategies discussed next, which are essential for generating the diverse interaction data required to train and validate such transparent reasoning capabilities.

5.3 Simulation Platforms and Synthetic Data Generation Strategies

Scalable simulation platforms and synthetic data generation strategies constitute the bedrock of modern embodied AI research, addressing the critical scarcity of real-world interaction data required to train robust Vision-Language-Action (VLA) models. The field has transitioned from fragmented, task-specific environments to unified, photorealistic frameworks capable of supporting complex, long-horizon reasoning. Prominent simulators such as Habitat, AI2-THOR, and AllenAct have established standardized interfaces for embodied tasks, enabling reproducible benchmarking across navigation and manipulation domains [123; 124]. These platforms are increasingly distinguished by their ability to render high-fidelity visual observations while providing precise, noise-free ground-truth annotations for state variables, object semantics, and physical dynamics—attributes essential for supervising multi-modal policy learning and diagnosing failure modes in reasoning chains.

A pivotal advancement in this domain is the shift towards automated, generative pipelines for synthetic scenario construction, designed to stress-test agents against rare edge cases and prevent overfitting to static training distributions. Recent methodologies leverage large language models and diffusion models to procedurally generate diverse instruction-scenario pairs, effectively scaling the complexity of evaluation benchmarks beyond human curation limits. For instance, frameworks

like AutoRT utilize foundation models to propose novel tasks and orchestrate data collection across fleets of robots in simulated and real environments, significantly enhancing the diversity of training corpora [94]. Similarly, specialized benchmarks such as CALVIN employ synthetic generation to create long-horizon, language-conditioned manipulation sequences that challenge an agent’s temporal consistency and compositional generalization [110]. Furthermore, approaches like CounterCurate demonstrate the efficacy of using generative models to create counterfactual examples, specifically targeting physical and semantic compositional reasoning gaps that standard datasets often miss [55].

Despite these advances, the “reality gap”—the discrepancy between simulated dynamics and physical deployment—remains a formidable barrier. Effective sim-to-real transfer necessitates sophisticated domain randomization strategies that perturb visual textures, lighting conditions, and physical parameters (e.g., friction, mass) during training to enforce invariant feature learning. Empirical studies in Vision-and-Language Navigation (VLN) indicate that while agents can achieve high success rates in simulation, performance degrades significantly in unseen physical environments without rigorous randomization and subgoal modeling [101]. Emerging solutions integrate neural implicit representations and generative world models to bridge this divide, allowing agents to practice in “imagined” trajectories that closely mimic real-world physics before execution [103]. Moreover, the integration of 3D spatial reasoning data, as seen in datasets like SceneVerse and SpatialVLM, is crucial for endowing models with metric awareness and geometric intuition that 2D-centric training fails to capture [125; 126].

Looking forward, the trajectory of simulation research points toward fully autonomous, self-improving ecosystems where agents not only learn within static simulators but actively contribute to the generation of their own training curricula. The convergence of neuro-symbolic planning with photorealistic simulation offers a promising pathway to verify the logical consistency of agent behaviors against formal constraints before physical deployment. Ultimately, the fidelity and diversity of synthetic environments will dictate the ceiling of embodied reasoning capabilities, necessitating continued innovation in physics engine accuracy, semantic scene generation, and the seamless integration of sim-to-real transfer protocols to realize trustworthy autonomous agents.

5.4 Real-World Deployment Protocols and Robustness Assessment

Building upon the high-fidelity synthetic environments and sim-to-real transfer strategies discussed previously, the transition to physical deployment necessitates a fundamental paradigm shift in evaluation: moving beyond static success rates in controlled settings to dynamic assessments of robustness against environmental perturbations and distribution shifts. While simulation platforms provide scalable ground truth, they often fail to capture the full stochasticity of real-world sensor noise, lighting variations, and unstructured object arrangements. To bridge this gap, recent protocols emphasize stress-testing agents under diverse physical conditions. For instance, frameworks like VL-CheckList systematically analyze model capabilities across objects, attributes, and relations to quantify resilience to visual degradations [127]. Similarly, while domain randomization during training is critical for sim-to-real transfer, evaluation must explicitly verify performance in completely unseen environments where prior mapping is unavailable—a setting where success rates can drop precipitously compared to mapped scenarios [101].

A central challenge in real-world assessment lies in evaluating generalization to novel objects and open-vocabulary instructions without task-specific fine-tuning. Benchmarks such as CALVIN probe an agent’s ability to decompose abstract commands and adapt to new semantic concepts in a zero-

shot manner [110]. However, static validation sets risk overfitting, as models may memorize dataset biases rather than learning genuine reasoning capabilities. To mitigate this, dynamic evaluation strategies like NPHardEval4V propose monthly-updated benchmarks that prevent saturation and ensure an authentic assessment of reasoning faculties over time [128]. Furthermore, the reliance on binary success metrics often obscures the quality of intermediate reasoning; thus, emerging protocols incorporate video-language models as behavior critics to detect undesirable trajectories—even those that technically achieve the goal—thereby aligning evaluation with human preferences and safety constraints [85].

Robustness also extends to multimodal input integrity, particularly when dealing with erroneous Automatic Speech Recognition (ASR) transcripts or ambiguous visual prompts. Effective deployment requires agents to recover from such noise using contextual visual grounding, a capability tested by frameworks that introduce realistic input corruptions [129]. Moreover, the limitations of purely language-based reasoning in physical spaces have spurred the development of embodied chain-of-thought protocols, where agents must ground their reasoning steps in precise spatial coordinates and sensor observations before executing actions [12]. Approaches like RePLan further enhance robustness by evaluating an agent’s capacity for online replanning when initial strategies fail due to unforeseen physical obstacles, effectively bridging the gap between high-level planning and low-level control fidelity [44].

Looking forward, the field must converge on holistic evaluation suites that integrate perception, cognition, and action into a continuous feedback loop. The emergence of embodied foundation models trained on diverse, in-the-wild robot data, such as AutoRT, suggests that future robustness assessments should leverage large-scale, heterogeneous real-world episodes rather than curated synthetic scenarios [94]. Ultimately, rigorous real-world deployment protocols must prioritize adaptability, safety, and interpretability, ensuring that vision-language agents can operate reliably in the face of the inherent unpredictability of physical environments.

6 Critical Challenges in Reliability, Safety, and Real-World Integration

6.1 Safety Alignment and Ethical Decision-Making in Physical Spaces

The deployment of Vision-Language Models (VLMs) in embodied systems necessitates a paradigm shift from statistical alignment to rigorous physical safety and ethical grounding. Unlike static digital agents, embodied VLMs operate in continuous state spaces where hallucinations or misaligned objectives can result in catastrophic physical consequences, including property damage and personal injury. The core challenge lies in bridging the gap between abstract linguistic safety constraints and the geometric realities of dynamic environments. Recent frameworks address this by integrating physical risk awareness directly into the planning loop, moving beyond reactive filtering to proactive hazard mitigation. For instance, [22] introduces a mechanism for detecting and recovering from plan-execution misalignments by leveraging VLMs to monitor constraint violations continuously, thereby preventing unsafe actions before they culminate in failure. Similarly, [44] utilizes physical grounding to adapt robot actions online when initial plans encounter unforeseen obstacles, ensuring that safety is maintained through dynamic replanning rather than rigid pre-computation.

A critical dimension of safety alignment involves the integration of ethical reasoning and social norms into decision-making processes. Embodied agents must not only avoid physical collisions but also adhere to human preferences regarding object placement and interaction propriety. The [130] benchmark illustrates this by evaluating agents on their ability to tidy environments based on implicit human norms rather than explicit instructions, highlighting the necessity for models

to internalize commonsense reasoning about social appropriateness. Furthermore, [131] and [132] emphasize that safe interaction requires a deep understanding of object affordances and potential consequences; an agent must reason not just about *how* to grasp an object, but *whether* doing so violates physical stability or social conventions. These approaches suggest that safety is not merely a constraint satisfaction problem but a semantic understanding of the environment’s physical and social topology.

However, purely neural approaches often lack the verifiability required for high-stakes applications. Consequently, neuro-symbolic hybrids have emerged as a promising avenue for ensuring interpretability and auditability. [17] and [47] demonstrate how translating natural language into executable code or 3D value maps allows for the explicit encoding of safety constraints and geometric reasoning, making the decision process transparent and debuggable. This contrasts with end-to-end policies where safety logic is latent and opaque. The [9] framework further enhances this by incorporating environmental feedback into a closed-loop reasoning process, allowing the agent to verbally articulate its safety checks and self-correct based on observed outcomes. Despite these advances, significant challenges remain in quantifying uncertainty and handling distribution shifts. As noted in [78], miscalibrated confidence can lead agents to act prematurely in uncertain scenarios, posing severe risks. Future research must therefore focus on developing confidence-calibrated planning mechanisms that inherently modulate agent behavior based on epistemic uncertainty, ensuring that when an agent cannot verify the safety of an action against its world model, it defaults to conservative strategies or seeks human intervention. Ultimately, achieving trustworthy embodied intelligence requires a unified architecture that seamlessly fuses semantic ethical reasoning with rigorous geometric constraint satisfaction.

6.2 Adversarial Robustness and Security Vulnerabilities

While the previous section established the necessity of rigorous safety alignment and uncertainty quantification to prevent catastrophic physical failures, the practical realization of these safeguards is confronted by a critical security paradox: the very multimodal flexibility that enables zero-shot generalization in Vision-Language Models (VLMs) also significantly expands the attack surface for adversarial exploitation. Unlike static perception systems, embodied agents operate in continuous feedback loops where adversarial perturbations in visual or linguistic modalities can cascade into physically dangerous actions, directly undermining the trustworthiness sought through ethical grounding. Recent benchmarks reveal that while models excel in standard tasks, they remain brittle against coordinated cross-modality attacks. For instance, [133] demonstrates that state-of-the-art multimodal models struggle with core perception tasks like depth estimation and forensics detection, suggesting an inherent vulnerability to visual deceptive inputs that could mislead navigation or manipulation policies. Similarly, [134] highlights significant performance drops when models are forced to integrate vision and text simultaneously, exposing gaps in robust reasoning under adversarial pressure that complicate the reliable execution of safety constraints.

A primary threat vector involves cross-modality jailbreaks, where adversaries manipulate the alignment between visual and textual streams to bypass the safety filters discussed previously. Attacks can leverage “text inertia,” a phenomenon where models rely excessively on linguistic priors while ignoring visual evidence, allowing attackers to inject malicious instructions that override visual safety cues [118]. This is particularly perilous in embodied settings where an agent might ignore a physical obstacle due to a conflicting textual command, effectively nullifying geometric constraint satisfaction. Furthermore, the integration of VLMs into autonomous driving and robotics has shown susceptibility to hallucinations and misgrounding, which can be intentionally exacerbated by ad-

versarial prompts [135; 136]. The [55] framework underscores the neglect of physically grounded reasoning in current models, noting near-chance performance in counting and position understanding when faced with counterfactual or perturbed inputs, thereby enabling attackers to trick agents into invalid spatial deductions that violate the very physical laws required for safe operation.

Real-world perturbations present another formidable challenge, specifically through typographic attacks and environmental modifications that exploit the perceptual ambiguities inherent in open-world scenarios. Adversaries can alter physical signs or object textures to induce catastrophic planning failures, capitalizing on the models’ inability to distinguish between benign variations and malicious setups. [137] and [138] illustrate that even generative models harbor deep biases and spatial inconsistencies, which can be exploited to generate confusing visual contexts for embodied agents. Moreover, [139] reveals a substantial gap between model and human performance in context-sensitive reasoning, implying that agents may fail to distinguish between benign environmental variations and malicious adversarial setups in text-rich visual scenes. The [56] benchmark further confirms that models often disregard counterfactual presuppositions in favor of visual data, or vice versa, creating a dichotomy that attackers can manipulate to force erroneous decision-making.

To mitigate these vulnerabilities, defensive strategies are evolving from passive robustness training to active verification mechanisms, though these solutions introduce new challenges relevant to the computational constraints discussed in the subsequent section. Approaches like [140] propose training objectives that enforce alignment between visual explanations and influential image regions, mitigating reliance on superficial statistical correlations. Additionally, [85] advocates for using VLMs as behavior critics to iteratively verify agent actions against video evidence, catching undesirable behaviors that simple reward models miss. However, these defenses often incur significant computational overhead, creating a direct trade-off between latency and security in real-time control loops [141]. This tension highlights that achieving robustness against adversarial duress cannot come at the expense of the real-time responsiveness required for physical control. Consequently, future research must prioritize the development of unified robustness frameworks that jointly optimize for perceptual fidelity and logical consistency while remaining computationally efficient, potentially leveraging neuro-symbolic verification to ensure that action sequences remain safe even under attack. Without such advancements, the integration of VLMs into safety-critical physical systems remains precarious, as the computational cost of ensuring security may ultimately preclude the deployment of these models on resource-constrained robotic hardware.

6.3 Computational Efficiency and Edge Deployment Constraints

The deployment of Vision-Language-Action (VLA) models on embodied agents is fundamentally constrained by the computational disparity between massive foundation models and the resource-limited hardware typical of robotic platforms. While models like [5] and [10] demonstrate emergent reasoning capabilities through internet-scale pre-training, their parameter counts and memory footprints often preclude real-time inference on edge devices, necessitating a reliance on cloud infrastructure that introduces unacceptable latency and connectivity vulnerabilities for safety-critical control loops. Addressing this bottleneck requires a multifaceted approach involving model compression, architectural innovation, and strategic trade-offs between generalization and efficiency.

Model compression techniques, particularly quantization and low-rank adaptation, have emerged as primary strategies for reducing the computational burden of VLAs without significant performance degradation. Recent work with [11] illustrates that 7-billion parameter models can be fine-tuned on consumer-grade GPUs and served efficiently via quantization, maintaining downstream success rates

comparable to much larger closed-source counterparts. Similarly, [14] employs fine-tuned adapters to preserve the inherent commonsense reasoning of Multimodal Large Language Models (MLLMs) while equipping them for manipulation, significantly reducing the number of trainable parameters. However, aggressive quantization risks impairing the model’s nuanced understanding of physical affordances and spatial relationships, which are critical for precise motor control. Consequently, adaptive token reduction mechanisms offer a complementary pathway; [142] proposes an adaptive visual token reduction approach that merges spatially redundant tokens, compressing visual input by up to 18 times while retaining reasoning fidelity. This is particularly vital for high-resolution inputs where quadratic attention complexity becomes prohibitive.

Beyond compression, architectural re-design is essential for achieving the linear inference complexity required for high-frequency control. The integration of State Space Models (SSMs) represents a paradigm shift from standard Transformer architectures. [16] leverages the Mamba architecture to achieve inference speeds seven times faster than existing robot MLLMs, enabling efficient reasoning and manipulation with minimal fine-tuning. This efficiency allows for the deployment of complex policies on embedded systems where traditional attention-based models would fail to meet temporal constraints. Furthermore, [143] demonstrates that lightweight adapter modules can bridge the gap between image encoders and LLMs, enabling joint optimization with drastically reduced computational expenditure compared to full model retraining.

A critical challenge remains in balancing the tension between the broad generalization capabilities of large foundation models and the strict latency requirements of dynamic physical interaction. While [126] and [144] enhance spatial perception through dynamic resolution and extensive training, these improvements often come at the cost of increased inference time. The trade-off is non-linear; slight increases in latency can destabilize closed-loop control, leading to oscillatory behaviors or collision risks. Future research must therefore focus on “anytime” algorithms that can dynamically adjust model depth or token density based on real-time computational availability and task urgency. Additionally, the development of hybrid neuro-symbolic architectures, as hinted at in [17], may offer a route where heavy reasoning is offloaded to efficient code generation, leaving only lightweight execution to the edge device. Ultimately, realizing trustworthy embodied intelligence depends on co-designing algorithms and hardware to ensure that the cognitive richness of VLAs can operate within the stringent physical constraints of the real world.

6.4 Handling Uncertainty and Dynamic Environmental Adaptation

While computational efficiency addresses the hardware constraints of embodied agents, the inherent unpredictability of unstructured real-world environments demands equally robust mechanisms to manage perceptual ambiguity, dynamic distribution shifts, and execution failures. Unlike static perception tasks, embodied reasoning requires continuous alignment between internal world models and evolving physical states, a challenge that persists even when latency requirements are met. A primary strategy to address this involves explicit uncertainty estimation to calibrate decision-making confidence. Recent approaches leverage conformal prediction to quantify the reliability of vision-language model (VLM) outputs, enabling agents to determine optimal stopping criteria for exploration tasks rather than relying on miscalibrated probability scores [78]. Furthermore, distinguishing between epistemic uncertainty (model ignorance) and aleatoric uncertainty (environmental noise) is critical; methods like ExACT introduce conceptual reasoning modules to estimate semantic uncertainty in dynamic action recognition, thereby reducing errors in ambiguous temporal contexts [145].

When uncertainty exceeds predefined thresholds or execution deviates from planned trajectories, active adaptation and failure recovery become paramount. Closed-loop frameworks increasingly utilize VLMs as “inner monologue” systems that process environmental feedback to detect discrepancies and trigger replanning without human intervention [9]. For instance, the RePlan framework employs VLMs to perceive state changes caused by unforeseen obstacles, allowing for online plan revision in long-horizon tasks where open-loop policies typically fail [44]. Similarly, COME-robot demonstrates that integrating GPT-4V for situated reasoning and progress verification significantly enhances recovery rates in open-vocabulary manipulation by iteratively refining action policies based on visual feedback [86]. These systems exemplify a shift from reactive error handling to proactive self-correction, where the agent treats failure as an informational signal to update its semantic map or task decomposition.

Handling dynamic environmental changes also requires robust generalization to novel objects and layouts that diverge from training distributions. Approaches such as VoxPoser compose 3D value maps using LLM-derived affordances, enabling zero-shot trajectory synthesis that remains robust to dynamic perturbations and unseen object configurations [47]. Moreover, visual prompting techniques like MOKA and PIVOT convert spatial reasoning into iterative question-answering problems, allowing pre-trained VLMs to localize keypoints and infer affordances for unseen objects without task-specific fine-tuning [146; 147]. While these methods improve adaptability, they often incur higher computational latency compared to end-to-end policies, reintroducing the efficiency challenges discussed previously. Consequently, hybrid architectures that balance the semantic flexibility of large models with the efficiency of reactive controllers represent a critical frontier. Future research must focus on developing unified world models that seamlessly integrate uncertainty quantification with rapid, low-level control adjustments, ensuring safety and reliability even under severe distribution shifts and partial observability.

7 Conclusion

The trajectory of Vision-Language Models (VLMs) for embodied reasoning has shifted decisively from modular, perception-decoupled pipelines toward unified, end-to-end architectures that treat action as a first-class token. This synthesis reveals a fundamental paradigm change: where early approaches relied on explicit symbolic grounding and separate planning modules [1; 148], contemporary foundation models such as RT-2 and PaLM-E demonstrate that co-fine-tuning on robotic trajectory data alongside internet-scale vision-language corpora enables emergent semantic generalization and zero-shot transfer to novel objects [10; 5]. By tokenizing continuous sensorimotor streams, these Vision-Language-Action (VLA) models effectively bridge the gap between abstract linguistic commands and low-level motor primitives, allowing agents to perform rudimentary logical deduction without task-specific retraining [11].

However, despite these architectural strides, a critical disconnect remains between statistical correlation and genuine physical understanding. While models like [7] successfully constrain language generation via pretrained skills to ensure feasibility, purely neural approaches often lack the rigorous guarantees required for safe, long-horizon deployment. The reliance on next-token prediction can lead to hallucinations in dynamic environments where physical constraints are non-negotiable. Consequently, neuro-symbolic hybrids and “inner monologue” mechanisms have emerged as vital corrections, leveraging environmental feedback to close the loop between planning and execution [9; 149]. These systems illustrate that while end-to-end policies excel in flexibility, interpretable planning frameworks remain essential for verifying safety constraints and managing complex temporal dependencies [150].

Looking forward, the field must transcend static policy learning to embrace unified world models that facilitate counterfactual simulation and lookahead planning. Current limitations in handling distribution shifts and unseen physical dynamics necessitate agents capable of constructing internal representations to predict future states before acting [6; 151]. The integration of 3D geometric encoding and neural implicit fields offers a promising pathway to enhance spatial awareness, moving beyond the 2D-centric biases of standard VLMs [32]. Furthermore, the transition from episodic training to lifelong learning is imperative; agents must continuously refine their policies through interactive human feedback and autonomous exploration to adapt to evolving environments [104; 94].

Ultimately, achieving trustworthy embodied intelligence requires a holistic alignment of perception, cognition, and action with rigorous safety guarantees. Future research must prioritize the development of robust verification mechanisms, such as video-language critics that detect undesirable behaviors beyond simple task success [85]. Only by unifying scalable foundation models with physically grounded world simulations and ethical alignment frameworks can we realize agents that are not merely reactive pattern matchers, but capable, safe, and generalizable partners in the physical world.

References

- [1] A Joint Model of Language and Perception for Grounded Attribute Learning
- [2] Learning the Semantics of Manipulation Action
- [3] Learning Cooperative Visual Dialog Agents with Deep Reinforcement Learning
- [4] Grounded Language Learning in a Simulated 3D World
- [5] PaLM-E An Embodied Multimodal Language Model
- [6] 3D-VLA A 3D Vision-Language-Action Generative World Model
- [7] Do As I Can, Not As I Say Grounding Language in Robotic Affordances
- [8] Learning Interpretable Spatial Operations in a Rich 3D Blocks World
- [9] Inner Monologue Embodied Reasoning through Planning with Language Models
- [10] RT-2 Vision-Language-Action Models Transfer Web Knowledge to Robotic Control
- [11] OpenVLA: An Open-Source Vision-Language-Action Model
- [12] Robotic Control via Embodied Chain-of-Thought Reasoning
- [13] EmbodiedGPT Vision-Language Pre-Training via Embodied Chain of Thought
- [14] ManipLLM Embodied Multimodal Large Language Model for Object-Centric Robotic Manipulation
- [15] LLARVA: Vision-Action Instruction Tuning Enhances Robot Learning
- [16] RoboMamba: Multimodal State Space Model for Efficient Robot Reasoning and Manipulation
- [17] Code as Policies Language Model Programs for Embodied Control
- [18] ViperGPT Visual Inference via Python Execution for Reasoning

- [19] Learning a natural-language to LTL executable semantic parser for grounded robotics
- [20] DeepSym Deep Symbol Generation and Rule Learning from Unsupervised Continuous Robot Interaction for Planning
- [21] JARVIS A Neuro-Symbolic Commonsense Reasoning Framework for Conversational Embodied Agents
- [22] DoReMi Grounding Language Model by Detecting and Recovering from Plan-Execution Misalignment
- [23] Neuro-Symbolic Video Search
- [24] Abstract Spatial-Temporal Reasoning via Probabilistic Abduction and Execution
- [25] Is A Picture Worth A Thousand Words? Delving Into Spatial Reasoning for Vision Language Models
- [26] What’s up with vision-language models Investigating their struggle with spatial reasoning
- [27] SpatialBot: Precise Spatial Understanding with Vision Language Models
- [28] SpatialRGPT: Grounded Spatial Reasoning in Vision Language Model
- [29] 3D-LLM Injecting the 3D World into Large Language Models
- [30] ShapeLLM Universal 3D Object Understanding for Embodied Interaction
- [31] SAT 2D Semantics Assisted Training for 3D Visual Grounding
- [32] Visual Language Maps for Robot Navigation
- [33] LLM-Grounder Open-Vocabulary 3D Visual Grounding with Large Language Model as an Agent
- [34] Embodied Visual Recognition
- [35] Scene-LLM Extending Language Model for 3D Visual Understanding and Reasoning
- [36] Cross-modal Map Learning for Vision and Language Navigation
- [37] SASRA Semantically-aware Spatio-temporal Reasoning Agent for Vision-and-Language Navigation in Continuous Environments
- [38] TopViewRS: Vision-Language Models as Top-View Spatial Reasoners
- [39] MultiPLY A Multisensory Object-Centric Embodied Large Language Model in 3D World
- [40] AlphaBlock Embodied Finetuning for Vision-Language Reasoning in Robot Manipulation
- [41] LLM as A Robotic Brain Unifying Egocentric Memory and Control
- [42] Ask, Attend and Answer Exploring Question-Guided Spatial Attention for Visual Question Answering
- [43] NavGPT Explicit Reasoning in Vision-and-Language Navigation with Large Language Models
- [44] RePLan Robotic Replanning with Perception and Language Models
- [45] Visually-Grounded Planning without Vision Language Models Infer Detailed Plans from High-level Instructions

- [46] Instruct2Act Mapping Multi-modality Instructions to Robotic Actions with Large Language Model
- [47] VoxPoser Composable 3D Value Maps for Robotic Manipulation with Language Models
- [48] Language-Conditioned Robotic Manipulation with Fast and Slow Thinking
- [49] Reasoning with Language Model is Planning with World Model
- [50] Dynamic Planning with a LLM
- [51] Do Embodied Agents Dream of Pixelated Sheep Embodied Decision Making using Language Guided World Modelling
- [52] MIRA Mental Imagery for Robotic Affordances
- [53] NavCoT Boosting LLM-Based Vision-and-Language Navigation via Learning Disentangled Reasoning
- [54] CLEVRER CoLLision Events for Video REpresentation and Reasoning
- [55] CounterCurate Enhancing Physical and Semantic Visio-Linguistic Compositional Reasoning via Counterfactual Examples
- [56] Eyes Can Deceive Benchmarking Counterfactual Reasoning Abilities of Multi-modal Large Language Models
- [57] Visual Sketchpad: Sketching as a Visual Chain of Thought for Multimodal Language Models
- [58] Mind’s Eye Grounded Language Model Reasoning through Simulation
- [59] Do Neural Language Representations Learn Physical Commonsense
- [60] PROST Physical Reasoning of Objects through Space and Time
- [61] Robot Object Retrieval with Contextual Natural Language Queries
- [62] Open-World Object Manipulation using Pre-trained Vision-Language Models
- [63] Physically Grounded Vision-Language Models for Robotic Manipulation
- [64] Octopi: Object Property Reasoning with Large Tactile-Language Models
- [65] Can Language Models Understand Physical Concepts
- [66] From Representation to Reasoning Towards both Evidence and Commonsense Reasoning for Video Question-Answering
- [67] Dynamic Visual Reasoning by Learning Differentiable Physics Models from Video and Language
- [68] Visual Grounding of Learned Physical Models
- [69] PIGLeT Language Grounding Through Neuro-Symbolic Interaction in a 3D World
- [70] Cross-Modal Causal Relational Reasoning for Event-Level Visual Question Answering
- [71] Quantifying and Mitigating Unimodal Biases in Multimodal Large Language Models A Causal Perspective
- [72] ComPhy Compositional Physical Reasoning of Objects and Events from Videos

- [73] NEWTON Are Large Language Models Capable of Physical Reasoning
- [74] From Word Models to World Models Translating from Natural Language to the Probabilistic Language of Thought
- [75] Bayesian Inference in Model-Based Machine Vision
- [76] A Multi-World Approach to Question Answering about Real-World Scenes based on Uncertain Input
- [77] Can an Embodied Agent Find Your Cat-shaped Mug LLM-Guided Exploration for Zero-Shot Object Navigation
- [78] Explore until Confident Efficient Exploration for Embodied Question Answering
- [79] Vision-based Navigation with Language-based Assistance via Imitation Learning with Indirect Intervention
- [80] MACAROON: Training Vision-Language Models To Be Your Engaged Partners
- [81] Chat with the Environment Interactive Multimodal Perception Using Large Language Models
- [82] HAZARD Challenge Embodied Decision Making in Dynamically Changing Environments
- [83] Building Cooperative Embodied Agents Modularly with Large Language Models
- [84] Towards Efficient LLM Grounding for Embodied Multi-Agent Collaboration
- [85] Task Success is not Enough Investigating the Use of Video-Language Models as Behavior Critics for Catching Undesirable Agent Behaviors
- [86] Closed-Loop Open-Vocabulary Mobile Manipulation with GPT-4V
- [87] Enhancing Visual-Language Modality Alignment in Large Vision Language Models via Self-Improvement
- [88] MIMIC-IT Multi-Modal In-Context Instruction Tuning
- [89] SVIT Scaling up Visual Instruction Tuning
- [90] LLaRA: Supercharging Robot Learning Data for Vision-Language Policy
- [91] Kosmos-2 Grounding Multimodal Large Language Models to the World
- [92] Groma Localized Visual Tokenization for Grounding Multimodal Large Language Models
- [93] ManipVQA Injecting Robotic Affordance and Physically Grounded Information into Multi-Modal Large Language Models
- [94] AutoRT Embodied Foundation Models for Large Scale Orchestration of Robotic Agents
- [95] Going Beyond Nouns With Vision & Language Models Using Synthetic Data
- [96] Vision-Language Models are Zero-Shot Reward Models for Reinforcement Learning
- [97] LIV Language-Image Representations and Rewards for Robotic Control
- [98] DRESS Instructing Large Vision-Language Models to Align and Interact with Humans via Natural Language Feedback

- [99] ViGoR Improving Visual Grounding of Large Vision Language Models with Fine-Grained Reward Modeling
- [100] Calibrated Self-Rewarding Vision Language Models
- [101] Sim-to-Real Transfer for Vision-and-Language Navigation
- [102] Natural Language Can Help Bridge the Sim2Real Gap
- [103] WorldGPT: Empowering LLM as Multimodal World Model
- [104] Grounding Language with Visual Affordances over Unstructured Data
- [105] Aligning Large Multimodal Models with Factually Augmented RLHF
- [106] Fine-Tuning Large Vision-Language Models as Decision-Making Agents via Reinforcement Learning
- [107] Vision-and-Language Navigation Interpreting visually-grounded navigation instructions in real environments
- [108] Speaker-Follower Models for Vision-and-Language Navigation
- [109] ALFRED A Benchmark for Interpreting Grounded Instructions for Everyday Tasks
- [110] CALVIN A Benchmark for Language-Conditioned Policy Learning for Long-Horizon Robot Manipulation Tasks
- [111] Touchdown Natural Language Navigation and Spatial Reasoning in Visual Street Environments
- [112] SQA3D Situated Question Answering in 3D Scenes
- [113] Mementos A Comprehensive Benchmark for Multimodal Large Language Model Reasoning over Image Sequences
- [114] STAR: A Benchmark for Situated Reasoning in Real-World Videos
- [115] AGQA A Benchmark for Compositional Spatio-Temporal Reasoning
- [116] Evaluating Object Hallucination in Large Vision-Language Models
- [117] Analyzing and Mitigating Object Hallucination in Large Vision-Language Models
- [118] Paying More Attention to Image: A Training-Free Method for Alleviating Hallucination in LVLMs
- [119] Logical Closed Loop Uncovering Object Hallucinations in Large Vision-Language Models
- [120] Benchmarking Multi-Image Understanding in Vision and Language Models: Perception, Knowledge, Reasoning, and Multi-Hop Reasoning
- [121] Prism: A Framework for Decoupling and Assessing the Capabilities of VLMs
- [122] MM-Vet Evaluating Large Multimodal Models for Integrated Capabilities
- [123] AllenAct A Framework for Embodied AI Research
- [124] A Survey of Embodied AI From Simulators to Research Tasks
- [125] SceneVerse Scaling 3D Vision-Language Learning for Grounded Scene Understanding

- [126] SpatialVLM Endowing Vision-Language Models with Spatial Reasoning Capabilities
- [127] VL-CheckList Evaluating Pre-trained Vision-Language Models with Objects, Attributes and Relations
- [128] NPHardEval4V A Dynamic Reasoning Benchmark of Multimodal Large Language Models
- [129] MultiModal-GPT A Vision and Language Model for Dialogue with Humans
- [130] Housekeep Tidying Virtual Households using Commonsense Reasoning
- [131] AffordanceLLM Grounding Affordance from Vision Language Models
- [132] Learning to Act Properly Predicting and Explaining Affordances from Images
- [133] BLINK Multimodal Large Language Models Can See but Not Perceive
- [134] MMMU-Pro: A More Robust Multi-discipline Multimodal Understanding Benchmark
- [135] Driving with LLMs Fusing Object-Level Vector Modality for Explainable Autonomous Driving
- [136] DriveLM Driving with Graph Visual Question Answering
- [137] DALL-Eval Probing the Reasoning Skills and Social Biases of Text-to-Image Generation Models
- [138] Benchmarking Spatial Relationships in Text-to-Image Generation
- [139] ConTextual Evaluating Context-Sensitive Text-Rich Visual Reasoning in Large Multimodal Models
- [140] Self-Critical Reasoning for Robust Visual Question Answering
- [141] Efficient Multimodal Learning from Data-centric Perspective
- [142] LLaVA-PruMerge Adaptive Token Reduction for Efficient Large Multimodal Models
- [143] Cheap and Quick Efficient Vision-Language Instruction Tuning for Large Language Models
- [144] Qwen2-VL: Enhancing Vision-Language Model’s Perception of the World at Any Resolution
- [145] ExACT Language-guided Conceptual Reasoning and Uncertainty Estimation for Event-based Action Recognition and More
- [146] MOKA Open-Vocabulary Robotic Manipulation through Mark-Based Visual Prompting
- [147] PIVOT Iterative Visual Prompting Elicits Actionable Knowledge for VLMs
- [148] ALFWorld Aligning Text and Embodied Environments for Interactive Learning
- [149] Large Language Models Are Neurosymbolic Reasoners
- [150] LLM-Planner Few-Shot Grounded Planning for Embodied Agents with Large Language Models
- [151] Language Models Meet World Models Embodied Experiences Enhance Language Models