

Symbiotic Scholarship: A Comprehensive Survey of Human-AI Collaboration in Scientific Writing and Research Workflows

1 Introduction

The integration of artificial intelligence into scientific inquiry marks a paradigmatic shift from static computational support to dynamic, symbiotic partnerships. Historically, computational tools functioned as passive instruments for statistical analysis or rule-based editing, yet the advent of large language models (LLMs) has catalyzed a transition toward systems capable of hypothesis formulation and co-creative drafting [1]. This survey delineates the scope of this evolution, distinguishing between traditional automation, which optimizes existing workflows, and emerging collaborative frameworks where human and machine agency are distributed across the research lifecycle. The motivation for this examination lies in the exponential growth of scientific literature and the increasing complexity of interdisciplinary research, which necessitate tools that do more than merely accelerate tasks; they must enhance cognitive capacity and foster novel discovery patterns [2].

Central to our theoretical foundation is the concept of human-machine teaming, where the objective is not the replacement of the researcher but the optimization of combined performance through complementarity. Recent studies demonstrate that systems trained to recognize instances difficult for humans while offloading routine cognitive loads can outperform either agents operating in isolation [3]. However, this synergy is contingent upon the calibration of trust and the alignment of mental models. Empirical evidence suggests that the efficacy of such collaborations varies significantly based on user expertise; while novices may benefit from high levels of scaffolding, mid-level performers often struggle with integrating AI advice, highlighting the non-linear relationship between automation levels and team performance [4; 5]. Furthermore, the risk of “automation complacency” and the potential for AI to induce an “illusion of competence” underscore the necessity for interfaces that promote critical engagement rather than blind reliance [6; 7].

To systematize these diverse interactions, this survey proposes a unified taxonomy categorizing collaborative modes based on initiative flow, task complexity, and the degree of human oversight. This classification spans from human-led assistance, where AI acts as a reactive tool for stylistic refinement or data extraction [8], to fully autonomous discovery loops where agents iteratively design and execute experiments [9]. A critical dimension of this taxonomy involves the spectrum of autonomy in literature synthesis, ranging from semantic retrieval engines that break down disciplinary silos [10] to multi-agent systems that autonomously construct knowledge graphs and identify interdisciplinary gaps [11].

The objectives of this survey are threefold: first, to critically analyze the technical mechanisms enabling these collaborative paradigms across writing, analysis, and discovery; second, to evaluate the ethical and epistemological implications of AI-generated content, particularly regarding authorship attribution and the integrity of the scientific record [12; 13]; and third, to identify gaps in current evaluation metrics that fail to capture the nuances of collaborative synergy. By synthesizing insights from recent advancements in argument mining for peer review [14] and the development of process-based verification frameworks [13], we aim to provide a comprehensive roadmap for future research. Ultimately, this work posits that the future of scientific progress depends not on the autonomy of machines, but on the design of adaptive, transparent, and ethically grounded interfaces that preserve human intuition while leveraging the scalable reasoning capabilities of artificial intelligence.

2 Artificial Intelligence Augmentation of the Scientific Writing Lifecycle

2.1 Generative Support for Ideation and Structural Conceptualization

The prewriting phase of scientific inquiry is undergoing a paradigmatic shift from static, linear outlining toward dynamic, agent-based conceptualization. Historically, computational support was limited to structural templates; however, contemporary Large Language Models (LLMs) now function as active collaborators in divergent thinking and argumentative scaffolding. This transformation is epitomized by systems that simulate expert personas to critique early-stage ideas and propose novel interdisciplinary intersections. For instance, [15] introduces a framework that explicitly optimizes for novelty by retrieving “inspirations” from past literature and iteratively refining idea suggestions, thereby mitigating the tendency of standard models to generate low-technical-depth concepts. Similarly, [16] demonstrates how LLM-based agents can co-create research questions through breadth-first and depth-first interaction designs, revealing that specific interaction modalities significantly influence perceived creativity and user trust. These systems move beyond simple text generation to facilitate a cognitive partnership where the AI acts as a “rubber duck that talks back,” aiding in sensemaking and hypothesis formulation [17].

Beyond ideation, the structural conceptualization of manuscripts increasingly relies on multi-perspective synthesis and graph-based reasoning. The [18] system exemplifies this by simulating conversations between writers carrying different perspectives to construct comprehensive topic outlines, resulting in significantly more organized and broad coverage compared to baseline retrieval-augmented methods. This approach aligns with the broader vision of including computers as partners in the scientific process by utilizing machine-readable semantics to structure debate and provenance [19]. Furthermore, the integration of knowledge graphs allows for the automatic extraction of entities and relationships to form the backbone of research arguments. [20] leverages background knowledge graphs to predict links and incrementally draft key elements of a paper, such as abstracts and future work, achieving human-preferred outputs in blind Turing tests. Advanced frameworks like [21] and [22] extend this capability by employing graph mining and multi-agent systems to autonomously generate and refine hypotheses, uncovering hidden interdisciplinary relationships that surpass traditional human-driven discovery scales.

However, the efficacy of these generative tools is contingent upon the balance between automation and human agency. While high-scaffolding approaches, such as next-paragraph suggestions, can significantly improve writing quality and productivity, they may concurrently reduce the user’s sense of text ownership [5]. Comparative analyses suggest that while LLMs excel at generating novel ideas, these outputs are often judged as less feasible than those proposed by human experts, highlighting a critical trade-off between novelty and practical implementability [23]. Moreover, the risk of “filter bubbles” persists if recommendation algorithms prioritize relevance over serendipity; systems like [24] address this by facilitating the discovery of scholars with diverse citation profiles to boost innovation. Future directions must therefore focus on mixed-initiative co-creativity systems that allow both humans and AI to initiate content changes, ensuring that the structural conceptualization process remains a synergistic endeavor rather than a passive consumption of algorithmic outputs [25]. Ultimately, the transition to agent-based prewriting demands rigorous evaluation frameworks that assess not only the coherence of the generated structures but also the preservation of the researcher’s intellectual signature and the epistemic validity of the proposed scientific claims.

2.2 Interactive Co-Drafting and Stylistic Adaptation

Building upon the structural conceptualization and hypothesis formulation established in the prewriting phase, the drafting stage represents the operationalization of these ideas into coherent scientific narratives through iterative, symbiotic co-creation. This phase marks a fundamental transformation from solitary, linear composition to a dynamic process where intelligent systems function as active collaborators rather than passive editing tools. Central to this shift are human-in-the-loop revision cycles that leverage continuous feedback loops to refine manuscript quality. Systems such as Read, Revise, Repeat (R3) exemplify this approach, enabling researchers to accept, reject, or modify AI-suggested edits iteratively, thereby achieving high-quality revisions with reduced cognitive load compared to one-shot generation models [26]. Such iterative frameworks are critical because writing is inherently adaptive; modeling the entire editing process, rather than isolated steps, significantly improves the alignment between algorithmic suggestions and human intent [27; 28].

A central challenge within this collaborative dynamic is balancing the preservation of authorial voice with the need to adapt stylistic elements to meet the rigorous conventions of specific academic audiences. Advanced models are increasingly capable of context-aware stylistic refinement, adjusting technical density and rhetorical tone to match venue-specific expectations, such as those of the CHI community, without eroding the researcher’s unique intellectual signature. However, the efficacy of these adaptations depends heavily on the level of scaffolding provided. Empirical studies indicate a U-shaped relationship between AI assistance and writing quality, where high-level scaffolding (e.g., paragraph suggestions) significantly boosts productivity for non-expert writers but may inadvertently reduce the user’s sense of text ownership and satisfaction [5]. Furthermore, the interface design profoundly influences this balance; suggestion lists tend to preserve perceived authorship better than continuous text generation, which can lead to more passive user engagement [29]. To mitigate the risk of voice homogenization, personalized systems like GhostWriter are emerging, which implicitly learn a user’s writing style to generate content that aligns with their specific linguistic patterns, thereby enhancing agency and control [30].

Despite these advances, the integration of AI in drafting introduces significant risks regarding content diversity and factual integrity that necessitate the critical oversight mechanisms discussed in subsequent evidence synthesis workflows. The widespread adoption of feedback-tuned models has been shown to statistically reduce lexical and content diversity across co-written essays, raising concerns about the potential homogenization of scientific discourse [31]. Moreover, while AI tools can produce fluent text, they often exhibit a distinct “writing style” gap characterized by subtle factual inaccuracies and a lack of deep argumentative logic compared to human experts [32]. This is particularly problematic for non-native English speakers (NNESs), who benefit disproportionately from AI assistance in lexical precision but face heightened risks of adopting “tortured phrases” or synthetic writing styles that may trigger scrutiny during peer review [33; 34]. Consequently, the future of interactive co-drafting lies not in full automation but in developing “leaky abstractions” and explainable interfaces, such as ConvXAI, that allow researchers to interrogate AI reasoning and maintain critical oversight over the generative process [35]. Ultimately, sustaining the integrity of scientific writing requires frameworks that treat AI as a scaffold for human creativity rather than a replacement for scholarly judgment, laying the groundwork for the rigorous verification and citation practices essential to the next phase of research.

2.3 Evidence Integration, Citation Integrity, and Claim Verification

The integration of artificial intelligence into evidence synthesis represents a paradigm shift from passive retrieval to active verification, fundamentally altering how scientific claims are grounded in literature. Modern systems increasingly rely on Retrieval-Augmented Generation (RAG) architectures to mitigate the hallucination risks inherent in standalone large language models. By coupling generative capabilities with real-time database access, tools like [36] and [37] facilitate dynamic literature discovery, utilizing vector embeddings to recommend conceptually relevant papers and integrate citations seamlessly during the drafting process. This approach ensures that generated text remains anchored in up-to-date empirical evidence rather than static training data, addressing the critical challenge of maintaining factual accuracy in rapidly evolving fields.

Beyond simple retrieval, the verification of scientific claims requires sophisticated reasoning over heterogeneous evidence sources. Emerging frameworks employ multi-task learning to simultaneously assess claim veracity and locate supporting rationales. For instance, [38] demonstrates the efficacy of adapting sequence-to-sequence models for abstract retrieval, sentence selection, and label prediction, outperforming baselines in the biomedical domain. Similarly, [39] extends these capabilities to open-domain settings, revealing that systems trained on smaller corpora often struggle to generalize when faced with the scale of 500,000 research abstracts, highlighting the necessity for robust pooling techniques in evidence collection. The complexity of this task is further underscored by benchmarks like [40], which demand compositional reasoning over scientific tables, exposing the limitations of current models in grounding claims within structured data.

A persistent threat to citation integrity is the generation of fabricated references, a phenomenon exacerbated by probabilistic text generators. The emergence of “tortured phrases” and non-existent citations in reputable journals signals a growing crisis in attribution accuracy [34]. To counter this, advanced models leverage the structural properties of citation graphs. [41] illustrates that augmenting text representations with citation graph context significantly improves information extraction, while [42] shows that jointly training models to predict both citation function and provenance yields synergistic performance gains. Furthermore, [43] provides empirical evidence that advanced RAG implementations can reduce hallucination rates by nearly 42% compared to standard generative models, particularly when augmented with relevant metadata.

Despite these advancements, significant challenges remain in achieving fully autonomous claim attribution. The [44] benchmark reveals a stark disparity between human and machine performance in identifying referenced papers, with large language models achieving only up to 18.5% accuracy compared to nearly 70% for humans. This gap suggests that while AI can assist in brainstorming coarse-grained citation groupings, it often fails at detailed synthesis without human intervention [45]. Consequently, the future of evidence integration lies not in full automation but in mixed-initiative workflows where AI handles retrieval and preliminary verification, while human experts retain oversight for complex attribution tasks. Future research must prioritize the development of interpretable verification trails and standardized benchmarks like [46] to ensure that AI-assisted writing enhances, rather than erodes, the epistemic integrity of the scientific record.

2.4 AI-Mediated Peer Review Simulation and Response Strategy

Bridging the gap between evidence verification and the broader ethical implications of AI authorship, the integration of artificial intelligence into the peer review lifecycle represents a critical transition from passive text editing to the active simulation of review dynamics and strategic response formulation. This subsection examines how Large Language Models (LLMs) are increasingly deployed

to emulate the critical evaluation processes of human reviewers, offering authors pre-submission diagnostics and assisting in the complex task of drafting rebuttals. Recent empirical studies indicate that LLM-generated feedback can achieve a level of constructiveness and aspect coverage comparable to human reviewers, particularly for identifying weaknesses in early-stage drafts [47]. Systems such as Reviewer2 have advanced this capability by modeling the distribution of review aspects, thereby generating feedback that better captures the diversity of opinions found in human peer reviews [48]. Furthermore, multi-agent frameworks like MARG leverage internal discussions among specialized AI agents to overcome context window limitations, significantly reducing the prevalence of generic comments and enhancing the specificity of feedback on experimental validity and impact [49].

Beyond generating initial critiques, AI systems are being utilized to deconstruct existing reviewer comments through argument mining and sentiment analysis, a necessary step for maintaining the epistemic rigor discussed in prior sections on claim verification. By parsing reviews into elementary argumentative propositions—distinguishing between factual corrections, requests for additional experiments, and subjective evaluations—these tools help authors prioritize revision tasks and understand the underlying logic of reviewer concerns [14]. This structural understanding is critical for formulating effective rebuttals. However, current models face challenges in aligning edits with the implicit intent of reviewer comments rather than their literal wording; for instance, while GPT-4 can address surface-level requests, it often struggles to capture the “spirit” of indirect feedback or incorporate necessary technical nuances [50]. To mitigate this, emerging frameworks employ iterative human-in-the-loop revision cycles, allowing researchers to refine AI-suggested responses through continuous feedback, analogous to the R3 system used for general text revision [26].

A critical frontier in this domain is the simulation of the social and psychological dynamics inherent in peer review, which directly informs the transparency and authorship debates explored in subsequent sections. Traditional analysis is often limited by the opacity of reviewer identities and latent biases. The AgentReview framework addresses this by utilizing LLM agents to simulate multivariate review scenarios, disentangling factors such as authority bias, altruism fatigue, and social influence to predict how these variables impact decision outcomes [51]. Such simulations reveal that reviewer bias can account for significant variation in paper decisions, suggesting that AI-mediated strategies could help authors anticipate and counteract potential biases in their response letters. Moreover, the rise of AI-assisted reviewing itself introduces new strategic considerations; evidence suggests that AI-assisted reviews tend to assign higher scores and may inadvertently increase acceptance probabilities for borderline papers, creating a “review lottery” dynamic that authors must navigate [52].

Despite these advancements, significant challenges remain regarding the reliability and ethical implications of AI-mediated review interactions, reinforcing the need for the governance frameworks discussed in the following section. While tools like ReviewFlow provide intelligent scaffolding to help novice reviewers adhere to community standards, there is ongoing concern that over-reliance on AI for both reviewing and responding may lead to the homogenization of scientific discourse or the erosion of critical human judgment [53]. Additionally, the detectability of AI-generated rebuttals and the potential for “gaming” the review process through automated persuasion strategies pose risks to the integrity of the scientific record [54]. Future research must therefore focus on developing robust evaluation metrics for AI-generated rebuttals and establishing governance frameworks that ensure these tools augment, rather than replace, the nuanced intellectual dialogue essential to scientific progress.

2.5 Transparency, Authorship Attribution, and Process Integrity

The integration of generative artificial intelligence into scientific writing necessitates a fundamental re-evaluation of transparency protocols, authorship attribution, and the mechanisms ensuring process integrity. As AI transitions from a passive editing tool to an active co-creator, the traditional binary distinction between human and machine contribution becomes increasingly untenable. Current discourse emphasizes the adoption of standardized disclosure frameworks, such as “AI Usage Cards,” which provide a structured mechanism for researchers to declare the scope, prompts, and extent of AI involvement across the research lifecycle [55]. Similarly, initiatives like “PaperCard” aim to distinguish between assistive utility and generative contribution, fostering a culture of accountable disclosure rather than concealment. These frameworks address the ethical imperative of transparency, particularly as studies indicate a rising prevalence of AI-generated text in abstracts and manuscripts, often correlated with authors from non-English speaking regions seeking linguistic equity [56].

However, reliance on post-hoc detection methodologies to enforce integrity has proven fundamentally flawed. Empirical evaluations reveal that existing binary classifiers struggle to reliably differentiate between human-written and AI-generated scientific text, often exhibiting high false-positive rates, particularly for non-native speakers [57]. The phenomenon of “tortured phrases”—bizarre semantic substitutions indicative of AI rewriting—highlights the unique artifacts of synthetic text, yet detection remains an arms race where obfuscation techniques easily evade current tools [34]. Furthermore, the emergence of “mixtext,” comprising interleaved human and machine segments, renders traditional detectors ineffective, as they fail to capture the nuanced stylistic adaptations inherent in iterative co-writing [58]. Consequently, the field is pivoting toward process-based verification models that monitor the “writer’s integrity” through behavioral footprints and revision histories, rather than analyzing the final product alone. This approach aligns with findings that human collaboration significantly enhances the detection of deepfake texts by focusing on coherence and consistency gaps that automated systems miss [59].

The ambiguity of authorship is further complicated by the “AI Ghostwriter Effect,” where users psychologically disown AI-generated content yet publicly claim authorship, creating a discrepancy between perceived ownership and actual contribution [60]. This psychological disconnect challenges legal and ethical norms regarding intellectual property, especially as high-end editing services increasingly blur the line between language correction and substantive intellectual input [61]. The stakes are elevated by evidence suggesting that AI-assisted peer reviews may induce score inflation and alter acceptance dynamics, potentially compromising the fairness of the scientific record [52]. Moreover, the deployment of autonomous agents capable of generating full papers raises critical questions about the verifiability of results and the potential for hallucinated references to permeate the literature [62].

Future directions must therefore prioritize the development of robust, process-oriented governance frameworks that safeguard scientific integrity without stifling innovation. This includes the implementation of dynamic provenance tracking systems that link claims to their generative origins and the establishment of community norms that value human oversight as a critical component of authorship. As the scientific ecosystem adapts, the focus must shift from detecting AI presence to verifying the rigor of the human-AI collaborative process, ensuring that the acceleration of discovery does not come at the cost of epistemic trust.

3 Intelligent Systems for Literature Synthesis and Knowledge Discovery

3.1 Semantic Retrieval and Context-Aware Recommendation Architectures

The paradigm of literature discovery is undergoing a fundamental shift from sparse lexical matching to dense semantic understanding, driven by the need to navigate exponentially growing scientific corpora. Traditional keyword-based systems often fail to capture latent conceptual relationships, particularly when terminology diverges across disciplinary boundaries. In response, modern architectures leverage vector embeddings to map documents into high-dimensional semantic spaces, enabling the retrieval of conceptually similar works regardless of surface-level lexical overlap. For instance, the Science Concierge A fast content-based recommendation system for scientific publications demonstrated that content-based algorithms utilizing such semantic similarity metrics significantly outperform keyword-based suggestions in correlating with human judgments of relevance. This transition is further advanced by hybrid models that integrate co-citation networks with textual content, as seen in the large-scale system described in A Scalable Hybrid Research Paper Recommender System for Microsoft Academic, which effectively balances paper novelty against authority while mitigating cold-start problems through the Microsoft Academic Graph.

Beyond static retrieval, the integration of Retrieval-Augmented Generation (RAG) has emerged as a critical architecture for grounding language models in verified scientific evidence. By coupling large language models with real-time database retrieval, systems like PaperQA Retrieval-Augmented Generative Agent for Scientific Research can synthesize answers from full-text articles, providing transparent justification trails that mitigate the hallucination risks inherent in parametric knowledge. This approach ensures that recommendations are not merely statistically probable but evidentially supported. Furthermore, the fusion of textual semantics with structured knowledge graphs enables more sophisticated reasoning. The OAG-BERT Towards A Unified Backbone Language Model For Academic Knowledge Services exemplifies this by pre-training on heterogeneous entity knowledge within the Open Academic Graph, creating a unified backbone that enhances zero-shot inference for diverse academic services without expensive human labeling. Similarly, domain-specific applications, such as the multi-modal models in A Molecular Multimodal Foundation Model Associating Molecule Graphs with Natural Language, bridge chemical structures with natural language descriptions, facilitating cross-modal retrieval that unlocks novel interdisciplinary connections.

To combat the “filter bubble” effect where algorithms reinforce existing disciplinary silos, recent architectures explicitly optimize for serendipity and cross-domain synthesis. The Bursting Scientific Filter Bubbles Boosting Innovation via Novel Author Discovery introduces faceted author representations to identify commonalities between disparate scientific communities, fostering innovation through non-obvious connections. Context-awareness is further refined by incorporating metadata and social signals; From Who You Know to What Read Augmenting Scientific Recommendations with Implicit Social Networks illustrates how augmenting recommendations with textual relevance messages derived from implicit social networks increases user engagement without biasing towards highly cited authors. Additionally, workflow integration has become paramount, with tools embedding semantic search directly into reading environments to reduce context switching. The Semantic Reader Project Augmenting Scholarly Documents through AI-Powered Interactive Reading Interfaces and Threddy An Interactive System for Personalized Thread-based Exploration and Organization of Scientific Literature demonstrate how interactive interfaces can scaffold the synthesis of research threads by leveraging authors’ existing summarizations and providing just-in-time definitions.

Despite these advances, challenges remain in balancing precision with novelty and ensuring the interpretability of vector-based recommendations. While systems like SciMON Scientific Inspiration Machines Optimized for Novelty attempt to iteratively optimize for novelty by comparing generated ideas against prior literature, the trade-off between technical depth and disruptive potential requires further investigation. Future directions point toward autonomous, multi-agent systems capable of dynamic graph reasoning, as envisioned in SciAgents: Automating scientific discovery through multi-agent intelligent graph reasoning, which utilize ontological knowledge graphs to uncover hidden interdisciplinary relationships. Ultimately, the evolution of these architectures signifies a move toward symbiotic scholarship, where intelligent systems do not merely retrieve information but actively participate in the conceptualization and synthesis of scientific knowledge.

3.2 Automation of Systematic Review Pipelines and Data Extraction

Bridging the gap between the semantic discovery of relevant literature and the dynamic sense-making required for synthesis, the automation of systematic literature reviews (SLRs) represents a critical frontier in managing the exponential growth of scientific output. This stage transitions research workflows from manual, labor-intensive protocols to semi-autonomous intelligent pipelines capable of handling the rigorous screening and precise data extraction necessary for large-scale analysis. Traditional SLRs require the evaluation of thousands of documents, a process inherently prone to human error and fatigue. Recent advancements leverage machine learning to optimize these stages, beginning with automated screening and eligibility classification. Supervised models have demonstrated the capacity to filter abstracts against inclusion criteria with high precision, allowing human reviewers to focus on borderline cases [8]. While early approaches relied on static classifiers, contemporary systems increasingly employ active learning frameworks where algorithms iteratively prioritize uncertain instances for human validation, thereby optimizing the trade-off between automation efficiency and methodological fidelity [63]. This human-in-the-loop paradigm ensures that rigorous standards are maintained while significantly reducing the cognitive load on researchers [64].

Once relevant literature is identified, the extraction of structured data from unstructured full-text articles remains a formidable challenge due to the complexity of scientific discourse and the variability of reporting standards. Domain-adapted large language models (LLMs) and Named Entity Recognition (NER) systems are now being fine-tuned to identify specific experimental parameters, statistical outcomes, and clinical variables with increasing accuracy [65]. However, the performance of these models often degrades when faced with cross-document dependencies or nuanced contextual cues. To address this, multi-task learning frameworks that jointly identify entities, relations, and coreference clusters have shown superior performance over pipelined approaches by mitigating cascading errors [66]. Furthermore, leveraging the citation graph as a complementary source of context has proven effective in enhancing information extraction, as the network of referential links provides implicit semantic signals that raw text alone may lack [41]. These extracted structures serve as the essential foundational data that subsequent argumentative dialogue systems rely upon to resolve conflicts and synthesize findings.

The integration of crowdsourcing mechanisms offers a scalable solution for generating the high-quality annotated corpora necessary to train these sophisticated models. Microtask platforms enable the rapid creation of gold-standard datasets for entity annotation, such as disease mentions or software usage, at a fraction of the cost and time required by expert-only annotation [67; 68]. These crowdsourced resources are instrumental in developing robust extractors capable of handling domain-specific terminology across diverse fields, from biomedicine to software engineering

[69]. Despite these advances, challenges persist regarding the reliability of LLMs in “human-out-of-the-loop” scenarios. Recent evaluations indicate that while models like GPT-4 can approach human-level performance in screening tasks under highly reliable prompting, their data extraction capabilities often suffer from moderate accuracy and susceptibility to hallucination without strict validation protocols [70].

Future directions for SLR automation lie in the synthesis of retrieval-augmented generation (RAG) with formal verification steps to ensure traceability and factual grounding, directly supporting the conversational engagement described in the following section. The emergence of multi-agent systems promises to further decentralize the review workflow, assigning distinct agents to search, filter, summarize, and critique evidence, effectively simulating a team of digital researchers [64]. Moreover, the field is moving toward updating reporting guidelines, such as PRISMA, to explicitly incorporate AI-driven processes, ensuring transparency and reproducibility in automated syntheses [71]. Ultimately, the goal is not the full replacement of human judgment but the creation of symbiotic workflows where intelligent systems handle volumetric processing and pattern recognition, providing the structured, verified data necessary for researchers to engage in higher-order synthesis and critical interpretation.

3.3 Interactive Sensemaking and Argumentative Dialogue Systems

The paradigm of literature synthesis is rapidly shifting from static retrieval to dynamic, conversational engagement, where intelligent systems function as active collaborators in the sensemaking process. This transition addresses the fundamental cognitive bottleneck researchers face when navigating exponentially growing corpora: the difficulty of maintaining context while synthesizing disparate, often conflicting, findings. Modern argumentative dialogue systems leverage Retrieval-Augmented Generation (RAG) architectures to ground natural language interactions in verified scientific evidence, thereby mitigating the hallucination risks inherent in standalone large language models. For instance, agents like PaperQA have demonstrated the capacity to match expert human performance in answering complex questions by performing information retrieval across full-text articles and assessing source relevance before synthesis [72]. Such systems move beyond simple question-answering to facilitate deep diagnostic querying, allowing researchers to probe the logical structure of arguments within and across documents.

A critical capability of these systems is the automated resolution of scientific conflicts. Unlike traditional search engines that return disparate results, advanced dialogue agents can identify contradictions or consensus across multiple studies, presenting side-by-side comparisons of methodologies and outcomes. This is exemplified by frameworks like Climinator, which employs a Mediator-Advocate architecture to synthesize varying scientific perspectives on climate change, iteratively debating diverse viewpoints to converge on evidence-based conclusions [73]. Similarly, the integration of argument mining techniques enables the decomposition of peer reviews and scientific texts into structured propositions, allowing users to trace the evolution of hypotheses and rebuttals over time [14]. By visualizing these argumentative graphs, systems help researchers navigate discordant claims and understand the provenance of specific scientific assertions.

Furthermore, interactive sensemaking tools are increasingly incorporating “human-in-the-loop” mechanisms to refine hypothesis generation and validation. Systems such as ResearchAgent utilize iterative feedback loops where reviewing agents critique generated ideas against scientific literature, mirroring the collaborative refinement process of human research teams [74]. This dialogic approach extends to data analysis workflows, where generative AI acts as a “rubber duck” that talks back,

assisting users in formulating queries and verifying results through participatory prompting [17]. However, challenges remain in ensuring the fidelity of these interactions. Studies indicate that while LLMs can structure vast knowledge bases to propose testable hypotheses, they still struggle with predicting novel empirical patterns absent from their training data, necessitating rigorous human oversight [75; 62].

Future directions for argumentative dialogue systems lie in enhancing explainability and reducing the cognitive load of verification. The development of conversational explainable AI (XAI) agents, such as ConvXAI, offers heterogeneous explanations through dialogue, allowing researchers to drill down into the reasoning behind AI suggestions [35]. Additionally, the integration of tool-augmented reasoning, as seen in SciAgent, shifts the focus from omniscient solvers to proficient tool-users that can retrieve and execute domain-specific functions to validate claims [76]. As these systems evolve, the ultimate goal is to create symbiotic partnerships where AI handles the scalability of information retrieval and conflict detection, while humans retain agency over critical interpretation and creative synthesis, thereby accelerating the pace of discovery without compromising scientific rigor.

3.4 Knowledge Graph Construction and Autonomous Hypothesis Generation

Building directly on the argumentative grounding and conflict resolution capabilities established in conversational systems, the next evolutionary step involves transforming unstructured scientific literature into machine-actionable knowledge graphs (KGs). This shift marks a transition from passive information retrieval to active, autonomous discovery, necessitating the precise extraction of entities, relations, and causal mechanisms to construct robust ontological frameworks. Foundational efforts in this domain have established multi-task learning frameworks, such as SciIE, which jointly identify entities, relations, and coreference clusters to mitigate cascading errors during graph construction [66]. Similarly, domain-specific initiatives like SoftwareKG and ORKG-Leaderboards demonstrate the viability of extracting structured procedural and empirical data from heterogeneous sources, enabling the tracking of software usage and state-of-the-art performance metrics respectively [68; 77]. However, recent analyses indicate that purely text-based extraction often yields shallow synthesis, lacking the deep interconnectedness required for rigorous hypothesis generation without human intervention [45].

To overcome these limitations and leverage the structured insights prepared by dialogue systems, advanced frameworks are increasingly employing graph reasoning and multi-agent architectures to autonomously identify hidden patterns and propose novel research directions. The PaperRobot framework exemplifies this trajectory by constructing comprehensive background KGs to predict linkages between entities, thereby generating new paper titles and abstracts that simulate human creativity [20]. Building upon this, the AGATHA system employs deep learning over citation graphs to prioritize plausible term-pairs, effectively winnowing vast search spaces to recommend high-probability research targets [21]. More recently, the SciAgents framework has introduced a “swarm intelligence” paradigm, where multi-agent systems traverse large-scale ontological graphs to uncover interdisciplinary relationships in materials science that were previously obscured by domain silos [22]. These agents do not merely retrieve information but engage in iterative critique and refinement of hypotheses, mirroring the peer review process to enhance the validity of generated ideas [78].

Despite these advances, significant challenges persist regarding the completeness and formal semantics of constructed graphs. While some approaches successfully map shallow semantic structures in materials synthesis [79], others argue for the necessity of formal logic representations, such as

the “super-pattern” in higher-order logic, to enable automated consistency checking of high-level scientific claims [80; 81]. Furthermore, the integration of Large Language Models (LLMs) as “soft” knowledge graphs offers a complementary pathway, where pre-trained models provide latent knowledge for link prediction in sparse data regimes, particularly in biomedical domains [82]. Emerging probabilistic frameworks like GFlowNets further enhance this capability by sampling diverse, high-reward candidates from vast hypothesis spaces, addressing the exploration-exploitation trade-off inherent in autonomous discovery [83].

Future directions must address the tension between automated scale and semantic fidelity, bridging the gap between the graph-based discovery engines described here and the fully autonomous workflows anticipated in subsequent sections. While fully autonomous agents like The AI Scientist can generate and evaluate papers end-to-end [84], the risk of hallucinated connections remains a critical barrier. Consequently, the next generation of systems will likely adopt hybrid human-in-the-loop architectures, where LLMs propose hypotheses based on graph inference, but human experts or specialized verification agents validate the causal mechanisms, ensuring that the acceleration of discovery does not compromise scientific rigor [85]. This symbiotic approach promises to unlock the full potential of KGs as dynamic engines for interdisciplinary innovation, serving as the structural backbone for the next generation of automated research pipelines.

4 Integration of Generative Models into Experimental Design and Data Analysis Workflows

4.1 Natural Language Interfaces for Code Synthesis and Pipeline Automation

The integration of Large Language Models (LLMs) into computational research workflows represents a paradigm shift from static scripting environments to dynamic, intent-driven programming interfaces. By translating natural language descriptions into executable code, these systems significantly lower the barrier to entry for domain scientists who possess deep theoretical knowledge but limited software engineering expertise. Recent advancements have moved beyond simple code completion to sophisticated “intent-to-code” translation, where models generate complex Python, R, or SQL scripts for data cleaning, statistical testing, and visualization based on high-level researcher prompts [86]. This capability is particularly transformative in interactive notebook environments, such as Jupyter, which serve as the primary locus for exploratory data analysis. Systems like ARCADE demonstrate that effective code synthesis in these contexts requires models to comprehend rich multi-modal contexts, including existing cell states and execution history, to maintain logical coherence across multi-turn interactions [86].

However, the transition from single-line autocompletion to full pipeline automation introduces significant challenges regarding reliability and verification. While tools like GitHub Copilot have been shown to enhance productivity in collaborative settings, particularly for iterative maintenance tasks, they can inadvertently shift the researcher’s role from author to evaluator, potentially leading to productivity losses if the cognitive load of verifying generated code exceeds the time saved in writing it [87; 88]. To mitigate these risks, emerging frameworks employ “chaining” strategies, where complex tasks are decomposed into modular sub-steps, allowing users to inspect intermediate results and calibrate model expectations before proceeding [89]. This modular approach enhances controllability and transparency, enabling researchers to “unit-test” specific components of a generated analysis pipeline. Furthermore, specialized agents are being developed to handle domain-specific complexities; for instance, MatPlotAgent utilizes visual feedback mechanisms to iteratively debug and refine code for scientific data visualization, bridging the gap between textual intent and graph-

ical output [90].

The orchestration of multi-step workflows also benefits from natural language interfaces that can manage dependencies and execution flows. Frameworks like AutoFlow automate the generation of agent workflows, optimizing the sequence of operations required to solve complex tasks without manual intervention [91]. Similarly, systems such as Iris leverage conversational strategies to compose sequences of commands, allowing users to weave together disparate operations—such as log-transforming data followed by histogram plotting—through natural language dialogue [92]. Despite these advances, empirical evaluations suggest that current multimodal agents still struggle with fine-grained, knowledge-intensive tasks in professional data engineering environments, achieving low success rates on end-to-end workflow automation [93]. Moreover, the integrity of generated code remains a critical concern, as LLMs may produce syntactically correct but logically flawed scripts that lead to incorrect scientific conclusions [94]. Consequently, the future of natural language interfaces for code synthesis lies not in full autonomy, but in symbiotic collaboration where AI handles routine syntactic construction and debugging [95], while humans retain oversight over experimental design and logical validation. Future research must focus on developing robust verification layers and interactive debugging tools that allow researchers to trace the provenance of generated code, ensuring that the acceleration of computational workflows does not come at the expense of scientific rigor.

4.2 Autonomous Experimental Design and Closed-Loop Optimization

Building upon the computational workflows and verification layers discussed previously, the integration of generative artificial intelligence into experimental design marks a paradigmatic shift from static, human-led planning to dynamic, closed-loop optimization systems capable of autonomous discovery. This subsection examines architectures where Large Language Models (LLMs) function not merely as passive assistants for code synthesis but as active agents that formulate hypotheses, navigate high-dimensional parameter spaces, and iteratively refine experimental protocols based on real-time feedback. The core mechanism driving this autonomy is the synthesis of generative reasoning with probabilistic optimization frameworks, such as Bayesian optimization, to maximize information gain while minimizing resource expenditure, effectively extending the “intent-to-code” capabilities of the prior section into the physical realm.

Recent advancements demonstrate the efficacy of LLM-driven agents in proposing novel experimental configurations by synthesizing prior literature and domain constraints. For instance, [96] illustrates how LLMs can generate untrained yet validated hypotheses from testing literature, effectively breaking down interdisciplinary information barriers. When coupled with active learning loops, these systems transcend simple suggestion engines to become robust optimization engines. [97] exemplifies this by integrating LLMs with Bayesian optimization to translate extracted literature knowledge into actionable parameters for catalyst synthesis, significantly accelerating the optimization of ammonia production protocols. Similarly, [98] shows that agents leveraging biological knowledge can outperform traditional Bayesian optimization baselines in detecting desired phenotypes, achieving an 18% improvement without explicit acquisition function design.

The realization of fully autonomous “self-driving laboratories” represents the culmination of these closed-loop capabilities, where AI agents control robotic infrastructure to execute physical experiments, bridging the gap between digital planning and tangible results. [99] showcases an intelligent agent system successfully performing catalyzed cross-coupling reactions through autonomous design, planning, and execution. In materials science, [100] details a cooperative workflow where

GPT-4 provides detailed synthesis instructions and learns from human feedback on experimental outcomes via in-context learning, enabling the discovery of isorecticular Metal-Organic Frameworks (MOF). Furthermore, [101] highlights the role of natural language interfaces in allowing robots to reason over user inputs, plan long-horizon tasks, and adapt to diverse chemistry experiments, thereby reducing physical workload and improving user experience.

However, the transition to full autonomy introduces critical challenges regarding safety, verification, and the balance between exploration and exploitation, echoing the reliability concerns raised in computational scripting but with higher stakes. The potential for LLM agents to design harmful experiments or bypass safety regulations necessitates rigorous governance frameworks. [102] emphasizes a triadic framework involving human regulation, agent alignment, and environmental feedback to mitigate vulnerabilities inherent in autonomous scientific agents. Moreover, while systems like [84] demonstrate the feasibility of end-to-end research cycles—from idea generation to paper writing—they also reveal limitations in novelty and feasibility compared to human experts, particularly in complex domains.

Future directions point toward hybrid human-AI models where intermittent expert steering guides asynchronous autonomous loops, ensuring alignment with scientific values while maintaining efficiency. [103] argues that for the foreseeable future, human operators will remain essential to monitor experiment progression in feature space and tune agent policies, suggesting that “human-aware” automation [104] which predicts human predictions and avoids crowd biases, offers a more robust path to discovery than purely autonomous agents. Ultimately, the convergence of generative modeling, robotic automation, and safety-aware governance defines the next frontier of experimental science. These autonomous physical workflows generate rich multimodal data, setting the stage for the subsequent discussion on how multimodal large language models interpret, visualize, and validate these complex experimental outcomes.

4.3 Multimodal Data Interpretation and Visual Analytics

The integration of multimodal large language models (MLLMs) into scientific workflows marks a paradigm shift from static visualization tools to dynamic, conversational agents capable of joint reasoning over text, tabular data, and complex figures. Unlike traditional visual analytics systems that require explicit user manipulation of graphical interfaces, emerging MLLMs facilitate deep data interrogation by interpreting abstract scientific plots and generating semantic insights autonomously. Recent benchmarks, such as [105] and [106], underscore the critical gap in current models’ ability to interpret domain-specific figures, revealing that while general-purpose vision-language models excel in natural scenes, they often struggle with the nuanced semantics of academic charts without specialized fine-tuning. To address this, frameworks like [107] have demonstrated that aligning foundation models with scientific multimodal instructions significantly enhances performance, surpassing human baselines in tasks requiring visual and linguistic coherence.

A primary advancement in this domain is the transition toward agentic visualization systems that automate the generation and refinement of scientific graphics. Systems such as [90] leverage iterative debugging and visual feedback loops to translate high-level natural language queries into executable plotting code, effectively bridging the gap between researcher intent and technical implementation. This capability is further augmented by tools like [108], which autonomously design versatile interfaces for data management and visualization, transforming raw heterogeneous data into visual narratives that match user intent. However, the reliance on purely natural language interfaces can limit exploratory flexibility; consequently, hybrid approaches like [109] have emerged, blend-

ing natural language commands with dynamically synthesized UI widgets. This mixed-initiative interaction allows researchers to specify edits verbally while retaining the granular control and immediate visual feedback inherent to graphical user interfaces, thereby reducing the cognitive load associated with complex visualization authoring [110].

Beyond generation, multimodal AI is revolutionizing the verification and documentation of analytical findings. The challenge of ensuring alignment between textual claims and graphical evidence is being addressed through cross-modal synthesis techniques. For instance, [111] introduces a sketch-based interaction mechanism within computational notebooks, enabling users to highlight specific data subsets to trigger automated finding documentation, thus ensuring that narrative explanations remain grounded in visual evidence. Similarly, [112] utilizes conditional parallel coordinates to open the “black box” of automated machine learning workflows, visualizing the model generation process to enhance user trust and understanding. In the biomedical domain, [113] exemplifies the potential of grounded report generation, where pixel-level visual clues from MRI scans are directly linked to textual findings, improving diagnostic accuracy and reporting consistency.

Despite these advances, significant challenges remain regarding the robustness of multimodal reasoning in open-domain scientific contexts. Evaluations using datasets like [114] and [115] reveal that current models often fail to capture complex logical relationships within multi-turn dialogues about graphs or to synthesize information across disparate modalities without hallucination. Furthermore, the automation of data storytelling workflows must balance efficiency with the preservation of human agency, as evidenced by studies indicating that while AI can streamline slide creation [116; 117], full automation risks oversimplifying nuanced scientific narratives. Future research must therefore focus on developing more rigorous alignment strategies that enforce logical consistency between data, visuals, and text, while fostering collaborative interfaces that support iterative sensemaking rather than mere output generation.

4.4 Provenance Tracking and Reproducibility Automation

Building upon the imperative for rigorous alignment between data, visuals, and text established in multimodal reasoning, the integration of generative models into experimental design necessitates a parallel evolution in provenance tracking to ensure the integrity of the scientific record. As AI agents increasingly automate code synthesis, parameter optimization, and data interpretation, the traditional manual documentation of methods becomes insufficient. The core challenge lies in capturing the dynamic, often non-deterministic interactions between human intent and algorithmic execution. Intelligent agents must therefore transition from passive tools to active stewards of reproducibility, automatically logging execution metadata, managing dependency trees, and constructing verifiable lineage graphs. This shift is critical because, as noted in the broader discourse on computational science, an article describing results is merely advertising; the actual scholarship comprises the full software environment, code, and data that produced the finding [118].

Current approaches to provenance automation vary in granularity and interoperability. Foundational ontologies like PAV (Provenance, Authoring and Versioning) provide a lightweight framework for distinguishing between creators, authors, and curators, extending general-purpose standards like W3C PROV-O to specific scientific needs [13]. However, static ontologies often struggle with the fluidity of AI-driven workflows where agents iteratively refine code. Systems like AiiDA address this by embedding provenance recording directly into high-throughput workflow engines, automatically preserving the full data lineage in queryable relational databases [119]. Similarly, YesWorkflow offers a language-independent mechanism to recover workflow information from standard scripting

languages via annotated comments, bridging the gap between flexible scripting and rigid workflow management systems [120]. Despite these advances, empirical studies indicate significant gaps in practice; for instance, automated analyses of Jupyter notebooks in biomedical research reveal that a vast majority fail to reproduce due to missing dependencies or environment mismatches, highlighting the limitations of current manual or semi-automated documentation practices [121].

Emerging trends leverage generative AI not just to execute tasks, but to audit and document them autonomously, extending the concept of automated visualization to the entire research lifecycle. Frameworks like DataDreamer enable researchers to construct reproducible LLM workflows by programmatically capturing every step of synthetic data generation and model interaction, ensuring that the stochastic nature of LLMs does not obscure the research path [122]. Furthermore, fully autonomous “data-to-paper” systems demonstrate the potential for end-to-end traceability, where agents back-track information flow from raw data to final manuscript, creating inherently verifiable research artifacts [123]. This aligns with the concept of the “Whole Tale,” which advocates for coupling data products with research articles to create living, executable publications [124]. Nevertheless, challenges remain in scaling these solutions. Benchmarks like CORE-Bench reveal that even state-of-the-art agents struggle with the complex task of setting up and executing existing research repositories, achieving low success rates in reproducing results without human intervention [125].

To bridge this gap, future systems must integrate co-audit mechanisms where AI agents actively verify their own outputs and provenance logs against ground truth or peer constraints [126]. The convergence of containerization technologies with intelligent workflow managers promises to encapsulate the entire computational context, mitigating the “it works on my machine” phenomenon [127]. Ultimately, the goal is to move from post-hoc verification to real-time provenance construction, where every algorithmic decision is timestamped, justified, and linked to its input data. This requires a cultural and technical shift towards treating provenance not as an administrative overhead, but as a first-class citizen in the architecture of AI-augmented science, ensuring that as automation accelerates discovery, it does not erode the trustworthiness of the resulting knowledge.

5 Human-Centric Evaluation, Trust, and Cognitive Dynamics in Collaborative Research

5.1 Metrics and Frameworks for Assessing Collaborative Synergy

Evaluating the efficacy of human-AI collaboration in scientific research necessitates a paradigm shift from assessing isolated agent performance to measuring the emergent properties of the combined team. Traditional metrics focusing solely on task completion rates or final output accuracy fail to capture the nuanced dynamics of co-creation, particularly in open-ended domains where ground truth is often ambiguous or unavailable. Consequently, recent scholarship has pioneered frameworks that quantify “collaborative synergy”—the extent to which the joint performance of human and machine exceeds the sum of their individual capabilities. A primary challenge in this domain is disentangling the specific intellectual contributions of human versus machine agents. While information-theoretic approaches, such as mutual information metrics, offer a theoretical basis for distinguishing proportional input in co-generated manuscripts, practical application remains complex due to the iterative nature of modern co-writing systems [128]. Furthermore, studies indicate that automated metrics often diverge significantly from human evaluations, particularly in tasks requiring deep synthesis, such as multi-document literature reviews, where standard n-gram similarity scores may be anti-correlated with perceived quality [129].

To address these limitations, the field is moving toward comprehensive benchmark suites designed to evaluate agents across multi-step research workflows rather than static prediction tasks. Emerging benchmarks like SciBench and OlympicArena rigorously test collegiate-level scientific problem-solving and cognitive reasoning, revealing that even state-of-the-art models struggle with the complex, interdisciplinary logic required for genuine discovery [130; 131]. Similarly, domain-specific evaluations such as LAB-Bench assess practical research capabilities, including protocol planning and data interpretation, providing a more granular view of an AI’s utility as a research assistant [132]. These benchmarks highlight a critical trade-off: while AI agents can generate ideas rated as more novel than those of human experts, these suggestions often suffer from lower feasibility, underscoring the need for metrics that balance originality with practical implementability [23].

Beyond output quality, rigorous evaluation must account for the interaction patterns that define successful teamwork. Frameworks like Parachute propose integrative views of interaction, categorizing metrics based on fluidity, turn-taking, and the system’s ability to recover from joint errors [133]. The concept of “mixed-initiative” co-creativity further expands the design space, suggesting that systems allowing both human and AI to initiate content changes yield higher user satisfaction and creative output than simple prompt-response models [25]. However, the “ironies of automation” warn that poorly designed interactions can lead to productivity losses, where users shift from production to evaluation roles without gaining efficiency [88]. Therefore, future metrics must incorporate dimensions of cognitive load and workflow restructuring to ensure that AI augmentation truly amplifies human potential rather than introducing friction. Ultimately, establishing robust standards for collaborative synergy requires a synthesis of quantitative benchmarks, qualitative user studies, and process-oriented logging to fully capture the transformative potential of human-AI scientific partnerships.

5.2 Trust Calibration, Reliance Strategies, and Mental Models

While robust evaluation metrics provide the quantitative foundation for assessing collaborative synergy, the effective integration of artificial intelligence into scientific workflows ultimately hinges on the psychological alignment between researcher expectations and system capabilities. This alignment is governed by trust calibration—the dynamic adjustment of reliance based on perceived performance—yet it is frequently disrupted by cognitive biases and opaque system behaviors. A primary challenge is automation complacency, where researchers, particularly novices, uncritically accept AI outputs, leading to performance degradation despite high tool accuracy. Empirical studies in software engineering and data analysis reveal that while AI assistance can accelerate task completion, it often induces a false sense of competence, causing users to overlook errors or fail to verify generated code and statistical interpretations [134; 135]. This phenomenon is exacerbated when AI systems provide confident but incorrect explanations, which can distort a researcher’s decision confidence more severely than silent errors, as misleading rationalizations reinforce flawed mental models of the underlying scientific logic [136].

These trust dynamics are further complicated by the interplay between user domain expertise and specific AI model tuning. Research indicates a non-linear relationship between expertise and trust calibration: while experts often maintain appropriate skepticism, mid-level performers exhibit the greatest variability in reliance, sometimes deferring to AI even when their own judgment is superior, a dynamic observed in medical imaging and blood vessel labeling tasks [137; 4]. This suggests that optimal collaboration requires “complementary tuning,” where AI systems are explicitly designed to handle instances where human performance is weakest, rather than attempting to outperform humans across all domains [3]. However, current AutoML and data science tools often fail to com-

municate these boundary conditions, leading to a mismatch between the desired level of automation and the actual transparency provided, particularly for complex, open-ended research questions [138; 139]. Without such calibration, the immediate productivity gains risk obscuring a decoupling of performance from metacognitive accuracy, setting the stage for the “illusion of competence” and potential skill atrophy discussed subsequently.

To mitigate these risks and foster the transparency necessary for long-term researcher development, the alignment of mental models through transparent system design is imperative. Researchers must develop accurate internal representations of an AI’s limitations to prevent both over-trust and under-utilization. Approaches such as “behavior descriptions”—detailing how a model performs across specific subgroups of data—have been shown to significantly improve human-AI team accuracy by helping users identify failure modes and calibrate reliance accordingly [140]. Furthermore, interactive documentation like “Interactive Model Cards” and structured reporting frameworks such as “AI Usage Cards” offer mechanisms to bridge the gap between technical specifications and user understanding, fostering a culture of accountable transparency [141; 55]. Future directions must move beyond static explanations toward dynamic, context-aware feedback loops that allow researchers to “unit-test” AI reasoning chains, thereby transforming passive consumers of AI output into active, critical collaborators capable of steering intelligent agents through complex scientific discovery pipelines [89]. Only by rigorously addressing these cognitive dynamics can the scientific community harness the full potential of human-AI teaming, ensuring that trust calibration serves as a safeguard for research integrity rather than a precursor to the developmental challenges outlined in the following section.

5.3 Cognitive Dynamics, Creativity, and Skill Acquisition Trajectories

The integration of generative artificial intelligence into scientific workflows precipitates a complex reconfiguration of cognitive dynamics, where the immediate augmentation of productivity often obscures longitudinal shifts in skill acquisition and creative agency. A primary concern is the decoupling of performance from metacognitive accuracy, creating an “illusion of competence” wherein researchers achieve high-quality outputs without deepening their underlying conceptual understanding. Empirical evidence suggests that while AI assistance can inflate task performance and user confidence, it may simultaneously degrade the ability to self-assess one’s own understanding, particularly when the AI handles foundational tasks like coding or statistical analysis [142]. This phenomenon is exacerbated by the “ironies of automation,” where the shift from production to evaluation roles can lead to productivity losses if users lack the expertise to effectively critique generated content [88].

The impact on creativity reveals a nuanced tension between individual and collective divergence. While AI tools can serve as catalysts for ideation by expanding the search space for novel hypotheses and simulating interdisciplinary personas [23], there is a documented risk of content homogenization. Studies indicate that co-writing with feedback-tuned models can statistically reduce lexical and content diversity, leading to a convergence of writing styles and argumentative structures across the scientific community [31]. Furthermore, the “flare and focus” dynamics of creative work are altered; while AI aids in generating volume, it may constrain the “scrutability” required for deep creative iteration unless interfaces are explicitly designed to support mixed-initiative co-creativity [25; 143]. The “dearth of the author” concept further posits that when users produce text without making commensurate creative decisions, the resulting output lacks expressive intent, potentially eroding the development of a distinctive scholarly voice [144].

For early-career researchers, these dynamics pose significant risks to skill acquisition trajectories. Heavy reliance on AI for routine tasks such as data cleaning, code debugging, and literature synthesis may impede the mastery of foundational technical competencies essential for independent inquiry [138]. This “skill atrophy” is particularly critical in domains where debugging and iterative refinement are primary mechanisms for learning domain logic. However, AI also offers opportunities for scaffolding; for instance, non-regular writers and less tech-savvy users often benefit significantly from high levels of AI scaffolding, which can accelerate their progression toward proficiency without imposing excessive cognitive load [5]. The challenge lies in designing “human-aware” systems that predict not only scientific outcomes but also the distribution of human expertise, ensuring that AI interventions complement rather than substitute the cognitive struggle necessary for deep learning [104].

Ultimately, the trajectory of skill acquisition in an AI-augmented era depends on the balance between automation and active engagement. Future research must prioritize longitudinal studies that track the evolution of researcher competencies over time, distinguishing between the efficient completion of tasks and the genuine internalization of scientific reasoning. Developing adaptive interfaces that modulate assistance based on user expertise—fading support as competence grows—will be crucial to mitigating the risks of over-reliance while harnessing AI’s potential to elevate the frontier of scientific creativity [4].

5.4 Interface Design for Agency, Explainability, and Conflict Resolution

Building on the necessity for adaptive interfaces that modulate assistance to prevent skill atrophy, the transition from static AI tools to dynamic scientific partners requires interface designs that rigorously preserve human agency while providing transparent access to algorithmic reasoning. Current paradigms often rely on simplistic prompt-response mechanisms that obscure the decision-making process, exacerbating the “illusion of competence” where researchers overestimate their understanding of AI-generated outputs. To counteract this, emerging frameworks advocate for interaction patterns that support iterative refinement and bidirectional clarification, moving beyond mere suggestion generation to true co-creation. For instance, the concept of “Chaining” allows users to decompose complex scientific tasks into modular steps, modifying intermediate results to calibrate model expectations and debug unexpected outputs, thereby significantly enhancing system transparency and controllability [89]. Similarly, conversational agents like ConvXAI demonstrate that facilitating heterogeneous explanations through dialogue improves human-perceived understanding and writing quality compared to static graphical interfaces [35].

A critical design tension exists between abstracting away technical complexity and exposing sufficient internal logic to foster trust, directly addressing the need to balance automation with active engagement. While traditional usability principles favor seamless abstractions, scientific workflows often benefit from “leaky abstractions” that reveal low-level reasoning traces, data provenance, and uncertainty estimates. This approach aligns with findings that behavior descriptions—details of how AI systems perform on specific subgroups—help researchers identify AI failures and appropriately calibrate reliance [140]. Furthermore, the workflow sequencing of AI interventions profoundly impacts decision-making; requiring researchers to register provisional hypotheses before viewing AI inferences can mitigate anchoring effects and preserve critical evaluation skills [137]. In contrast, fully autonomous agents risk eroding intellectual ownership unless mechanisms for conditional delegation are implemented, allowing humans to define trustworthy regions of operation rather than blindly accepting global model outputs [145].

Conflict resolution between human intuition and algorithmic recommendations necessitates interfaces that explicitly visualize sources of divergence and support constructive skepticism, ensuring the “scrutability” required for deep creative iteration. Systems that employ mixed-initiative interactions, where both human and AI can propose changes, have been shown to outperform rigid, user-led or system-led configurations in creative tasks [25]. However, consistent with the need to tailor support to varying expertise levels, the degree of scaffolding must be personalized; while high-level scaffolding aids novice writers, it may reduce the sense of text ownership among experts, suggesting a need for adaptive interfaces that modulate automation levels based on user expertise [5]. Moreover, the integration of dynamic UI widgets synthesized from natural language commands, as seen in DynaVis, offers a middle ground by combining the flexibility of language with the precision of direct manipulation, enhancing editing confidence through immediate visual feedback [109].

Ultimately, effective interface design for scientific collaboration must prioritize scrutability—the ability for users to probe and test model boundaries—and ensure that the human remains the primary agent of discovery, thereby securing the trajectory of genuine skill internalization. Future research should focus on developing standardized metrics for evaluating not just output quality, but the quality of the interactive process itself, including measures of user enjoyment, ownership, and cognitive load [146]. By embedding principles of conditional delegation, leaky abstractions, and adaptive scaffolding, next-generation interfaces can transform AI from a black-box oracle into a transparent, accountable partner that amplifies rather than supplants human scientific ingenuity.

6 Ethical Frameworks, Governance, and the Future of Scientific Integrity

6.1 Redefining Authorship and Intellectual Property in the Age of Generative Agents

The integration of generative agents into scientific workflows precipitates a fundamental crisis in the ontology of authorship, challenging the centuries-old convention that intellectual ownership resides exclusively with human agents. As artificial intelligence transitions from passive editorial assistance to active participation in hypothesis formulation and experimental design, the boundary between tool and co-author dissolves. Traditional metrics for attributing credit, often based on linear contributions to writing or data collection, fail to capture the nuanced, iterative nature of human-AI co-creation [12]. The emergence of the “AI Ghostwriter Effect,” where users psychologically disown AI-generated text while publicly claiming authorship, highlights a dangerous disconnect between perceived ownership and actual creative input [60]. This phenomenon is exacerbated by interface designs that obscure the degree of machine initiative, leading to a “dearth of the author” where expressive intent is diluted by automated generation [144].

Legally, the landscape remains fraught with ambiguity. Current copyright frameworks generally deny authorship to non-human entities, creating a precarious status for discoveries where the generative agent serves as the primary intellectual architect [147]. While some argue that high-end editing services have long blurred these lines, the autonomous capacity of modern agents to synthesize novel research questions pushes existing guidelines to their breaking point [61]. The inability of current legal structures to accommodate “mixed-initiative” systems, where both human and AI can initiate content changes, necessitates a re-evaluation of intellectual property rights to prevent the erosion of accountability [25]. Furthermore, the rise of fully autonomous agents capable of executing end-to-end research cycles—from idea generation to paper writing—challenges the very notion of human oversight required for patentability and credit allocation [84].

To address these complexities, the community is shifting toward process-oriented attribution models

rather than product-based detection. Frameworks such as the Provenance, Authoring and Versioning (PAV) ontology offer a structured approach to distinguishing between contributors, authors, and curators, enabling granular tracking of agent roles in digital artifacts [13]. Complementing this, standardized reporting mechanisms like “AI Usage Cards” propose a three-dimensional model of transparency, integrity, and accountability, allowing researchers to explicitly declare the scope of AI involvement in ideation, drafting, and analysis [55]. These tools aim to replace the futile arms race of binary detection with a culture of responsible disclosure, acknowledging that “mixtext”—a hybrid of human and machine generation—is becoming the norm [58].

Looking forward, the concept of scholarship itself may evolve toward “late-binding” knowledge artifacts, where content is dynamically rendered by AI models at the point of consumption rather than existing as a static record. This paradigm shift demands new governance structures that prioritize verifiable provenance chains over fixed bylines, ensuring that the lineage of ideas remains transparent even as the medium of expression becomes fluid [123]. Ultimately, redefining authorship requires a symbiotic framework that recognizes algorithmic agency without abdicating human responsibility, fostering an ecosystem where credit is assigned based on intellectual stewardship rather than mere textual production.

6.2 Detection Methodologies and the Epistemology of Synthetic Scientific Text

While the previous section establishes the necessity of process-oriented attribution to redefine authorship, the practical enforcement of these standards relies heavily on the capacity to distinguish between human and machine contributions. Consequently, the detection of AI-generated and AI-modified scientific text has evolved from a binary classification problem into a complex epistemological challenge, necessitating a shift from punitive policing to nuanced verification frameworks. Early methodologies relied heavily on statistical stylometry and perplexity metrics to distinguish human authorship from machine generation. However, the efficacy of these binary detectors has degraded significantly as generative models achieve fluency indistinguishable from human writing. Empirical evaluations reveal that current detection tools often suffer from high false-positive rates, particularly penalizing non-native English speakers, and demonstrate poor robustness against simple obfuscation techniques or mixed-initiative editing workflows [57; 148]. Furthermore, the prevalence of “mixtext”—documents where human-written content is iteratively refined by Large Language Models (LLMs)—renders traditional detectors ineffective, as they struggle to identify subtle semantic shifts or style adaptations introduced by human-AI collaboration [58; 149].

Beyond mere authorship attribution, the integrity of scientific discourse is increasingly threatened by specific epistemological artifacts inherent to synthetic text, which complicate the verification of the provenance chains discussed earlier. A critical phenomenon is the emergence of “tortured phrases,” where standard technical terminology is replaced by bizarre synonyms (e.g., “counterfeit consciousness” for “artificial intelligence”) to evade plagiarism filters or optimize for specific algorithmic rewards [34]. These linguistic anomalies often co-occur with hallucinated citations and shallow knowledge synthesis, where LLMs generate plausible-sounding but factually unsupported claims or reference non-existent literature [62; 45]. While advanced frameworks like CheckGPT attempt to capture deep semantic patterns to identify such content, the arms race between generation and detection suggests that purely post-hoc technical solutions are insufficient for guaranteeing scientific rigor [150].

Recognizing these limitations, the field is pivoting toward transparency-oriented governance mechanisms that prioritize process integrity over product detection, directly supporting the attribution

models proposed in the preceding section. Innovative frameworks such as “AI Usage Cards” and “PaperCards” propose standardized reporting structures where authors explicitly disclose the scope, prompts, and extent of AI involvement in their research lifecycle [55; 151]. These mechanisms align with broader calls for “scienceography,” which advocates for analyzing the source code and revision history of manuscripts to understand the collaborative dynamics between human and machine agents, rather than relying solely on the final static output [152]. Moreover, emerging methods like Raidar leverage the tendency of LLMs to modify human text less than AI text during rewriting tasks, offering a more robust, model-agnostic detection signal that complements disclosure protocols [153].

Future directions must integrate these technical and procedural approaches into a unified epistemology of synthetic science, laying the groundwork for the broader ecosystem integrity discussed in the following section. This involves developing mixed-initiative detection workflows that combine automated flagging of structural anomalies—such as citation graph inconsistencies or uncertainty cue deviations—with human expert verification [154; 155]. Ultimately, preserving scientific integrity requires moving beyond the futile quest for perfect detectors toward a culture of accountable transparency, where the provenance of ideas is as rigorously tracked as the data supporting them, ensuring that the biases and evaluation challenges inherent in peer review do not undermine the trust established through transparent authorship.

6.3 Algorithmic Bias, Fairness, and the Integrity of Peer Review Systems

The integration of artificial intelligence into peer review ecosystems introduces complex vectors for bias propagation that threaten the epistemic integrity of scientific publishing. While algorithmic systems promise efficiency in managing the exponential growth of submissions, they risk codifying and amplifying existing societal and institutional inequities. A primary concern lies in literature synthesis and reviewer matching, where recommender systems often rely on historical citation data that inherently favors established paradigms and dominant institutions. [156] demonstrate that such systems can obscure dissenting viewpoints and reinforce “rich-get-richer” dynamics, effectively marginalizing novel or interdisciplinary research that deviates from trained distributions. Furthermore, the automation of reviewer assignment is susceptible to adversarial manipulation; [157] illustrate how authors can subtly modify manuscript text to game topic models, selectively inviting sympathetic reviewers and compromising the impartiality of the evaluation process.

The deployment of Large Language Models (LLMs) as active participants in the review loop presents acute challenges regarding score validity and fairness. Empirical evidence suggests a systematic bias toward leniency in AI-assisted reviews. [52] report that submissions receiving AI-generated or AI-modified reviews are significantly more likely to achieve higher scores and acceptance, with a relative probability increase of 14.4% compared to purely human evaluations. This score inflation creates an uneven playing field, potentially rewarding those who leverage generative tools for review drafting over those who adhere to traditional, rigorous critique. Moreover, [51] reveals that algorithmic simulations of review dynamics can exhibit substantial variance in decision outcomes driven by latent biases such as authority bias and altruism fatigue, suggesting that AI agents may inadvertently replicate the sociological flaws of human committees without the mitigating factor of conscious ethical reflection.

Beyond score distortion, there is a profound risk of cognitive homogenization. As researchers increasingly utilize AI tools like [53] to structure their critiques or [49] to generate feedback, the diversity of critical perspectives may converge toward a standardized, algorithmic average. This convergence

threatens the pluralism essential for scientific progress, as unique, idiosyncratic insights—often the source of paradigm shifts—are smoothed over by models trained on consensus. Additionally, the opacity of these systems complicates accountability; while [54] estimates that up to 16.9% of review text in major conferences may be LLM-generated, the inability to distinguish between constructive AI assistance and wholesale delegation undermines trust in the review record.

Addressing these challenges requires a shift from purely accuracy-driven optimization to fairness-aware algorithmic design. Future frameworks must incorporate adversarial testing to detect manipulation in reviewer matching and develop calibration mechanisms that correct for the inherent positivity bias of generative models. Crucially, governance policies must mandate transparency not only in authorship but in the reviewing process itself, ensuring that the “black box” of algorithmic evaluation does not become a vehicle for reinforcing institutional stratification. Without rigorous intervention, the automation of peer review risks transforming from a tool for scaling scientific quality control into a mechanism that systematically entrenches bias and stifles the diversity of thought necessary for discovery.

6.4 Governance Frameworks for Autonomous Discovery and Research Safety

Transitioning from the risks of bias and homogenization in semi-automated peer review, the advent of fully autonomous scientific agents and “self-driving laboratories” marks a definitive paradigm shift from human-in-the-loop assistance to open-ended, machine-directed discovery. While systems like [84] and [22] demonstrate the capacity for independent hypothesis generation, experimental execution, and paper writing, this autonomy introduces profound safety and governance challenges that traditional ethical frameworks are ill-equipped to address. The primary concern lies in the potential for dual-use misuse, where autonomous agents could inadvertently or maliciously design harmful substances or bypass regulatory controls. Recent evaluations indicate that large language models can already provide detailed protocols for generating pandemic-class pathogens to non-experts, highlighting an urgent need for robust safeguarding mechanisms [158]. To mitigate these risks, emerging governance structures propose a triadic framework involving human regulation, agent alignment, and environmental feedback loops [102].

A critical component of this governance is the implementation of proactive safety filters specifically designed for scientific domains. The [159] study introduces “SciGuard,” a system designed to prevent AI models from generating harmful chemical synthesis routes or circumventing safety regulations without compromising their utility in benign research contexts. Complementary to this is the development of red-teaming benchmarks like “SciMT-Safety,” which rigorously assess the vulnerability of autonomous agents to jailbreaking attempts aimed at eliciting dangerous experimental designs [159]. Furthermore, the integration of such safety protocols must be intrinsic to the agent’s architecture rather than an external add-on, ensuring that risk assessment occurs dynamically during the planning and execution phases of autonomous workflows [102].

Beyond physical safety, the integrity of the scientific record demands rigorous computational reproducibility standards for AI-driven discovery, extending the provenance concerns raised in earlier discussions of authorship to the entirety of the research lifecycle. The complexity of autonomous workflows, which often involve iterative code generation, tool usage, and data analysis, creates significant barriers to verification. Frameworks such as AiiDA and Globus Automation Services offer foundational infrastructure for tracking data provenance and managing complex workflows across distributed resources, yet they must be adapted to handle the stochastic nature of generative agents [119; 160]. The [123] initiative underscores this by demonstrating a “data-to-paper”

platform that programmatically back-traces information flow, ensuring that every claim in an AI-generated manuscript is linked to its underlying data and methodological steps. However, as goal complexity increases, fully autonomous cycles still struggle with accuracy, necessitating a “human co-pilot” model for critical validation stages [123].

Governance must also address the epistemological risks of open-ended discovery loops, where agents might optimize for metrics that do not align with scientific truth or human values, thereby exacerbating the incentive distortions discussed in the context of autonomous publishing. The [161] paper argues for a vision of augmented AI that incorporates diverse approaches to causality analysis to prevent biased searches across the hypothesis space. Moreover, the deployment of virtual environments like [162] provides a safe sandbox for benchmarking agent behaviors and testing safety protocols before real-world deployment. Ultimately, effective governance requires a synthesis of technical safeguards, standardized reproducibility benchmarks [163], and evolving policy frameworks that mandate transparency and human accountability. As autonomous agents transition from theoretical prototypes to active participants in the scientific ecosystem, the community must prioritize safety and verifiability to ensure that the acceleration of discovery does not come at the cost of scientific integrity or public safety, setting the stage for the broader policy harmonization and incentive realignment required in the subsequent section.

6.5 Evolving Policy Landscapes and the Future of Academic Incentive Structures

The integration of generative artificial intelligence into scientific workflows has precipitated a fragmented policy landscape, characterized by a “Babel Tower Effect” of disparate journal guidelines and institutional mandates that risk confusing rather than standardizing ethical practice [164]. Current governance mechanisms largely rely on disclosure frameworks, such as AI Usage Cards, which aim to codify transparency, integrity, and accountability as the triad of responsible AI deployment [55]. However, the efficacy of these self-reporting mechanisms is contested by empirical evidence suggesting that geographic restrictions on AI access are ineffective and that detection algorithms struggle to distinguish between human and machine contributions with high fidelity, particularly for non-native speakers [165; 34]. Consequently, the field is shifting from punitive detection models toward process-oriented verification, where the provenance of intellectual contribution is tracked through iterative workflow logs rather than static text analysis.

This regulatory evolution is inextricably linked to the crisis in academic incentive structures. Traditional metrics, which prioritize publication volume and citation counts, are increasingly misaligned with an era where AI can exponentially accelerate manuscript production, potentially exacerbating the “strain on scientific publishing” without a commensurate increase in scientific validity [166]. The advent of autonomous agents capable of generating full research cycles—from hypothesis formulation to paper drafting—challenges the foundational assumptions of authorship and merit [84]. If AI systems can produce publishable units at near-zero marginal cost, incentive systems based on output quantity risk inducing severe score inflation and homogenization of scientific discourse, as evidenced by studies showing that AI-assisted peer reviews often yield higher scores and increased acceptance probabilities [52].

To mitigate these distortions, future incentive frameworks must transition from evaluating output volume to assessing the quality of human-AI collaborative synergy and the rigor of human oversight. This requires new evaluation metrics that quantify the distinct intellectual value added by human researchers, such as the novelty of conceptual framing, the ethical robustness of experimental design, and the critical synthesis of AI-generated evidence. Preliminary attempts to automate quality

assessment using machine learning on bibliometric and metadata inputs show promise but reveal significant disciplinary variances, indicating that a one-size-fits-all algorithmic evaluation is insufficient [167]. Furthermore, the tension between efficiency and integrity necessitates a re-evaluation of peer review itself; while AI can scaffold the review process for novices and reduce turnaround times, it introduces risks of bias propagation and “automation complacency” that may undermine the gatekeeping function of science [53; 54].

Ultimately, the future of scientific integrity depends on harmonizing global policies to prevent regulatory arbitrage while redesigning reward systems that prioritize “slow science” virtues—reproducibility, deep critical thinking, and ethical stewardship—over velocity. As AI tools become ubiquitous, the academic community must foster a culture where the responsible integration of these technologies is rewarded, ensuring that the acceleration of discovery does not come at the expense of the epistemic foundations of research [168]. This paradigm shift requires not only technical solutions for transparency but also a sociological realignment of how scientific contribution is defined, measured, and valued in an augmented intelligence ecosystem.

7 Conclusion

This subsection synthesizes the paradigmatic shift from static computational aids to dynamic, symbiotic partnerships in scientific inquiry, underscoring that the integration of artificial intelligence represents not merely an acceleration of existing workflows but a fundamental reconfiguration of the epistemic landscape. The survey reveals a trajectory where AI has evolved from passive tooling for data management to active agents capable of hypothesis formulation, experimental design, and autonomous execution [84]. While early computational support focused on isolated tasks, contemporary systems demonstrate the capacity to navigate the entire research lifecycle, from ideation to peer review simulation, effectively functioning as “cognitive microscopes” that extend human intellectual reach [169]. However, this transition is characterized by a critical tension between the promise of accelerated discovery and the imperative of maintaining rigorous human oversight.

A central insight emerging from this analysis is the nuanced nature of human-AI complementarity, which relies heavily on appropriate trust calibration and the alignment of mental models. Empirical evidence suggests that while AI can significantly augment decision-making, particularly in domains requiring the synthesis of vast literature, the benefits are contingent upon the user’s expertise level and the system’s ability to communicate uncertainty [4]. The phenomenon of “automation complacency” poses a substantial risk, where over-reliance on generative models may lead to an “illusion of competence,” potentially eroding foundational scientific skills among early-career researchers [88]. Furthermore, the deployment of AI in peer review and writing workflows has introduced complex ethical challenges, including the proliferation of “tortured phrases,” hallucinated citations, and the potential for algorithmic bias to homogenize scientific discourse [34; 52].

Despite these challenges, the potential for AI to drive interdisciplinary breakthroughs remains profound. Systems leveraging large-scale knowledge graphs and multi-agent reasoning have demonstrated the ability to uncover hidden connections across disparate fields, proposing novel hypotheses that might elude human intuition alone [22]. Yet, the current state of the field is marked by a significant gap in standardized evaluation metrics and governance frameworks. Existing benchmarks often fail to capture the complexity of open-ended scientific tasks, and detection methods for AI-generated content remain unreliable, necessitating a shift from product-based detection to process-oriented transparency mechanisms such as “AI Usage Cards” [55].

Looking forward, the strategic agenda for research must prioritize the development of human-centric interfaces that foster true collaborativity rather than mere automation. This involves designing systems that support bidirectional communication, explainable reasoning, and dynamic role allocation, ensuring that humans retain intellectual ownership and accountability [35]. Future efforts should also focus on creating robust safety protocols for autonomous agents to prevent misuse in sensitive domains, alongside establishing global policy harmonization to address the evolving landscape of authorship and intellectual property [102]. Ultimately, the future of augmented science depends on our ability to cultivate a symbiotic ecosystem where AI serves as a catalyst for human creativity, guided by unwavering ethical standards and a commitment to the integrity of the scientific record [104].

References

- [1] A Computational Inflection for Scientific Discovery
- [2] Artificial Intelligence for Literature Reviews Opportunities and Challenges
- [3] Learning to Complement Humans
- [4] Advancing Human-AI Complementarity The Impact of User Expertise and Algorithmic Tuning on Joint Decision Making
- [5] Shaping Human-AI Collaboration Varied Scaffolding Levels in Co-writing with Language Models
- [6] ChatGPT is a Remarkable Tool – For Experts
- [7] The Value, Benefits, and Concerns of Generative AI-Powered Assistance in Writing
- [8] Automated Paper Screening for Clinical Reviews Using Large Language Models
- [9] MLR-Copilot: Autonomous Machine Learning Research based on Large Language Models Agents
- [10] Relatedly Scaffolding Literature Reviews with Existing Related Work Sections
- [11] Empowering Biomedical Discovery with AI Agents
- [12] Patterns of authors contribution in scientific manuscripts
- [13] PAV ontology Provenance, Authoring and Versioning
- [14] Argument Mining for Understanding Peer Reviews
- [15] SciMON Scientific Inspiration Machines Optimized for Novelty
- [16] CoQuest Exploring Research Question Co-Creation with an LLM-based Agent
- [17] “It’s like a rubber duck that talks back”: Understanding Generative AI-Assisted Data Analysis Workflows through a Participatory Prompting Study
- [18] Assisting in Writing Wikipedia-like Articles From Scratch with Large Language Models
- [19] Theoretical And Technological Building Blocks For An Innovation Accelerator
- [20] PaperRobot Incremental Draft Generation of Scientific Ideas
- [21] AGATHA Automatic Graph-mining And Transformer based Hypothesis generation Approach

- [22] SciAgents: Automating scientific discovery through multi-agent intelligent graph reasoning
- [23] Can LLMs Generate Novel Research Ideas? A Large-Scale Human Study with 100+ NLP Researchers
- [24] Bursting Scientific Filter Bubbles Boosting Innovation via Novel Author Discovery
- [25] Beyond Prompts Exploring the Design Space of Mixed-Initiative Co-Creativity Systems
- [26] Read, Revise, Repeat A System Demonstration for Human-in-the-loop Iterative Text Revision
- [27] Learning to Model Editing Processes
- [28] Understanding Iterative Revision from Human-Written Text
- [29] Suggestion Lists vs. Continuous Generation Interaction Design for Writing with Generative Models on Mobile Devices Affect Text Length, Wording and Perceived Authorship
- [30] GhostWriter Augmenting Collaborative Human-AI Writing Experiences Through Personalization and Agency
- [31] Does Writing with Language Models Reduce Content Diversity
- [32] AI vs. Human – Differentiation Analysis of Scientific Content Generation
- [33] Examining Scientific Writing Styles from the Perspective of Linguistic Complexity
- [34] Tortured phrases A dubious writing style emerging in science. Evidence of critical issues affecting established journals
- [35] ConvXAI Delivering Heterogeneous AI Explanations via Conversations to Support Human-AI Scientific Writing
- [36] LitLLM A Toolkit for Scientific Literature Review
- [37] PubTator 3.0 an AI-powered Literature Resource for Unlocking Biomedical Knowledge
- [38] Scientific Claim Verification with VERT5ERINI
- [39] SciFact-Open Towards open-domain scientific claim verification
- [40] SCITAB A Challenging Benchmark for Compositional Reasoning and Claim Verification on Scientific Tables
- [41] CitationIE Leveraging the Citation Graph for Scientific Information Extraction
- [42] Neural Multi-Task Learning for Citation Function and Provenance
- [43] REASONS: A benchmark for REtrieval and Automated citationS Of scieNtific Sentences using Public and Proprietary LLMs
- [44] CiteME: Can Language Models Accurately Cite Scientific Claims?
- [45] Shallow Synthesis of Knowledge in GPT-Generated Texts A Case Study in Automatic Related Work Composition
- [46] CLIMATE-FEVER A Dataset for Verification of Real-World Climate Claims
- [47] Can We Automate Scientific Reviewing
- [48] Reviewer2 Optimizing Review Generation Through Prompt Generation

- [49] MARG Multi-Agent Review Generation for Scientific Papers
- [50] ARIES A Corpus of Scientific Paper Edits Made in Response to Peer Reviews
- [51] AgentReview: Exploring Peer Review Dynamics with LLM Agents
- [52] The AI Review Lottery: Widespread AI-Assisted Peer Reviews Boost Paper Scores and Acceptance Rates
- [53] ReviewFlow Intelligent Scaffolding to Support Academic Peer Reviewing
- [54] Monitoring AI-Modified Content at Scale A Case Study on the Impact of ChatGPT on AI Conference Peer Reviews
- [55] AI Usage Cards Responsibly Reporting AI-generated Content
- [56] Towards the relationship between AIGC in manuscript writing and author profiles evidence from preprints in LLMs
- [57] Testing of Detection Tools for AI-Generated Text
- [58] LLM-as-a-Coauthor Can Mixed Human-Written and Machine-Generated Text Be Detected
- [59] Does Human Collaboration Enhance the Accuracy of Identifying LLM-Generated Deepfake Texts
- [60] The AI Ghostwriter Effect When Users Do Not Perceive Ownership of AI-Generated Text But Self-Declare as Authors
- [61] Ethics of using language editing services in an era of digital communication and heavily multi-authored papers
- [62] ChatGPT as Research Scientist: Probing GPT’s Capabilities as a Research Librarian, Research Ethicist, Data Generator and Data Predictor
- [63] Artificial intelligence to automate the systematic review of scientific literature
- [64] System for systematic literature review using multiple AI agents Concept and an empirical evaluation
- [65] Large Language Models for Scientific Information Extraction An Empirical Study for Virology
- [66] Multi-Task Identification of Entities, Relations, and Coreference for Scientific Knowledge Graph Construction
- [67] Microtask crowdsourcing for disease mention annotation in PubMed abstracts
- [68] Investigating Software Usage in the Social Sciences A Knowledge Graph Approach
- [69] A large dataset of software mentions in the biomedical literature
- [70] Can large language models replace humans in the systematic review process Evaluating GPT-4’s efficacy in screening and extracting data from peer-reviewed and grey literature in multiple languages
- [71] Automating Research Synthesis with Domain-Specific Large Language Model Fine-Tuning
- [72] PaperQA Retrieval-Augmented Generative Agent for Scientific Research
- [73] Automated Fact-Checking of Climate Change Claims with Large Language Models

- [74] ResearchAgent Iterative Research Idea Generation over Scientific Literature with Large Language Models
- [75] Can ChatGPT be used to generate scientific hypotheses
- [76] SciAgent Tool-augmented Language Models for Scientific Reasoning
- [77] ORKG-Leaderboards A Systematic Workflow for Mining Leaderboards as a Knowledge Graph
- [78] Towards Reasoning in Large Language Models via Multi-Agent Peer Review Collaboration
- [79] The Materials Science Procedural Text Corpus Annotating Materials Synthesis Procedures with Shallow Semantic Structures
- [80] Expressing High-Level Scientific Claims with Formal Semantics
- [81] Underspecified Scientific Claims in Nanopublications
- [82] Scientific Language Models for Biomedical Knowledge Base Completion An Empirical Study
- [83] GFlowNets for AI-Driven Scientific Discovery
- [84] The AI Scientist: Towards Fully Automated Open-Ended Scientific Discovery
- [85] From human experts to machines An LLM supported approach to ontology and knowledge graph construction
- [86] Natural Language to Code Generation in Interactive Data Science Notebooks
- [87] The Impact of Large Language Models on Open-source Innovation: Evidence from GitHub Copilot
- [88] Ironies of Generative AI Understanding and mitigating productivity loss in human-AI interactions
- [89] AI Chains Transparent and Controllable Human-AI Interaction by Chaining Large Language Model Prompts
- [90] MatPlotAgent Method and Evaluation for LLM-Based Agentic Scientific Data Visualization
- [91] AutoFlow: Automated Workflow Generation for Large Language Model Agents
- [92] Iris A Conversational Agent for Complex Tasks
- [93] Spider2-V: How Far Are Multimodal Agents From Automating Data Science and Engineering Workflows?
- [94] How Do Analysts Understand and Verify AI-Assisted Data Analyses
- [95] LLMs for Science Usage for Code Generation and Data Analysis
- [96] Large Language Models are Zero Shot Hypothesis Proposers
- [97] An Artificial Intelligence (AI) workflow for catalyst design and optimization
- [98] BioDiscoveryAgent: An AI Agent for Designing Genetic Perturbation Experiments
- [99] Emergent autonomous scientific research capabilities of large language models
- [100] A GPT-4 Reticular Chemist for Guiding MOF Discovery

- [101] ORGANA A Robotic Assistant for Automated Chemistry Experimentation and Characterization
- [102] Prioritizing Safeguarding Over Autonomy Risks of LLM Agents for Science
- [103] Human-in-the-loop The future of Machine Learning in Automated Electron Microscopy
- [104] Accelerating science with human-aware artificial intelligence
- [105] SciFIBench: Benchmarking Large Multimodal Models for Scientific Figure Interpretation
- [106] Multimodal ArXiv A Dataset for Improving Scientific Comprehension of Large Vision-Language Models
- [107] SCITUNE Aligning Large Language Models with Scientific Multimodal Instructions
- [108] Data-Copilot Bridging Billions of Data and Humans with Autonomous Workflow
- [109] DynaVis Dynamically Synthesized UI Widgets for Visualization Editing
- [110] Towards Natural Language-Based Visualization Authoring
- [111] InkSight Leveraging Sketch Interaction for Documenting Chart Findings in Computational Notebooks
- [112] AutoAIViz Opening the Blackbox of Automated Artificial Intelligence with Conditional Parallel Coordinates
- [113] AutoRG-Brain: Grounded Report Generation for Brain MRI
- [114] SciGraphQA A Large-Scale Synthetic Multi-Turn Question-Answering Dataset for Scientific Graphs
- [115] SPIQA: A Dataset for Multimodal Question Answering on Scientific Papers
- [116] Telling Stories from Computational Notebooks AI-Assisted Presentation Slides Creation for Presenting Data Science Work
- [117] DOC2PPT Automatic Presentation Slides Generation from Scientific Documents
- [118] Sustainable computational science the ReScience initiative
- [119] AiiDA 1.0, a scalable computational infrastructure for automated reproducible workflows and data provenance
- [120] YesWorkflow A User-Oriented, Language-Independent Tool for Recovering Workflow Information from Scripts
- [121] Computational reproducibility of Jupyter notebooks from biomedical publications
- [122] DataDreamer A Tool for Synthetic Data Generation and Reproducible LLM Workflows
- [123] Autonomous LLM-driven research from data to human-verifiable research papers
- [124] Computing Environments for Reproducibility Capturing the Whole Tale
- [125] CORE-Bench: Fostering the Credibility of Published Research Through a Computational Reproducibility Agent Benchmark
- [126] Co-audit tools to help humans double-check AI-generated content

- [127] Computational and informatics advances for reproducible data analysis in neuroimaging
- [128] CoAuthor Designing a Human-AI Collaborative Writing Dataset for Exploring Language Model Capabilities
- [129] Automated Metrics for Medical Multi-Document Summarization Disagree with Human Evaluations
- [130] SciBench Evaluating College-Level Scientific Problem-Solving Abilities of Large Language Models
- [131] OlympicArena: Benchmarking Multi-discipline Cognitive Reasoning for Superintelligent AI
- [132] LAB-Bench: Measuring Capabilities of Language Models for Biology Research
- [133] Parachute Evaluating Interactive Human-LM Co-writing Systems
- [134] Take It, Leave It, or Fix It Measuring Productivity and Trust in Human-AI Collaboration
- [135] How Do Data Analysts Respond to AI Assistance A Wizard-of-Oz Study
- [136] Human-AI Collaboration for UX Evaluation Effects of Explanation and Synchronization
- [137] Who Goes First Influences of Human-AI Workflow on Decision Making in Clinical Imaging
- [138] How Much Automation Does a Data Scientist Want
- [139] Fits and Starts Enterprise Use of AutoML and the Role of Humans in the Loop
- [140] Improving Human-AI Collaboration With Descriptions of AI Behavior
- [141] Interactive Model Cards A Human-Centered Approach to Model Documentation
- [142] AutoDS Towards Human-Centered Automation of Data Science
- [143] AI Song Contest Human-AI Co-Creation in Songwriting
- [144] The Dearth of the Author in AI-Supported Writing
- [145] Human-AI Collaboration via Conditional Delegation A Case Study of Content Moderation
- [146] Evaluating Human-Language Model Interaction
- [147] ChatGPT and a New Academic Reality Artificial Intelligence-Written Research Papers and the Ethics of the Large Language Models in Scholarly Publishing
- [148] Perception, performance, and detectability of conversational artificial intelligence across 32 university courses
- [149] Towards an Understanding and Explanation for Mixed-Initiative Artificial Scientific Text Detection
- [150] On the Detectability of ChatGPT Content Benchmarking, Methodology, and Evaluation through the Lens of Academic Writing
- [151] Techniques for supercharging academic writing with generative AI
- [152] Scienceography the study of how science is written
- [153] Raidar geneRative AI Detection viA Rewriting

- [154] A Scalable and Adaptive Method for Finding Semantically Equivalent Cue Words of Uncertainty
- [155] Structural Scaffolds for Citation Intent Classification in Scientific Publications
- [156] Biases in Scholarly Recommender Systems Impact, Prevalence, and Mitigation
- [157] No more Reviewer 2 Subverting Automatic Paper-Reviewer Assignment using Adversarial Learning
- [158] Can large language models democratize access to dual-use biotechnology
- [159] Control Risk for Potential Misuse of Artificial Intelligence in Science
- [160] Globus Automation Services Research process automation across the space-time continuum
- [161] The Future of Fundamental Science Led by Generative Closed-Loop Artificial Intelligence
- [162] DISCOVERYWORLD: A Virtual Environment for Developing and Evaluating Automated Scientific Discovery Agents
- [163] BLADE: Benchmarking Language Model Agents for Data-Driven Science
- [164] Development of the ChatGPT, Generative Artificial Intelligence and Natural Large Language Models for Accountable Reporting and Use (CANGARU) Guidelines
- [165] Where there's a will there's a way: ChatGPT is used more for science in countries where it is prohibited
- [166] The strain on scientific publishing
- [167] Predicting article quality scores with machine learning The UK Research Excellence Framework
- [168] AI speeds up science, but scientists should guide its direction
- [169] On scientific understanding with artificial intelligence