

Synthetic 3D Data Generation: Paradigms, Methodologies, and Future Directions

APRIL Lab Team

APRIL Lab, Zhejiang University

Synthetic three-dimensional data serves as a critical resource for training embodied agents and autonomous systems where real-world acquisition remains costly or dangerous. Recent advancements have transitioned the field from rigid, rule-based procedural pipelines to flexible neural generative models capable of producing high-fidelity assets from sparse inputs. Despite these gains, a persistent disconnect exists between visual realism and physical plausibility, often yielding geometrically inconsistent structures unsuitable for rigorous simulation. This survey presents a unified taxonomy categorizing generation methodologies into four distinct paradigms: procedural algorithmic synthesis, direct neural 3D modeling, two-dimensional prior-guided lifting, and language-driven compositional workflows. Analysis reveals that current evaluation protocols disproportionately favor perceptual quality over topological validity, creating a bottleneck for deploying generated data in physics-grounded tasks. Furthermore, emerging hybrid architectures suggest that integrating explicit geometric constraints with latent diffusion processes offers the most viable path toward simulation-ready content. Future progress demands a fundamental shift from purely visual metrics to standardized benchmarks assessing semantic alignment, dynamic interaction, and structural integrity.

Keywords: Synthetic 3D Data Generation, Neural Generative Modeling, Procedural Synthesis, 2D-to-3D Lifting, Embodied AI Simulation

 **Date:** August 8, 2026



1 Introduction

The scarcity of high-quality annotated three-dimensional data has driven a fundamental transition from manual asset creation to automated algorithmic and neural synthesis [79, 137]. Contemporary approaches now span a broad spectrum, ranging from rule-based procedural systems to foundation model-driven architectures capable of constructing complex scenes and dynamic environments [79, 168]. This technological evolution specifically targets the critical gaps between visual fidelity, physical plausibility, and semantic controllability required for training embodied AI and autonomous systems [137, 168].

The core methodologies can generally be categorized into several primary paradigms. 1) Native 3D generation utilizes direct supervision on volumetric or implicit representations to synthesize geometry without intermediate projections [74, 87, 153]. 2) 2D prior-guided lifting transfers knowledge from pre-trained image models into three-dimensional space through score distillation sampling, effectively mitigating data scarcity [74, 87]. 3) Hybrid architec-

tures integrate multi-view consistency constraints with procedural rules to enforce geometric validity while managing scene complexity [87, 153]. This classification spans static objects, dynamic humans, and unbounded scenes, offering a unified reference for navigating diverse representation formats and generative pipelines.

A fundamental tension persists between the demand for high-fidelity visual synthesis and the severe computational constraints inherent in current generative architectures [139, 154]. This adversarial dynamic intensifies as increasing scene complexity and semantic controllability directly exacerbate training instability and mode collapse within latent space optimization. Systematic analyses further reveal that rigid workflows and limited dataset diversity create barriers to generalization, preventing models from adapting to unseen or dynamically evolving environments.

Existing systematic reviews frequently fragment the field by isolating specific generative paradigms or representation types, failing to capture the converging trends between procedural rules and neural priors [10, 71]. This taxonomic separation obscures a uni-

fied framework capable of simultaneously addressing the trade-offs between geometric fidelity, semantic controllability, and physical plausibility across diverse synthesis methods. Furthermore, the absence of standardized metrics for assessing simulation readiness and downstream utility hinders the reliable deployment of generated assets in embodied AI and autonomous systems. Therefore, the primary motivation of this survey is to provide a unified taxonomy that bridges these gaps, serving as a foundational reference for advancing content creation pipelines and machine learning applications.

Contributions. In this survey, we make three key contributions: **1) Unified Taxonomy:** We establish a cohesive framework that bridges the gap between visual fidelity, physical plausibility, and semantic controllability across diverse generation paradigms. **2) Methodological Categorization:** The work provides a comprehensive classification of techniques into four distinct conceptual pillars, analyzing their specific architectural choices and optimization strategies as detailed in [Sec. 3](#). **3) Critical Synthesis and Roadmap:** We critically examine persistent limitations regarding geometric consistency and evaluation standards while mapping concrete future directions for embodied AI integration in [Sec. 4](#).

Scope. To maintain a rigorous focus on automated synthetic data creation, this survey delineates its boundaries around four primary methodological paradigms: procedural algorithmic synthesis, neural 3D generative modeling, 2D prior-guided lifting, and language-driven composition, as systematically categorized in [Sec. 3](#). We explicitly exclude manual artistic modeling workflows and non-generative acquisition pipelines such as standard photogrammetry, thereby concentrating exclusively on techniques capable of programmatic or learned generation without direct human geometric intervention. Furthermore, the analysis prioritizes methods designed to support downstream machine learning applications, including embodied AI training and autonomous system simulation, rather than approaches optimized solely for visual entertainment or static asset storage. This selection criterion ensures that the reviewed techniques address critical requirements for semantic controllability and physical plausibility necessary for robust model training.

Survey Pipeline. Following the foundational context established in [Sec. 2](#), which formalizes problem definitions and traces technological evolution, the core of this survey presents a systematic deconstruction of generation methodologies in [Sec. 3](#). This central segment categorizes approaches into four distinct

paradigms: procedural algorithmic synthesis, neural 3D generative modeling, 2D prior-guided lifting, and language-driven composition, offering a granular analysis of their respective architectural mechanisms. The narrative culminates in [Sec. 4](#), where we critically examine persistent systemic limitations regarding geometric consistency and physical plausibility before outlining a forward-looking research agenda for scalable, simulation-ready data generation.

2 Background

This chapter establishes the foundational context for 3D dataset generation by defining core problem formulations, tracing technological evolution, and outlining evaluation protocols. [Sec. 2.1](#) formalizes the mathematical frameworks for mapping latent spaces to diverse 3D representations while delineating objectives for fidelity and annotation accuracy. The historical trajectory from procedural content generation to modern neural generative models is examined in [Sec. 2.2](#), highlighting key shifts toward data-driven and physics-grounded approaches. [Sec. 2.3](#) summarizes existing benchmarks and metrics, noting current gaps in standardized measures for physical plausibility and semantic consistency. Finally, [Sec. 2.4](#) maps adjacent research domains such as computer graphics and robotics simulation to identify critical overlaps in representation learning and multimodal alignment.

2.1 Problem Definition

Generative Mapping and Equivariant Constraints.

Formalizing the synthesis of 3D geometry requires defining a stochastic mapping that transforms a latent code z and a reference set $\mathcal{R}$ into an output set of primitives $\mathcal{S}$, where the cardinality and ordering of elements must remain invariant to input permutations [158]. This structural requirement is governed by the condition that the generator function satisfies (F, l) -equivariance, decomposing the generation process into distinct creation and update phases to handle unordered data representations:

$$\begin{aligned} \mathcal{S} &= f_{\theta}(z, \mathcal{R}) \\ \text{s.t. } \pi(\mathcal{S}) &= f_{\theta}(z, \pi(\mathcal{R})) \end{aligned} \quad (1)$$

where π denotes a permutation operator acting on the reference set. To address the failure modes of standard independent and identically distributed sampling, which often yield structurally incoherent geometries, the framework replaces random selection with a dynamic mechanism that maximizes the cosine similarity between latent projections and learnable angular parameters θ [158]. This approach employs

a hard selection strategy via argsort operations combined with FiLM modulation to ensure gradient propagation while strictly conditioning the deterministic choice of relevant reference points on the specific latent code. Consequently, theoretical constraints on exchangeability are relaxed; the system maintains stability during Evidence Lower Bound (ELBO) optimization by ensuring parameter updates remain independent of the output symmetry, thereby enhancing extrapolation capabilities for complex topological structures without enforcing rigid output symmetries [158].

Optimization Objectives and Structural Priors.

The generative process is governed by a composite objective function balancing geometric reconstruction accuracy against distribution entropy, formally expressed as:

$$\begin{aligned}\mathcal{L}_{total} = & \lambda_{fid}\mathcal{L}_{rec} \\ & - \lambda_{div}\mathcal{H}(p(z)) \\ & + \Omega(\Phi_{prog}),\end{aligned}\quad (2)$$

where $\mathcal{L}_{rec}$ quantifies fidelity between the generated point cloud $\mathcal{P}_{gen}$ and ground truth primitives $\mathcal{P}_{gt}$. This term utilizes metrics such as the L_2 Chamfer Distance to bypass explicit semantic annotations during pretraining [20]. To satisfy the diversity term $\mathcal{H}(p(z))$, the framework employs stochastic sampling over unlimited combinations of elementary shapes (e.g., cubes, spheres) and affine transformation parameters, ensuring the latent space z captures broad structural variations without inheriting human-crafted biases. Crucially, the regularization term $\Omega(\Phi_{prog})$ encodes structural priors directly through the logic of the procedural program Φ_{prog} , where deterministic rule application inherently enforces manifold validity and physical plausibility rather than relying on learned soft constraints, although this rigidity can introduce scalability bottlenecks in highly complex scenes.

2.2 Technological History and Roadmap

Evolution from Rules to Data. The trajectory of 3D acquisition has shifted from rigid procedural modeling and manual CAD construction toward flexible, data-driven pipelines capable of bridging domain gaps with minimal real-world supervision [107]. Although high-fidelity laser scanning and image-based reconstruction provide superior initial feature representations, contemporary analyses indicate that generative text-to-3D methods currently lag in geometric fidelity within specific agricultural domains. Strategic synthetic pretraining followed by fine-tuning on

extremely sparse real samples effectively neutralizes these performance disparities between disparate acquisition modalities, validating a hybrid synthesis roadmap that achieves parity using as few as 4 real images.

Data Scarcity and Evolutionary Stages. The maturation of generative AI in text and image domains follows a sequential dependency where Professionally Generated Content establishes initial quality baselines before User Generated Content provides the scale necessary for foundation model training [73]. Three-dimensional vision currently lacks a unified large-scale pretrained model comparable to language counterparts due to the absence of a consumer hardware catalyst capable of generating sufficient volumetric user data. Consequently, deep learning progression in spatial domains adheres to a distinct three-stage evolution from task-specific architectures to general-task systems, contingent upon overcoming current data insufficiency through future hardware-driven content explosions.

Neural Representations and Diffusion Paradigms.

Contemporary pipelines have largely abandoned traditional mesh-based workflows in favor of neural implicit fields and explicit Gaussian representations, which serve as differentiable backbones for high-fidelity human synthesis [5, 163]. By integrating 2D diffusion priors through Score Distillation Sampling, these frameworks enable text-controlled generation, although the process frequently introduces geometric artifacts such as the Janus effect and suffers from significant optimization latency. To address dynamic coherence, recent architectural designs decouple semantic motion planning from trajectory execution by coupling Large Language Models with vector-quantized variational autoencoders [5]. This separation allows for LLM-driven scheduling that effectively reduces issues like foot sliding while maintaining consistent physical trajectories.

2.3 Datasets, Metrics, and Benchmarks

Geometric Fidelity and Physical Plausibility Metrics.

Contemporary evaluation protocols have transitioned from simple pixel-wise error calculations to no-reference frameworks that decompose scene quality into spatial view consistency and angular surface fidelity [116]. These approaches often employ architectures such as EfficientNet V2 or PointNet alongside metrics like Pointwise Normalized Spherical Gradient (PNSG) and Just-Objectable-Difference (JOD) to quantify inter-view coherence and angular quality dynamics. Beyond assessing visual realism, attribution benchmarks now treat cross-view geometric

inconsistencies and frequency-domain structural artifacts as stable fingerprints for identifying generative model sources [98]. By utilizing Hierarchical Multi-View Multi-Modal Transformers to analyze FFT frequency features and normal maps, these methods detect generator-specific patterns inherent in cross-view inconsistency. To support such rigorous assessment, synthetic dataset construction increasingly relies on formula-driven procedural generation to provide analytical ground truth for complex geometric properties like symmetry planes and axes [141]. This strategy involves generating parametric planar curves via Bézier extrusion with stochastic perturbations, thereby establishing precise rotation probabilities and symmetry detection baselines that circumvent the scarcity of labeled real-world data.

Perceptual Quality and Semantic Consistency.

Moving beyond static geometric fidelity checks, modern evaluation protocols now integrate non-visual sensor streams to assess dynamic perceptual quality within egocentric contexts. Frameworks such as Aria-NeRF and NeuralDiff extend traditional metrics like PSNR, SSIM, and LPIPS by incorporating temporal harmonic coding alongside IMU integration and audio cues to achieve robust dynamic foreground disentanglement [147]. Ensuring high-fidelity semantic control further requires automated pipelines that generate dense multi-view quadruples, verifying that textual descriptions align with geometric attributes across diverse viewpoints. Systems like UniG3D utilize CLIP-ViT, BLIP, and SVD heuristics during voxel sampling and Z-circular rendering of RGBAD images to enforce text-image consistency and viewpoint robustness [149]. Contemporary benchmarks have subsequently shifted from simple retrieval success metrics toward adversarial sampling strategies that quantify semantic grounding accuracy and hallucination rates. By employing random and adversarial sampling on instance-level annotations within USD formats, recent approaches leverage Mask3D proposals and LoRA fine-tuning to measure object overlap rates and generation success more rigorously [186, 219].

Downstream Task Benchmarks and Sim-to-Real Protocols.

Initial efforts to overcome the labor constraints of real-world data collection prioritized raw scale and automated annotation density through procedural spawners and multi-camera array simulations. These early frameworks generated pixel-level surface normals, occlusion modeling, and stereo image pairs to support specific perception tasks [55, 61]. Subsequent methodologies shifted toward parametric steerability and matched-pair experimental designs, utilizing Poisson-disc camera sampling and DAG-

based cue generation to enable rigorous evaluation across diverse mid-level vision cues [36]. Contemporary protocols now simulate full sensor reconstruction chains, incorporating stacked virtual LiDARs, Monte Carlo surface sampling, and GIS-driven terrain layouts to validate sim-to-real transfer under complex operational dynamics [16, 216]. Furthermore, recent object-centric composition strategies decouple geometric layout from category priors via mask-area-weighted blending and diffusion-based harmonization, effectively breaking spurious correlations to enhance generalization in low-data regimes [57].

Physical and Semantic Validation Gaps. Standardized semantic hierarchies frequently remain absent from physical property collection protocols, causing ground truth invalidity when illumination conditions diverge between targets and scenes [38]. Such discrepancies arise from mixed-light estimation errors and sensor incompatibilities that necessitate expensive manual labor for exposure compensation, as observed in systems utilizing SpyderCube 3D targets. Concurrently, large-scale asset repositories often harbor uncurated noise and subjective aesthetic inconsistencies that impede generative model convergence without rigorous multi-view quality filtering [83]. Although methods employing ResNet50-LSTM classifiers and four-tier quality rubrics attempt to mitigate these issues through metadata fusion, they still face limitations in detecting geometric watertightness and aligning with user preference rates.

2.4 Related Research Domains

Synthetic Pre-training for Geometric Consistency.

Integrating computer graphics pipelines into representation learning addresses the structural deficit inherent in purely 2D pre-training objectives by enforcing geometric consistency through procedural mesh generation and Gaussian process deformations [47]. Such synthetic domains, constructed from parametric primitives as seen in ProcSynthDB and MorphSynthDB, facilitate the systematic isolation of viewpoint and lighting factors to better approximate biological ventral stream hierarchies, yielding improved neural predictivity in primate V4/IT alignment via representational similarity analysis. This methodological shift reduces reliance on manual annotation by prioritizing algorithmic label generation, thereby enabling scalable supervised fine-tuning on natural image benchmarks through precise domain statistics alignment.

Multimodal Alignment and Semantic Grounding via Large Language Models.

Replacing sparse manual metadata with dense, holistic semantic descrip-

tions, recent pipelines integrate large multimodal models to resolve language scarcity in 3D pre-training [181]. These approaches utilize automated caption generation and multi-view rendering to construct rich annotations, thereby enhancing label diversity without human intervention. To establish a unified cross-modal representation manifold, tri-modal contrastive objectives align raw geometric encoders with frozen image and text embedding spaces [181]. By employing backbones such as Point-BERT alongside vision transformers like OpenCLIP ViT-G/14, these systems optimize dual contrastive losses for both point-to-image and point-to-text alignment. This self-supervised paradigm eliminates the dependency on human-curated triplets by algorithmically constructing scalable training data from unannotated 3D geometry alone [181]. Leveraging datasets like Objaverse-LVIS, such methods enable zero-shot transfer capabilities by ranking raw point cloud inputs against CLIP similarity metrics, effectively scaling semantic grounding to massive unlabeled corpora.

Embodied Utility and Simulation Readiness. Moving beyond visual fidelity, the synthesis of 3D assets for robotics necessitates strict adherence to simulation readiness criteria that integrate geometric validity with physical parameterization and kinematic executability [192]. This shift from appearance-oriented generation to functional utility ensures compatibility with downstream simulators by enforcing constraints on mass distribution, friction coefficients, and joint limits. Generative digital twins and physics-aware data augmentation serve as primary mechanisms to bridge the domain gap between synthetic asset creation and real-world policy deployment [192]. By embedding these physical priors directly into the generation process, recent approaches facilitate zero-shot policy transfer and improve task success rates in complex environments such as those defined by RoboCasa and PartNet-Mobility benchmarks.

Cross-Domain Stylization and Texture Synthesis Mechanisms. Bridging the dimensional gap between 2D diffusion priors and 3D geometric representations relies on optimization-based score distillation or direct feed-forward inference [21]. These frameworks enforce view-independent appearance constraints and multi-view consistency losses to prevent texture flooding while aligning CLIP features for semantic fidelity. The objective function typically balances a style loss term, derived from Gram matrix statistics or diffusion scores, against a content preservation component to maintain underlying geometry. Such approaches enable the transfer of artistic styles onto complex 3D assets without requiring paired training data across domains.

2.5 Discussion

Despite advances in generative capacity, prevailing evaluation protocols remain disproportionately fixated on geometric fidelity and perceptual scores, systematically neglecting physical plausibility and semantic consistency [21, 47, 60]. Metrics such as Chamfer Distance, FID, and CLIPScore quantify visual alignment yet fail to capture the absence of physical constraints or the scarcity of reliable semantic annotations, creating a pronounced disconnect between rendered realism and downstream utility. This evaluation gap exacerbates the inherent tension between controllability and diversity, where rigid structural priors like (F, l) -equivariance or SMPL template reliance ensure stability but limit generalization, whereas unconstrained generation yields variety at the cost of topological errors and Janus artifacts [20, 158, 163]. Consequently, bridging the domain divide necessitates unified frameworks that integrate geometric, semantic, and physical constraints rather than optimizing rendering quality or representation learning in isolation [20, 47, 107]. Such integration is essential for supporting robust real-world deployment, moving beyond procedural bias-free data generation toward systems capable of handling affine transformations, Boolean augmentations, and high-fidelity synthetic-to-real transfer with minimal fine-tuning.

3 Methods

This section categorizes 3D dataset generation methodologies into four distinct paradigms based on their underlying technical mechanisms. Sec. 3.1 examines rule-based and engine-driven approaches that utilize explicit algorithms, scene graphs, and physics simulators for data synthesis. Sec. 3.2 addresses neural generative models that directly produce 3D representations through architectures such as GANs, diffusion processes, and autoregressive transformers. Sec. 3.3 investigates techniques that lift 2D inputs into 3D structures by leveraging pretrained foundation models and enforcing multi-view consistency. Finally, Sec. 3.4 explores frameworks where large language models and agentic workflows orchestrate scene composition, asset retrieval, and procedural code generation.

3.1 Procedural and Algorithmic Synthesis

Procedural and algorithmic synthesis establishes a foundational framework for generating 3D data through explicit, deterministic rules rather than learned distributions. This approach begins with

formal grammars and shape languages that enforce structural validity, as detailed in [Sec. 3.1.1](#), before extending to graph-based representations that optimize spatial layouts for semantic coherence in [Sec. 3.1.2](#). To bridge the gap between abstract geometry and visual reality, [Sec. 3.1.3](#) examines pipelines utilizing game engines and ray tracers to produce photorealistic renderings with precise ground truth. The scope further narrows to hardware-specific fidelity in [Sec. 3.1.4](#), which models sensor physics and noise characteristics for LiDAR and depth cameras, while [Sec. 3.1.5](#) addresses generalization by systematically varying environmental parameters through domain randomization. Together, these methodologies provide a controlled mechanism for creating diverse, annotated datasets essential for training and evaluating perception systems.

3.1.1 Rule-Based and Grammar-Driven Generation

Syntactic Enforcement of Topological Integrity.

Procedural frameworks enforce topological validity by decomposing objects into executable structure programs governed by explicit syntactic rules rather than statistical priors [146]. These systems distinguish between macro spatial structure programs and micro geometric details, utilizing continuous, discrete, or advanced primitive alterations to ensure point-wise correspondence and analytic label alignment. Constructive Solid Geometry extends this logic through randomized boolean operation trees that assemble primitive shapes, where leaf count saturation and Maximum Mean Discrepancy distillation guide the generation process [80]. This approach guarantees geometric diversity while maintaining automatic part-level annotation traces via primitive origin tracking. For tasks prioritizing variation over semantic realism, minimalist rule sets generate abstract NURBS surfaces through lofting profile curves and 3D random walks [99]. Such methods apply micro-geometry displacement and boolean-combined procedural patterns to create hybrid object distributions suitable for multi-view stereo evaluation.

Stochastic Parameterization and Formula-Driven Diversity Mechanisms. Continuous surface geometries replace discrete mesh assets by defining radial distance functions in spherical coordinates, ensuring topological consistency during stochastic placement through analytical constraints [77]. Hierarchical complexity scales via the extrusion or revolution of parametric planar curves, such as Bézier splines or lemniscates, where ground truth symmetry planes derive directly from the generating algebraic definitions rather than post-hoc analysis [141]. Beyond static

forms, physiological and kinematic diversity emerges from direct manipulation of skeletal and morphological parameters within character engines, bypassing manual keyframing to automate mask derivation and segmentation indices [11]. These formula-driven approaches collectively shift the generation burden from asset retrieval to the sampling of latent mathematical domains.

Local Rule Propagation and Semantic Consistency.

By replacing global autoregressive sampling with iterative local transition kernels, generative cellular automata exploit volumetric sparsity through sparse convolutional backbones that update voxel states based on neighborhood connectivity [203]. This locality principle extends to scene synthesis, where complex arrangements decompose into sequential relational tuples governed by learned support constraints and functional anchors rather than intractable global joint distributions [119]. Convergence toward structurally valid configurations emerges from hybrid pipelines that integrate probabilistic rule prediction with physics-based collision resolution and gravity simulation, effectively filtering invalid assemblies through rejection sampling and ground truth mixing [119, 203].

Computational Trade-offs in Recursive Assembly.

Sequential assembly navigates combinatorial explosion by employing surrogate-model guidance, such as Bayesian optimization with Gaussian Process regression, to direct primitive placement within an occupiability grid [66]. This approach integrates multi-objective acquisition functions to balance competing structural goals while utilizing score threshold reversion mechanisms to rollback locally suboptimal decisions when greedy selection fails. In contrast, gridless constructive algorithms achieve real-time performance by enforcing visual connector constraints through pin-and-color matrices, thereby sacrificing global backtracking capabilities for immediate topological validation [35]. These systems rely on swappable selection strategies—ranging from arena and corridor modes to star and branch configurations—validated via box collider overlap detection and random point sampling to maintain the maximum area metric without exhaustive search.

Constraints on Global Coherence. Hierarchical solvers attempt to mitigate combinatorial explosion by decomposing global layout optimization into sequential stages, yet this decoupling inherently prevents the simultaneous resolution of distant geometric conflicts [118]. Similarly, agent-based chaining enforces strict workflow order through finite state machines, creating rigid dependency chains where

Table 1 Taxonomy of representative methods in **Procedural and Algorithmic Synthesis**. This table compares procedural synthesis methods by generation logic and rendering engine, showing diverse algorithmic strategies, and highlights the dominance of hybrid simulation pipelines.

Method	Venue	Generation Logic	Rendering Engine	Highlight
Rule-Based and Grammar-Driven Generation (Sec. 3.1.1)				
Infinigen Indoors [118]	CVPR'24	Constraint Solver	Blender Nodes	Declarative DSL defines geometric constraints; hierarchical solver enforces topological validity across scales. Structure programs apply primitive alterations to generate diverse shapes while preserving analytic label correspondence. Agent chains translate text to parameters; finite state machines enforce spatial adjacency rules during construction.
Arti-PG [146]	ICCV'24	Structure Programs	Analytic Mapping	
FilmSceneDesigner [175]	AAAI'25	Agent Chaining	SketchUp API	
Scene Graph and Layout Optimization (Sec. 3.1.2)				
DeBaRA [100]	NeurIPS'24	Score-Based Diffusion	PointNet Encoder	Semantic-aware Chamfer loss prevents category mismatches during permutation-invariant spatial layout optimization. Multi-instance attention layers capture cross-instance dependencies to optimize simultaneous multi-object denoising. Graph Attention Networks transform sparse scene graphs into dense embeddings for precise 2D-to-3D layout synthesis.
MIDI [58]	CVPR'25	Multi-Instance Diffusion	Latent Denoiser	
Scene Graph Guided 3D Gen [94]	ICCV'25	Graph Diffusion	3D-UNet	
Engine-Based Photorealistic Rendering (Sec. 3.1.3)				
Hypersim Pipeline [123]	ICCV'21	Asset Reuse	V-Ray Cloud	Triangle-density saliency guides collision-free walks; V-Ray renders disentangled lighting components for precise ground truth. Blender physics engine renders twenty-one mid-level cues via DAGs, synchronizing geometric primitives with dense semantic labels. Node transpilation randomizes Blender graphs; custom OpenGL extraction bypasses renderer artifacts to ensure noise-free ground truth. Differentiable rasterization of separated Gaussian models yields photorealistic views with automatic pixel-wise segmentation masks. Hybrid pipeline merges 3DGS backgrounds with ray-traced meshes in Isaac Sim, aligning coordinates via ArUco markers.
Omnidata Pipeline [36]	ICCV'21	Parametric Rendering	Blender Physics	
Infinigen [117]	CVPR'23	Mathematical Rules	Blender Cycles	
3DGS Surgical Fusion [200]	MICCAI'24	Gaussian Fusion	Differentiable Rasterizer	
RE3SIM [49]	arXiv'25	Hybrid Reconstruction	Isaac Sim	
Sensor-Specific Simulation (Sec. 3.1.4)				
DR.CPO [138]	AAAI'23	Point Merging	LiDAR Simulator	Applies Hidden Point Removal algorithms to enforce realistic self-occlusion and distance-dependent point density. Utilizes occlusion-aware composition losses with segmentation masks to train robust neural volumetric rendering. Integrates procedural foliage generation with physics-based LiDAR ray-tracing to simulate realistic point cloud sparsity.
GINA-3D [133]	CVPR'23	Vector Quantization	Neural Volumetric	
CAMP3D Pipeline [132]	IJCV'25	Foliage Procedural	HELIOS LiDAR	
Domain Randomization and Variation (Sec. 3.1.5)				
LRM-Zero [174]	NeurIPS'24	Primitive Compositing	GS-LRM	Procedural geometric primitives with targeted shape augmentations drive robust reconstruction generalization. Non-semantic scene synthesis controls geometry and lighting to align with real-world statistics. NURBS lofting and hybrid object distributions maximize viewpoint diversity for robust training.
MegaSynth [63]	CVPR'25	Primitive Composition	GS-LRM	
SimpleProc [99]	arXiv'26	NURBS Lofting	MVSAnywhere	

early-stage parameter errors propagate irreversibly to downstream spatial arrangements [175]. Although declarative constraint APIs successfully separate rule specification from execution, they frequently rely on local move acceptance criteria that fail to capture holistic scene semantics requiring long-range contextual awareness [118, 175]. Consequently, these methods struggle to maintain structural validity when local rules conflict with global ergonomic or cinematic requirements.

3.1.2 Scene Graph and Layout Optimization

Graph-to-Layout Encoding and Dense Embedding Transformation. Converting sparse scene graph structures into dense spatial representations requires mechanisms that aggregate node features while resolving ambiguous relational constraints. Recent approaches employ Graph Attention Networks to compute context-aware node embeddings, which are subsequently mapped to Bird’s Eye View grids via probabilistic allocation strategies such as Gumbel Softmax sampling [94]. This sparse-to-dense transfor-

mation facilitates conditional diffusion processes that generate physically plausible layouts from abstract semantic inputs. To maintain geometric consistency during end-to-end synthesis, other frameworks encode layout geometry and object shape features within a shared latent distribution [32]. By utilizing parallel Graph Convolutional Network encoders and variational autoencoders, these methods jointly predict 3D bounding boxes and decode implicit shape representations, ensuring that generated objects adhere to both global arrangement logic and local structural fidelity.

Simultaneous Multi-Instance Optimization. By processing multiple object latents through shared attention layers, recent frameworks enable direct modeling of cross-instance spatial dependencies while avoiding the error accumulation typical of sequential generation pipelines [58]. This architecture employs a diffusion transformer backbone where token querying across the entire scene facilitates synchronous denoising, ensuring that geometric adjustments in one region immediately inform updates elsewhere. To further enhance layout realism under floor plan constraints, other approaches disentangle categorical semantics from geometric attributes using permutation-invariant matching losses [100]. Such methods utilize semantic-aware Chamfer distances and unordered set processing to improve data efficiency, allowing noise-conditioned denoisers to learn robust distributions without relying on fixed object ordering. Collectively, these joint optimization strategies mitigate alignment errors inherent in retrieval-based or multi-stage compositional workflows by preserving global scene context throughout the latent refinement process [58, 100].

Relational Priors for Ambiguity Resolution. Decomposing global scene synthesis into sequential local relational tuples defined by support and functional anchors circumvents the intractability of modeling joint distributions for complex layouts [119]. This hierarchical relational serialization employs breadth-first or depth-first traversal to construct support trees, optimizing Negative Log-Likelihood to ensure structural coherence. Integrating fine-grained point cloud geometry via specialized attention mechanisms enables precise physical affordance reasoning during procedural assembly [119]. Architectures utilizing Point-MAE encoders combined with geometry-aware cross-attention and relational self-attention modules effectively model mixture of logistics distributions while respecting gravity simulation constraints. Distinct from purely geometric approaches, employing generative Bayesian networks to model realistic scene-object distributions yields transferable representa-

tions that resolve ambiguous spatial configurations where multiple valid placements exist for a single object [165]. These probabilistic models utilize relaxed point pair matching and coarse-to-fine contrastive learning to enforce object-level consistency, thereby selecting contextually appropriate positions in dense environments without relying on explicit visual occlusion cues.

Hierarchical Decomposition and Stochastic Search. To address the combinatorial explosion inherent in large-scale spatial arrangement, recent protocols decompose global optimization into tractable sub-problems solved via stochastic search. Methods such as ArrangeIt employ hierarchical numerical optimization where simulated annealing iteratively refines object poses to satisfy ergonomic and aesthetic constraints while maintaining global context conditioning [90]. Parallel approaches automate valid asset distribution by coupling explicit 3D scene graphs with structured sub-graph extraction, thereby eliminating manual spatial tuning across diverse environments. By integrating OpenStreetMap data to define container object types, these pipelines ensure that placements adhere to structural validity without requiring per-scene reconfiguration [126].

3.1.3 Engine-Based Photorealistic Rendering

Rendering Pipeline Integration for Procedural Assets. Engine-based photorealistic rendering pipelines ingest algorithmically generated geometry to facilitate high-fidelity visual synthesis with precise ground truth annotations [55, 117]. These systems prioritize the real-time ingestion of procedural meshes, where the rendering engine maps dynamic geometry directly into shader graphs and manages buffer allocation for immediate rasterization or ray tracing. Hybrid architectures integrate such geometry with learned priors, employing the engine’s core to enforce structural fidelity while synthesizing complex surface appearances [68]. Frameworks utilize sampling strategies within the rendering loop to maintain geometric consistency during image generation, ensuring that synthetic outputs align with physical light transport models. Furthermore, texture mapping techniques operate at the material stage, decoupling surface complexity from underlying mesh topology by applying displacement maps during the shading pass [140]. By integrating external displacement data, these methods achieve mesh-independent variation essential for semantic city generation while preserving the integrity of the rendering pipeline’s output buffers.

Synchronization of Visual Output with Semantic

Ground Truth. By explicitly manipulating scene representations, pipelines decouple dynamic foreground entities from static backgrounds to yield pixel-perfect segmentation masks without retraining underlying models [200]. Ray-casting algorithms further establish exact bijective mappings between 2D image coordinates and 3D mesh vertices, generating dense correspondence labels that manual annotation cannot match [69]. Directed Acyclic Graph architectures derive complex mid-level vision cues directly from geometric primitives during the rendering pass, ensuring intrinsic consistency across multi-task datasets [36]. These frameworks synchronize automated metadata streams with the rendering loop to maintain strict alignment between generated imagery and their corresponding semantic labels [109].

Computational Efficiency and Scalable Rendering Architectures. Bypassing the latency inherent in embedded script interpreters requires decoupling scene graph management from the rendering kernel through external interfaces that facilitate asynchronous state replication across multi-GPU clusters [109]. This architectural separation enables hardware-accelerated path tracing backends, such as those utilizing OptiX Wrapper Libraries, to operate independently of Python-based Entity Component Systems while maintaining synchronized data uploads. Tighter integration strategies merge physics-based arrangement engines directly with rasterization pipelines, allowing online data synthesis within training loops to eliminate storage overhead for large-scale cluttered environments [130]. Such systems couple PhysX simulations with OpenGL contexts to support render-and-compare workflows essential for lifetime learning without intermediate disk I/O. To further optimize compute utilization, adaptive object-centric sampling employs learned semantic filters that maximize usable render yield per unit while mitigating the prohibitive costs of full path tracing [157]. These approaches combine orbital camera placement with fill-fraction scaling and vision-language model captionability checks, discarding low-information samples based on semantic alignment scores before embedding them into foundation model training sets to streamline batch processing.

Artifact Mitigation and Rendering Fidelity. To reconcile the conflict between visual continuity and physical interaction, hybrid architectures decouple background reconstruction from foreground dynamics by merging neural point-based representations with explicit mesh geometry [49]. This synthesis allows ray-traced meshes to enforce rigid object constraints and collision accuracy while neural splats optimize high-frequency texture fidelity, often vali-

dated through post-processing alignment protocols. Beyond geometric hybridization, disentangled image factorization isolates diffuse reflectance and illumination from non-diffuse residuals to explicitly model view-dependent lighting effects that standard rasterization pipelines obscure [123]. Such separation facilitates inverse rendering support and high-dynamic-range export by treating triangle-density saliency as a distinct optimization target rather than a baked texture attribute. Complementing these architectural refinements, extreme density randomization of photo-realistic distractor objects compels feature extractors to disregard spurious background correlations in favor of invariant object characteristics [124, 185]. By enforcing edge-map preservation under high-density occlusion within engines like Cycles, these methods generate training sets structured to support domain contrastive losses and pseudo-correspondences, ensuring segmentation models remain stable against out-of-distribution background selections.

Trajectory Optimization for Viewpoint Robustness. To maximize geometric coverage while preventing viewpoint clustering, discrete orbital sampling strategies position virtual cameras along trajectories defined by object bounding boxes with adaptive distance scaling [157]. This object-centric placement enforces optimal fill-fractions, ensuring that the yield of diverse views remains high without redundant angular overlap. Beyond static distribution, systematic manipulation of camera extrinsics generates extreme angular configurations and rapid motion sequences to explicitly address self-occlusion and trajectory drift [6, 213]. By synthesizing homogeneous projection matrices across egocentric and allocentric frames, these methods reduce correspondence ambiguity and minimize per-joint error during fast viewpoint changes.

3.1.4 Sensor-Specific Simulation

Empirical Calibration of Distance-Dependent Noise. Accurate sensor simulation demands fitting statistical distributions to physical variables such as range and surface incidence angle, replacing uniform perturbation models with distance-correlated noise profiles [111]. This approach typically models lateral and axial standard deviations using second-degree polynomials that explicitly account for surface angle dependency. To address unmodeled environmental interference, optimal domain adaptation often requires training with noise magnitudes scaled slightly above empirical measurements, where a scaling factor of 1.25 has been shown to maximize IoU performance through implicit robustness optimization [111]. Beyond direct ray casting, replicating the full photogrammetric reconstruction pipeline from rendered

images generates point cloud artifacts and density variations that more faithfully mimic real-world sensor outputs [16]. Such procedural environment generation and artifact emulation yield superior mIoU results by capturing the complex degradation patterns inherent in actual data acquisition workflows.

Geometric Visibility and Density Modulation. Algorithmic enforcement of self-occlusion replaces uniform point sampling by applying spherical flipping and convex hull checks to simulate strict line-of-sight constraints [138]. To generate stratified sparsity patterns without separate scene renders, multi-density architectures stack virtual sensors with varying channel counts at identical coordinates [216]. Physics-based ray-tracing simulators further improve ecological validity over simple mesh nodal sampling by modeling specific scanner beam divergence and flight paths [132]. This approach couples distance-dependent sparsity with geometric completion, ensuring synthesized objects maintain realistic fall-off profiles consistent with fixed-viewpoint LiDAR characteristics [132, 138].

Temporal Alignment and Multi-Sensor Synchronization. Maintaining frame-to-frame consistency in dynamic pipelines requires simulating environmental migration sequences that explicitly test place recognition robustness against short-term structural variations [216]. To address multi-sensor coherence, generative frameworks enforce strict temporal registration across modalities by aligning discrete latent structures with synchronized timestamp streams, ensuring geometric consistency despite occlusion and lighting shifts [133]. This architectural separation enables self-supervised adversarial training to validate synchronization strategies without relying on static scene assumptions.

Material Reflectivity and Geometric Fidelity Constraints. Achieving detector robustness on reflective industrial surfaces necessitates randomizing Physically Based Rendering (PBR) parameters, thereby forcing reliance on geometric cues rather than brittle texture correlations [3]. However, current automated 2D-to-3D reconstruction pipelines often fail to provide the geometric consistency required for high-precision sensor simulation, exhibiting artifacts that manual CAD modeling or 3D Gaussian Splatting baselines avoid. These asset granularity bottlenecks further restrict physics-based LiDAR simulations to binary semantic labels, creating an ecological validity gap where accurate point density modeling exists without fine-grained multi-class segmentation [132].

3.1.5 Domain Randomization and Variation

Geometric and Spatial Parameter Randomization. Object identity decouples from background context through stochastic 3D insertion and edge-conditioned composition [2, 199]. This approach integrates HDRI lighting consistency with SE(3) transformations, utilizing ControlNet edge conditioning to quantify view-point diversity without relying on specific rendering engines. Hierarchical probabilistic sampling in local spherical coordinates further generates biologically plausible skeletal variations while preserving kinematic constraints through parent-joint conditioning [222]. Such conditioning ensures scaleless multi-view consistency and reduces Mean Per-Joint Position Error (MPJPE). To address partial observation limits, pipelines synthesize complete whole-body shapes by merging partitioned point clouds [63, 99, 138], focusing on geometric completion rather than sensor-specific noise modeling or surface reconstruction techniques.

Photometric Variability and Material Property Perturbation. Replacing static baked textures with physically-based material definitions utilizing Principled BSDF models narrows the reality gap for reflective surfaces by randomizing metallic and roughness parameters within defined physical bounds [3, 109]. To enforce detector invariance against illumination shifts, pipelines sample lighting intensity from exponential distributions while perturbing RGB color channels, a strategy that maintains physical plausibility amidst high-density distractors and occlusion variance [3, 124]. Rather than detailing generative architectures, recent methods employ edge-conditioned networks to reskin synthetic CAD renders, aligning output texture distributions with real-world statistics through schematic representation learning [28]. Hybrid approaches further automate surface appearance generation by coupling procedural depth constraints with synthesis models, where style adapters function strictly as mechanisms to map randomized material attributes onto geometric fidelity without altering the underlying domain randomization strategy [68].

Bounded Randomization Strategies and Physical Plausibility Constraints. Preventing overfitting to specific training distributions requires defining strict parameter bounds rather than relying on unbounded stochasticity. Adversarial trajectory generation expands state coverage by injecting perturbations specifically at bottleneck poses, compelling policies to execute recovery maneuvers within physically constrained simulation environments [34]. To mimic real-world solidity and improve shape generalization, pipelines replace uniform placement with

Gaussian mixture distributions and apply stochastic mesh extrusions that enforce non-uniform object clustering [99, 108]. These strategies prioritize the calculation of valid sampling limits and the enforcement of physical constraints on object placement and density to ensure generated variations remain within the manifold of plausible reality, distinct from the synthesis of novel geometric primitives.

Decoupling Semantics from Domain Artifacts.

Strategic disruption of spurious correlations is achieved by decoupling task semantics from domain-specific background features through stochastic environment randomization and targeted feature suppression. Techniques such as OOD-aware background randomization and edge-conditioned style transfer force the model to ignore non-invariant contextual cues, while mesh material modification further severs reliance on specific surface textures [105, 185]. Complementing this, SHAP-guided common feature disruption and contrastive feature invariance objectives explicitly penalize attention to static environmental priors, ensuring that learned representations remain stable across shifting distributions. To enforce object-level discriminative convergence, subsequent approaches eliminate contextual priors entirely via blank-background synthesis and replace semantic complexity with geometric primitives [67, 174]. Methods utilizing null-context synthesis and randomized object distribution prevent reliance on global scene statistics by stitching objects onto empty canvases or compositing non-semantic primitives, effectively acting as a visual dropout mechanism that compels dependence on local geometric cues rather than holistic scene context.

3.1.6 Discussion

Strict adherence to syntactic and physical constraints within procedural pipelines ensures geometric validity and noise-free ground truth, yet this rigidity often curtails visual diversity and limits automation scalability [20, 117, 123, 140]. While integrating scene graphs with probabilistic distributions improves semantic plausibility and spatial coherence, such hybridization introduces fragility when underlying graph extraction or distribution models falter under imperfect conditions [126, 161, 165]. Consequently, recent architectural shifts increasingly combine deterministic rules with neural components to balance controllable variation against learned realism, moving beyond pure randomization toward context-aware synthesis that preserves label fidelity [61, 112, 132]. This evolution reflects a broader trend where the procedural paradigm serves not as a standalone solution but as a structured backbone for neural generative pro-

cesses, mitigating the trade-off between precise physical grounding and the stochastic richness required for robust model generalization.

3.2 Neural 3D Generative Modeling

This section surveys deep generative frameworks that synthesize 3D geometry and appearance directly, bypassing the need for 2D image priors or intermediate projections. We first examine adversarial architectures designed to stabilize training across discrete voxels, point clouds, and continuous implicit fields in Sec. 3.2.1. The discussion then shifts to diffusion-based approaches that model complex 3D data distributions through iterative denoising, including recent advancements in latent space compression detailed in Sec. 3.2.2. Subsequently, Sec. 3.2.3 analyzes autoregressive transformers that treat 3D structures as token sequences to capture long-range spatial dependencies via attention mechanisms. We further distinguish between methods generating continuous neural representations for high-fidelity rendering in Sec. 3.2.4 and those producing explicit meshes or Gaussian splats optimized for real-time editing in Sec. 3.2.5.

3.2.1 3D Generative Adversarial Networks

Architectural Inductive Biases for Volumetric Consistency. Achieving coherent inter-slice gradients in medical synthesis often necessitates replacing standard generator backbones with structures featuring explicit skip connections or style-based modulation [53, 136]. These designs, such as 3D U-Nets or style-based 3D convolutions, map latent codes z to intermediate spaces w to preserve global-local structural dependencies across the volume. Distinct from single-network approaches, certain frameworks enforce geometric coherence through dual-discriminator setups that simultaneously optimize rendered view realism and internal volumetric feature distributions [65]. This strategy typically employs differentiable volume tracing on discrete voxel grids to constrain foreground density while satisfying multi-view consistency. Alternatively, compositional scene generation avoids the computational burden of native 3D convolutions by utilizing efficient 2D-to-3D extrusion operators [4]. Such methods impose strong geometric priors via semantic layout conditioning and trilinear interpolation, ensuring top-down rendering consistency during direct voxel assembly.

Latent Disentanglement and Geometric Control.

Enforcing cyclic consistency between latent neighbors and their reconstructed 3D outputs densifies valid regions on the manifold, effectively suppress-

Table 2 Taxonomy of representative methods in **Neural 3D Generative Modeling**. This table presents a taxonomy of direct 3D generative methods, shows representation formats and generative paradigms, and highlights architectural diversity.

Method	Venue	Rep. Format	Gen. Paradigm	Highlight
3D Generative Adversarial Networks (Sec. 3.2.1)				
SceneDreameer [24]	TPAMI’23	Hash Grid	BEV GAN	BEV height and semantic fields with hash grids enable unbounded scene synthesis via semantic-aware discrimination.
CC3D [4]	ICCV’23	Neural Fields	GAN Extrusion	2D-to-3D extrusion operator and segmentation-auxiliary discriminator enforce geometric consistency from semantic layouts.
TPA3D [14]	ECCV’24	Triplane GAN	Word-Refine	TriPlane Attention blocks and pseudo-caption discriminators align fine-grained text features with 3D geometry.
3D Diffusion and Latent Models (Sec. 3.2.2)				
Frankenstein [184]	SIGGRAPH Asia’24	Tri-Plane SDF	Latent Diffusion	Compresses tri-planes via VAE for layout-conditioned diffusion decoding multiple SDF fields.
MIDI [58]	CVPR’25	Multi-Instance UV	Rectified Flow	Integrates multi-instance attention to synchronously denoise latent representations of correlated 3D objects.
SIGMAN [190]	ICCV’25	Gaussians Structured	Latent Diffusion	Maps multi-view renderings to UV-structured latent tokens on SMPL-X space for pose-driven generation.
WorldGrow [78]	AAAI’25	Latents	Inpainting	Adapts SLAT with occlusion-aware aggregation for context-guided coarse-to-fine scene block inpainting.
WorldFlow3D [64]	arXiv’26	TUDF Volume	Flow Matching	Formulates latent-free volumetric flow matching on TUDF using chunk-aware velocity averaging for unbounded scenes.
Autoregressive and Transformer-Based Synthesis (Sec. 3.2.3)				
G3PT [205]	IJCAI’24	Point Tokens	Next-Scale	Octree-based window attention and depth-wise teacher-forcing masks enable parallel multi-token prediction while preserving hierarchical dependencies.
SAR3D [22]	CVPR’25	Triplane Tokens	Next-Scale	Next-scale prediction paradigm replaces token-level steps, significantly accelerating autoregressive 3D synthesis via hierarchical token maps.
MAR-3D [15]	CVPR’25	Point Clouds	Autoregressive	Cascaded Masked Auto-Regressive stages with condition augmentation mitigate error propagation during coarse-to-fine resolution upscaling.
Direct3D-S2 [170]	NeurIPS’25	Sparse SDF	Rectified Flow	Spatial Sparse Attention partitions tokens by 3D coordinates to preserve geometric coherence and reduce self-attention computational costs.
Implicit Neural Representations (Sec. 3.2.4)				
SceneDreameer [24]	TPAMI’23	Hash Grid	BEV GAN	Generative neural hash grids parameterize space-varied features for unbounded BEV scene synthesis.
GINA-3D [133]	CVPR’23	Tri-Plane Codes	MaskGIT	Vector quantization discretizes tri-plane embeddings to decouple representation learning from generative modeling.
Frankenstein [184]	SIGGRAPH Asia’24	Tri-Plane SDF	Latent Diffusion	Single tri-plane tensor decodes multiple class-specific SDF fields preserving complete compositional shapes.
Explicit Representation Generation (Sec. 3.2.5)				
Gaussian Shell Maps (GSM) [1]	CVPR’24	Shell Gaussians	Texture GAN	Anchors Gaussians on deformable shells, predicting attributes via texture maps for articulated rendering.
DreamScene [76]	ECCV’24	3D Gaussians	Score Distill	Filters redundant Gaussians and samples cameras incrementally to ensure stable scene consistency.
Can3Tok [43]	ICCV’25	3D Gaussians	VAE	Tokenizes unstructured Gaussians using canonical voxel queries for structured latent representation learning.

ing volumetric discontinuities through geometric continuity constraints [46]. To enable controlled non-linear deformation while preserving identity, domain-specific PCA subspaces restrict latent trajectories using statistical face space metrics that balance character exaggeration with cosine similarity constraints [193]. Hybrid architectures further decouple shape and texture by employing linear parametric models for demographic offsets alongside non-linear adversarial networks dedicated to high-frequency render map synthesis [30]. Finally, analysis-by-synthesis pipelines replace global additive vectors with learned non-linear residual controllers, utilizing synthetic paired datasets derived from latent navigation to achieve fine-grained attribute manipulation via conditional parameter shifts [26].

Latent Regularization and Feature Matching Strategies. Stabilizing adversarial training on diverse, unposed data necessitates averaging discrimination losses across multiple rendered views to counteract distributional mismatches arising from uniform camera sampling [134]. To prevent geometric mode collapse in single-scene modeling without relying on dual-stream architectures, methods employ feature matching objectives that minimize the distance between intermediate discriminator statistics of real and generated volumes, thereby enforcing structural fidelity directly within the latent manifold. In unbounded domains, rejection strategies based on depth statistics and semantic entropy filter invalid camera trajectories, maintaining model stability where naive random sampling fails [24]. Furthermore, text-guided discriminators trained on synthetic

pseudo-captions enforce fine-grained semantic alignment through cross-word attention, effectively mitigating mode collapse in unconditional voxel synthesis despite limitations in global sentence embedding [14].

Conditional Generation and Semantic Layout Control.

Structural topology transfer operates by utilizing conditional voxel grids to predict node and link probability distributions, which procedural graph extraction algorithms subsequently convert into explicit geometry [143]. Distinct from this probabilistic approach, scene graph conditioning enforces semantic layout constraints by embedding relational graphs into a shared latent space that jointly decodes bounding box arrangements and object shapes [32]. In the domain of molecular synthesis, variational encoding of atomic density grids enables the generation of ligand structures that satisfy specific 3D spatial interaction constraints within receptor binding sites [101]. Expanding beyond fixed categorical labels, open-vocabulary text control injects dynamic noise into text embeddings to create noisy fields that align view-invariant image latents with semantic prompts [56].

3.2.2 3D Diffusion and Latent Models

Hierarchical and Contextual Noise Adaptation.

By factorizing the joint distribution across resolution scales, diffusion models can synthesize sparse voxel hierarchies through sequential denoising where coarser levels condition the generation of finer details [122]. This coarse-to-fine strategy often employs VDB-based indexing to manage conditional probability modeling efficiently during level-by-level sampling. To maintain geometric continuity when expanding scenes beyond fixed grid boundaries, sliding-window strategies synchronize denoising trajectories across overlapping blocks via cross-plane attention and feature injection [172]. Such mechanisms facilitate tri-plane extrapolation while ensuring non-rigid registration aligns features across disjoint regions. Furthermore, context-aware block inpainting separates global layout prediction from local structure refinement to support iterative scene growth [78]. This two-stage approach utilizes occlusion-aware aggregation and boundary-optimized decoding to adapt sparse latent representations dynamically without compromising structural coherence at expansion frontiers.

Continuous Latent Compression via Triplanes and UV Maps.

To maintain structural integrity during the compression of colored point clouds into tri-plane latents, recent approaches replace fixed pooling with learnable projection operators and 3D-aware attention mechanisms [201]. These systems utilize

feature volume querying within a dual-conditioning strategy to guide FlexiCubes decoding, ensuring that geometric details survive the dimensionality reduction. Addressing the ill-posed mapping ambiguities common in high-frequency human asset generation, other methods define UV-structured latent tokens on parametric body priors to decouple appearance from pose [190]. By operating in SMPL-X UV space with Conv-Atten dual branches, such frameworks model distribution shifts and handle loose clothing deformation through differentiable linear blend skinning.

Conditioning via Cross-Attention and Layout Embeddings.

Synchronous feature querying across multiple instances utilizes shared attention layers to capture inter-object spatial dependencies, effectively mitigating error accumulation in global scene contexts [58]. For discrete pipelines, sparse structural descriptions transform into Bird’s Eye View maps through graph-to-dense embedding modules that employ Graph Attention Networks and Gumbel Softmax sampling for semantic map refinement [94]. Distinct class-specific shapes are modeled simultaneously under layout constraints by decoding a single tri-plane representation through separate implicit field heads, utilizing semantic-aware point sampling within a VAE compressed latent space [184]. Furthermore, direct injection of layout embeddings into the denoising network enables instruction-driven editing while preserving geometric identity without quantization loss [23].

Convergence Stability and Volumetric Scalability.

Replacing stochastic diffusion trajectories with deterministic optimal transport paths via flow matching ensures structural stability when conditioning on extremely sparse categorical inputs [64, 96]. This approach utilizes simulation-free vector field regression optimized through Mean Squared Error to mitigate representation collapse, effectively managing coarse-to-fine distribution transport over Truncated Unsigned Distance Fields. To address boundary discontinuities and noise propagation in large-scale synthesis, spatially adaptive inference mechanisms employ attention-gated skip connections alongside chunk-aware velocity averaging [81]. Such architectures integrate 3D Spatial Attention Gates within decoder pathways to generate gating signals that facilitate feather-weighted blending across overlapping chunks, while Poisson disk resampling supports consistent feed-forward feature extraction without requiring global re-computation.

3.2.3 Autoregressive and Transformer-Based Synthesis

Spatial Serialization and Tokenization Strategies.

Resolving the ambiguity inherent in flattening volumetric grids requires projecting 3D data onto lower-dimensional manifolds or latent vectors to establish a tractable serialization order [114]. Hierarchical spatial structures improve upon this by utilizing tree-based traversal orders, such as octree Z-ordering, which preserve geometric locality more effectively than rasterization while enabling parallel token generation across scales through depth-wise teacher-forcing masks and 3D Rotary Positional Embeddings (RoPE3D) [159, 167]. Adaptive strategies further decouple structural partitioning from content generation by allocating resolution based on geometric saliency, employing techniques like Saliency-guided Centroidal Voronoi Tessellation (SCVT) to create variable-sized supervoxels that optimize sequence length [92]. Finally, cross-scale querying mechanisms replace fixed spatial orders with learnable attention queries, utilizing Cross-scale Vector Quantization (CVQ) within coarse-to-fine causal transformers to model natural sequential relationships across levels of detail without relying on predetermined mapping sequences [205].

Hierarchical Scale Prediction and Residual Quantization.

Reframing 3D synthesis as next-scale prediction rather than strict next-token generation reduces autoregressive steps, lowering computational complexity compared to flat token generation while preserving global structural priors [22, 102]. This coarse-to-fine sequence utilizes residual quantization to accelerate inference without sacrificing high-frequency detail fidelity. Variational constraints further enable latent token distributions to self-organize across scales, mapping inputs to Gaussian subspaces that mitigate exposure bias and maintain probabilistic manifold consistency within compact representations [206]. Such variational vector quantization achieves significant compression ratios, directly addressing the exponential growth in token counts and associated gradient instability inherent in deep multi-scale generative stacks.

Architectural Mechanisms for Error Mitigation and Variable-Length Handling.

To suppress error propagation during resolution upscaling, cascaded generation stages decouple structural synthesis from detail refinement through condition augmentation and spatial autoregression [15]. Addressing variable-complexity asset sequences, graph traversal re-ordering enforces local connectivity while speculative decoding strategies reduce latency via draft

model acceleration and discrete token tuples [221]. Test-time search algorithms further correct autoregressive drift by optimizing geometric alignment rewards, such as silhouette overlap, without requiring differentiable rendering gradients [209]. Finally, interleaved cross-modal training cycles impose implicit structural constraints to stabilize generation under conditions of severe 3D data scarcity by unifying heterogeneous datasets within a dual-stream transformer framework [191].

Sparse Attention and Representation Compression.

To mitigate the quadratic scaling of self-attention with respect to token count, spatially adaptive block partitioning replaces dense sequence indexing to preserve geometric continuity while reducing complexity to linear growth [170]. This approach utilizes 3D coordinate-based blocking and top-k relevant block selection to filter background tokens, effectively gating global-local interactions through sparse compression mechanisms. Complementing this structural sparsity, tri-plane feature encoding projects volumetric data onto orthogonal planes, transforming cubic memory requirements into quadratic complexity [197]. By compressing the attention matrix footprint, this representation supports large-scale autoregressive transformers via vector quantization and discrete codebook indices without prohibitive memory overhead.

3.2.4 Implicit Neural Representations

Discretization Strategies for Continuous Field Parameterization.

Decoupling continuous geometric representation from generative sampling via discrete latent quantization facilitates robust asset synthesis from uncurated data distributions [133]. This approach typically employs vector quantization on tri-plane embeddings, utilizing codebook projection alongside occlusion-aware composition losses and pseudo-labeled segmentation masks to maintain structural fidelity. To address memory constraints in high-resolution generation, hybrid planar factorization replaces volumetric convolution with axis-aligned feature maps [169]. Such architectures integrate voxel pyramids and position-invariant MLP decoders, often trained with hierarchical generators, Gaussian noise augmentation, and local patch discriminators to refine fine-grained details. For modeling unbounded scenes, semantic-conditioned hash grids parameterize geometry by disentangling height and semantic fields [24]. These systems leverage Bird’s-Eye-View representations initialized with simplex noise, applying rejection sampling strategies, depth mean filtering, and semantic entropy metrics to scale effectively across infinite domains.

Enforcing Surface Smoothness and Topological Correctness via SDF Priors. Decomposing a unified latent tensor into multiple class-specific implicit fields enforces geometric separation and prevents semantic entanglement during scene composition [184]. This factorization typically employs tri-plane representations with a shared MLP decoder but distinct signed distance function (SDF) heads, utilizing semantic-aware point sampling to guide coarse-to-fine optimization. Mapping discrete surface definitions to a continuous latent manifold further regularizes topology by constraining optimization trajectories within a learned prior of valid geometries [33]. Such approaches often instantiate an SDF Variational Autoencoder, where latent feature tuning and parameter sculpting on synthetic random geometry training data incorporate multi-view normal priors to maintain structural integrity. Integrating these neural surface representations with differentiable mesh extraction pipelines ensures watertight outputs suitable for physics-based simulation and automated dataset synthesis [128]. By coupling neural implicit surfaces with adaptive spatial encoding and differentiable iso-surface extraction, systems achieve refined rendering capabilities that support domain randomization for downstream applications.

View Consistency Mechanisms in Single-Image and Unsupervised Regimes. To enforce geometric coherence without multi-view supervision, dual-discriminator frameworks simultaneously optimize rendered image realism and volumetric feature consistency through differentiable volume tracing [65]. This approach couples a 2D patch discriminator with a 3D volumetric discriminator acting on ReLU-based field representations, effectively mitigating mode collapse in single-scene reconstructions. Distinct from purely adversarial volume regularization, parametric disentanglement strategies embed shape, expression, and appearance into a native 3D latent space to guarantee view-consistent avatar synthesis across novel poses [223]. By utilizing hex-plane factorization and bilinear parameter mapping, these methods achieve full 360-degree renderability even when trained on limited data sequences.

Gradient Propagation and Caching in Sparse Representations. While sparse hash grids and axis-aligned tri-planes reduce static memory footprint, their primary computational trade-off lies in the overhead of dynamic index lookups and non-contiguous gradient accumulation during backpropagation [24, 169]. Unlike dense tensors that benefit from coalesced memory access, these factorized structures require frequent random access to feature caches, which can stall pipeline throughput despite

linear scaling with respect to sample count. Efficient synthesis in unbounded domains therefore depends on optimizing voxel pyramid traversal and Bird’s-Eye-View height field updates to minimize cache misses rather than solely maximizing resolution capacity.

Generalization to Unseen Categories. Synthesis across uncured, long-tail distributions relies on regularization strategies that decouple shape priors from specific training domains rather than reusing discrete quantization mechanics [133]. By optimizing for shape diversity and category-agnostic fidelity metrics through pseudo-labeled segmentation masks, these systems achieve category agnosticism without explicit scene modeling constraints. Alternatively, restricting optimization trajectories to a continuous latent manifold trained exclusively on synthetic random geometries removes dependency on scarce public datasets while preserving topological validity for novel object classes [33]. Such methods yield improved geometric consistency scores and semantic alignment on unseen categories by enforcing structural coherence through multi-view normal priors rather than dataset-specific feature memorization.

3.2.5 Explicit Representation Generation

Structured Priors for Joint Topology and Geometry. Anchoring explicit primitives to deformable parametric surfaces bypasses volumetric optimization bottlenecks while preserving articulation consistency through fixed-location anchors and linear blend skinning [1]. Unified latent representations subsequently encode mesh topology and texture jointly within diffusion frameworks, utilizing transformer-based scene encoders and flow matching to eliminate domain gaps inherent in 2D-to-3D lifting pipelines [91]. Sequential modular architectures further decouple geometry synthesis from material estimation via rectified flow diffusion and coordinate-conditioned UV inpainting, ensuring watertight manifolds suitable for downstream physics simulation [42].

Geometric Constraints and Collision Avoidance. To satisfy conflicting visual and physical simulation requirements, hybrid rendering pipelines decouple background photorealism from foreground geometric fidelity by combining 3D Gaussian Splatting for environments with ray-traced meshes for interactive objects [49]. This architectural separation ensures precise collision accuracy and supports rigid body dynamics within simulators such as Isaac Sim without compromising the high-frequency detail of the surrounding scene. In biological contexts, global structural coherence during the assembly of novel geometries is maintained through graph-constrained

skeletal priors that prevent inter-part intersection [95]. These methods optimize layout via semantic sub-skeletons and employ skinning-weight driven fusion to suppress interface artifacts that typically arise from voxel-level interpolation.

Stabilization Strategies for Gaussian Parameter Optimization. To counteract initialization sparsity and prevent optimization collapse in under-sampled regions, recent approaches employ density-weighted selection coupled with perturbation modules that enforce uniform spatial coverage [17]. Architectural decoupling further stabilizes amortized training by separating geometry and texture prediction into distinct network branches, thereby avoiding local minima inherent to intertwined attribute updates [177]. Temporal consistency is maintained through multi-timestep gradient aggregation and contribution-based filtering, which balance semantic richness against shape coherence to mitigate scene-wide visual artifacts [76]. Finally, direct manipulation of mean positions and covariance matrices enables precise pose editing and object fusion without necessitating global re-optimization of the entire scene representation [200].

Hierarchical Decoding for Scale. Native hierarchical decoding architectures synthesize multi-resolution Gaussian structures directly from compact latents, eliminating the bottlenecks associated with post-hoc downsampling for planetary-scale streaming [113]. These systems employ inherent multi-level-of-detail decoders alongside sliding-window inference to manage trillion-scale primitive streaming while maintaining seamless spatial blending and real-time rendering latency. To resolve inconsistencies in unbounded scenes, canonical query initialization coupled with global scale normalization enables stable variational compression of massive Gaussian sets [43]. This approach utilizes regular voxel-grid queries and cross-attention tokenization with Fourier positional encoding, minimizing KL divergence through scene-level normalization to ensure coherent representation across vast spatial extents.

Explicit Texture Mapping and Material Disentanglement. To prevent specular baking, two-stream attention mechanisms within diffusion backbones enforce strict separation between intrinsic albedo and metallic-roughness properties [42, 121]. These architectures employ dual-encoder conditioning and rectified flow objectives to isolate diffuse components from illumination effects during PBR parameter estimation. Addressing topological continuity, coordinate-conditioned inpainting operates directly in UV space to resolve occlusion artifacts and ensure seamless tex-

ture synthesis across manifold boundaries [42, 103]. This approach utilizes differentiable rasterization to guide occlusion-aware completion, enabling explicit domain editing without global re-projection. Furthermore, semantic segmentation of UV layouts facilitates localized material optimization by partitioning surfaces into independent editable layers [103, 121]. Such region-specific refinement allows for landmark-guided reconstruction and asset matching, supporting strand-based hair synthesis and distinct surface modifications without requiring full model re-synthesis.

3.2.6 Discussion

Resolving the fundamental tension between high-fidelity synthesis and real-time rendering efficiency has driven a decisive migration from pure implicit volumetric representations toward hybrid and explicit geometries [75, 113, 223]. This structural evolution facilitates critical capabilities such as NeRF-to-Gaussian transitions, UV-mapped albedo disentanglement, and screen-space shadow prediction while ensuring 360° renderability and efficient LOD streaming. Concurrently, autoregressive and diffusion paradigms increasingly operate within compact 3D latent spaces rather than on 2D view sequences to enforce global spatial consistency and eliminate multi-view redundancy [82, 102, 166, 206]. By employing techniques like view-decoupled token aggregation and flow matching in latent space, these approaches achieve significant complexity reductions, often approaching $O(n^6)$ savings, while mitigating exposure bias through residual quantization across scales. Furthermore, the strategic decoupling of geometry and appearance via parametric priors or semantic masks enables robust generalization from limited 3D data without sacrificing editability for downstream applications [25, 37, 204]. Methods utilizing silhouette-based consistency losses, density-prior guided rendering, and FLAME mesh alignment demonstrate how two-stage cascaded generation and semantic volumetric back-projection can maintain topological integrity even under sliding-window inference constraints.

3.3 2D Prior-Guided Lifting

This section examines the paradigm of lifting 2D visual data into 3D representations by exploiting priors from pretrained foundation models. We first analyze optimization strategies that distill diffusion knowledge into 3D geometries through gradient-based guidance in Sec. 3.3.1. Subsequently, Sec. 3.3.2 addresses the enforcement of geometric coherence across synthesized views using epipolar constraints and shared latent spaces. The discussion then extends to specific

input modalities, covering single-image reconstruction techniques in Sec. 3.3.3 and the synthesis of dynamic 4D scenes from video sequences in Sec. 3.3.4. Finally, Sec. 3.3.5 investigates methods that decouple texture and geometry optimization to enhance control over surface appearance and structural fidelity.

3.3.1 Score Distillation Sampling

Gradient Approximation via 2D Priors. The optimization framework circumvents the intractability of direct 3D score estimation by distilling gradients from pre-trained 2D diffusion priors into explicit volumetric representations such as NeRF and 3D Gaussian Splatting [211]. To mitigate instability arising from sparse initializations, hybrid pipelines incorporate coarse geometric structures generated by 3D diffusion models like Point-E or Shap-E before refining details through 2D score guidance [17, 194]. These approaches employ localized densification and perturbation strategies, including density-weighted selection and noisy point growing, to ensure uniform spatial coverage and stable gradient propagation. Such mechanisms facilitate the recovery of high-frequency textures by enforcing edge sharpening and densest-color assignment within specific regions, effectively bridging the gap between coarse 3D priors and fine-grained 2D supervision.

Gradient Aggregation for Variance Reduction. To suppress instability during early optimization stages, distillation signal stability is enhanced by aggregating gradients across multiple timesteps and applying spatial filtering, which reduces high-frequency noise that typically drives estimator variance [76]. This formation pattern sampling combined with 3D Gaussian filtering lowers gradient variance, facilitating more stable reconstructive generation. Convergence reliability is further improved via progressive guidance strategies that dynamically adjust updates based on visible surface regions [76, 104]. Such mechanisms rely on incremental three-stage sampling to modulate gradient magnitude without invoking external geometric constraints or collision detection.

Regularization and Noise Scheduling for Structural Stability. Stabilizing the optimization trajectory requires explicit geometric constraints that function as static regularization terms, allowing the distillation process to prioritize high-frequency texture refinement without diverging into structural artifacts [110]. Techniques such as disparity regularization, normal consistency, and proxy mesh grounding serve to anchor the underlying structure, effectively separating geometry updates from appearance learn-

ing. Complementing this spatial stabilization, deterministic noise scheduling via inversion techniques replaces stochastic sampling to maintain gradient coherence across iterations [17, 110]. By employing fixed-step noise schedules derived from DDIM inversion or sparse point cloud initializations, these methods ensure that texture updates remain aligned with the established coarse geometry, thereby sharpening edges while preserving global form.

Variance Reduction in Sparse-View Score Estimation. Under sparse view conditions, the Score Distillation Sampling gradient estimator suffers from high variance and mode collapse due to insufficient angular coverage, causing optimization to stagnate in local minima. To counteract this instability without relying on external geometric templates, recent methods modify the distillation objective itself to reduce gradient noise inherent in the score function. Techniques such as voxel-aligned feature aggregation integrate sparse observations into a shared latent space through average pooling, effectively smoothing the gradient landscape across unseen viewpoints by altering the score approximation rather than reconstructing geometry [220]. This architecture incorporates learnable residuals within a sparse 3D decoder to suppress estimation noise while preserving high-frequency details. Concurrently, dual self-supervised objectives balance reconstruction fidelity against adversarial realism, ensuring stable convergence by regularizing the score distribution specifically, instead of enforcing rigid structural constraints typical of multi-view consistency pipelines [220].

Semantic Alignment in Distillation. Conditioning the diffusion prior on explicit geometric depth maps enforces strict structural correspondence between the rendered 3D proxy and the synthesized 2D guidance signal, effectively grounding the generation process in physical shape constraints [125]. To mitigate the risk of over-constraint artifacts while preserving this alignment, dynamic attenuation mechanisms modulate visual conditioning strength based on feature magnitude, allowing the optimization to adaptively balance fidelity and flexibility [62]. Complementing these geometric controls, multi-agent textual decomposition frameworks ensure that generated prompts remain semantically consistent with specific geometric attributes, thereby preventing the introduction of hallucinated content during the distillation phase. Furthermore, disentangling identity embeddings from shape parameters enables the optimization to maintain geometric fidelity independent of texture variations within the distilled gradient field, reducing bias when handling multiple identities per mesh [125].

Table 3 Taxonomy of representative methods in **2D Prior-Guided Lifting**. This table presents a taxonomy of 2D prior-guided lifting methods, shows distinct reconstruction strategies and optimization objectives, and highlights the shift toward explicit scene representations.

Method	Venue	Lifting Strategy	Optimization Paradigm	Highlight
Score Distillation Sampling (Sec. 3.3.1)				
GaussianDreamer [194]	CVPR'24	Point Cloud Init	Score Distillation	Initializes Gaussians via 3D diffusion priors and noisy point growing to accelerate SDS convergence.
DreamScene [76]	ECCV'24	Object-Environment Split	Score Distillation	Formation Pattern Sampling aggregates multi-timestep gradients with Gaussian filtering to stabilize scene optimization.
GGDS [110]	AAAI'25	Proxy Mesh Distill	DDIM Inversion	Geometry-Grounded Distillation Sampling employs DDIM inversion for consistent noise scheduling across non-overlapping views.
Multi-View Consistency Optimization (Sec. 3.3.2)				
One-2-3-45++ [88]	CVPR'24	Multi-View Tiling	Supervised Diffusion	Tiles sparse views for shared attention, enforcing inter-view consistency via projected 3D feature volumes.
LayerPano3D [187]	SIGGRAPH	Depth Layering	LoRA Fine-Tuning	Decomposes panoramas into depth layers, resolving occlusion conflicts through iterative alignment and Gaussian selection.
See3D [97]	CVPR'25	Visual Warping	Pose-Free Diffusion	Derives time-dependent visual conditions from pixel noise to learn 3D priors without explicit camera poses.
Single-Image to 3D Reconstruction (Sec. 3.3.3)				
Compress3D [201]	ECCV'24	Triplane Latents	Conditional Diffusion	Learnable 3D-to-2D projection and dual-conditioning resolve depth ambiguity for high-fidelity mesh reconstruction.
IDOL [224]	CVPR'24	UV-Token Mapping	Feed-Forward Regression	UV-space Gaussian prediction on canonical templates hallucinates unseen geometry while ensuring immediate animatability.
Single-Image to 3D World [129]	ICCV'25	Panorama Lifting	In-Context Learning	Anchored panorama synthesis and forward-backward warping resolve occlusions to lift single images into consistent scenes.
Video-to-3D Dynamic Synthesis (Sec. 3.3.4)				
See3D [97]	CVPR'25	Visual Warping	Pose-Free Diffusion	Time-dependent visual conditions replace explicit poses to learn 3D priors from unposed video sequences.
WonderVerse [41]	arXiv'25	Video Extension	Zero-Shot Inference	Camera trajectory monitoring detects geometric inconsistencies, triggering regeneration to ensure temporal coherence in extended scenes.
Lyra 2.0 [135]	arXiv'26	Autoregressive Cache	Flow Matching	Decoupled geometry caching and canonical coordinate warping mitigate autoregressive drift while preserving temporal consistency.
Texture and Geometry Decoupling (Sec. 3.3.5)				
Gen2Det [150]	arXiv'23	Grounded Inpainting	Synthetic Augmentation	Grounded inpainting regenerates instances within layouts, separating object appearance from background structure via dual-level filtering.
Dreamland [106]	arXiv'25	Layered Abstraction	Sim2Real Fine-Tuning	Layered World Abstraction decouples traffic and map geometry from backgrounds, allowing instructional editing to refine texture independently.

3.3.2 Multi-View Consistency Optimization

Epipolar-Driven Discrimination. Recent frameworks enforce strict alignment across synthesized perspectives by embedding epipolar constraints directly into adversarial objectives and attention mechanisms [4, 39]. These architectures integrate epipolar line projections into discriminator loss functions, penalizing feature activations that deviate from geometrically consistent trajectories across views. Attention maps are further restricted through epipolar masking, ensuring that feature aggregation occurs only along valid corresponding lines rather than arbitrary spatial regions. Explicit structural operators validate rendering outputs against epipolar planes, rejecting generations that violate multi-view geometric consistency. Such direct integration of epipolar geometry effectively mitigates object omission and stabilizes

cross-view coherence, yielding measurable improvements in structural similarity metrics through intrinsic geometric regularization.

Latent Space Mechanisms for Texture Coherence. Synchronized joint diffusion architectures enforce texture coherence by iteratively refining 2D reverse trajectories through explicit 3D geometric feedback rendered at every timestep, utilizing asymmetric U-Net transformers to inject shape priors and correct inconsistencies [182]. To globally constrain inter-view feature consistency, multi-view tiling strategies consolidate sparse viewpoints into a single composite input that enables shared self-attention mechanisms across the entire visual field [88]. This spatial conditioning framework further projects local 2D patch features into a dense 3D feature volume via sparse convolution, offering superior geometric fidelity com-

pared to global embedding methods. Furthermore, aggregating multi-view information into fixed-length view-invariant latent tokens eliminates cross-view redundancy, allowing flow matching to synchronize texture attributes across views within a unified latent domain while supporting spatial extrapolation [166].

Pose Refinement and Depth Alignment. Discontinuities in estimated camera trajectories derived from Structure-from-Motion pipelines serve as primary indicators of geometric inconsistency within generated sequences. By monitoring these trajectories, systems can trigger segment regeneration when positional shifts or rotational deviations exceed predefined tolerances, effectively stabilizing the view synthesis process [41]. Complementing this trajectory monitoring, iterative sparse pixel-wise depth alignment addresses warping-induced artifacts without relying on explicit camera pose annotations during training. This approach utilizes offline depth estimators within a correction loop to resolve scale ambiguities inherent in warping-based generation, ensuring coherent geometry across views [97].

Occlusion and Scalability Constraints. Layered decomposition strategies mitigate semantic drift in occluded regions through panoptic segmentation and K-Means clustering, yet introduce geometric artifacts when pre-trained depth priors fail to align distinct scene planes during Flux-Fill inpainting [187]. Recursive expansion frameworks maintain global coherence via point cloud updates but struggle to recover valid geometry in highly complex occlusion scenarios due to overly restrictive depth consistency checks that induce false negative geometry exclusion [39]. Both paradigms exhibit scalability bottlenecks where synthesis quality degrades as the number of required layers or recursive steps exceeds optimal thresholds, limiting exploration ranges and imposing significant computational overhead [39, 187].

3.3.3 Single-Image to 3D Reconstruction

Architectural Shifts to Feed-Forward Inference. Replacing iterative optimization with end-to-end feed-forward networks eliminates the need for score distillation sampling, compressing inference latency from minutes to sub-second scales [178, 180]. These supervised frameworks often employ rectified flow objectives and frozen triplane decoders to achieve speeds around 0.148s, representing a roughly $600\times$ acceleration over prior methods. To stabilize amortized training and prevent attribute divergence, cascaded pipelines decouple geometric structure prediction from texture refinement using hybrid intermediate representations [177, 180]. Such architectures inte-

grate point-voxel convolutions and learnable 3D anchor queries at fixed canonical scales, occasionally followed by latent space super-resolution modules. However, the architectural capacity required for real-time generation introduces bottlenecks; fixed query counts or frozen backbones inherently constrain representation fidelity [177, 178, 180]. While these design choices enable memory efficiency and mitigate input-view bias in compact models (e.g., 280KB), they simultaneously impose hard limits on the maximum reconstructible Gaussian count and rely heavily on pre-trained large reconstruction model dependencies.

Data-Centric Strategies for Resolving Depth Ambiguity and Backside Hallucination. Scaling training corpora through high-fidelity synthetic generation from video diffusion models effectively resolves pose-dependent depth ambiguities while enhancing generalization to diverse clothing geometries [151, 224]. By leveraging pose-controllable video diffusion and multi-view skeletal annotation, frameworks such as HuGe100K ensure 24-view consistency and utilize semantic-segmented background projection to create robust non-articulated models derived from PIFu architectures. Complementing this, lifting single views to pseudo-3D point clouds via pre-trained depth estimators enables scalable self-supervised pre-training even without ground-truth multi-view geometry [19]. This approach employs sparse weight indexing and voxel-based tokenizers optimized through hybrid reconstruction losses, including Chamfer Distance and binary cross-entropy for occupancy, following offline depth distillation. To further bridge the domain gap between synthetic isolation and real-world occlusion patterns, iterative model-in-the-loop data engines combine retrieval ensembles with human preference ranking [155]. These systems integrate Direct Preference Optimization and cascaded training recipes that incorporate semi-synthetic occlusion, human ranking feedback, and sparse latent flow matching anchored by DINOv2 image encodings. Finally, replacing hand-crafted assets with deep learning-reconstructed humans mitigates the realism discrepancy inherent in traditional CGI while automating pose annotation for non-articulated meshes [151]. Such pipelines utilize pixel-aligned implicit functions and OpenGL perspective projection alongside OpenPose integration within mixed real-synthetic training sets to significantly boost face verification precision.

Compositional Pipelines for Large-Scale Scene Synthesis. Hybrid generative-procedural frameworks decouple 2D layout synthesis from 3D assembly by projecting casual sketches into isometric reference views to enable precise footprint estimation through sketch-aware losses and step-unrolled denoising [179].

To address object-to-scene domain shifts, modular approaches decompose top-down inputs into overlapping latent segments conditioned on monocular spatial priors, utilizing masked rectified flow and sequential latent fusion for zero-shot inference [214]. Global coherence in large environments is further enforced via two-stage strategies that employ anchored panorama heuristics and forward-backward warping before lifting pseudo-3D point clouds through metric depth alignment and trainable grid distortion modules [129]. Hierarchical voxel enhancement pipelines complement these methods by aligning 2D segmentation trees with frozen diffusion backbones via attention-based cross-modal matching, facilitating coarse-to-fine geometric retrieval through recursive refinement and shape-based scale estimation [198].

Latent Space Compression and Representation Trade-offs.

To mitigate geometric information loss during latent space compression, recent approaches replace fixed pooling operations with learnable 3D-to-2D projection layers within triplane autoencoders [201]. These architectures utilize point-to-triangle mapping and 3D-aware cross-attention to query feature volumes, effectively compressing colored point clouds while preserving structural fidelity. Distinct from pure visual encoding, other methods resolve simulation stability issues by decomposing scene representations into separate physical collision meshes and visual Gaussian splatting layers [208]. This separation enables alternating SDF-physics refinement and penetration loss minimization, ensuring gravity-consistent frame alignment through layered rendering compositing rather than unified latent encodings.

3.3.4 Video-to-3D Dynamic Synthesis

Disentangling Camera and Object Dynamics via Explicit 3D Representations. Structured volumetric decompositions facilitate the separation of foreground motion from background deformation by encoding scenes into distinct latent components. Methods employing deformable UV-sphere meshes alongside voxel grids utilize global latent vectors and presence indicators to isolate temporal pose dynamics within a unified framework [51]. To bypass the need for explicit camera annotations, alternative approaches replace static conditioning with time-dependent visual cues, learning generic 3D priors directly from unposed video sequences. These pipelines process masked noisy latents with timestep-varying weights and pixel-space hints, effectively bridging domain gaps through warping-based generation mechanisms [97].

Temporal Coherence and Drift Correction. To mit-

igate accumulated drift in autoregressive pipelines, recent approaches inject self-corrupted historical latents during training, forcing the model to recover from prediction errors via flow-matching noise corruption and denoised prediction feedback [135]. This self-augmentation strategy effectively reduces train-test discrepancies and suppresses artifact propagation across extended sequences. Beyond diffusion-based corrections, reinforcement learning objectives optimize bidirectional optical flow cycle consistency, enabling geometry-grounded agents to correct hallucinations through rendered supervision without explicit ground-truth reconstruction [148]. Complementary iterative generative repair mechanisms further refine texture resolution and suppress view-inconsistent artifacts emerging from initial panoramic synthesis steps by employing robustness modules within a FLUX-based optimization loop.

Trajectory Continuity Verification. To maintain geometric fidelity across extended temporal horizons, recent approaches monitor camera trajectories derived from Structure-from-Motion pipelines for abrupt discontinuities that signal reconstruction failures [41]. This verification process automatically flags aberrant segments when positional deviations or rotational shifts exceed predefined qualitative thresholds, indicating a break in motion consistency. Upon detection, the system triggers an iterative regeneration loop specifically for the identified faulty intervals, thereby correcting trajectory drift without requiring full sequence re-computation.

3.3.5 Texture and Geometry Decoupling

Architectural Disentanglement via Layered Abstractions. Structural decoupling is achieved by partitioning the scene into independent semantic strata, separating static map geometry from dynamic traffic and editable background elements to enable component-wise optimization [106]. This framework preserves geometric fidelity by retaining physics-based simulator outputs for structural layers while restricting generative priors to surface appearance through an instructional editing module that bridges the domain gap without altering underlying forms. A latent projection adapter subsequently integrates these disentangled multi-modal conditions directly into the noise space of a pretrained diffusion model, ensuring that the synthesized output adheres strictly to the imposed structural constraints.

Artifact Suppression via Selective Masking. To mitigate the coupling of spurious geometry with surface texture, recent approaches modify the optimization objective to exclude regions prone to hallucina-

tion. Gen2Det implements a background ignore loss that dynamically masks high-confidence non-ground-truth areas during gradient computation, effectively preventing the model from fitting synthetic noise to geometric structures [150]. This confidence-based thresholding isolates valid signal from artifacts, ensuring that texture updates do not erroneously deform the underlying mesh in ambiguous regions.

3.3.6 Discussion

Despite the prevalence of diffusion-based priors, the fundamental mismatch between 2D image statistics and 3D geometric reality often manifests as structural artifacts, including texture blurring and normal map inconsistencies that iterative refinement struggles to fully correct [17, 84, 183]. To address these fidelity gaps and the prohibitive computational cost of optimization, recent architectures have pivoted toward feed-forward mechanisms that predict bundled 3D representations in a single pass, utilizing specialized attention modules to enforce multi-view coherence without explicit gradient descent [39, 97, 155]. While hybrid strategies incorporating video foundations or simulation data offer a potential route to 3D-native consistency, they remain constrained by domain gaps and the difficulty of aligning dynamic motion with static scene assumptions [41, 52, 106, 150].

3.4 Language and Agent-Driven Composition

This section examines the integration of Large Language Models and agentic workflows to orchestrate complex 3D scene construction through high-level semantic directives. We first analyze methods that translate natural language into precise spatial configurations and bounding box layouts in Sec. 3.4.1, followed by an exploration of multi-agent systems where specialized entities collaborate to plan and verify scene consistency in Sec. 3.4.2. The discussion then shifts to approaches utilizing executable code and API calls for controlling 3D engines and procedural tools as detailed in Sec. 3.4.3, alongside techniques for retrieving and assembling existing assets from databases based on semantic queries in Sec. 3.4.4. Finally, we address interfaces that enable users to refine generated environments through iterative natural language feedback or direct manipulation in Sec. 3.4.5. Together, these subsections illustrate a transition from static generation to dynamic, controllable, and reasoning-capable 3D content creation pipelines.

3.4.1 Text-to-Layout and Spatial Reasoning

Semantic-to-Geometric Mapping via Intermediate Representations. Decoupling semantic parsing from geometric synthesis allows large language models to generate explicit bounding box parameters rather than relying on implicit latent codes. By enforcing structured intermediate formats, such as CSS-styled strings or coordinate tuples, these pipelines translate natural language descriptions directly into spatial layouts with clear semantic labels [85, 226]. Alternative approaches bypass lossy textual abstraction by injecting raw point cloud geometry directly into the model context through native 3D tokenization. Techniques utilizing SPFormer encoders or PointBERT prefixes reduce the modality gap by applying geometry-aware attention masks to the input stream, while simultaneously employing hierarchical mechanisms that separate global anchoring from local refinement via discrete anchor tokens and residual regression vectors [160]. This coarse-to-fine decoding strategy stabilizes autoregressive generation without requiring redundant architectural descriptions.

Constraint Satisfaction and Collision Avoidance.

Spatial relational tuples are resolved through depth-first search solvers and autoregressive mechanisms that explicitly enforce collision avoidance via GPT-4 prompting strategies [8, 189]. This formulation treats layout synthesis as a problem of satisfying soft and hard logical rules, applying hard-constraint searches on 2D centroid arrangements to guarantee functional quantity constraints prior to asset placement. Furthermore, hierarchy-enhanced metadata parsing facilitates precise part-level functional filtering, ensuring strict adherence to quantitative and topological requirements before the final geometric solving stage [29]. Such validation relies on explicit geometric overlap checks and coordinate consistency verification, maintaining focus on spatial constraints without invoking semantic embedding metrics or external retrieval protocols.

Learning Functional Placement Priors. Supervised fine-tuning on datasets containing instance-level natural language annotations allows models to internalize functional spatial priors, dictating context-appropriate object positioning such as placing lamps on tables rather than floors [219]. This data-driven approach implicitly encodes commonsense relational reasoning during training, enabling the direct synthesis of semantically coherent arrangements by learning distributional patterns of typical usage instead of relying on explicit geometric search algorithms [218]. Structured chain-of-thought supervision complements these approaches by forcing explicit deduction of

Table 4 Taxonomy of representative methods in **Language and Agent-Driven Composition**. This table compares generation methods by control mode and optimization strategy, showing diverse agent architectures and highlighting the shift toward iterative verification loops.

Method	Venue	Control Mode	Optimization Strategy	Highlight
Text-to-Layout and Spatial Reasoning (Sec. 3.4.1)				
3D-GRAND [186]	CVPR’25	Scene Graph Grounding	Supervised Fine-Tuning	LLM-driven pipeline generates densely-grounded 3D annotations to suppress object hallucination in spatial grounding tasks.
SceneTeller [226]	ECCV’24	CSS Layout Coding	Diffusion Stylization	LLM generates CSS-formatted bounding boxes via in-context learning to enforce global spatial consistency before geometry synthesis.
HOLODECK 2.0 [8]	arXiv’25	VLM-Guided Layout	Iterative Constraint Refinement	VLM-guided DFS solver iteratively refines spatial constraints to ensure physically coherent layouts and prevent collisions.
ASSIST-3D [218]	AAAI’25	Synthetic Data Layout	Collision-Free DFS	LLM-guided depth-first search strategy ensures collision-free placement and physical plausibility during sequential object layout generation.
NaLA [160]	arXiv’26	End-to-End Tokens	Coarse-to-Fine Regression	Native 3D token injection with coarse-to-fine pose prediction prevents autoregressive confusion and ensures precise collision avoidance.
Multi-Agent Collaborative Workflows (Sec. 3.4.2)				
FilmSceneDesigner [175]	AAAI’25	Agent Chaining	Finite State Machine	Manager agent coordinates sequential sub-agents via finite state machine to enforce spatial constraints.
SimWorlds [86]	arXiv’26	Staged Code Generation	Deterministic Verification	Staged coder-verifier-reviewer agents enforce layered scene protocols through deterministic runtime state inspection.
EmbodiedGen V2 [164]	arXiv’26	Affordance Labeling	Physics Validation	Hierarchical quality gating and task-conditioned routing decouple semantic reasoning from geometric constraint solving.
Code-Based and API-Driven Generation (Sec. 3.4.3)				
T2P [212]	CVPR’23	Parameter Translation	Evolutionary Search	Evolutionary search optimizes discrete assets while neural imitators approximate rendering for differentiable parameter tuning.
3D-GPT [144]	3DV’25	Function Dispatch	Zero-Shot Chaining	Three-agent decomposition separates planning, enrichment, and scripting to generate executable Blender Python code without fine-tuning.
RoboSimGS [210]	RA-L’25	Hybrid Reconstruction	Photometric Gradient Descent	MLLMs infer kinematic structures from images to generate articulated URDF assets alongside hybrid 3DGS-mesh scene representations.
Modular Multi-LLM [195]	CAI’26	Modular Retrieval	RAG-Enhanced Scripting	Modular three-stage pipeline synthesizes validated Unreal Engine Python scripts using retrieval and domain-specific knowledge injection.
Retrieval-Augmented Asset Assembly (Sec. 3.4.4)				
HOLODECK [189]	CVPR’24	Constraint Generation	Depth-First Search	Retrieves Objaverse assets via semantic search, optimizing placement with LLM-generated relational constraints.
SceneX [217]	arXiv’24	Script Synthesis	Sequential Prompting	Orchestrates hybrid retrieval using CLIP embeddings for static models and text matching for procedural APIs.
SCAN2SIM [196]	CVPR’25	Retrieval Alignment	Contrastive Matching	Aligns multi-modal queries to 3D point clouds via contrastive learning, finalizing assembly through heuristic pose adjustment.
Interactive and Controllable Editing (Sec. 3.4.5)				
LI3D [85]	arXiv’23	Layout Interpretation	Visual Feedback Loop	LLM parses text to layouts; LLaVA verifies renders to iteratively refine bounding boxes via rule-based operations.
Scene Copilot [115]	arXiv’24	Code Translation	Human-in-the-Loop	Scene Codex translates text to Python commands while BlenderGPT facilitates iterative editing through textual scene condensation.
HOLODECK 2.0 [8]	arXiv’25	VLM-Guided Layout	Iterative Constraint Refinement	VLM-driven DFS solver refines spatial constraints for layout edits, supporting global rearrangement and local style replacement.

situational constraints prior to localization, demonstrating that reasoning quality often outweighs data volume for complex spatial queries [54].

Relational Reasoning Fragility. Although dense grounding supervision effectively curbs random object hallucination, model precision deteriorates sharply when queries involve adversarial sampling or frequent object co-occurrences [186]. Evaluations on the 3D-POPE benchmark reveal that while performance remains stable under random splits, accuracy drops significantly in adversarial and popular splits where models rely on statistical biases rather than

genuine spatial understanding. This vulnerability indicates that current noun-phrase grounding mechanisms struggle to disentangle true relational constraints from common contextual priors in complex scenes.

Stochastic Diversity in Diffusion Priors. Injecting stochastic semantic distributions directly into the diffusion prior enables the synthesis of divergent layout topologies from identical textual inputs, effectively decoupling prompt specificity from geometric output [31]. Scaling such variance requires automated scene graph generation, where programmatic conversion of

single prompts into massive batches of unique spatial configurations drives training environment variety [27]. Empirical analysis confirms a direct correlation between the volume of generated scenes and policy robustness, indicating that prompt-induced geometric variance serves as the primary mechanism for reducing the in-distribution to out-of-distribution performance gap in narrow settings.

3.4.2 Multi-Agent Collaborative Workflows

Architectural Decomposition and Role Specialization. Functional decomposition segregates high-level intent interpretation from low-level geometric parameterization to mitigate semantic ambiguity inherent in procedural generation [144, 195]. This separation enables distinct agents to handle task dispatch, conceptualization, and attribute enrichment independently. Domain-specific workflow emulation further structures these interactions into sequential stages that enforce professional design constraints and historical fidelity [175]. By utilizing finite state machines and turn-taking protocols, such systems implement four-stage chaining to maintain structural coherence and cinematic fidelity without relying on monolithic control loops. Specialized visual-textual roles within these multi-agent frameworks align semantic prompts with geometric features to prevent conditioning artifacts [62]. Architectures deploying seven-agent VLM configurations assign specific duties such as subject identification, pose extraction, and scene composition, optimizing perceptual similarity through metrics like DreamSim.

Iterative Refinement via Visual Feedback and Auto-Verification. Closing the generation loop requires visual perception modules to analyze rendered outputs, identifying geometric and material discrepancies against input specifications through panoramic image analysis and anatomical accuracy checks [89, 145, 207]. These perceptual critiques drive auto-verification agents that translate findings into executable corrective actions, utilizing Chain-of-Thought reasoning to adjust procedural parameters or latent variables without manual intervention [89, 90, 207]. To manage computational complexity during this cycle, hierarchical refinement policies decouple global scene consistency checks from local asset placement optimization, applying distinct strategies for scene-level coherence versus receptacle surface constraints [90, 145]. Subsequent self-improvement mechanisms utilize these high-reward action sequences to fine-tune policy parameters, balancing alignment maximization against latency trade-offs for future tasks.

Deterministic State Verification and Conflict Resolution. Bypassing visual heuristics, certain protocols directly inspect engine runtime states, such as modifier stacks and physics caches, to validate mechanism correctness through non-visual critique [86]. To ensure simulation readiness, hierarchical quality gating enforces contracts via generate-verify-retry loops that couple metric geometry with physical validity before exporting assets in formats like URDF or USD [164]. When spatial conflicts arise during assembly within a fixed layout, numerical optimization solvers operate strictly on these verified runtime states to minimize placement errors against hard geometric constraints by adjusting scaling and rotation parameters [31].

Coordination Protocols for Global Consistency. Distributed modular frameworks enforce scene-wide geometric coherence by decoupling object motion trajectories from rigid kinematic constraints through specialized agent communication protocols [162]. Rather than relying on internal generative loss functions, these systems prioritize translation-aware motion planning while maintaining point-wise distance preservation across generated instances via inter-agent constraint negotiation. The integration of Cross-instance Proximity Contact Alignment further ensures that local interactions do not compromise the global structural integrity of the assembled environment, securing consistency through coordinated parameter updates rather than monolithic model optimization.

3.4.3 Code-Based and API-Driven Generation

Architectural Protocols for API Encapsulation. To unify heterogeneous procedural plugins under a single management protocol, dynamic API conversion layers standardize interactions across diverse engine interfaces [207, 217]. This structural encapsulation resolves plugin heterogeneity by mapping distinct function signatures to a universal PCG protocol, ensuring consistent execution regardless of the underlying toolset. Systematic prompt templates further enforce strict parameter mapping by embedding role definitions, documentation constraints, and format specifications directly into the generation context [195, 217]. By integrating few-shot examples and dependency management rules within these schemas, such approaches facilitate the synthesis of executable scripts that adhere to rigid syntactic requirements. Modular decomposition isolates domain-specific knowledge injection from code synthesis to prevent attribute hallucination during parameter assignment [195].

Deterministic Verification via Runtime State In-

spection. Verification fidelity shifts from visual critique to direct inspection of engine runtime states, including modifier stacks and physics caches, to distinguish functional mechanisms from visually similar fakes [86]. Within this framework, staged construction pipelines isolate errors to specific build phases, enabling localized repair triggers based on protocol violations before advancing to subsequent generation steps. Beyond static checks, auto-verification agents establish dynamic feedback loops that iteratively tune procedural parameters to correct discrepancies between generated anatomical structures and input specifications [89].

Error-Driven Code Regeneration and Logic Repair. Natural language feedback loops enable localized code regeneration to rectify syntax errors or logical mismatches without re-initializing the entire generation pipeline [115, 144]. To mitigate failures arising from sparse user instructions, the system parses runtime error logs and exception traces to trigger targeted script repairs, bypassing the need for external planning modules. Beyond syntactic correction, dynamic auto-verification cycles evaluate functional constraints directly within the execution environment, driving iterative code refinement protocols that modify control flow or function arguments until the script executes without violation [90]. These mechanisms collectively shift the generation paradigm from single-pass execution to a closed-loop process where script validity is restored through successive, log-guided revisions rather than numerical parameter tuning.

Code-Driven Parameter Optimization. To resolve modality-specific discontinuities within executable scripts, hybrid frameworks decouple continuous parameter tuning via differentiable approximation from discrete argument selection through gradient-free evolutionary search [212]. These systems employ a neural imitator to align latent representations with prompt tuning, utilizing snapshot ensemble initialization to mitigate gradient discontinuities while maximizing alignment scores. Distinct from direct mesh synthesis, multi-modal large language models infer latent physical properties and kinematic structures to generate code that instantiates explicit simulation objects, thereby enabling high-fidelity interaction logic [210]. This approach facilitates the programmatic construction of URDF definitions by estimating density and inferring joint constraints, often integrated within a Real2Sim2Real pipeline that refines generated control scripts through iterative simulation feedback rather than rendering-based optimization.

3.4.4 Retrieval-Augmented Asset Assembly

Multi-modal Semantic Alignment and Indexing.

Aligning heterogeneous query modalities with 3D geometric representations requires contrastive learning frameworks that optimize retrieval accuracy across large-scale candidate pools [196, 217]. These approaches typically employ supervised contrastive losses alongside frozen vision-language encoders to map learnable point cloud features into hybrid text-to-text and text-to-image embedding spaces, where dot-product similarity scoring dictates selection. To resolve semantic ambiguity inherent in sparse textual descriptions, structured knowledge graphs integrate causal inference mechanisms that filter irrelevant shape priors through backdoor adjustment for confounding noise removal [215]. Such systems further utilize hierarchical transformer-based prior fusion and multi-entity search strategies to condition implicit decoders on refined context.

Functional and Hierarchical Constraint Filtering.

Beyond global semantic matching, retrieval systems parse hierarchical metadata to enforce fine-grained functional requirements, utilizing large language models to verify part counts and filter the candidate pool based on explicit object attributes rather than generating spatial layouts [29]. To resolve viewpoint inconsistencies when grounding open-vocabulary affordance queries to specific geometric regions, automated pipelines integrate chain-of-thought reasoning with multi-view voting aggregation, converting 2D masks into dense heatmaps for precise interaction point prediction [72]. Physical plausibility is further secured by evaluating geometric compatibility scores against pre-retrieved assets, ensuring candidates possess valid support surfaces and non-interfering bounds prior to any scene assembly [189].

Geometric Adaptation for Memory Constraints.

To satisfy strict hardware memory constraints for immediate deployment, real-time adaptation pipelines employ quadric edge collapse algorithms that reduce mesh vertex counts to targets near 1,000 while balancing visual fidelity against GPU memory reduction [7]. This simplification process ensures sub-minute loading latency essential for AR rendering optimization by prioritizing topological compression over aesthetic harmonization.

3.4.5 Interactive and Controllable Editing

Latent and Parameter Space Refinement. High-level natural language intents drive direct updates to an existing scene representation by modifying latent codes or specific object parameters rather than

generating initial layouts [85, 115]. These frameworks bridge the modality gap between vague user instructions and precise engine-specific attributes by computing gradients or offsets within the generative manifold, ensuring physical consistency during modification without serializing new structural definitions. To further refine these edits without continuous human intervention, automated feedback loops utilize vision-language models like LLaVA to verify rendered outputs against original descriptions through iterative adjustment cycles [85].

Latency Optimization via Incremental Mesh Updates. Achieving real-time interactivity within resource-constrained augmented reality environments necessitates architectures that bypass the inherent latency of full scene regeneration. These systems couple aggressive geometric simplification, specifically employing quadric edge collapse to reduce vertex counts to approximately 1,000, with differential rendering pipelines that compute only local changes [7]. By restricting computational updates to modified regions rather than re-evaluating global illumination or geometry, the pipeline shifts the burden from on-the-fly synthesis to efficient incremental mesh refinement. This architectural choice supports rapid update cycles, ensuring that user refinements remain responsive even on mobile hardware.

Graph-Based Structural and Latent Refinement. To resolve spatial ambiguities during iterative editing, specific architectures employ graph traversal strategies that update relational constraints directly without invoking procedural planning solvers [8]. Distinct from these topological adjustments, latent graph manipulation networks facilitate intent alignment by modulating shared node embeddings while utilizing adversarial relationship discriminators to correct deviations between user instructions and modified attributes [32]. This separation allows frameworks to decouple global structural reconfiguration from local asset style replacement, ensuring stylistic coherence without requiring full scene regeneration [8]. Furthermore, end-to-end graph-based approaches mitigate one-to-many mapping issues in scene manipulation by enforcing triplet recall metrics and constraint accuracy through the parallel encoding of layout and shape features [32].

3.4.6 Discussion

By shifting from static retrieval to generative and code-driven workflows, the field has effectively closed the semantic alignment gap, enabling open-domain scene synthesis that adheres to complex functional constraints through executable script generation and

hierarchical parsing [8, 29, 195]. This evolution toward agentic autonomy further refines physical plausibility, as multi-agent collaboration and iterative self-correction loops produce layouts with superior ergonomic validity compared to single-pass inference [31, 85, 90]. Despite these gains in controllability, persistent deficiencies in spatial reasoning and high hallucination rates necessitate hybrid architectures that couple symbolic logic with neural generation to guarantee geometric consistency [7, 8, 196].

3.5 Discussion

The persistent tension between strict procedural controllability and high neural fidelity has catalyzed a shift toward hybrid architectures that explicitly decouple geometric structure from appearance synthesis. By separating these concerns, modern pipelines employ node transpilation and layered world abstractions to maintain valid scene graphs while utilizing SDF-based meshing and UV-mapped albedo maps for detailed surface generation [75, 106, 117]. This structural disentanglement facilitates sim-to-real editing, allowing developers to modify physical properties without compromising the photorealistic textures generated by neural components. Concurrently, integrating large language models as spatial reasoners addresses the semantic alignment gaps found in purely generative or retrieval-based systems. Through iterative depth-first search solvers and multi-agent verification loops, these frameworks translate natural language queries into executable scripts that populate scenes with assets filtered via hybrid retrieval-augmented generation metadata [8, 90, 161, 195]. Such code-driven approaches ensure that spatial relationships adhere to logical constraints while maintaining high semantic relevance. Furthermore, efficiency bottlenecks inherent in score distillation and iterative optimization are being mitigated by transitioning from global image-level guidance to localized, object-centric adaptation strategies. Techniques such as density-weighted perturbation and mask-guided inpainting restrict computational overhead to relevant regions, while occlusion-aware reconstruction and gradient profile prior losses refine convergence rates [17, 112, 165]. These targeted optimizations, often supported by LoRA fine-tuning, enable rapid generation of complex scenes without sacrificing the consistency required for downstream simulation tasks.

4 Challenges and Prospects

This chapter critically examines the systemic barriers hindering the reliability and utility of generated 3D

data, moving beyond isolated technical failures to address broader macro-dimensions. Sec. 4.1 clusters these persistent limitations into distinct categories that affect evaluation rigor, physical grounding, and overall scalability. Following this analysis, Sec. 4.2 outlines a concrete research roadmap designed to overcome these hurdles and advance the field toward robust, simulation-ready data generation. Together, these subsections provide a structured perspective on both the current bottlenecks and the high-impact directions necessary for future progress.

4.1 Challenges

The reliability and utility of generated 3D data are currently constrained by a confluence of structural, visual, and functional limitations that hinder widespread adoption. We first examine fundamental geometric defects, including manifold violations and self-intersections, in Sec. 4.1.1, followed by an analysis of the domain gap between synthetic appearances and real-world distributions in Sec. 4.1.2. Significant barriers regarding computational overhead and memory demands are addressed in Sec. 4.1.3, while Sec. 4.1.4 investigates the persistent difficulties in achieving precise semantic alignment and controllability. Furthermore, we discuss the absence of physical plausibility and interaction capabilities in Sec. 4.1.5, concluding with a critique of the current scarcity of standardized evaluation metrics in Sec. 4.1.6. Collectively, these dimensions define the critical boundaries of current 3D generation methodologies.

4.1.1 Geometric and Topological Consistency

Neural generators frequently output meshes that violate manifold constraints, exhibiting self-intersections and disconnected components because optimization objectives prioritize visual fidelity over explicit topological regularization [134, 136]. The reliance on volumetric or implicit field representations often introduces high-frequency noise and ReLU-induced artifacts, leading to non-watertight surfaces with G0 continuity errors that break downstream simulation pipelines [65, 225].

- 1) **Manifold Violations from Implicit Biases.** Optimization processes operating on unsigned distance fields or density volumes lack inherent guarantees for surface watertightness, resulting in floating artifacts and blobby geometry where the zero-level set fails to form a closed shell [134, 136]. These topological defects stem from the inability of standard loss functions to penalize local field inconsistencies, leaving generated

assets riddled with holes and self-intersecting faces that require extensive post-processing [65].

- 2) **Global Incoherence in Multi-View Lifting.**

Score distillation strategies struggle to maintain geometric consistency across unseen viewpoints, causing hallucinated surfaces and depth collapse in occluded regions [211, 220]. The ambiguity inherent in lifting 2D priors to 3D space leads to structural contradictions where pseudo-ground-truth noise creates pseudo-surfaces that do not align with the global object structure [127, 129].

- 3) **Topological Smoothing via Data Sparsity.**

Domain-specific training sets with limited categorical diversity force models to average fine-scale structural discontinuities, erasing essential micro-scale details required for physical validity [96, 143]. Fixed feature resolutions and procedural extraction limits induce representation collapse, causing part decomposition errors and object splitting that distort the intended topology of complex assemblies [51, 202].

4.1.2 Photorealism and Domain Gap

The visual fidelity of synthetic data often fails to bridge the domain gap due to fundamental mismatches in texture statistics and lighting distributions that confuse downstream perception models. While geometric structures may appear coherent, the decoupling of appearance from layout introduces spurious correlations that degrade model generalization when transferring to real-world scenarios.

- 1) **Fragile Appearance-Geometry Decoupling.**

Forcing photorealistic textures onto synthetic layouts frequently degrades fine-grained structural adherence, as low-rank residual adapters struggle to separate domain identity from control signals [142, 185]. This sensitivity creates lighting biases and texture variability that preserve edge maps only superficially while introducing artifacts that mislead feature extractors.

- 2) **Upstream Error Propagation in Proxies.**

Monocular foundation models employed as 3D proxies inject upstream estimation noise into coherence assessments, which limits the discrimination variance for motion and geometry cues in generated sequences [13, 131]. Features from models like DINOv2 or UniDepth often suffer from detector noise and silhouette corruption, particularly failing during self-occlusion events where flow estimates become unreliable.

- 3) **Context Loss in Domain-Agnostic Representations.** Abstracting visual appearance through

point-based representations sacrifices critical scene context and view-dependent information required for robust policy transfer in cluttered environments [48, 152]. Camera-aligned sampling strategies often lead to body-part query confusion and left-right limb ambiguity, while sparse convolution limits prevent vision-language models from reliably interpreting minimal clothing constraints.

4) Calibration Instability in Hybrid Pipelines.

Hybrid rendering pipelines that composite physics-simulated masks with reconstructed backgrounds struggle to maintain dynamic object properties without manual metric calibration [44]. Dependencies on calibration boards and static scene reconstruction often result in generative geometric discrepancies, such as Gaussian ellipsoid distortion, when aligning SfM points via scaled ICP.

4.1.3 Computational Efficiency and Scalability

The rapid expansion of 3D generative modeling faces severe constraints imposed by the computational overhead of processing high-dimensional volumetric data, which restricts both training throughput and inference resolution. As architectures evolve to handle complex scenes, the demand for memory and processing power grows non-linearly, creating a bottleneck that prevents scaling to realistic environment sizes without sacrificing fidelity or speed.

1) Quadratic Scaling in Volumetric Attention.

Autoregressive and diffusion frameworks struggle with prohibitive costs when attending to long sequences of volumetric tokens, as attention mechanisms scale quadratically with sequence length while memory requirements grow cubically. This complexity forces practitioners to rely on compressed representations like tri-planes or sparse window partitioning, yet these approximations often introduce bottlenecks that limit the effective resolution and detail of generated assets [15, 102, 170, 197].

2) Latency in Unbounded Scene Expansion.

Methods designed for infinite or large-scale scene synthesis frequently depend on iterative block-wise inpainting and hierarchical pyramid cascades, which accumulate significant latency during the generation process. The necessity to fuse overlapping chunks or average velocity fields across coarse-to-fine stages leads to error propagation and visible seam artifacts at boundaries, fundamentally undermining the coherence of the expanded context [64, 78, 93, 172].

3) Capacity Limits of Efficient Alternatives. Alternative paradigms employing feed-forward sparse queries or procedural pre-training often fail to generalize due to fixed capacity constraints and inherent biases toward specific input views. These approaches frequently suffer from optimization instability and insufficient semantic grounding, particularly when reconstructing complex geometries that deviate from the convexity assumptions embedded in their bird’s-eye-view representations [24, 174, 180, 211].

4.1.4 Semantic Alignment and Controllability

Precise semantic alignment in 3D generation remains elusive due to fundamental disconnects between high-level linguistic intent and low-level geometric realization. Current architectures struggle to maintain attribute fidelity across the synthesis pipeline, often resulting in outputs that deviate significantly from user prompts or exhibit collapsed diversity. These failures stem from inherent biases in training data, limitations in spatial reasoning modules, and information loss during representation lifting.

1) Distributional Bias in Generative Backbones.

Text-to-3D models frequently exhibit severe category collapse, restricting output diversity to dominant classes found in training sets while failing on out-of-distribution concepts [121, 156]. This bias manifests as ethnic homogeneity in human avatars and rigid pose restrictions, effectively limiting the scale of identifiable identities [176, 223].

2) Deficient Spatial Reasoning in Language Models.

Large Language Models lack intrinsic geometric understanding, leading to physically implausible layouts when translating natural language into 3D coordinates [85, 217]. Such systems often hallucinate object relationships, violate common sense constraints, and fail to maintain consistent allocentric directions during scene graph extraction [161].

3) Information Loss in Representation Lifting.

The process of lifting intermediate 2D representations to 3D introduces deterministic bottlenecks that degrade multi-modal attribute synchronization [30, 45]. Linear geometric constraints and texture artifact propagation further exacerbate self-occlusion errors, causing a collapse in semantic distribution [183].

4) Fragility of Consistency Filtering. Automated mechanisms designed to enforce semantic consistency remain highly sensitive to threshold set-

tings, often necessitating strict volume limits that suppress layout diversity [150, 156]. These filters frequently fail to detect subtle semantic drifts, leading to degraded performance on small objects and unstable numerical predictions in downstream tasks [50, 59].

4.1.5 Physical Plausibility and Interaction

Current generative pipelines frequently decouple visual synthesis from physical grounding, treating mass, friction, and articulation as secondary attributes rather than intrinsic geometric constraints. This architectural separation results in assets that appear visually coherent yet fail to satisfy basic mechanical laws required for simulation or robotic interaction. The disconnect manifests across multiple stages of the generation process, from latent space formulation to final scene assembly.

- 1) **Decoupled Latent Parameterization.** Generative architectures often encode physical parameters as independent latent variables instead of deriving them from intrinsic geometry, leading to inconsistent correlations between mass, friction, and articulation properties [9, 164]. Methods employing dual-branch diffusion or conditional flow matching struggle to enforce these coupled constraints, while generate-verify-retry loops introduce computational overhead without guaranteeing kinematic diversity.
- 2) **Representation Discontinuity.** Hybrid scene representations typically segregate high-fidelity visual backgrounds from interactive foreground objects, creating sharp discontinuities at their interface [44, 208, 210]. Approaches combining 3D Gaussian Splatting for static environments with explicit meshes for dynamic entities face significant challenges during URDF export and metric scaling, often leaving floating artifacts that disrupt physics engine initialization.
- 3) **Simplified Stability Heuristics.** Procedural assembly and optimization-based methods rely on coarse stability checks that cannot capture complex material stress or multi-body dynamic constraints during synthesis [40, 66, 119]. Although techniques like Bayesian optimization and occupiability grids attempt to model gravity and context-aware segmentation, they frequently lack the fidelity to predict yield constraints or handle rollback strategies under realistic load conditions.
- 4) **Neglected Environmental Priors.** Data augmentation pipelines prioritize geometric oc-

clusion patterns while ignoring environment-dependent physical priors, resulting in object distributions that are synthetically valid but physically improbable [138]. Algorithms utilizing stochastic iterative merging or uniform random placement often fail to account for distance-dependent density adjustments, producing scenes where shape variability contradicts natural accumulation or stacking behaviors.

4.1.6 Evaluation Metrics and Benchmarking

The absence of unified evaluation standards creates a fundamental disconnect where high scores on automated metrics frequently fail to correlate with perceptual quality or functional utility in downstream applications. Current assessment protocols often prioritize easily computable statistical distances over holistic structural validity, leading to optimization trajectories that satisfy numerical benchmarks while producing geometrically incoherent or visually artifacts-laden outputs. This misalignment is exacerbated by the reliance on 2D proxies that cannot adequately capture the unique constraints of three-dimensional consistency, resulting in a fragmented landscape where methodological progress is difficult to verify objectively.

- 1) **Metric Blindness to Topology.** Standard distributional metrics and slice-wise comparisons lack differentiable mechanisms to penalize non-manifold edges or anatomical discontinuities, allowing models to generate assets that score well statistically despite failing basic structural sanity checks [53, 225]. These approaches project complex 3D structures into lower-dimensional latent spaces where critical continuity errors remain undetected because the evaluation pipeline itself cannot compute topological validity.
- 2) **Multi-View Inconsistency Blind Spots.** Optimization objectives derived from monocular priors frequently induce style dilution and view-dependent artifacts that standard image-based scores like CLIPScore cannot quantify within a volumetric context [13, 21]. The reliance on contrastive learning gaps and sparse view sampling creates a deceptive signal where temporal-spatial incoherence persists despite converging loss values, leaving the true 3D consistency of the generated asset unverified.
- 3) **Downstream Utility Disconnect.** Proxy metrics optimized for detection accuracy often neglect the specific physical constraints required for simulation, creating a scenario where high instance embedding contrast does not translate

to successful policy transfer [70, 98]. Layout optimization biases and frequency domain cues may improve specific benchmark numbers while simultaneously degrading occlusion handling and scale variance, rendering the generated data ineffective for training robust real-world systems.

4.2 Prospects

The future trajectory of 3D data generation points toward integrated architectures capable of unifying diverse tasks and representations within single frameworks, as detailed in Sec. 4.2.1. Concurrently, the field must pivot from producing static assets to generating simulation-ready content that supports embodied AI training and dynamic 4D scene evolution, which are examined in Sec. 4.2.2 and Sec. 4.2.3 respectively. Advancing these technical capabilities requires new collaborative workflows that embed human feedback directly into the creation loop, a topic explored in Sec. 4.2.4. Finally, realizing this potential demands rigorous attention to ethical safeguards and the establishment of standardized evaluation protocols to ensure reliability and safety, as discussed in Sec. 4.2.5 and Sec. 4.2.6.

4.2.1 Unified 3D Foundation Models

Next-generation frameworks will likely decouple understanding and generation modules to preserve distinct pretrained priors rather than forcing a single architecture to master both objectives [23]. By treating heterogeneous data types as unified discrete sequences within shared vocabularies, interleaved cross-modal training strategies can enforce geometric consistency across diverse inputs without requiring signal quantization [114, 191]. Furthermore, scaling laws for these foundation models depend on transforming volumetric representations into compact latent vectors, effectively reducing the computational burden of sequence length growth from cubic to quadratic scaling while maintaining structural fidelity through capacity expansion in decoder-only transformers.

Scalability for next-generation architectures will likely depend on repurposing the vast accumulation of user-generated two-dimensional content as a surrogate dataset, effectively bypassing the scarcity of native three-dimensional datasets through cross-modal transfer learning [73]. Parallel to this data strategy, analytical synthesis pipelines are expected to generate universal volumetric data by systematically varying structural forms and surface properties to achieve zero-shot domain adaptation without manual annotation [12]. Frozen video diffusion backbones may subsequently function as implicit geometric oracles,

where iterative refinement processes activate latent spatial priors to extract multi-view consistency for static understanding tasks without requiring explicit three-dimensional supervision [171]. To harmonize these heterogeneous inputs, unified world models will probably adopt canonical spatial primitives that enforce geometric cycle consistency through rendered supervision from reconstructed assets and iterative texture repair mechanisms [148].

Overcoming the information loss inherent in projecting three-dimensional structures onto two-dimensional planes requires synthesizing diverse category anchors directly from textual descriptions. True architectural unification demands processing heterogeneous representations such as meshes and point clouds within a single latent space, thereby eliminating the need for domain-specific adapters. This approach enables zero-shot transfer across distinct 3D domains by aligning geometric distributions through shared semantic priors without explicit retraining, facilitating rotation-invariant feature extraction across unseen categories [173]. Parallel advancements in scalable foundation models point toward end-to-end autoregressive coordinate prediction that eliminates intermediate vector quantization steps. By framing explicit coordinate serialization as an autoregressive generation task within neural coordinate fields, these systems preserve high-frequency geometric fidelity through implicit neural embeddings rather than discretized representations [18].

4.2.2 Simulation-Ready and Embodied AI Integration

Moving beyond purely geometric synthesis, future generative architectures will likely employ unified representational frameworks that jointly model structural topology and intrinsic physical parameters [9, 42]. This shift enables the direct prediction of cross-domain correlations, such as disentangled albedo-metallic estimates and kinematic properties, through integrated latent modeling. Ensuring immediate simulation compatibility requires treating watertight manifold constraints and standardized export formats as explicit generation contracts rather than relying on post-hoc repairs [42, 164]. Techniques like geometry-aware texture synthesis and multi-stage validation pipelines will become essential to guarantee URDF or USD portability and valid collision geometry from the outset. Scalable asset creation for embodied agents will subsequently depend on agentic verification loops that autonomously recover missing physical attributes and validate interaction semantics [9, 164]. These systems will employ VLM-driven property recovery and affordance autolabeling

within generate-verify-retry cycles, effectively supervising fine-tuning with embedded physics priors to eliminate manual intervention.

Scaling the diversity of generated interactive environments directly correlates with improved zero-shot generalization in robotic policies, effectively overcoming the overfitting inherent in narrow manual scene design through language-driven scene graph generation and parallel reinforcement learning fine-tuning [27]. Separating object motion prediction from low-level robot control via synthesized 3D goal states mitigates simulation inaccuracies, enabling robust transfer without real-world trajectory data by enforcing geometric consistency across varying contact configurations [162]. Hybrid rendering pipelines that merge photorealistic backgrounds with precise mesh-based foregrounds optimize both visual fidelity and collision accuracy for effective behavior cloning, utilizing 3D Gaussian Splatting for background recovery alongside ray-traced foreground integration [49]. Adopting domain-agnostic 3D point cloud representations abstracts away visual appearance discrepancies, allowing policies trained on camera-aligned synthetic data to transfer directly to real-world settings through automated keypoint extraction and unified sim-to-real point encoding [48].

Automating the extraction of functional digital twins from single-view inputs will soon facilitate rapid sim-to-real adaptation without extensive manual scene engineering [120]. This shift relies on modular foundation model integration to lift RGB-D data while preserving object affordances and correlating them with sub-task evaluation metrics. To ensure dynamic stability in these one-shot generated worlds, standard pipelines must enforce iterative consistency checks between geometric representations and physical simulators [208]. Such refinement resolves interpenetration issues through layered scene decomposition and gravity-consistent frame alignment, effectively decoupling visual context from contact conflict resolution. Bridging the gap between photorealism and collision accuracy further requires hybrid representations that combine Gaussian splatting for static backgrounds with explicit meshes for interactive objects [210]. Driven by multimodal large language models, these architectures facilitate kinematic inference tailored for embodied agent control and articulated URDF synthesis, allowing holistic domain randomization to reduce visual discrepancies during training.

4.2.3 Dynamic and 4D Scene Evolution

Moving beyond static fusion, next-generation architectures will likely abandon global representations in favor of decoupled spatial memory systems that retrieve per-frame geometry to mitigate error accumulation during long-horizon synthesis [135]. Such frameworks operate through an autoregressive retrieve-generate-update loop, utilizing spatial alignment techniques and visibility scoring to maintain a persistent 3D geometry cache while actively correcting drift. Concurrently, generative pipelines are expected to transition from latent-compressed denoising toward direct volumetric flow matching across hierarchical distributions, enabling unbounded scene expansion without boundary artifacts [64]. This shift employs conditional flow matching on continuous volumetric representations with aggregated motion modeling to achieve direct generation without intermediate latent compression. To enforce strict temporal consistency, future models will inject dense geometric correspondences directly into attention mechanisms rather than relying on pixel-level history conditioning [64, 135]. By integrating aligned canonical maps and independent learned velocity fields into attention operations guided by geometric tracking signals, these systems decouple geometric tracking from appearance synthesis to optimize transport between distributions.

Replacing visual-only feedback loops with deterministic runtime state verification will become essential for maintaining physical mechanism correctness over extended temporal horizons. Future architectures must integrate layered scene protocols that perform runtime state verification and physical consistency checks to detect deviations before they propagate [86]. This approach shifts the generative objective from optimizing for visual similarity metrics to enforcing strict engine-state checking, enabling localized error repair without collapsing the entire sequence. Structural integrity in evolving four-dimensional worlds further demands staged construction pipelines that isolate assembly faults to specific build phases rather than relying on unstable global optimization strategies. By executing ordered construction stages through strategic sequential planning and script-based agent actions, systems can preserve editing success rates while providing clear quantitative success indicators for each developmental step [86].

Advancing beyond static assembly requires pipelines that enforce structural alignment prior to appearance synthesis through sequential geometry-to-appearance diffusion frameworks. By integrating three-dimensional semantic occupancy priors with geometry-guided image synthesis, these architectures

utilize advanced attention mechanisms for spatial coherence to preserve consistency across temporal steps [188]. Realizing multi-instance dependency modeling demands a transition from single-object reconstruction to architectures that capture cross-instance interactions via simultaneous refinement of multiple objects focused strictly on temporal alignment. Such systems employ unified generative passes and optimized flow-based generation to model inter-object dynamics within a single coordinated operation [58]. Achieving comprehensive video-level understanding further necessitates the integration of retrieval-augmented large language models to transform coarse temporal prompts into structured, evolving scene graphs. This approach enables high-level temporal control and consistency-aware outpainting by converting enriched prompts into layout constraints for controllable four-dimensional expansion [188].

4.2.4 Human-in-the-Loop and Collaborative Creation

Future data engines will transition from static dataset curation to iterative self-training loops where human ranking signals directly optimize generative policies via preference alignment [155]. This shift enables cascaded training recipes that bridge synthetic-to-real domains by incorporating scalable annotation pipelines for direct preference optimization. By relying on explicit feedback rather than automated attribution maps, these systems allow users to guide targeted geometric and material modifications to resolve specific model failure modes [105]. Such an approach facilitates precise manipulation of mesh properties while suppressing confounding artifacts through reinforcement mechanisms driven by human input to distinguish between valid and invalid outputs. Consequently, hybrid workflows will emerge where humans act as high-level verifiers for complex edge cases while models handle bulk generation and initial refinement steps [105, 155]. These collaborative iteration cycles address manual intervention bottlenecks by automating feature steering based on ranked preferences while retaining human oversight for semi-synthetic occlusion handling.

Bridging the gap between abstract user intent and executable scene logic will likely depend on retrieval-augmented generation frameworks that synthesize procedural code from natural language queries [115]. Such systems enable non-experts to manipulate complex scene graphs by translating textual metadata condensation into direct asset modifications without requiring manual scripting. To ensure spatial precision, future models will internalize explicit logi-

cal deductions through structured reasoning supervision derived from programmatic generation tasks [54]. This approach reduces dependence on massive annotated datasets while improving bounding box localization and spatial relationship adherence. A dual-stage pipeline combining large language model-driven code synthesis with human verification loops will become essential to correct hallucinations and enforce physical consistency [54, 115]. These iterative workflows allow users to intervene specifically where automated generation fails to satisfy geometric or semantic constraints.

Future annotation pipelines will evolve into collaborative verification loops where chained foundation models propose dense semantic ground truth for explicit human correction rather than operating without intervention. By employing chain-of-thought prompting alongside open-vocabulary query generation, these systems convert sparse point inputs into precise masks that users refine to resolve ambiguity before aggregating them via voting-based multi-view strategies [72]. Direct integration of such human-in-the-loop data engines into training cycles enables models to achieve superior cross-dataset generalization by learning from corrected errors instead of relying solely on automated outputs. This approach effectively bridges the synthetic-to-real domain gap while preserving frozen encoder capabilities for fine-grained part-level localization through iterative feedback [72]. Eliminating closed-vocabulary constraints through this interactive scale robustness allows architectures to adapt dynamically to novel object categories as users guide the system through complex scene compositions.

4.2.5 Ethical and Safety Considerations

To prevent malicious applications such as the creation of deceptive 3D deepfakes, next-generation pipelines will embed imperceptible neural watermarks directly into mesh topology and texture maps [217]. Such cryptographic signatures will enable forensic verification of asset origin, distinguishing between human-created content and algorithmically synthesized scenes. Policy enforcement mechanisms will subsequently block the rendering or export of unwatermarked models in sensitive simulation environments [8].

Deploying synthetic datasets with latent biases in safety-critical domains, such as autonomous driving simulation, creates direct physical hazards where skewed semantic distributions corrupt agent policy training. Future certification protocols must mandate adversarial stress testing that isolates visual fidelity

from ethical alignment, ensuring object libraries do not propagate stereotypes regardless of photorealistic quality [49]. Regulatory frameworks will require explicit decoupling metrics to verify that generated scenes maintain cultural diversity standards before entering embodied learning loops [66]. These safeguards prevent the injection of discriminatory behaviors into downstream systems by blocking assets that fail strict fairness thresholds during pre-deployment audits [165].

Responsible deployment policies require automated safety auditing loops that enforce ethical guidelines by rejecting harmful content before it enters production environments. These governance frameworks utilize critic-agent feedback not for optimization efficiency, but to verify compliance with predefined safety and utility standards during synthesis. By integrating constraint satisfaction checks directly into the generation pipeline, systems can automatically discard outputs that violate ergonomic or functional constraints without human intervention. This architectural approach shifts ethical oversight from post-hoc filtering to real-time preventative control, ensuring only compliant assets proceed to downstream applications [90].

4.2.6 Standardized Evaluation Protocols

Current heterogeneity in evaluation metrics fails to capture structural coherence across volumetric, mesh, and implicit representations, leading to slice-wise distribution mismatches and projection ambiguities [53, 60]. Existing benchmarks frequently rely on domain-specific proxies that lack transferability between medical imaging, avatar reconstruction, and procedural pre-training contexts, thereby exacerbating synthetic-to-real domain shifts [47, 53]. A viable pathway toward representation-agnostic quality assurance involves deploying automated filtering protocols and ensemble consistency checks that operate without manual ground truth [47, 60]. Such mechanisms integrate standard deviation filtering with systematic illumination variation and noise model integration to enforce double-check consistency across diverse output formats. Future standards must prioritize pose disentanglement and manifold violation detection to ensure metrics reflect geometric fidelity rather than rendering artifacts.

Organized challenge competitions offer a pragmatic pathway to unify evaluation standards by validating proxy optimization strategies that correlate directly with downstream task performance without requiring semantic labels. By focusing on layout parameter optimization and occlusion proxies, these events can establish scale distribution metrics and instance count

variables that drive community consensus [70]. Such frameworks encourage geometric invariance learning and contrastive pretraining alignment, ensuring that benchmarked improvements reflect genuine utility rather than overfitting to specific dataset artifacts. This competitive structure forces the field to prioritize measurable outcomes that translate effectively across diverse application domains.

Creating a viable ecosystem for industry standards demands a multi-stakeholder consortium to define unified benchmarking suites that reconcile the divergent evaluation needs of procedural pre-training and neural stylization pipelines [21, 47]. Such cross-paradigm metric alignment must enforce representation-agnostic quality assurance while quantifying synthetic-to-real domain shifts and style-content disentanglement scores. Current fragmentation in assessment protocols necessitates a shared governance framework where academic researchers and industrial practitioners jointly validate metrics against downstream utility rather than isolated visual fidelity [21, 47]. This collaborative validation should prioritize geometric consistency constraints, optimization convergence rates, and feed-forward inference latency benchmarks over subjective aesthetic judgments or biologically inspired predictivity metrics. Standardized reporting mandates must further evolve to include detailed provenance of generative priors and illumination conditions, ensuring reproducibility across diverse 3D asset creation workflows [21, 47]. Rigorous documentation requirements will specifically cover generative prior provenance, geometric consistency logs, deformation parameter records, and systematic pose variation reports to substantiate any claim of generative advancement.

5 Conclusion

Overcoming the chronic scarcity of annotated 3D data requires a synthesis where the structural exactness of procedural algorithms merges with the visual diversity of neural generative models. This hybrid approach yields simulation-ready assets possessing both geometric fidelity and semantic richness, effectively connecting static content creation to dynamic utility. Future progress depends on embedding physical priors directly into generative frameworks, ensuring synthesized worlds exhibit the dynamic plausibility necessary for training embodied agents rather than merely approximating visual appearance. To sustain this trajectory, the community must replace fragmented benchmarking practices with unified evaluation standards that simultaneously quantify geometric accuracy, semantic alignment, and physical

consistency. Rigorous standardization combined with physics-grounded generation will establish synthetic data as a foundational resource for machine learning.

References

- [1] Rameen Abdal, Wang Yifan, Zifan Shi, Yinghao Xu, Ryan Po, Zhengfei Kuang, Qifeng Chen, Dit-Yan Yeung, and Gordon Wetzstein. Gaussian Shell Maps for Efficient 3D Human Generation. *2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023. [12](#), [15](#)
- [2] V. Agarwal, J. Guo, Declan Hoban, Si Zhang, Nicholas Moran, Peter Cho, Srilakshmi Pattabiraman, and S. Joshi. Object-Centric Data Synthesis for Category-level Object Detection. *arXiv preprint arXiv:2511.23450*, 2025. <https://arxiv.org/abs/2511.23450>. [10](#)
- [3] Jose Moises Araya-Martinez, Thushar Tom, Adrián Sanchis Reig, Pablo Rey Valiente, Jens Lambrecht, and Jörg Krüger. SynthRender and IRIS: Open-Source Framework and Dataset for Bidirectional Sim-Real Transfer in Industrial Object Perception. *arXiv preprint arXiv:2602.21141*, 2026. <https://arxiv.org/abs/2602.21141>. [10](#)
- [4] Sherwin Bahmani, J. Park, Despoina Paschalidou, Xingguang Yan, Gordon Wetzstein, L. Guibas, and Andrea Tagliasacchi. CC3D: Layout-Conditioned Generation of Compositional 3D Scenes. *2023 IEEE/CVF International Conference on Computer Vision (ICCV)*, 2023. [11](#), [12](#), [18](#)
- [5] Song Bai and Jie Li. Progress and Prospects in 3D Generative AI: A Technical Overview including 3D human. *arXiv preprint arXiv:2401.02620*, 2024. <https://arxiv.org/abs/2401.02620>. [3](#)
- [6] Renato Baptista, Alexandre Saint, Kassem Al Ismaeil, and D. Aouada. Towards Generalization of 3D Human Pose Estimation In The Wild. *ICPR Workshops*, 2020. [9](#)
- [7] Majid Behravan. Generative AI Framework for 3D Object Generation in Augmented Reality. *arXiv preprint arXiv:2502.15869*, 2025. <https://arxiv.org/abs/2502.15869>. [24](#), [25](#)
- [8] Zixuan Bian, Ruohan Ren, Yue Yang, and Christopher Callison-Burch. HOLODECK 2.0: Vision-Language-Guided 3D World Generation with Editing. *arXiv preprint arXiv:2508.05899*, 2025. <https://arxiv.org/abs/2508.05899>. [21](#), [22](#), [25](#), [31](#)
- [9] Ziang Cao, Zhaoxi Chen, Liang Pan, and Ziwei Liu. PhysX-3D: Physical-Grounded 3D Asset Generation. *Neural Information Processing Systems*, 2025. [28](#), [29](#)
- [10] Nicolás Esleyder Cayturo-Silva and I. Sipiran. 3D Shape Generation: A Survey. *arXiv preprint arXiv:2506.22678*, 2025. <https://arxiv.org/abs/2506.22678>. [1](#)
- [11] P. Cañas, J. Ortega, M. Nieto, and O. Otaegui. Virtual passengers for real car solutions: synthetic datasets. *arXiv preprint arXiv:2205.06556*, 2022. <https://arxiv.org/abs/2205.06556>. [6](#)
- [12] Satrajit Chakrabarty, Sourya Sengupta, Gopal Avinash, and Ravi Soni. Synthetic Volumetric Data Generation Enables Zero-Shot Generalization of Foundation Models in 3D Medical Image Segmentation. *2026 IEEE 23rd International Symposium on Biomedical Imaging (ISBI)*, 2026. [29](#)
- [13] Chirui Chang, Zhengzhe Liu, Xiaoyang Lyu, and Xiaojuan Qi. How Far are AI-Generated Videos from Simulating the 3D Visual World: A Learned 3D Evaluation Approach. *2025 IEEE/CVF International Conference on Computer Vision (ICCV)*, 2024. [26](#), [28](#)
- [14] Hong-En Chen, Bin Wu, Sheng-Yu Huang, and Yu-Chiang Frank Wang. TPA3D: Triplane Attention for Fast Text-to-3D Generation. *European Conference on Computer Vision*, 2023. [12](#), [13](#)
- [15] Jinnan Chen, Lingting Zhu, Zeyu Hu, Shengju Qian, Yugang Chen, Xin Wang, and Gim Hee Lee. MAR-3D: Progressive Masked Auto-regressor for High-Resolution 3D Generation. *2025 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025. [12](#), [14](#), [27](#)
- [16] Meida Chen, Qingyong Hu, Hugues Thomas, Andrew Feng, Yu Hou, Kyle McCullough, and L. Soibelman. STPLS3D: A Large-Scale Synthetic and Real Aerial Photogrammetry 3D Point Cloud Dataset. *British Machine Vision Conference*, 2022. [4](#), [10](#)
- [17] Shen Chen, Jiale Zhou, Zhongyu Jiang, Tianfang Zhang, Zongkai Wu, Jenq-Neng Hwang, and Lei Li. ScalingGaussian: Enhancing 3D Content Creation with Generative Gaussian Splatting. *arXiv preprint arXiv:2407.19035*, 2024. <https://arxiv.org/abs/2407.19035>. [16](#), [17](#), [21](#), [25](#)
- [18] Sijin Chen, Xin Chen, Anqi Pang, Xianfang Zeng, Wei Cheng, Yijun Fu, Fukun Yin, Yanru Wang, Zhibin Wang, C. Zhang, Jingyi Yu, Gang Yu, Bin Fu, and Tao Chen. MeshXL: Neural Coordinate Field for Generative 3D Foundation Models. *Neural Information Processing Systems*, 2024. [29](#)
- [19] Xue Chen, Ying Chen, Jialin Li, Qiang Nie, Hanqiu Deng, Yong Liu, Qixing Huang, and Yang Li. P3P: Pseudo-3D Pre-training for Scaling 3D Voxel-based Masked Autoencoders. *arXiv preprint arXiv:2408.10007*, 2024. <https://arxiv.org/abs/2408.10007>. [19](#)

- [20] Xuweiyi Chen and Zezhou Cheng. Semantic-Free Procedural 3D Shapes are Surprisingly Good Teachers. *2026 International Conference on 3D Vision (3DV)*, 2024. 3, 5, 11
- [21] Yingshu Chen, Guocheng Shao, Ka Chun Shum, Binh-Son Hua, and Sai-Kit Yeung. Advances in 3D Neural Stylization: A Survey. *International Journal of Computer Vision*, 2023. 5, 28, 32
- [22] Yongwei Chen, Yushi Lan, Shangchen Zhou, Tengfei Wang, and Xingang Pan. SAR3D: Autoregressive 3D Object Generation and Understanding via Multi-scale 3D VQVAE. *2025 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. 12, 14
- [23] Yongwei Chen, Tianyi Wei, Yushi Lan, Zhaoyang Lyu, Shangchen Zhou, Xudong Xu, and Xingang Pan. PnP-U3D: Plug-and-Play 3D Framework Bridging Autoregression and Diffusion for Unified Understanding and Generation. *arXiv preprint arXiv:2602.03533*, 2026. <https://arxiv.org/abs/2602.03533>. 13, 29
- [24] Zhaoxi Chen, Guangcong Wang, and Ziwei Liu. SceneDreamer: Unbounded 3D Scene Generation From 2D Image Collections. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2023. 12, 14, 15, 27
- [25] Yu Chi, Fangneng Zhan, Sibow Wu, C. Theobalt, and A. Kortylewski. DatasetNeRF: Efficient 3D-aware Data Factory with Generative Radiance Fields. *European Conference on Computer Vision*, 2023. 16
- [26] Wonwoong Cho, Inyeop Lee, and David I. Inouye. Enhanced 3DMM Attribute Control via Synthetic Dataset Creation Pipeline. *arXiv preprint arXiv:2011.12833*, 2020. <https://arxiv.org/abs/2011.12833>. 12
- [27] Andrew Choi, Xinjie Wang, Zhizhong Su, and Wei Xu. Scaling Sim-to-Real Reinforcement Learning for Robot VLAs with Generative 3D Worlds. *arXiv preprint arXiv:2603.18532*, 2026. <https://arxiv.org/abs/2603.18532>. 23, 30
- [28] Nathan L. Clement, Alan Schoen, Arnold P. Boedhardjo, and A. Jenkins. Synthetic Data and Hierarchical Object Detection in Overhead Imagery. *ACM Transactions on Multimedia Computing, Communications and Applications*, 2021. 10
- [29] Jaime Corsetti, Francesco Giuliani, Davide Boscaini, P. Hermosilla, Andrea Pilzer, Guofeng Mei, Alexandros Delitzas, Francis Engelmann, and Fabio Poiesi. Action-guided generation of 3D functionality segmentation data. *arXiv preprint arXiv:2511.23230*, 2025. <https://arxiv.org/abs/2511.23230>. 21, 24, 25
- [30] Debayan Deb, Suvidha Tripathi, and Pranit Puri. MUNCH: Modelling Unique 'N' Controllable Heads. *Proceedings of the 16th ACM SIGGRAPH Conference on Motion, Interaction and Games*, 2023. 12, 27
- [31] Jie Deng, Wenhao Chai, Junsheng Huang, Zhonghan Zhao, Qixuan Huang, Mingyan Gao, Jianshu Guo, Shengyu Hao, Wenhao Hu, Jenq-Neng Hwang, Xi Li, and Gaoang Wang. CityCraft: A Real Crafter for 3D City Generation. *arXiv preprint arXiv:2406.04983*, 2024. <https://arxiv.org/abs/2406.04983>. 22, 23, 25
- [32] Helisa Dharmo, Fabian Manhardt, Nassir Navab, and Federico Tombari. Graph-to-3D: End-to-End Generation and Manipulation of 3D Scenes Using Scene Graphs. *2021 IEEE/CVF International Conference on Computer Vision (ICCV)*, 2021. 8, 13, 25
- [33] Huanning Dong, Yinyu Huang, Fan Li, and Ping Kuang. Single Image to High-Quality 3D Object via Latent Features. *arXiv preprint arXiv:2511.19512*, 2025. <https://arxiv.org/abs/2511.19512>. 15
- [34] Junyi Dong, Haotian Luo, Ziwei Xu, S. Bian, Heng Zhang, Sitong Mao, Jing Guo, Yang Xu, Wenhao Chen, Qiuyu Feng, Yao Mu, Ping Luo, Shunbo Zhou, and Xiaodong Wu. HyperSim: A Holistic Sim-To-Real Framework For Robust Robotic Manipulation. *arXiv preprint arXiv:2605.26638*, 2026. <https://arxiv.org/abs/2605.26638>. 10
- [35] Rafael C. e Silva, Nuno Fachada, Diogo de Andrade, and Nélío Códices. Procedural Generation of 3D Maps With Snappable Meshes. *IEEE Access*, 2021. 6
- [36] Ainaz Eftekhari, Alexander Sax, Roman Bachmann, J. Malik, and Amir Zamir. Omnidata: A Scalable Pipeline for Making Multi-Task Mid-Level Vision Datasets from 3D Scans. *2021 IEEE/CVF International Conference on Computer Vision (ICCV)*, 2021. 4, 7, 9
- [37] Paul Engstler, Iro Laina, Christian Rupprecht, and Andrea Vedaldi. SynCity 3000: Bootstrapping Scene-Scale 3D Diffusion. *arXiv preprint arXiv:2607.05392*, 2026. <https://arxiv.org/abs/2607.05392>. 16
- [38] E. Ershov, A. Belokopytov, and A. Savchik. Problems of dataset creation for light source estimation. *arXiv preprint arXiv:2006.02692*, 2020. <https://arxiv.org/abs/2006.02692>. 4
- [39] Shaoheng Fang, Chaohui Yu, Fan Wang, and Qi-Xing Huang. MVRoom: Controllable 3D Indoor Scene Generation with Multi-View Diffusion Models. *arXiv preprint arXiv:2512.04248*, 2025. <https://arxiv.org/abs/2512.04248>. 18, 19, 21
- [40] Faraz Faruqi, Yingtao Tian, Vrushank Phadnis, Varun Jampani, and Stefan Müller. Shaping Realities: Enhancing 3D Generative AI with Fabrication

- Constraints. *arXiv preprint arXiv:2404.10142*, 2024. <https://arxiv.org/abs/2404.10142>. 28
- [41] Hao Feng, Zhi Zuo, Jiaji Pan, Ka-Hei Hui, Qi Dou, Jingyu Hu, and Zhengzhe Liu. WonderVerse: Extendable 3D Scene Generation with Video Generative Models. *arXiv preprint arXiv:2503.09160*, 2025. <https://arxiv.org/abs/2503.09160>. 18, 19, 20, 21
- [42] Jiashi Feng, Xiu Li, Jingyu Lin, Jiahang Liu, Gao-hong Liu, Weiqiang Lou, Suizhi Ma, Guang Shi, Qinlong Wang, Jun Wang, Zhongcong Xu, Xuanyu Yi, Zihao Yu, Jianfeng Zhang, Yifan Zhu, Rui Chen, Jinxin Chi, Zixian Du, Li Han, Lixin Huang, Kaihua Jiang, Yuhan Li, Guan Luo, Shuguang Wang, Qianyi Wu, Fan Yang, Junyang Zhang, and Xuanmeng Zhang. Seed3D 1.0: From Images to High-Fidelity Simulation-Ready 3D Assets. *arXiv preprint arXiv:2510.19944*, 2025. <https://arxiv.org/abs/2510.19944>. 15, 16, 29
- [43] Quankai Gao, Iliyan Georgiev, Tuanfeng Y. Wang, Krishna kumar Singh, Ulrich Neumann, and Jae Shin Yoon. Can3Tok: Canonical 3D Tokenization and Latent Modeling of Scene-Level 3D Gaussians. *2025 IEEE/CVF International Conference on Computer Vision (ICCV)*, 2025. 12, 16
- [44] Zhengqing Gao, Ziwen Li, Xin Wang, Jiaxin Huang, Zhenyang Ren, Ming Shao, Hanlue Zhang, Tianyu Huang, Yongkang Cheng, Yandong Guo, Runqi Lin, Yuanyuan Wang, Tongliang Liu, Kun Zhang, and Mingming Gong. Mirage2Matter: A Physically Grounded Gaussian World Model from Video. *arXiv preprint arXiv:2602.00096*, 2026. <https://arxiv.org/abs/2602.00096>. 27, 28
- [45] Téo Guichoux, Laure Soulier, Nicolas Obin, and Catherine Pelachaud. Investigating the impact of 2D gesture representation on co-speech gesture generation. *arXiv preprint arXiv:2406.15111*, 2024. <https://arxiv.org/abs/2406.15111>. 27
- [46] Jiarong Guo, Xiaogang Xu, and Hengshuang Zhao. Self-supervised Learning for Enhancing Geometrical Modeling in 3D-Aware Generative Adversarial Network. *arXiv preprint arXiv:2312.11856*, 2023. <https://arxiv.org/abs/2312.11856>. 12
- [47] Shubh Gupta, Thomas P. O’Connell, and Bernhard Egger. Beyond Flatland: Pre-training with a Strong 3D Inductive Bias. *arXiv preprint arXiv:2112.00113*, 2021. <https://arxiv.org/abs/2112.00113>. 4, 5, 32
- [48] Siddhant Haldar, Lars Johannsmeier, Lerrel Pinto, Abhishek Gupta, Dieter Fox, Yashraj S. Narang, and A. Mandlekar. Point Bridge: 3D Representations for Cross Domain Policy Learning. *arXiv preprint arXiv:2601.16212*, 2026. <https://arxiv.org/abs/2601.16212>. 27, 30
- [49] Xiaoshen Han, Minghuan Liu, Yilun Chen, Junqiu Yu, Xiaoyang Lyu, Yang Tian, Bolun Wang, Weinan Zhang, and Jiangmiao Pang. Re3Sim: Generating High-Fidelity Simulation Data via 3D-Photorealistic Real-to-Sim for Robotic Manipulation. *arXiv preprint arXiv:2502.08645*, 2025. <https://arxiv.org/abs/2502.08645>. 7, 9, 15, 30, 32
- [50] Liu He, Xiao Zeng, Yizhi Song, Albert Y. C. Chen, Lu Xia, Shashwat Verma, Sankalp Dayal, Min Sun, Cheng-Hao Kuo, and Daniel G. Aliaga. Advancing Multimodal LLMs by Large-Scale 3D Visual Instruction Dataset Generation. *2026 IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, 2025. 28
- [51] Paul Henderson and Christoph H. Lampert. Unsupervised object-centric video generation and decomposition in 3D. *Neural Information Processing Systems*, 2020. 20, 26
- [52] Nick Heppert, Minh Nguyen, and A. Valada. Scaling Single Human Demonstrations for Imitation Learning using Generative Foundational Models. *arXiv preprint arXiv:2602.12734*, 2026. <https://arxiv.org/abs/2602.12734>. 21
- [53] Sun-Beom Hong, Razvan V. Marinescu, Adrian V. Dalca, A. Bonkhoff, Martin Bretzner, N. Rost, and P. Golland. 3D-StyleGAN: A Style-Based Generative Adversarial Network for Generative Modeling of Three-Dimensional Medical Images. *DGM4MICCAI/DALI@MICCAI*, 2021. 11, 28, 32
- [54] Hsiang-Wei Huang, Kuang-Ming Chen, Wenhao Chai, Cheng-Yen Yang, Jen-Hao Cheng, and Jenq-Neng Hwang. Reasoning Matters for 3D Visual Grounding. *arXiv preprint arXiv:2601.08811*, 2026. <https://arxiv.org/abs/2601.08811>. 22, 31
- [55] Julia Huang, T. Smith, Aloukika Patro, and V. Chhabra. A Novel Approach For Generating Customizable Light Field Datasets for Machine Learning. *2022 IEEE MIT Undergraduate Research Technology Conference (URTC)*, 2022. 4, 8
- [56] Tianyu Huang, Yihan Zeng, Bowen Dong, Hang Xu, Songcen Xu, Rynson W. H. Lau, and Wangmeng Zuo. TextField3D: Towards Enhancing Open-Vocabulary 3D Generation with Noisy Text Fields. *International Conference on Learning Representations*, 2023. 13
- [57] Weikai Huang, Jieyu Zhang, Taoyang Jia, Chenhao Zheng, Ziqi Gao, J. Park, Winson Han, and Ranjay Krishna. Synthetic Object Compositions for Scalable and Accurate Learning in Detection, Segmentation, and Grounding. *arXiv preprint arXiv:2510.09110*, 2025. <https://arxiv.org/abs/2510.09110>. 4
- [58] Zehuan Huang, Yuan-Chen Guo, Xingqiao An, Yu nuo Yang, Yangguang Li, Zi-Xin Zou, Ding Liang, Xihui Liu, Yan-Pei Cao, and Lu Sheng. MIDI: Multi-Instance Diffusion for Single Image

- to 3D Scene Generation. *2025 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. 7, 8, 12, 13, 31
- [59] Zhening Huang, Xiaoyang Wu, Xi Chen, Hengshuang Zhao, Lei Zhu, and Joan Lasenby. OpenIns3D: Snap and Lookup for 3D Open-vocabulary Instance Segmentation. *European Conference on Computer Vision*, 2023. 28
- [60] Zihao Huang, Shou-Yong Hu, Guangcong Wang, Tianqi Liu, Yuhang Zang, Zhiguo Cao, Wei Li, and Ziwei Liu. WildAvatar: Learning In-the-wild 3D Avatars from the Web. *2025 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. 5, 32
- [61] Mona Jalal, J. Spjut, Ben Boudaoud, and Margrit Betke. SIDOD: A Synthetic Image Dataset for 3D Object Pose Recognition With Distractors. *2019 IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, 2019. 4, 11
- [62] Eric Ji, Qiran Hu, Wufei Ma, Sarthak Jain, Yingying Li, Minh N. Do, and Yaoyao Liu. AC3S: Adaptive Conditioning for 3D-Aware Synthetic Data Generation. *arXiv preprint arXiv:2606.31204*, 2026. <https://arxiv.org/abs/2606.31204>. 17, 23
- [63] Hanwen Jiang, Zexiang Xu, Desai Xie, Ziwen Chen, Haian Jin, Fujun Luan, Zhixin Shu, Kai Zhang, Sai Bi, Xin Sun, Jiuxiang Gu, Qixing Huang, Georgios Pavlakos, and Hao-Tian Tan. MegaSynth: Scaling Up 3D Scene Reconstruction with Synthesized Data. *2025 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. 7, 10
- [64] Amogh Joshi, Julian Ost, and Felix Heide. WorldFlow3D: Flowing Through 3D Distributions for Unbounded World Generation. *arXiv preprint arXiv:2603.29089*, 2026. <https://arxiv.org/abs/2603.29089>. 12, 13, 27, 30
- [65] Animesh Karnewar, Oliver Wang, Tobias Ritschel, and N. Mitra. 3inGAN: Learning a 3D Generative Model from Images of a Self-similar Scene. *2022 International Conference on 3D Vision (3DV)*, 2022. 11, 15, 26
- [66] Jungtaek Kim, H. Chung, Minsu Cho, and Jaesik Park. Combinatorial 3D Shape Generation via Sequential Assembly. *arXiv preprint arXiv:2004.07414*, 2020. <https://arxiv.org/abs/2004.07414>. 6, 28, 32
- [67] A. Kushwaha, Sarthak Gupta, Anish Bhanushali, and T. R. Dastidar. Rapid Training Data Creation by Synthesizing Medical Images for Classification and Localization. *2020 IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, 2020. 11
- [68] A. I. Károly and P. Galambos. A Scalable Pipeline Combining Procedural 3D Graphics and Guided Diffusion for Photorealistic Synthetic Training Data Generation in White Button Mushroom Segmentation. *arXiv preprint arXiv:2512.08747*, 2025. <https://arxiv.org/abs/2512.08747>. 8, 10
- [69] Mithun Lal, Anthony Paprocki, N. Habili, L. Petersson, Olivier Salvado, and C. Fookes. Learning Dense Correspondence from Synthetic Environments. *2022 International Conference on Digital Image Computing: Techniques and Applications (DICTA)*, 2022. 9
- [70] Hei Law and Jia Deng. Label-Free Synthetic Pre-training of Object Detectors. *2024 IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, 2022. 29, 32
- [71] Han-Hung Lee, M. Savva, and Angel X. Chang. Text-to-3D Shape Generation. *Computer Graphics Forum*, 2024. 1
- [72] Junha Lee, Eunha Park, Chunghyun Park, Dahyun Kang, and Minsu Cho. Affogato: Open-Vocabulary Affordance Grounding with Automated Data Generation at Scale. *arXiv preprint arXiv:2506.12009*, 2025. <https://arxiv.org/abs/2506.12009>. 24, 31
- [73] Chenghao Li and Chaoning Zhang. When ChatGPT for Computer Vision Will Come? From 2D to 3D. *arXiv preprint arXiv:2305.06133*, 2023. <https://arxiv.org/abs/2305.06133>. 3, 29
- [74] Chenghao Li, Chaoning Zhang, Atish Waghvase, Lik-Hang Lee, François Rameau, Yang Yang, S. Bae, and Choong-Seon Hong. Generative AI meets 3D: A Survey on Text-to-3D in AIGC Era. *arXiv preprint arXiv:2305.06131*, 2023. <https://arxiv.org/abs/2305.06131>. 1
- [75] Guohao Li, Hongyu Yang, Yifang Men, Di Huang, Weixin Li, Ruijie Yang, and Yunhong Wang. Generating Editable Head Avatars with 3D Gaussian GANs. *ICASSP 2025 - 2025 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2024. 16, 25
- [76] Haoran Li, Haolin Shi, Wenli Zhang, Wenjun Wu, Yong Liao, Lin Wang, Lik-Hang Lee, and Pengyuan Zhou. DreamScene: 3D Gaussian-based Text-to-3D Scene Generation via Formation Pattern Sampling. *European Conference on Computer Vision*, 2024. 12, 16, 17, 18
- [77] Lanxiao Li and M. Heizmann. Randomized 3D Scene Generation for Generalizable Self-supervised Pre-training. *arXiv preprint arXiv:2306.04237*, 2023. <https://arxiv.org/abs/2306.04237>. 6
- [78] Sikuang Li, Chen Yang, Jiemin Fang, Taoran Yi, Jia Lu, Jiazhong Cen, Lingxi Xie, Wei Shen, and Qi Tian. WorldGrow: Generating Infinite 3D World. *AAAI Conference on Artificial Intelligence*, 2025. 12, 13, 27

- [79] Xiaoyu Li, Qi Zhang, Di Kang, Weihao Cheng, Yiming Gao, Jingbo Zhang, Zhihao Liang, Jing Liao, Yanpei Cao, and Ying Shan. Advances in 3D Generation: A Survey. *arXiv preprint arXiv:2401.17807*, 2024. <https://arxiv.org/abs/2401.17807>. 1
- [80] Xinke Li, Henghui Ding, Zekun Tong, Yuwei Wu, and Y. Chee. Primitive3D: 3D Object Dataset Synthesis from Randomly Assembled Primitives. *2022 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022. 6
- [81] Zuoyue Li, Zhenqiang Li, Zhaopeng Cui, M. Pollefeys, and Martin R. Oswald. Sat2Scene: 3D Urban Scene Generation from Satellite Images with Diffusion. *2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. 13
- [82] Hao li Feng, Yifan Niu, Zehuan Huang, Yang-Tian Sun, Chunchao Guo, Yuxin Peng, and Lu Sheng. Repurposing 3D Generative Model for Autoregressive Layout Generation. *arXiv preprint arXiv:2604.16299*, 2026. <https://arxiv.org/abs/2604.16299>. 16
- [83] Chendi Lin, Heshan Liu, Qunshu Lin, Zachary Bright, Shitao Tang, Yihui He, Minghao Liu, Ling Zhu, and Cindy X. Le. Objaverse++: Curated 3D Object Dataset with Quality Annotations. *2025 IEEE/CVF International Conference on Computer Vision Workshops (ICCVW)*, 2025. 4
- [84] Jiantao Lin, Xin Yang, Meixi Chen, Yingjie Xu, Dongyu Yan, Leyi Wu, Xinli Xu, Lie Xu, Shunsi Zhang, and Ying-Cong Chen. Kiss3DGen: Repurposing Image Diffusion Models for 3D Asset Generation. *2025 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025. 21
- [85] Yiqi Lin, Hao Wu, Ruichen Wang, H. Lu, Xiaodong Lin, Hui Xiong, and Lin Wang. Towards Language-guided Interactive 3D Generation: LLMs as Layout Interpreter with Generative Feedback. *arXiv preprint arXiv:2305.15808*, 2023. <https://arxiv.org/abs/2305.15808>. 21, 22, 25, 27
- [86] Chunjiang Liu, Xiaoyuan Wang, Haoyu Chen, Yizhou Zhao, Ming-Hsuan Yang, and László A. Jeni. SimWorlds: A Multi-Agent System for Dynamic 3D Scene Creation. *arXiv preprint arXiv:2607.01766*, 2026. <https://arxiv.org/abs/2607.01766>. 22, 23, 24, 30
- [87] Jian Liu, Xiaoshui Huang, Tianyu Huang, Lu Chen, Yuenan Hou, Shixiang Tang, Ziwei Liu, Wanli Ouyang, Wangmeng Zuo, Junjun Jiang, and Xianming Liu. A Comprehensive Survey on 3D Content Generation. *arXiv preprint arXiv:2402.01166*, 2024. <https://arxiv.org/abs/2402.01166>. 1
- [88] Minghua Liu, Ruoxi Shi, Linghao Chen, Zhuoyang Zhang, Chao Xu, Xinyue Wei, Hansheng Chen, Chong Zeng, Jiayuan Gu, and Hao Su. One-2-3-45++: Fast Single Image to 3D Objects with Consistent Multi-View Generation and 3D Diffusion. *2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023. 18
- [89] Xinhang Liu, Yu-Wing Tai, and Chi-Keung Tang. Multimodal Generation of Animatable 3D Human Models with AvatarForge. *arXiv preprint arXiv:2503.08165*, 2025. <https://arxiv.org/abs/2503.08165>. 23, 24
- [90] Xinhang Liu, Chi-Keung Tang, and Yu-Wing Tai. WorldCraft: Photo-Realistic 3D World Creation and Customization via LLM Agents. *arXiv preprint arXiv:2502.15601*, 2025. <https://arxiv.org/abs/2502.15601>. 8, 23, 24, 25, 32
- [91] Yibo Liu, Ziwei Zhang, Haozhou Pang, Menghao Li, Lanshan He, and Gan Qi. Native3D: End-to-End 3D Scene Generation via Unified Mesh-Texture Modeling and Semantic Alignment. *arXiv preprint arXiv:2606.07117*, 2026. <https://arxiv.org/abs/2606.07117>. 15
- [92] Yuan Liu, Cong Zhang, Xifeng Gao, and Xiaohu Guo. SuperVoxelGPT: Adaptive and Ordered 3D Tokenization for Autoregressive Shape Generation. *arXiv preprint arXiv:2605.29655*, 2026. <https://arxiv.org/abs/2605.29655>. 14
- [93] Yuheng Liu, Xinke Li, Xueting Li, Lu Qi, Chongshou Li, and Ming-Hsuan Yang. Pyramid Diffusion for Fine 3D Large Scene Generation. *European Conference on Computer Vision*, 2023. 27
- [94] Yuheng Liu, Xinke Li, Yuning Zhang, Lu Qi, Xin Li, Wenping Wang, Chongshou Li, Xueting Li, and Ming-Hsuan Yang. Controllable 3D Outdoor Scene Generation via Scene Graphs. *2025 IEEE/CVF International Conference on Computer Vision (ICCV)*, 2025. 7, 13
- [95] Hexiao Lu, Xiaokun Sun, Zeyu Cai, Hao-Xiang Guo, Ying Tai, Jian Yang, and Zhenyu Zhang. Muses: Designing, Composing, Generating Nonexistent Fantasy 3D Creatures without Training. *arXiv preprint arXiv:2601.03256*, 2026. <https://arxiv.org/abs/2601.03256>. 16
- [96] Zhixiang Lu, Mengqi Han, Peixin Guo, T. Bai, Jionglong Su, Fei Fang, and Sifan Song. Attention-Guided Flow-Matching for Sparse 3D Geological Generation. *arXiv preprint arXiv:2604.09700*, 2026. <https://arxiv.org/abs/2604.09700>. 13, 26
- [97] Baorui Ma, Huachen Gao, Haoge Deng, Zhengxiong Luo, Tiejun Huang, Lulu Tang, and Xinlong Wang. You See it, You Got it: Learning 3D Creation on Pose-Free Videos at Scale. *2025 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. 18, 19, 20, 21
- [98] Sihan Ma, Siyuan Liang, and Dacheng Tao. Who Generated This 3D Asset? Learning Source Attribution for Generative 3D Models. *arXiv preprint*

- arXiv:2605.18132*, 2026. <https://arxiv.org/abs/2605.18132>. 4, 29
- [99] Zeyu Ma, Alexander R. E. Raistrick, and Jia Deng. SimpleProc: Fully Procedural Synthetic Data from Simple Rules for Multi-View Stereo. *arXiv preprint arXiv:2604.04925*, 2026. <https://arxiv.org/abs/2604.04925>. 6, 7, 10, 11
- [100] L'eopold Maillard, Nicolas Sereyjol-Garros, Tom Durand, and M. Ovsjanikov. DeBaRA: Denoising-Based 3D Room Arrangement Generation. *Neural Information Processing Systems*, 2024. 7, 8
- [101] Tomohide Masuda, Matthew Ragoza, and D. Koes. Generating 3D Molecular Structures Conditional on a Receptor Binding Site with Deep Generative Models. *arXiv preprint arXiv:2010.14442*, 2020. <https://arxiv.org/abs/2010.14442>. 13
- [102] Tejaswini Medi, Arianna Rampini, P. Reddy, P. Jayaraman, and Margret Keuper. 3D-WAG: Hierarchical Wavelet-Guided Autoregressive Generation for High-Fidelity 3D Shapes. *arXiv preprint arXiv:2411.19037*, 2024. <https://arxiv.org/abs/2411.19037>. 14, 16, 27
- [103] Yifang Men, Biwen Lei, Yuan Yao, Miaomiao Cui, Zhouhui Lian, and Xuansong Xie. En3D: An Enhanced Generative Model for Sculpting 3D Humans from 2D Synthetic Data. *2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. 16
- [104] Sandeep Mishra, Oindrila Saha, and A. Bovik. C3DAG: Controlled 3D Animal Generation using 3D pose guidance. *arXiv preprint arXiv:2406.07742*, 2024. <https://arxiv.org/abs/2406.07742>. 17
- [105] N. Mital, Simon Malzard, Richard Walters, C. M. D. Melo, Raghuveer Rao, and Victoria Nockles. Improving Object Detection by Modifying Synthetic Data with Explainable AI. *Artificial Intelligence for Security and Defence Applications III*, 2024. 11, 31
- [106] Sicheng Mo, Ziyang Leng, Leo Liu, Weizhen Wang, Honglin He, and Bolei Zhou. Dreamland: Controllable World Creation with Simulator and Generative Models. *arXiv preprint arXiv:2506.08006*, 2025. <https://arxiv.org/abs/2506.08006>. 18, 20, 21, 25
- [107] Steven Moonen, Rob Salaets, Kenneth Batstone, A. Bey-Temsamani, and Nick Michiels. A Comparative Study of 3D Model Acquisition Methods for Synthetic Data Generation of Agricultural Products. *arXiv preprint arXiv:2601.03784*, 2026. <https://arxiv.org/abs/2601.03784>. 3, 5
- [108] Enric Moreu, Kevin McGuinness, Diego Ortego, and N. O'Connor. Domain Randomization for Object Counting. *Irish Conference on Artificial Intelligence and Cognitive Science*, 2022. 11
- [109] N. Morrical, Jonathan Tremblay, Yunzhi Lin, Stephen Tyree, Stan Birchfield, Valerio Pascucci, and I. Wald. NViSII: A Scriptable Tool for Photorealistic Image Generation. *arXiv preprint arXiv:2105.13962*, 2021. <https://arxiv.org/abs/2105.13962>. 9, 10
- [110] Julian Ost, Andrea Ramazzina, Amogh Joshi, Maximilian Bömer, Mario Bijelic, and Felix Heide. LSD-3D: Large-Scale 3D Driving Scene Generation with Geometry Grounding. *AAAI Conference on Artificial Intelligence*, 2025. 17, 18
- [111] Katar'ina Osvaldov'a, Luk'avs Gajdovsech, Viktor Kocur, and Martin Madaras. Enhancement of 3D Camera Synthetic Training Data with Noise Models. *arXiv preprint arXiv:2402.16514*, 2024. <https://arxiv.org/abs/2402.16514>. 9
- [112] Alvaro Patricio, Atabak Dehban, and Rodrigo Ventura. FLORA: Efficient Synthetic Data Generation for Object Detection in Low-Data Regimes via finetuning Flux LoRA. *arXiv preprint arXiv:2508.21712*, 2025. <https://arxiv.org/abs/2508.21712>. 11, 25
- [113] Ming Qian, Tianjian Ouyang, Mingchao Sun, Zijian Wang, Jincheng Xiong, Jiarong Han, Yongchang Zhang, Jiawei Zhang, Xu Wang, Yu Liu, Luyang Tang, Fei Yu, Zengye Ge, Mengmeng Du, Yuan Liu, Nianfei Fan, Song jian Wang, Yingliang Peng, Chun Jia, Yang Liu, S. Zeng, Haozhe Shi, Junnan Lai, Hongyu Pan, Zhenglian Wu, Ning Guo, Mu Xu, and Han Zhang. ABot-Earth 0.5: Generative 3D Earth Model. *arXiv preprint arXiv:2606.09967*, 2026. <https://arxiv.org/abs/2606.09967>. 16
- [114] Xue-Qing Qian, Yu Wang, Simian Luo, Yinda Zhang, Ying Tai, Zhenyu Zhang, Chengjie Wang, Xiangyang Xue, Bo Zhao, Tiejun Huang, Yunsheng Wu, and Yanwei Fu. Pushing Auto-regressive Models for 3D Shape Generation at Capacity and Scalability. *arXiv preprint arXiv:2402.12225*, 2024. <https://arxiv.org/abs/2402.12225>. 14, 29
- [115] Zhaofang Qian, Abolfazl Sharifi, T. Carroll, and Ser-Nam Lim. Scene Co-pilot: Procedural Text to Video Generation with Human in the Loop. *arXiv preprint arXiv:2411.18644*, 2024. <https://arxiv.org/abs/2411.18644>. 22, 24, 25, 31
- [116] Qiang Qu, Hanxue Liang, Xiaoming Chen, Yuk Ying Chung, and Yiran Shen. NeRF-NQA: No-Reference Quality Assessment for Scenes Generated by NeRF and Neural View Synthesis Methods. *IEEE Transactions on Visualization and Computer Graphics*, 2024. 3
- [117] Alexander R. E. Raistrick, Lahav Lipson, Zeyu Ma, Lingjie Mei, Mingzhe Wang, Yiming Zuo, Karhan Kayan, Hongyu Wen, Beining Han, Yihan Wang, Alejandro Newell, Hei Law, Ankit Goyal, Kaiyu Yang, and Jia Deng. Infinite Photorealistic Worlds

- Using Procedural Generation. *2023 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023. 7, 8, 11, 25
- [118] Alexander R. E. Raistrick, Lingjie Mei, Karhan Kayan, David Yan, Yiming Zuo, Beining Han, Hongyu Wen, Meenal Parakh, Stamatis Alexandropoulos, Lahav Lipson, Zeyu Ma, and Jia Deng. Infinigen Indoors: Photorealistic Indoor Scenes using Procedural Generation. *2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. 6, 7
- [119] Xingjian Ran, Shujie Zhang, Weipeng Zhong, L. Luo, and Bo Dai. Pair2Scene: Learning Local Object Relations for Procedural Scene Generation. *arXiv preprint arXiv:2604.11808*, 2026. <https://arxiv.org/abs/2604.11808>. 6, 8, 28
- [120] Nadun Ranawaka, Josiah Wong, Wei-Lin Pai, Wei-Teng Chu, Tianyuan Dai, Masoud Moghani, Hang Yin, Yunfan Jiang, Wesley Durbano, Brandon Huynh, Linxi Fan, Danfei Xu, Ruohan Zhang, Fei-Fei Li, Bowen Wen, A. Mandlekar, and Yuke Zhu. SimFoundry: Modular and Automated Scene Generation for Policy Learning and Evaluation. *arXiv preprint arXiv:2606.28276*, 2026. <https://arxiv.org/abs/2606.28276>. 30
- [121] Jianqiang Ren, Chao He, Lin Liu, Jiahao Chen, Yutong Wang, Yafei Song, Jianfang Li, Tangli Xue, Siqi Hu, Tao Chen, Kunkun Zheng, Jianjing Xiang, and Liefeng Bo. Make-A-Character: High Quality Text-to-3D Character Generation within Minutes. *arXiv preprint arXiv:2312.15430*, 2023. <https://arxiv.org/abs/2312.15430>. 16, 27
- [122] Xuanchi Ren, Jiahui Huang, Xiaohui Zeng, Ken Museth, Sanja Fidler, and Francis Williams. XCube: Large-Scale 3D Generative Modeling using Sparse Voxel Hierarchies. *2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023. 13
- [123] Mike Roberts and Nathan Paczan. Hypersim: A Photorealistic Synthetic Dataset for Holistic Indoor Scene Understanding. *2021 IEEE/CVF International Conference on Computer Vision (ICCV)*, 2020. 7, 9, 11
- [124] Roey Ron and G. Elbaz. Detection and Segmentation of Custom Objects using High Distraction Photorealistic Synthetic Data. *arXiv: Computer Vision and Pattern Recognition*, 2020. 9, 10
- [125] William Rowan, P. Huber, Nick E. Pears, and Andrew Keeling. Fake It Without Making It: Conditioned Face Generation for Accurate 3D Face Shape Estimation. *arXiv preprint arXiv:2307.13639*, 2023. <https://arxiv.org/abs/2307.13639>. 17
- [126] Mehrnaz Sabet, Praveen Palanisamy, and Sakshi Mishra. Scalable Modular Synthetic Data Generation for Advancing Aerial Autonomy. *Robotics Auton. Syst.*, 2022. 8, 11
- [127] Katharina Schmid, Nicolas von Lützwow, Jozef Hladký, Angela Dai, and Matthias Nießner. GenRecon: Bridging Generative Priors for Multi-View 3D Scene Reconstruction. *arXiv preprint arXiv:2605.23888*, 2026. <https://arxiv.org/abs/2605.23888>. 26
- [128] P. Schulz, T. Hempel, and A. Al-Hamadi. Automating 3D Dataset Generation with Neural Radiance Fields. *International Conference on Robotics, Computer Vision and Intelligent Systems*, 2025. 15
- [129] Katja Schwarz, Denys Rozumnyi, S. R. Bulò, L. Porzi, and P. Kotschieder. A Recipe for Generating 3D Worlds from a Single Image. *2025 IEEE/CVF International Conference on Computer Vision (ICCV)*, 2025. 18, 20, 26
- [130] Max Schwarz and Sven Behnke. Stilleben: Realistic Scene Synthesis for Deep Learning in Robotics. *2020 IEEE International Conference on Robotics and Automation (ICRA)*, 2020. 9
- [131] A. Sengupta, Ignas Budvytis, and R. Cipolla. Synthetic Training for Accurate 3D Human Pose and Shape Estimation in the Wild. *British Machine Vision Conference*, 2020. 26
- [132] Yihang She, Andrew Blake, David A. Coomes, and Srinivasan Keshav. Scaling Up Forest Vision with Synthetic Data. *International Journal of Computer Vision*, 2025. 7, 10, 11
- [133] Bokui Shen, Xinchun Yan, C. Qi, Mahyar Najibi, Boyang Deng, L. Guibas, Yin Zhou, and Drago Anguelov. GINA-3D: Learning to Generate Implicit Neural Assets in the Wild. *2023 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023. 7, 10, 12, 14, 15
- [134] Qijia Shen and Guangrun Wang. Geometry aware 3D generation from in-the-wild images in ImageNet. *arXiv preprint arXiv:2402.00225*, 2024. <https://arxiv.org/abs/2402.00225>. 12, 26
- [135] Tianchang Shen, Sherwin Bahmani, Kai He, Sangeetha Grama Srinivasan, Tianshi Cao, Jiawei Ren, Ruilong Li, Zian Wang, Nicholas Sharp, Zan Gojcic, Sanja Fidler, Jiahui Huang, Huan Ling, Jun Gao, and Xuanchi Ren. Lyra 2.0: Explorable Generative 3D Worlds. *arXiv preprint arXiv:2604.13036*, 2026. <https://arxiv.org/abs/2604.13036>. 18, 20, 30
- [136] Yu Shi, Hannah Tang, M. Baine, M. A. Hollingsworth, Huijing Du, Dandan Zheng, Chi Zhang, and Hongfeng Yu. 3DGAUnet: 3D Generative Adversarial Networks with a 3D U-Net Based Generator to Achieve the Accurate and Effective Synthesis of Clinical Tumor Image Data for Pancreatic Cancer. *Cancers*, 2023. 11, 26

- [137] Zifan Shi, Sida Peng, Yinghao Xu, Andreas Geiger, Yiyi Liao, and Yujun Shen. Deep Generative Models on 3D Representations: A Survey. *arXiv preprint arXiv:2210.15663*, 2022. <https://arxiv.org/abs/2210.15663>. 1
- [138] Jungwook Shin, Jaeill Kim, Kyungeun Lee, Hyunghun Cho, and Wonjong Rhee. DR.CPO: Diversified and Realistic 3D Augmentation via Iterative Construction, Random Placement, and HPR Occlusion. *AAAI Conference on Artificial Intelligence*, 2023. 7, 10, 28
- [139] Miguel V. Silva and A. Carvalho. AI-powered Contextual 3D Environment Generation: A Systematic Review. *arXiv preprint arXiv:2506.05449*, 2025. <https://arxiv.org/abs/2506.05449>. 1
- [140] Chahat Deep Singh, R. Kumari, Cornelia Fermuller, N. Sanket, and Y. Aloimonos. WorldGen: A Large Scale Generative Simulator. *2023 IEEE International Conference on Robotics and Automation (ICRA)*, 2022. 8, 11
- [141] I. Sipiran, Gustavo Santelices, Lucas Oyarzún, Andrea Ranieri, C. Romanengo, Silvia Biasotti, and Bianca Falcidieno. Symmetria: A Synthetic Dataset for Learning in Point Clouds. *International Journal of Computer Vision*, 2025. 4, 6
- [142] Ido Sobol, Kihyuk Sohn, Yoav Blum, Egor Zakharov, Max Bluvstein, Andrea Vedaldi, and O. Litany. Realiz3D: 3D Generation Made Photorealistic via Domain-Aware Learning. *arXiv preprint arXiv:2605.13852*, 2026. <https://arxiv.org/abs/2605.13852>. 26
- [143] Luca Stornaiuolo, Nima Dehmamy, A. Barab'asi, and Mauro Martino. 3D Topology Transformation with Generative Adversarial Networks. *ICCC*, 2020. 13, 26
- [144] Chunyi Sun, Junlin Han, Weijian Deng, Xinlong Wang, Z. Qin, and Stephen Gould. 3D-GPT: Procedural 3D Modeling with Large Language Models. *2025 International Conference on 3D Vision (3DV)*, 2023. 22, 23, 24
- [145] Fan-Yun Sun, Shengguang Wu, Christian Jacobsen, Thomas Yim, Haoming Zou, Alex Zook, Shangru Li, Yu-Hsin Chou, Ethem Can, Xunlei Wu, Clemens Eppner, Valts Blukis, Jonathan Tremblay, Jiajun Wu, Stanley T. Birchfield, and Nick Haber. 3D-Generalist: Self-Improving Vision-Language-Action Models for Crafting 3D Worlds. *arXiv preprint arXiv:2507.06484*, 2025. <https://arxiv.org/abs/2507.06484>. 23
- [146] Jianhua Sun, Yuxuan Li, Jiude Wei, Longfei Xu, Nange Wang, Yining Zhang, and Cewu Lu. Arti-PG: A Toolbox for Procedurally Synthesizing Large-Scale and Diverse Articulated Objects with Rich Annotations. *2025 IEEE/CVF International Conference on Computer Vision (ICCV)*, 2024. 6, 7
- [147] Jiankai Sun, Jianing Qiu, Chuanyang Zheng, Johnathan Tucker, Javier Yu, and Mac Schwager. Aria-NeRF: Multimodal Egocentric View Synthesis. *arXiv preprint arXiv:2311.06455*, 2023. <https://arxiv.org/abs/2311.06455>. 4
- [148] Mingchao Sun, Luyang Tang, Yu Liu, Xu Yan, Zhanqi Li, Yunwei Zhang, Fei Yu, Zengye Ge, Yumin Liu, Jiachen Zhang, Yongchang Zhang, Jiawei Zhang, Zhicheng Liu, Zhong-Hao Sun, Tianjian Ouyang, Wenzheng Chen, Shixing Yang, Nianfei Fan, Guo Sun, Huan Li, Zhengfu Zhou, Yongze Li, Yingliang Peng, Mengmeng Du, Yuan Liu, Haozhe Shi, Chunnuo Gong, Chen Yu, Chun Jia, Yang Liu, S. Zeng, Junnan Lai, Hang Zhang, Ning Guo, Baoquan Chen, Mu Xu, and Hongyu Pan. ABot-3DWorld 0: A Universal World Model to Explore Any 3D Space. *arXiv preprint arXiv:2607.11673*, 2026. <https://arxiv.org/abs/2607.11673>. 20, 29
- [149] Qinghong Sun, Yangguang Li, Zexia Liu, Xiaoshui Huang, Fenggang Liu, Xihui Liu, Wanli Ouyang, and Jing Shao. UniG3D: A Unified 3D Object Generation Dataset. *arXiv preprint arXiv:2306.10730*, 2023. <https://arxiv.org/abs/2306.10730>. 4
- [150] Saksham Suri, Fanyi Xiao, Animesh Sinha, S. Culatana, Raghuraman Krishnamoorthi, Chenchen Zhu, and Abhinav Shrivastava. Gen2Det: Generate to Detect. *arXiv preprint arXiv:2312.04566*, 2023. <https://arxiv.org/abs/2312.04566>. 18, 21, 28
- [151] C. Symeonidis, P. Nousi, P. Tosidis, K. Tsampazis, N. Passalis, A. Tefas, and N. Nikolaidis. Proceedings of the 22nd Engineering Applications of Neural Networks Conference - EANN 2021, Crete, Greece, 25-27 June, 2021. *International Conference on Engineering Applications of Neural Networks*, 2021. 19
- [152] Ayca Takmaz, Jonas Schult, Irem Kaftan, Mertcan Akccay, R. Sumner, Bastian Leibe, Francis Engelmann, and Siyu Tang. 3D Segmentation of Humans in Point Clouds with Synthetic Data. *2023 IEEE/CVF International Conference on Computer Vision (ICCV)*, 2022. 27
- [153] Xiang Tang, Ruotong Li, and Xiaopeng Fan. Recent Advances in 3D Object and Scene Generation: A Survey. *arXiv preprint arXiv:2504.11734*, 2025. <https://arxiv.org/abs/2504.11734>. 1
- [154] S. Tata and Subhankar Mishra. 3D GANs and Latent Space: A comprehensive survey. *arXiv preprint arXiv:2304.03932*, 2023. <https://arxiv.org/abs/2304.03932>. 1
- [155] S. Team, Xingyu Chen, Fu-Jen Chu, Pierre Gleize, Kevin J. Liang, Alexander Sax, Hao Tang, Weiyao Wang, Michelle Guo, Thibaut Hardin, Xiang Li, A. Lin, Jiawei Liu, Ziqi Ma, Anushka Sagar, Bowen Song, Xiaodong Wang, Jianing Yang, Bowen Zhang, Piotr Doll'ar, Georgia Gkioxari, Matt Feiszli, and

- Jitendra Malik. SAM 3D: 3Dfy Anything in Images. *arXiv preprint arXiv:2511.16624*, 2025. <https://arxiv.org/abs/2511.16624>. 19, 21, 31
- [156] Kohei Torimi, Ryosuke Yamada, Daichi Otsuka, Kensho Hara, Yuki M. Asano, Hirokatsu Kataoka, and Yoshimitsu Aoki. Text-guided Synthetic Geometric Augmentation for Zero-shot 3D Understanding. *arXiv preprint arXiv:2501.09278*, 2025. <https://arxiv.org/abs/2501.09278>. 27, 28
- [157] Thejas Venkatesh and S. Velury. BlendFusion - Scalable Synthetic Data Generation for Diffusion Model Training. *arXiv preprint arXiv:2604.09022*, 2026. <https://arxiv.org/abs/2604.09022>. 9
- [158] Clément Vignac and P. Frossard. Top-N: Equivariant set and graph generation without exchangeability. *International Conference on Learning Representations*, 2021. 2, 3, 5
- [159] Nicolas von Lützwow, Barbara Rössle, Katharina Schmid, and Matthias Niekner. GaussianGPT: Towards Autoregressive 3D Gaussian Scene Generation. *arXiv preprint arXiv:2603.26661*, 2026. <https://arxiv.org/abs/2603.26661>. 14
- [160] Cheng Wan, Yongsan Mao, Wenzheng Wu, Yuxuan Xie, C. Xiang, Runze Wang, Xiang Zhang, Zhongyuan Liu, Rushi Dai, and Yuan Liu. NaLA: A 3D Native LLM Layout Agent for High-quality 3D Scene Generation. *arXiv preprint arXiv:2606.29395*, 2026. <https://arxiv.org/abs/2606.29395>. 21, 22
- [161] Austin T. Wang, Dongchen Yang, and Angel X. Chang. Scaling Diverse Language Generation for 3D Visual Grounding. *arXiv preprint arXiv:2606.20946*, 2026. <https://arxiv.org/abs/2606.20946>. 11, 25, 27
- [162] Dexin Wang, F. Chang, and Chunsheng Liu. ForeRobo: Unlocking Infinite Simulation Data for 3D Goal-driven Robotic Manipulation. *arXiv preprint arXiv:2511.04381*, 2025. <https://arxiv.org/abs/2511.04381>. 23, 30
- [163] Ruihe Wang, Yukang Cao, Kai Han, and K. K. Wong. A Survey on 3D Human Avatar Modeling - From Reconstruction to Generation. *arXiv preprint arXiv:2406.04253*, 2024. <https://arxiv.org/abs/2406.04253>. 3, 5
- [164] Xinjie Wang, Liu Liu, Taojun Ding, Andrew Choi, Chaodong Huang, Mengao Zhao, Ziang Li, Jackson Jiang, Chunlei Yu, Shengxiang Liu, Wei Xu, and Zhizhong Su. EmbodiedGen V2: An Agentic, Simulation-Ready 3D World Engine for Embodied AI. *arXiv preprint arXiv:2607.07459*, 2026. <https://arxiv.org/abs/2607.07459>. 22, 23, 28, 29
- [165] Yunsong Wang, Na Zhao, and Gim Hee Lee. Enhancing Generalizability of Representation Learning for Data-Efficient 3D Scene Understanding. *2024 International Conference on 3D Vision (3DV)*, 2024. 8, 11, 25, 32
- [166] Dongxu Wei, Qi Xu, Zhiqi Li, Hangning Zhou, Congpei Qiu, Hailong Qin, Mu Yang, Zhaopeng Cui, and Peidong Liu. Any 3D Scene is Worth 1K Tokens: 3D-Grounded Representation for Scene Generation at Scale. *arXiv preprint arXiv:2604.11331*, 2026. <https://arxiv.org/abs/2604.11331>. 16, 19
- [167] Si-Tong Wei, Rui-Huan Wang, Chuanpei Zhou, Baoquan Chen, and Peng-Shuai Wang. OctGPT: Octree-based Multiscale Autoregressive Models for 3D Shape Generation. *Proceedings of the Special Interest Group on Computer Graphics and Interactive Techniques Conference Conference Papers*, 2025. 14
- [168] Beichen Wen, Haozhe Xie, Zhaoxi Chen, Fangzhou Hong, and Ziwei Liu. 3D Scene Generation: A Survey. *arXiv preprint arXiv:2505.05474*, 2025. <https://arxiv.org/abs/2505.05474>. 1
- [169] Rundui Wu and Changxi Zheng. Learning to Generate 3D Shapes from a Single Example. *ACM Transactions on Graphics (TOG)*, 2022. 14, 15
- [170] Shuang Wu, Youtian Lin, Feihu Zhang, Yifei Zeng, Yikang Yang, Yajie Bao, Jiachen Qian, Siyu Zhu, Xun Cao, Philip H. S. Torr, and Yao Yao. Direct3D-S2: Gigascale 3D Generation Made Easy with Spatial Sparse Attention. *Neural Information Processing Systems*, 2025. 12, 14, 27
- [171] Xian Wu, Dingkan Liang, Tianrui Feng, Kui Xia, Yumeng Zhang, Xiaofan Li, X. Tan, and Xiangyu Bai. Generation Models Know Space: Unleashing Implicit 3D Priors for Scene Understanding. *arXiv preprint arXiv:2603.19235*, 2026. <https://arxiv.org/abs/2603.19235>. 29
- [172] Zhennan Wu, Yang Li, Han Yan, Taizhang Shang, Weixuan Sun, Senbo Wang, Ruikai Cui, Weizhe Liu, Hiroyuki Sato, Hongdong Li, and Pan Ji. BlockFusion: Expandable 3D Scene Generation using Latent Tri-plane Extrapolation. *ACM Transactions on Graphics (TOG)*, 2024. 13, 27
- [173] Xin Xia, Weiguang Zhao, Yuyao Yan, Guanyu Yang, Rui Zhang, Kaizhu Huang, and Xi Yang. Towards Training-Free Open-World Classification with 3D Generative Models. *Proceedings of the 33rd ACM International Conference on Multimedia*, 2025. 29
- [174] Desai Xie, Sai Bi, Zhixin Shu, Kai Zhang, Zexiang Xu, Yi Zhou, Soren Pirk, Arie E. Kaufman, Xin Sun, and Hao Tan. LRM-Zero: Training Large Reconstruction Models with Synthesized Data. *Neural Information Processing Systems*, 2024. 7, 11, 27
- [175] Zhifeng Xie, Keyi Zhang, Yiye Yan, Yuling Guo, Fan Yang, Jiting Zhou, and Mengtian Li. FilmSceneDesigner: Chaining Set Design for Procedural

- Film Scene Generation. *AAAI Conference on Artificial Intelligence*, 2025. 7, 22, 23
- [176] Zhangyang Xiong, Di Kang, Derong Jin, Weikai Chen, Linchao Bao, and Xiaoguang Han. Get3DHuman: Lifting StyleGAN-Human into a 3D Generative Model using Pixel-aligned Reconstruction Priors. *2023 IEEE/CVF International Conference on Computer Vision (ICCV)*, 2023. 27
- [177] Dejie Xu, Ye Yuan, Morteza Mardani, Sifei Liu, Jiaming Song, Zhangyang Wang, and Arash Vahdat. AGG: Amortized Generative 3D Gaussians for Single Image to 3D. *Trans. Mach. Learn. Res.*, 2024. 16, 19
- [178] Hongbin Xu, Weitao Chen, Zhipeng Zhou, Feng Xiao, Baigui Sun, M. Shou, and Wenxiong Kang. ControlLRM: Fast and Controllable 3D Generation via Large Reconstruction Model. *arXiv preprint arXiv:2410.09592*, 2024. <https://arxiv.org/abs/2410.09592>. 19
- [179] Yongzhi Xu, Yonhon Ng, Yifu Wang, Inkyu Sa, Yunfei Duan, Yang Li, Pan Ji, and Hongdong Li. Sketch2Scene: Automatic Generation of Interactive 3D Game Scenes from User's Casual Sketches. *arXiv preprint arXiv:2408.04567*, 2024. <https://arxiv.org/abs/2408.04567>. 19
- [180] Zhiyuan Xu, Jiuming Liu, Yuxin Chen, Masayoshi Tomizuka, Chenfeng Xu, and Chensheng Peng. Rethinking Image-to-3D Generation with Sparse Queries: Efficiency, Capacity, and Input-View Bias. *arXiv preprint arXiv:2604.13905*, 2026. <https://arxiv.org/abs/2604.13905>. 19, 27
- [181] Le Xue, Ning Yu, Shu Zhang, Junnan Li, Roberto Mart'in-Mart'in, Jiajun Wu, Caiming Xiong, Ran Xu, Juan Carlos Niebles, and S. Savarese. ULIP-2: Towards Scalable Multimodal Pre-Training for 3D Understanding. *2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023. 5
- [182] Yuxuan Xue, Xianghui Xie, R. Marin, and Gerard Pons-Moll. Gen-3Diffusion: Realistic Image-to-3D Generation via 2D & 3D Diffusion Synergy. *IEEE transactions on pattern analysis and machine intelligence*, 2024. 18
- [183] Yuxuan Xue, Xianghui Xie, Margaret Kostyrko, and Gerard Pons-Moll. InfiniHuman: Realistic 3D Human Creation with Precise Control. *Proceedings of the SIGGRAPH Asia 2025 Conference Papers*, 2025. 21, 27
- [184] Han Yan, Yang Li, Zhennan Wu, Shenzhou Chen, Weixuan Sun, Taizhang Shang, Weizhe Liu, Tian Chen, Xiqiang Dai, Chao Ma, Hongdong Li, and Pan Ji. Frankenstein: Generating Semantic-Compositional 3D Scenes in One Tri-Plane. *SIGGRAPH Asia 2024 Conference Papers*, 2024. 12, 13, 15
- [185] Jiahao Yang, Wufei Ma, Angtian Wang, Xiaoding Yuan, A. Yuille, and Adam Kortylewski. Robust Category-Level 3D Pose Estimation from Synthetic Data. *arXiv preprint arXiv:2305.16124*, 2023. <https://arxiv.org/abs/2305.16124>. 9, 11, 26
- [186] Jianing Yang, Xuwei Chen, Nikhil Madaan, Madhavan Iyengar, Shengyi Qian, David F. Fouhey, and Joyce Chai. 3D-GRAND: A Million-Scale Dataset for 3D-LLMs with Better Grounding and Less Hallucination. *2025 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. 4, 22
- [187] Shuai Yang, Jing Tan, Mengchen Zhang, Tong Wu, Gordon Wetzstein, Ziwei Liu, and Dahua Lin. LayerPano3D: Layered 3D Panorama for Hyper-Immersive Scene Generation. *Proceedings of the Special Interest Group on Computer Graphics and Interactive Techniques Conference Conference Papers*, 2024. 18, 19
- [188] Yu Yang, Alan Liang, Jianbiao Mei, Yukai Ma, Yong Liu, and Gim Hee Lee. X-Scene: Large-Scale Driving Scene Generation with High Fidelity and Flexible Controllability. *Neural Information Processing Systems*, 2025. 31
- [189] Yue Yang, Fan-Yun Sun, Luca Weihs, Eli Vanderbilt, Alvaro Herrasti, Winson Han, Jiajun Wu, Nick Haber, Ranjay Krishna, Lingjie Liu, Chris Callison-Burch, Mark Yatskar, Aniruddha Kembhavi, and Christopher Clark. Holodeck: Language Guided Generation of 3D Embodied AI Environments. *2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023. 21, 22, 24
- [190] Yuhang Yang, Fengqi Liu, Yixing Lu, Qin Zhao, Ping Wu, Wei Zhai, Ran Yi, Yang Cao, Lizhuang Ma, Zhengjun Zha, and Juntong Dong. SIGMAN: Scaling 3D Human Gaussian Generation with Millions of Assets. *2025 IEEE/CVF International Conference on Computer Vision (ICCV)*, 2025. 12, 13
- [191] Chongjie Ye, Chen Cao, Chuanyu Pan, Yiming Hao, Yihao Zhi, Yuanming Hu, and Xiaoguang Han. Omni123: Exploring 3D Native Foundation Models with Limited 3D Data by Unifying Text to 2D and 3D Generation. *arXiv preprint arXiv:2604.02289*, 2026. <https://arxiv.org/abs/2604.02289>. 14, 29
- [192] Tianwei Ye, Yifan Mao, Minwen Liao, Jian Liu, Chunchao Guo, Dazhao Du, Quanxin Shou, Fangqi Zhu, and Song Guo. 3D Generation for Embodied AI and Robotic Simulation: A Survey. *arXiv preprint arXiv:2604.26509*, 2026. <https://arxiv.org/abs/2604.26509>. 5
- [193] Zipeng Ye, Mengfei Xia, Yanan Sun, Ran Yi, Mingjing Yu, Juyong Zhang, Yu-Kun Lai, and Yong-Jin Liu. 3D-CariGAN: An End-to-End Solution to 3D Caricature Generation from Normal Face Photos.

- IEEE Transactions on Visualization and Computer Graphics*, 2020. 12
- [194] Taoran Yi, Jiemin Fang, Junjie Wang, Guanjuan Wu, Lingxi Xie, Xiaopeng Zhang, Wenyu Liu, Qi Tian, and Xinggang Wang. GaussianDreamer: Fast Generation from Text to 3D Gaussians by Bridging 2D and 3D Diffusion Models. *2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023. 17, 18
 - [195] Arafa Yoncalik, W. Jansen, Nico Huebel, Mohammad Rahmani, and Jan Steckel. LLM-Driven 3D Scene Generation of Agricultural Simulation Environments. *2026 IEEE Conference on Artificial Intelligence (CAI)*, 2026. 22, 23, 25
 - [196] Huangyue Yu, Baoxiong Jia, Yixin Chen, Yandan Yang, Puhao Li, Rongpeng Su, Jiabin Li, Qing Li, Wei Liang, Song-Chun Zhu, Tengyu Liu, and Siyuan Huang. MetaScenes: Towards Automated Replica Creation for Real-world 3D Scans. *2025 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025. 22, 24, 25
 - [197] Wangbo Yu, Xuelin Qian, Jingyang Huo, Tiejun Huang, Bo Zhao, and Yanwei Fu. Pushing the Limits of 3D Shape Generation at Scale. *arXiv preprint arXiv:2306.11510*, 2023. <https://arxiv.org/abs/2306.11510>. 14, 27
 - [198] Binggui Zang, Wenting Zheng, Xiaoliang Luo, Zhiyuan Fang, Shi Li, Lvchun Wang, Wei Yu, Yi Zhao, Tian Xie, Yuchi Huo, and Rengan Xie. HIVE-3D: Hierarchical Voxel Enhancement for High-Quality 3D Scene Generation. *arXiv preprint arXiv:2607.13468*, 2026. <https://arxiv.org/abs/2607.13468>. 20
 - [199] F. G. Zanjani, Davide Abati, A. Wiggers, Dimitris Kalatzis, Jens Petersen, Hong Cai, and A. Habibi. Gaussian Splatting is an Effective Data Generator for 3D Object Detection. *arXiv preprint arXiv:2504.16740*, 2025. <https://arxiv.org/abs/2504.16740>. 10
 - [200] Tianle Zeng, Gerardo Loza Galindo, Junlei Hu, Pietro Valdastri, and Dominic Jones. Realistic Surgical Image Dataset Generation Based On 3D Gaussian Splatting. *International Conference on Medical Image Computing and Computer-Assisted Intervention*, 2024. 7, 9, 16
 - [201] Bowen Zhang, Tianyu Yang, Yu Li, Lei Zhang, and Xi Zhao. Compress3D: a Compressed Latent Space for 3D Generation from a Single Image. *European Conference on Computer Vision*, 2024. 13, 18, 20
 - [202] Bowen Zhang, Chunchao Guo, Dong Guo, Haolin Liu, Hongyu Yan, Huiwen Shi, Jiaao Yu, Jiachen Xu, Jingwei Huang, Kunhong Li, Lifu Wang, Linus, Penghao Wang, Qin Lin, Ruining Tang, Xianguhui Yang, Yang Li, Yirui Guan, Yunfei Zhao, Yunhan Yang, Zeqiang Lai, Zhihao Liang, and Zibo Zhao. HY3D-Bench: Generation of 3D Assets. *arXiv preprint arXiv:2602.03907*, 2026. <https://arxiv.org/abs/2602.03907>. 26
 - [203] Dongsu Zhang, Changwoon Choi, Jeonghwan Kim, and Y. Kim. Learning to Generate 3D Shapes with Generative Cellular Automata. *International Conference on Learning Representations*, 2021. 6
 - [204] Jichao Zhang, E. Sangineto, Hao Tang, Aliaksandr Siarohin, Zhun Zhong, N. Sebe, and Wei Wang. 3D-Aware Semantic-Guided Generative Model for Human Synthesis. *European Conference on Computer Vision*, 2021. 16
 - [205] Jinzhi Zhang, Feng Xiong, and Mu Xu. G3PT: Unleash the power of Autoregressive Modeling in 3D Generation via Cross-scale Querying Transformer. *International Joint Conference on Artificial Intelligence*, 2024. 12, 14
 - [206] Jinzhi Zhang, Feng Xiong, and Mu Xu. 3D representation in 512-Byte: Variational tokenizer is the key for autoregressive 3D generation. *arXiv preprint arXiv:2412.02202*, 2024. <https://arxiv.org/abs/2412.02202>. 14, 16
 - [207] Shougao Zhang, Mengqi Zhou, Yuxi Wang, Chuanchen Luo, Rong Wang, Yiwei Li, Xucheng Yin, Zhaoxiang Zhang, and Junran Peng. CityX: Controllable Procedural Content Generation for Unbounded 3D Cities. *arXiv preprint arXiv:2407.17572*, 2024. <https://arxiv.org/abs/2407.17572>. 23
 - [208] Shujie Zhang, Jingkun Yi, Weipeng Zhong, Zirui Zhou, Yangkun Zhu, Hanqing Wang, Xudong Xu, Weinan Zhang, and Chunhua Shen. RoboSnap: One-Shot Real-to-Sim Scene Generation for Generalizable Robot Learning and Evaluation. *arXiv preprint arXiv:2607.06699*, 2026. <https://arxiv.org/abs/2607.06699>. 20, 28, 30
 - [209] Xinyi Zhang, Daoyi Gao, Naiqi Li, and Angela Dai. ProcGen3D: Learning Neural Procedural Graph Representations for Image-to-3D Reconstruction. *arXiv preprint arXiv:2511.07142*, 2025. <https://arxiv.org/abs/2511.07142>. 14
 - [210] Haoyu Zhao, Cheng Zeng, Linghao Zhuang, Yaxi Zhao, Shengke Xue, Hao Wang, Xingyue Zhao, Zhongyu Li, Kehan Li, Siteng Huang, Mingxiu Chen, Xin Li, Deli Zhao, and Hua Zou. High-Fidelity Simulated Data Generation for Real-World Zero-Shot Robotic Manipulation Learning With Gaussian Splatting. *IEEE Robotics and Automation Letters*, 2025. 22, 24, 28, 30
 - [211] Ke Zhao and A. Larsen. Challenges and Opportunities in 3D Content Generation. *arXiv preprint arXiv:2405.15335*, 2024. <https://arxiv.org/abs/2405.15335>. 17, 26, 27
 - [212] Rui Zhao, Wei Li, Zhipeng Hu, Lincheng Li,

- Zhengxia Zou, Z. Shi, and Changjie Fan. Zero-Shot Text-to-Parameter Translation for Game Character Auto-Creation. *2023 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023. 22, 24
- [213] Weiguang Zhao, Haoran Xu, Xingyu Miao, Qin Zhao, Rui Zhang, Kai Huang, Ning Gao, Peizhou Cao, Mingze Sun, Mulin Yu, Tao Lu, Linning Xu, Junting Dong, and Jiangmiao Pang. SynthVerse: A Large-Scale Diverse Synthetic Dataset for Point Tracking. *Proceedings of the Special Interest Group on Computer Graphics and Interactive Techniques Conference Conference Papers*, 2026. 9
- [214] Kaizhi Zheng, Ruijian Zha, Zishuo Xu, Jing Gu, Jie Yang, and Xin Eric Wang. Constructing a 3D Scene from a Single Image. *arXiv preprint arXiv:2505.15765*, 2025. <https://arxiv.org/abs/2505.15765>. 20
- [215] Wei zhi Nie, Ruidong Chen, Weijie Wang, B. Lepri, and N. Sebe. T2TD: Text-3D Generation Model Based on Prior Knowledge Guidance. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2023. 24
- [216] Jiahao Zhou, Chen Long, Yue Xie, Jialiang Wang, Conglang Zhang, Boheng Li, Haiping Wang, Zhe Chen, and Zhen Dong. WHU-Synthetic: A Synthetic Perception Dataset for 3-D Multitask Model Research. *arXiv preprint arXiv:2402.19059*, 2024. <https://arxiv.org/abs/2402.19059>. 4, 10
- [217] Mengqi Zhou, Yuxi Wang, Jun Hou, Shougao Zhang, Yiwei Li, Chuanchen Luo, Junran Peng, and Zhaoxiang Zhang. SceneX: Procedural Controllable Large-scale Scene Generation. *arXiv preprint arXiv:2403.15698*, 2024. <https://arxiv.org/abs/2403.15698>. 22, 23, 24, 27, 31
- [218] Shengchao Zhou, Jiehong Lin, Jiahui Liu, Shizhen Zhao, Chirui Chang, and Xiaojuan Qi. ASSIST-3D: Adapted Scene Synthesis for Class-Agnostic 3D Instance Segmentation. *AAAI Conference on Artificial Intelligence*, 2025. 21, 22
- [219] Wenxu Zhou, Kaixuan Nie, Hang Du, Dong Yin, Wei Huang, Siqi Guo, Xiaobo Zhang, and Peng Hu. IL3D: A Large-Scale Indoor Layout Dataset for LLM-Driven 3D Scene Generation. *arXiv preprint arXiv:2510.12095*, 2025. <https://arxiv.org/abs/2510.12095>. 4, 21
- [220] Hao Zhu, Miaomiao Zhao, Di yue Liu, Runze Tian, Yan Zhang, Jie Guo, and Fenggen Yu. GeoDiff3D: Self-Supervised 3D Scene Generation with Geometry-Constrained 2D Diffusion Guidance. *arXiv preprint arXiv:2601.19785*, 2026. <https://arxiv.org/abs/2601.19785>. 17, 26
- [221] Lingting Zhu, Shengju Qian, Haidi Fan, Jiayu Dong, Zhenchao Jin, Siwei Zhou, Gen Dong, Xin Wang, and Lequan Yu. AssetFormer: Modular 3D Assets Generation with Autoregressive Transformer. *arXiv preprint arXiv:2602.12100*, 2026. <https://arxiv.org/abs/2602.12100>. 14
- [222] Yue Zhu and David Picard. Decanus to Legatus: Synthetic training for 2D-3D human pose lifting. *Asian Conference on Computer Vision*, 2022. 10
- [223] Yi Zhuang, Yuxiao He, Jiawei Zhang, Yanwen Wang, Jiahe Zhu, Yao Yao, Siyu Zhu, Xun Cao, and Hao Zhu. Towards Native Generative Model for 3D Head Avatar. *arXiv preprint arXiv:2410.01226*, 2024. <https://arxiv.org/abs/2410.01226>. 15, 16, 27
- [224] Yi Zhuang, Jiaxi Lv, Hao Wen, Qing Shuai, Ailing Zeng, Hao Zhu, Shifeng Chen, Yujiu Yang, Xun Cao, and Wei Liu. IDOL: Instant Photorealistic 3D Human Creation from a Single Image. *2025 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. 18, 19
- [225] R. Ziatdinov and Rifkat I. Nabiyev. Generating Digital Models Using Text-to-3D and Image-to-3D Prompts: Critical Case Study. *2025 International Russian Smart Industry Conference (SmartIndustryCon)*, 2025. 26, 28
- [226] Basak Melis Öcal, Maxim Tatarchenko, Sezer Karaoglu, and Theo Gevers. SceneTeller: Language-to-3D Scene Generation. *European Conference on Computer Vision*, 2024. 21, 22